

Introduction to the LLVM Compiler Framework

SS 2012

Jun.-Prof. Dr. Christian Plessl

Custom Computing
University of Paderborn

- brief overview of a state of the art compiler framework
 - we are using LLVM in our research
 - we will use it as an example in the lecture and in the exercises
- outline
 - overview of the LLVM compiler framework
 - compilation tool flows
 - LLVM intermediate representations
 - optimizations
 - code generation

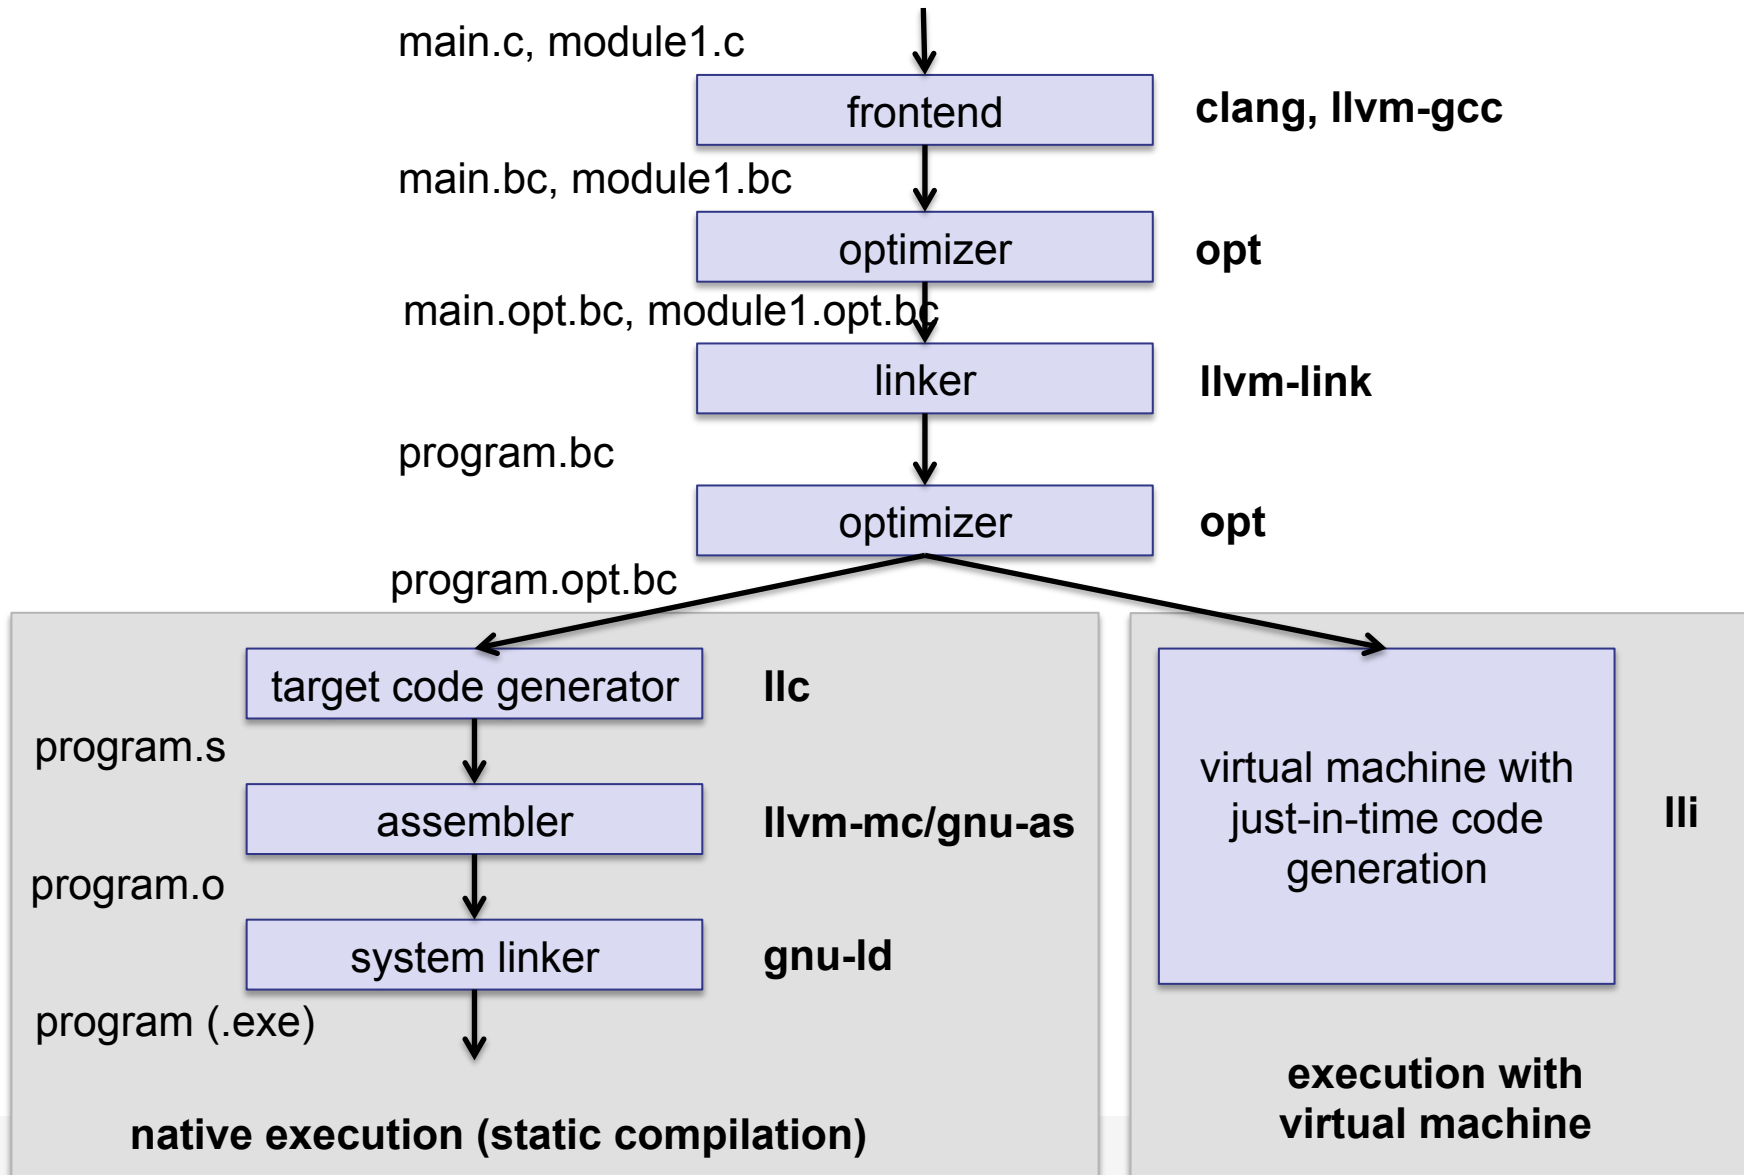
What is the LLVM Compiler Framework

- modern open-source compiler infrastructure
 - implemented in C++
 - modular and extensible design
 - combines a static compilation tool flow with a virtual machine
- many supported front-ends/languages
 - C, C++, Objective-C (Clang, GCC/dragonegg)
 - Ruby (Rubinius, MacRuby)
 - Python (unladen swallow)
 - and many more
- many supported CPU architectures in backend
 - ARM, Alpha, Intel x86, Microblaze, MIPS, PowerPC, SPARC, ...
- very popular and widely used
 - Apple, AMD, NVidia, Cray, Google, ...

LLVM Design Principle

- separation of the compilation process in frontend / analysis and transformation / backend
- LLVM intermediate representation (LLVM IR) plays a central role in this process
 - all code optimizations are implemented as “LLVM IR to LLVM IR transformation passes”
 - code analysis is also implemented as pass, generated results can be shared between passes
- all target processor-specific optimizations are handled in the backend

Static LLVM Compilation Toolflow

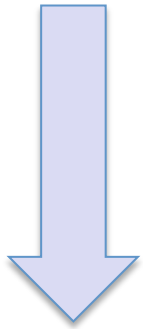


LLVM Intermediate Representation

- basis for all LLVM optimization passes
- low-level assembly language for a “virtual” processor
 - load/store architecture
 - infinite amount of named registers
 - each register is assigned exactly once (static single assignment, SSA)
- exists in three equivalent representations
 - in-memory C++ data structures
 - binary files (LLVM bitcode) (file extension: .bc)
 - human readable assembly notation (file extension: .ll)
- reference
 - <http://llvm.org/docs/LangRef.html>

```
double circle_area(double radius) {  
    return radius * radius * 3.14159265;  
}
```

C code



compiling with Clang frontend and LLVM

```
clang -S -emit-llvm circle_area.c -o circle_area.ll  
opt -mem2reg -instsimplify -S circle_area.ll -o circle_area.opt.ll
```

```
define double @circle_area(double %radius)  
nounwind ssp {  
    %1 = fmul double %radius, %radius  
    %2 = fmul double %1, 0x400921FB53C8D4F1  
    ret double %2  
}
```

LLVM IR code

note:

- strongly typed data types
- each variable only assigned once (static single assignment)

Running optimizations passes on LLVM IR

C code

```
int test1(int x, int y, int z) {  
    return (x & z) ^ (y & z);  
}
```

example

- compile this code to LLVM IR
- run series of optimization passes to iteratively improve code

Running optimizations passes on LLVM IR 2

Compiled with Clang (no optimizations)

clang -S -emit-llvm boolean_optimization.c -o boolean_optimization.ll

```
define i32 @test1(i32 %x, i32 %y, i32 %z)
nounwind ssp {
    %1 = alloca i32, align 4
    %2 = alloca i32, align 4
    %3 = alloca i32, align 4
    store i32 %x, i32* %1, align 4
    store i32 %y, i32* %2, align 4
    store i32 %z, i32* %3, align 4
    %4 = load i32* %1, align 4
    %5 = load i32* %3, align 4
    %6 = and i32 %4, %5
    %7 = load i32* %2, align 4
    %8 = load i32* %3, align 4
    %9 = and i32 %7, %8
    %10 = xor i32 %6, %9
    ret i32 %10
}
```

properties of the
generated code

- correct, but obviously very inefficient
- each function argument is written to the stack first ...
- ... and read from the stack later on

Running optimizations passes on LLVM IR 3

First optimization (mem2reg, instsimplify)

opt -mem2reg -instsimplify -S boolean_optimization.ll -o boolean_optimization.opt.ll

```
define i32 @test1(i32 %x, i32 %y, i32 %z)
nounwind ssp {
    %1 = and i32 %x, %z
    %2 = and i32 %y, %z
    %3 = xor i32 %1, %2
    ret i32 %3
}
```

properties of the generated code:

- removed redundant instructions
- used registers instead of stack memory
- instructions of the actual computation remain unchanged

can this code be simplified any further?

Running optimizations passes on LLVM IR 4

Second optimization: instcombine

opt -mem2reg -instcombine -S boolean_optimization.ll -o boolean_optimization.opt.ll

```
define i32 @test1(i32 %x, i32 %y, i32 %z)
nounwind ssp {
    %1 = xor i32 %x, %y
    %2 = and i32 %1, %z
    ret i32 %2
}
```

properties of the generated code:

- further simplification of the code
- instcombine not only removes redundant instructions but changes instructions
- optimization pass did understand the semantics of the boolean operations and figured out that $(x \text{ and } z) \text{ xor } (y \text{ and } z) == z \text{ and } (x \text{ xor } y)$

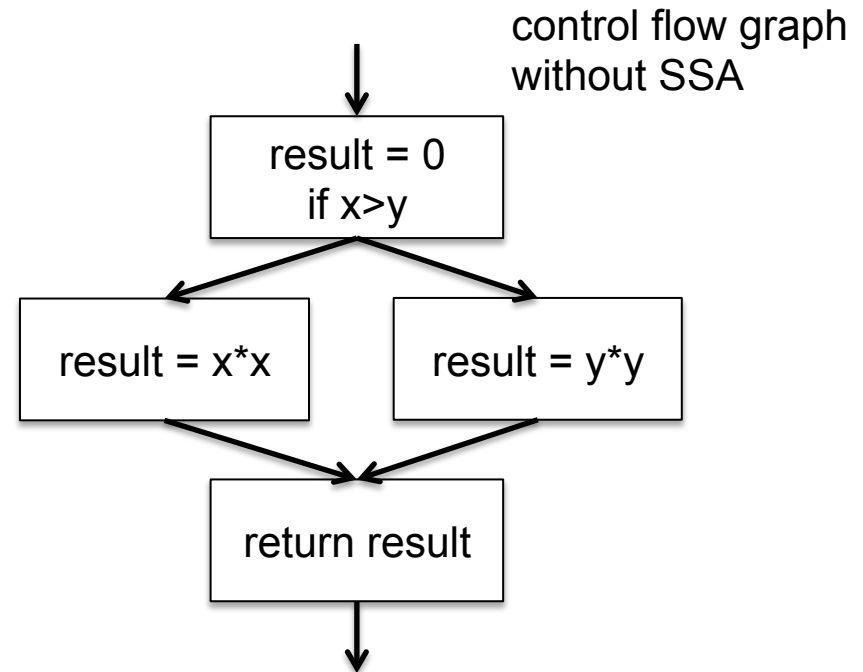
numerous additional optimizations available, consult opt manual page for details

Static Single Assignment 1

- LLVM IR uses static single assignment (SSA) form
 - each virtual register is assigned only once
 - allows to easily track define-use chains, i.e. what values are used by which instructions (useful e.g. for dead code elimination)
- what happens if we need to assign a register several times, e.g. in a loop or in branches?
- example

```
int max_square(int x, int y)
{
    int result = 0;

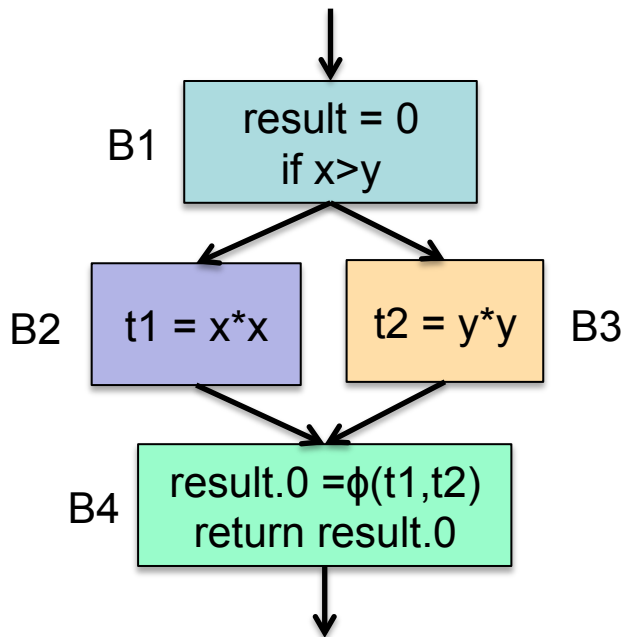
    if (x>y){
        result = x*x;
    } else {
        result = y*y;
    }
    return result;
}
```



Static Single Assignment 2

- use of phi-nodes/instructions (ϕ)
 - phi nodes keep track which control-flow path was taken and use the corresponding value (like a multiplexer)
 - not actually implemented, compiler just makes sure that the virtual registers are mapped to the same physical register

control flow graph with SSA



```
define i32 @max_square(i32 %x, i32 %y) {  
    %1 = icmp sgt i32 %x, %y  
    br i1 %1, label %2, label %4
```

```
; <label>:2 ← beginning and name of basic block  
    %3 = mul nsw i32 %x, %x  
    br label %6
```

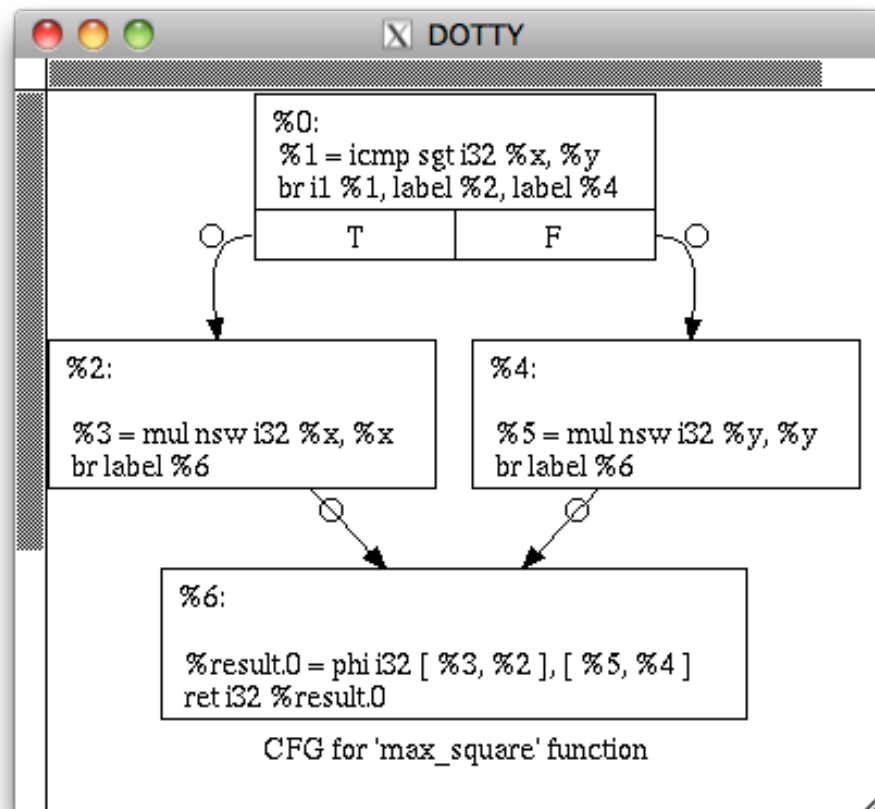
```
; <label>:4  
    %5 = mul nsw i32 %y, %y  
    br label %6
```

```
; <label>:6  
    %result.0 = phi i32 [ %3, %2 ], [ %5, %4 ]  
    ret i32 %result.0  
}
```

choose reg %3 if control flow enters from BB %2,
choose reg %5 if control flow enters from BB %4

Visualizing Control Flow Graphs

- LLVM has built-in support for visualizing various steps in the compilation process



`opt -view-cfg -S phi.opt.ll`

- example: using different backends, compilation for MIPS and ARM instruction set

```
max_square:
# BB#0:
    addiu    $sp, $sp, -16
    slt      $2, $5, $4
    beq      $2, $zero, $BB0_2
    nop
# BB#1:
    mult     $4, $4
    j        $BB0_3
    nop
$BB0_2:
    mult     $5, $5
$BB0_3:
    mflo     $2
    addiu    $sp, $sp, 16
    jr       $ra
    nop
```

MIPS assembler code

```
_max_square:    @ @max_square
@ BB#0:
    cmp      r0, r1
    mulle    r2, r1, r1
    mulgt    r2, r0, r0
    mov      r0, r2
    bx       lr
```

ARM assembler code

```
llc --march=mips phi.ll -o phi.mips.s
llc --march=arm phi.ll -o phi.arm.s
```

- LLVM is a modern open source compiler framework
 - very powerful and easy to use
 - human readable IR allows for following optimization steps
 - modular design allows adding own functionality
- LLVM may also be of practical use for you
 - as a replacement for GCC
 - for generating code for embedded processors
 - for learning about compilers and optimizations
 - building your own programming language (frontend) that uses LLVM as a backend (search the web for inspiration)
- acknowledgement
 - this presentation is based partly on materials that have been kindly provided by Tobias Grosser (<http://grosser.es/>), visit his website for more information on LLVM

- 2012-04-24 (v.1.1.0)
 - updated for SS2012
- 2011-05-05 (v1.0.1)
 - fix a couple of minor typos