4 Myths about in-memory databases busted

Yiftach Shoolman – Co-Founder & CTO @ Redis Labs

@yiftachsh, @redislabsinc

redislabs

Background - Redis

Created by Salvatore Sanfilippo (@antirez)

OSS, in-memory NoSQL k/v database/data-structure engine

1

Fastest growing DB
2013-01 – to date:
DB-Engines

2

Most popular database on containers:
@DevOps.com
& ClusterHQ

3

Top NoSQL databases:
DB-Engines

12

Top tools developers love:
StackShare

Backgroud – Redis Labs

Founded in 2011. HQ in Mountain View CA, R&D in Tel-Aviv IL

The largest commercial company behind OSS Redis

- **5000+ paying customers**
- **30,000+ free users**
- **100,000+ databases under management**
- **±200 new databases/day**

Provide enterpsie class Redis deployement

- **As a service – Redis Cloud**
- **On-premises – Redis Labs Enterprise Cluster (RLEC)**

$28MM VC funding

Why an in-memory operational DBMS ?

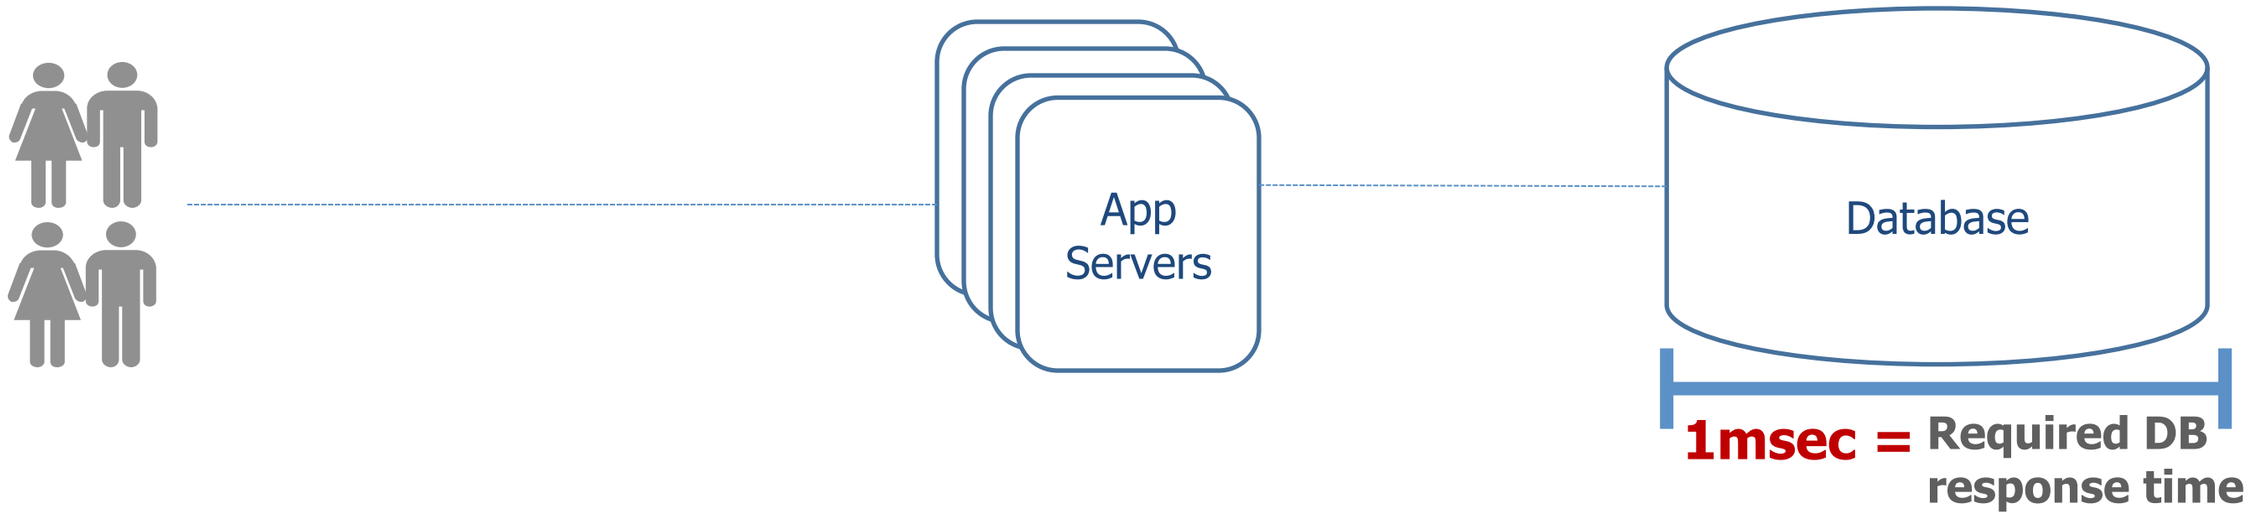

redislabs

Why an in-memory analytics DBMS ?

Query time:

Days/Hours

Disk-based

Query time:

Mins/Secs

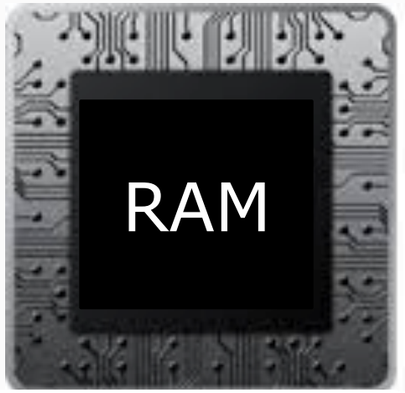

RAM-based

Myth #1

all in-memory databases are equally fast

How many ops/sec can a single EC2 instance do?

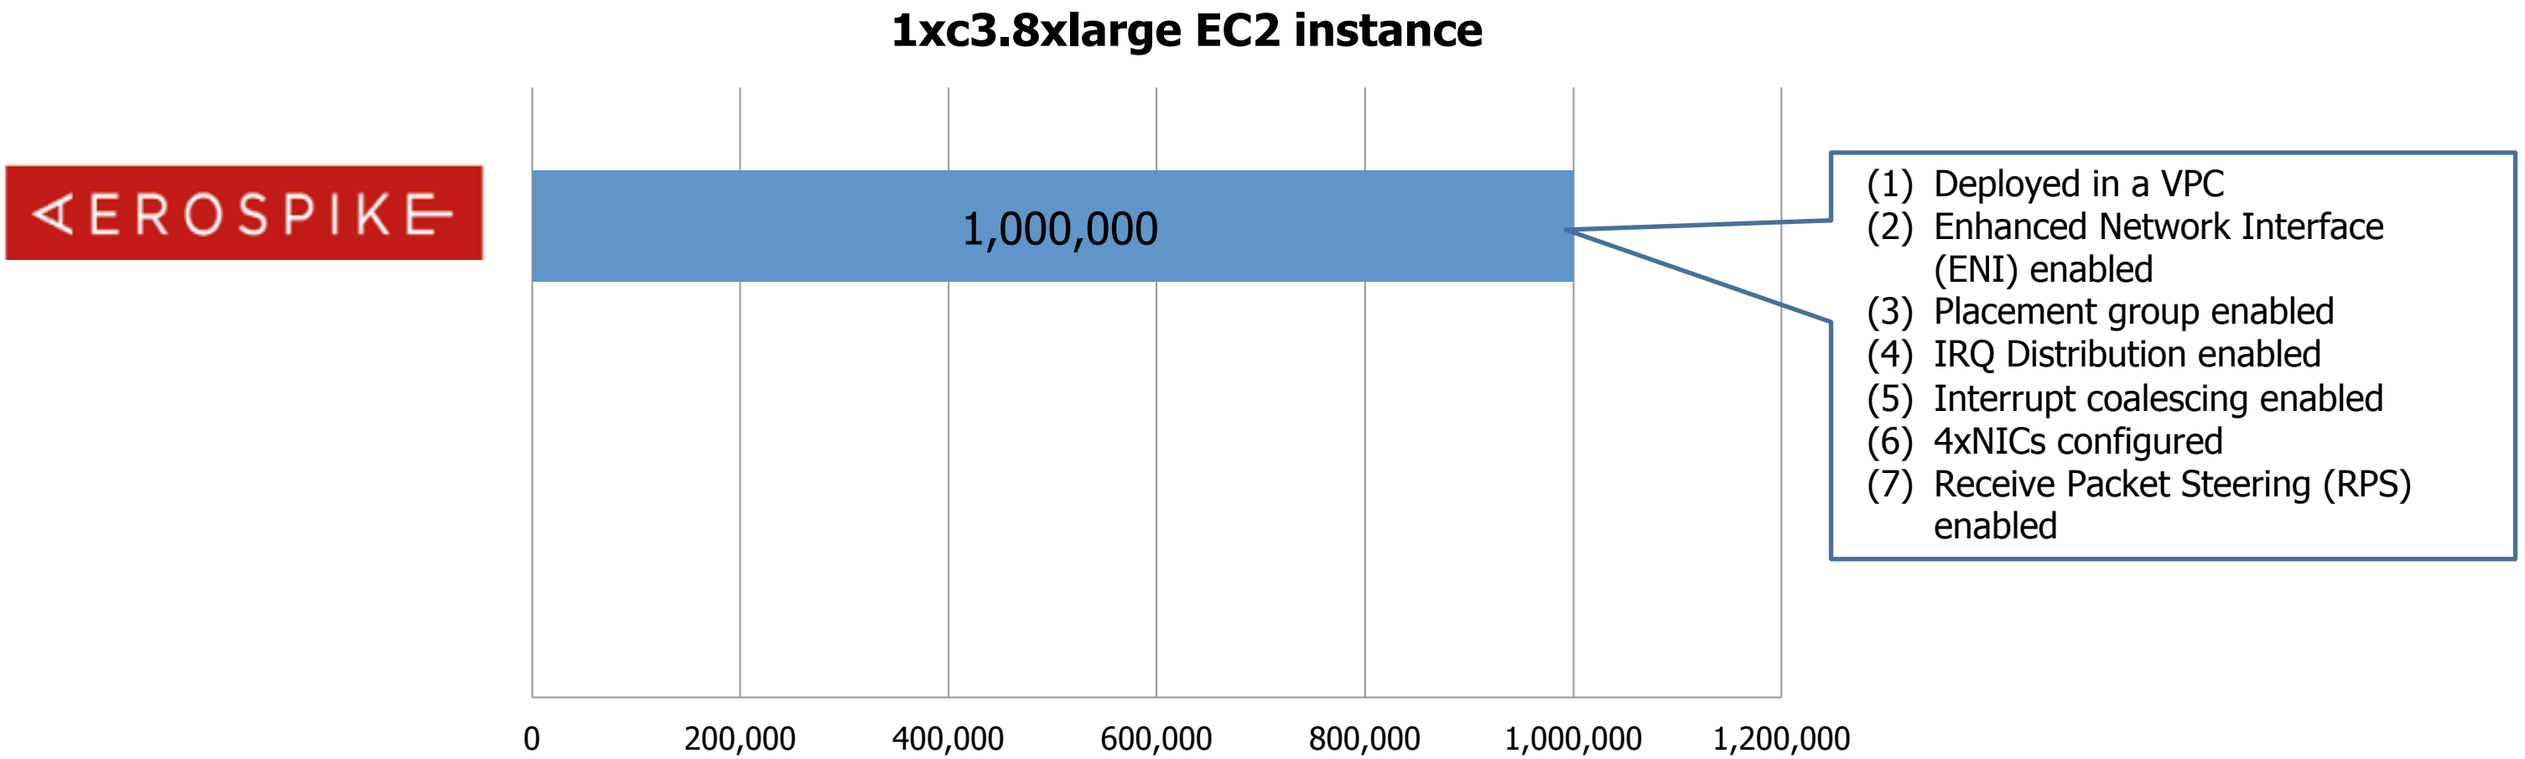

redislabs

How many ops/sec can a single EC2 instance do?

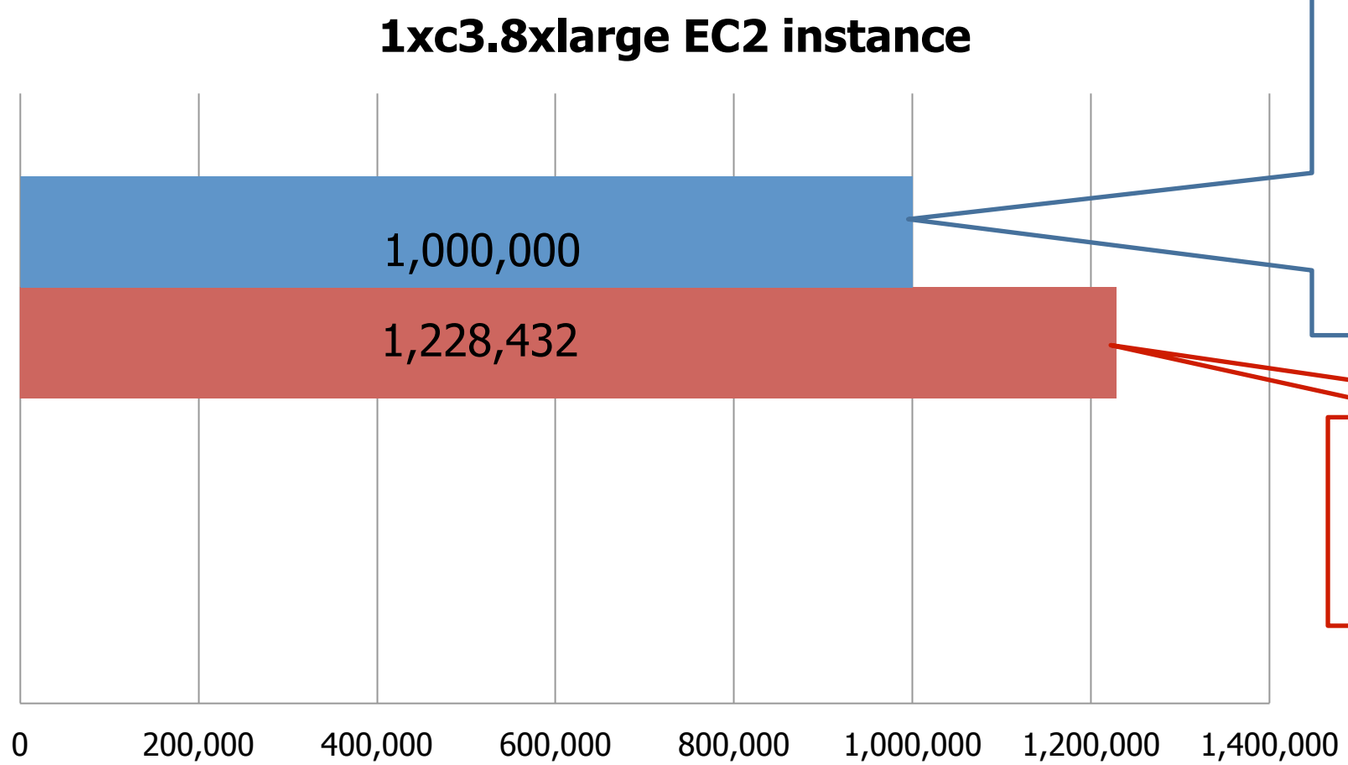

Fully optimized
(1) Deployed in a VPC
(2) ENI enabled
(3) Placement group enabled
(4) IRQ Distribution enabled
(5) Interrupt coalescing enabled
(6) 4xNICs configured
(7) RPS enabled

Not optimized
(1) Deployed on Classic EC2
(2) No tuning
(3) Reached packets/sec limit

How many ops/sec can a single EC2 instance do?

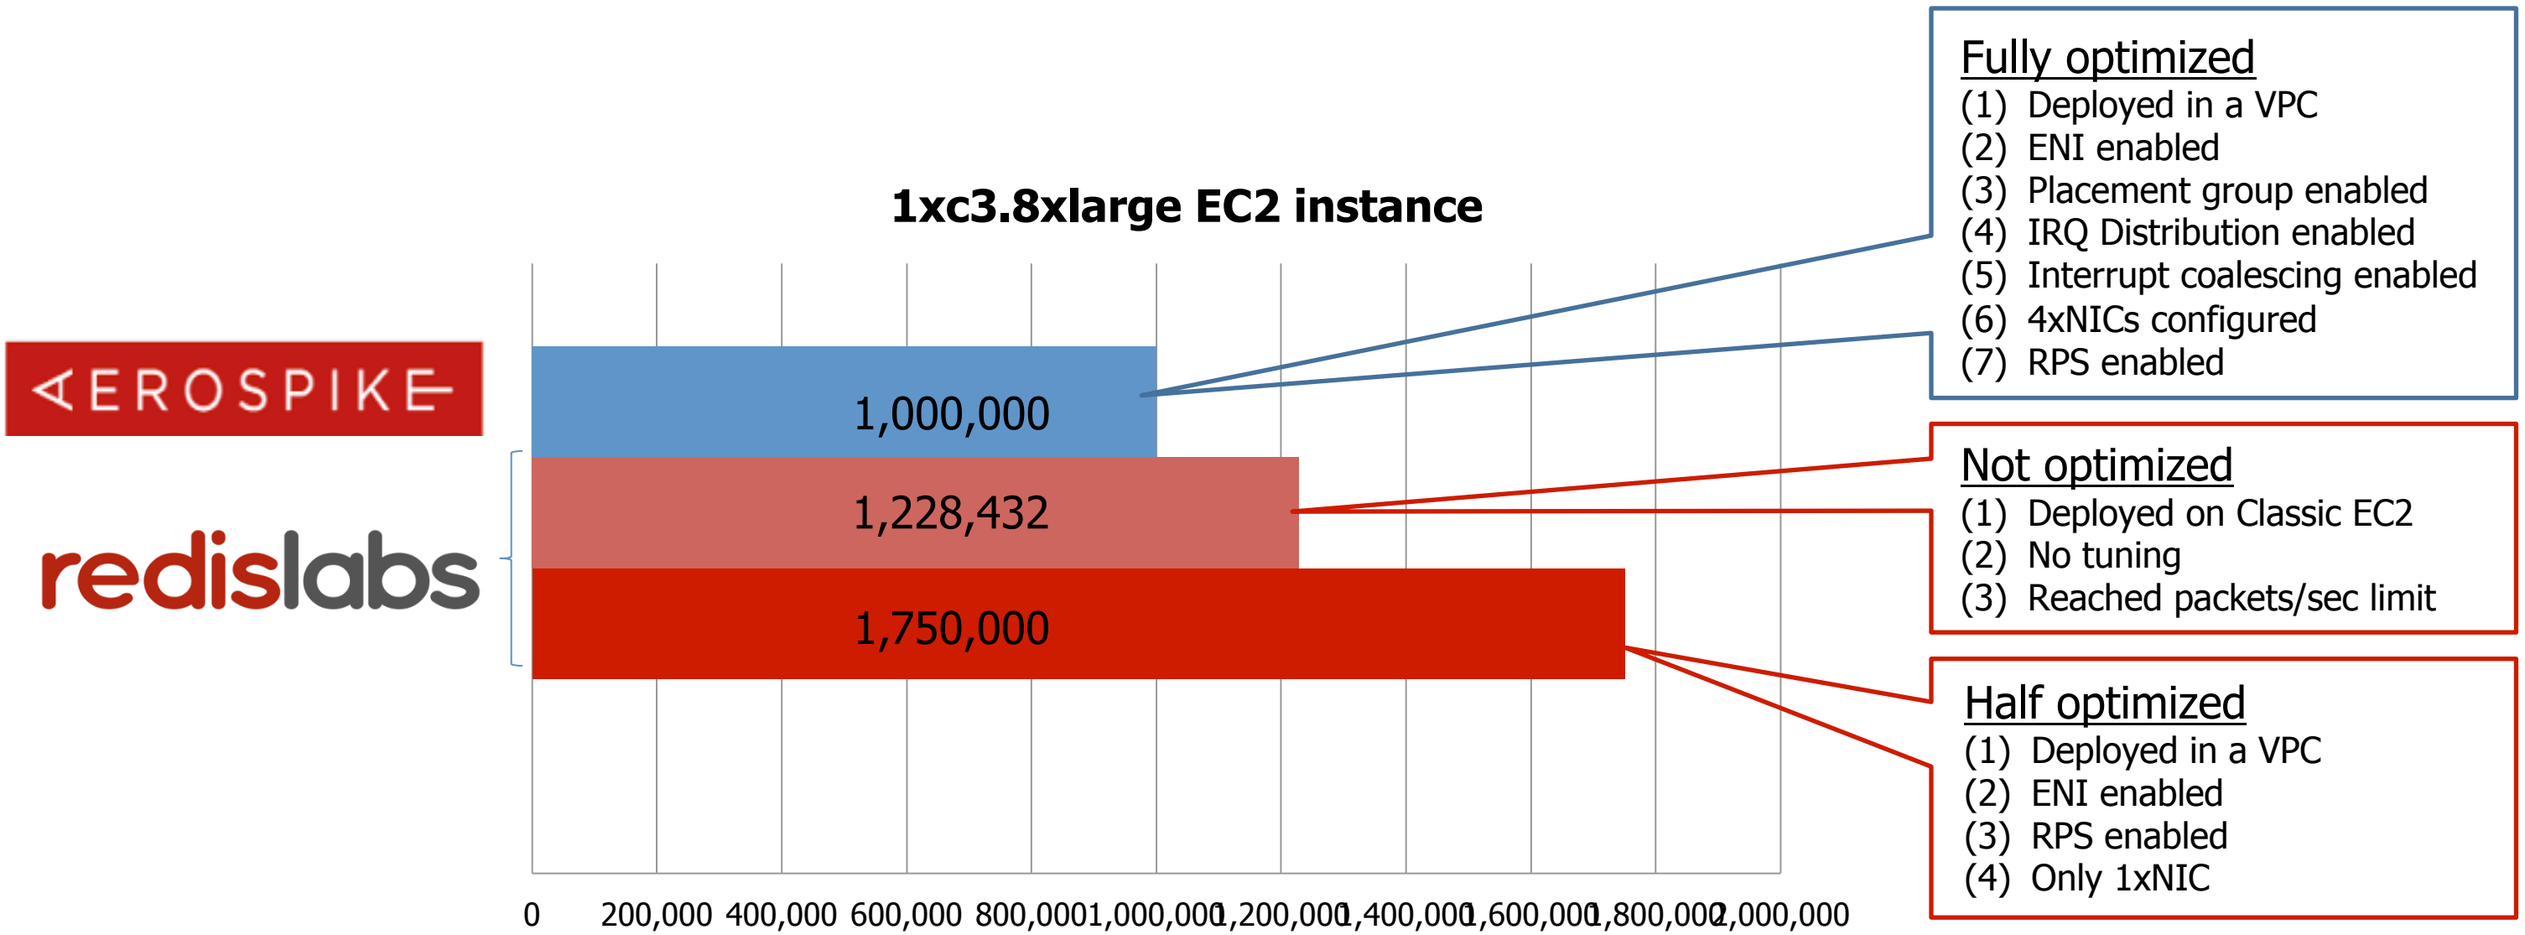

redislabs

How many ops/sec can a single EC2 instance do?

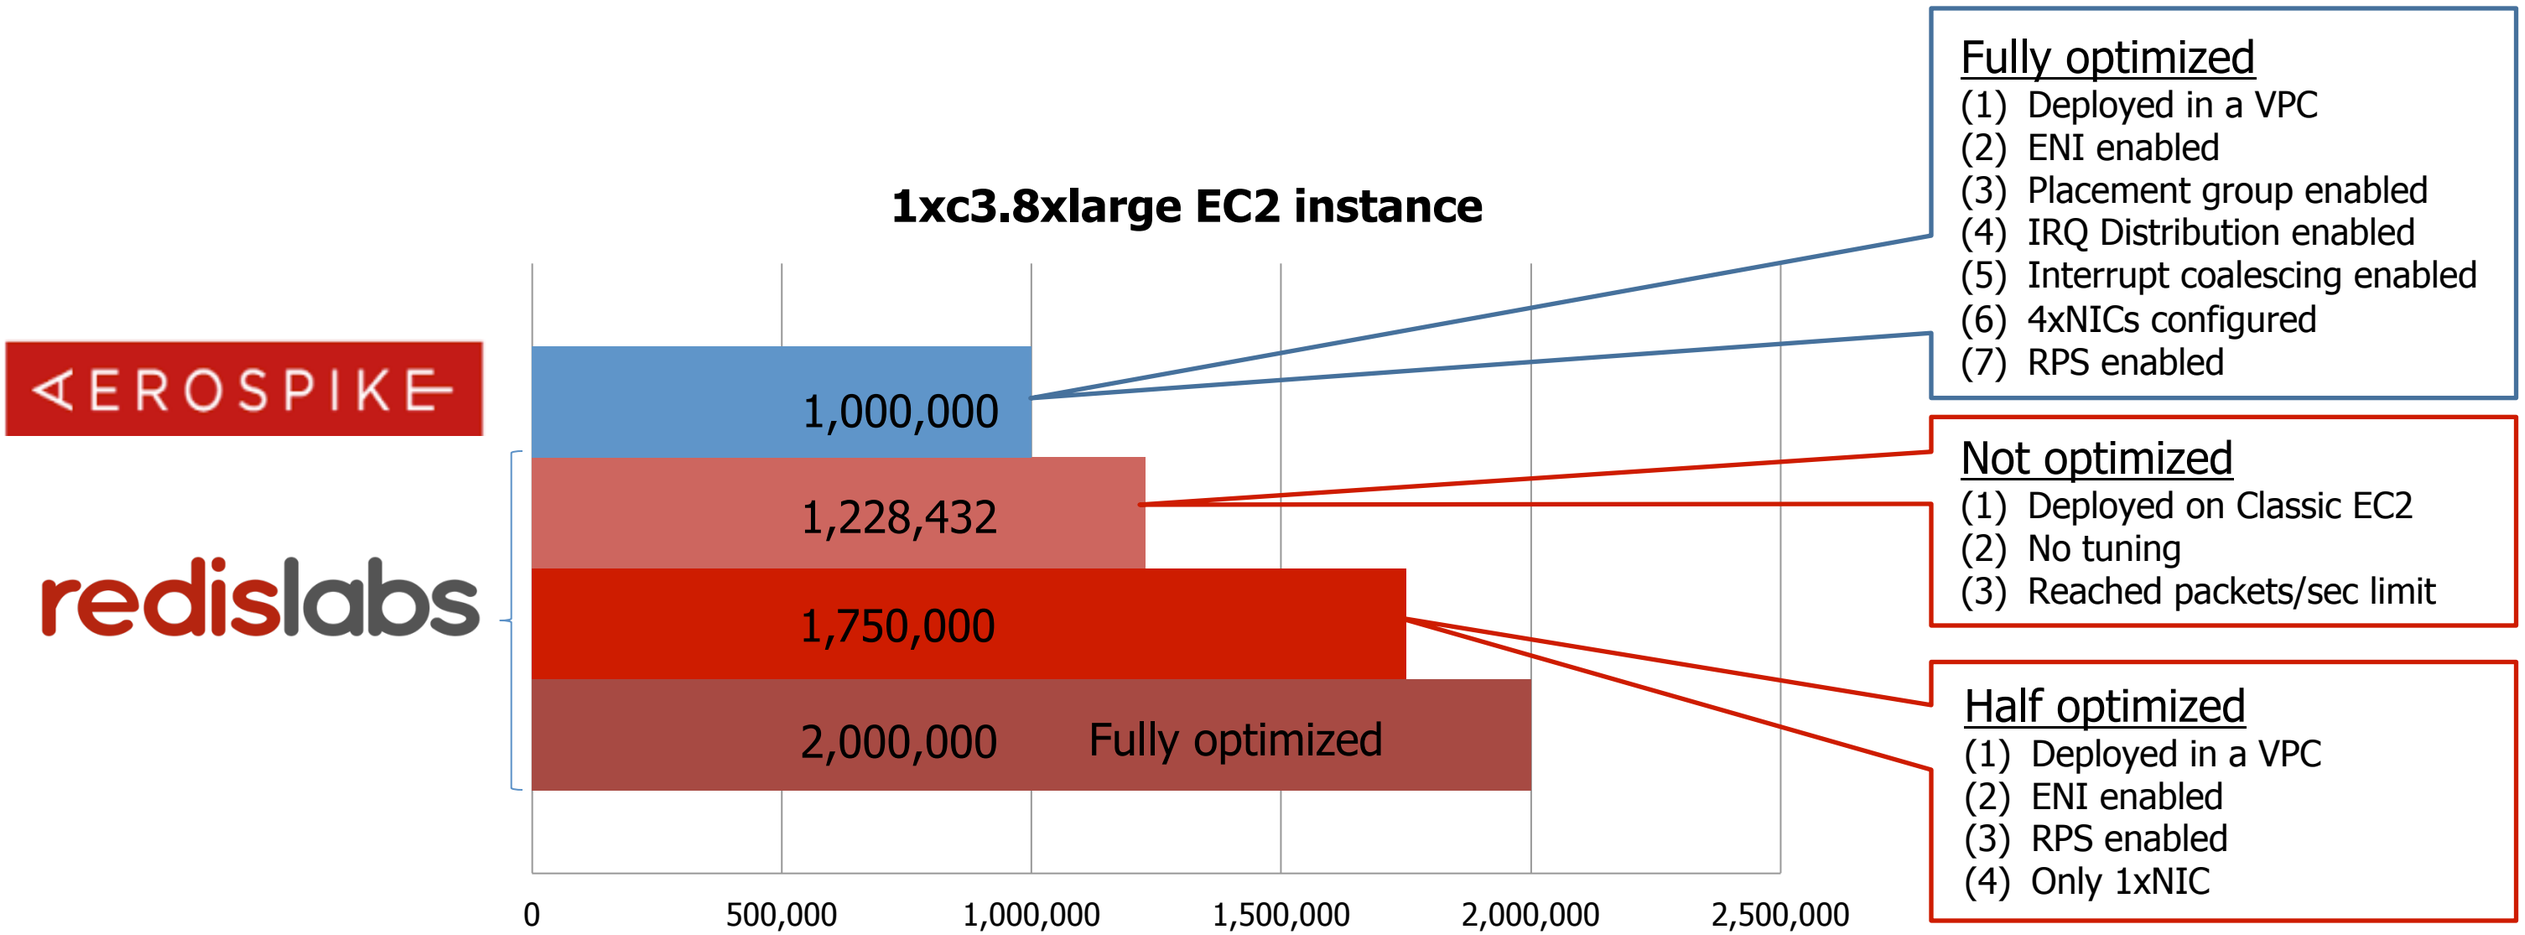

How many servers to get 1M writes/sec on GCE?

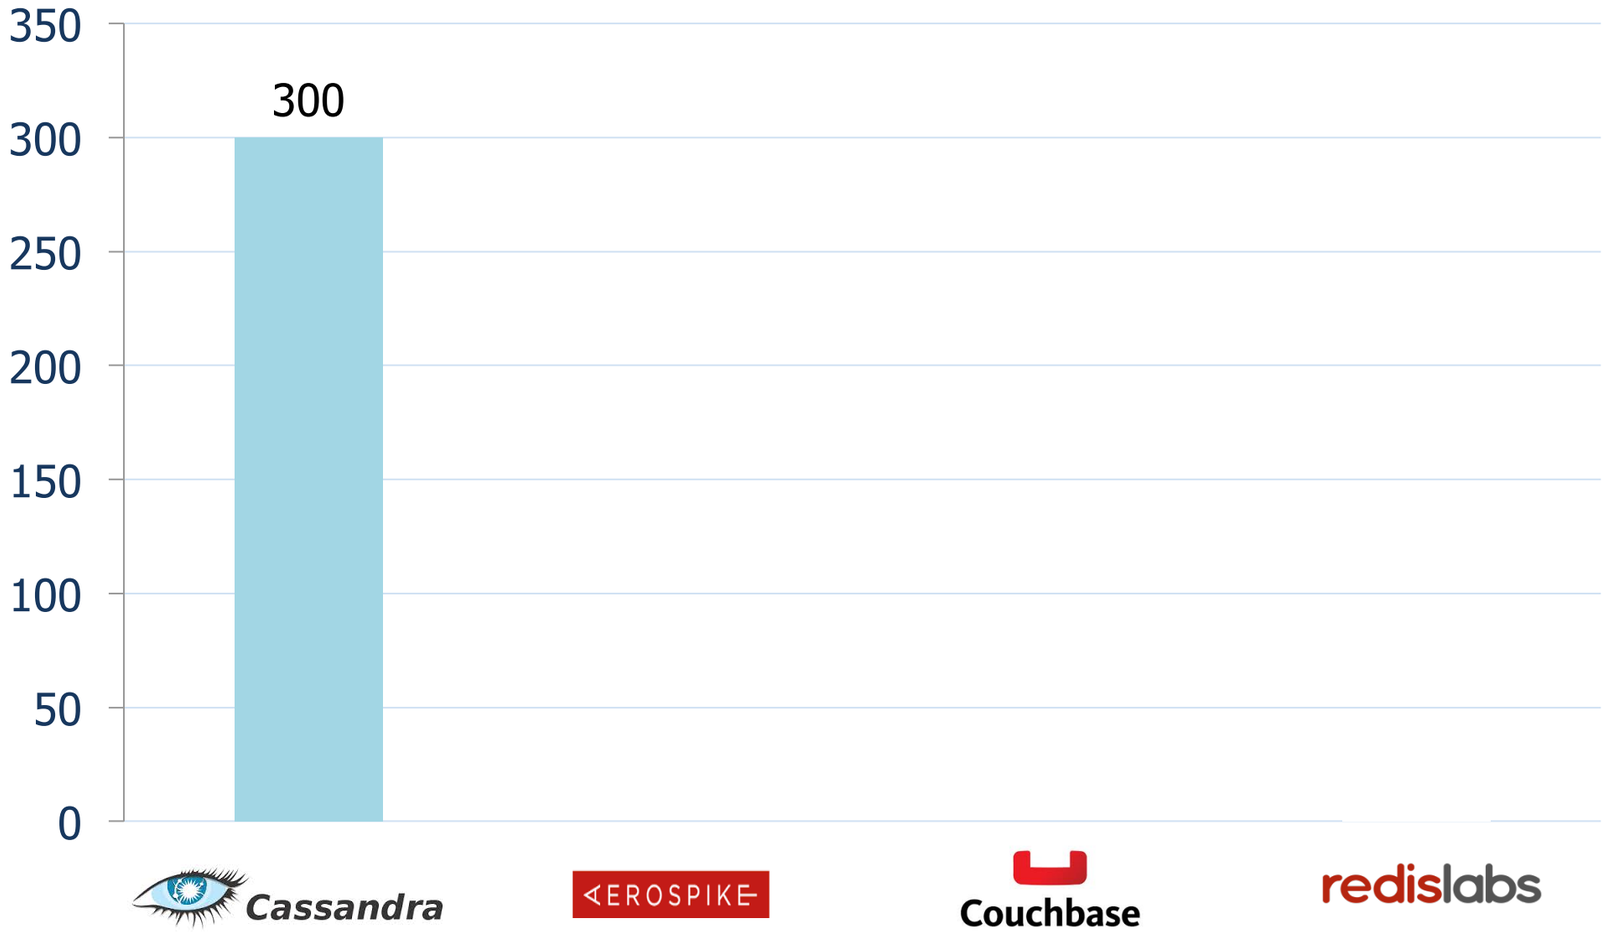

How many servers to get 1M writes/sec on GCE?

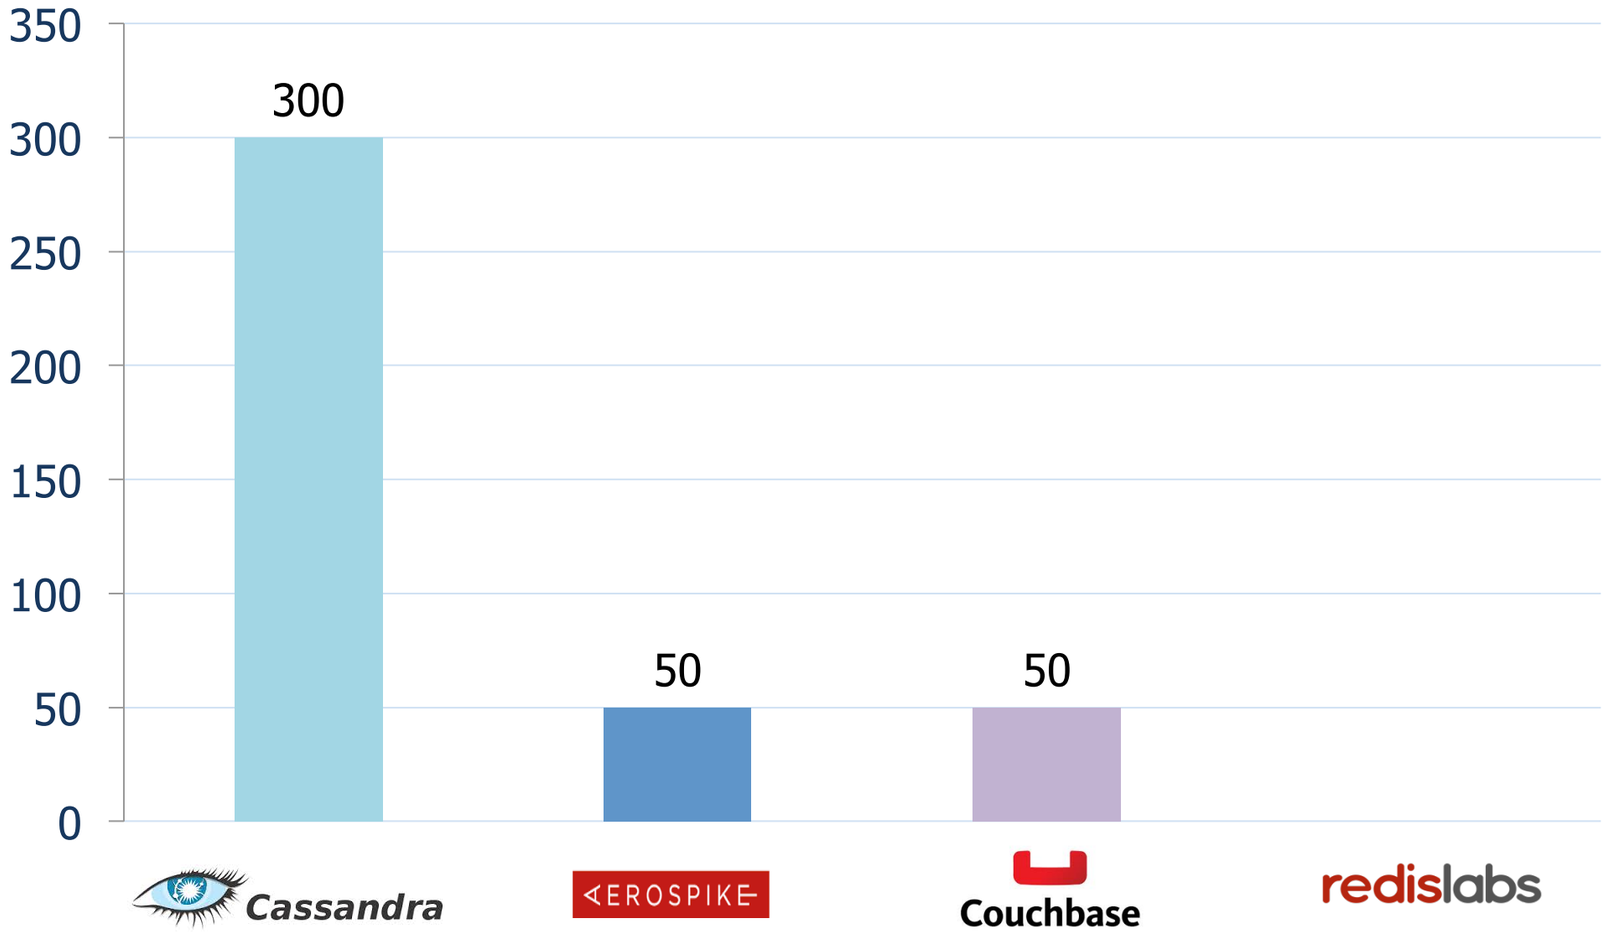

How many servers to get 1M writes/sec on GCE?

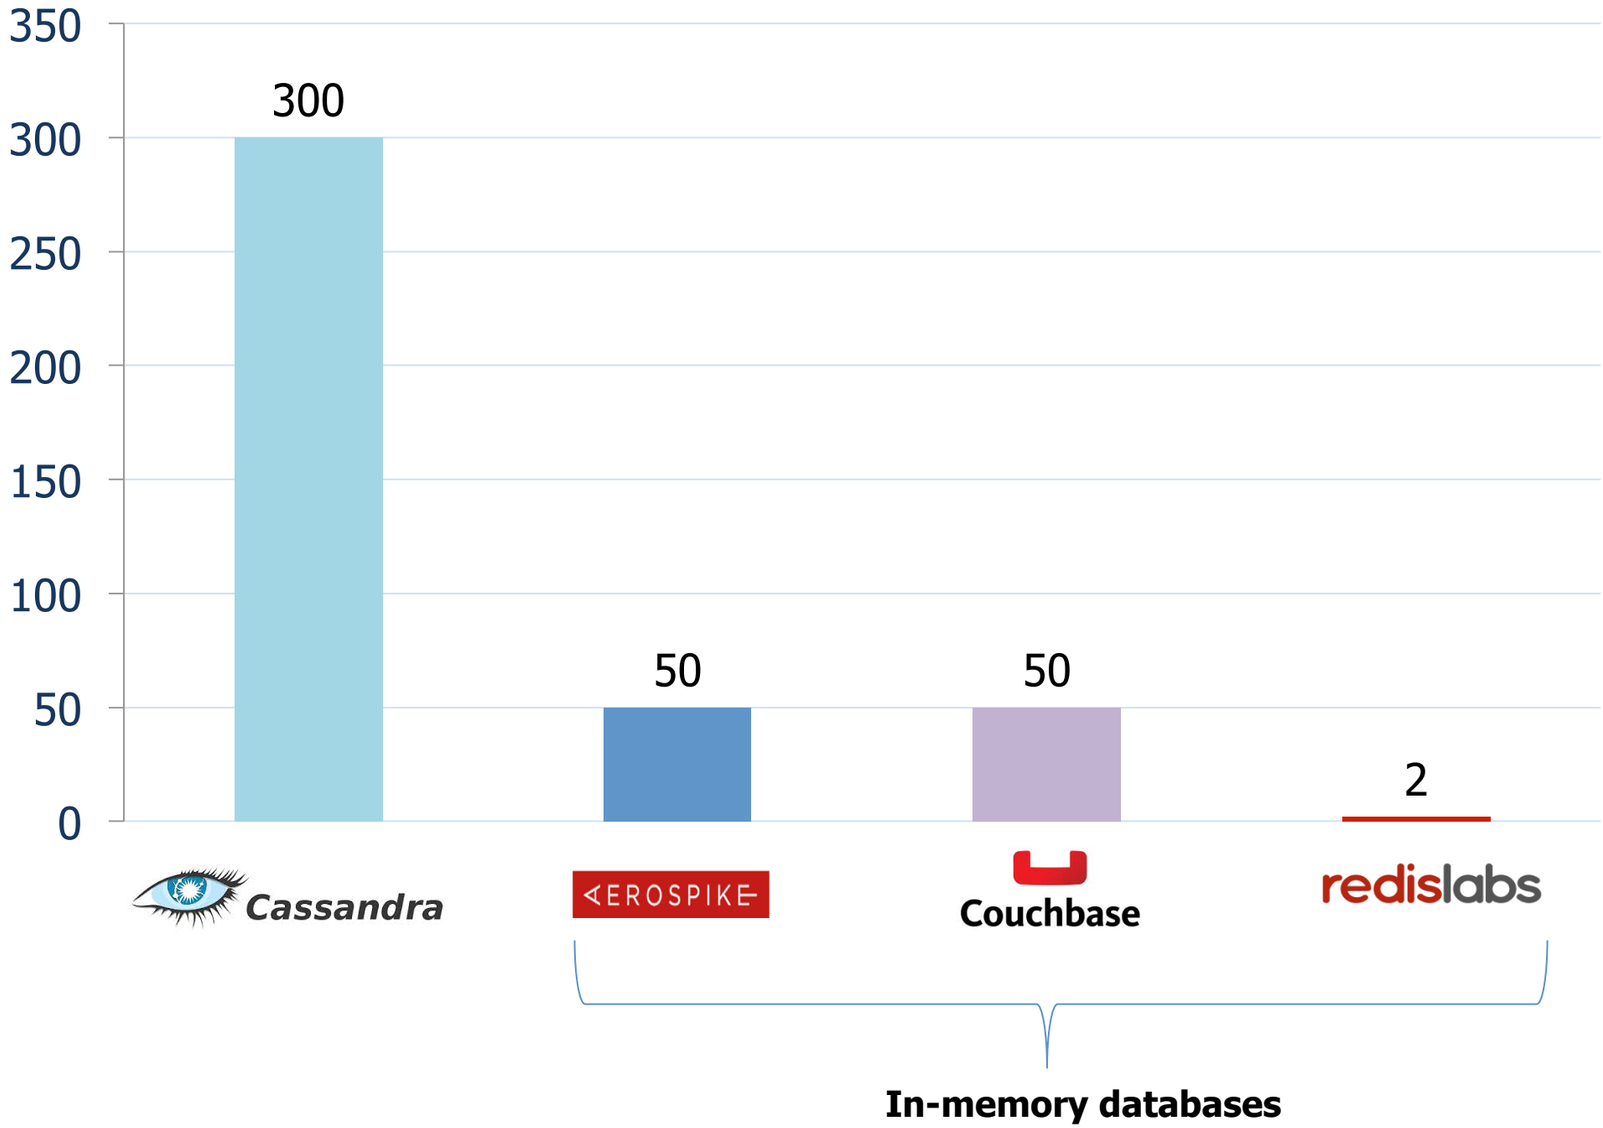

redislabs

Real-world write intensive app

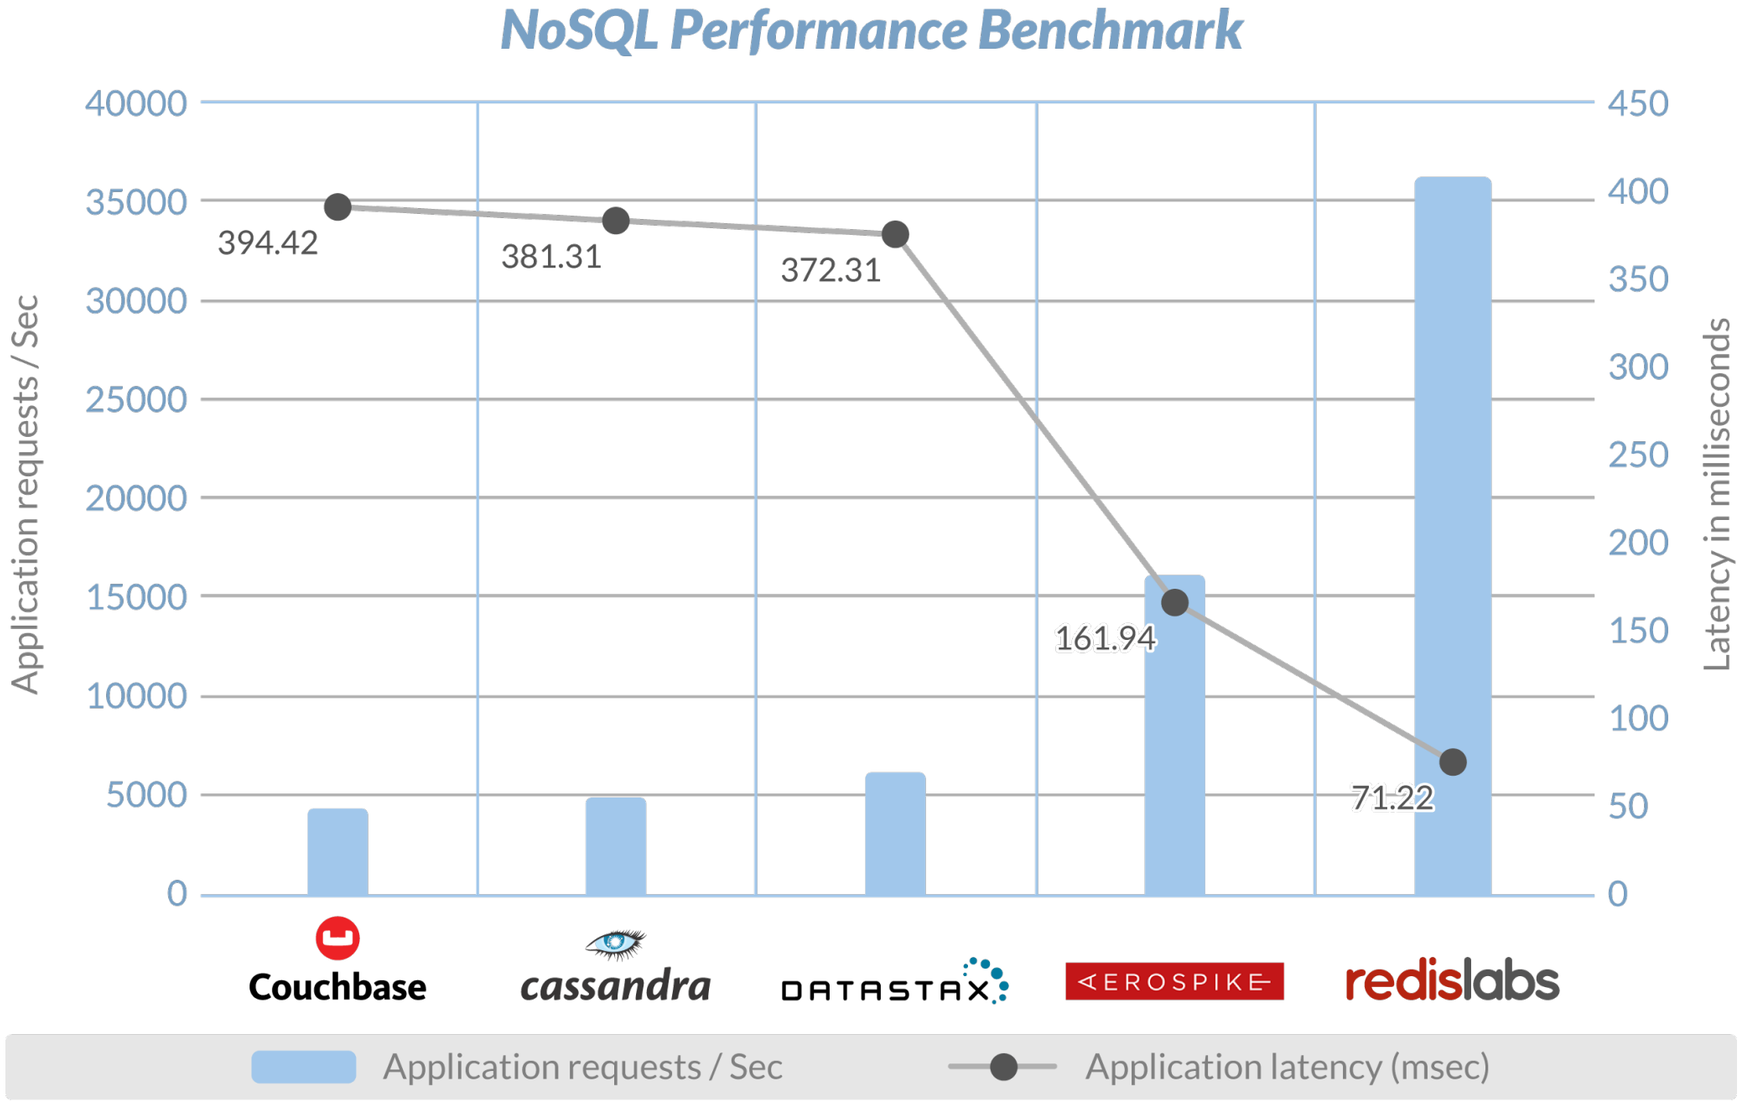

redislabs

Hbase+Internal Cache vs. Hbase+Redis

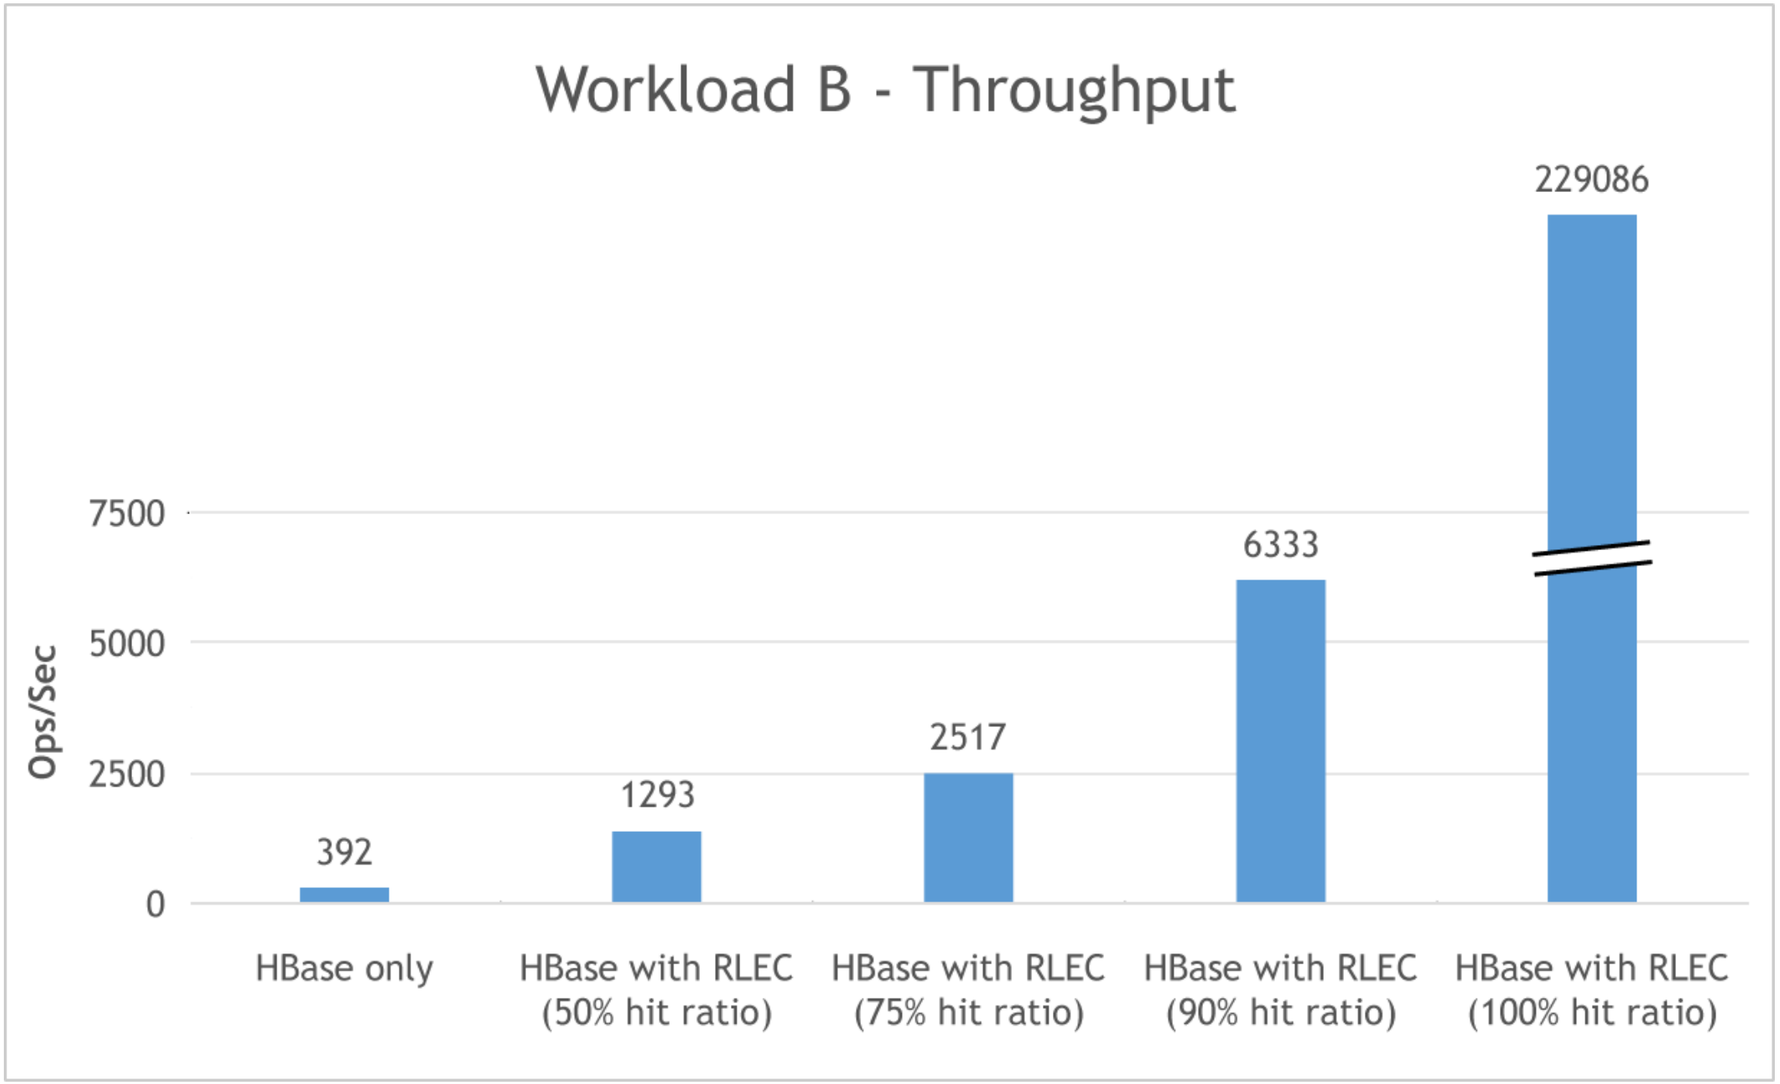

redislabs

Same Redis core, same HW, different performance

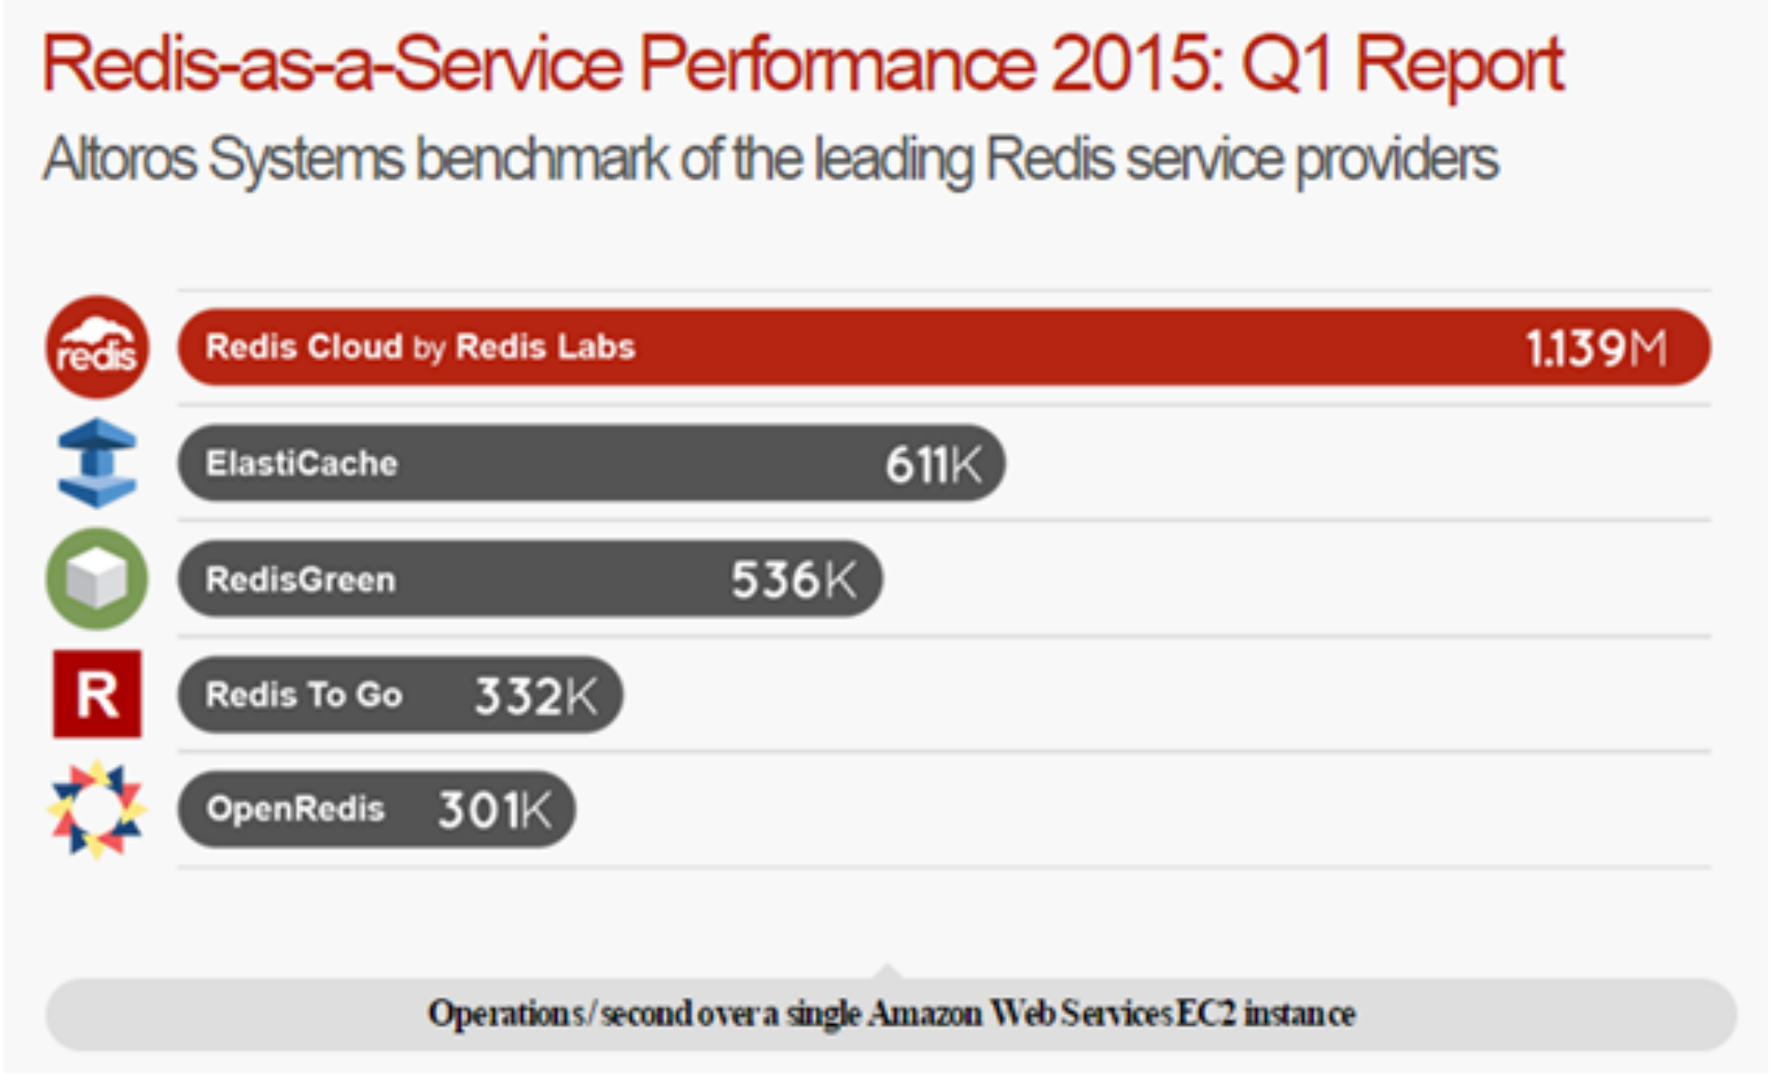

redislabs

So why aren't in-memory DBs equally fast?

Most are written in C/C++…. but programming language isn't the only thing to consider

redislabs

What affects in-memory DB performance?

(1) Complexity of processing commands

→ How many lines of code per command ? What is the computation complexity (e.g. in Redis most commands are O(1))?

(2) Query efficiency

→ Is it limited to blob queries? Can you query a discrete value?

(3) Pipelining

→ Can you send multiple requests at once to get lower latency and less context switches?

What affects in-memory DB performance?

(4) Protocol efficiency

→ How long it takes to parse a request or to serialize a response

(5) TCP overhead

→ Long-lived (connection pool) vs. short-lived connections

What affects in-memory DB performance?

(6) Single-threaded or multi-threaded architecture

→ Lock-free vs. parallel computing

(7) Shared-nothing (the best) vs. shared-something vs. shared-everything

(8) Built-in acceleration components

Myth #2

A single node is not a cluster

The truth:

A single node can be a cluster but not a HA cluster

In the new containers/VMs world a cluster is:

A bunch of processes that together look like one big process

A real-world example

A Binary Option platform; 400MB dataset; 1,000,000 ops/sec

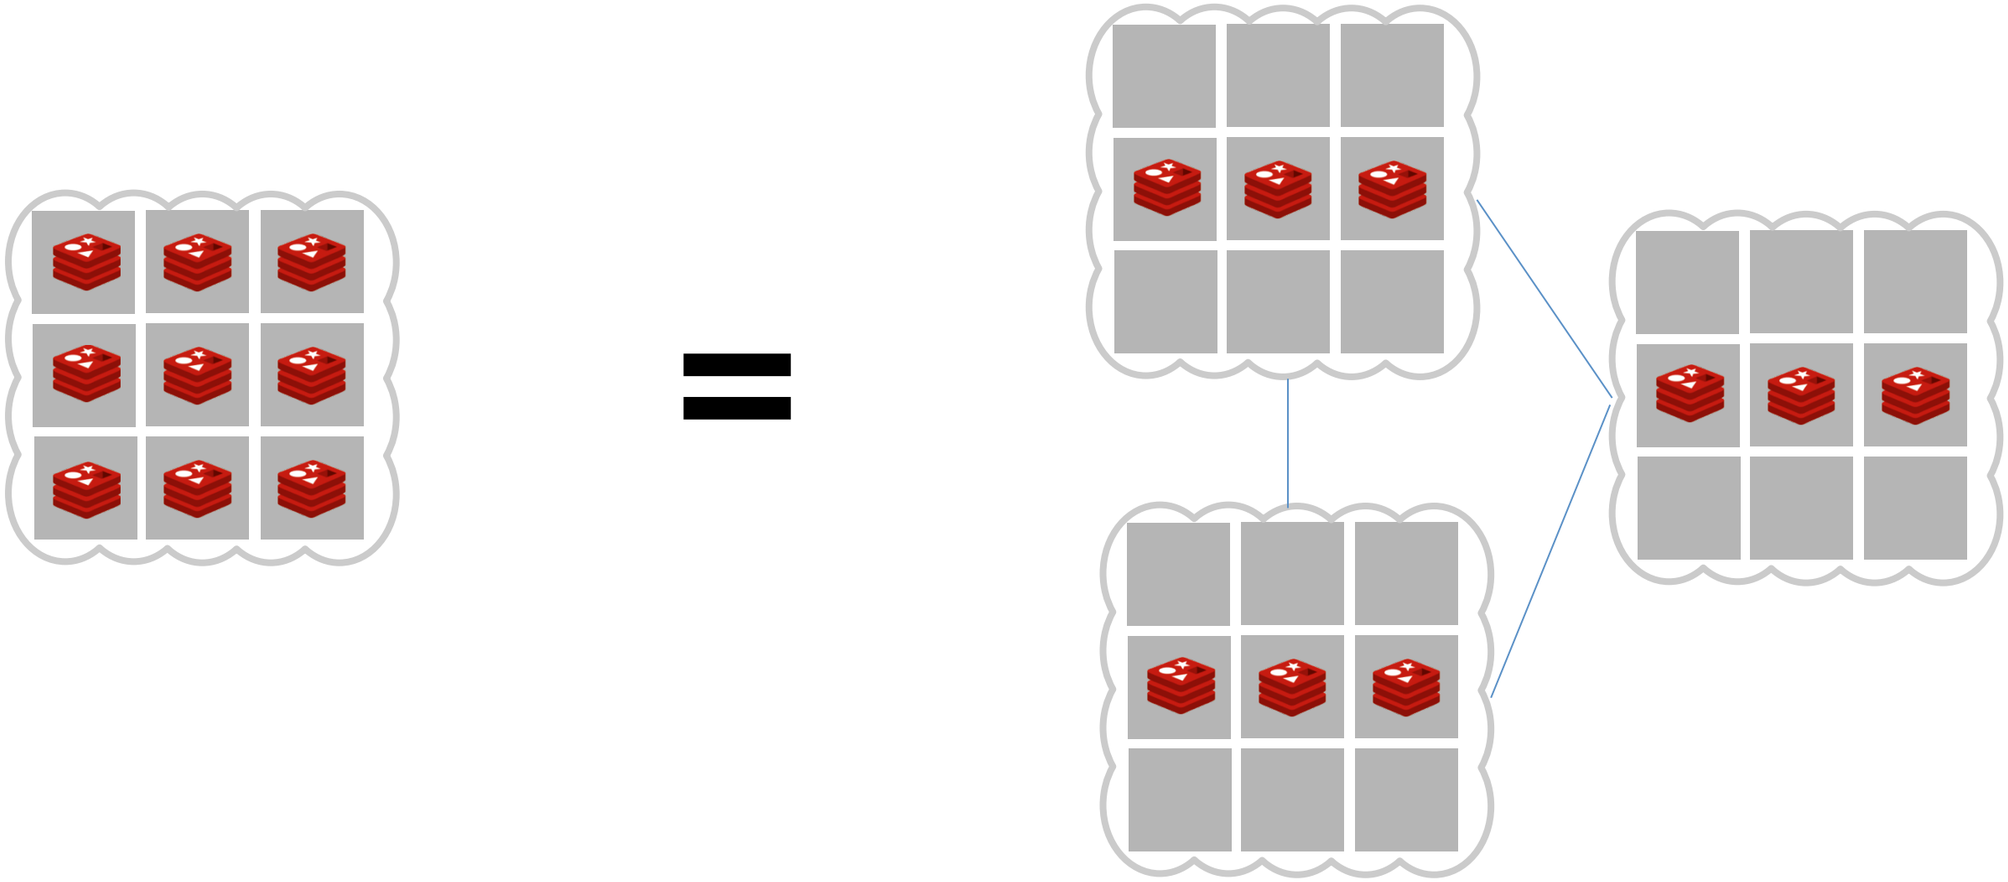

A single node, 9 shards cluster

3 nodes, 9 shards cluster

Myth #3

In-memory databases are inconsistent and unreliable

A few facts/questions about consistency

Almost all NoSQL databases (not just in-memory) *ack* the client before commiting to disk

Almost all in-memory databases can commit to disk before they *ack* the client

However, even if you *ack* after everything is committed:

- Is your driver memory buffer persistent and consistent?
- Is your storage system cache persistent and consistent?

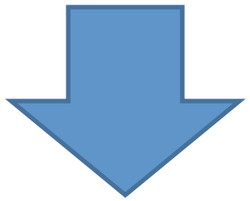

Most databases are NOT bulletproof

Most in-memory databases are async most of the time

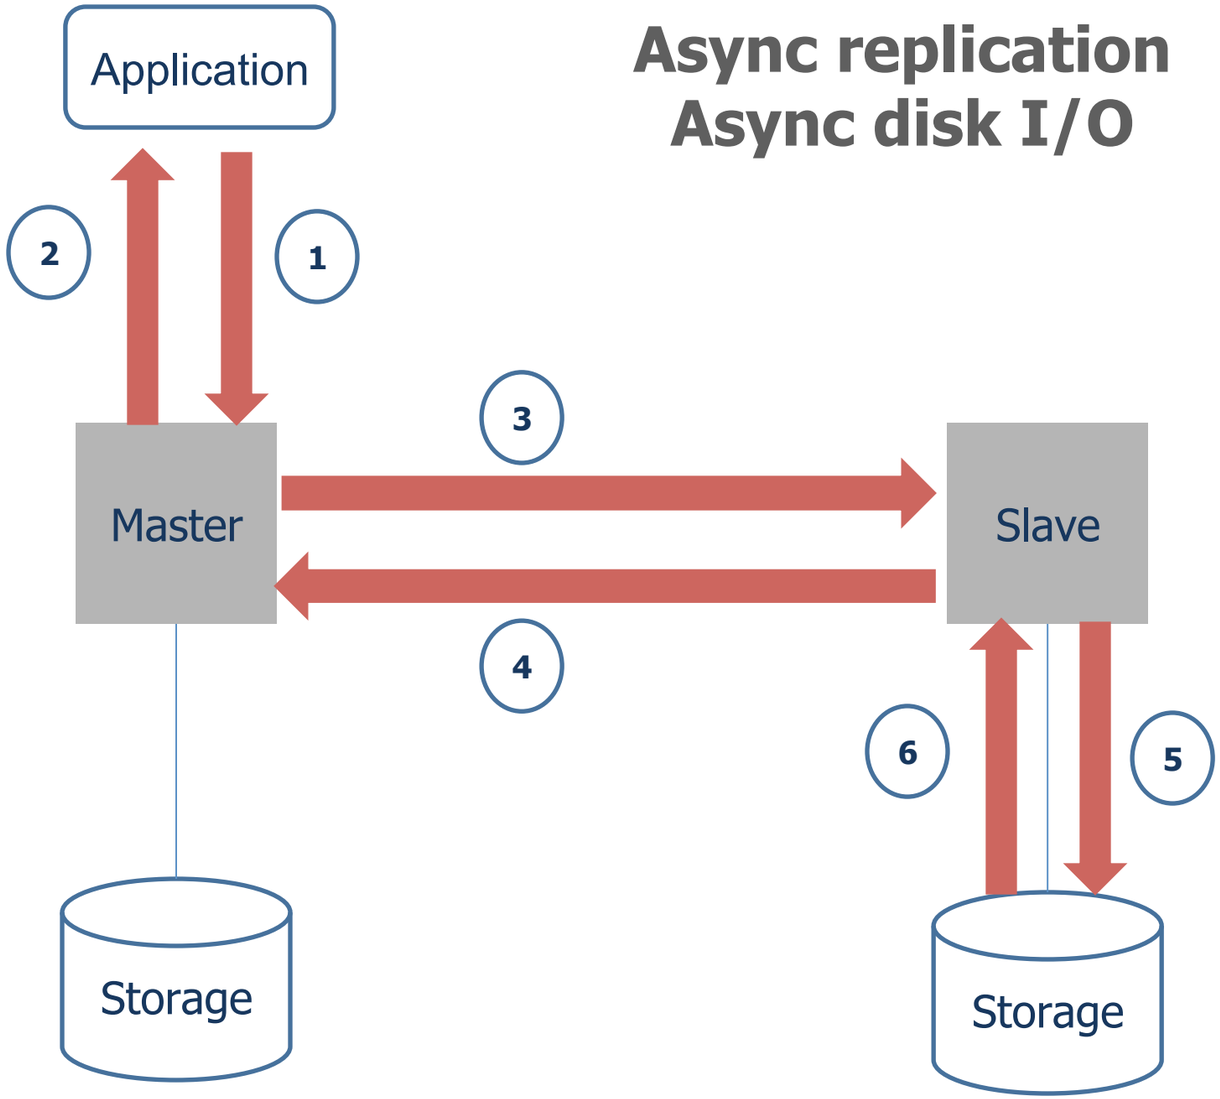

Some of them can partially sync

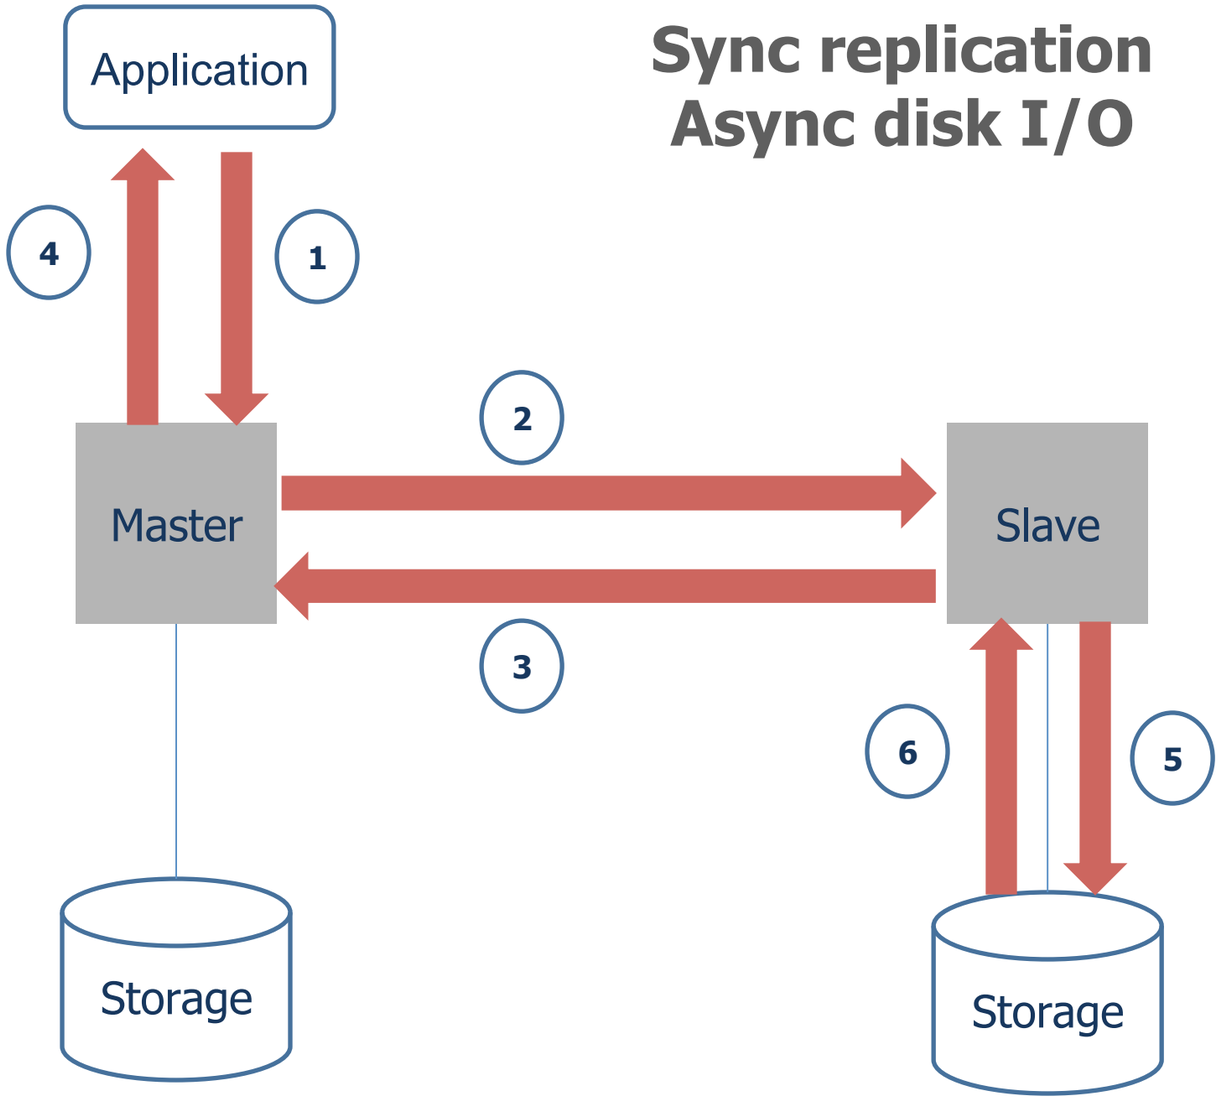

A few of them fully sync

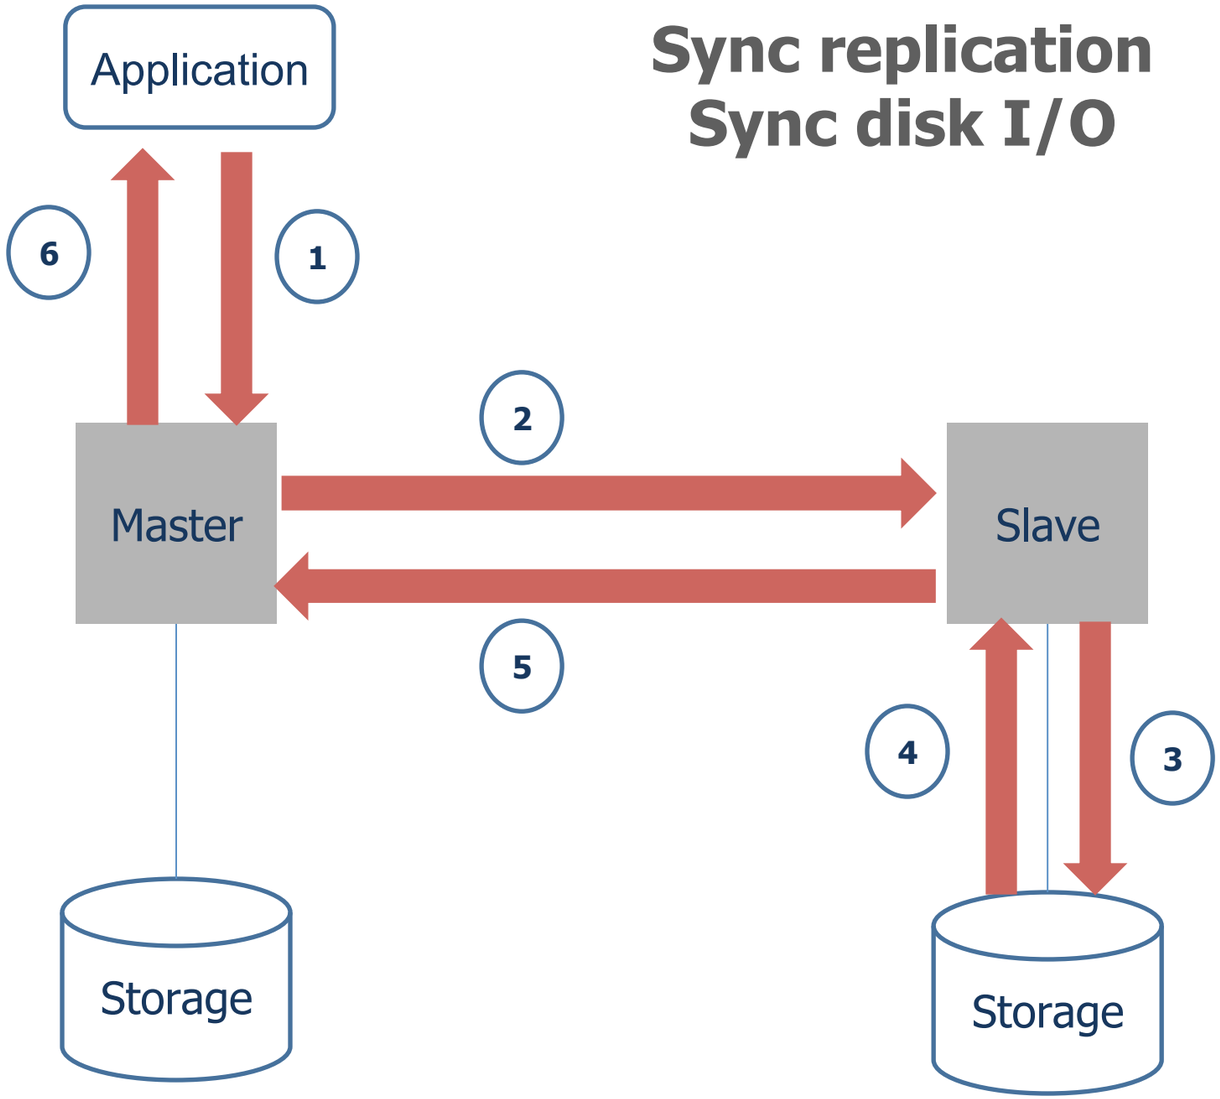

CAP and in-memory databases

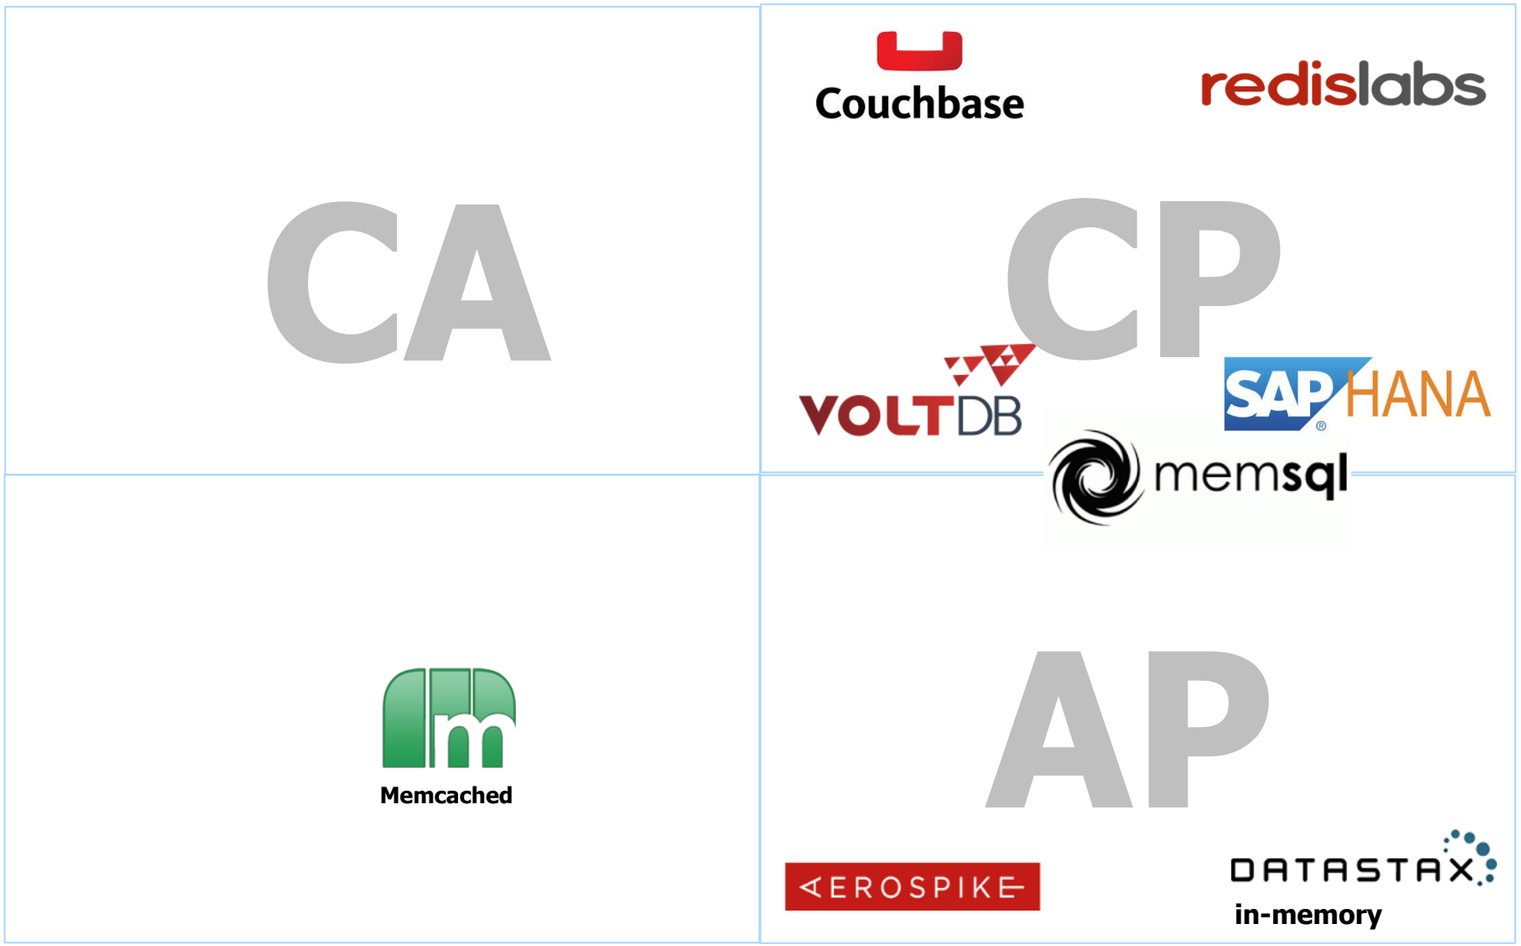

Behavior during network splits

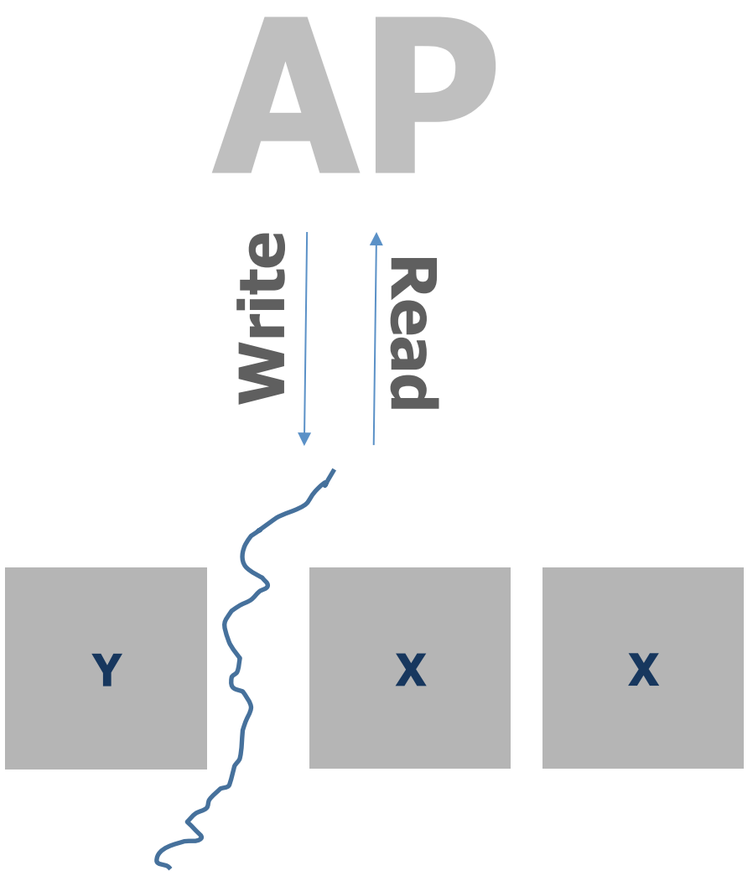

Inconsistency

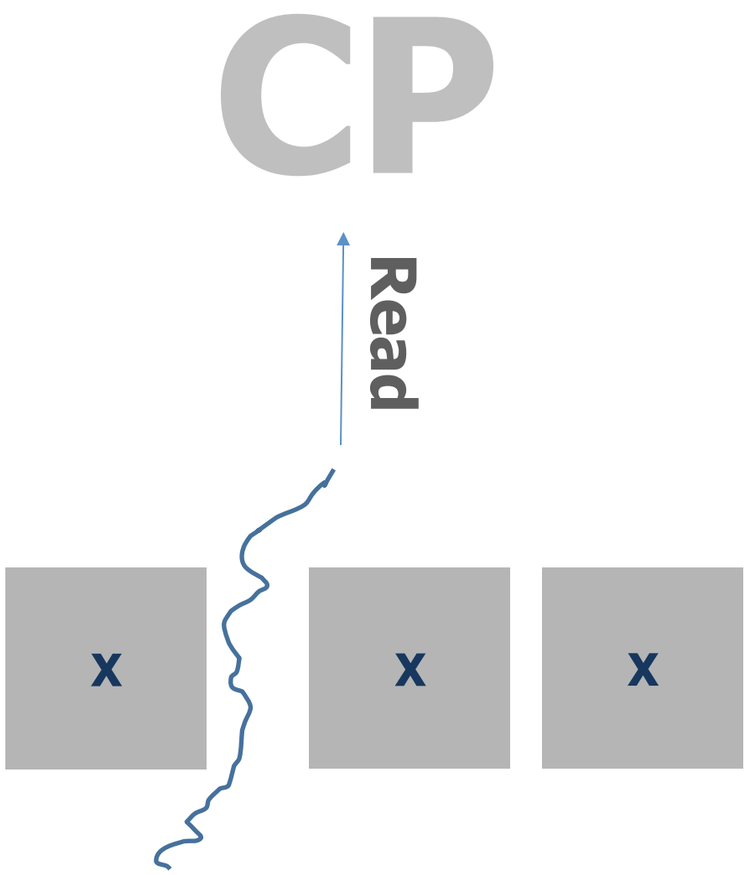

Full-consistency

Are in-memory databases reliable?

Redis Labs facts:

- Provisioned 100s of TBs of RAM
- **500+ node failure events → 1 failure every 2 days**
- ~30 complete data-center outages → 1 outage every month
- **Users with high availability (HA) features enabled haven't lost a single byte of data**

Myth #4
In-memory databases are expensive

Which one costs more?

Real-world use case:

- **500+GB**
- **400K writes/sec**
- **1500 reads/sec**
- **37.5KB average object size**

	Others
1.5Gbps	120Gbps
No extra work at app level	Tons of work at app level
6-node cluster	**30+ node cluster**

Which one costs more (2)?

Serving Over 1,000,000 Ops/s
with Google Compute Engine

	Redis Labs	Aerospike/ Couchbase	Cassandra
Read			
Write		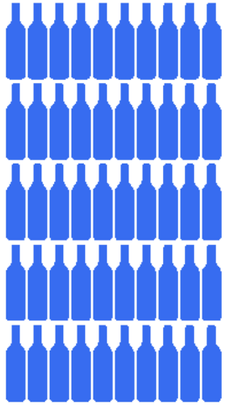	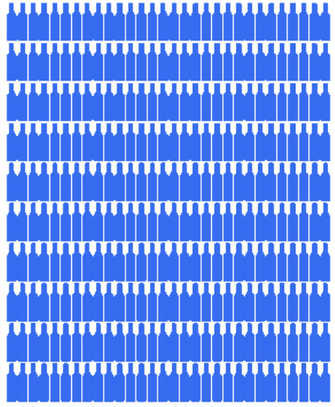
Cost			

redislabs

Sometimes in-memory can be very expensive

HA deployment of 10TB in-memory dataset on EC2

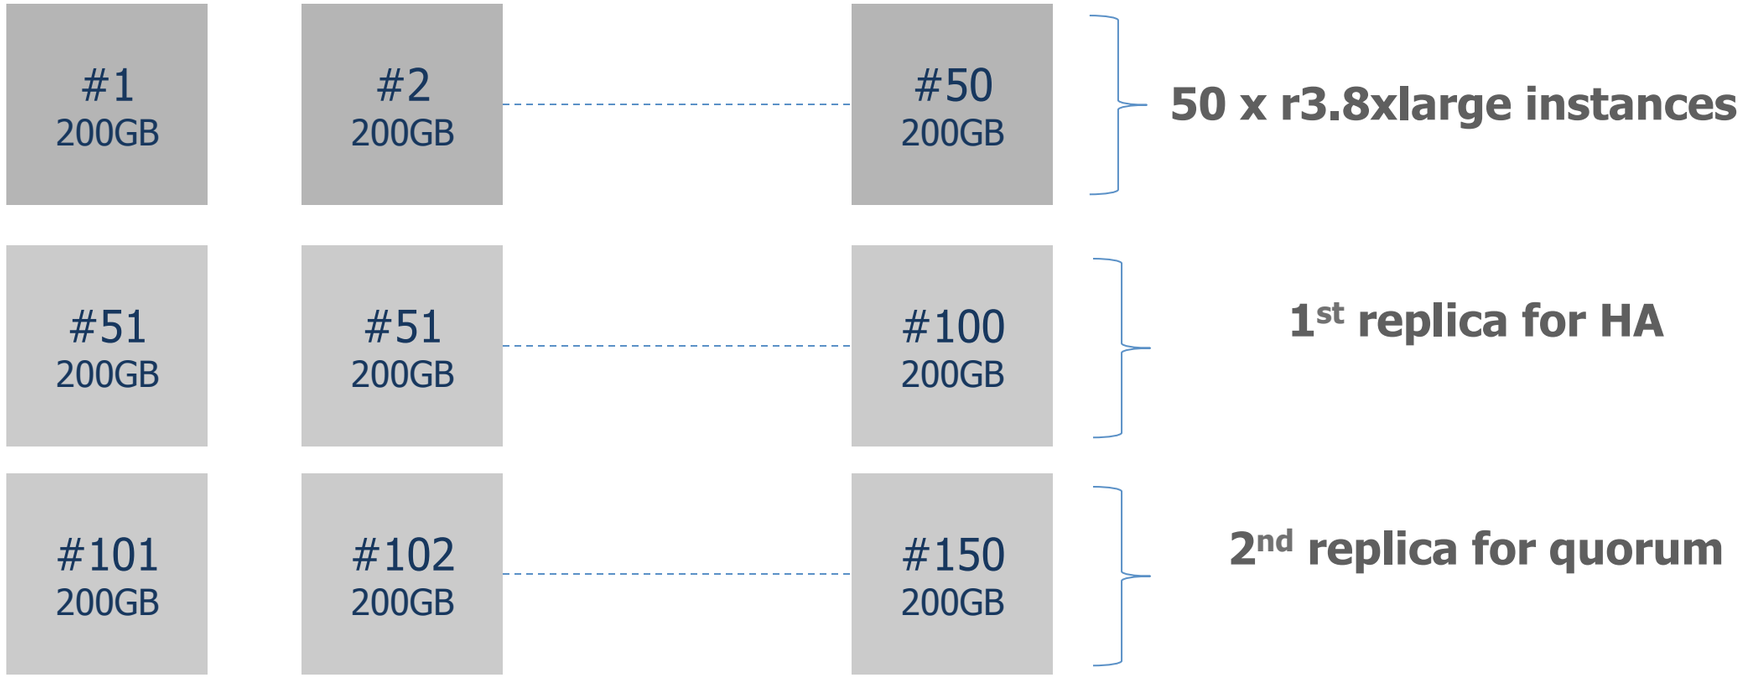

Total cost (reserved instances) = $2,132,250/yr

redislabs

Do we really need 2 replicas?

Efficient HA deployment of 10TB in-memory dataset on EC2

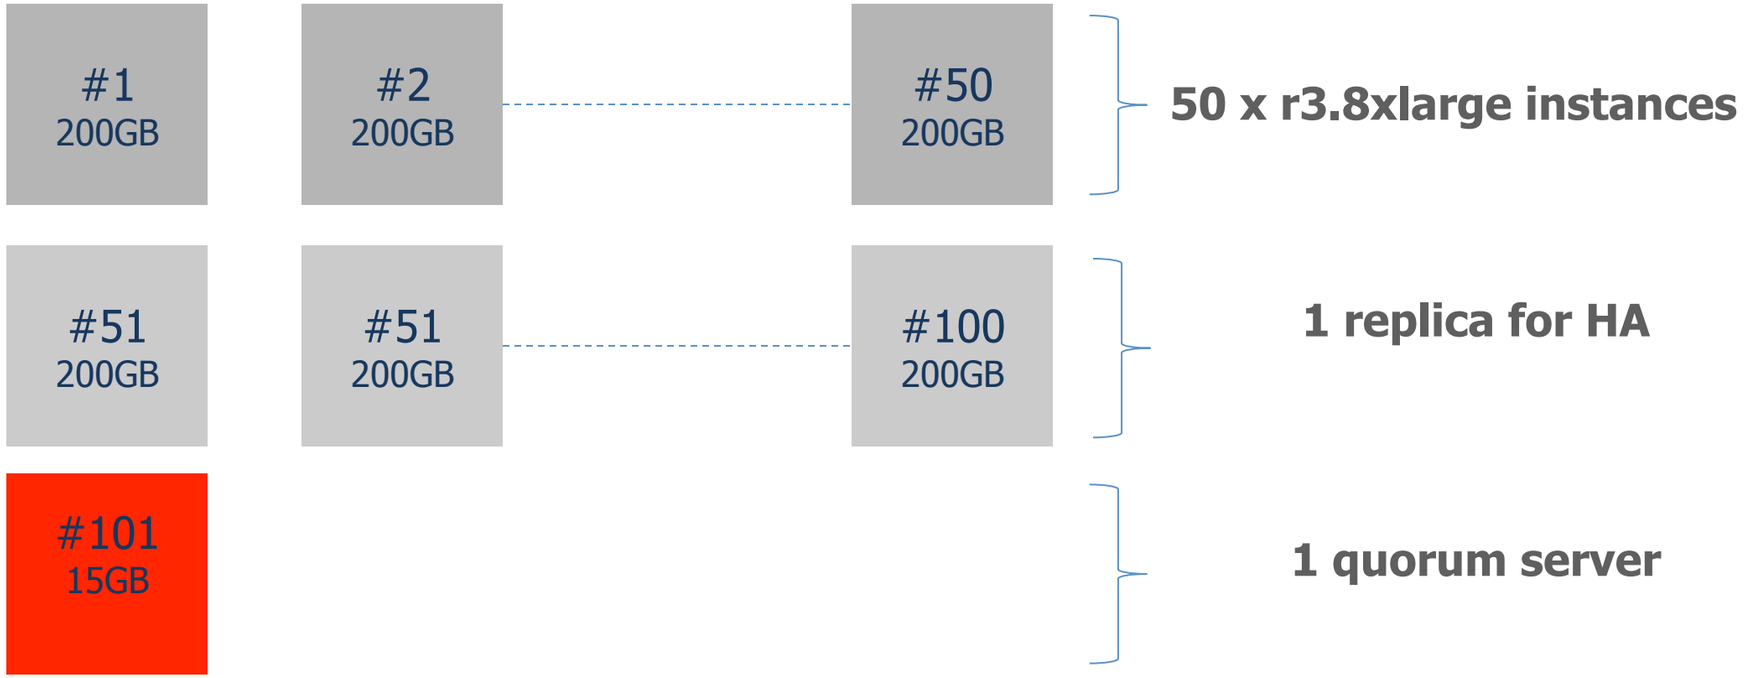

Total cost (reserved instances) = $1,421,500/yr **Savings = $710,750/yr**

Redis Labs Patent Pending Technology

redislabs

Can we save more?

Price/performance of memory technologies

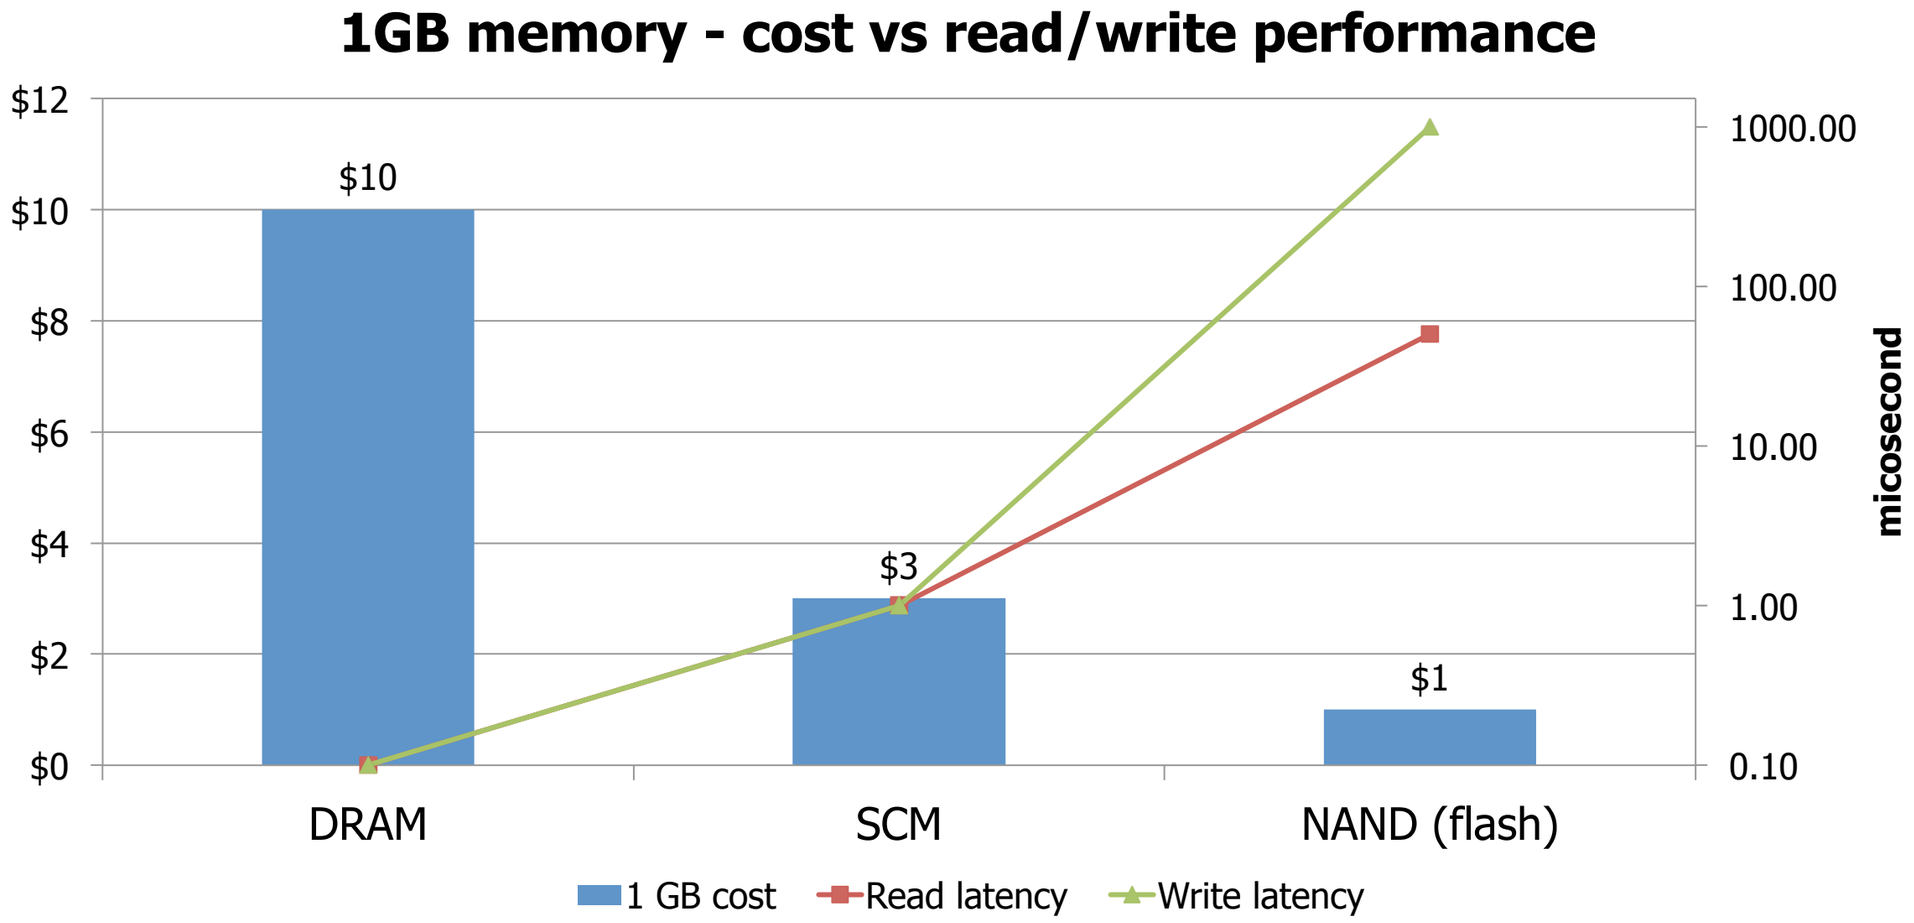

redislabs

Redis on Flash

Flash used as a RAM extender and NOT as persistent storage

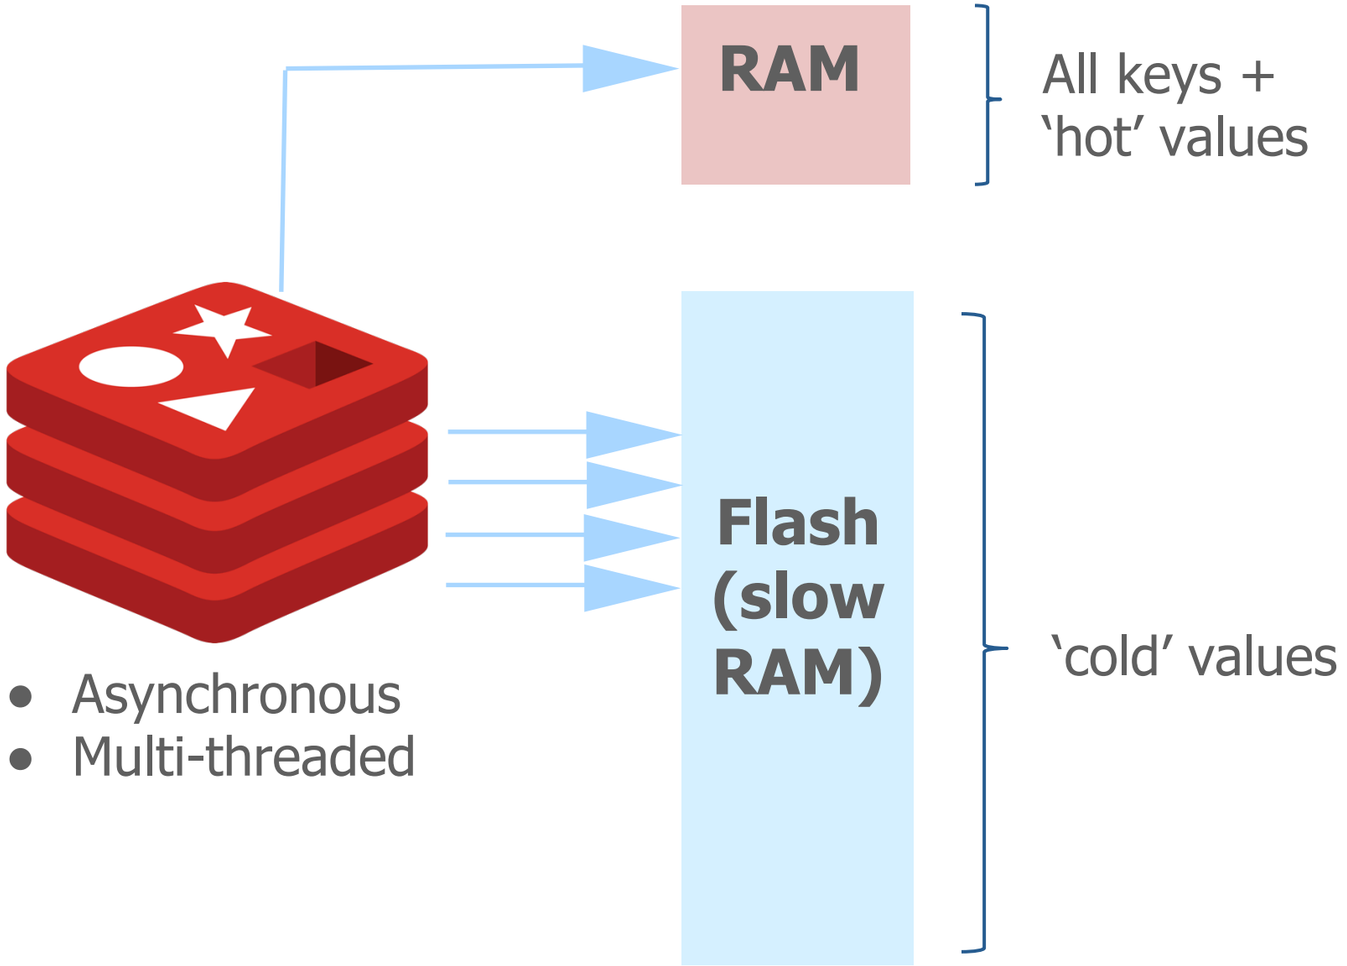

How to achieve optimal price/performance

By dynamically setting RAM/Flash ratio

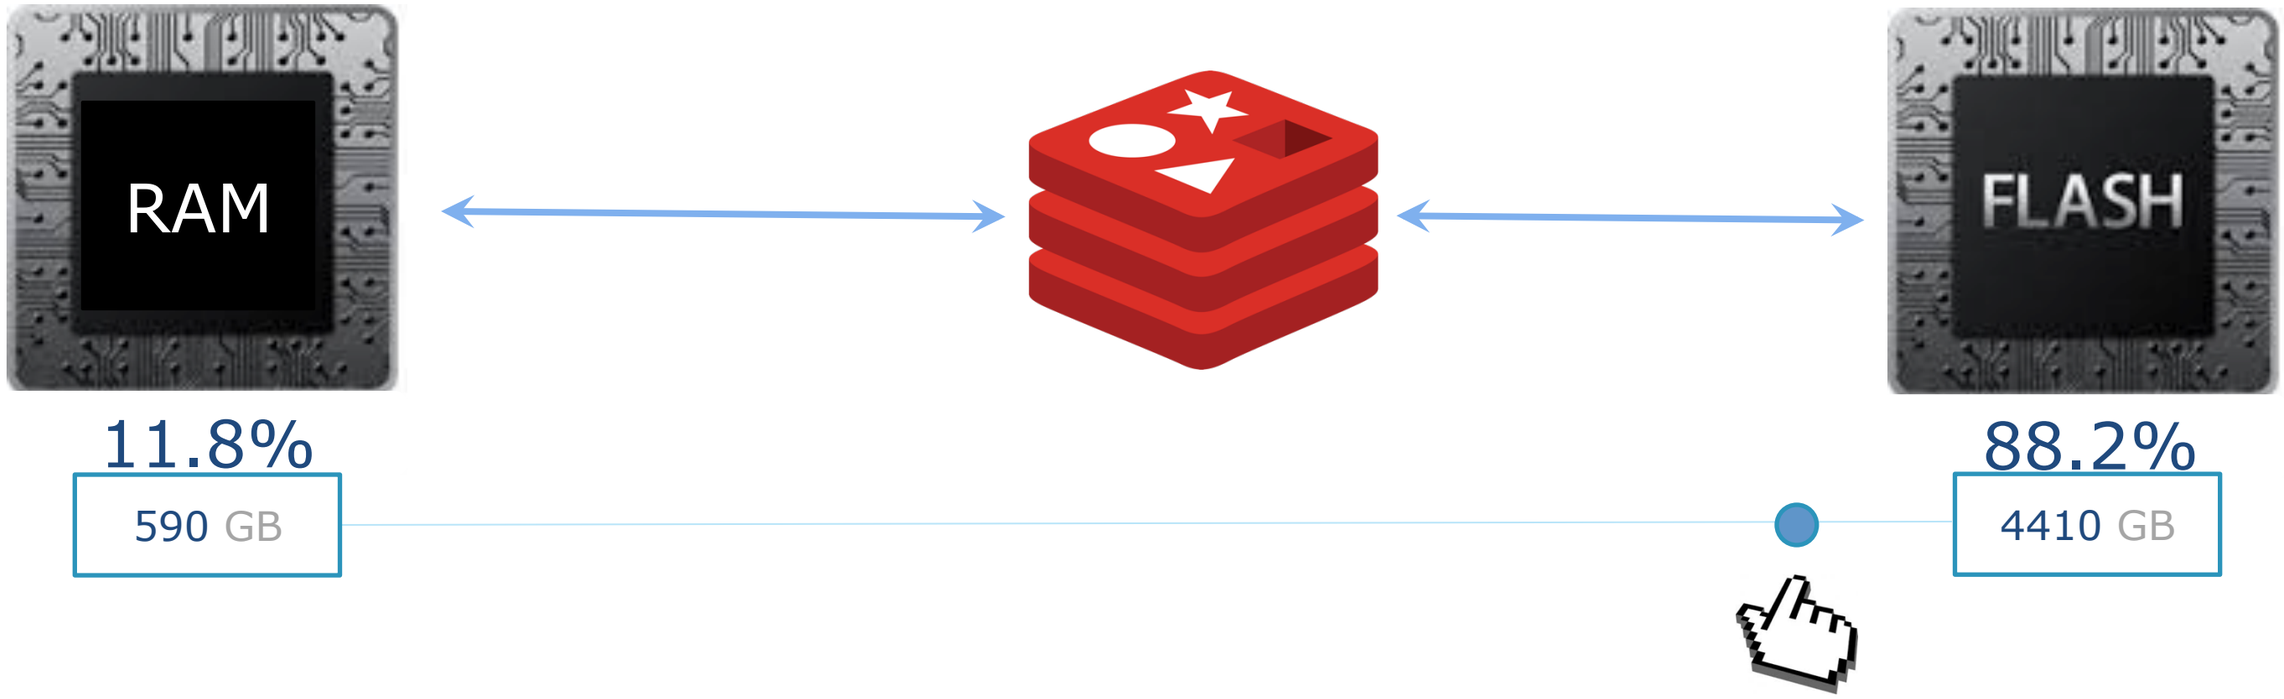

redislabs

Single server performance - 10% in RAM / 90% in Flash

RAM Hits Ratio	Ops/Sec	Latency
Low latency scenarios		
100%	1.35M	1.00 msec
80%	340K	1.07 msec
50%	200K	0.96 msec
20%	160K	1.00 msec
High throughput scenarios		
100%	2.00M	2.40 msec
80%	671K	6.20 msec
50%	483K	10.00 msec
20%	366K	14.50 msec

**IBM P8 with IBM FlashSystem*

10TB Redis deployment on EC2

	Redis (on RAM) 2 replicas	Redis (on RAM) 1 replicas	Redis on Flash 1 replica
Instance type	r3.8xlarge	r3.8xlarge	i2.8xlarge
# of instances	150	100	10
RAM	30TB	20TB	2TB
Flash	-	-	64TB
Persistent storage (EBS)	150TB	100TB	80TB
1yr costs (reserved instances)	**$2,132,250**	**$1,421,500**	**$318,090**
Yearly savings	-	$710,750	$1,814,160
Savings %	-	33.33%	**85.08%**

Summary

4 myths about in-memory databases busted

All in-memory databases are NOT equally fast

You can create a single node in-memory cluster

In-memory databases can be consistent and reliable

With the right technology, in-memory databases are not expensive

Thank you

Click to get more info about
Redis, Redis Labs, Redis Cloud and RLEC

redislabs