



Introduction to Spark SQL & Catalyst

Takuya UESHIN

Spark Meetup 2014/09/08(Mon)

Who am I?

- Takuya UESHIN
 - @ueshin
 - github.com/ueshin
- Nautilus Technologies, Inc.
- A Spark contributor




Agenda

- What is Spark SQL?
- Catalyst in depth
- SQL core in depth
- Interesting issues
- How to contribute

What is Spark SQL?

What is Spark SQL?

- Spark SQL is one of Spark components.



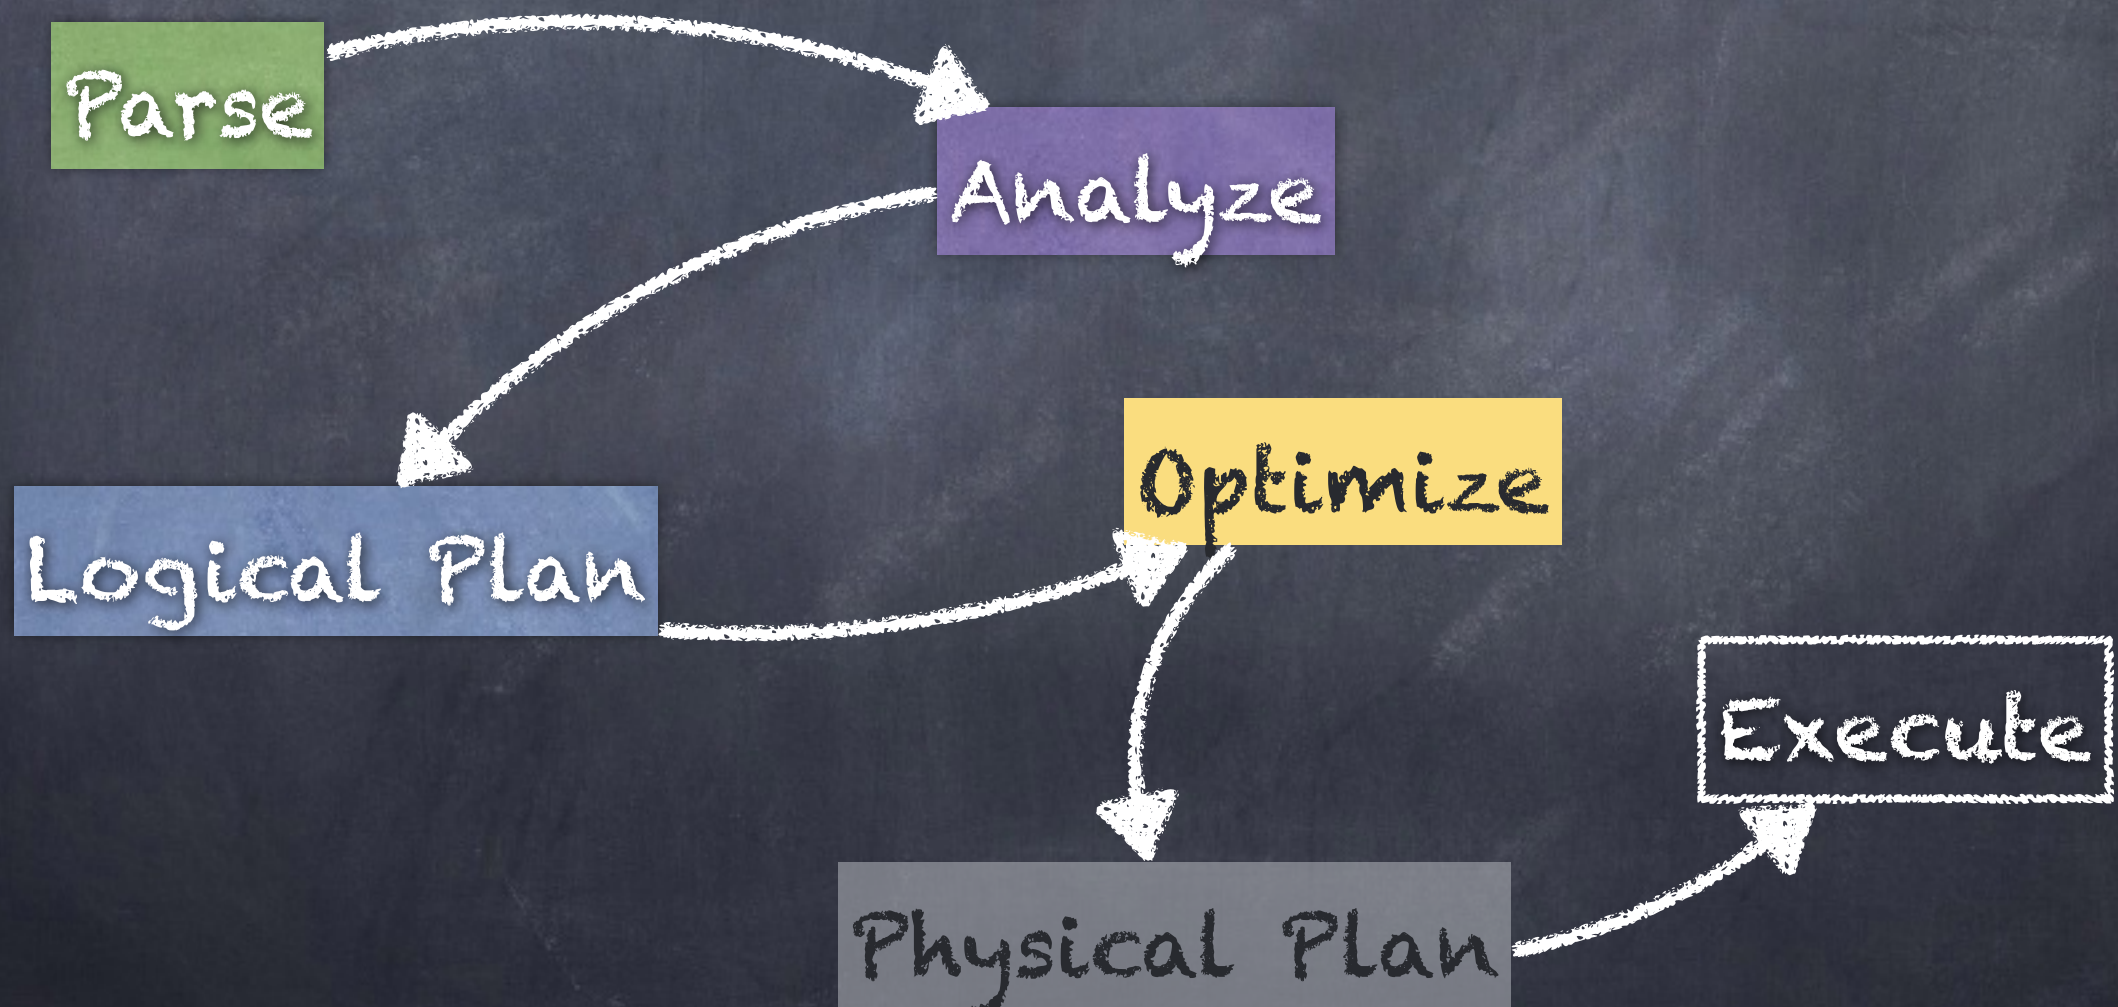
- Executes SQL on Spark
- Builds SchemaRDD like LINQ
- Optimizes execution plan.

What is Spark SQL?

- Catalyst provides a execution planning framework for relational operations.
- Including:
 - SQL parser & analyzer
 - Logical operators & general expressions
 - Logical optimizer
 - A framework to transform operator tree.

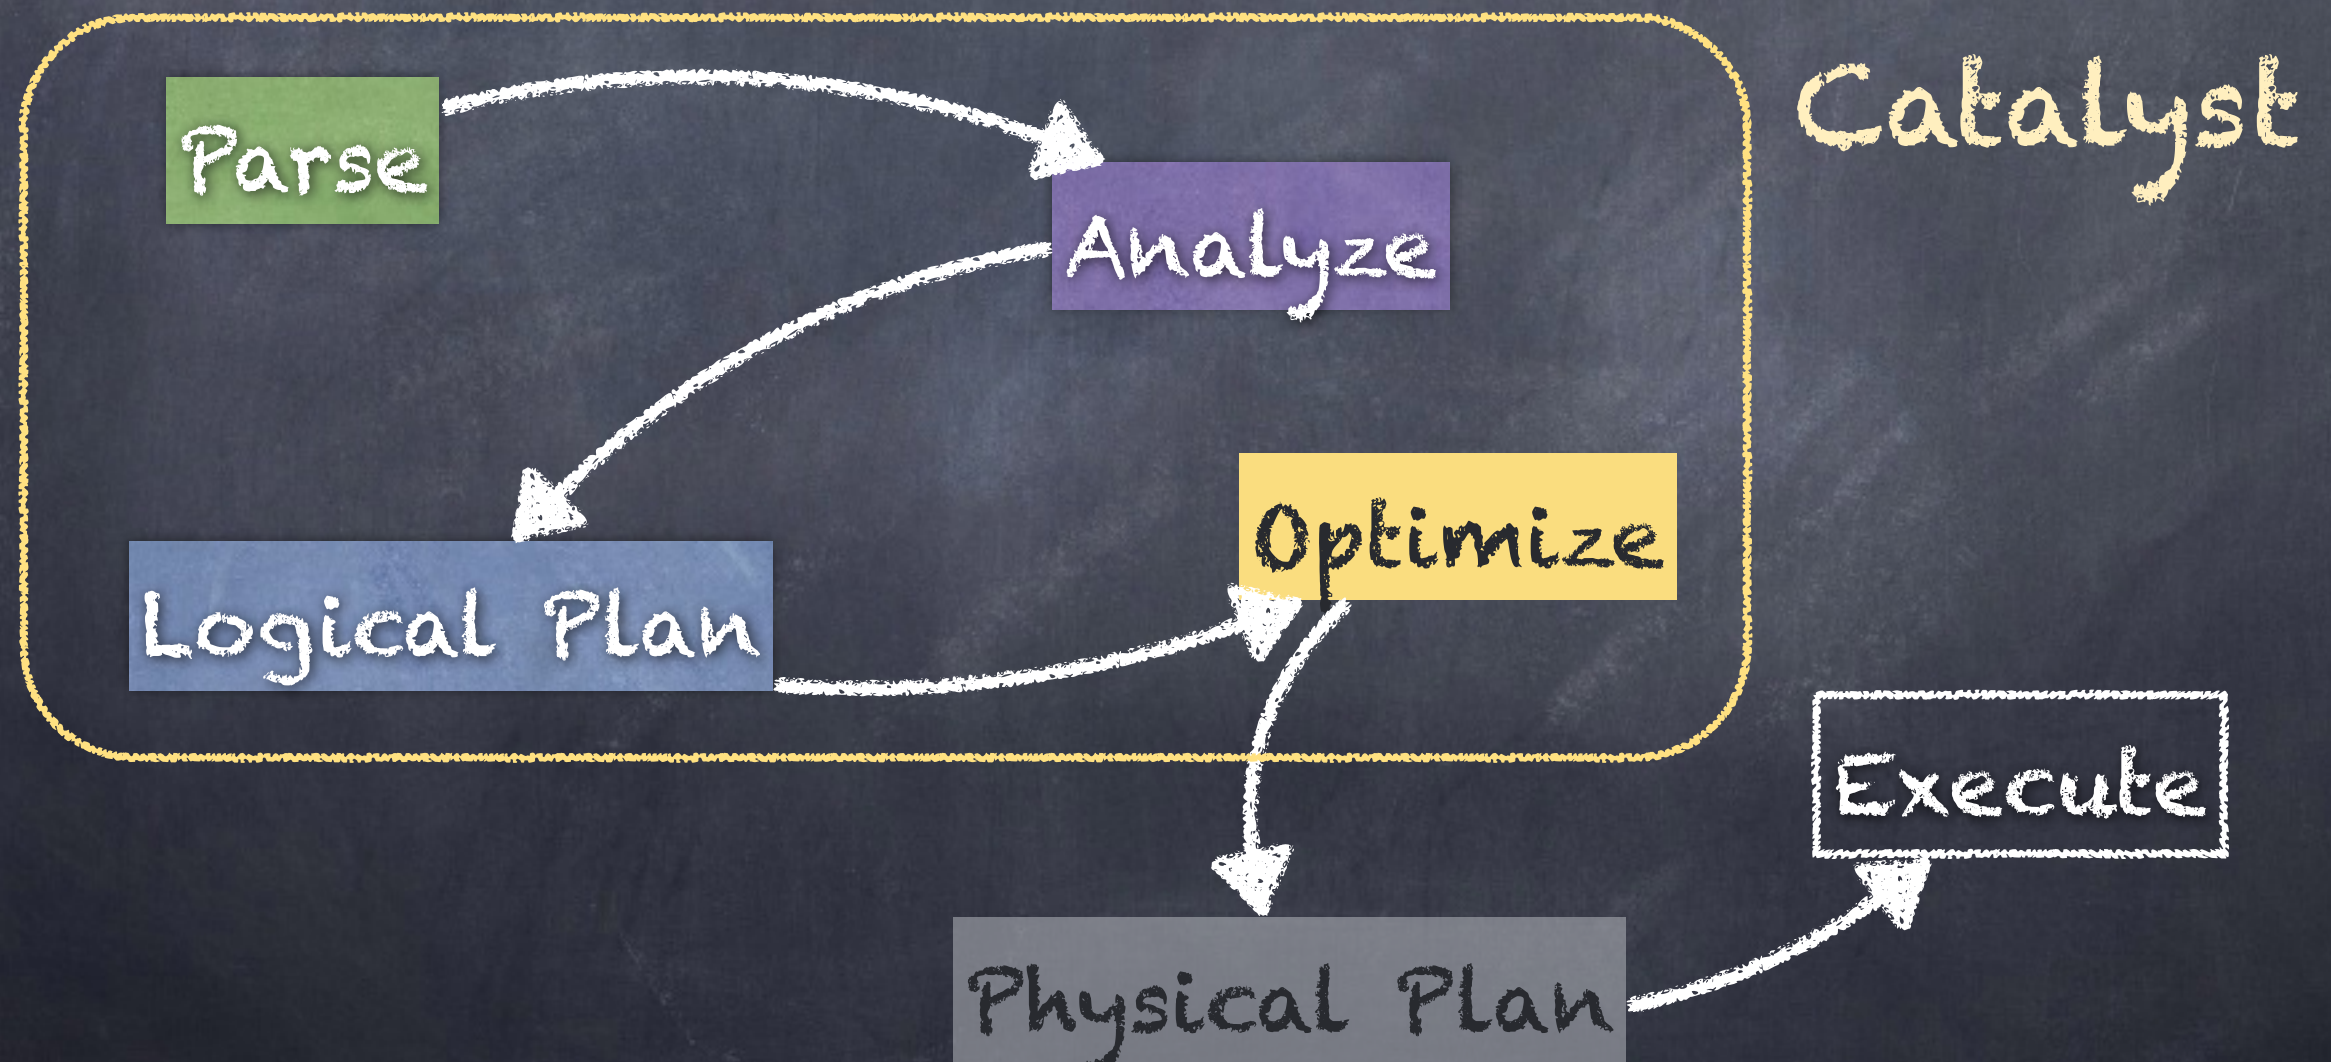
What is Spark SQL?

- To execute query needs some steps.



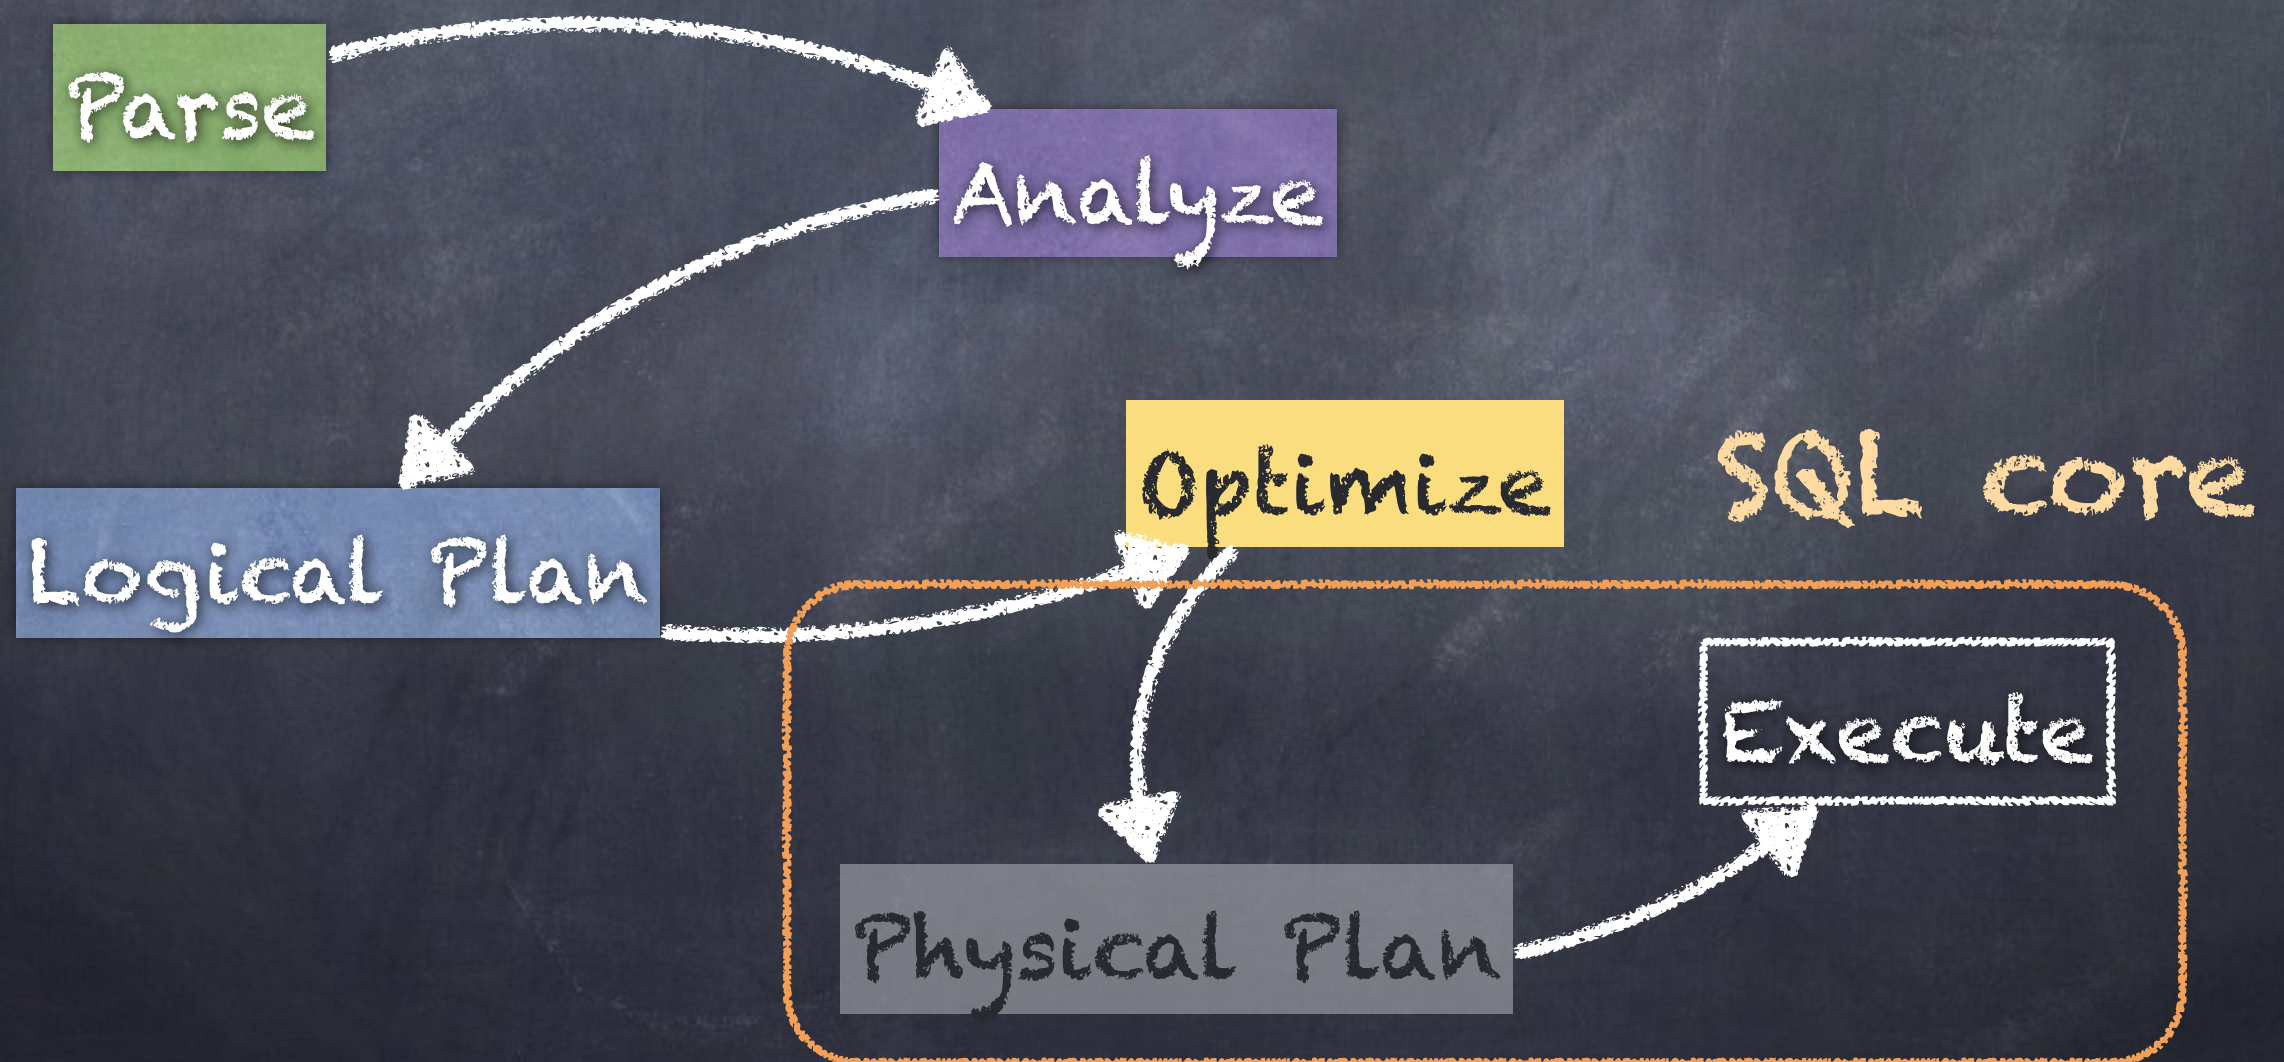
What is Spark SQL?

- To execute query needs some steps.



What is Spark SQL?

- To execute query needs some steps.



Catalyst in depth

Catalyst in depth

- Provides a execution planning framework for relational operations.
 - Row & Data Type's
 - Trees & Rules
 - Logical Operators
 - Expressions
 - Optimizations

Row & Data Type's

- `o.a.s.sql.catalyst.types.DataType`
 - Long, Int, Short, Byte, Float, Double, Decimal
 - String, Binary, Boolean, Timestamp
 - Array, Map, Struct
- `o.a.s.sql.catalyst.expressions.Row`
 - Represents a single row.
 - Can contain complex types.

Trees & Rules


- `o.a.s.sql.catalyst.trees.TreeNode`
- Provides transformations of tree.
 - `foreach`, `map`, `flatMap`, `collect`
 - `transform`, `transformUp`,
`transformDown`
- Used for operator tree, expression tree.

Trees & Rules

- `o.a.s.sql.catalyst.rules.Rule`
 - Represents a tree transform rule.
- `o.a.s.sql.catalyst.rules.RuleExecutor`
 - A framework to transform trees based on rules.


Logical Operators

- Basic Operators
 - Project, Filter, ...
- Binary Operators
 - Join, Except, Intersect, Union, ...
- Aggregate
- Generate, Distinct
- Sort, Limit
- InsertInto, WriteToFile



Expressions

- Literal
- Arithmetics
 - UnaryMinus, Sqrt, MaxOf
 - Add, Subtract, Multiply, ...
- Predicates
 - EqualTo, LessThan, LessThanOrEqual, GreaterThan, GreaterThanOrEqual
 - Not, And, Or, In, If, CaseWhen



Expressions

- Cast
- GetItem, GetField
- Coalesce, IsNull, IsNotNull
- StringOperations
 - Like, Upper, Lower, Contains, StartsWith, EndsWith, Substring, ...

Optimizations

- ConstantFolding
 - NullPropagation
 - ConstantFolding
 - BooleanSimplification
 - SimplifyFilters
- FilterPushdown
 - CombineFilters
 - PushPredicateThroughProject
 - PushPredicateThroughJoin
 - ColumnPruning

Optimizations

- NullPropagation, ConstantFolding
 - Replace expressions that can be evaluated with some literal value to the value.
 - ex)
 - $1 + \text{null} \Rightarrow \text{null}$
 - $1 + 2 \Rightarrow 3$
 - $\text{Count}(\text{null}) \Rightarrow 0$

Optimizations

- BooleanSimplification

- Simplifies boolean expressions that can be determined.

- ex)

- $\text{false AND } \$right \Rightarrow \text{false}$

- $\text{true AND } \$right \Rightarrow \$right$

- $\text{true OR } \$right \Rightarrow \text{true}$

- $\text{false OR } \$right \Rightarrow \$right$

- $\text{If}(\text{true}, \$then, \$else) \Rightarrow \$then$

Optimizations

- SimplifyFilters

- Removes filters that can be evaluated trivially.

- ex)

- $\text{Filter}(\text{true}, \text{child}) \Rightarrow \text{child}$

- $\text{Filter}(\text{false}, \text{child}) \Rightarrow \text{empty}$

Optimizations

- CombineFilters

- Merges two filters.

- ex)

- $\text{Filter}(\$fc, \text{Filter}(\$nc, \text{child})) \Rightarrow \text{Filter}(\text{AND}(\$fc, \$nc), \text{child})$

Optimizations

- PushPredicateThroughProject
 - Pushes Filter operators through Project operator.
 - ex)
 - $\text{Filter('i' === 1, Project('i, 'j, child))}$
 \Rightarrow
 $\text{Project('i, 'j, Filter('i' === 1, child))}$

Optimizations

- PushPredicateThroughJoin
 - Pushes Filter operators through Join operator.
 - ex)
 - $\text{Filter}(\text{"left.i.attr"} == 1, \text{Join}(\text{left}, \text{right})) \Rightarrow$
 $\text{Join}(\text{Filter}(\text{'i'} == 1, \text{left}), \text{right})$

Optimizations

- Column Pruning

- Eliminates the reading of unused columns.

- ex)

- $\text{Join}(\text{left}, \text{right}, \text{LeftSemi},$
 $\text{"left.id".attr} === \text{"right.id".attr}) \Rightarrow$
 $\text{Join}(\text{left}, \text{Project('id', right)}, \text{LeftSemi})$

SQL core in depth

SQL core in depth

- Provides:
 - Physical operators to build RDD
 - Conversion from Existing RDD of Product to SchemaRDD support
 - Parquet file read/write support
 - JSON file read support
 - Columnar in-memory table support

SQL core in depth

- `o.a.s.sql.SchemaRDD`
 - Extends `RDD[Row]`.
 - Has logical plan tree.
 - Provides LINQ-like interfaces to construct logical plan.
 - `select`, `where`, `join`, `orderBy`, ...
 - Executes the plan.

SQL core in depth

- o.a.s.sql.execution.SparkStrategies
 - Converts logical plan to physical.
 - Some rules are based on statistics of the operators.

SQL core in depth

- Parquet read/write support
 - Columnar storage format for Hadoop
 - Reads existing Parquet files.
 - Converts Parquet schema to row schema.
 - Writes new Parquet files.
 - Currently DecimalType and TimestampType are not supported.

SQL core in depth

- JSON read support
 - Loads a JSON file (one object per line)
 - Infers row schema from the entire dataset.
 - Giving the schema is experimental.
 - Inferring the schema by sampling is also experimental.

SQL core in depth

- Columnar in-memory table support
 - Caches table like RDD.cache, but as columnar style.
 - Can prune unnecessary columns when read data.

Interesting issues

Interesting issues

- Support the GroupingSet/ROLLUP/CUBE
 - <https://issues.apache.org/jira/browse/SPARK-2663>
- Use statistics to skip partitions when reading from in-memory columnar data
 - <https://issues.apache.org/jira/browse/SPARK-2961>

Interesting issues

- Pluggable interface for shuffles

- <https://issues.apache.org/jira/browse/SPARK-2044>

- Sort-merge join

- <https://issues.apache.org/jira/browse/SPARK-2213>

- Cost-based join reordering

- <https://issues.apache.org/jira/browse/SPARK-2216>

How to contribute

How to contribute

- See: Contributing to Spark
- Open an issue on JIRA
- Send pull-request at GitHub
- Communicate with committers and reviewers
- Congratulations!

Conclusion

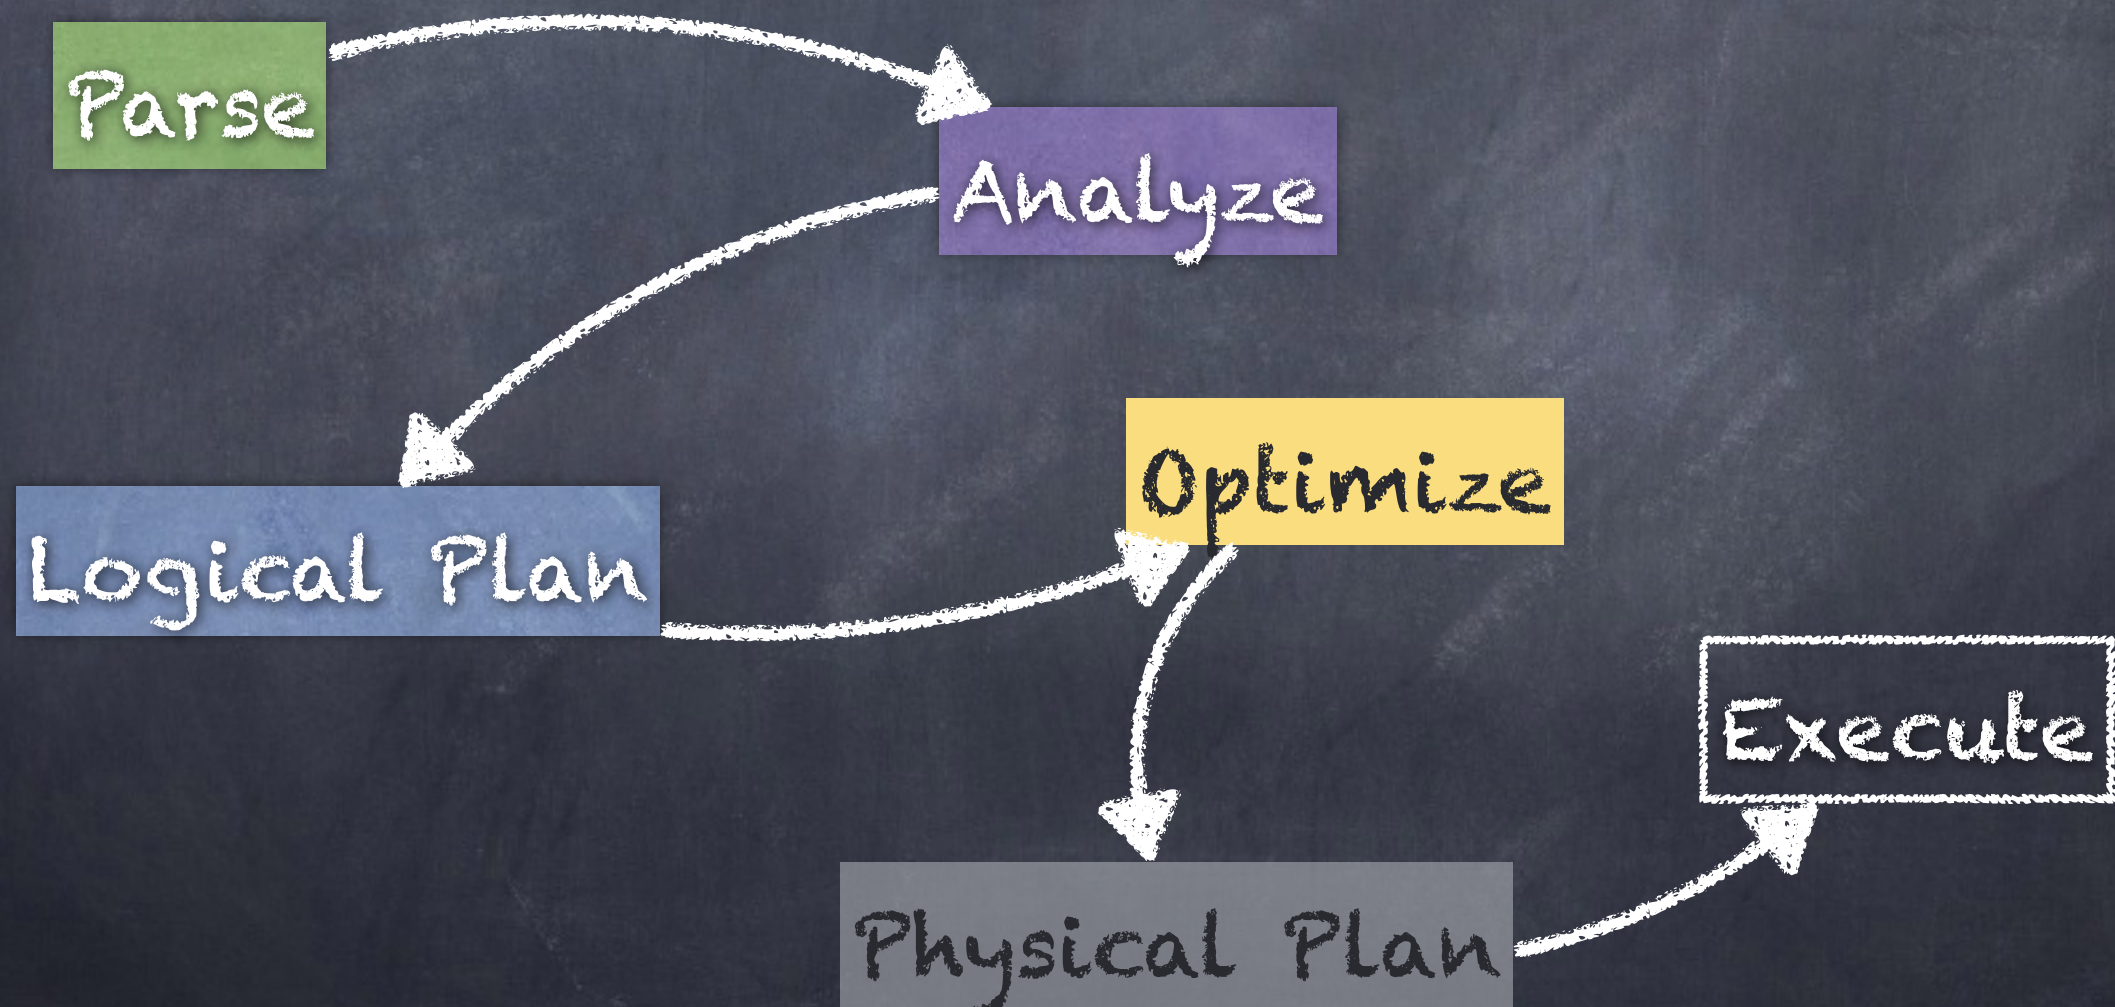
- Introduced Spark SQL & Catalyst
 - Now you know them very well!
- And you know how to contribute.
- Let's contribute to Spark & Spark SQL!!

an addition

What are we doing?

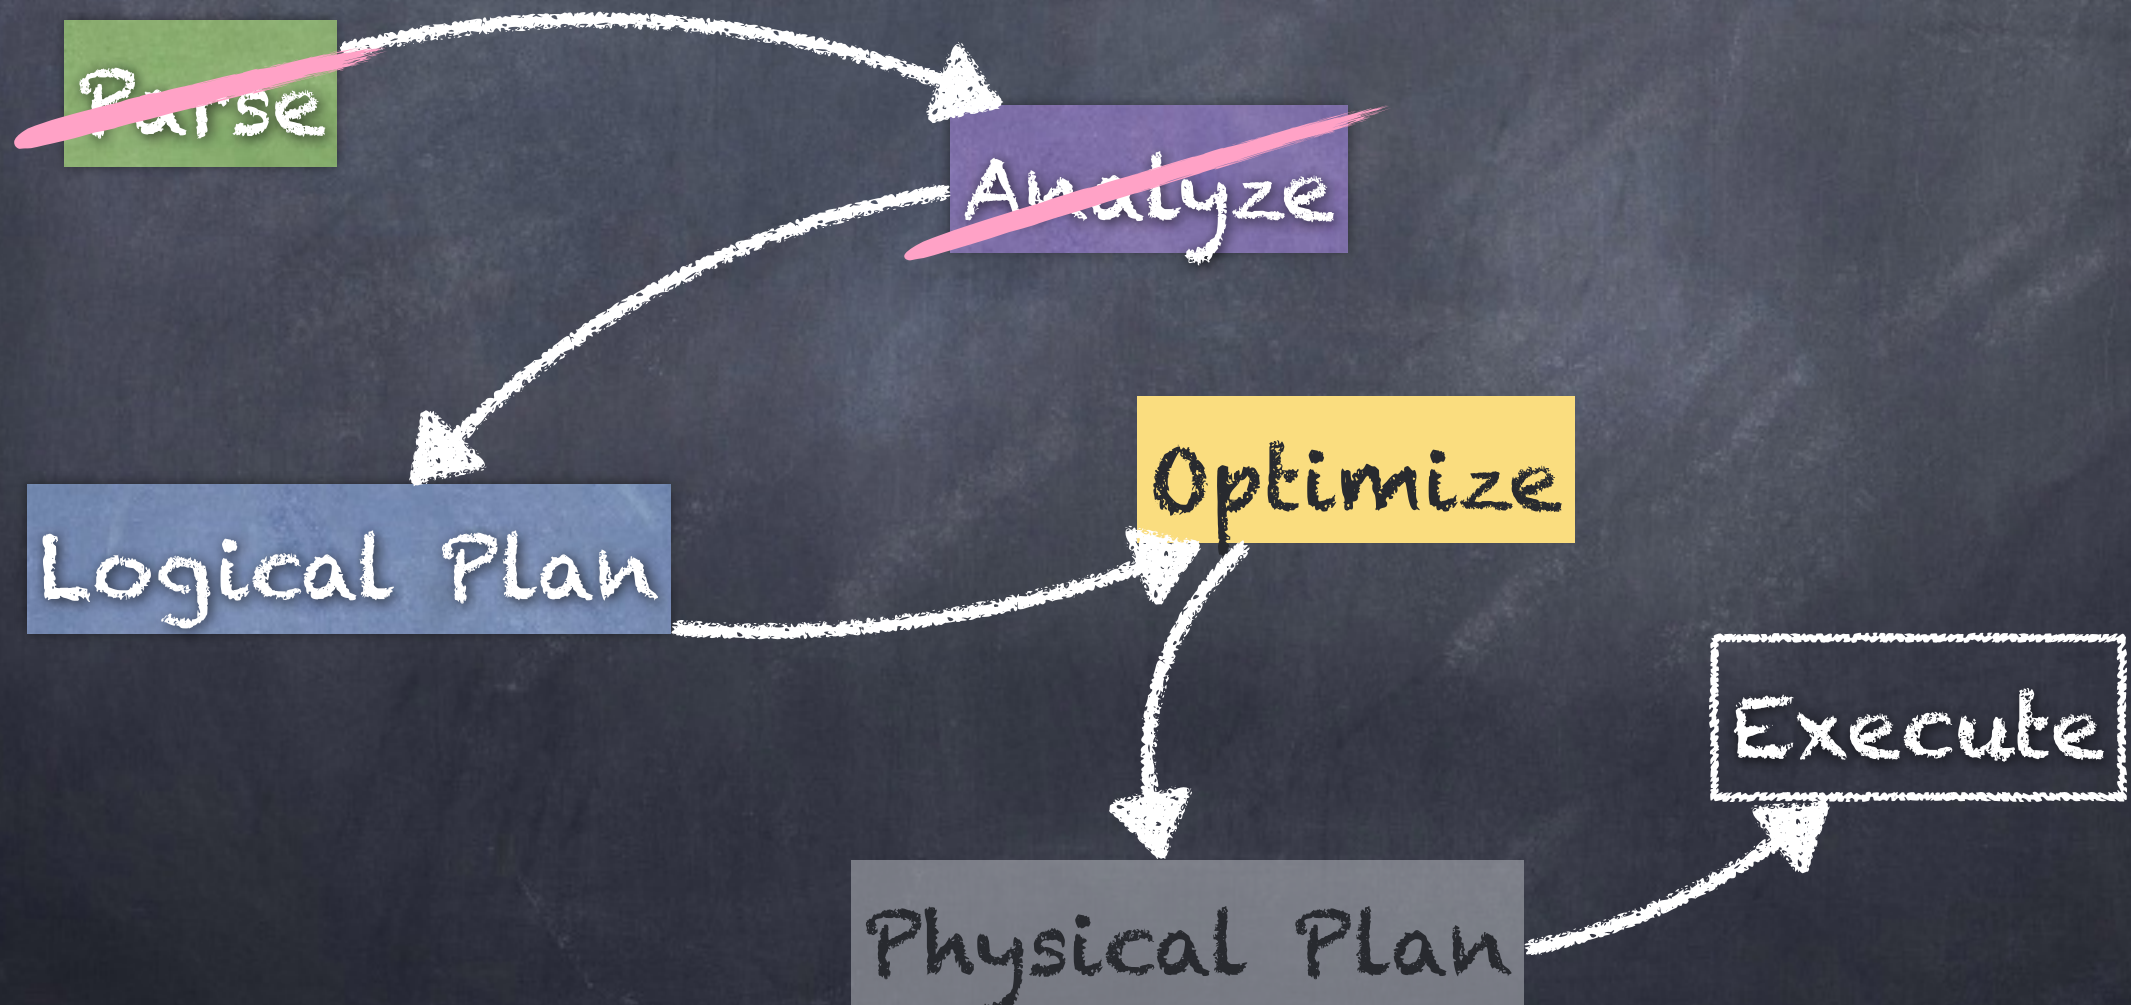
What are we doing?

- To execute query needs some steps.



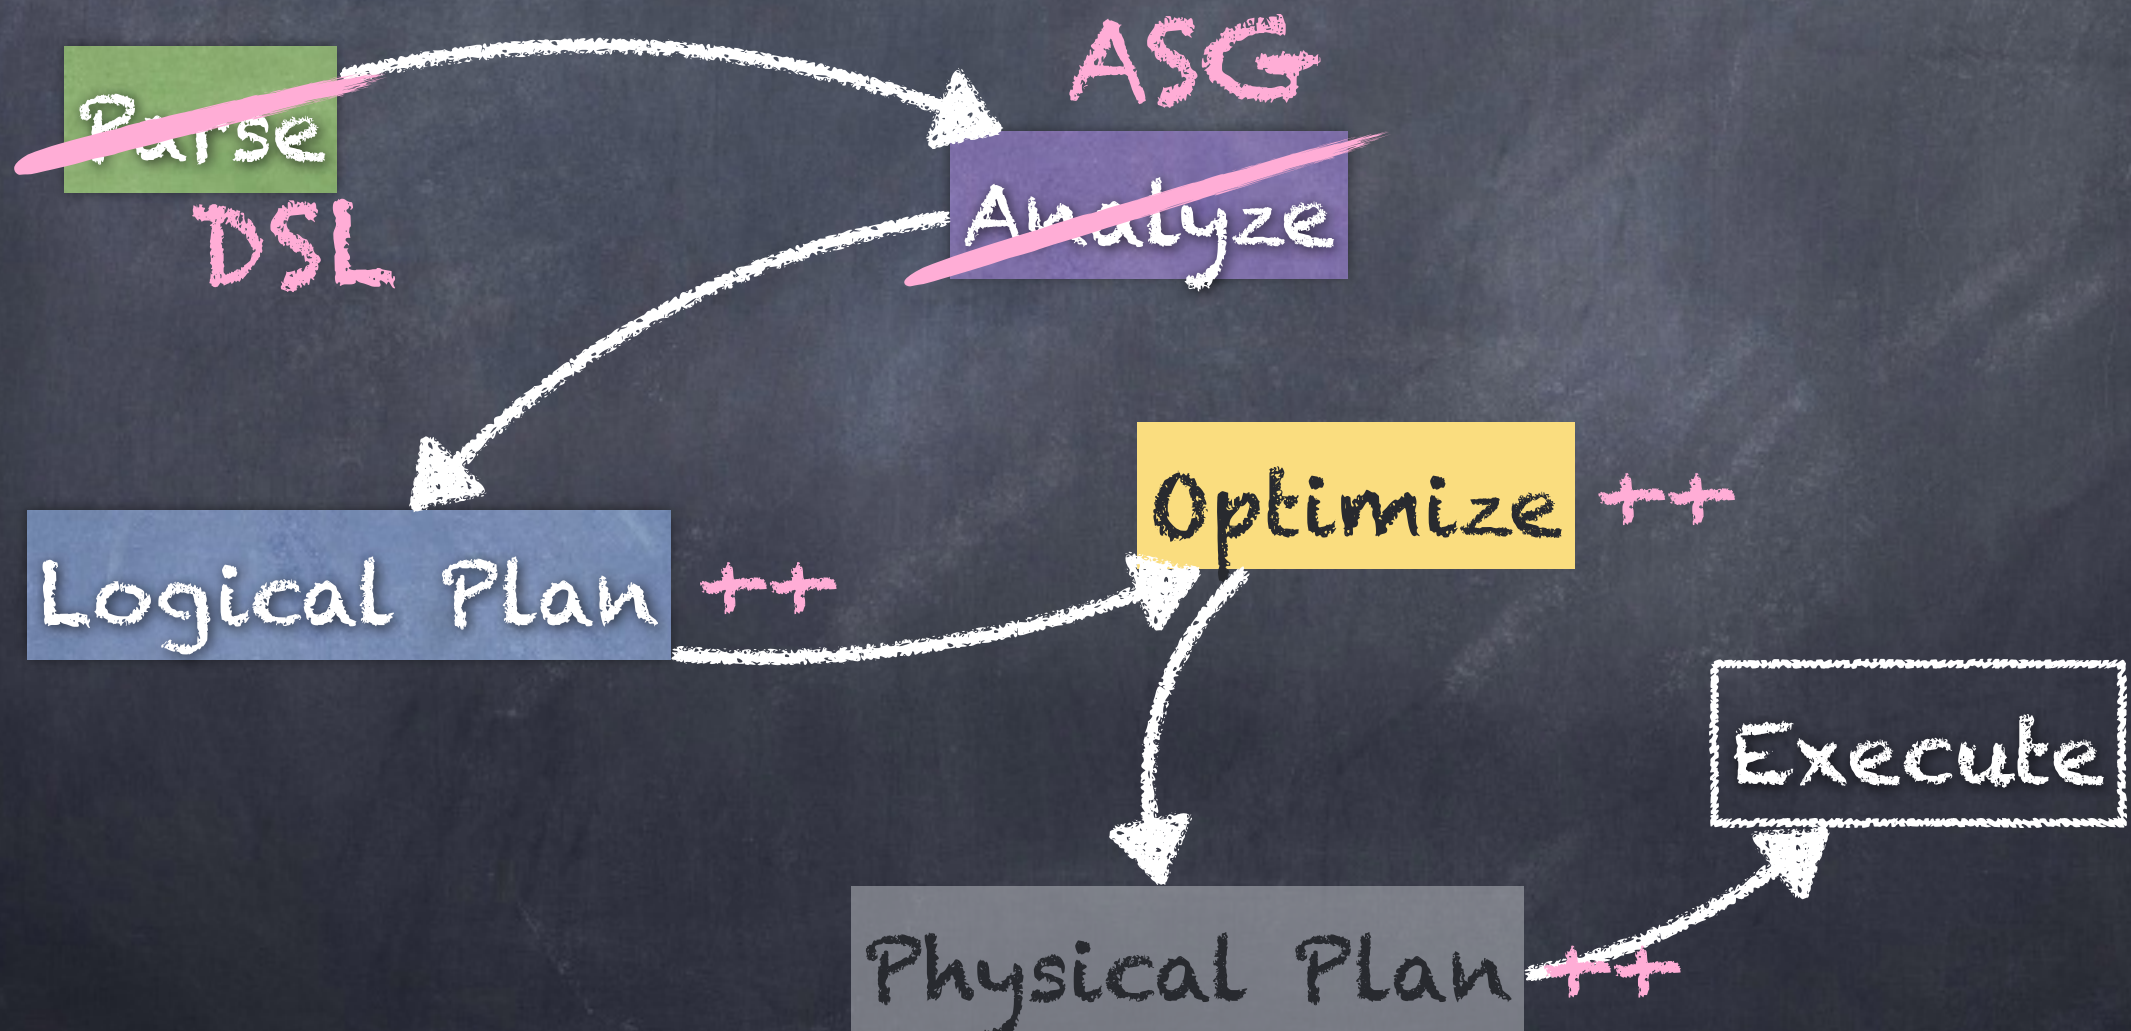
What are we doing?

- To execute ~~query~~ needs some steps.



What are we doing?

- business logic
• To execute ~~query~~ needs some steps.



Thanks!