Efficient In-situ Processing of Various Storage Types on Apache Tajo

Hadoop Summit 2015 San Jose

Hyunsik Choi, Gruter Inc.

Agenda

- Tajo Overview
- Various Storage Support
  - Motivation
  - Design Consideration
  - What we did/are doing

An overview of Apache Tajo

Tajo: A Data Warehouse System

- Data Warehouse System
- Apache Top-level project
- Low latency, and long running batch queries in a single system
  - ~100 ms up to several hours
  - Fault tolerance
- Features
  - ANSI SQL compliance
  - Mature SQL features: Joins, Group by, Sort, Multiple distinct aggregations and Window function
  - Partitioned table support
  - Java/Python UDF support
  - JDBC driver and Java-based asynchronous API
  - SQL data type and Nested type support
  - Direct JSON support

Tajo Overall Architecture

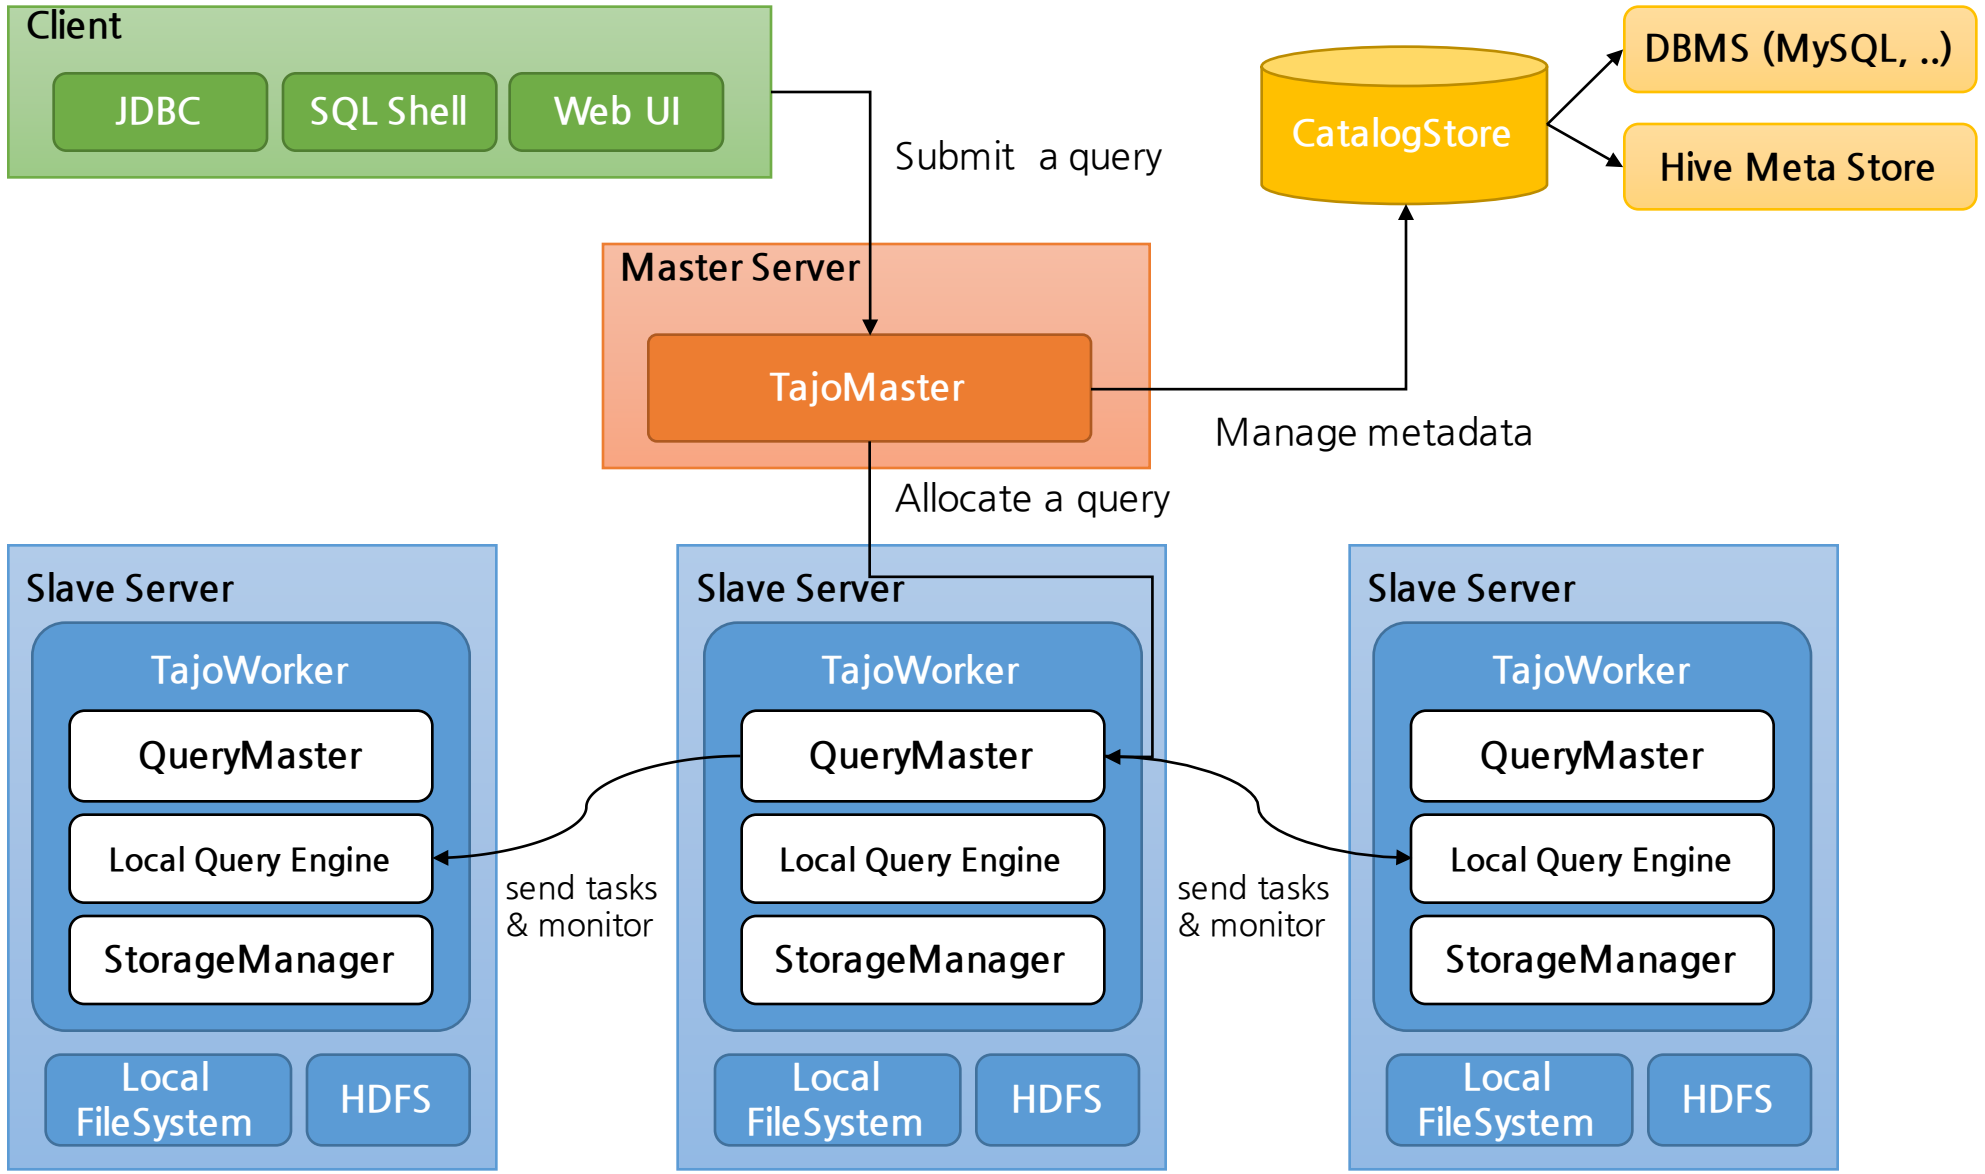

Background: Query Optimization Phases

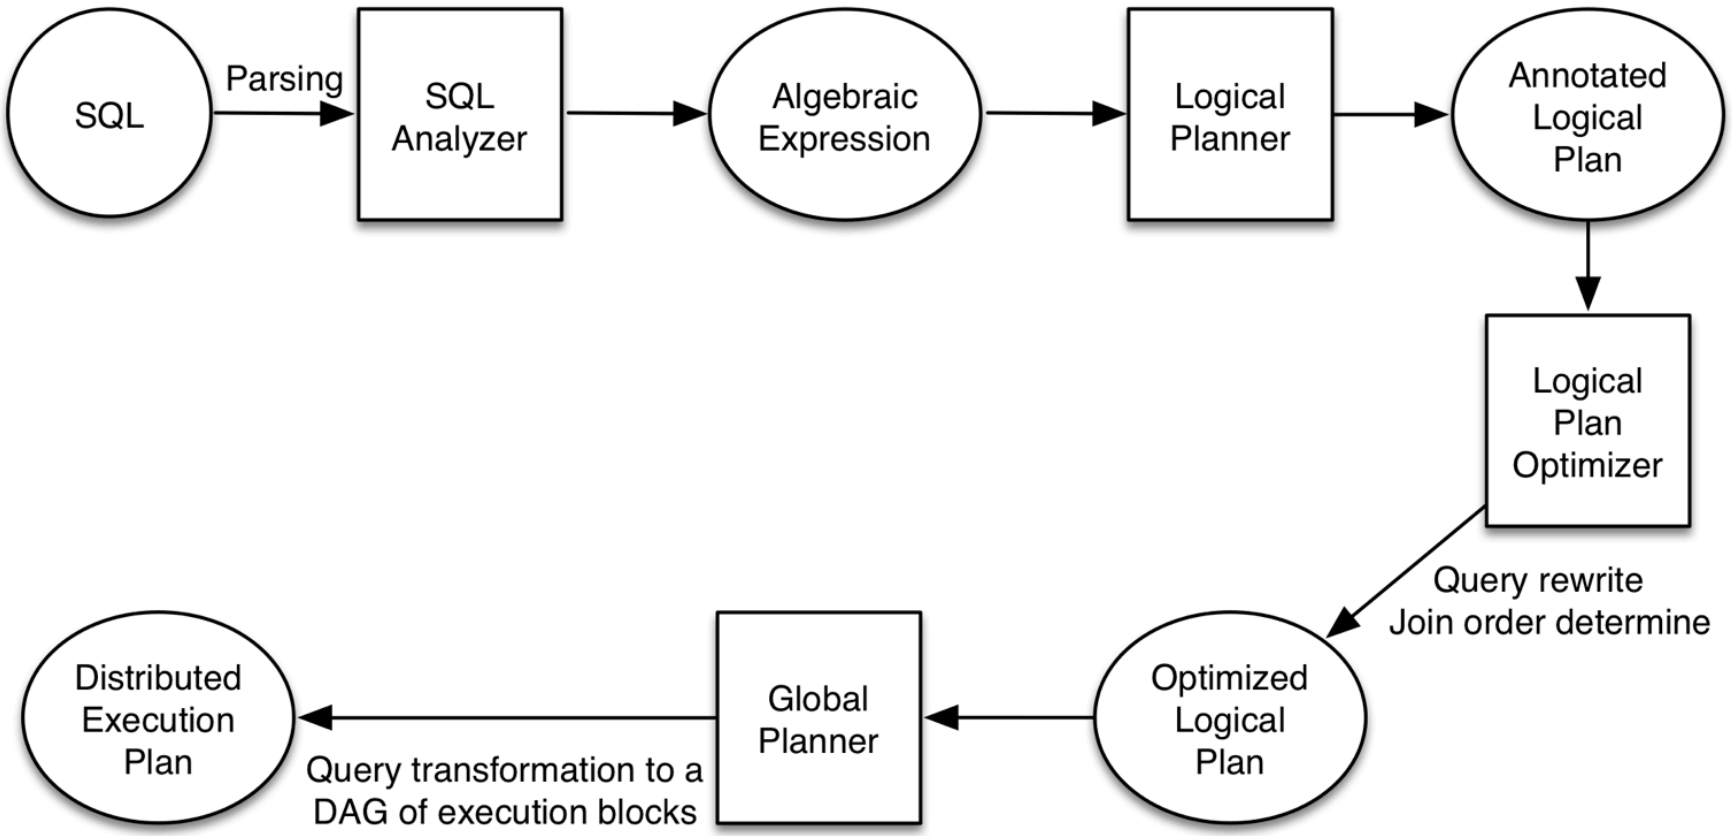

Background: Task Execution

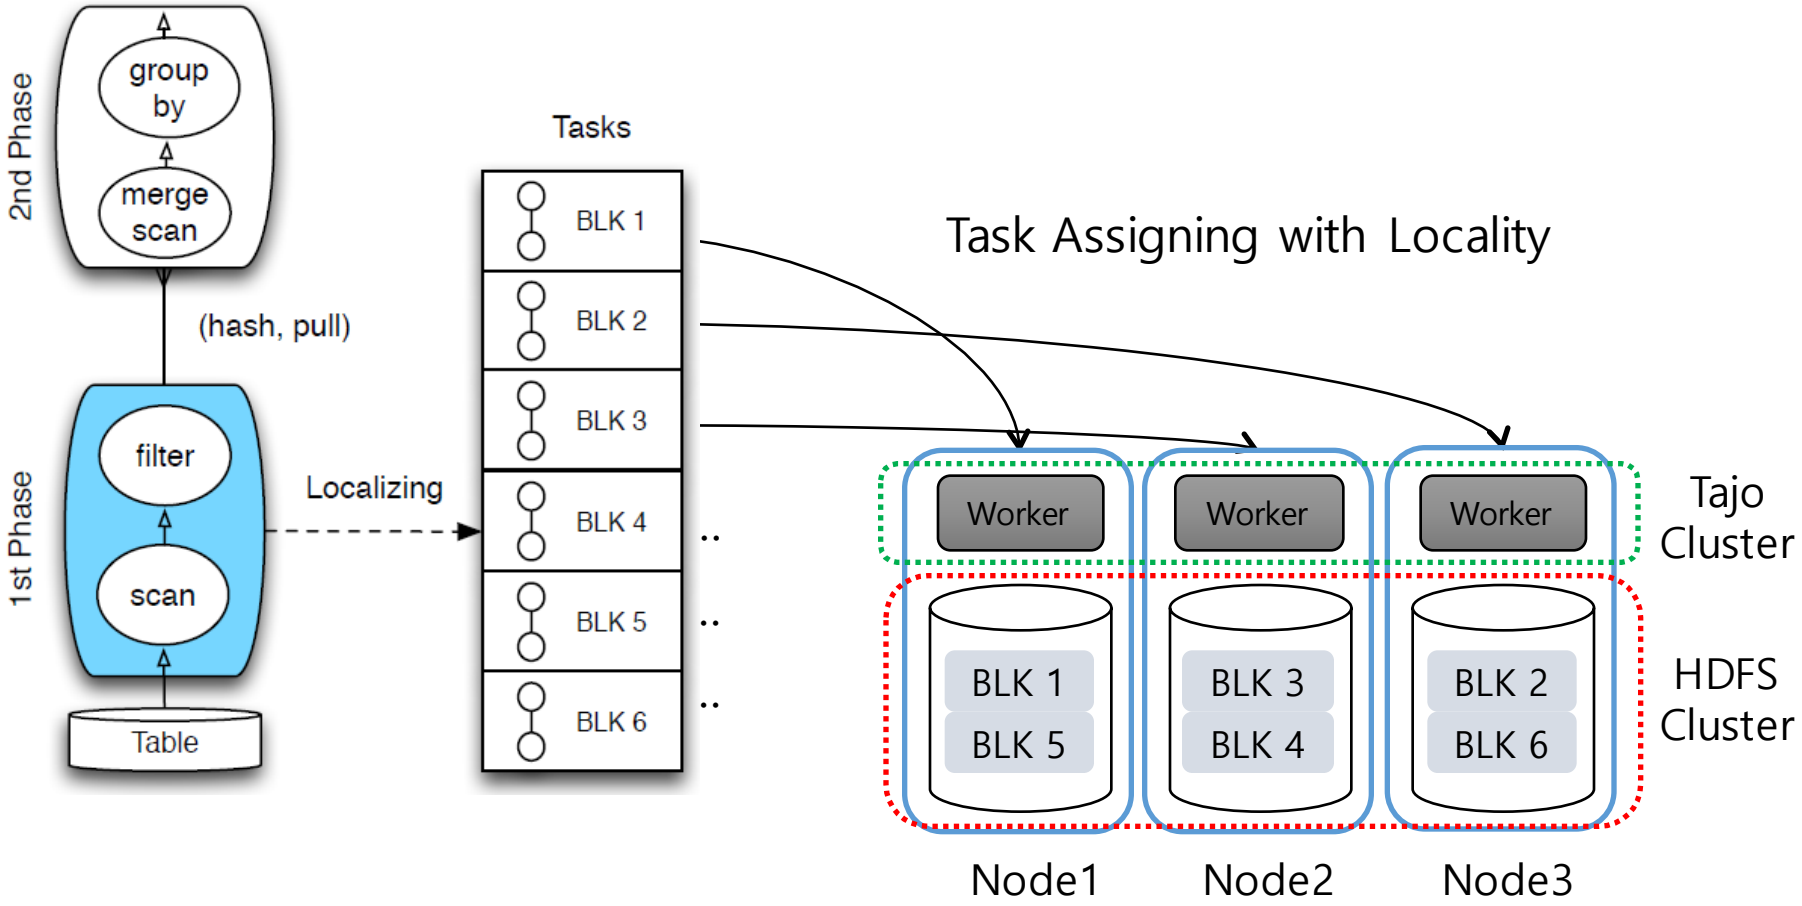

- Each task is assigned to a node according to its locality.

Background: Local Execution

- Physical operators are assembled into a tree and their execution pipelined in the same machine.
- Leaf operators must be scanners.
- Tajo provides abstraction scanner, allowing to read different physical tables.

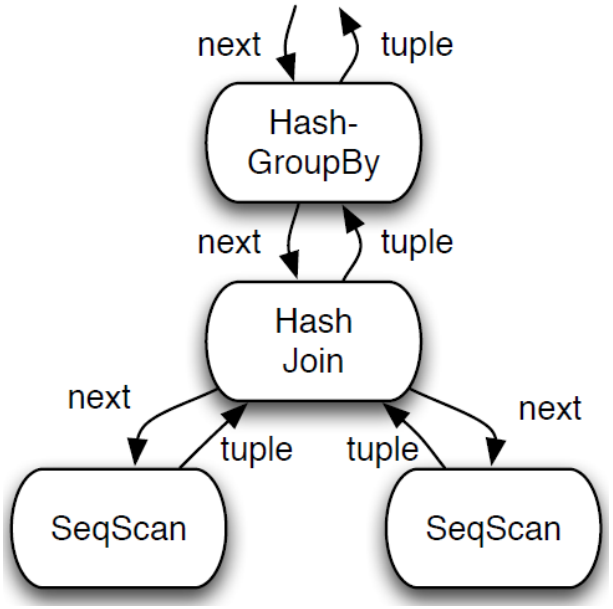

Various Storage Support

Motivation

- Unified Interface
- Data Integration
- In-situ Processing

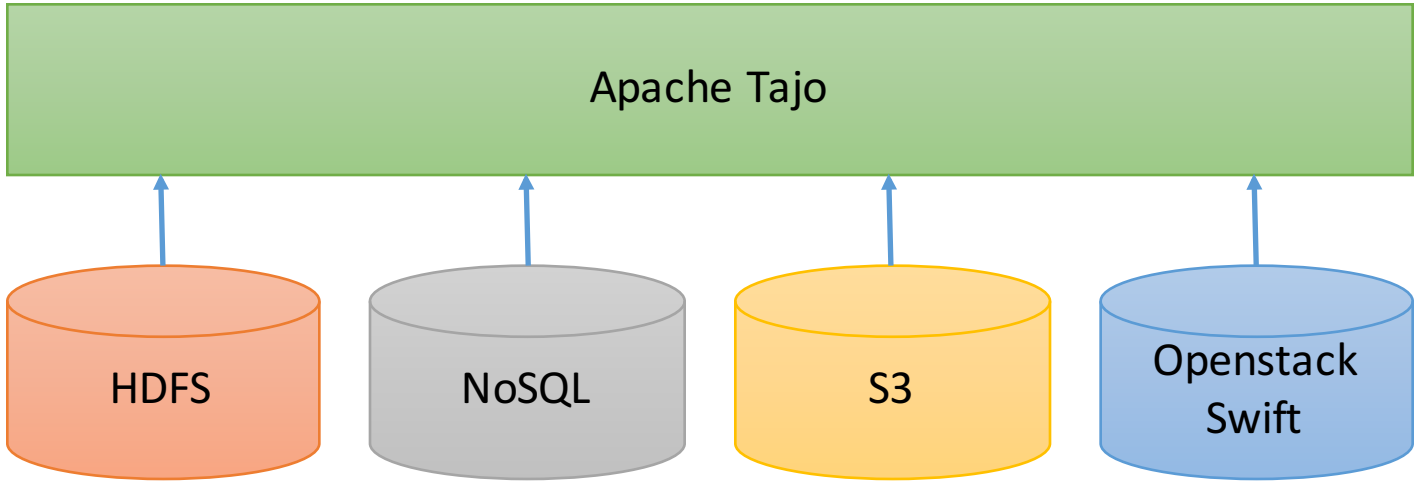

Datasets stored in Various Formats/Storages

Sequence File

Protocol Buffer

RCFile

Design Considerations

- More Storage Properties
  - Splittable, compressible (codecs), indexable, seekable, projectable, aggregatable, …
- Query Optimization
- Pluggable Storage and Data Format
- More operation pushdown

Separation between Storage and Format

Data Formats

Sequence File

RCFile

Protocol Buffer

Storage Types

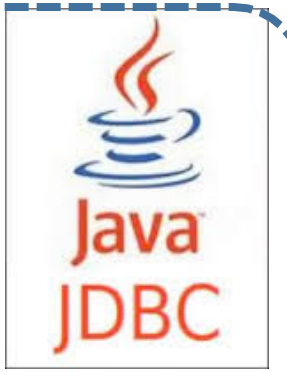

Relationships between Storage and Format

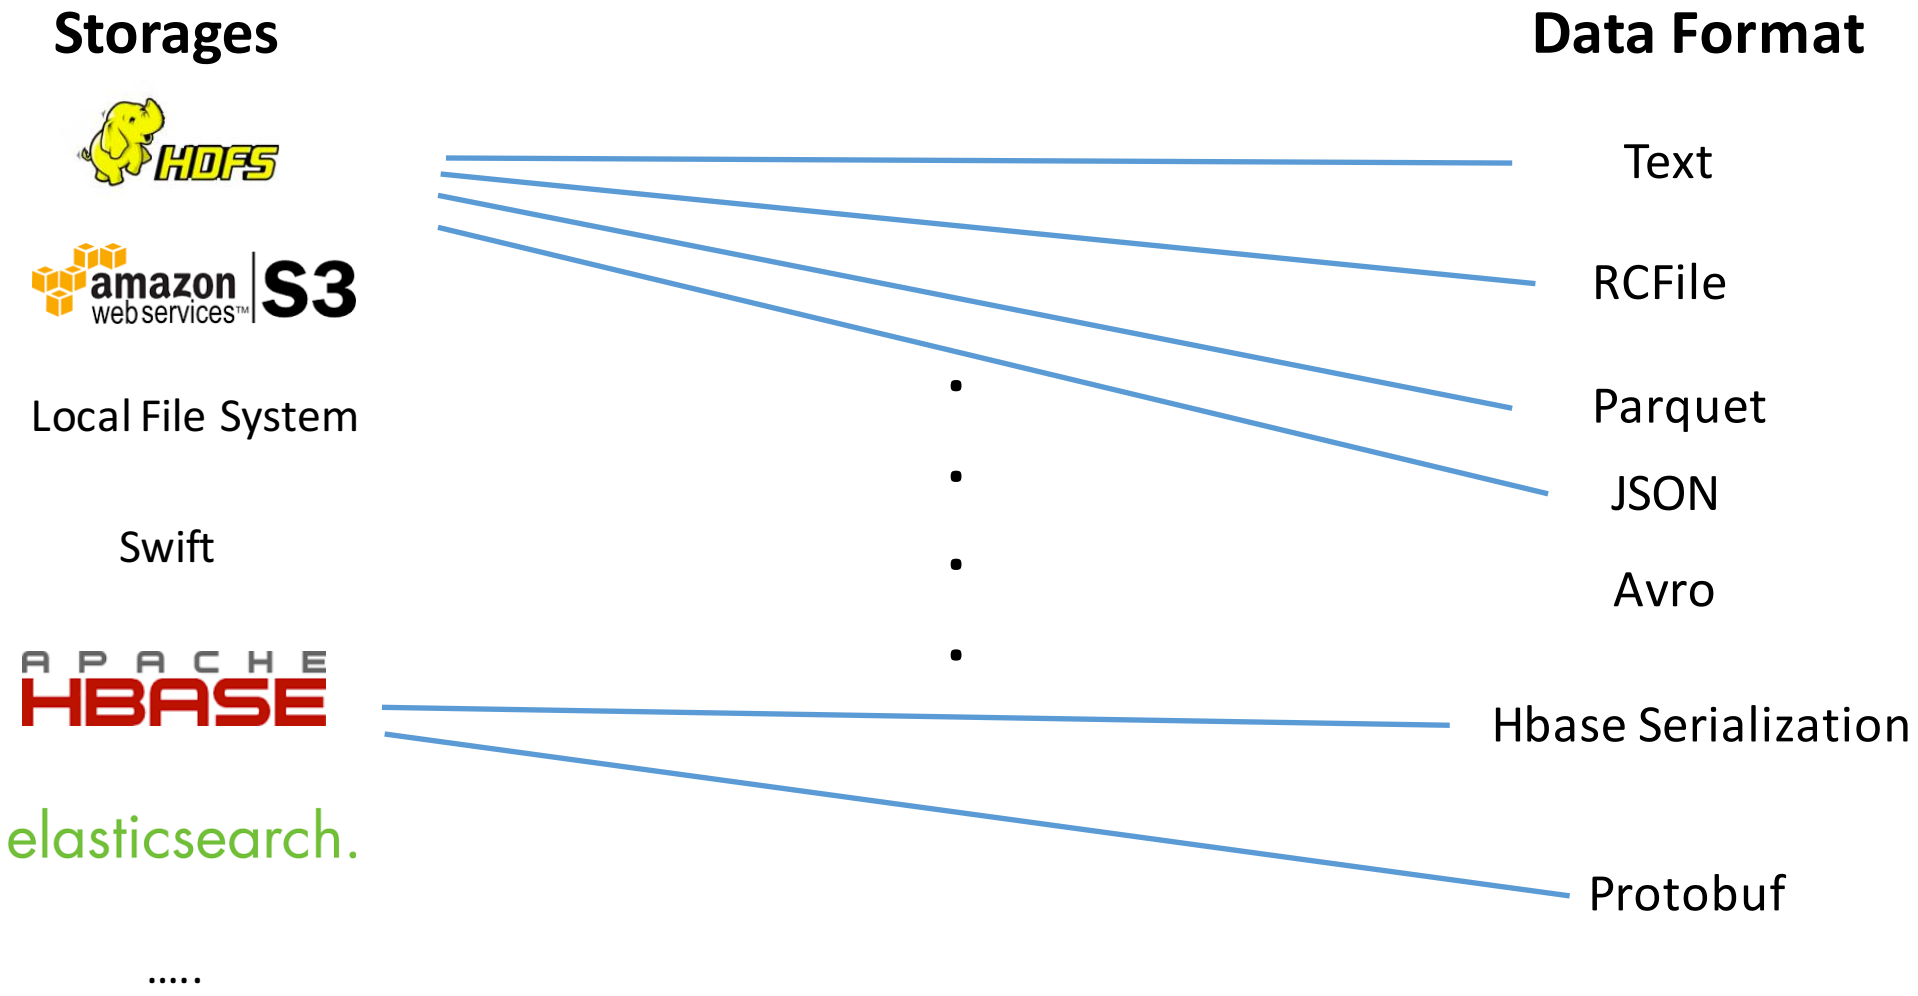

Tablespace

- Tablespace
  - Each table space is identified by a URI.
    - Hdfs://host:port/warehouse, hbase:zk://quorum1:2171, quorum2:2171, …
  - All tables in the same tablespace shares the same physical configuration.
  - URI scheme indicates storage type.
    - Hdfs, hbase, jdbc, …
  - Multiple tablespaces is possible in single storage namespace.
    - HDFS-2832: Enable support for heterogeneous in HDFS.
      - e.g.,
        - /warehouse/ (disk)
        - /today/ (ssd)

Storage Configuration

```
"storages": {
  "hdfs": {
    "handler": "org.apache.tajo.storage.HdfsTablespace",
    "default-format": "text"
  },
  "file": {
    "handler": "org.apache.tajo.storage.FileTablespace",
    "default-format": "text"
  },
  "hbase": {
    "handler": "org.apache.tajo.storage.hbase.HBaseTablespace",
    "default-format": "hbase"
  }
},
```

Storage Type Name and URI scheme

Storage Handler Class

Tablespace Configuration

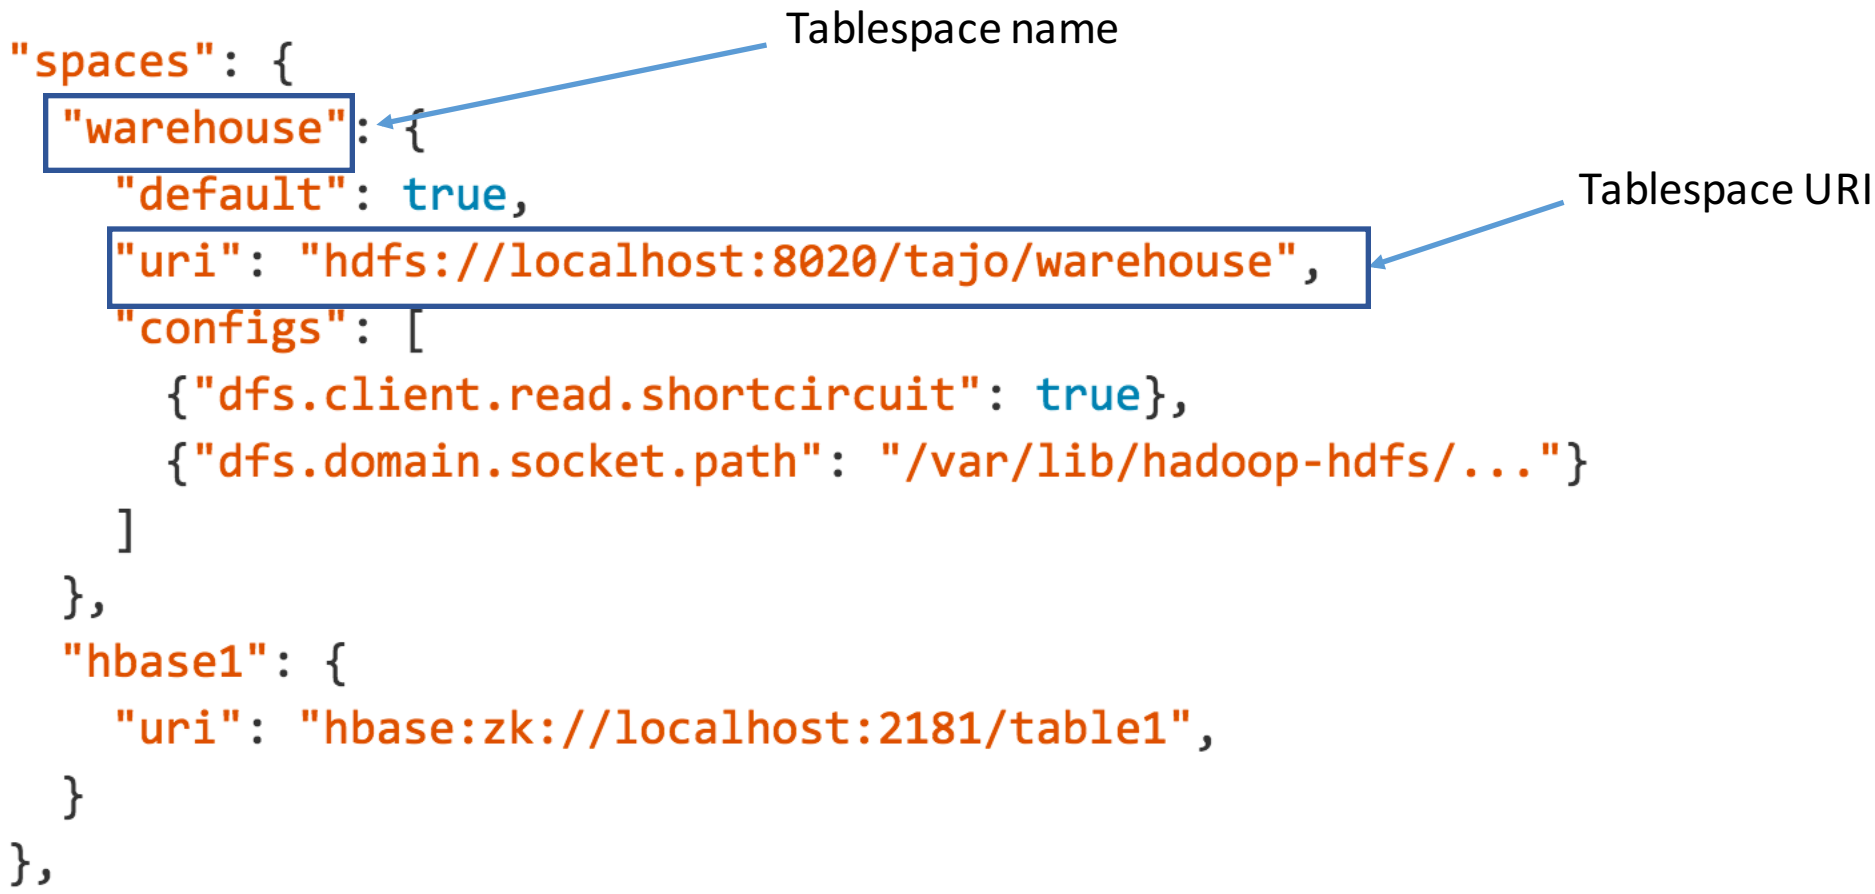

Format Configuration

```
"formats": {
  "avro": {
    "storage-support": ["hdfs", "file", "s3"],
    "handler": "org.apache.tajo.storage.AvroHandler"
  },
  "text": {
    "storage-support": ["hdfs", "file", "s3"],
    "handler": "org.apache.tajo.storage.TextHandler"
  },
  "hbase": {
    "storage-support": ["hbase"],
    "handler": "org.apache.tajo.storage.HbaseHandler"
  }
```

Format names

The relationship between formats and storages

CREATE Table using Tablespace

```
CREATE TABLE uptodate (key TEXT, …) TABLESPACE hbase1;
```

Tablespace Name

```
CREATE TABLE archive (l_orderkey bigint, …) TABLESPACE warehouse
USING text WITH ('text.delimiter' = '|');
```

Format name

Storage Layer Access over Query Lifecycle

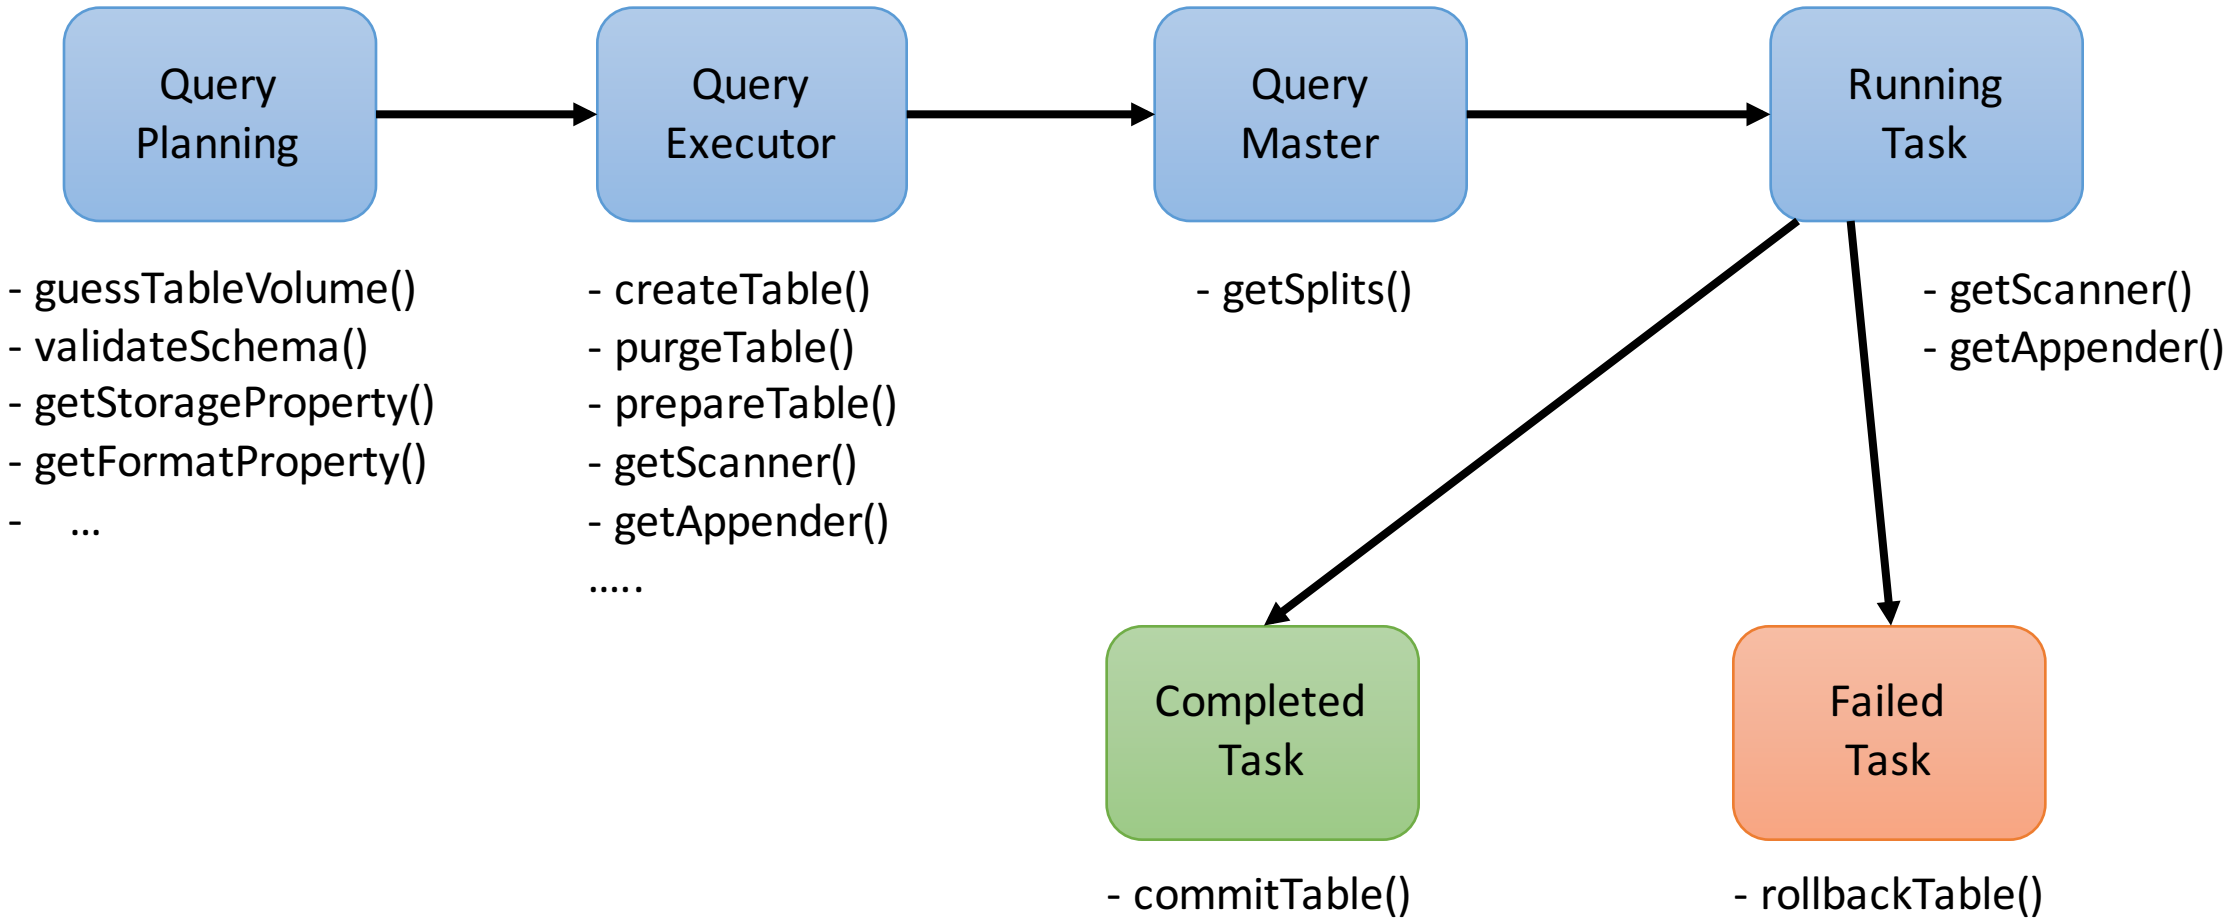

Query Rewrite for Specific Storages

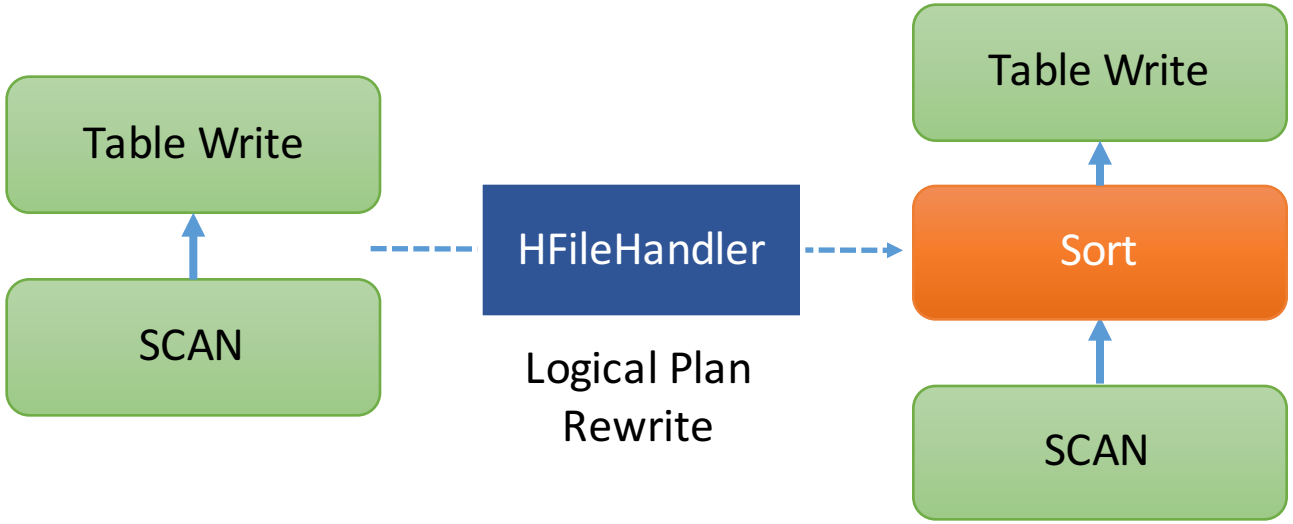


```
CREATE TABLE hbase_table (key TEXT, …)
INSERT INTO hbase_table SELECT id, name, …
```

Operation Push Down

```
SELECT
  X,
  SUM(Y)
FROM
  table1
WHERE
  x > 100
GROUP BY
  x
```

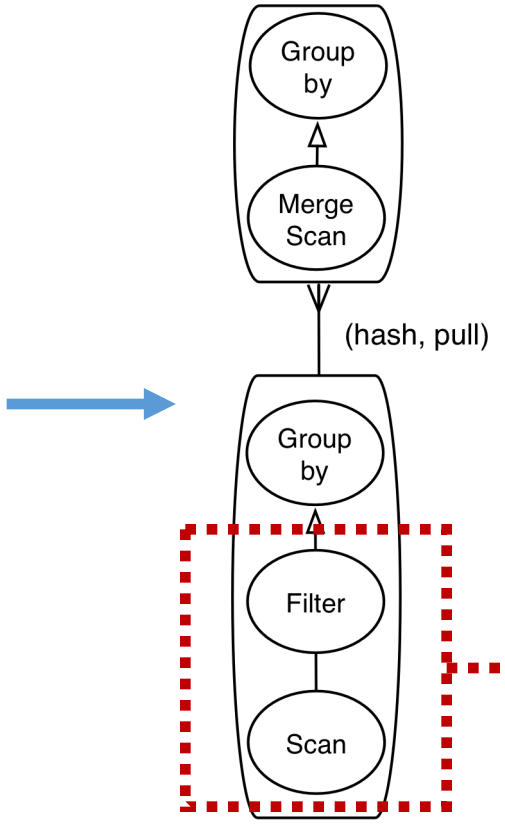

Filter and Projection can be pushed down into Underlying storages (like RDBMS, Hbase, Elasticsearch, …)

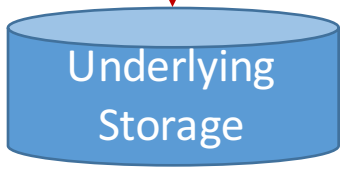

Current Status

- Storages:
  - HDFS support
  - Amazon S3 and Openstack Swift
  - Hbase Scanner and Writer - Hfile and Put Mode
  - JDBC-based Scanner and Writer (Working)
  - Kafka Scanner (Patch Available)
  - Elastic Search (Patch Available)
- Data Formats
  - Text, JSON, RCFile, SequenceFile, Avro, Parquet, and ORC (Patch Available)

Get Involved!

- We are recruiting contributors!

- General
  - http://tajo.apache.org

- Getting Started
  - http://tajo.apache.org/docs/0.10.0/getting_started.html

- Downloads
  - http://tajo.apache.org/downloads.html

- Jira – Issue Tracker
  - https://issues.apache.org/jira/browse/TAJO

- Join the mailing list
  - dev-subscribe@tajo.apache.org
  - issues-subscribe@tajo.apache.org

Q&A