

Beyond SQL: Speeding up Spark with DataFrames

Michael Armbrust - @michaelarmbrust

March 2015 – Spark Summit East





About Me and **SQL**

Spark SQL

- Part of the core distribution since Spark 1.0 (April 2014)

Graduated
from Alpha
in 1.3



About Me and **Spark SQL**

Spark SQL

- Part of the core distribution since Spark 1.0 (April 2014)
- Runs SQL / HiveQL queries, optionally alongside or replacing existing Hive deployments



```
SELECT COUNT(*)  
FROM hiveTable  
WHERE hive_udf(data)
```

About Me and **SQL**

Spark SQL

- Part of the core distribution since Spark 1.0 (April 2014)
- Runs SQL / HiveQL queries, optionally alongside or replacing existing Hive deployments
- Connect existing BI tools to Spark through JDBC



About Me and **SQL**

Spark SQL

- Part of the core distribution since Spark 1.0 (April 2014)
- Runs SQL / HiveQL queries, optionally alongside or replacing existing Hive deployments
- Connect existing BI tools to Spark through JDBC
- Bindings in Python, Scala, and Java



About Me and **Spark SQL**

Spark SQL

- Part of the core distribution since Spark 1.0 (April 2014)
- Runs SQL / HiveQL queries, optionally alongside or replacing existing Hive deployments
- Connect existing BI tools to Spark through JDBC
- Bindings in Python, Scala, and Java

[@michaelarmbrust](#) 


- Lead developer of Spark SQL [@databricks](#)

The not-so-secret truth...


Spark SQL
is not about SQL.

Execution Engine Performance

TPC-DS Performance



The not-so-secret truth...



is about more than SQL.

Spark SQL: The whole story

Creating and Running Spark Programs Faster:

- Write less code
- Read less data
- Let the optimizer do the hard work

DataFrame

noun – [dey-tuh-freym]

1. A distributed collection of rows organized into named columns.
2. An abstraction for selecting, filtering, aggregating and plotting structured data (*cf. R, Pandas*).
3. Archaic: Previously SchemaRDD (*cf. Spark < 1.3*).

Write Less Code: Input & Output

Spark SQL's Data Source API can read and write DataFrames using a variety of formats.

Built-In



External



Write Less Code: High-Level Operations

Common operations can be expressed concisely as calls to the DataFrame API:

- Selecting required columns
- Joining different data sources
- Aggregation (count, sum, average, etc)
- Filtering

Write Less Code: Compute an Average



```
private IntWritable one =
    new IntWritable(1)
private IntWritable output =
    new IntWritable()
protected void map(
    LongWritable key,
    Text value,
    Context context) {
    String[] fields = value.split("\t")
    output.set(Integer.parseInt(fields[1]))
    context.write(one, output)
}

IntWritable one = new IntWritable(1)
DoubleWritable average = new DoubleWritable()

protected void reduce(
    IntWritable key,
    Iterable<IntWritable> values,
    Context context) {
    int sum = 0
    int count = 0
    for(IntWritable value : values) {
        sum += value.get()
        count++
    }
    average.set(sum / (double) count)
    context.write(key, average)
}
```



```
data = sc.textFile(...).split("\t")
data.map(lambda x: (x[0], [x[1], 1])) \
    .reduceByKey(lambda x, y: [x[0] + y[0], x[1] + y[1]]) \
    .map(lambda x: [x[0], x[1][0] / x[1][1]]) \
    .collect()
```

Write Less Code: Compute an Average

Using RDDs

```
data = sc.textFile(...).split("\t")
data.map(lambda x: (x[0], [int(x[1]), 1])) \
    .reduceByKey(lambda x, y: [x[0] + y[0], x[1] + y[1]]) \
    .map(lambda x: [x[0], x[1][0] / x[1][1]]) \
    .collect()
```


Using DataFrames

```
sqlCtx.table("people") \
    .groupBy("name") \
    .agg("name", avg("age")) \
    .collect()
```

Full API Docs

- [Python](#)
- [Scala](#)
- [Java](#)

Not Just Less Code: Faster Implementations



Demo: Data Sources API

Using Spark SQL to read, write, and transform data in a variety of formats.

<http://people.apache.org/~marmbrus/talks/dataframe.demo.pdf>

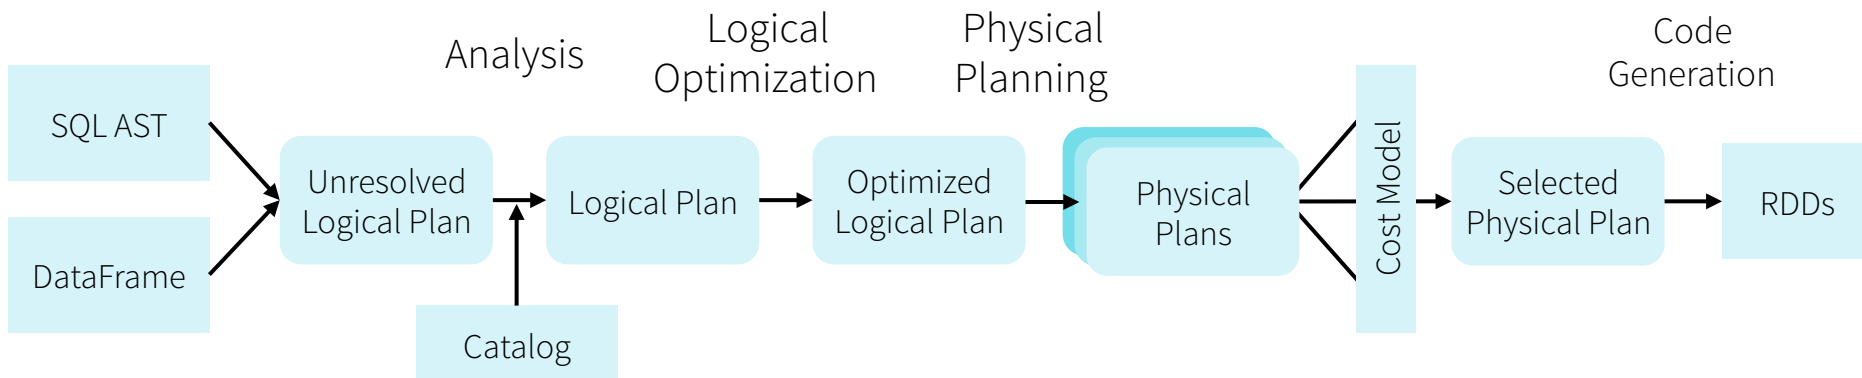
Read Less Data

The fastest way to process big data is to never read it.

Spark SQL can help you read less data automatically:

- Converting to more efficient formats
- Using columnar formats (i.e. parquet)
- Using partitioning (i.e., **/year=2014/month=02/...**)¹
- Skipping data using statistics (i.e., min, max)²
- Pushing predicates into storage systems (i.e., JDBC)

Plan Optimization & Execution




DataFrames and SQL share the same optimization/execution pipeline

Optimization happens as late as possible, therefore Spark SQL can optimize *even across functions*.


```
def add_demographics(events):
    u = sqlCtx.table("users")           # Load Hive table
    events \
        .join(u, events.user_id == u.user_id) \   # Join on user_id
        .withColumn("city", zipToCity(df.zip))     # Run udf to add city column

events = add_demographics(sqlCtx.load("/data/events", "json"))
training_data = events.where(events.city == "New York").select(events.timestamp).collect()
```

Logical Plan




Physical Plan




```
def add_demographics(events):
    u = sqlCtx.table("users")           # Load partitioned Hive table ←
    events \
        .join(u, events.user_id == u.user_id) \   # Join on user_id
        .withColumn("city", zipToCity(df.zip))      # Run udf to add city column

events = add_demographics(sqlCtx.load("/data/events", "parquet")) ←
training_data = events.where(events.city == "New York").select(events.timestamp).collect()
```


Logical Plan



Physical Plan

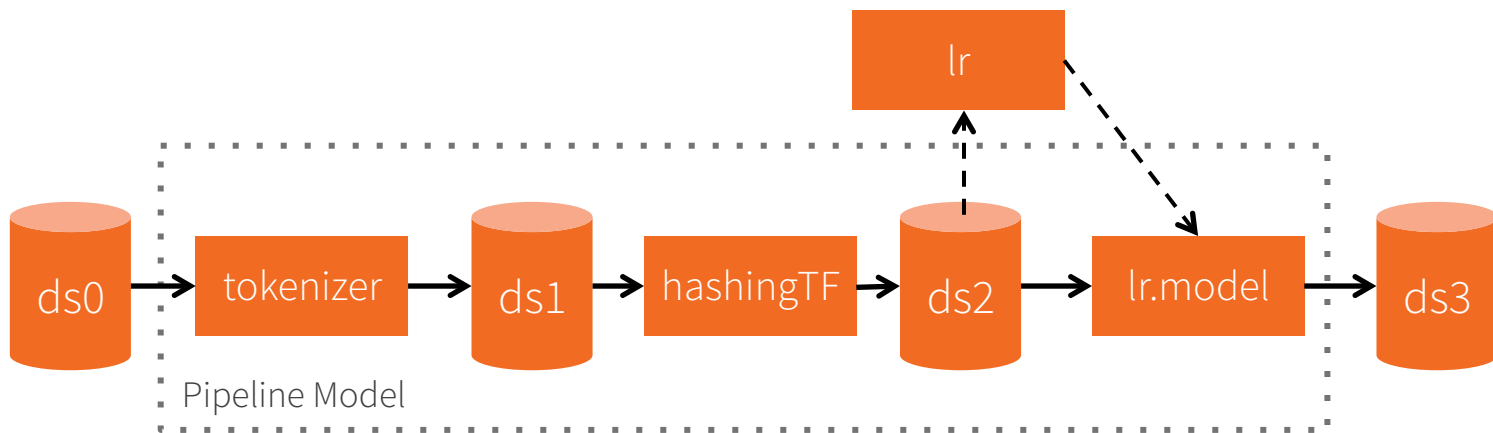


Physical Plan
with Predicate Pushdown
and Column Pruning



Machine Learning Pipelines

```
tokenizer = Tokenizer(inputCol="text", outputCol="words")  
hashingTF = HashingTF(inputCol="words", outputCol="features")  
lr = LogisticRegression(maxIter=10, regParam=0.01)  
pipeline = Pipeline(stages=[tokenizer, hashingTF, lr])  
  
df = sqlCtx.load("/path/to/data")  
model = pipeline.fit(df)
```





Create and Run Spark Programs Faster:

- Write less code
- Read less data
- Let the optimizer do the hard work



Questions?