


BlinkDB and G-OLA:

Supporting Approximate Answers in SparkSQL

Sameer Agarwal and Kai Zeng

Spark Summit | San Francisco, CA | June 15th 2015



About Us

1. Sameer Agarwal

- Software Engineer at Databricks
- PhD in Databases (UC Berkeley)
- Research on Approximate Query Processing (BlinkDB)

2. Kai Zeng

- Post-doc in AMP Lab/ Intern at Databricks
- PhD in Databases (UCLA)
- Research on Approximate Query Processing (ABM)

100 TB on 1000 machines

$\frac{1}{2}$ - 1 Hour



Hard Disks

1 - 5 Minutes



Memory

1 second



Continuous Query Execution on Samples of Data

Continuous Query Execution on Samples

ID	City	Latency
1	NYC	30
2	NYC	38
3	SLC	34
4	LA	36
5	SLC	37
6	SF	28
7	NYC	32
8	NYC	38
9	LA	36
10	SF	35
11	NYC	38
12	LA	34

What is the average latency in the table?

34.6667

Continuous Query Execution on Samples

ID	City	Latency
1	NYC	30
2	NYC	38
3	SLC	34
4	LA	36
5	SLC	37
6	SF	28
7	NYC	32
8	NYC	38
9	LA	36
10	SF	35
11	NYC	38
12	LA	34

What is the average latency in the table?

35

Continuous Query Execution on Samples

ID	City	Latency
1	NYC	30
2	NYC	38
3	SLC	34
4	LA	36
5	SLC	37
6	SF	28
7	NYC	32
8	NYC	38
9	LA	36
10	SF	35
11	NYC	38
12	LA	34

What is the average latency in the table?

35 ± 2.1

Continuous Query Execution on Samples

ID	City	Latency
1	NYC	30
2	NYC	38
3	SLC	34
4	LA	36
5	SLC	37
6	SF	28
7	NYC	32
8	NYC	38
9	LA	36
10	SF	35
11	NYC	38
12	LA	34

What is the average latency in the table?

35 ± 2.1
33.83 ± 1.3

Continuous Query Execution on Samples

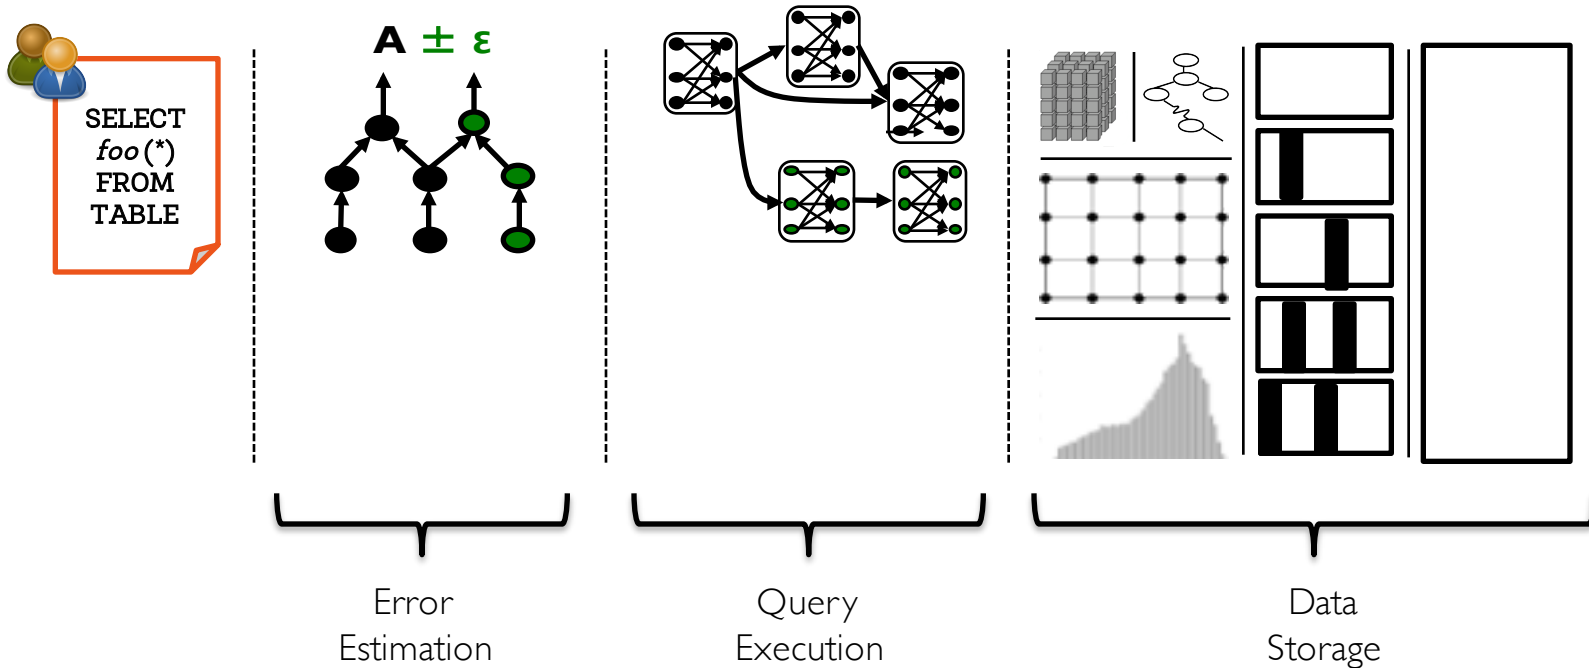
ID	City	Latency
1	NYC	30
2	NYC	38
3	SLC	34
4	LA	36
5	SLC	37
6	SF	28
7	NYC	32
8	NYC	38
9	LA	36
10	SF	35
11	NYC	38
12	LA	34

What is the average latency in the table?

35 ± 2.1
33.83 ± 1.3
34.6667 ± 0.0

Demo


Continuous Query Execution on Samples



Continuous Query Execution on Samples




SELECT
foo(*)
FROM
TABLE



$A \pm \epsilon$





BlinkDB

Error
Estimation



G-OLA

Query
Execution



Data
Storage

Interface

```
val dataframe =  
    sqlCtx.sql("select avg(latency) from log")  
  
// batch processing  
val result = dataframe.collect() // 34.6667
```

Interface

```
val dataframe =  
    sqlCtx.sql("select avg(latency) from log")  
  
// online processing  
val onlineDataFrame = dataframe.online  
onlineDataFrame.collectNext() // 35 ± 2.1  
onlineDataFrame.collectNext() // 33.83 ± 1.3
```

Interface

```
val dataframe =  
    sqlCtx.sql("select avg(latency) from log")  
  
// online processing  
val onlineDataFrame = dataframe.online  
while (onlineDataFrame.hasNext()) {  
    onlineDataFrame.collectNext()  
}
```

Interface

```
val dataframe =  
    sqlCtx.sql("select avg(latency) from log")  
  
// online processing  
val onlineDataFrame = dataframe.online  
while (onlineDataFrame.hasNext() &&  
    responseTime <= 10.seconds) {  
    onlineDataFrame.collectNext()  
}
```

Interface

```
val dataframe =  
    sqlCtx.sql("select avg(latency) from log")  
  
// online processing  
val onlineDataFrame = dataframe.online  
while (onlineDataFrame.hasNext() &&  
    errorBound >= 0.01) {  
    onlineDataFrame.collectNext()  
}
```


Interface

```
val dataframe =  
    sqlCtx.sql("select avg(latency) from log")  
  
// online processing  
val onlineDataFrame = dataframe.online  
while (onlineDataFrame.hasNext() &&  
    userEvent.cancelled()) {  
    onlineDataFrame.collectNext()  
}
```

Interface


```
val dataframe =  
    sqlCtx.sql("select avg(latency) from log")
```

```
// online processing
```


```
val onlineDataFrame = dataframe.online  
while (onlineDataFrame.hasNext() &&  
    userEvent.cancelled()) {  
    onlineDataFrame.collectNext()  
}
```

AGGREGATES/ UDAFs
JOINS/GROUP BYs
NESTED QUERIES

Continuous Query Execution on Samples



Continuous Query Execution on Samples



Sameer Agarwal, Barzan Mozafari, Aurojit Panda, Henry Milner, Samuel Madden, Ion Stoica. **BlinkDB: Queries with Bounded Errors and Bounded Response Times on Very Large Data.** In ACM EuroSys 2013.

Ariel Kleiner, Ameet Talwalkar, Sameer Agarwal, Ion Stoica, Michael Jordan. **A General Bootstrap Performance Diagnostic.** In ACM KDD 2013

Sameer Agarwal, Henry Milner, Ariel Kleiner, Ameet Talwalkar, Michael Jordan, Samuel Madden, Barzan Mozafari, Ion Stoica. **Knowing When You're Wrong: Building Fast and Reliable Approximate Query Processing Systems.** In ACM SIGMOD 2014.

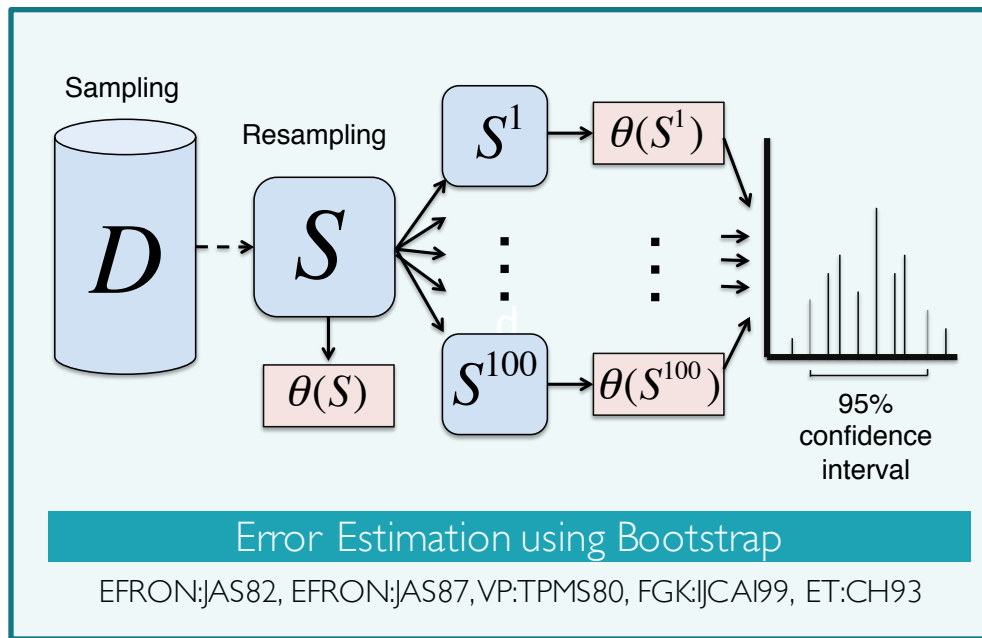
Error Estimation on a Sample of Data

Focused on estimating aggregate errors given *representative* samples

1. Count: $N(np, n(1 - p)p)$
2. Sum: $N(np\mu, np(\sigma^2 + (1 - p)\mu^2))$
3. Mean: $N(\mu, \sigma^2/n)$
4. Variance: $N(\sigma^2, (\mu_4 - \sigma^4)/n)$
5. Stddev: $N(\sigma, (\mu_4 - \sigma^4)/(4\sigma^2n))$

Central Limit Theorem (CLT)

HOE:ASTAT63,BIL: WILEY86, CGL:ASTAT83,PH:IBM96



Error Estimation using Bootstrap

ID	City	Latency
1	NYC	30
2	NYC	38
3	SLC	34
4	LA	36
5	SLC	37
6	SF	28
7	NYC	32
8	NYC	38
9	LA	36
10	SF	35
11	NYC	38
12	LA	34

What is the average latency in the table?

Error Estimation using Bootstrap

ID	City	Latency
1	NYC	30
2	NYC	38
3	SLC	34
4	LA	36

ID	City	Latency
1	NYC	30
2	NYC	38
3	SLC	34
4	SLC	34

ID	City	Latency
1	NYC	30
2	NYC	30
3	SLC	34
4	LA	36

...

ID	City	Latency
1	SLC	34
2	LA	36
3	SLC	34
4	LA	36

$$\theta_1 = 34$$

$$\theta_2 = 32.5$$

...

$$\theta_{100} = 35$$

$$34.5 \pm 2$$

Error Estimation using Bootstrap

ID	City	Latency
1	NYC	30
2	NYC	38
3	SLC	34
4	LA	36
5	SLC	37
6	SF	28
7	NYC	32
8	NYC	38
9	LA	36
10	SF	35
11	NYC	38
12	LA	34

ID	City	Latency
1	NYC	30
2	NYC	38
3	SLC	34
4	SLC	34

ID	City	Latency
1	NYC	30
2	NYC	30
3	SLC	34
4	LA	36

ID	City	Latency
1	SLC	34
2	LA	36
3	SLC	34
4	LA	36

$$\theta_1 = 34$$

$$\theta_2 = 32.5$$

...

$$\theta_{100} = 35$$

$$34.5 \pm 2$$

Error Estimation using Bootstrap

ID	City	Latency
1	NYC	30
2	NYC	38
3	SLC	34
4	LA	36
5	SLC	37
6	SF	28
7	NYC	32
8	NYC	38
9	LA	36
10	SF	35
11	NYC	38
12	LA	34

ID	City	Latency
1	NYC	30
2	NYC	38
3	SLC	34
4	SLC	34
5	SLC	37

$$\theta_1 = 34.6$$

ID	City	Latency
1	SLC	37
2	NYC	30
3	SLC	34
4	LA	36
5	NYC	30

$$\theta_2 = 33.4$$

ID	City	Latency
1	SLC	34
2	SLC	37
3	SLC	34
4	LA	36
5	LA	36

...

$$\theta_{100} = 35.4$$

$$35 \pm 1.6$$

Error Estimation in BlinkDB

ID	City	Latency
1	NYC	30
2	NYC	38
3	SLC	34
4	LA	36
5	SLC	37

6	SF	28
7	NYC	32
8	NYC	38
9	LA	36
10	SF	35
11	NYC	38
12	LA	34

What is the average latency in the table?

ID	City	Latency
1	NYC	30
2	NYC	38
3	SLC	34
4	SLC	34
5	SLC	37

Leverage Poissonized Resampling to generate samples with replacement

Error Estimation in BlinkDB

ID	City	Latency
1	NYC	30
2	NYC	38
3	SLC	34
4	LA	36
5	SLC	37
6	SF	28
7	NYC	32
8	NYC	38
9	LA	36
10	SF	35
11	NYC	38
12	LA	34

What is the average latency in the table?

ID	City	Latency	# ₁
1	NYC	30	2
2	NYC	38	1
3	SLC	34	0
4	SLC	34	1
5	SLC	37	1

Sample from a
Poisson (1) Distribution

$$\theta_1 = 33.8$$

Error Estimation in BlinkDB

ID	City	Latency
1	NYC	30
2	NYC	38
3	SLC	34
4	LA	36
5	SLC	37
6	SF	28
7	NYC	32
8	NYC	38
9	LA	36
10	SF	35
11	NYC	38
12	LA	34

What is the average latency in the table?

ID	City	Latency	# ₁
1	NYC	30	2
2	NYC	38	1
3	SLC	34	0
4	SLC	34	1
5	SLC	37	1
6	SF	28	2

Incremental Error Estimation

Error Estimation in BlinkDB

ID	City	Latency
1	NYC	30
2	NYC	38
3	SLC	34
4	LA	36
5	SLC	37
6	SF	28
7	NYC	32
8	NYC	38
9	LA	36
10	SF	35
11	NYC	38
12	LA	34

What is the average latency in the table?

ID	City	Latency	# ₁	# ₂
1	NYC	30	2	1
2	NYC	38	1	0
3	SLC	34	0	2
4	SLC	34	1	2
5	SLC	37	1	0
6	SF	28	2	1

Construct all Resamples in a Single Pass

0.2-0.5% additional overhead

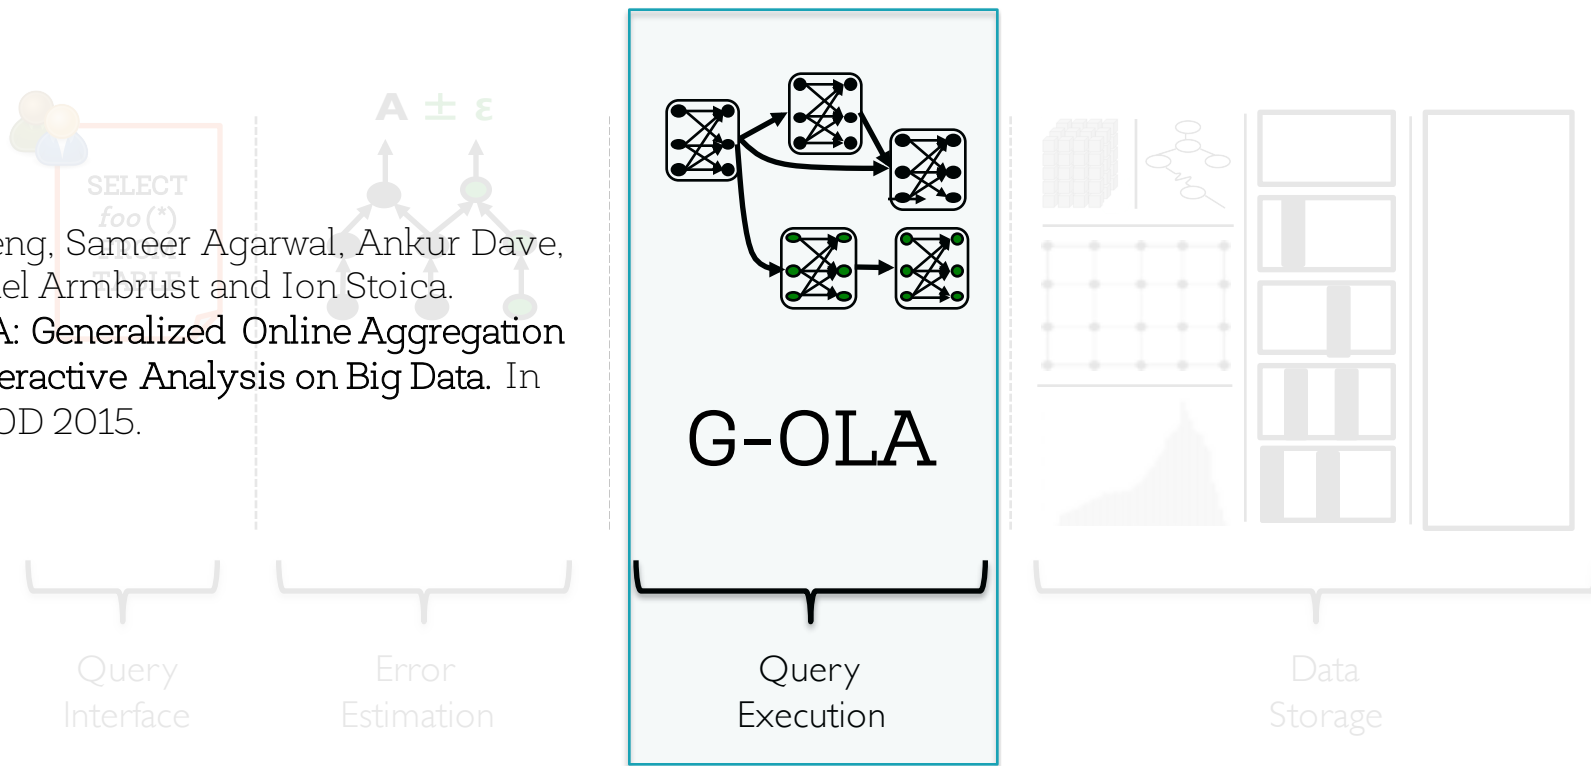
High Level Take-away:

Bootstrap and Poissonized Resampling Techniques are the key towards achieving quick and continuous error bars for a general set of queries


Continuous Query Execution on Samples

Kai Zeng, Sameer Agarwal, Ankur Dave,
Michael Armbrust and Ion Stoica.

G-OLA: Generalized Online Aggregation
for Interactive Analysis on Big Data. In
SIGMOD 2015.



Query Execution: Under The Hood



Query Execution: Under The Hood


Answer $\pm \epsilon$

A


A'

Query

Data






10 sec



10+10 sec

Query Execution: Under The Hood




Answer $\pm \epsilon$



Query Execution: Under The Hood



Overall Quadratic Cost!

Answer $\pm \epsilon$



Query Execution: Under The Hood

Answer $\pm \epsilon$




Query Execution: Under The Hood

Answer $\pm \epsilon$

A

Query




Data



10 sec


A'




10+10 sec

Query Execution: Under The Hood


Answer $\pm \epsilon$



Query Execution: Under The Hood



Query Execution: Under The Hood




Delta Update Queries

Answer $\pm \epsilon$

Query

Data




Delta Update Queries

Answer $\pm \epsilon$

Query

Data




Delta Update Queries

Answer $\pm \epsilon$

Query

Data




Delta Update Queries

Answer $\pm \epsilon$

Query

Data





Delta Update Queries

Answer $\pm \epsilon$

Query

Data





Delta Update Queries

Answer $\pm \epsilon$


Query

Data



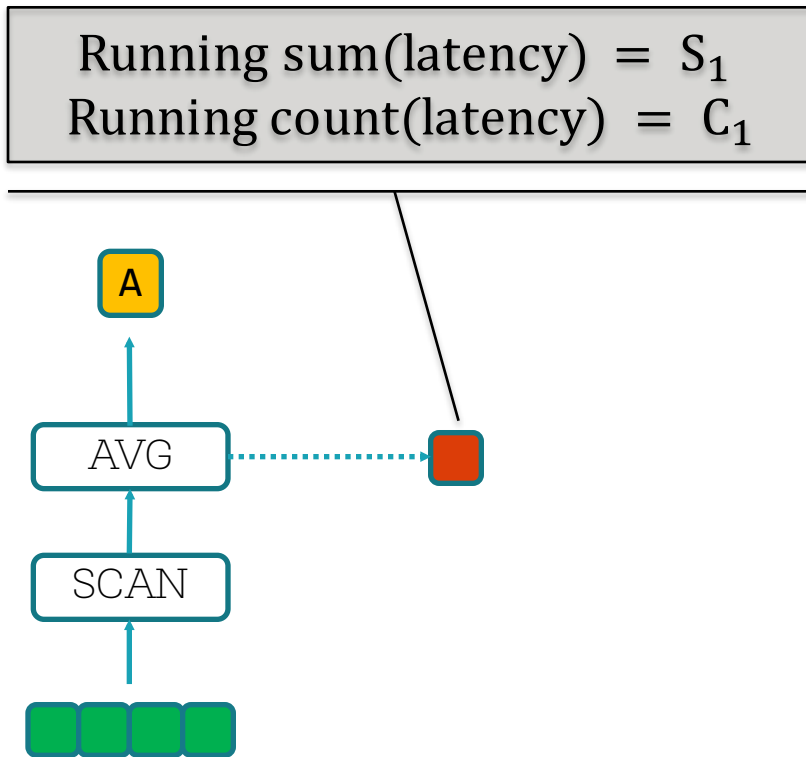
Delta Update: Simple Queries

```
SELECT avg(latency)  
FROM log
```



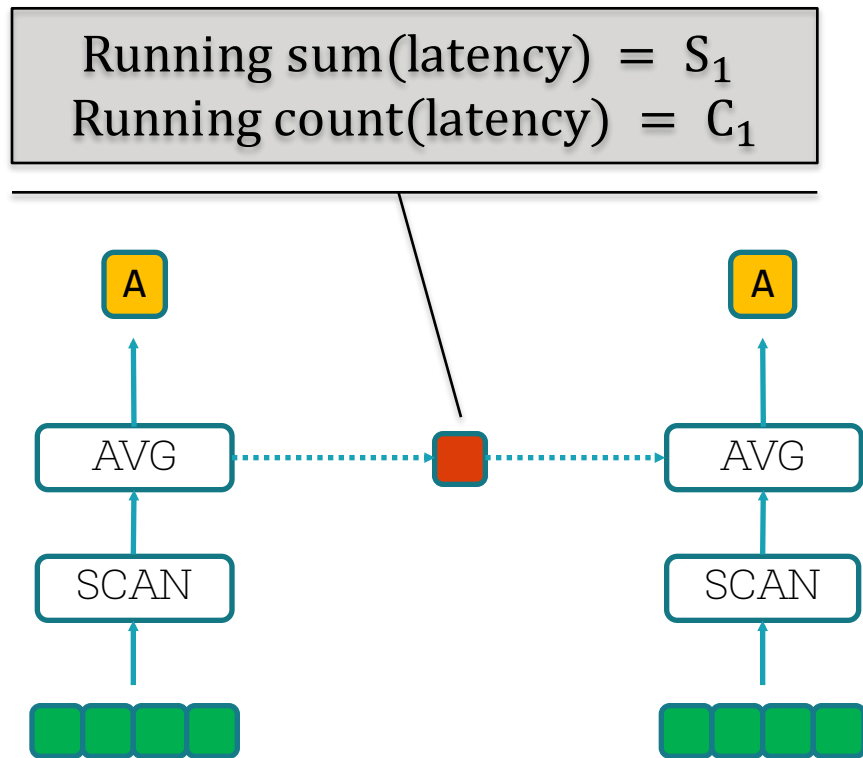
Delta Update: Simple Queries

```
SELECT avg(latency)  
FROM log
```



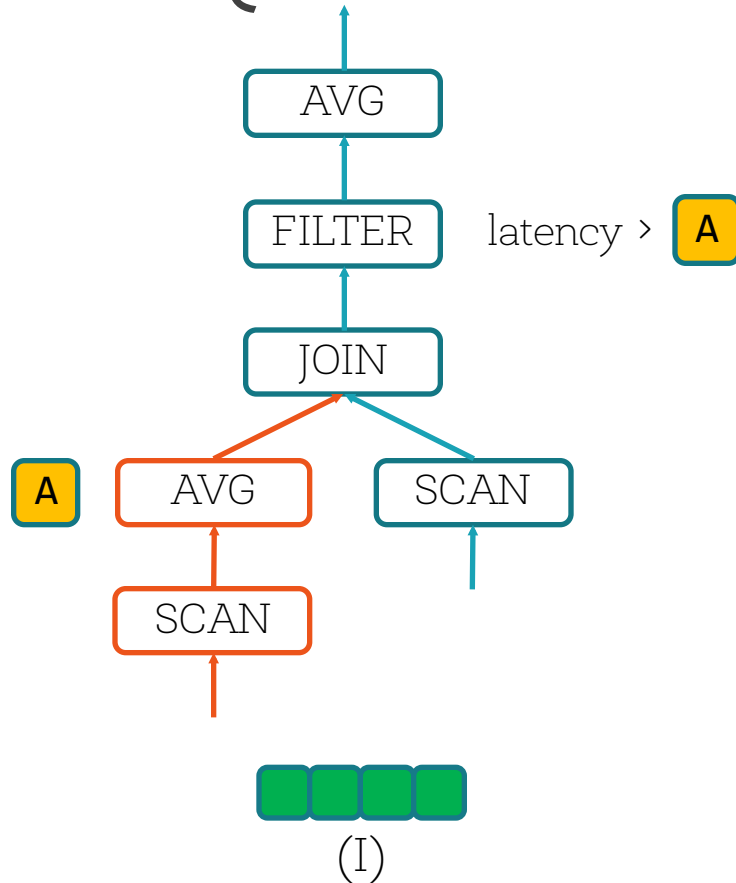
Delta Update: Simple Queries

```
SELECT avg(latency)  
FROM log
```



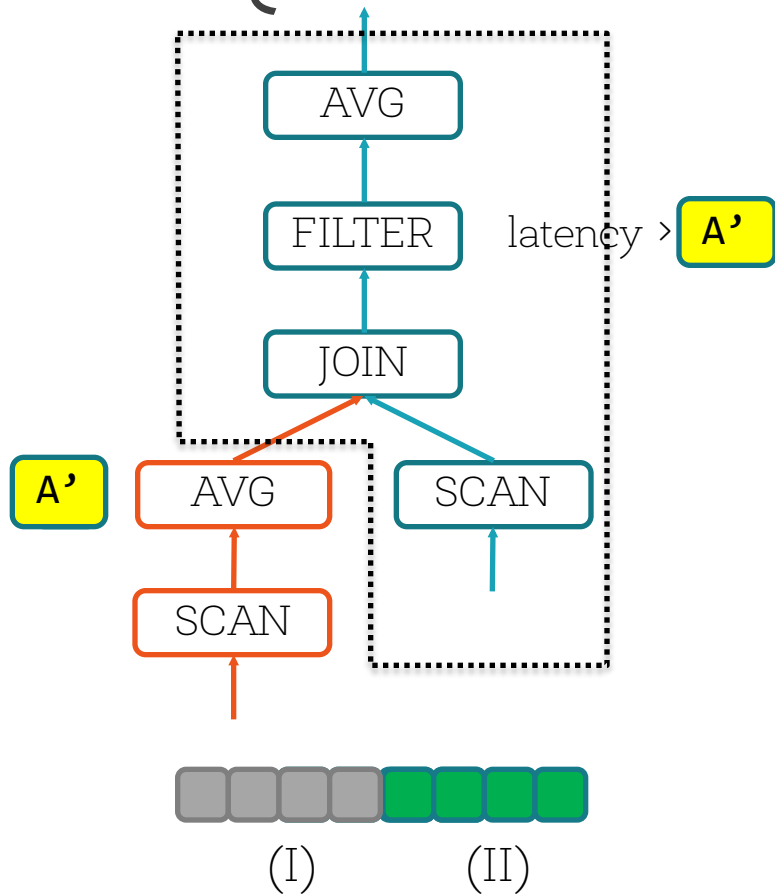
Delta Update: Nested Queries

```
SELECT avg(latency)
FROM log
WHERE latency >
(
  SELECT avg(latency)
  FROM log
)
```




Delta Update: Nested Queries

```
SELECT avg(latency)
FROM log
WHERE latency >
(
  SELECT avg(latency)
  FROM log
)
```



Delta Update: Nested Queries


```
SELECT avg(latency)
FROM log
WHERE latency >
(
  SELECT avg(latency)
  FROM log
)
```



Delta Update: Nested Queries

```
SELECT avg(latency)
FROM log
WHERE latency >
```


```
(
  SELECT avg(latency)
  FROM log
)
```



Delta Update: Nested Queries

```
SELECT avg(latency)
FROM log
WHERE latency >
```

```
(
  SELECT avg(latency)
  FROM log
)
```



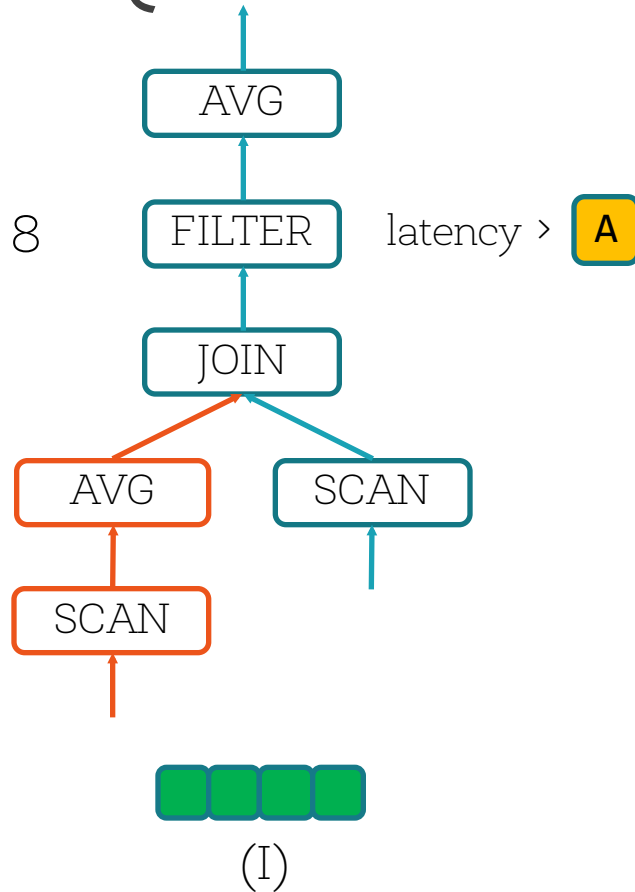
Delta Update: Nested Queries

```
SELECT avg(latency)
FROM log
WHERE latency >
```

~~latency < 8~~

```
(
  SELECT avg(latency)
  FROM log
)
```

10 ± 2



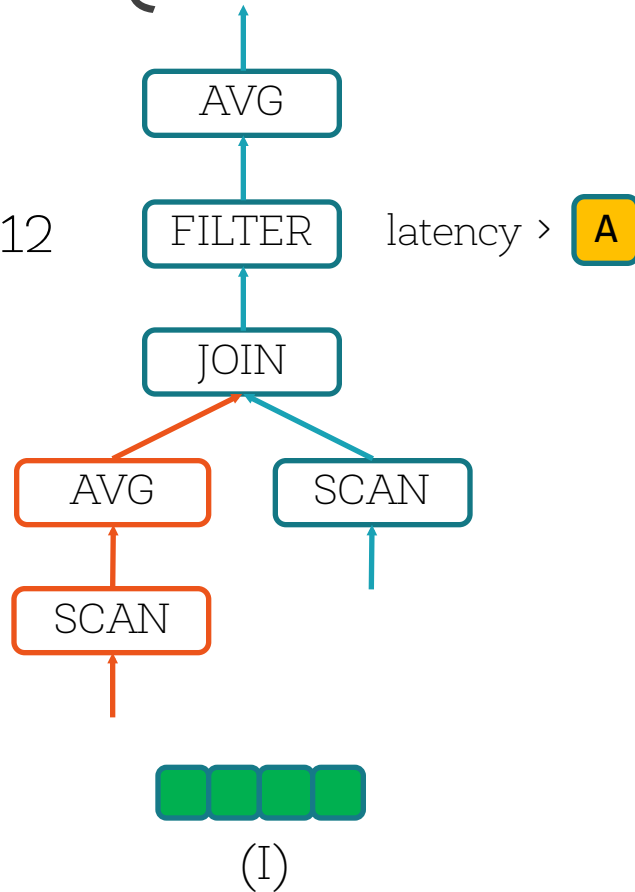
Delta Update: Nested Queries

```
SELECT avg(latency)
FROM log
WHERE latency >
```

✓ latency > 12

```
(
  SELECT avg(latency)
  FROM log
)
```

10 ± 2

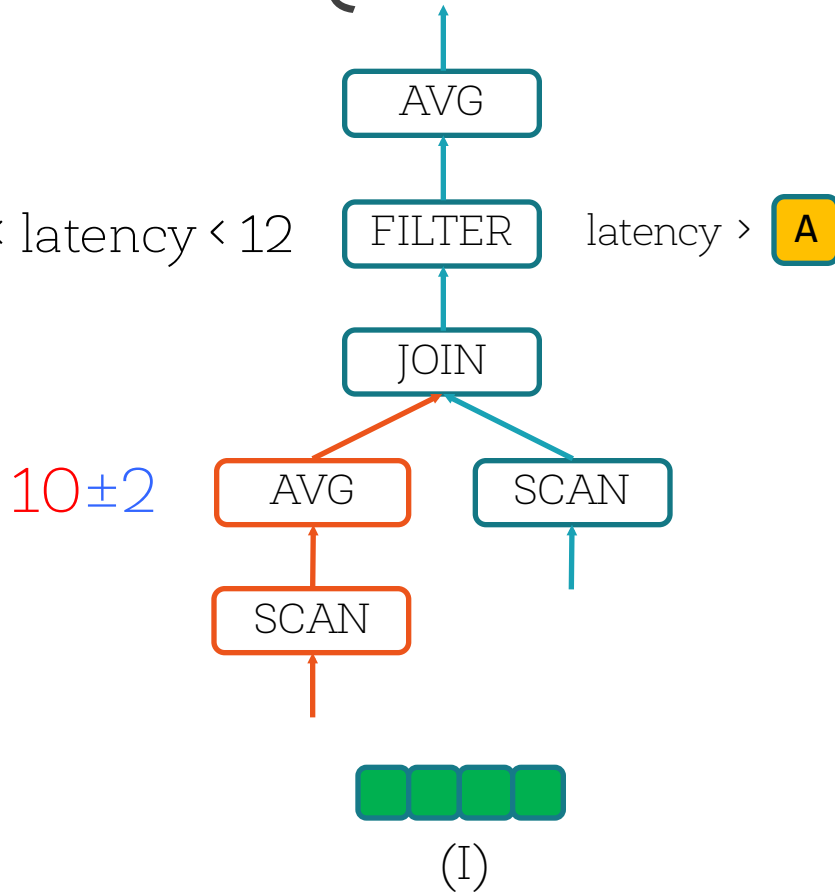


Delta Update: Nested Queries

```
SELECT avg(latency)
FROM log
WHERE latency >
```

? $8 < \text{latency} < 12$

```
(
  SELECT avg(latency)
  FROM log
)
```

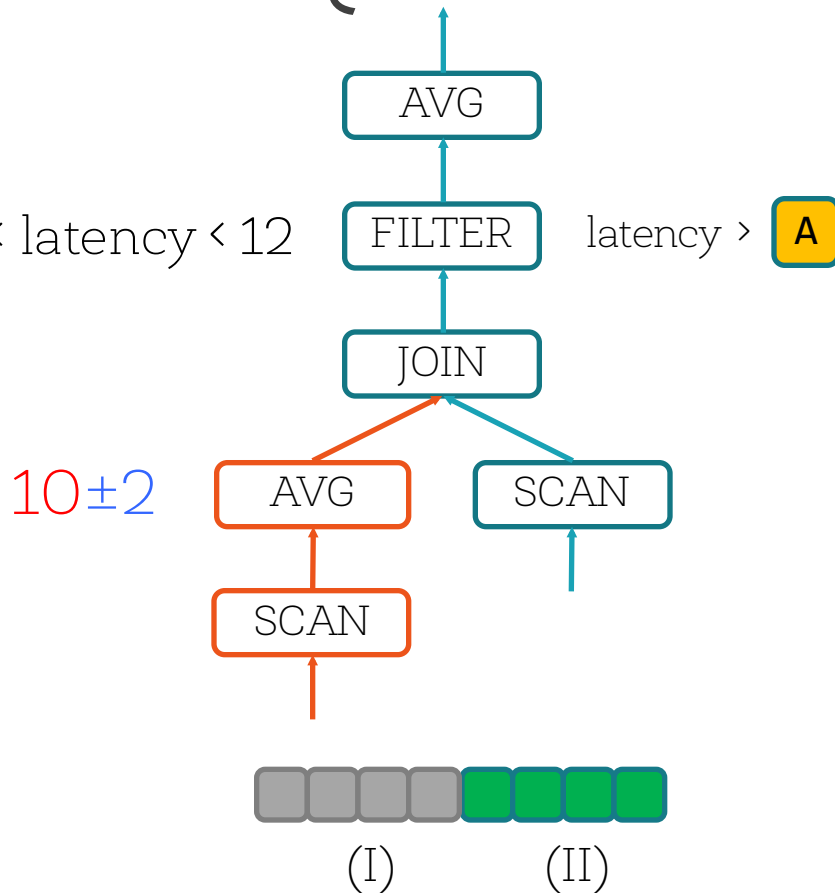


Delta Update: Nested Queries

```
SELECT avg(latency)
FROM log
WHERE latency >
```

? $8 < \text{latency} < 12$

```
(
  SELECT avg(latency)
  FROM log
)
```



High Level Take-away:

Introduce Delta Update Queries as a First Class Citizen in Query Execution

Check out our code!

1. **Code Preview:** <http://github.com/amplab/bootstrap-sql>.
Send us an email to kaizeng@cs.berkeley.edu and sameer@databricks.com to get access!
2. Spark Package in July'15
3. Gradual Native SparkSQL Integration in 1.5, 1.6 and beyond

Conclusion

1. Continuous Query Execution on Samples of Data is an important means to achieve **interactivity** in processing large datasets
2. New SparkSQL Libraries:
 - **BlinkDB** for Continuous Error Bars
 - **G-OLA** for Continuous Partial Answers

Thank you.

Sameer Agarwal (sameer@databricks.com)

Kai Zeng (kaizeng@cs.berkeley.edu)

