


# BlinkDB and G-OLA:

## Supporting Approximate Answers in SparkSQL

Sameer Agarwal and Kai Zeng

Spark Summit | San Francisco, CA | June 15<sup>th</sup> 2015



# About Us

## 1. Sameer Agarwal

- Software Engineer at Databricks
- PhD in Databases (UC Berkeley)
- Research on Approximate Query Processing (BlinkDB)

## 2. Kai Zeng

- Post-doc in AMP Lab/ Intern at Databricks
- PhD in Databases (UCLA)
- Research on Approximate Query Processing (ABM)

# 100 TB on 1000 machines

$\frac{1}{2}$  - 1 Hour



Hard Disks

1 - 5 Minutes



Memory

1 second



Continuous Query Execution on Samples of Data

# Continuous Query Execution on Samples

| ID | City | Latency |
|----|------|---------|
| 1  | NYC  | 30      |
| 2  | NYC  | 38      |
| 3  | SLC  | 34      |
| 4  | LA   | 36      |
| 5  | SLC  | 37      |
| 6  | SF   | 28      |
| 7  | NYC  | 32      |
| 8  | NYC  | 38      |
| 9  | LA   | 36      |
| 10 | SF   | 35      |
| 11 | NYC  | 38      |
| 12 | LA   | 34      |

What is the average latency in the table?

34.6667

# Continuous Query Execution on Samples

| ID | City | Latency |
|----|------|---------|
| 1  | NYC  | 30      |
| 2  | NYC  | 38      |
| 3  | SLC  | 34      |
| 4  | LA   | 36      |
| 5  | SLC  | 37      |
| 6  | SF   | 28      |
| 7  | NYC  | 32      |
| 8  | NYC  | 38      |
| 9  | LA   | 36      |
| 10 | SF   | 35      |
| 11 | NYC  | 38      |
| 12 | LA   | 34      |

What is the average latency in the table?

35

# Continuous Query Execution on Samples

| ID | City | Latency |
|----|------|---------|
| 1  | NYC  | 30      |
| 2  | NYC  | 38      |
| 3  | SLC  | 34      |
| 4  | LA   | 36      |
| 5  | SLC  | 37      |
| 6  | SF   | 28      |
| 7  | NYC  | 32      |
| 8  | NYC  | 38      |
| 9  | LA   | 36      |
| 10 | SF   | 35      |
| 11 | NYC  | 38      |
| 12 | LA   | 34      |

What is the average latency in the table?

35 ± 2.1

# Continuous Query Execution on Samples

| ID | City | Latency |
|----|------|---------|
| 1  | NYC  | 30      |
| 2  | NYC  | 38      |
| 3  | SLC  | 34      |
| 4  | LA   | 36      |
| 5  | SLC  | 37      |
| 6  | SF   | 28      |
| 7  | NYC  | 32      |
| 8  | NYC  | 38      |
| 9  | LA   | 36      |
| 10 | SF   | 35      |
| 11 | NYC  | 38      |
| 12 | LA   | 34      |

What is the average latency in the table?

35  $\pm 2.1$   
33.83  $\pm 1.3$

# Continuous Query Execution on Samples

| ID | City | Latency |
|----|------|---------|
| 1  | NYC  | 30      |
| 2  | NYC  | 38      |
| 3  | SLC  | 34      |
| 4  | LA   | 36      |
| 5  | SLC  | 37      |
| 6  | SF   | 28      |
| 7  | NYC  | 32      |
| 8  | NYC  | 38      |
| 9  | LA   | 36      |
| 10 | SF   | 35      |
| 11 | NYC  | 38      |
| 12 | LA   | 34      |

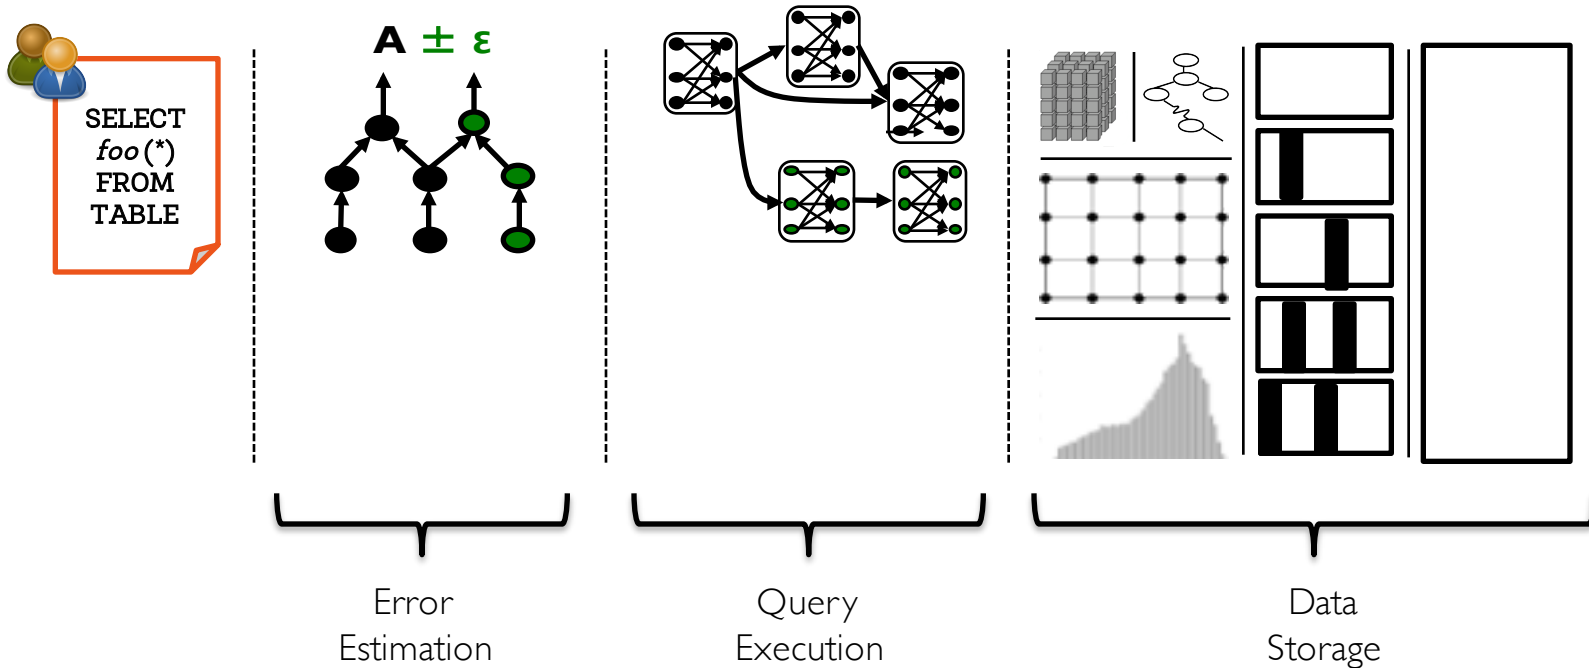
What is the average latency in the table?

35  $\pm 2.1$   
33.83  $\pm 1.3$   
34.6667  $\pm 0.0$



# Demo


# Continuous Query Execution on Samples



# Continuous Query Execution on Samples




SELECT  
*foo*(\*)  
FROM  
TABLE



$A \pm \epsilon$





BlinkDB

Error  
Estimation



G-OLA

Query  
Execution



Data  
Storage

# Interface

```
val dataframe =  
    sqlCtx.sql("select avg(latency) from log")  
  
// batch processing  
val result = dataframe.collect() // 34.6667
```

# Interface

```
val dataframe =  
    sqlCtx.sql("select avg(latency) from log")  
  
// online processing  
val onlineDataFrame = dataframe.online  
onlineDataFrame.collectNext() // 35 ± 2.1  
onlineDataFrame.collectNext() // 33.83 ± 1.3
```

# Interface

```
val dataframe =  
    sqlCtx.sql("select avg(latency) from log")  
  
// online processing  
val onlineDataFrame = dataframe.online  
while (onlineDataFrame.hasNext()) {  
    onlineDataFrame.collectNext()  
}
```

# Interface

```
val dataframe =  
    sqlCtx.sql("select avg(latency) from log")  
  
// online processing  
val onlineDataFrame = dataframe.online  
while (onlineDataFrame.hasNext() &&  
    responseTime <= 10.seconds) {  
    onlineDataFrame.collectNext()  
}
```

# Interface

```
val dataframe =  
    sqlCtx.sql("select avg(latency) from log")  
  
// online processing  
val onlineDataFrame = dataframe.online  
while (onlineDataFrame.hasNext() &&  
    errorBound >= 0.01) {  
    onlineDataFrame.collectNext()  
}
```



# Interface

```
val dataframe =  
    sqlCtx.sql("select avg(latency) from log")  
  
// online processing  
val onlineDataFrame = dataframe.online  
while (onlineDataFrame.hasNext() &&  
    userEvent.cancelled()) {  
    onlineDataFrame.collectNext()  
}
```

# Interface


```
val dataframe =  
    sqlCtx.sql("select avg(latency) from log")
```

```
// online processing
```


```
val onlineDataFrame = dataframe.online  
while (onlineDataFrame.hasNext() &&  
    userEvent.cancelled()) {  
    onlineDataFrame.collectNext()  
}
```

AGGREGATES/ UDAFs  
JOINS/GROUP BYs  
NESTED QUERIES

# Continuous Query Execution on Samples



# Continuous Query Execution on Samples



Sameer Agarwal, Barzan Mozafari, Aurojit Panda, Henry Milner, Samuel Madden, Ion Stoica. **BlinkDB: Queries with Bounded Errors and Bounded Response Times on Very Large Data.** In ACM EuroSys 2013.

Ariel Kleiner, Ameet Talwalkar, Sameer Agarwal, Ion Stoica, Michael Jordan. **A General Bootstrap Performance Diagnostic.** In ACM KDD 2013

Sameer Agarwal, Henry Milner, Ariel Kleiner, Ameet Talwalkar, Michael Jordan, Samuel Madden, Barzan Mozafari, Ion Stoica. **Knowing When You're Wrong: Building Fast and Reliable Approximate Query Processing Systems.** In ACM SIGMOD 2014.

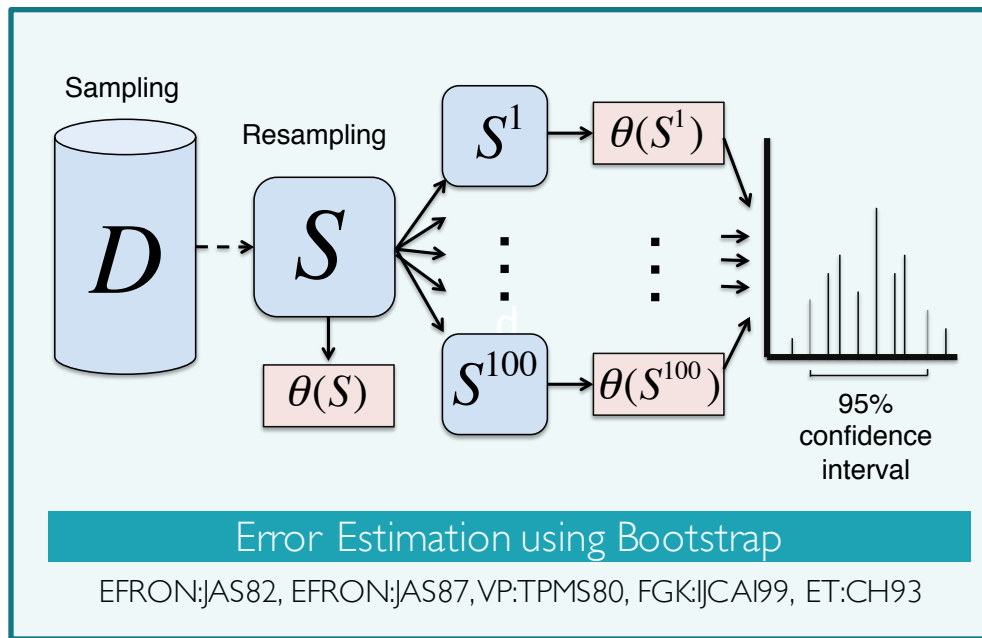
# Error Estimation on a Sample of Data

Focused on estimating aggregate errors given *representative* samples

1. Count:  $N(np, n(1 - p)p)$
2. Sum:  $N(np\mu, np(\sigma^2 + (1 - p)\mu^2))$
3. Mean:  $N(\mu, \sigma^2/n)$
4. Variance:  $N(\sigma^2, (\mu_4 - \sigma^4)/n)$
5. Stddev:  $N(\sigma, (\mu_4 - \sigma^4)/(4\sigma^2n))$

Central Limit Theorem (CLT)

HOE:ASTAT63,BIL: WILEY86, CGL:ASTAT83,PH:IBM96



# Error Estimation using Bootstrap

| ID | City | Latency |
|----|------|---------|
| 1  | NYC  | 30      |
| 2  | NYC  | 38      |
| 3  | SLC  | 34      |
| 4  | LA   | 36      |
| 5  | SLC  | 37      |
| 6  | SF   | 28      |
| 7  | NYC  | 32      |
| 8  | NYC  | 38      |
| 9  | LA   | 36      |
| 10 | SF   | 35      |
| 11 | NYC  | 38      |
| 12 | LA   | 34      |

What is the average latency in the table?

# Error Estimation using Bootstrap

| ID | City | Latency |
|----|------|---------|
| 1  | NYC  | 30      |
| 2  | NYC  | 38      |
| 3  | SLC  | 34      |
| 4  | LA   | 36      |

| ID | City | Latency |
|----|------|---------|
| 1  | NYC  | 30      |
| 2  | NYC  | 38      |
| 3  | SLC  | 34      |
| 4  | SLC  | 34      |

| ID | City | Latency |
|----|------|---------|
| 1  | NYC  | 30      |
| 2  | NYC  | 30      |
| 3  | SLC  | 34      |
| 4  | LA   | 36      |

...

| ID | City | Latency |
|----|------|---------|
| 1  | SLC  | 34      |
| 2  | LA   | 36      |
| 3  | SLC  | 34      |
| 4  | LA   | 36      |

$$\theta_1 = 34$$

$$\theta_2 = 32.5$$

...

$$\theta_{100} = 35$$

$$34.5 \pm 2$$

# Error Estimation using Bootstrap

| ID | City | Latency |
|----|------|---------|
| 1  | NYC  | 30      |
| 2  | NYC  | 38      |
| 3  | SLC  | 34      |
| 4  | LA   | 36      |
| 5  | SLC  | 37      |
| 6  | SF   | 28      |
| 7  | NYC  | 32      |
| 8  | NYC  | 38      |
| 9  | LA   | 36      |
| 10 | SF   | 35      |
| 11 | NYC  | 38      |
| 12 | LA   | 34      |

| ID | City | Latency |
|----|------|---------|
| 1  | NYC  | 30      |
| 2  | NYC  | 38      |
| 3  | SLC  | 34      |
| 4  | SLC  | 34      |

| ID | City | Latency |
|----|------|---------|
| 1  | NYC  | 30      |
| 2  | NYC  | 30      |
| 3  | SLC  | 34      |
| 4  | LA   | 36      |

| ID | City | Latency |
|----|------|---------|
| 1  | SLC  | 34      |
| 2  | LA   | 36      |
| 3  | SLC  | 34      |
| 4  | LA   | 36      |

$$\theta_1 = 34$$

$$\theta_2 = 32.5$$

...

$$\theta_{100} = 35$$

$$34.5 \pm 2$$



# Error Estimation using Bootstrap

| ID | City | Latency |
|----|------|---------|
| 1  | NYC  | 30      |
| 2  | NYC  | 38      |
| 3  | SLC  | 34      |
| 4  | LA   | 36      |
| 5  | SLC  | 37      |
| 6  | SF   | 28      |
| 7  | NYC  | 32      |
| 8  | NYC  | 38      |
| 9  | LA   | 36      |
| 10 | SF   | 35      |
| 11 | NYC  | 38      |
| 12 | LA   | 34      |

| ID | City | Latency |
|----|------|---------|
| 1  | NYC  | 30      |
| 2  | NYC  | 38      |
| 3  | SLC  | 34      |
| 4  | SLC  | 34      |
| 5  | SLC  | 37      |

$$\theta_1 = 34.6$$

| ID | City | Latency |
|----|------|---------|
| 1  | SLC  | 37      |
| 2  | NYC  | 30      |
| 3  | SLC  | 34      |
| 4  | LA   | 36      |
| 5  | NYC  | 30      |

$$\theta_2 = 33.4$$

| ID | City | Latency |
|----|------|---------|
| 1  | SLC  | 34      |
| 2  | SLC  | 37      |
| 3  | SLC  | 34      |
| 4  | LA   | 36      |
| 5  | LA   | 36      |

...

$$\theta_{100} = 35.4$$

$$35 \pm 1.6$$

# Error Estimation in BlinkDB

| ID | City | Latency |
|----|------|---------|
| 1  | NYC  | 30      |
| 2  | NYC  | 38      |
| 3  | SLC  | 34      |
| 4  | LA   | 36      |
| 5  | SLC  | 37      |
| 6  | SF   | 28      |
| 7  | NYC  | 32      |
| 8  | NYC  | 38      |
| 9  | LA   | 36      |
| 10 | SF   | 35      |
| 11 | NYC  | 38      |
| 12 | LA   | 34      |

What is the average latency in the table?

| ID | City | Latency |
|----|------|---------|
| 1  | NYC  | 30      |
| 2  | NYC  | 38      |
| 3  | SLC  | 34      |
| 4  | SLC  | 34      |
| 5  | SLC  | 37      |

Leverage Poissonized Resampling to generate samples with replacement

# Error Estimation in BlinkDB

| ID | City | Latency |
|----|------|---------|
| 1  | NYC  | 30      |
| 2  | NYC  | 38      |
| 3  | SLC  | 34      |
| 4  | LA   | 36      |
| 5  | SLC  | 37      |
| 6  | SF   | 28      |
| 7  | NYC  | 32      |
| 8  | NYC  | 38      |
| 9  | LA   | 36      |
| 10 | SF   | 35      |
| 11 | NYC  | 38      |
| 12 | LA   | 34      |

What is the average latency in the table?

| ID | City | Latency | # <sub>1</sub> |
|----|------|---------|----------------|
| 1  | NYC  | 30      | 2              |
| 2  | NYC  | 38      | 1              |
| 3  | SLC  | 34      | 0              |
| 4  | SLC  | 34      | 1              |
| 5  | SLC  | 37      | 1              |

Sample from a  
Poisson (1) Distribution

$$\theta_1 = 33.8$$

# Error Estimation in BlinkDB

| ID | City | Latency |
|----|------|---------|
| 1  | NYC  | 30      |
| 2  | NYC  | 38      |
| 3  | SLC  | 34      |
| 4  | LA   | 36      |
| 5  | SLC  | 37      |
| 6  | SF   | 28      |
| 7  | NYC  | 32      |
| 8  | NYC  | 38      |
| 9  | LA   | 36      |
| 10 | SF   | 35      |
| 11 | NYC  | 38      |
| 12 | LA   | 34      |

What is the average latency in the table?

| ID | City | Latency | # <sub>1</sub> |
|----|------|---------|----------------|
| 1  | NYC  | 30      | 2              |
| 2  | NYC  | 38      | 1              |
| 3  | SLC  | 34      | 0              |
| 4  | SLC  | 34      | 1              |
| 5  | SLC  | 37      | 1              |
| 6  | SF   | 28      | 2              |

Incremental Error Estimation

# Error Estimation in BlinkDB

| ID | City | Latency |
|----|------|---------|
| 1  | NYC  | 30      |
| 2  | NYC  | 38      |
| 3  | SLC  | 34      |
| 4  | LA   | 36      |
| 5  | SLC  | 37      |
| 6  | SF   | 28      |
| 7  | NYC  | 32      |
| 8  | NYC  | 38      |
| 9  | LA   | 36      |
| 10 | SF   | 35      |
| 11 | NYC  | 38      |
| 12 | LA   | 34      |

What is the average latency in the table?

| ID | City | Latency | # <sub>1</sub> | # <sub>2</sub> |
|----|------|---------|----------------|----------------|
| 1  | NYC  | 30      | 2              | 1              |
| 2  | NYC  | 38      | 1              | 0              |
| 3  | SLC  | 34      | 0              | 2              |
| 4  | SLC  | 34      | 1              | 2              |
| 5  | SLC  | 37      | 1              | 0              |
| 6  | SF   | 28      | 2              | 1              |

Construct all Resamples in a Single Pass

0.2-0.5% additional overhead

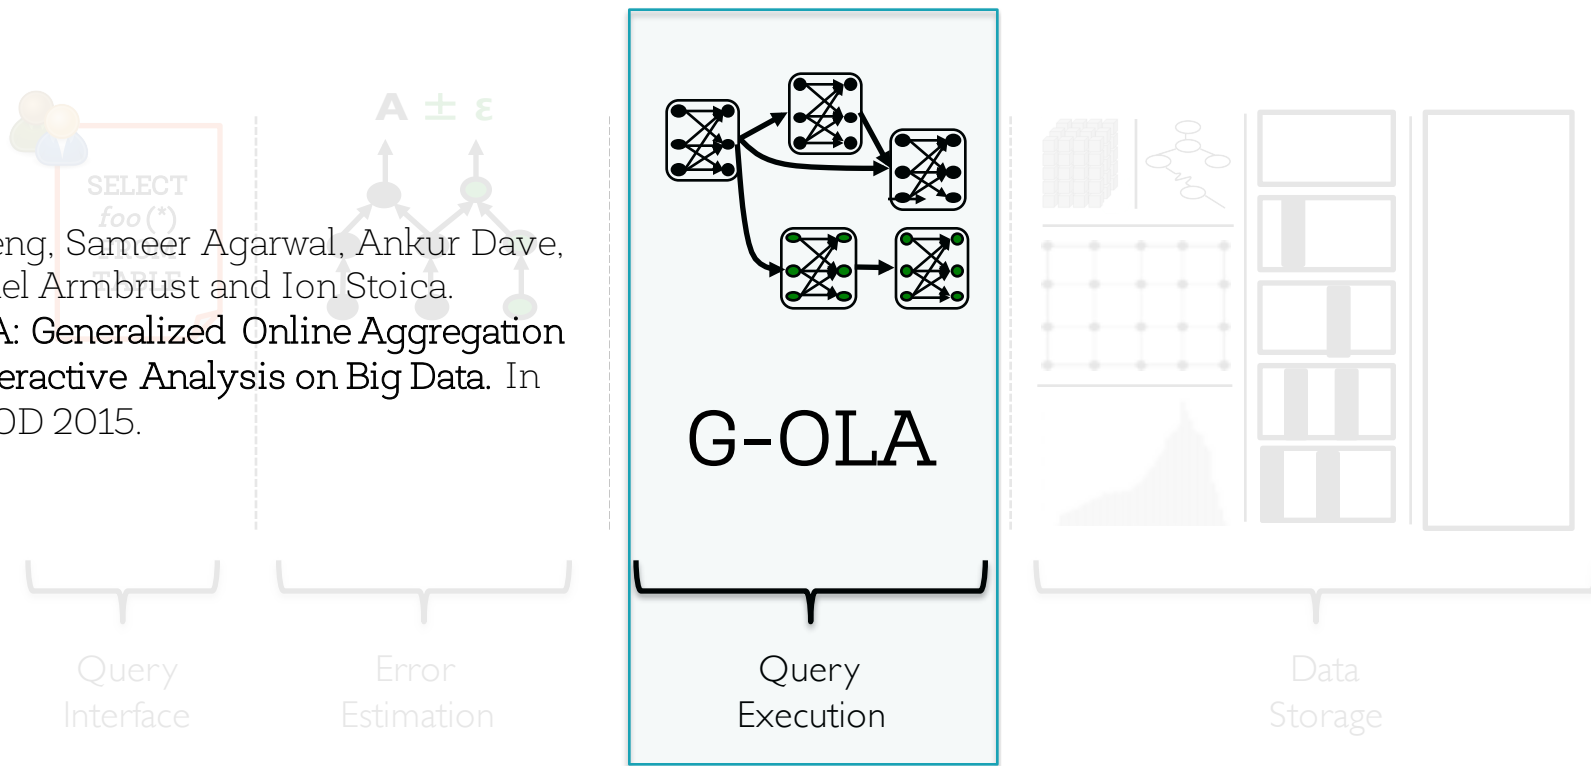
## High Level Take-away:

Bootstrap and Poissonized Resampling Techniques are the key towards achieving quick and continuous error bars for a general set of queries


# Continuous Query Execution on Samples

Kai Zeng, Sameer Agarwal, Ankur Dave,  
Michael Armbrust and Ion Stoica.

G-OLA: Generalized Online Aggregation  
for Interactive Analysis on Big Data. In  
SIGMOD 2015.



# Query Execution: Under The Hood





# Query Execution: Under The Hood


Answer  $\pm \epsilon$

A


A'

Query

Data






10 sec



10+10 sec

# Query Execution: Under The Hood




Answer  $\pm \epsilon$



# Query Execution: Under The Hood



Overall Quadratic Cost!

Answer  $\pm \epsilon$



# Query Execution: Under The Hood

Answer  $\pm \epsilon$




# Query Execution: Under The Hood

Answer  $\pm \epsilon$

A

Query




Data



10 sec


A'




10+10 sec

# Query Execution: Under The Hood


Answer  $\pm \epsilon$



# Query Execution: Under The Hood



# Query Execution: Under The Hood






# Delta Update Queries

Answer  $\pm \epsilon$

Query

Data




# Delta Update Queries

Answer  $\pm \epsilon$

Query

Data




# Delta Update Queries

Answer  $\pm \epsilon$

Query

Data




# Delta Update Queries

Answer  $\pm \epsilon$

Query

Data





# Delta Update Queries

Answer  $\pm \epsilon$

Query

Data





# Delta Update Queries

Answer  $\pm \epsilon$


Query

Data



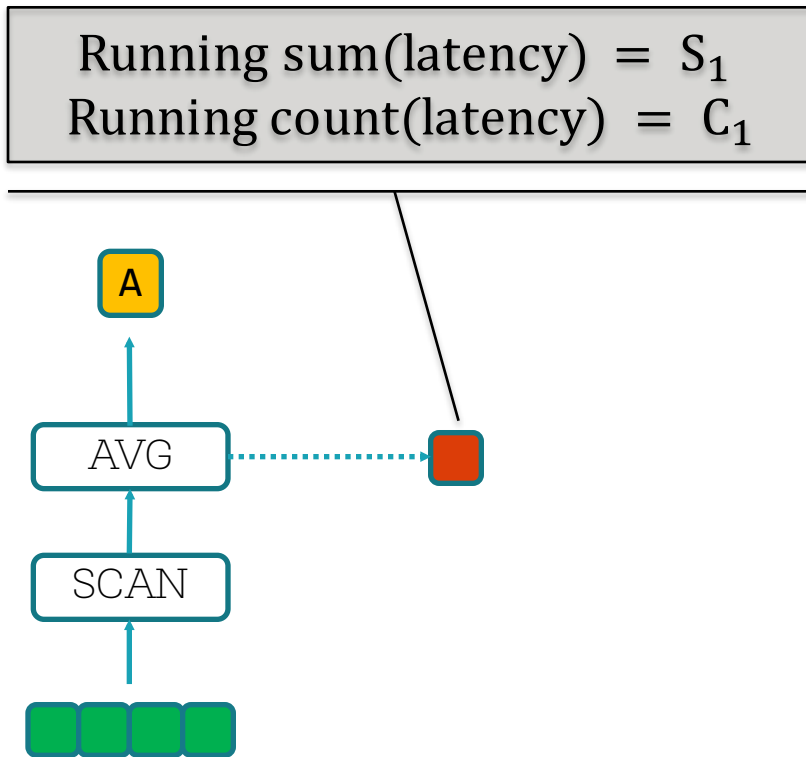
# Delta Update: Simple Queries

```
SELECT avg(latency)  
FROM log
```



# Delta Update: Simple Queries

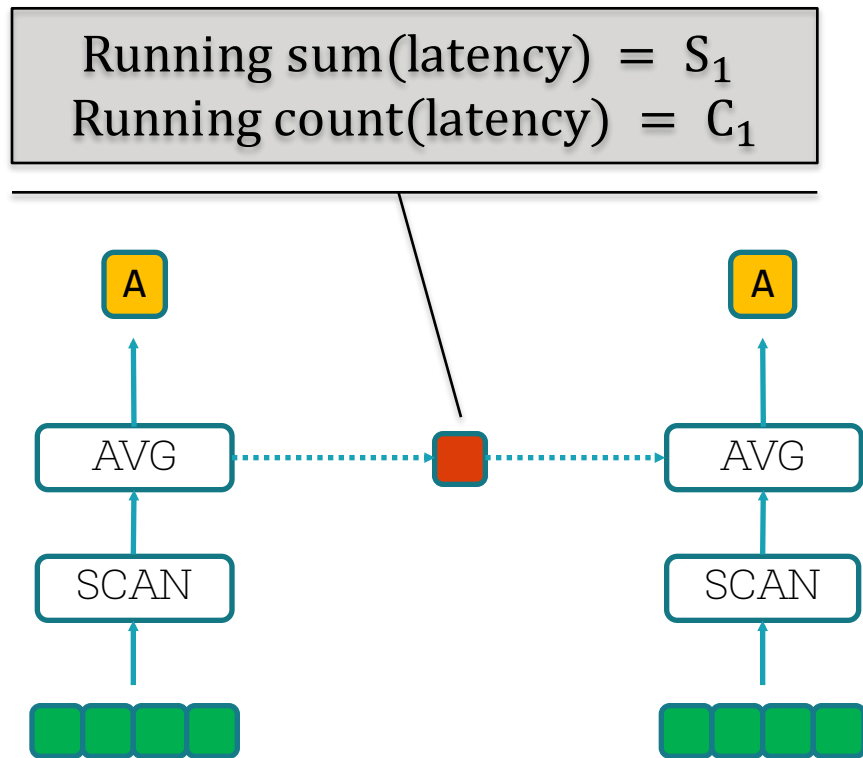
```
SELECT avg(latency)
FROM log
```





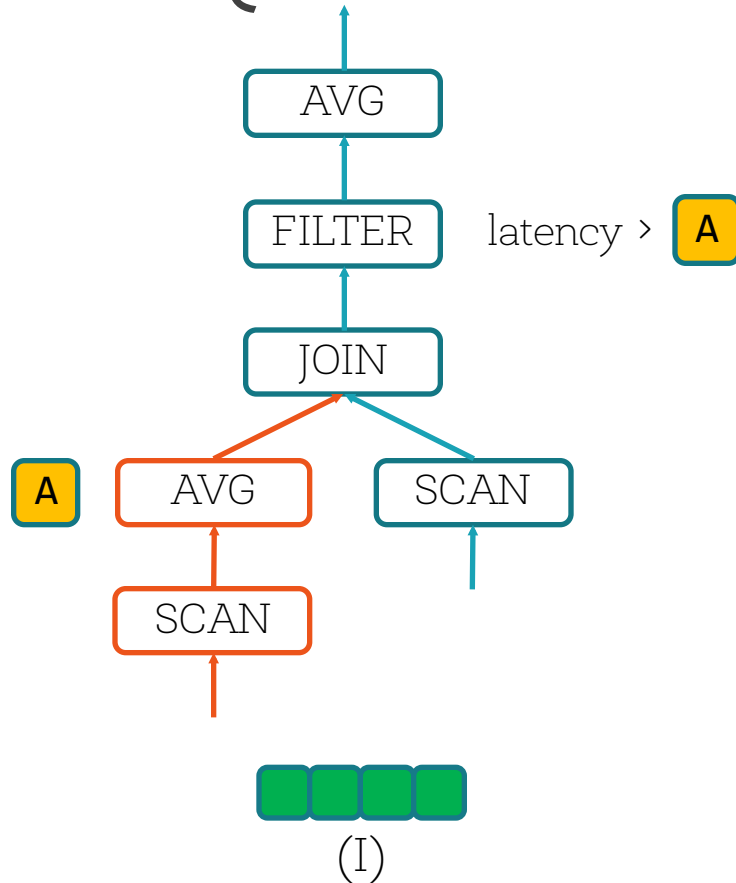
# Delta Update: Simple Queries

```
SELECT avg(latency)  
FROM log
```



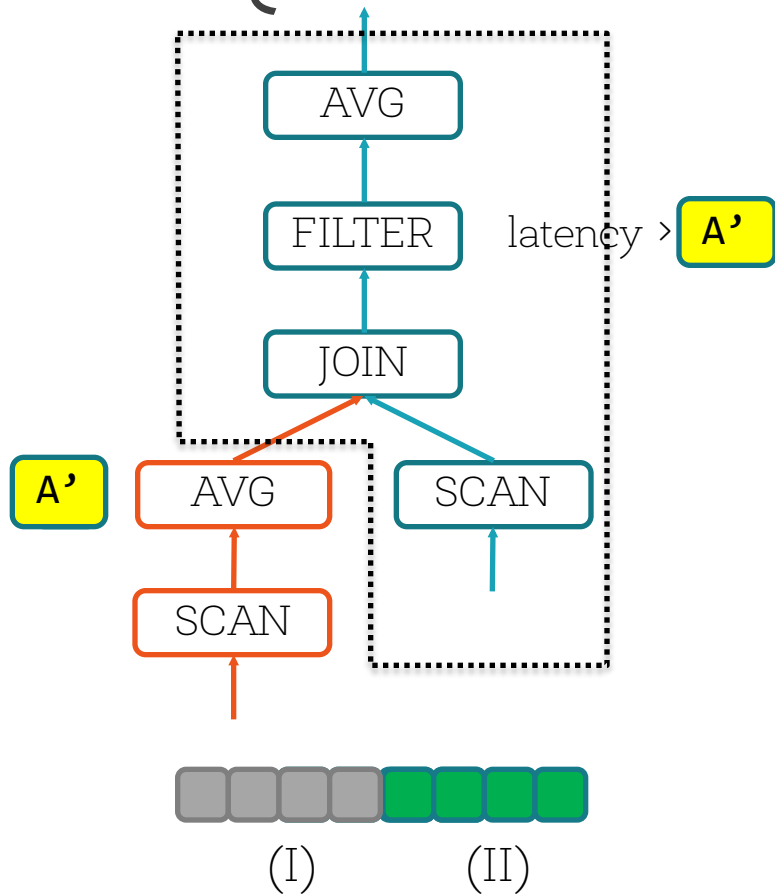
# Delta Update: Nested Queries

```
SELECT avg(latency)
FROM log
WHERE latency >
(
  SELECT avg(latency)
  FROM log
)
```




# Delta Update: Nested Queries

```
SELECT avg(latency)
FROM log
WHERE latency >
(
  SELECT avg(latency)
  FROM log
)
```



# Delta Update: Nested Queries


```
SELECT avg(latency)
FROM log
WHERE latency >
(
  SELECT avg(latency)
  FROM log
)
```



# Delta Update: Nested Queries

```
SELECT avg(latency)
FROM log
WHERE latency >
```


```
(
  SELECT avg(latency)
  FROM log
)
```



# Delta Update: Nested Queries

```
SELECT avg(latency)
FROM log
WHERE latency >
```

```
(
  SELECT avg(latency)
  FROM log
)
```



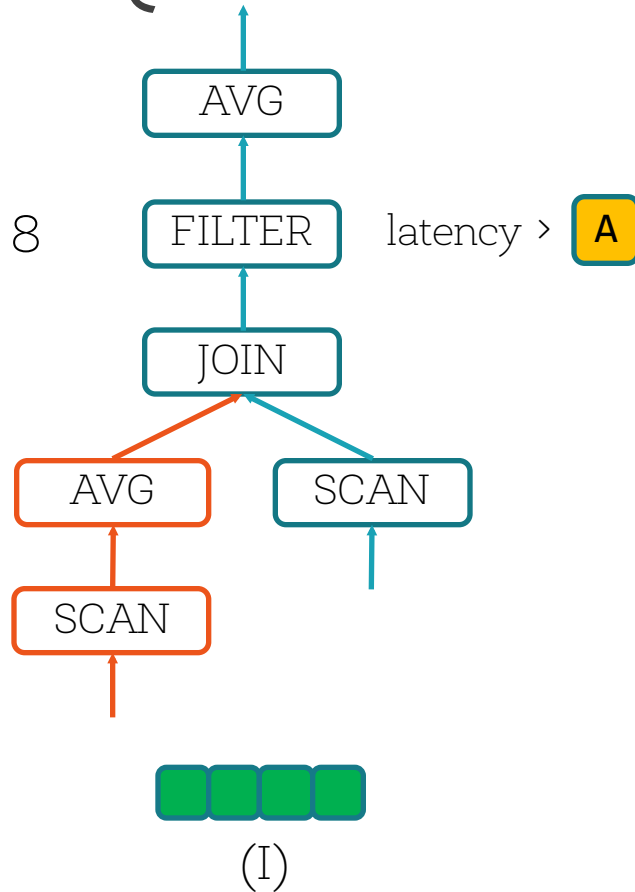
# Delta Update: Nested Queries

```
SELECT avg(latency)
FROM log
WHERE latency >
```

~~latency < 8~~

```
(
  SELECT avg(latency)
  FROM log
)
```

$10 \pm 2$



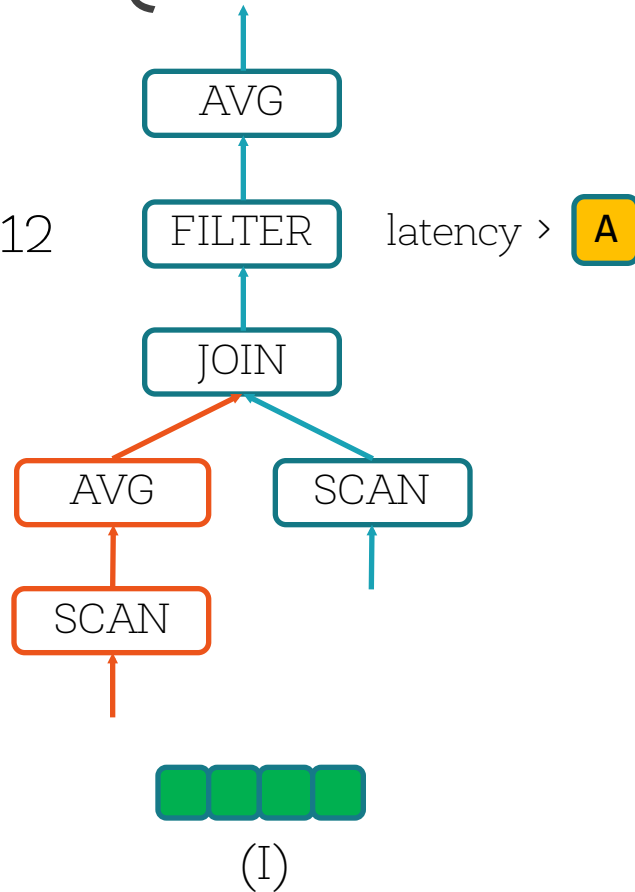
# Delta Update: Nested Queries

```
SELECT avg(latency)
FROM log
WHERE latency >
```

✓ latency > 12

```
(
  SELECT avg(latency)
  FROM log
)
```

$10 \pm 2$



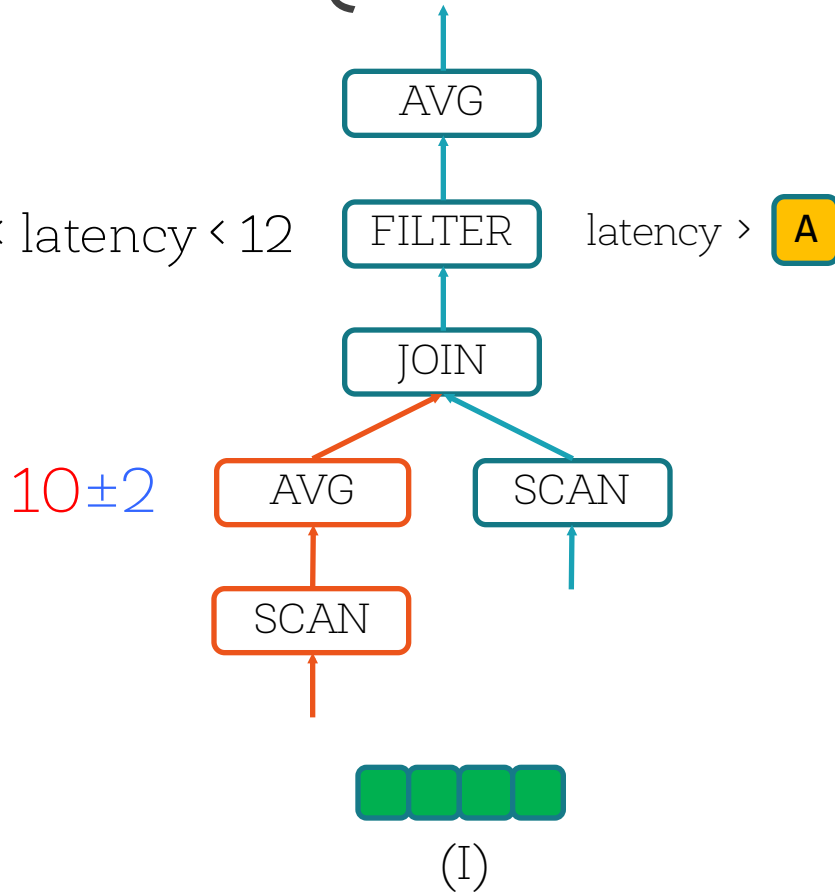


# Delta Update: Nested Queries

```
SELECT avg(latency)
FROM log
WHERE latency >
```

?  $8 < \text{latency} < 12$

```
(
  SELECT avg(latency)
  FROM log
)
```

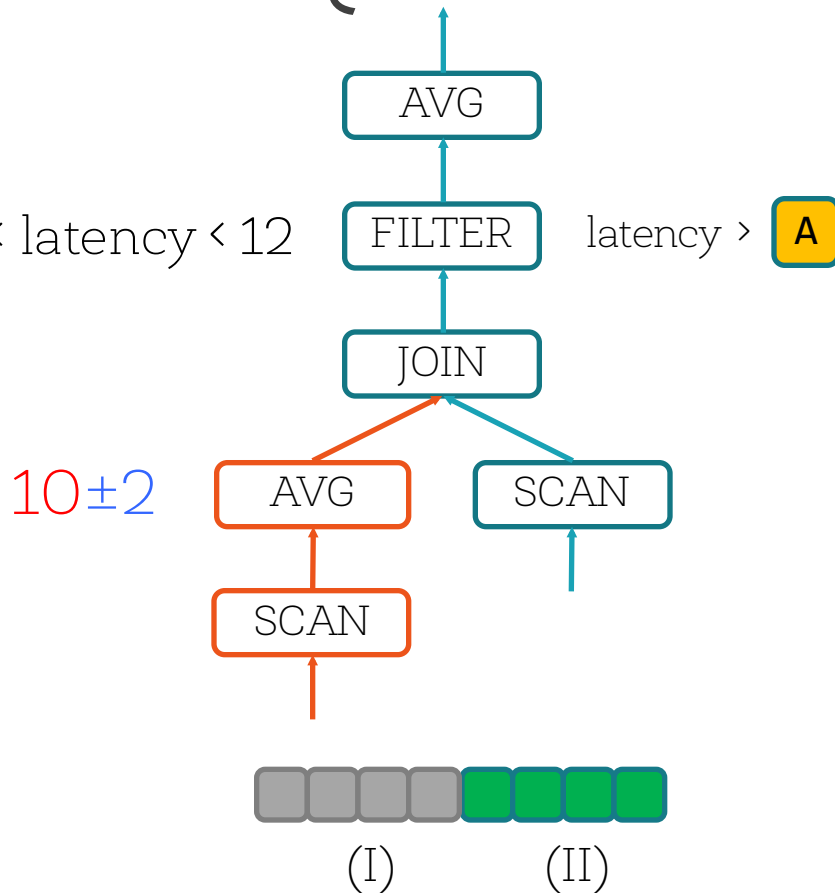


# Delta Update: Nested Queries

```
SELECT avg(latency)
FROM log
WHERE latency >
```

?  $8 < \text{latency} < 12$

```
(
  SELECT avg(latency)
  FROM log
)
```



High Level Take-away:

Introduce Delta Update Queries as a First Class Citizen in Query Execution

# Check out our code!

1. **Code Preview:** <http://github.com/amplab/bootstrap-sql>.  
Send us an email to [kaizeng@cs.berkeley.edu](mailto:kaizeng@cs.berkeley.edu) and [sameer@databricks.com](mailto:sameer@databricks.com) to get access!
2. Spark Package in July'15
3. Gradual Native SparkSQL Integration in 1.5, 1.6 and beyond

# Conclusion

1. Continuous Query Execution on Samples of Data is an important means to achieve **interactivity** in processing large datasets
2. New SparkSQL Libraries:
  - **BlinkDB** for Continuous Error Bars
  - **G-OLA** for Continuous Partial Answers

# Thank you.

Sameer Agarwal (sameer@databricks.com)

Kai Zeng (kaizeng@cs.berkeley.edu)

