

CUDA Performance Study on Hadoop MapReduce Clusters

CSE 930 Advanced Computer Architecture @ Fall 2010

Chen He

Peng Du

che@cse.unl.edu

pdu@cse.unl.edu


University of Nebraska-Lincoln

Overview

- Introduction
- Methodology
- Evaluation
- Conclusions
- Future work

Introduction

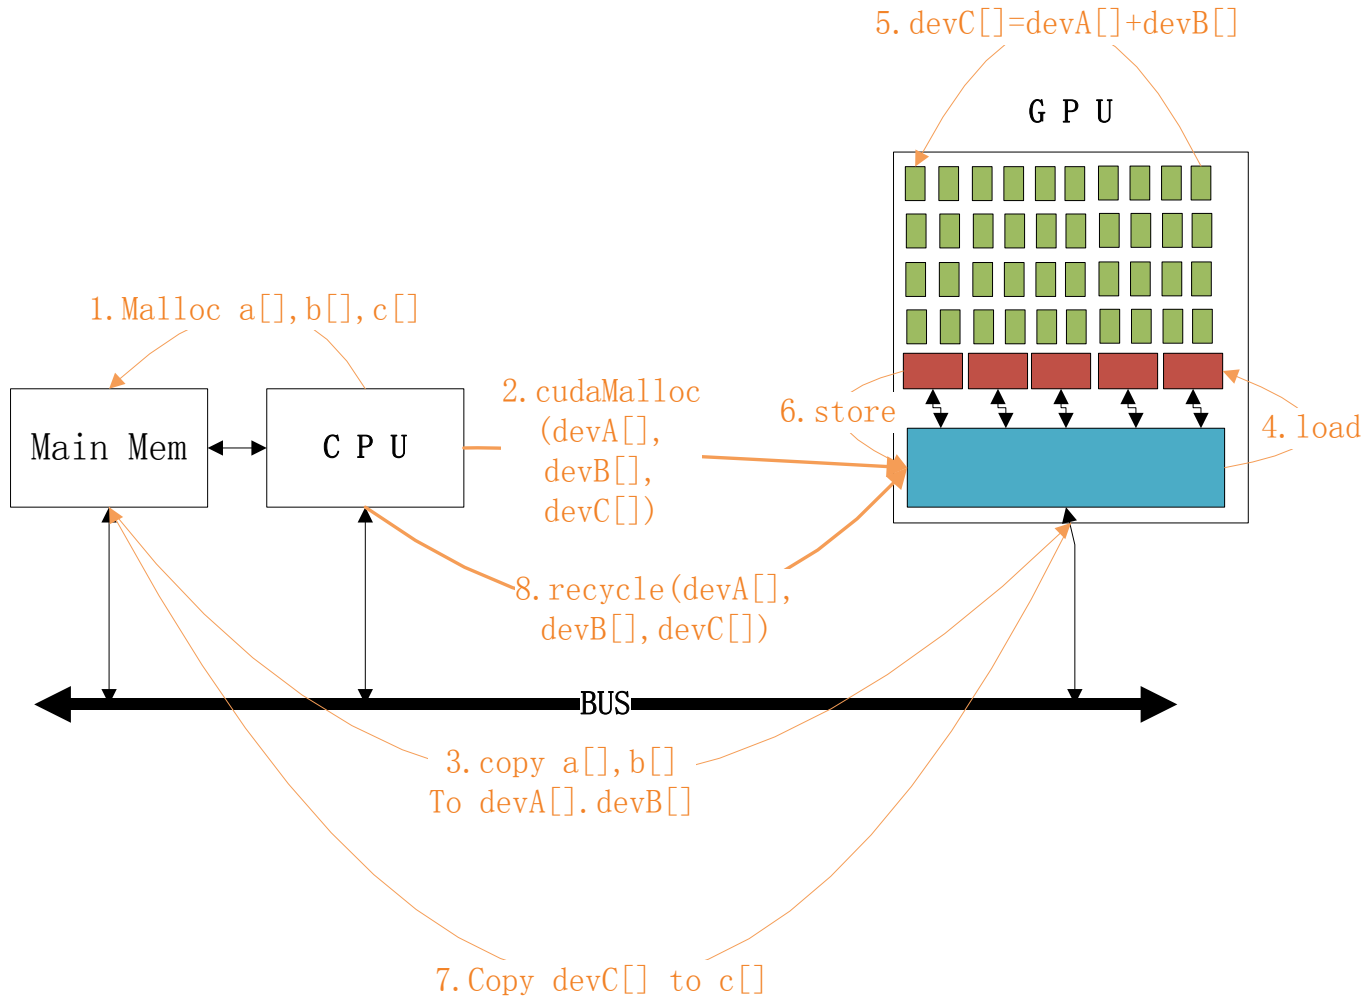
- Hadoop MapReduce



Original Mapper

Introduction

CPU+GPU Architecture





Introduction

- Questions
 - Can we introduce CUDA into Hadoop MapReduce Clusters?
 - Mechanism and implementation
 - Is this reasonable?
 - Effects and Costs

Methodology

- Question-1: Can we introduce CUDA into Hadoop ?



Accelerated
Mapper-2

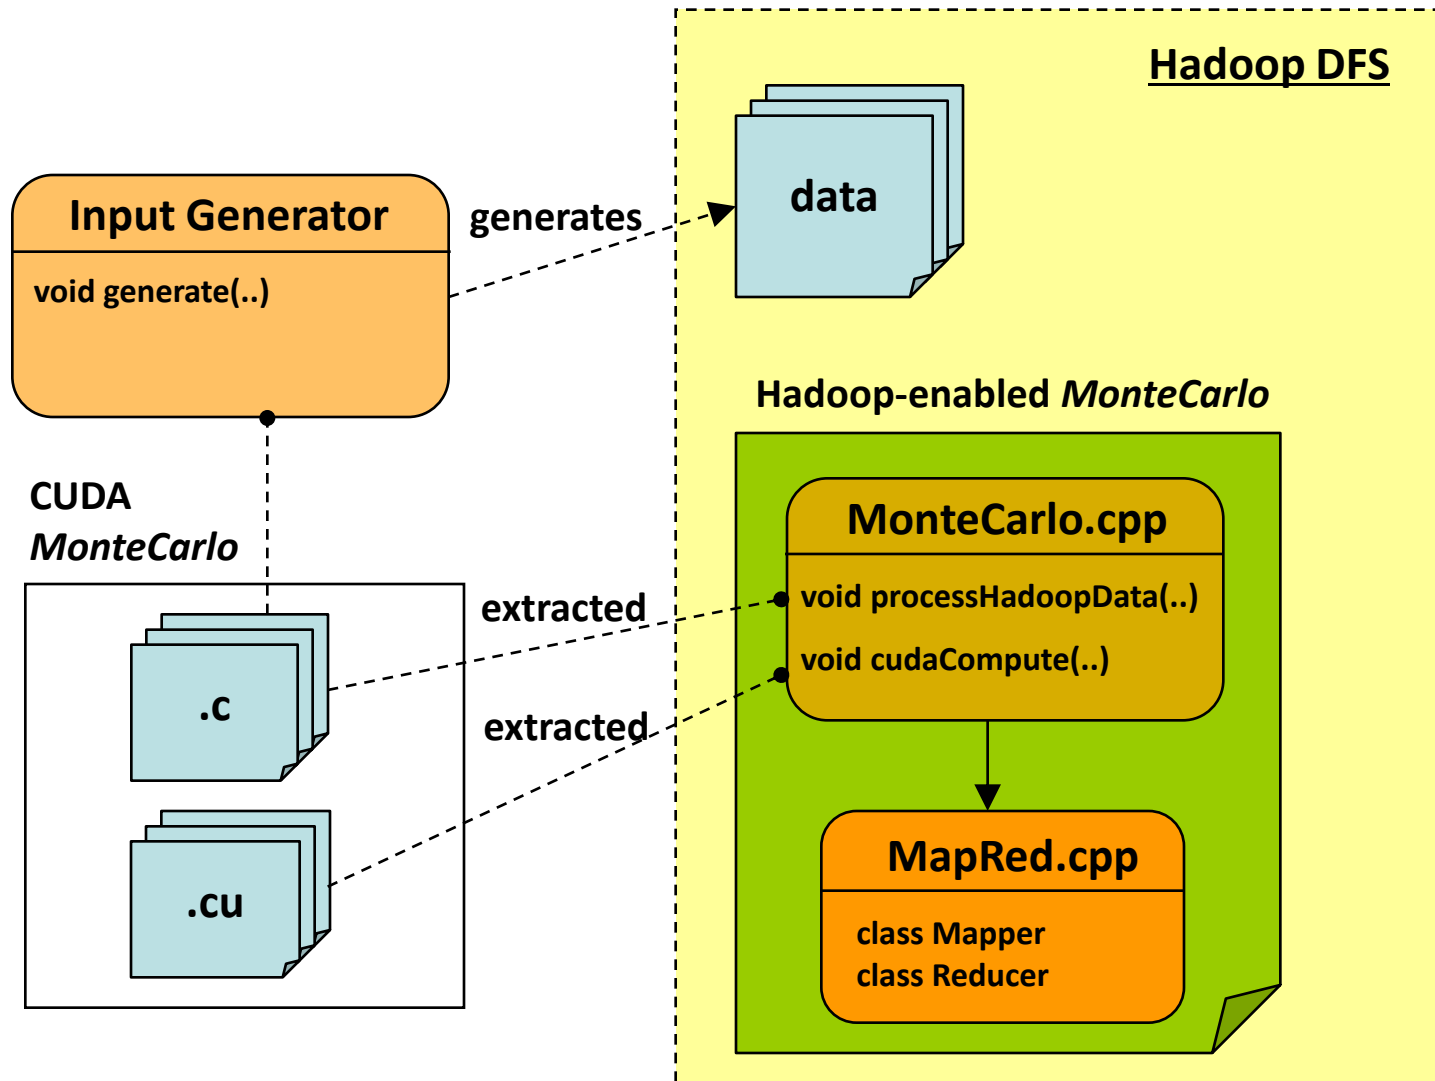
Methodology

- Test cases
 - SDK programs
 - Data intensive: Matrix Multiplication
 - Computation intensive: Monte Carlo
 - MDMR (Molecular Dynamics simulation based on MapReduce)
 - Pure Java program
 - Introduce [JCUDA](#)

Methodology

- Port CUDA programs onto Hadoop
 - GPU (CUDA-C) vs CPU (C)
 - Approach
 - MapRed (*processHadoopData* & *cudaCompute*)
 - Main (*Hadoop Pipes*)
 - Scripts (*runbase.sh*, *run-<prog>-CPU/GPU.sh*)
 - Input data generators

Methodology



Methodology

- MDMR (Molecular Dynamics simulation based on MapReduce)

- Time Complexity by using CPU

$$T(n) = c_1n^2 + c_2n + c_3$$

- We can simply employ GPU to parallel the n-square portion and reduce the time complexity to linear (within the limit of GPU threads)

$$T'(n) = c_1(dn) + c_2n + c_3 + c_4$$

Evaluation

- **Environment**


- Head: 2xAMD 2.2GHz, 4GB DDR400 RAM, 800GB HD
- Slaves: 3 PCs (AMD 2.3G CPU, 2G DDR2-667 RAM, 400GB HD, 1Gbps Ethernet)
- GPU: XFX 9400GT 64bit 512MB DDR3
- CUDA 3.2 Toolkit
- Hadoop 0.20.3
- ServerTech CWG-CDU power distribution unit (for the power consumption monitoring)

- **Factors**

- Speedup
- Power consumption
- Cost

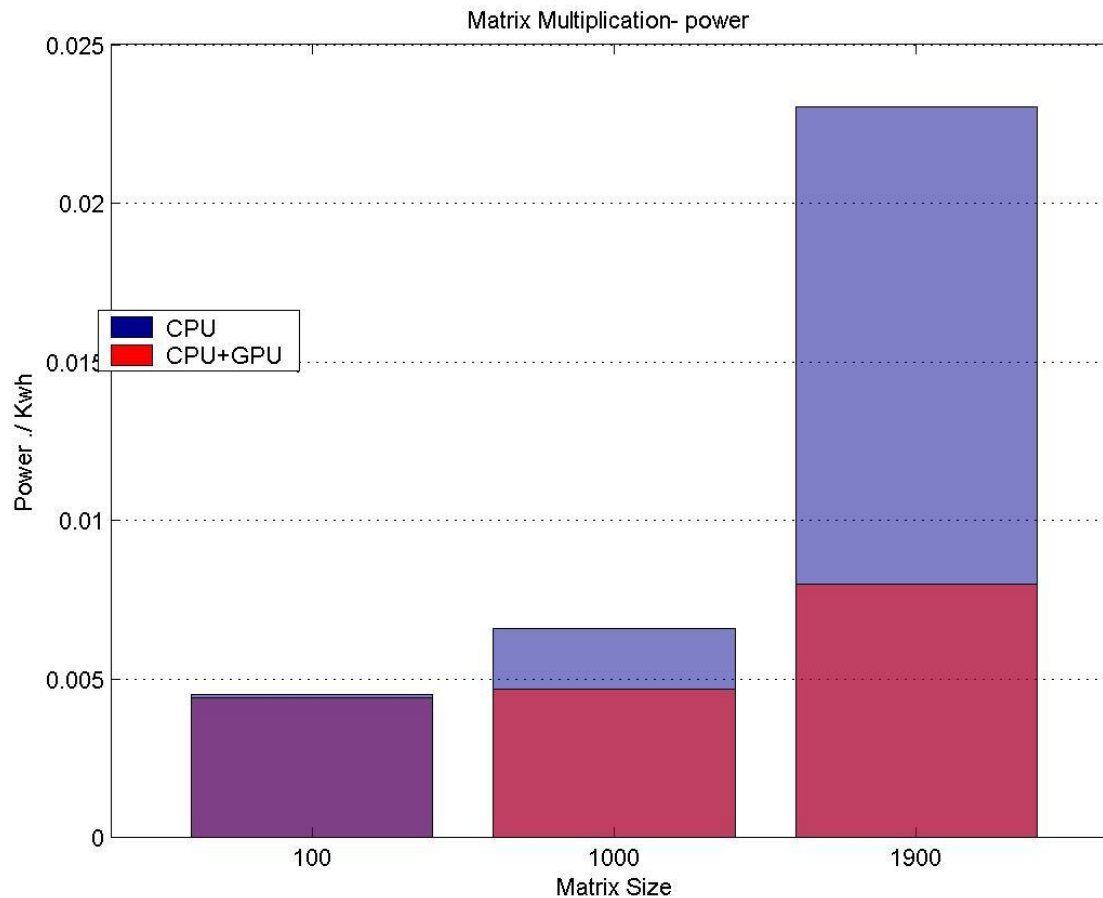
Evaluation

- Matrix Multiplication (Execution time)




Evaluation

- Matrix Multiplication (Power consumption)




Evaluation




Evaluation

MonteCarlo Power




Evaluation

- MDMR
 - Execution time



Evaluation

- MDMR
 - Power consumption



Conclusions

- Introduced GPU into MapReduce cluster and obtained up to **20 times** speedup.
- Reduced up to **19/20** power consumption with the current preliminary solution and work load.
- Compared with upgrading CPUs and adding more nodes, deploying GPU on Hadoop has **high cost-to-benefit ratio**.
- Provided **practical implementations** for people wanting to construct MapReduce clusters with GPUs.

Future Work

- Port more CUDA programs onto Hadoop.
- Incorporate reducers into the experiments
- Support heterogeneous clusters which mixed GPU-nodes and non-GPU nodes.

Reference

- nVIDIA CUDA
<http://developer.nvidia.com/object/cuda-3.2/>
- Hadoop, *<http://www.hadoop.com>*.
- J. Polo, D. Carrera, Y. Becerra, V. Beltran, J. Torres and E. Ayguadé *Performance Management of Accelerated MapReduce Workloads in Heterogeneous Clusters*, ICPP2010, (2010), 654-662.
- C. He, D. Swanson. *Molecular Dynamics simulation based on MapReduce*, poster section, LCI 2010, (2010).