

From DataFrames to Tungsten: A Peek into Spark's Future

Reynold Xin @rxin

Spark Summit, San Francisco

June 16th, 2015



DataFrame

noun

Making Spark accessible to everyone (data scientists, engineers, statisticians, ...)

Tungsten

noun

Making Spark faster & prepare for the next five years.

How do DataFrames and
Tungsten relate to each other?

Google Trends for “dataframe”

Single-node tabular data structure, with API for

relational algebra (filter, join, ...)

math and stats

input/output (CSV, JSON, ...)

ad infinitum



Data frame: lingua franca for “small data”

```
head(flights)
#> Source: local data frame [6 x 16]
#>
#>   year month day dep_time dep_delay arr_time arr_delay carrier tailnum
#> 1  2013     1   1     517         2     830         11      UA   N14228
#> 2  2013     1   1     533         4     850         20      UA   N24211
#> 3  2013     1   1     542         2     923         33      AA   N619AA
#> 4  2013     1   1     544        -1    1004        -18      B6   N804JB
#> ..   ...     ...   ...         ...         ...         ...     ...     ...
```

Spark DataFrame

Distributed data frame for Java, Python, R, Scala

Similar APIs as single-node tools (Pandas, dplyr), i.e. easy to learn

```
> head(filter(df, df$waiting < 50)) # an example in R
##   eruptions waiting
##1      1.750       47
##2      1.750       47
##3      1.867       48
```

Spark
DataFrame



Existing
Single-node
Data Frames



KB

MB

GB

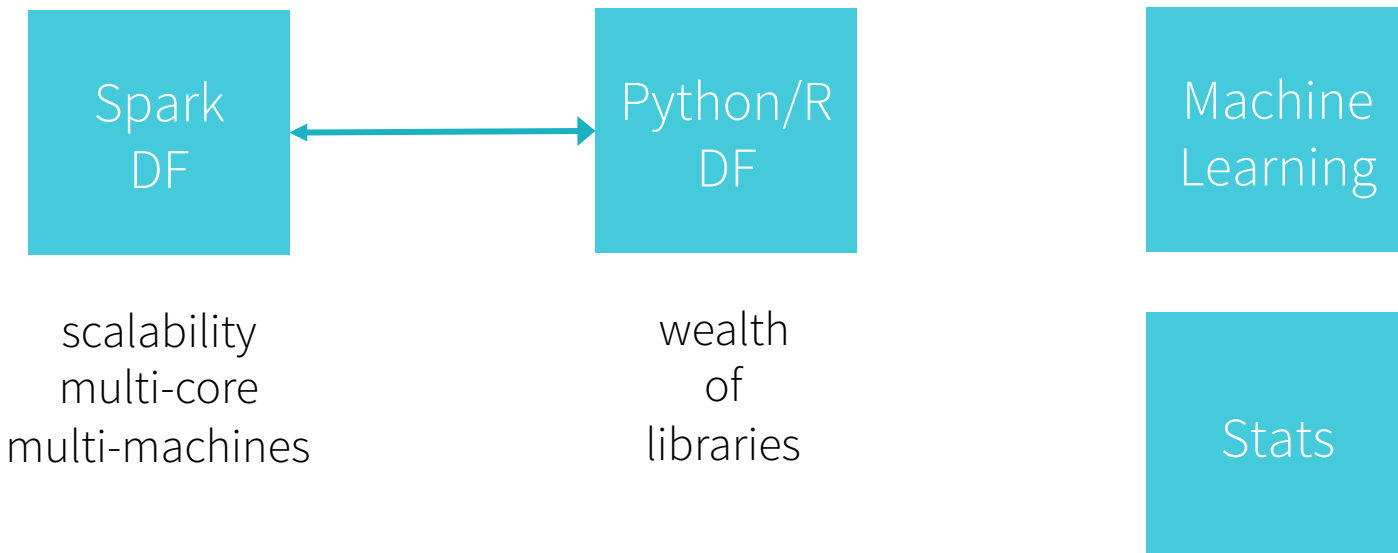
TB

PB

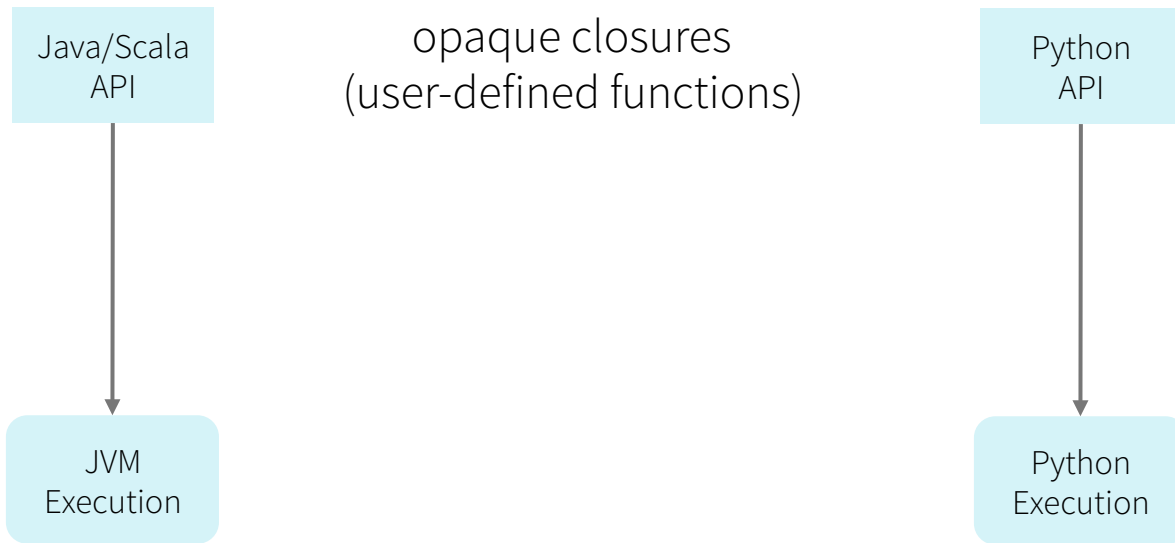
data size

It is not Spark vs Python/R,
but Spark and Python/R.

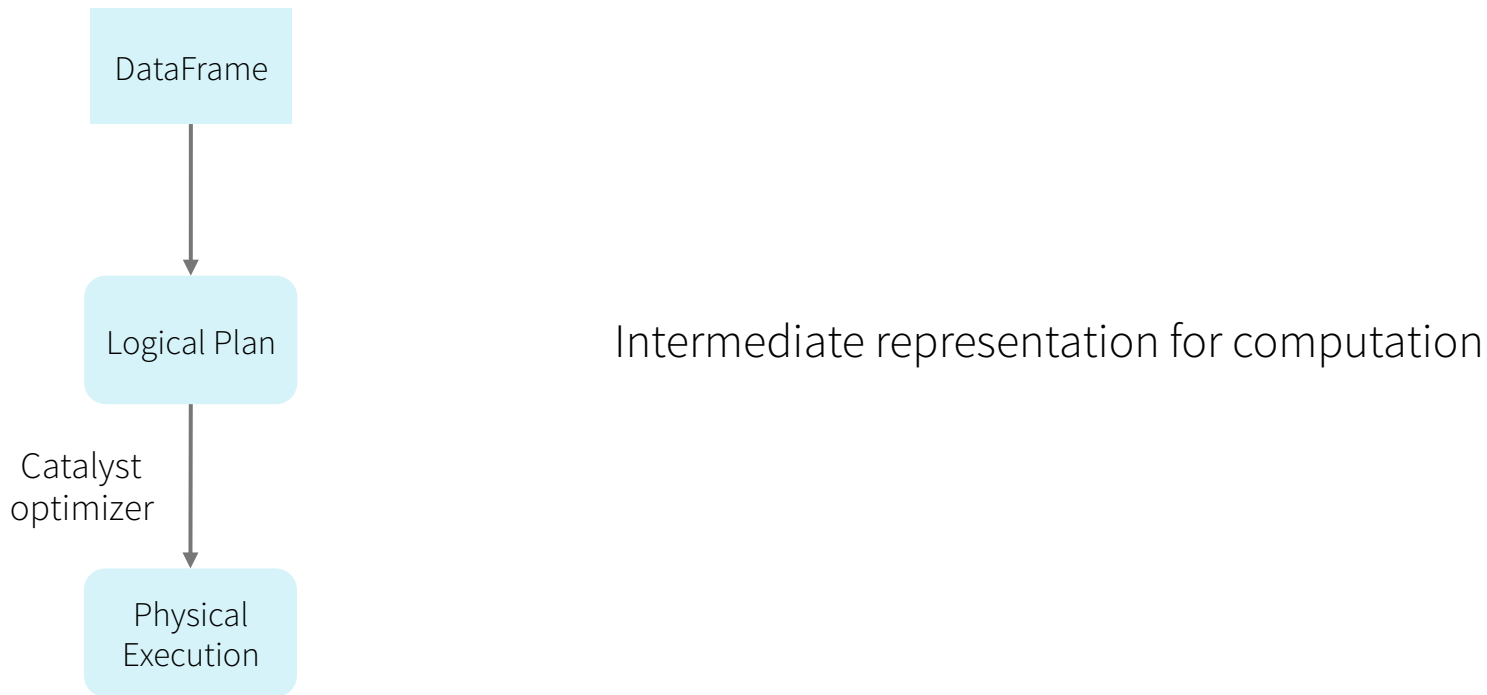
Spark *and* Python/R



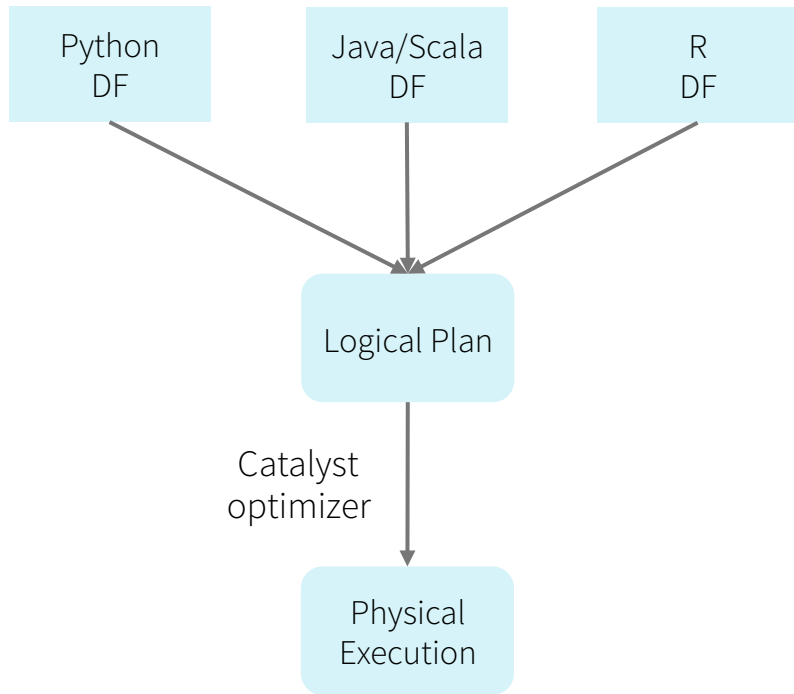
Spark RDD Execution



Spark DataFrame Execution



Spark DataFrame Execution



Simple wrappers to create logical plan

Intermediate representation for computation

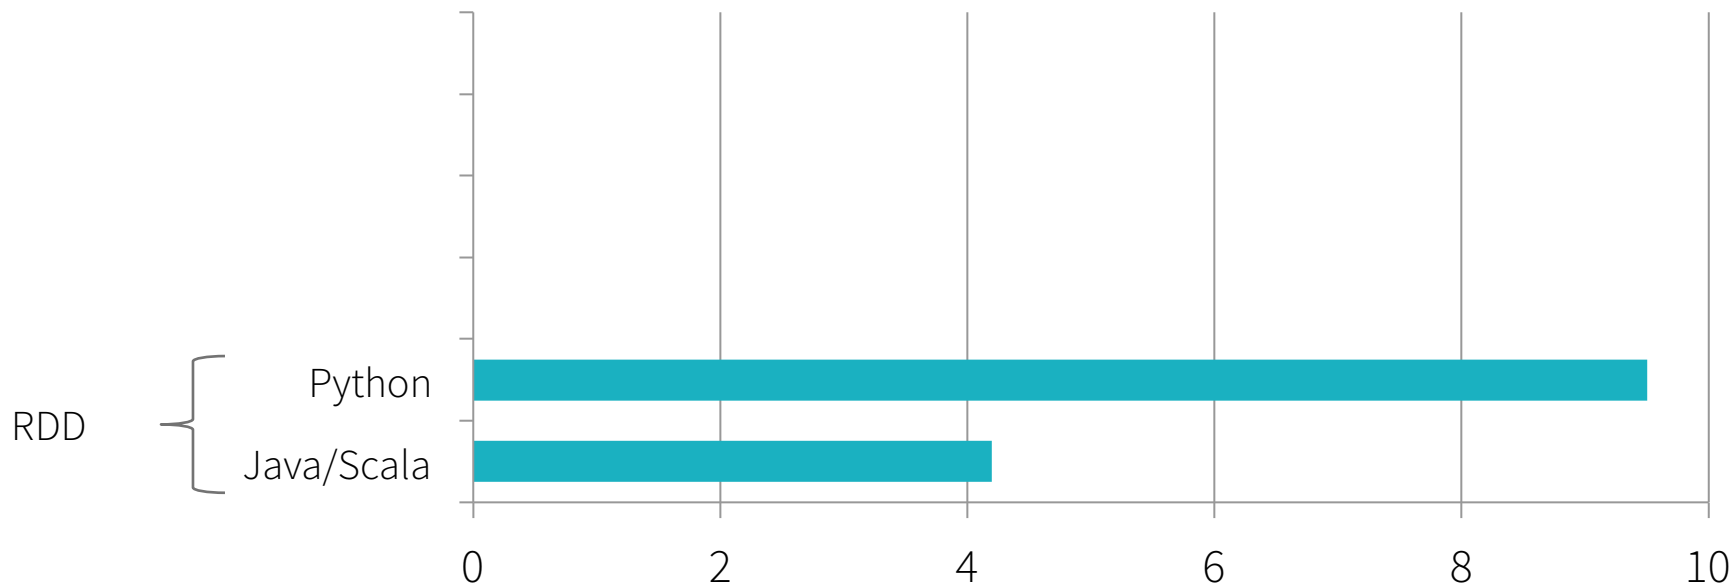
Benefit of Logical Plan: Simpler Frontend

Python : ~2000 line of code (built over a weekend)

R : ~1000 line of code

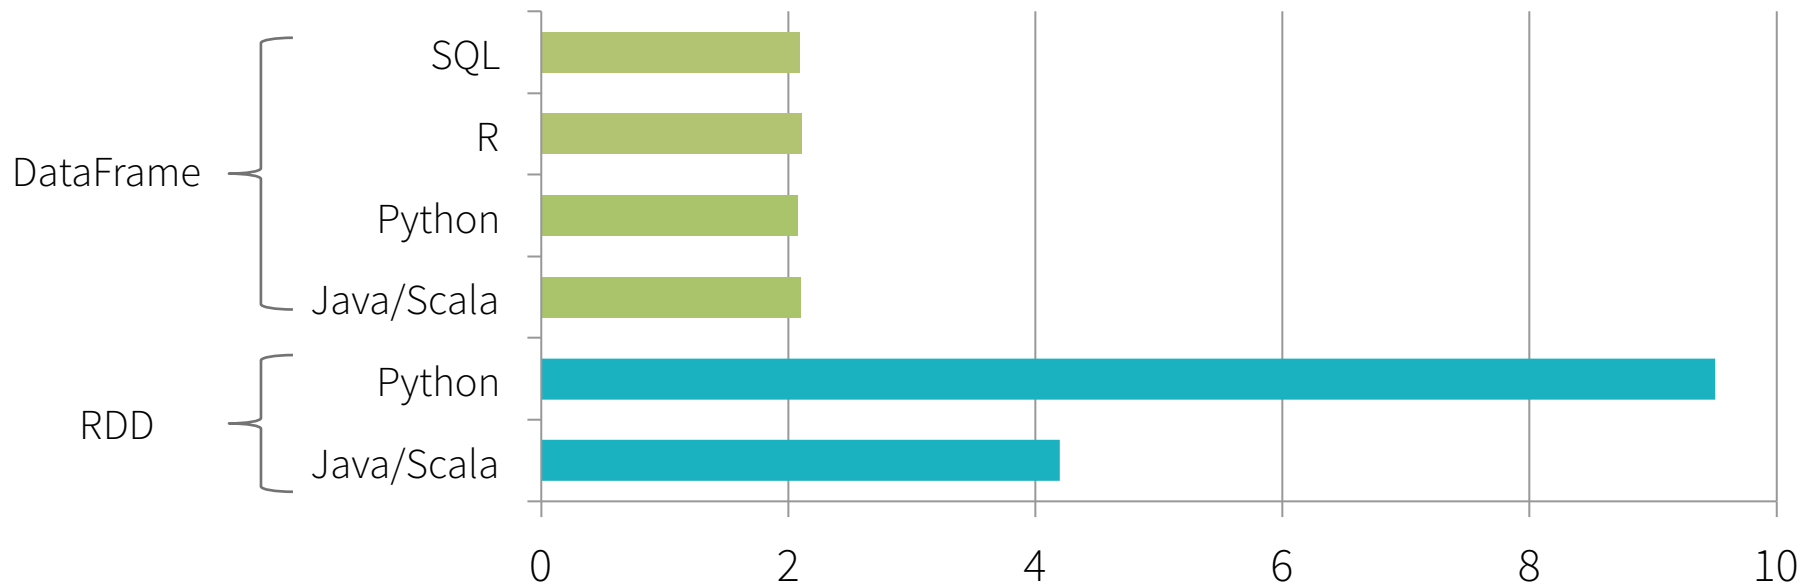
i.e. much easier to add new language bindings (Julia, Clojure, ...)

Performance



Runtime for an example aggregation workload

Benefit of Logical Plan: Performance Parity Across Languages



Runtime for an example aggregation workload (secs)

The background is a textured teal watercolor wash. It features darker, more saturated teal areas in the upper left and center, which blend into lighter, more translucent teal towards the bottom and right edges. The overall effect is organic and painterly.

What about Tungsten?

Hardware Trends

Storage

Network

CPU

Hardware Trends

2010

Storage 50+MB/s
(HDD)

Network 1Gbps

CPU ~3GHz

Hardware Trends

	2010	2015
Storage	50+MB/s (HDD)	500+MB/s (SSD)
Network	1Gbps	10Gbps
CPU	~3GHz	~3GHz

Hardware Trends

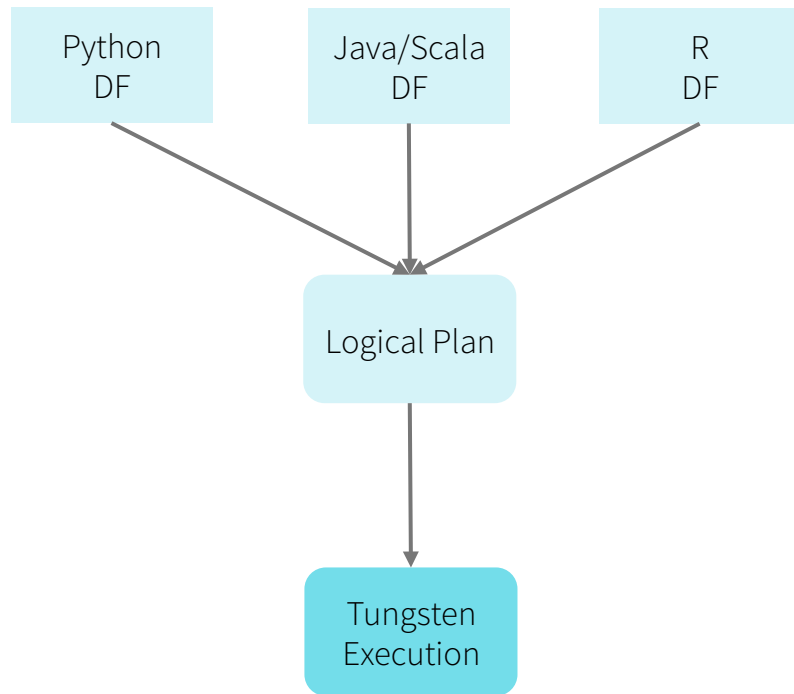
	2010	2015	
Storage	50+MB/s (HDD)	500+MB/s (SSD)	10X
Network	1Gbps	10Gbps	10X
CPU	~3GHz	~3GHz	☹️

Tungsten: Preparing Spark for Next 5 Years

Substantially speed up execution by optimizing CPU efficiency, via:

- (1) Runtime code generation
- (2) Exploiting cache locality
- (3) Off-heap memory management

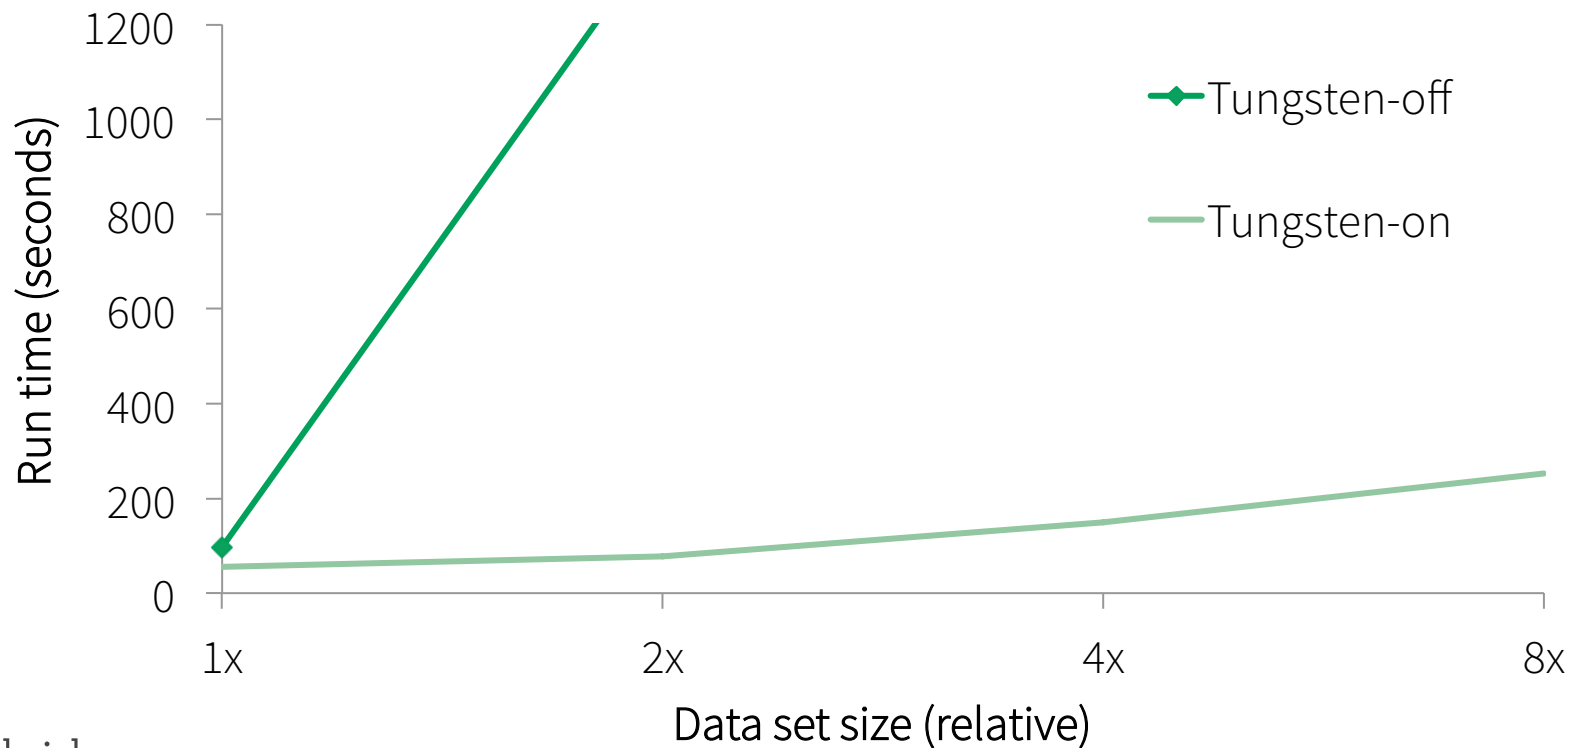
From DataFrame to Tungsten



5PM

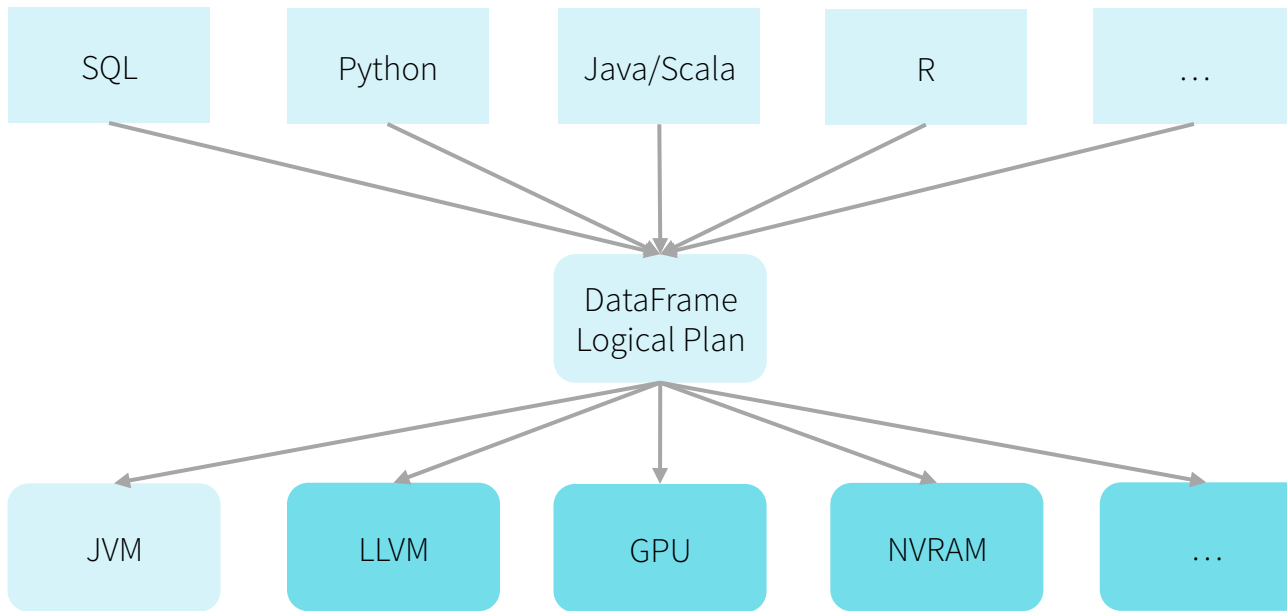
Deep Dive into Project Tungsten
Developer Track by Josh Rosen

Initial Performance Results



Unified API, One Engine, Automatically Optimized

language
frontend



Tungsten
backend

SQL

Python

R

Streaming

Advanced
Analytics

DataFrame

Tungsten Execution

Spark Office Hours Today

Topic Area	
1:00-1:45	Core, YARN, Ops
1:45-2:30	Core/SQL/Data Science
3:00-3:40	Streaming
3:40-4:15	Core, Python, R
4:30-5:15	Machine Learning
5:15-6:00	Matei Zaharia

Databricks booth A1