

Data Center Workload Characterization

Christina Delimitrou

Outline

- Introduction
- Related Work
- Implementation
- Proposal
- Future Work
- Suggestions

Outline

- Introduction
- Related Work
- Implementation
- Proposal
- Future Work
- Suggestions

Introduction


- No representative DC workloads publicly available
 - If available, only represent snapshot of the system
- We need to decouple the workload from the underlying system
 - What are the real characteristics of the workload?
 - How can we use it to improve performance/efficiency of the system?

Methodology	Platform Decoupled	System (hw/sw) evaluation	App characterization
Traces	No	(sort of)	No
Models	Yes	Yes	Yes

- Models offer more capabilities to **understand** and **use** a data center workload

Goal

Create a compressed but representative model
that captures the features of:

- Storage system
 - 3-tier system
 - A complete large scale application
- 

... without being as dependent to the underlying system as traces

Outline

- Introduction
- **Related Work**
- Implementation
- Proposal
- Future Work
- Suggestions

Related Work

- **Storage Characterization for Unstructured Data in Online Services Applications.** Sriram Sankar, Kushagra Vaid (Microsoft)
 - *Mantra:* Derive a **probabilistic model** for disk accesses in large scale applications
 - Start from traces for 3 popular 3-tier apps (hotmail, maps, user-client)
 - Collect stats on:
 - Number of I/Os (app load intensity-block size, type(rd/wr), randomness)
 - Spatial distribution of I/Os (disk blocks)
 - Inter-arrival times between I/Os
 - Outstanding I/O queues
 - Divide disk space in block ranges - assign one state per block range
 - Represent probabilities of changing states between I/Os w/ transitions

Outline

- Introduction
- Related Work
- **Implementation**
- Proposal
- Future Work
- Suggestions

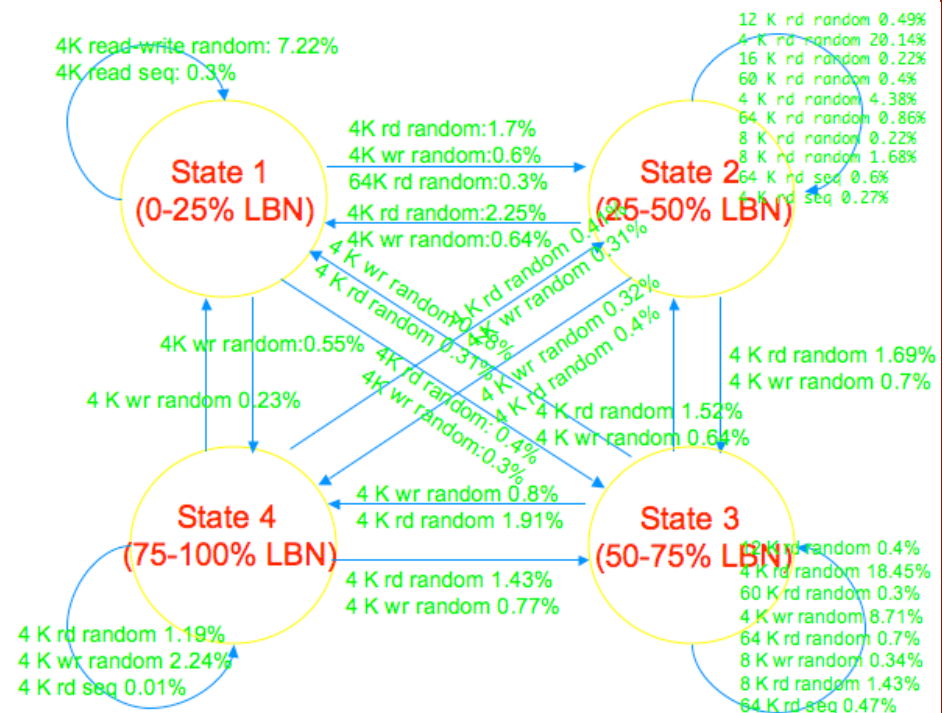
Impl : Storage Accesses via State Diagrams

- Model: probabilistic state transition diagrams

- **State:** Block range
- **Transition:** Probability of changing block range
- **Stats:** rd/wr, rd/seq, block size


- Single or Multiple Levels

- Coarse grained representation
- Hierarchical representation
- User Defined level of granularity



This state diagram from a multicore MS server that hosts **hotmail**

Impl : Storage Accesses via State Diagrams (contd)



1. Scripts to read the state diagram (of one of more levels) and define number of worker threads
2. Generator: IOmeter++
 - ✓ Modified open source storage workload generator
 - ✓ Can replay workload on one or more servers (user defined parameter)
 - ✓ Generates a workload that resembles the original in:
 - Probabilities of transitioning between states
 - Inter-arrival rates
 - Block sizes
 - Read/Write and Random/Sequential characteristics
 - ✓ Remains to validate on the response time of requests...

Outline

- Introduction
- Related Work
- Implementation
- **Proposal**
- Future Work
- Suggestions

Main Milestones

■ Workload modeling

- Use stats techniques to model behavior DC workloads
- Capture all tiers, whole system, interactions
 - Net, CPU, memory, storage, etc
- Discussion topics: tracing & modeling techniques, level of detail, target workloads, platform (in)dependence, metrics to capture, ...
- Will try to do that with help from REU/CURIS in the summer...

■ Workload generators

- Create generators that can replay workloads on other sites and other systems
- Share the generators and models with the community

Datacenter Workload Modeling

- First step: Setup some workloads locally
 - Analytics => Hadoop + (GridMix and PigMix)
 - Virtualized computation => Xen
 - SPECWeb 2009
 - Media streaming & mining => biocomp data
- Workload classes
 - Analytics (mapreduce, hadoop, ...)
 - HADI (large scale graph algorithm)
 - 3-tier (mail, maps, apps, ...)
 - Latency sensitive apps (search, facebook)
 - Virtualized computation (EC2) and storage (S3)
 - Streaming media (youtube, ...)
 - What else??

The Work Proposed

■ Traces - 2 - State Diagrams

- Create the probabilistic model from traces of real workloads (extract probabilities of states and transitions)
- Make the representation hierarchical and modular
 - Configurable sublevels of hierarchy
 - Inter-arrival time distributions
- Make the tool for creating the state diagrams publicly available (online library??)

■ State Diagrams - 2 - Synthetic Workloads

- Use the modified workload generator to read the state diagrams and create the synthetic workload
- Validate the accuracy of workloads against real Microsoft applications
- Use the synthetic workload to evaluate hardware/software options for efficient servers
- Make the workload generator publicly available

In Progress

- ✓ Modified Iometer to read the state diagram and create the synthetic workload (represents: states / probabilities / inter-arrival rates)
- ✓ Extended to create a workload from a hierarchical representation
- ✗ Waiting for the scripts that create the state diagrams
- ✗ Currently, validating the resemblance between original and synthetic workloads

Benefits & Limitations

- ✓ High portability
 - ✓ Not as “perishable” as traces - not as coupled to hardware
 - ✓ If validation succeeds :) they will offer a compressed version of a highly scalable application
 - ✓ Opens opportunities for detailed characterization (pattern analysis, learning techniques) for DC workloads
 - ✓ Can be used to evaluate and propose efficient solutions for DC design
-

- ✗ Some things are considered fixed:
 - ✗ App code
 - ✗ OS/software in general
 - ✗ If these change, the model is no longer representative of the application
- } Model assumptions

Outline

- Introduction
- Related Work
- Implementation
- Proposal
- **Future Work**
- Suggestions

The Fork...

■ Hybrid System Proposal

- Use the synthetic workload to evaluate different software - hardware options for efficiency and performance
- Would use of mobile computing components make sense?
 - Mobile RAM, Flash, SSD
- Would powering-off disks make sense?
- Propose a hybrid storage system that improves efficiency

■ Expand a similar methodology to other aspects of the system

- **CPU** Utilization
 - **Network** Traffic
 - **Memory** Utilization
 - What metrics would be interesting/useful to consider?
 - Again, validate with original workloads and propose system options for efficiency
 - Low power processors (Atom, Nvidia, ARM)
 - Network topologies
 - Memory technology
- Correlation???

Outline

- Introduction
- Related Work
- Implementation
- Proposal
- Future Work
- Suggestions

Suggestions

- Get workloads - traces through internships
 - Use the models to create our synthetic workloads
 - Use them to:
 - Evaluate systems - propose energy efficient hybrid systems for data center design
 - Expand this study to other aspects of the system
 - Is it efficient?
 - What overheads will it issue??
 - **Feedback needed**
 - What pattern analysis techniques??
 - Energy efficient memories
 - Network issues
 - Scale down issues
- } All of interest to EPIC

Suggestions

- What workloads are of interest to each team/would like to see synthetic versions of?
- In what projects could this be used?
- What is the time table for your workloads (initial stage, ready, finished?)
- Will it make a difference in the systems you are designing?
- What other tools/profiling/workload generation would be useful to you?

Questions??

Thank you