

SMART DATA FAST.™



VOLTDDB

FAST DATA – THE NEW BIG DATA

OVERVIEW

- Trends
- Fast vs Big
- Approaches
- Use Cases

DATA-FICATION OF LIFE

The 10 Trillion Device World

Computerworld, September 2015

"Smartness can be embedded everywhere," said Professor Sangiovanni-Vincentelli, EE/CS at University of California at Berkeley.

"The entire environment is going to be full of sensors of all kinds. Chemical sensors, cameras and microphones of all types and shapes. Sensors will check the quality of the air and temperatures. Microphones around your environment will listen to you giving commands."



Fast Data



Big Data

All data originates as fast data,
why wait to analyze and act on it?

FAST = ADVANTAGE



“Real-time” contextual offers
=
offer uptake rates 75%
data revenues by 15%.”

Source: Openet 2014 survey of 87 mobile operators




Perishable insights have exponentially more value than after-the-fact traditional historical analytics.

Fast (in motion)

Streaming Analytics:
*real time summary and
aggregation*

Transaction Processing:
*per-event decisions using
context + history*



Big (at rest)

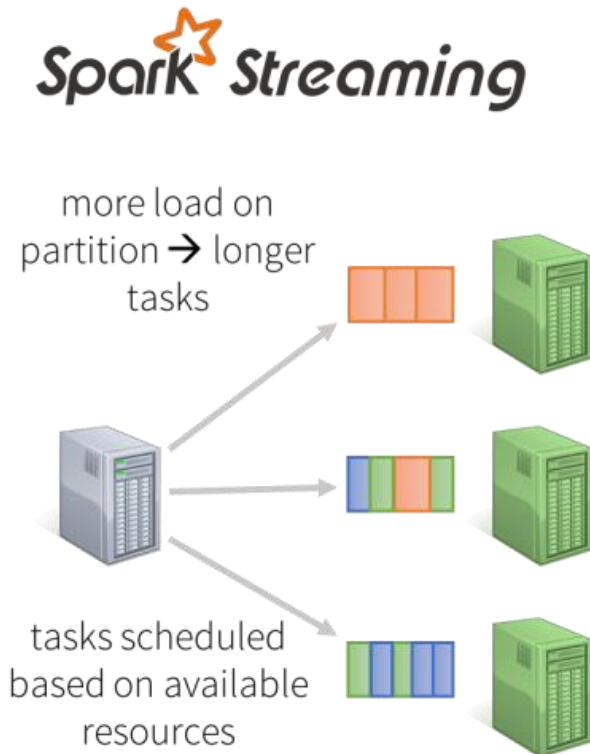
Exploration:
*data science, investigation of
large data sets*

Reporting:
*recommendation matrices,
search indexes, trend and BI*

APPROACHES

IN THE BEGINNING THERE WAS BATCH....


- Collect data, process it (used to be overnight), produce a report (output)
 - If batch job fails, delete the data, and start over
- Distributed systems made this better, more efficient
- Challenges
 - Response time (latency)
 - Processing events in order



NOSQL AND “EVENTUALLY CONSISTENT” SOLUTIONS

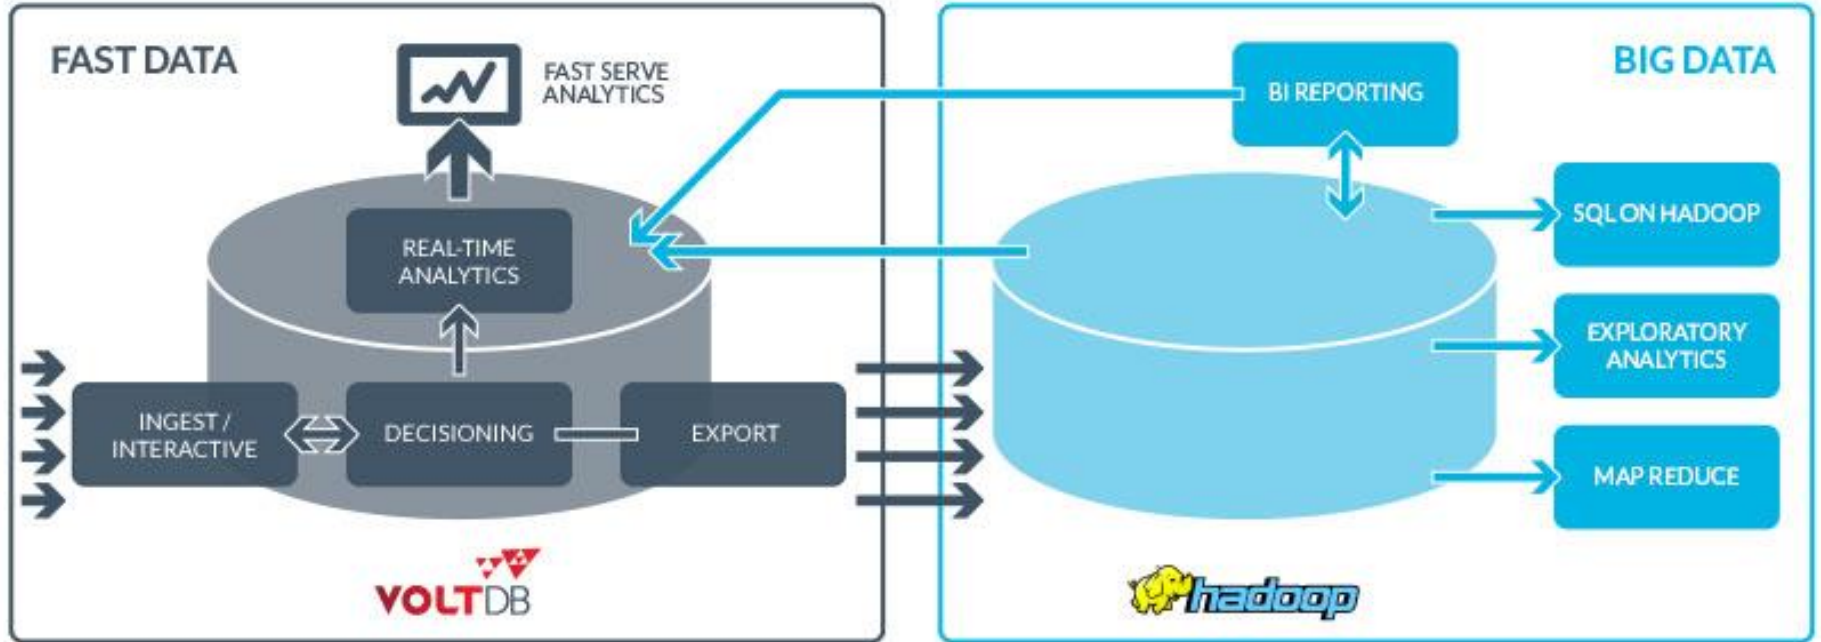
- Combine stream processing frameworks with NoSQL DBs
- Challenges
 - DiY requires building in reliability, code for ‘book keeping’ to ensure accuracy
 - Response time/latency goes up as components are added
 - Failure modes

Lambda Architecture





NEW ENTERPRISE ARCHITECTURE: FAST + BIG



ARCHITECTURE IS IMPORTANT....



Fast data requires
a different
architecture.

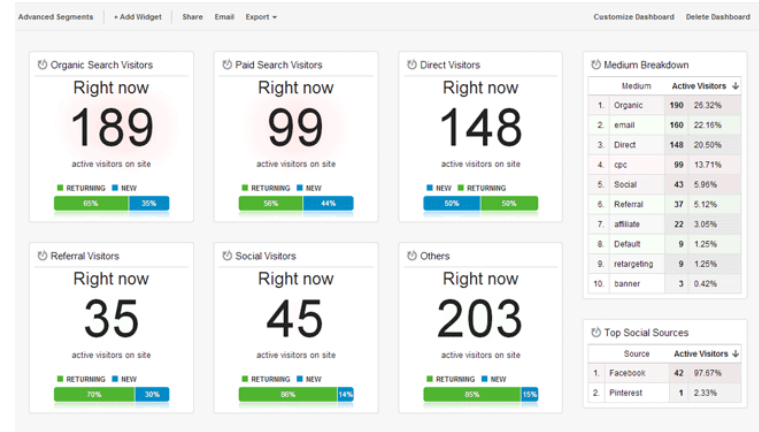
STREAMING ANALYTICS

What:

Filter, aggregate, enrich, and analyze a high throughput of data from live data sources


Why:

To identify patterns, detect urgent situations, and automate immediate actions in real-time



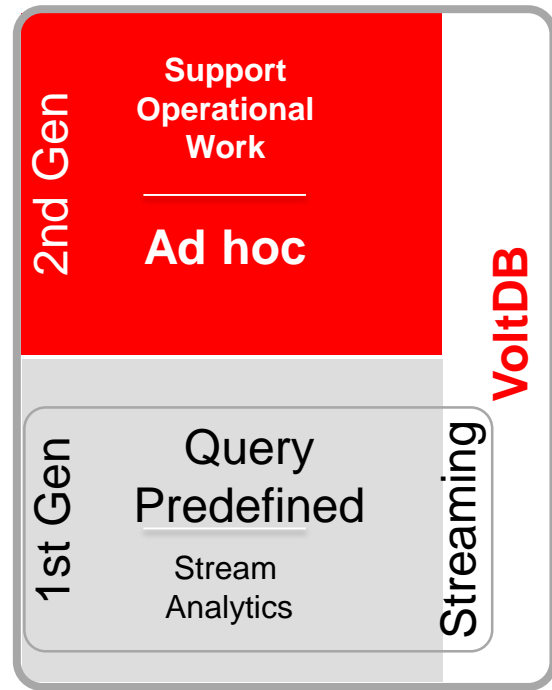
1ST GENERATION FAST DATA: STREAMING ANALYTICS

- Examples: Spark Streaming, Storm, Kinesis, TIBCO StreamBase, et al.
- Technical:
 - Lack “state” for transaction processing (operational)
 - Complex programming model
 - No ability to do ad hoc queries
- Functional:
 - 1st Gen only offers streaming analytics
 - Separate database required for any meaningful work
 - Proprietary interface is inconsistent with the rest of the data pipeline
 - Does not support applications requirement for interaction




2ND GENERATION FAST DATA: STREAMING ANALYTICS & OPERATIONAL WORK

- Streaming Analytics converges with the operational applications
 - Convergence is necessary to use data in real-time
 - Automated application interactions are informed by data
 - Brings the application into the “data analytics” world
- Streaming Analytics alone is *passive*, Fast Data is *interactive*



WHAT'S NEW HERE?



Combining streaming analytics and transactions allows you to act at the rate that you learn.



“By definition the only way to do streaming analytics is to do it **in-memory**. Don’t make the mistake of thinking that streaming is just about ingestion. Streaming analytics is about *analytics* more than it is about *ingestion*.”

“Spark Streaming is micro batch processing. That’s still batch processing but it does it in micro batches. I don’t consider that a true real-time streaming platform because it’s geared more for batch processing.”

A new category of databases is emerging we call **translytical databases**: streaming analytics with transactions in a single database.

FAST DATA REQUIRES ANALYTICS WITH (TRANS)ACTIONS

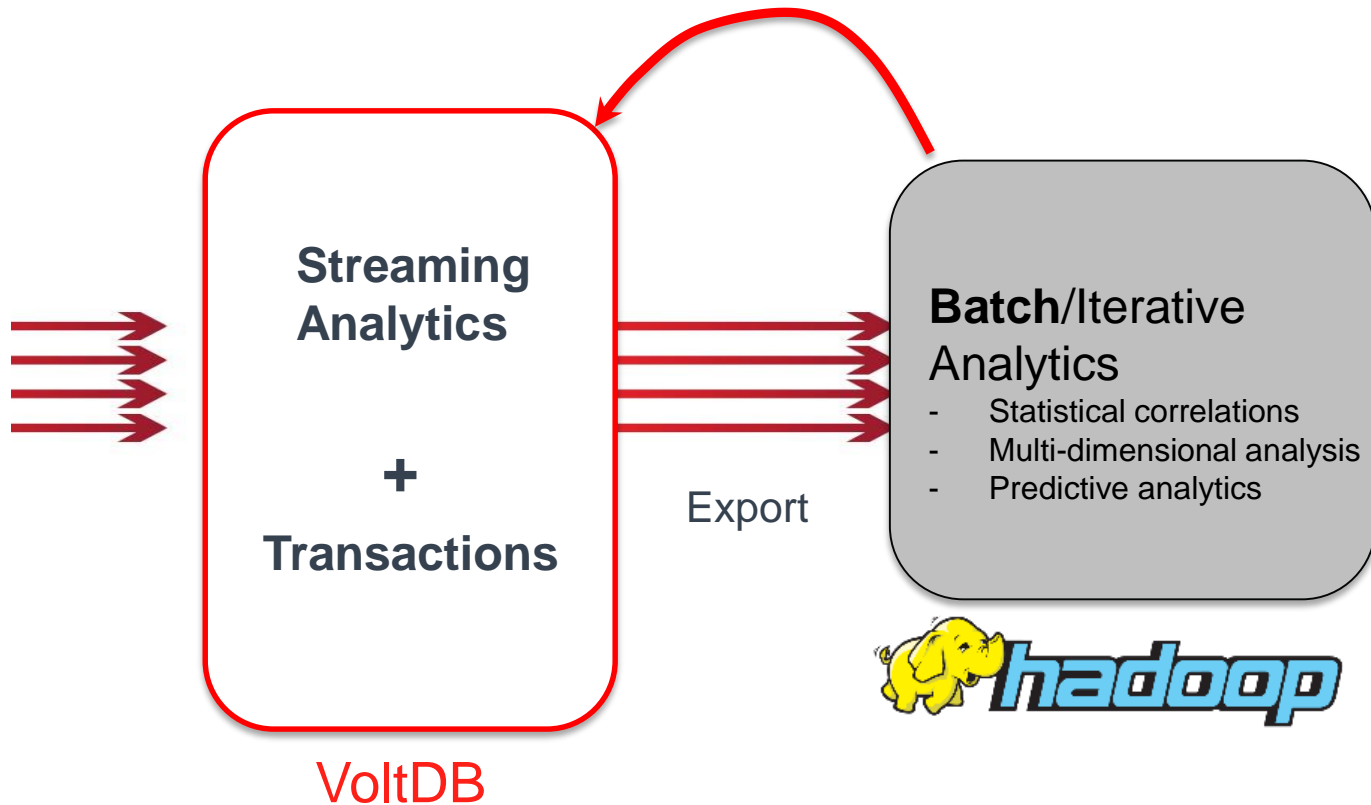
Customer-Facing

- Personalization
- Customer experience

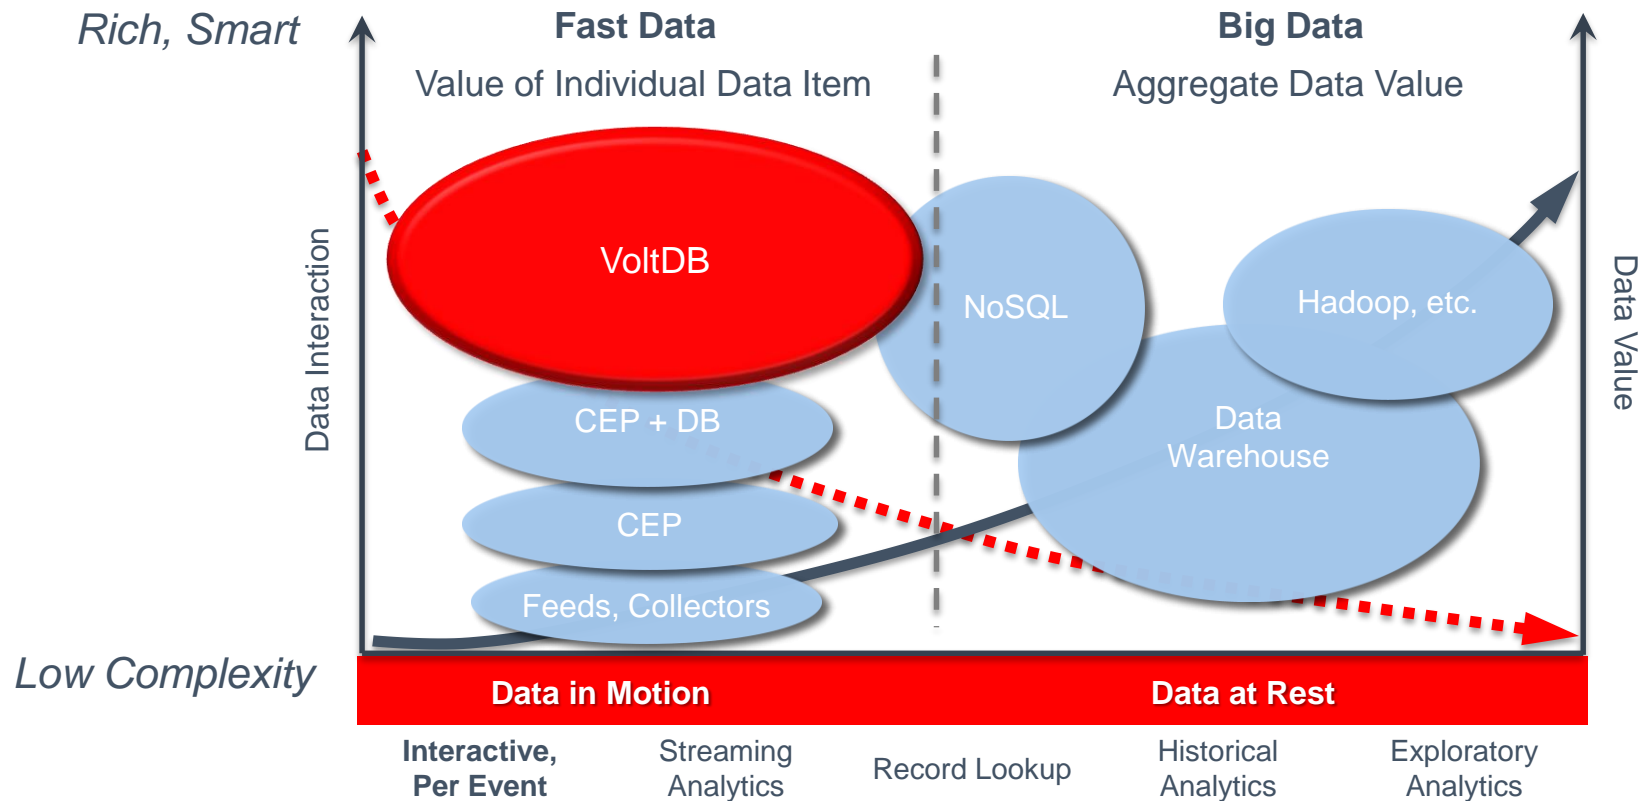


Operations-Facing

- Network optimization
- API monitoring
- Sensors



THE TIME VALUE OF DATA



VOLTDDB: A SUPERIOR ARCHITECTURE FOR FAST DATA

- ✓ In-Memory performance
- ✓ Scale-out, shared nothing
- ✓ ACID & SQL & Java
- ✓ Continuous, per event
- ✓ Reliability and fault tolerance
- ✓ Hadoop ecosystem integration



VoltDB is really different than everything else

THE SO WHAT

VoltDB allows companies to act on data in real-time, enabling new levels of application functionality and performance that drive new revenue streams while reducing infrastructure costs

USE CASES

USE CASE EXAMPLES: ANALYTICS + (TRANS)ACTIONS

	Streaming Analytics (Stream Proc. or OLTP)	(Trans)Actions (OLTP)
Mobile Usage	Count current usage minutes	Will current usage plus previous balance cause the customer to exceed his quota?
Gaming	Real-time stats on player effectiveness	Change game interaction to increase engagement of the player
Real-time Risk	Determine position values as prices and positions change	Does a new trade violate the defined risk tolerance? If “no,” place trade
Ad placement	With which segment is this user identified	Identify ad, check vendor quota balance, determine best network and place ad
Content Delivery Service	Count content views	Update log records in real time for accurate billing based on content views

USE CASES

Telco

- Subscriber Management
- Session Management
- OSS/BSS – policy, billing, routing
- SLA Management

Financial Services

- Risk Management (portfolio, trading)
- Fraud Detection
- Compliance (BB&O)
- Customer Engagement

Media and Entertainment

- Personalization
- Digital Advertising
- Content Delivery
- Gaming

IoT/Sensors

- Smart Energy
- Connected Home
- Patient Monitoring

SIMPLIFYING THE LAMBDA ARCHITECTURE



Content delivery network service provider

	Trident with Cassandra or HBase	VoltDB
Number of Environments to Manage	At least 3	1
Atomicity	Single-write	Multi-write
Unit of Atomicity	A single row	A single partition
Indexed Look-Up Requirements Per Micro-Batch	150,000	336
Transaction ID Space Requirements Per Micro-Batch	18 GBs	0.000012 GB

Implementing VoltDB and micro-batching with multi-write atomicity within the Lambda framework simplified management and improved performance, storage, scalability, and operations.

Use Case

- Counting “content” views in real time for billing and reporting

Why VoltDB?

- Real-time analytics + transactions w/scale
- Need for accuracy – chose VoltDB over Trident/Storm+Cassandra combination for real-time streaming aggregations with “exactly once” semantic



MaxCDN uses 1/10th
compute resources of
alternate solutions.

Behzad Pirvali
Performance Architect



- Mobile advertising service
- Managing over 150,000 applications

HYPERTARGET

Real-Time targeting = $f(\text{persona, interests, behaviors})$




Requirement:
Hundreds of
thousands of
concurrent
connections with
round-trip
latencies in
milliseconds



Before (MySQL)

100 servers



After (VoltDB)

7 servers






“Achieved a previously impossible level of budget management accuracy”

Dan Khasis
Chief Technology Officer

APPLICATIONS BUILT WITH VOLTDDB ARE:

- ✓ **Faster, more performant**
 - tps, latency
- ✓ **Simpler**
 - Fraction of components and coding vs. alternatives
 - Lower maintenance and support
- ✓ **Better**
 - Lower system risk
 - Correct results
 - Higher availability and reliability

WHY VOLTDB? Our customers realize exceptional business value



QUESTIONS?

- Use the chat window to type in your questions
- Try VoltDB yourself:
 - Free trial of the Enterprise Edition:
 - www.voltodb.com/Download
 - Open source version is available on github.com
- Use the chat tab to ask your questions.
- Join the conversation on Twitter #VoltDBFastData
- Download our latest report from O'Reilly in the resources window

