

Statistics Driven Workload Modeling for the Cloud

Archana Ganapathi, Yanpei Chen
Armando Fox, Randy Katz, David Patterson

SMDB 2010



Data analytics are moving to the cloud

Cloud computing → economy of scale

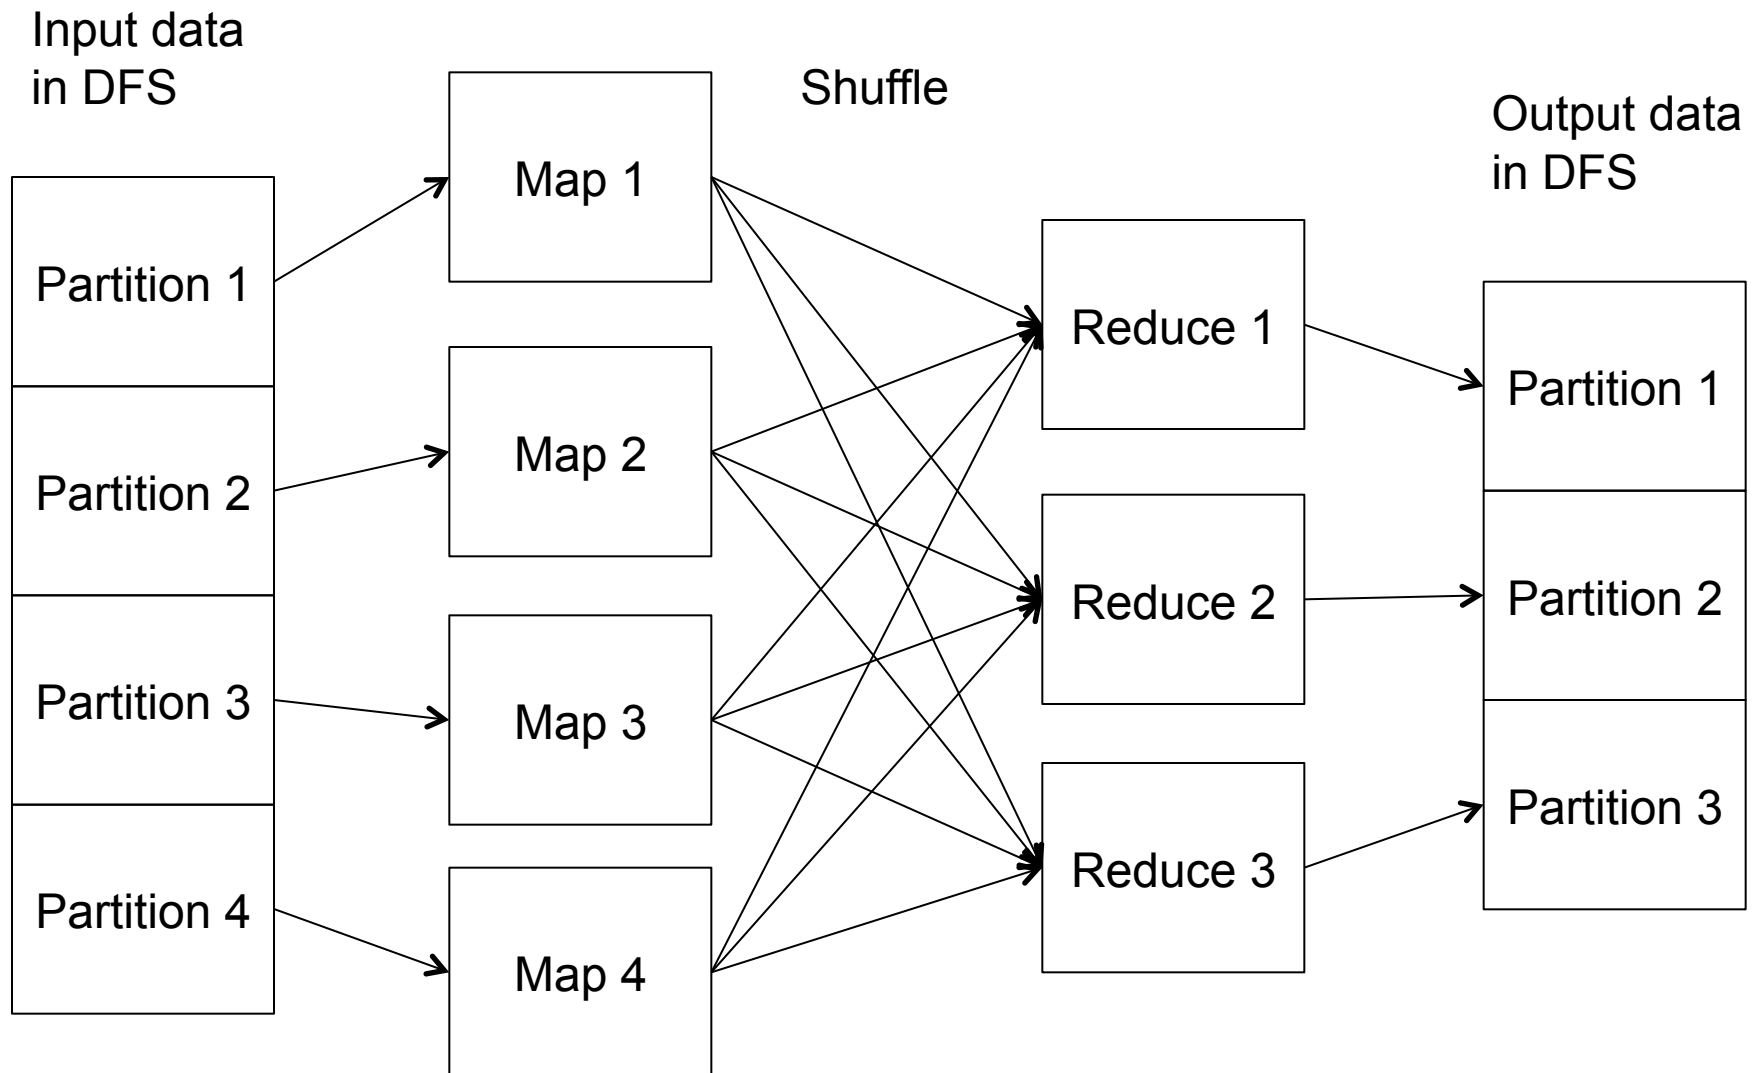
MapReduce for warehousing/analytics in the cloud

New challenges:

- Heterogeneous HW/SW/configuration/infrastructure
- Large variation in workload/software
- Ability to change resource consumption elastically

We present a partial solution

Brief MapReduce overview





Scheduling? Design trade-offs? Plan for the future?

Existing approaches:

- System simulation
- Benchmarks
- Trace replay
- Hardware/VM statistics

Our approach:

1. Predict performance on a fixed configuration
2. Workload synthesis & evaluation across configurations



Performance prediction via SML

Desirable properties:

- ✓ Predict job execution time and resource requirements in a single model
- ✓ Works equally well for SQL-like and traditional MapReduce jobs
- ✓ Generalizable across different hardware, software, configurations, etc.

We've done this before ...

A. Ganapathi, H. Kuno, U. Dayal , J. Wiener, A. Fox , M. Jordan , D. Patterson,
**Predicting Multiple Performance Metrics for Queries: Better Decisions
Enabled by Machine Learning**, ICDE 2009.

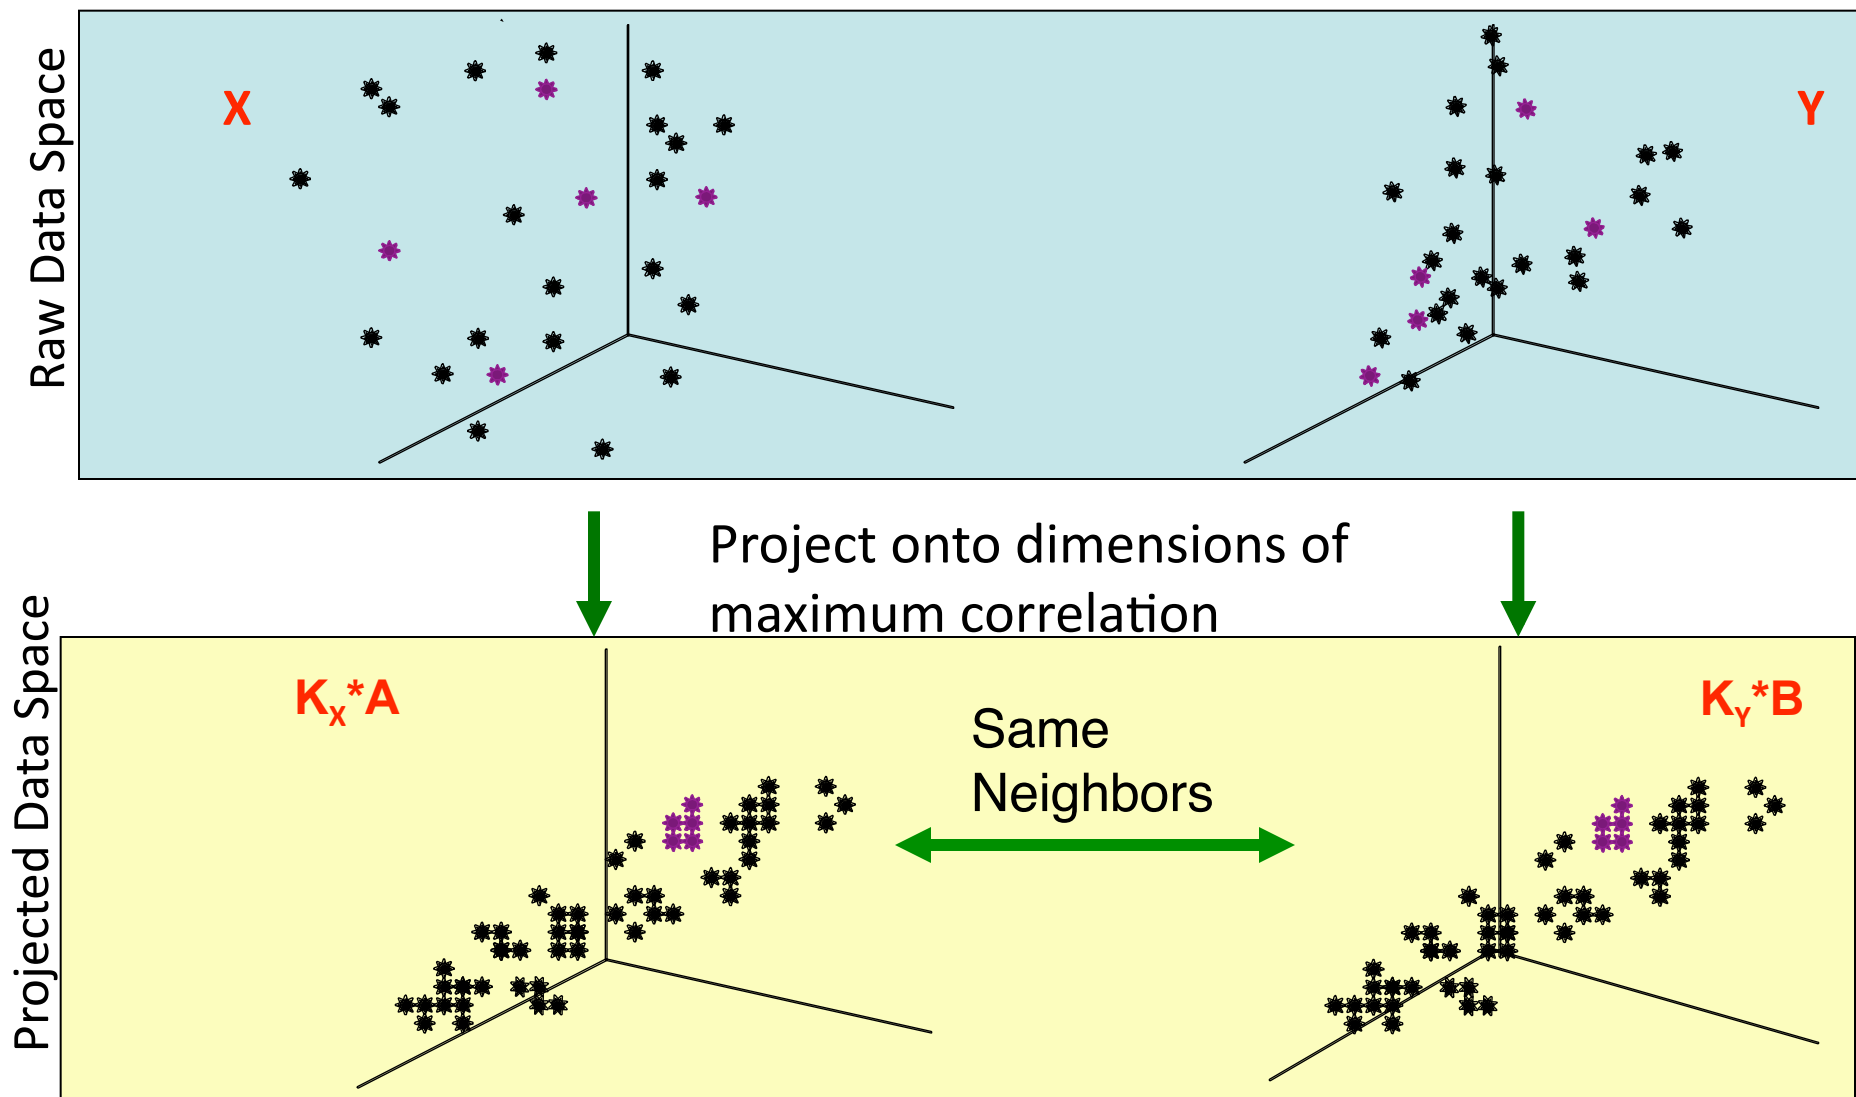
Silver bullet:

Kernel Canonical Correlation Analysis (KCCA)

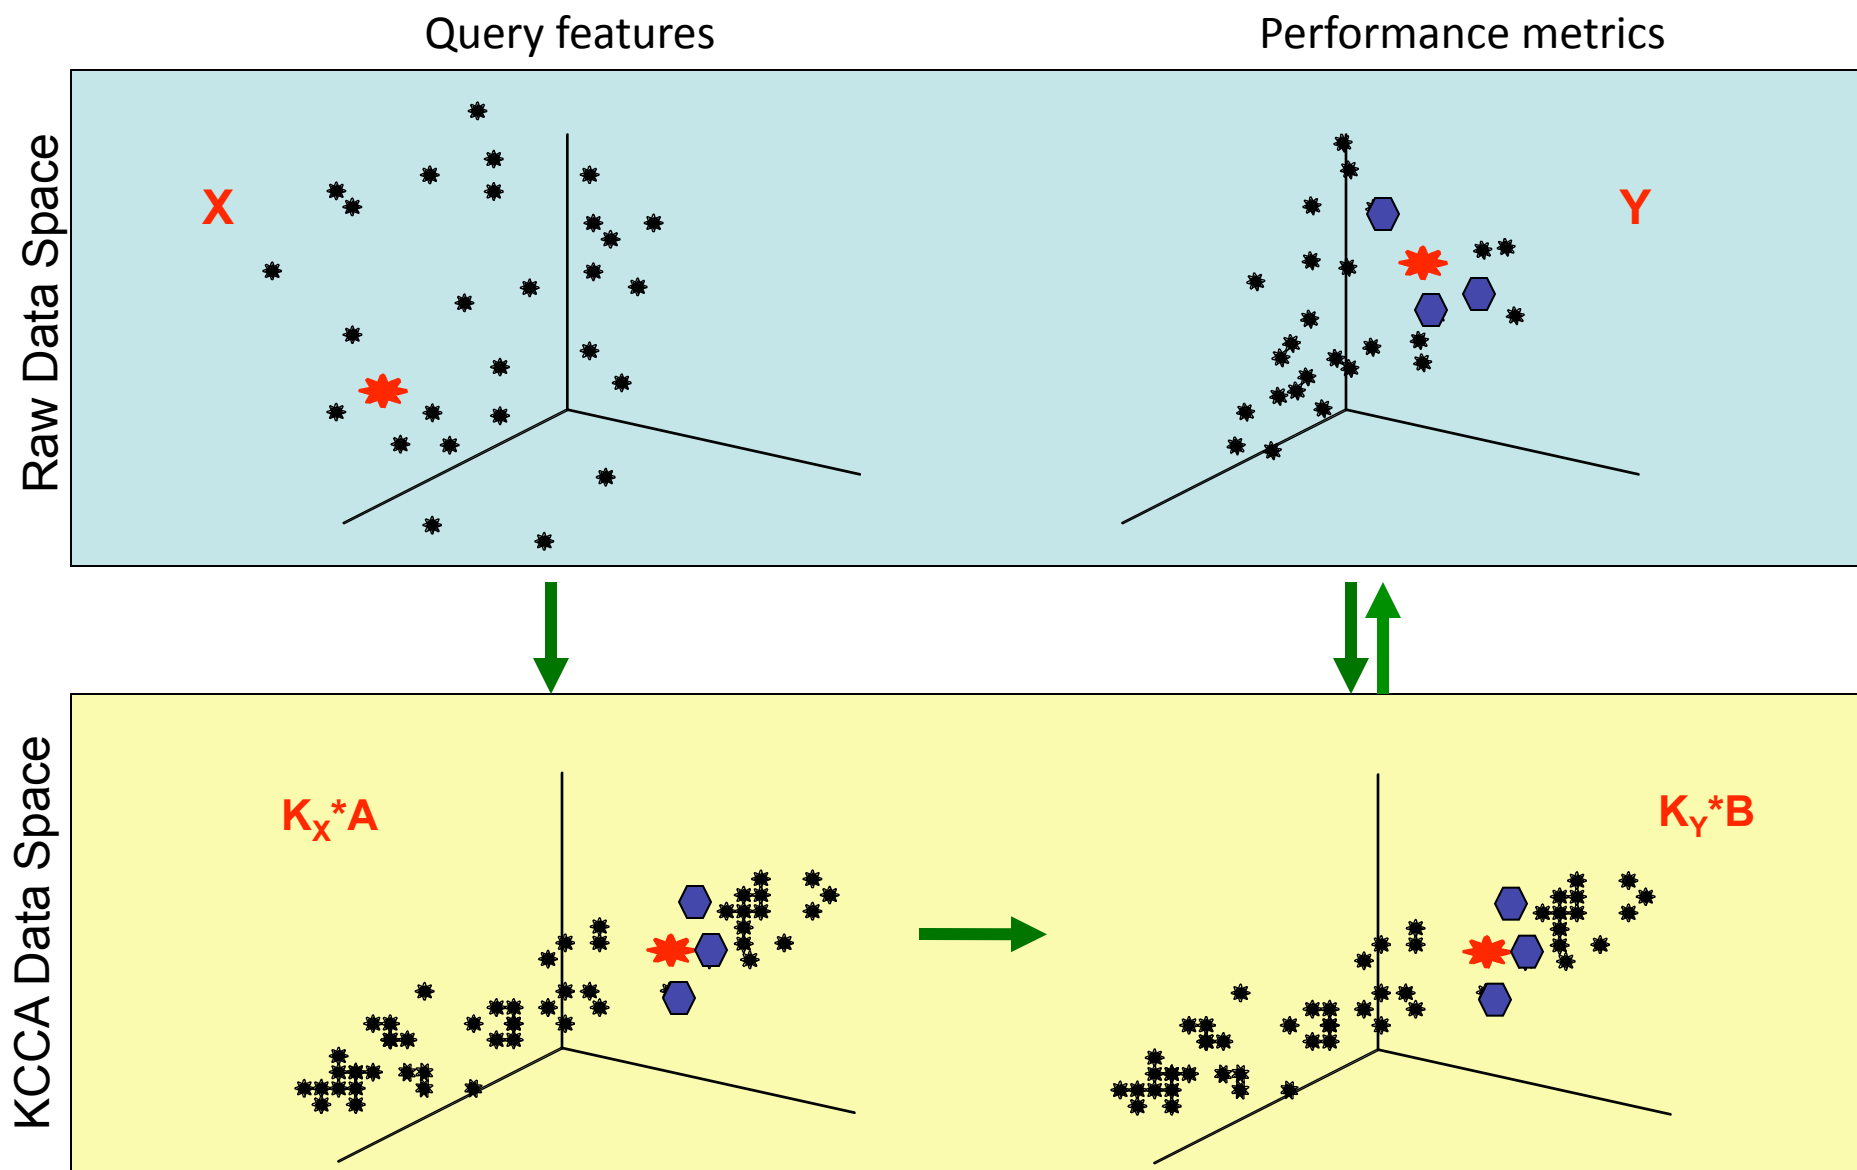
High level description of KCCA

Query features

Performance metrics



Prediction using KCCA model





Our particular experiment

Hadoop & Hive: open-source data analytics SQL interface

Production Hadoop Deployment at well-known social network

Multi-user environment

Nodes 1-300: 16 GB memory, 5 map slots, 5 reduce slots

Nodes 301-600: 8 GB memory, 5 map slots, 0 reduce slots

Training Set of 5000 Hive queries

Test Set of 1000 *unseen* Hive queries



SQL features predict poorly


Query plan features:

Count of

- Create Table
- Filter
- Forward
- Group By
- Join
- Move
- Reduce Output

System behavior metrics:

- Total Execution Time
- Map Time
- Reduce Time
- Map Output Bytes
- Local/HDFS Bytes Written





MapReduce features predict very well

Query plan features:

Job config parameters


- Number of Maps
- Number of Reduces

Data characteristics

- Map Input Bytes
- Local/HDFS Bytes Read

System behavior metrics:

- Total Execution Time
- Map Time
- Reduce Time
- Map Output Bytes
- Local/HDFS Bytes Written





Generalize to other MapReduce jobs

Query plan features:

Job config parameters

Number of Maps

Number of Reduces

Data characteristics

Map Input Bytes

Local/HDFS Bytes Read

System behavior metrics:

Total Execution Time


Map Time

Reduce Time

Map Output Bytes

Local/HDFS Bytes Written

E.g. Prediction for Extract-Transform-Load (ETL) jobs:





How do we use this tool?

- Given a particular setup, KCCA can predict performance
 - What if set up changes?
- Workload generator to ask “what-if” questions
 - Ideally, generalizes across hw/sw/config.
- KCCA identifies workload features that affect performance
 - Genesis for the MapReduce workload generator



Workload Generator Design

1. Collect statistics from real traces for these jobs features:
 - Inter-job arrival time
 - Per-job input data size
 - Per-job input:shuffle:output data ratio
2. Create synthetic workload that mimics real job stream
 - Compute approximate distributions
 - Sample from distribution to construct workload
3. Replay with low overhead



Workload Generator Design

1. Collect statistics from real traces for these jobs features:

- Inter-job arrival time
- Per-job input data size
- Per-job input:shuffle:output data ratio

2. Create synthetic workload that mimics real job stream

- Compute approximate distributions
- Sample from distribution to construct workload

3. Replay with low overhead

Design trade-offs:

- Why not direct replay? Burdened with design defects in original system
- Why ignore locality & data skew? Test varied placement/data schemes
- Why ignore compute? IO is usually the bottleneck, confidentiality issues



We have been using our prototype!

MapReduce energy efficiency


MapReduce schedulers

Multi-tenant resource provisioning and capacity planning



Preliminary look at latest results

Comparisons between different energy efficiency mechanisms



Most of the improvement comes from batching

Reverse existing design priorities – idle energy > active energy

Are all traces like this?



Useful beyond MapReduce

Same approach works in systems other than MapReduce

- Bootstrap model with production traces
- Predictive framework to identify workload characteristics
- Workload description focused on application semantics
- Workload synthesis by directly sampling empirical traces

Only the workload and performance feature vectors are specific to MapReduce

Ongoing work: network storage, wind power, etc.

Lessons learned

- Performance predictions need to be multi-dimensional
- Effective prediction features focus on application semantics
- Workload characterization should also focus on application semantics
- Workload synthesis by sampling production traces is effective

Questions?

archanag@eecs.berkeley.edu

yichen2@eecs.berkeley.edu
