

Gentlest Intro to Tensorflow

Khor, @neth_6, re:Culture

In collaboration with (Sam & Edmund):



Goal

- Perform **linear regression** in its most basic form through simplest example
 - Predict house price with just single feature, i.e., house size

What We Will Learn Together

- 4 steps of Machine Learning:
 - Create model
 - Define cost function
 - Collect data
 - Train
- 2 parts of Tensorflow:
 - Model linear regression & cost function as Tensorflow graph
 - Train Tensorflow graph with dataset

Machine Learning

Machine Learning Purpose

Predict something given a set of **features** that influences prediction:

- Predict house price given its size, location, configuration (features)

Machine Learning Training

Before predicting, we have to:

- Choose/Create a **model**
- **Train** the model to learn prediction with **data**

Problem

Predict House Price Given House Size

y = house price

x = house size (feature)


Solve:

- Given a house size, what should be the house price?

Model

What is the House Price?

House Price, \$



House Size, sqm

What is the House Price?

House Price, \$

Without data you cannot predict!

House Size, sqm

What is the House Price?

House Price, \$

We can answer house price for 10, 20, 30, 50, 60 sqm

10

20

30


50

60


House Size, sqm




What is the House Price?




What is the House Price?



What is the House Price?




What is the House Price?




Cost Function: Best-fit Prediction

Minimize difference between predicted & actual values


What is the House Price?



What is the House Price?



What is the House Price?



Cost function: How different is your model from reality

$$\text{sum}((y_ - y)^{**2})$$

Where:

- $y_:$ Actual value
- $y:$ Predicted value
- $**2:$ Power of 2

Linear Regression

Linear Regression

$$\mathbf{y} = \mathbf{W}.\mathbf{x} + \mathbf{b}$$

y: Predicted house price


x: House size from dataset

Find a good W, b


W: Gradient (Steepness)

$$W1 < W2 < W3$$


$$y = W1.x + b$$



$$y = W2.x + b$$




$$y = W3.x + b$$




b: Intersect

$$b1 < b2 < b3$$


$$y = W.x + b1$$



$$y = W.x + b2$$



$$y = W.x + b3$$



Train

Train: Find W , b

$W_best, b_best = 0, 0$

Loop J times:

Choose $current_W, current_b$

If $cost(current_W, current_b) < cost(W_best, b_best)$:

$W_best, b_best = current_W, current_b$

Train: Find W , b

$W_best, b_best = 0, 0$

Loop J times:


Choose $current_W, current_b$

What is the best way to choose
 $current_W, current_b$?

If $cost(current_W, current_b) < cost(W_best, b_best)$:

$W_best, b_best = current_W, current_b$

Finding Minimum Cost Function



How to Reach Minimum?



Choosing Training Method

- Random
 - Could take forever
- Gradient descent
 - From current viewpoint, move towards direction where steepness is greatest

Ready?

Recap:

- Model: Linear regression, $y = W.x + b$
- Cost: Least squared, $\text{cost} = \text{sum}((y_ - y)^{**2})$
- Train: Gradient descent

Step 1: Model $y = W.x + b$

Tensorflow Model

tf.placeholder

- Hold data to be feed into the model
- `x = tf.placeholder(tf.float32, [None, 1])`

tf.Variable

One output, house price

- Hold variables to be trained
- `W = tf.Variable(tf.zeros([1, 1]))`
- `b = tf.Variable(tf.zeros([1]))`

One feature, house size

The diagram consists of three red arrows pointing from the right side of the code snippets to a central point. One arrow originates from the red '1' in the placeholder definition `x = tf.placeholder(tf.float32, [None, 1])`. Another arrow originates from the red '1' in the weight variable definition `W = tf.Variable(tf.zeros([1, 1]))`. The third arrow originates from the red '1' in the bias variable definition `b = tf.Variable(tf.zeros([1]))`. These three arrows converge towards the right, pointing towards the text 'One feature, house size'. Additionally, a blue arrow points from the blue '1' in the weight variable definition `W = tf.Variable(tf.zeros([1, 1]))` towards the text 'One output, house price'.

Tensorflow Model

`tf.matmul(x,W)`

- Multiply 2 matrices
- `product = tf.matmul(x,W)`

`‘+’`

- `y = product + b`
- Expands to:
 - `y = tf.matmul(x,W) + b = W.x + b` **Done!**

Step 2: Cost Function $\text{sum}((y_ - y)^{**} 2)$

Best-fit: Minimize Difference Prediction and Actual

Actual values:

One output, house price

- `y_ = tf.placeholder(tf.float32, [None, 1])`

Minimize difference between actual and prediction:

- `cost = tf.reduce_sum(tf.pow((y_-y), 2))`

Step 3: Data

Faking Data

```
for i in range(100):  
    // Create fake data for actual data  
    xs = np.array([[i]])  
    ys = np.array([[2*i]])
```

- $xs = 0, 1, 2, \dots, 99$
- $ys = 0, 2, 4, \dots, 198$
- ys is always twice of xs
- **$y = W.x + b$ where $W = 2, b = 0 \Rightarrow$ Best fit: $y = 2x$**

Step 4: Train

Train using Gradient Descent

Train using Gradient Descent with steps in 0.000001

- `train_step = tf.train.GradientDescentOptimizer(0.00001).minimize(cost)`

Train Model

- `sess = tf.Session()`
- `init = tf.initialize_all_variables()`
- `sess.run(init)`
- `steps = 100`

Tensorflow Non-interactive Mode

Nothing is actually executed until `sess.run(...)`

Go!!!!

```
for i in range(steps):
```

```
    # Create fake data for  $y = W.x + b$  where  $W = 2$ ,  $b = 0$ 
```

```
    xs = np.array([[i]])
```

```
    ys = np.array([[2*i]])
```

```
    # Train
```

```
    feed = { x: xs, y_: ys }
```

```
    sess.run(train_step, feed_dict=feed)
```

```
print("After %d iteration:" % i)
```

```
print("W: %f" % sess.run(W))
```

```
print("b: %f" % sess.run(b))
```

Result

Output

After 0 iteration:

W: 0.000000

b: 0.000000

After 1 iteration:

W: 0.000040

b: 0.000040

...

...

After 98 iteration:

W: 1.997181

b: 0.051121

After 99 iteration:

W: 1.997632

b: 0.051126

After 100 iterations:
W ~ 2.0 and b ~ 0.0

Debugging

When Gradient Descent Go Awry?

Re-run:

```
sess.run(tf.initialize_all_variables())  
train_step = tf.train.GradientDescentOptimizer(0.01).minimize(cost)  
for i in range(steps):  
    ....
```


Output

After 0 iteration:

W: 0.000000

b: 0.000000

After 1 iteration:

W: 0.040000

b: 0.040000

After 2 iteration:

W: 0.195200

b: 0.117600

...

...

After 46 iteration:

W: 323066974878921632721935905090174976.000000


b: 7019517401500716765746276968955904.000000

After 47 iteration:


W: -inf

b: -296803829357474433113896004853170176.000000


Finding Minimum Cost Function




Finding Minimum Cost Function



Finding Minimum Cost Function



Finding Minimum Cost Function



Errors: Missing Placeholders

```
feed = { x: xs, y: ys }
```

```
sess.run(train_step, feed_dict=feed)
```

- **y** should be y_

*W tensorflow/core/common_runtime/executor.cc:1076] 0x229e510 Compute
status: Invalid argument: You **must feed a value for placeholder tensor**
'Placeholder_1' with dtype float*

[[Node: Placeholder_1 = Placeholder[dtype=DT_FLOAT, shape=[], _device="/job:localhost/replica:0/task:0/cpu:0"]()]]]

Placeholder Numbers

```
cost = tf.reduce_sum(tf.pow((y_1-y), 2))
```

```
y_1 = tf.placeholder(tf.float32, [None, 1])
```

Placeholder₁

```
y = product + b
```

```
product = tf.matmul(x, W)
```

```
x = tf.placeholder(tf.float32, [None, 1])
```

Placeholder₀

Assigning Names to Placeholders

```
cost = tf.reduce_sum(tf.pow((y_y), 2))
```

```
y_ = tf.placeholder(tf.float32, [None, 1],  
name='y-input')
```

Placeholder_1

```
y = product + b
```

```
product = tf.matmul(x, W)
```

```
x = tf.placeholder(tf.float32, [None, 1])
```

Placeholder_0

Errors: Wrong Shape

```
for i in range(steps):
```

```
    xs = np.array([i])
```

```
    ys = np.array([[2*i]])
```

- xs should be np.array([[i]])

Traceback (most recent call last):

File "<stdin>", line 7, in <module>

*File "/usr/local/lib/python2.7/dist-packages/tensorflow/python/client/session.py",
line 364, in run*

tuple(subfeed_t.get_shape().dims)))

*ValueError: Cannot feed value of shape (1,) for Tensor u'Placeholder:0', which has
shape (Dimension(None), Dimension(1))*

Placeholder Numbers

```
cost = tf.reduce_sum(tf.pow((y_y), 2))
```

```
y_ = tf.placeholder(tf.float32, [None, 1])
```

Placeholder_1

```
y = product + b
```

```
product = tf.matmul(x, W)
```

```
x = tf.placeholder(tf.float32, [None, 1])
```

Placeholder_0

Print: Simplest Debugger

```
print (sess.run(product), feed_dict={ x: np.array([1]) })
```

File "<stdin>", line 1

```
    print (sess.run(product), feed_dict={ x: np.array([1]) })
```

^

SyntaxError: invalid syntax

```
>>> print (sess.run(product), feed_dict={ x: np.array([1]) })
```

Traceback (most recent call last):

File "<stdin>", line 1, in <module>

File "/usr/local/lib/python2.7/dist-packages/tensorflow/python/client/session.py", line 364, in run
 tuple(subfeed_t.get_shape().dims)))

ValueError: Cannot feed value of **shape (1,)** for Tensor u'Placeholder:0', which has shape **(Dimension (None), Dimension(1))**

Same error so must be it!

Progressing to Tensorflow Tutorial

Difference

Parameters	Gentlest Tensorflow Tutorial	Tensorflow Tutorial
Purpose	Find house price (Float32)	Classify image
y	Each sample is 1 x 1; Single-outcome (house price). Feature value is Float32.	Each sample is 10 x 1: One-hot vector for 10 possible values. Vector element is binary.
x	Each sample is 1 x 1: Single-feature (house size) 1 x 1. Feature value is Float32.	Each sample is 1 x 784: 784 features (28 x 28 bits for each image). Feature value is binary
W	1 x 1: Single-outcome, single-feature	784 x 10: one for each 784 bits and 10 probability its one of 10 possible values

One-hot
vector for
image 4

$$\begin{pmatrix} 0 \\ 0 \\ 0 \\ 1 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \end{pmatrix}$$

Difference

Parameters	Gentlest Tensorflow Tutorial	Tensorflow Tutorial
b	1 x 1: Bias for single-outcome	10 x 1: Bias for 10 possible values of on-hot vector
cost	Square of difference	Entropy-based
gradient descent batch size	Single sample x, y	100 samples each batch

Summary Machine Learning

4 steps:

- Build **model** (linear regression, logistic regression, neural networks, etc.)
- Define **cost function** to determine how well current model fits actual
- Collect and prepare **data**
- **Train** model to tweak model based on cost function using data

Summary of Tensorflow

- Modelling

- `tf.zeros`
- `tf.placeholder`
- `tf.Variable`
- `tf.matmul`
- `tf.reduce_sum`
- `tf.pow`
- `tf.train.GradientDescentOptimizer(step_size).minimize(cost)`

- Train

- `init = tf.initialize_all_variables()`
- `sess = tf.session()`
- `sess.run(init)`
- `sess.run(train, feed_dict={ x: ..., y: ... })`

Questions

- Why not just use numpy.dot methods instead of tf.matmul?
 - tf.matmul is a Tensorflow Operation, and Tensorflow can easily assign each parallelizable operation to an available nodes to speed up operations

References

Code: https://github.com/nethsix/gentle_tensorflow/blob/master/code/linear_regression_one_feature.py