

Gentlest Intro to Tensorflow (Part 2)

Khor Soon Hin, @neth_6, re:Culture

In collaboration with Sam & Edmund



Next Steps

- Review of Tensorflow
 - In illustrations
 - `tf.placeholder` & feed
 - Training process
- Tensorboard: Visualize **cost**, **W**, **b**
- Training batch size: Stochastic/Mini-batch/Full-batch

Quick Review

Given house size, use machine learning to predict house price

Quick Review


House Price, \$



House Size, sqm

Quick Review

House Price, \$




Collect ('generate fake') data
 y (house price) = $2 \cdot x$ (house size)

House Size, sqm

Quick Review

House Price, \$




Model using linear regression
 $y = \mathbf{W} \cdot \mathbf{x} + \mathbf{b}$

House Size, sqm

Quick Review

House Price, \$




Model using linear regression
 $y = \mathbf{W} \cdot \mathbf{x} + \mathbf{b}$

Learn the best \mathbf{W} , \mathbf{b}

House Size, sqm

Quick Review

House Price, \$




Model using linear regression
 $y = \mathbf{W} \cdot \mathbf{x} + \mathbf{b}$

Learn the best \mathbf{W} , \mathbf{b}
by reducing **cost**

House Size, sqm

Quick Review



Quick Review

Modeling linear regression in Tensorflow

Tensorflow Graph

Model: $y = W.x + b$

Cost: `tf.reduce_mean(tf.square(y_ - y))`

Tensorflow Graph

prediction

Model: $\hat{y} = W.x + b$

Cost: `tf.reduce_mean(tf.square(y_ - \hat{y}))`

Tensorflow Graph

prediction

Model:

$$y = W \cdot x + b$$

Goal: Train

Cost: `tf.reduce_mean(tf.square(y_ - y))`

Tensorflow Graph

prediction

Goal: Train

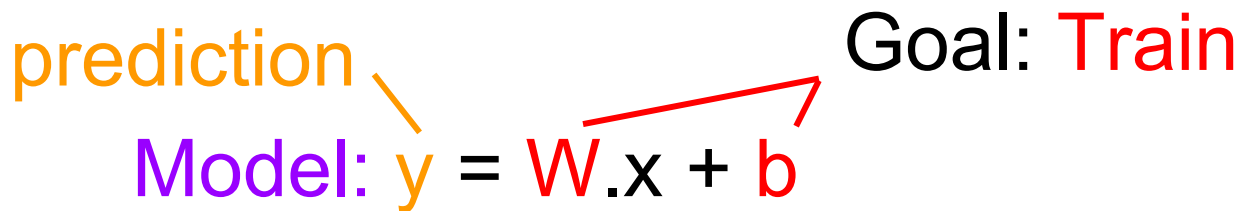
Model: $y = W \cdot x + b$

Cost: `tf.reduce_mean(tf.square(y_ - y))`

How: Minimizing actual vs. prediction

Tensorflow Graph

prediction
Model: $y = W \cdot x + b$ Goal: Train



Cost: `tf.reduce_mean(tf.square(y_ - y))`



How: Minimizing actual vs. prediction



Train: `tf.train.GradientDescentOptimizer(learning_rate).minimize(cost)`

What is **Train**?

W, b

Init with some values

What is **Train**?





Model: $y = W.x + b$

Cost: `tf.reduce_mean(tf.square(y_ - y))`

W, b

Init with some values

What is Train?

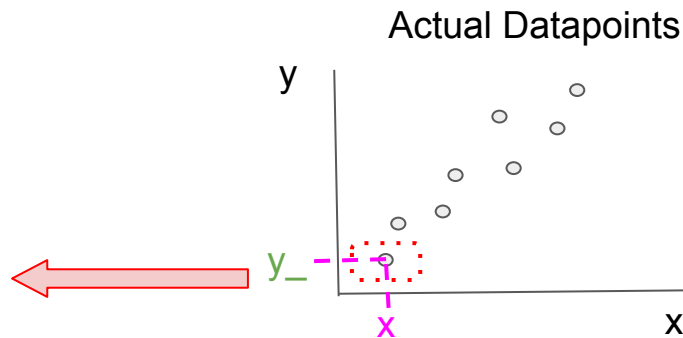


Model: $y = W \cdot x + b$




Cost: `tf.reduce_mean(tf.square(y_ - y))`

W, b

Init with some values



What is **Train**?




Model: $y = W.x + b$

Cost: `tf.reduce_mean(tf.square(y_ - y))`

W, b


Init with some values



Train: `tf.train.GradientDescentOptimizer(01).minimize(cost)`


Gradient Descent

Cost

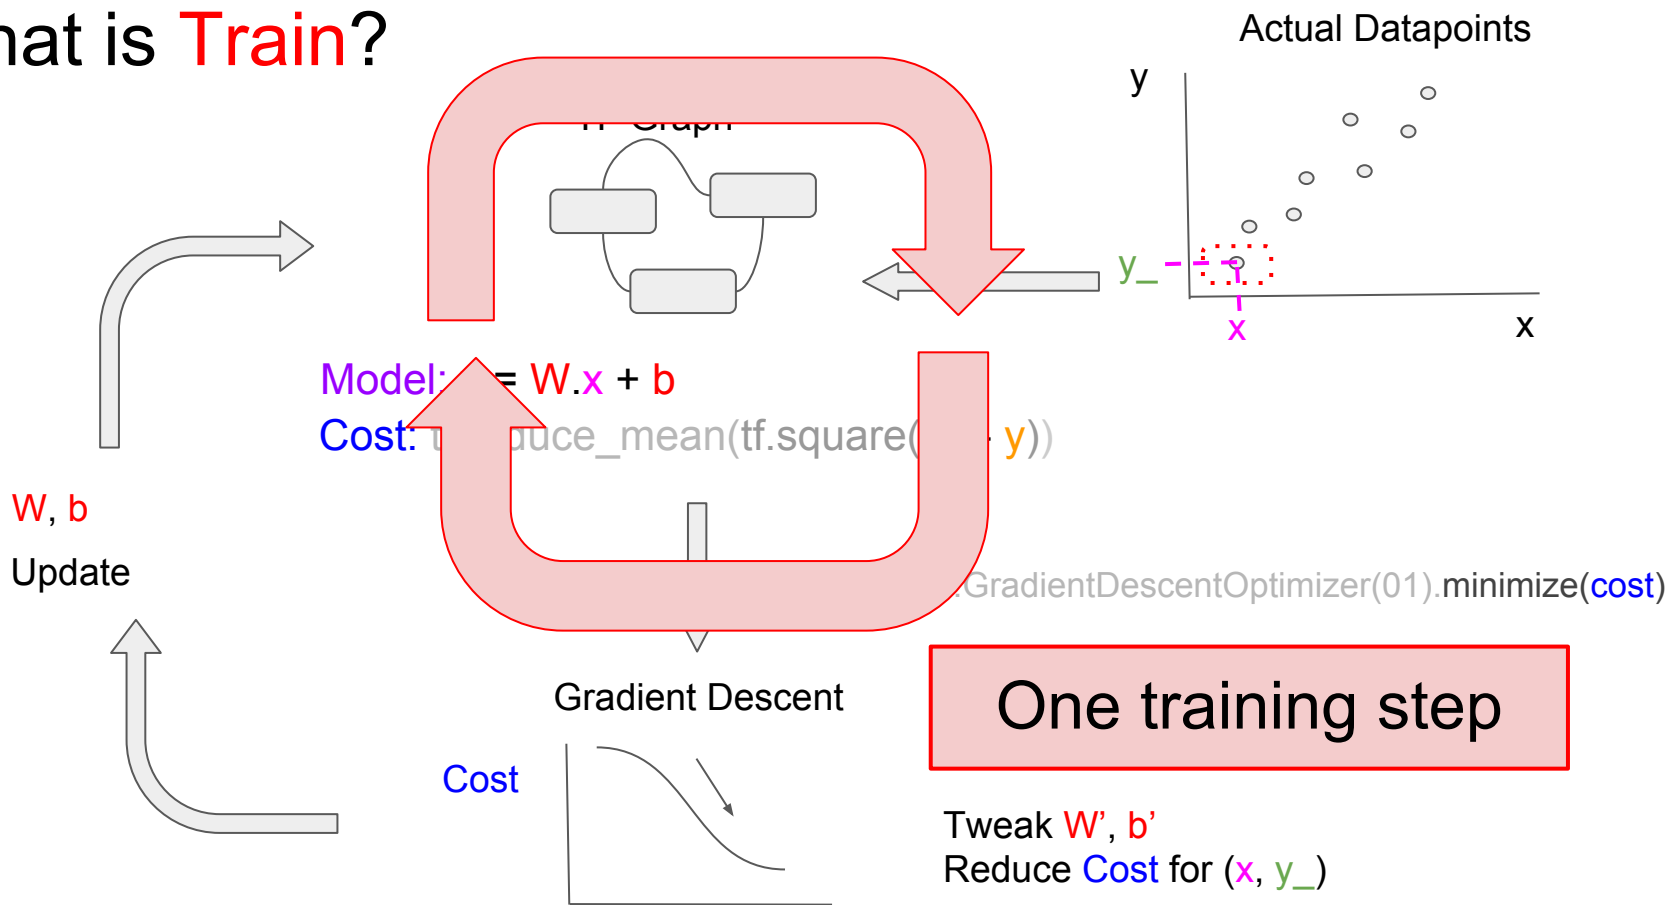


Tweak W' , b'
Reduce Cost for $(x, y_)$


What is **Train**?




What is Train?



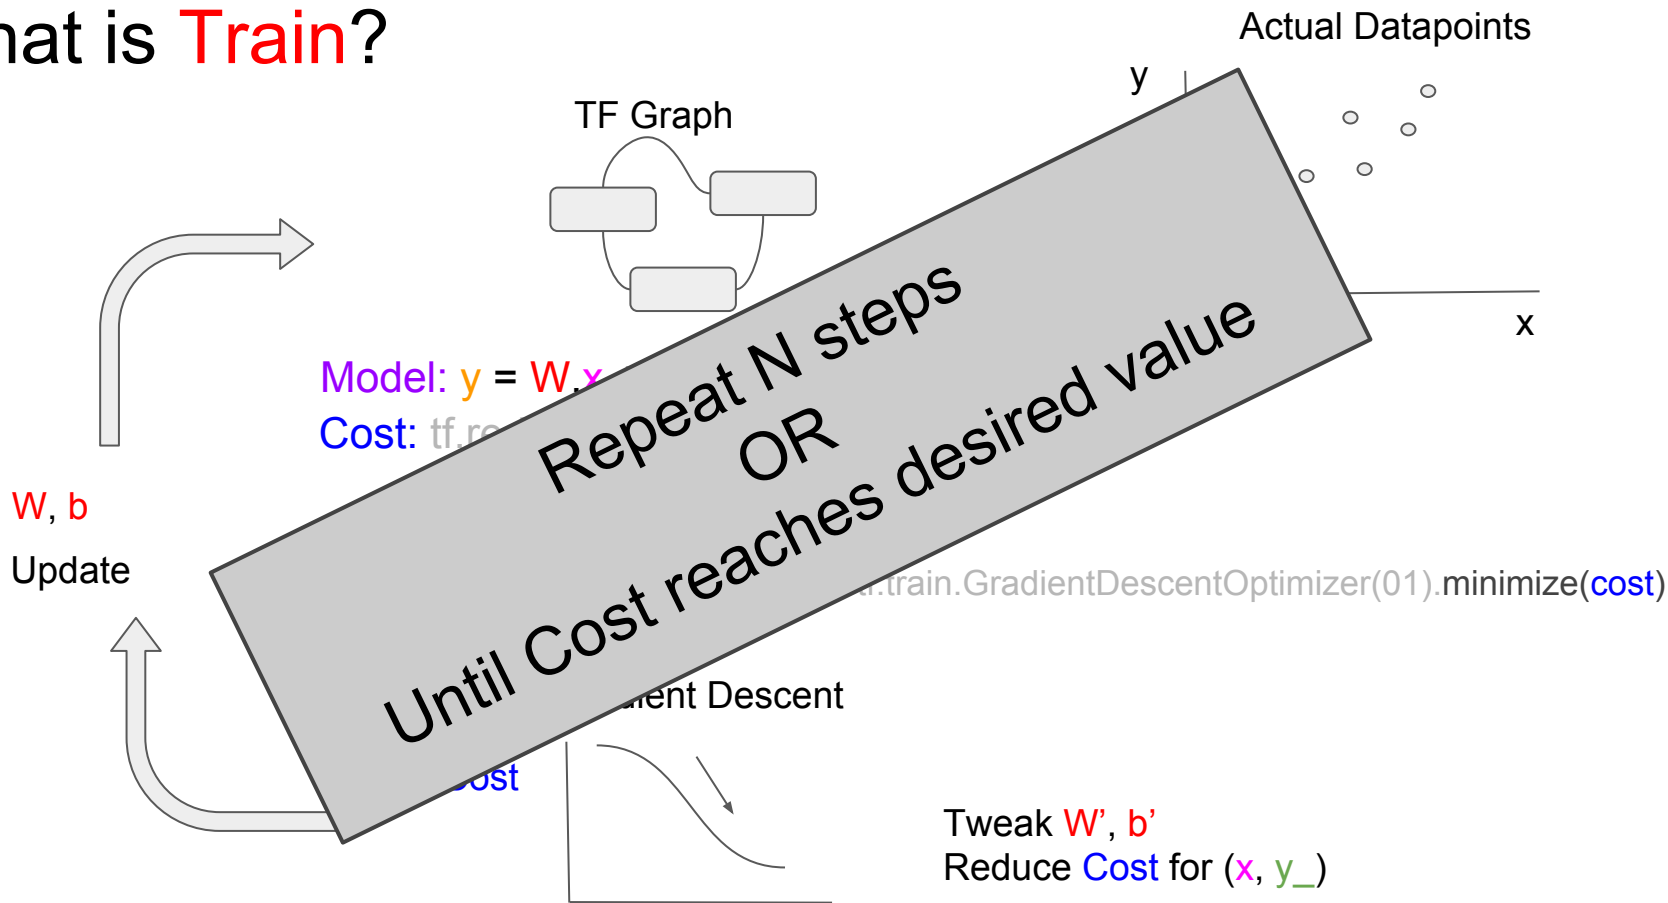
What is **Train**?



What is **Train**?



What is **Train**?




Creating Tensorflow Graph


Tensorflow Graph

Model in TF!

TF Graph



Actual Datapoints



Model: $y = W.x + b$

Cost: `tf.reduce_mean(tf.square(y_ - y))`


W, b

Update

Train: `tf.train.GradientDescentOptimizer(0.1).minimize(cost)`


Gradient Descent

Cost




Tweak W', b'
Reduce Cost for $(x, y_)$

Tensorflow Graph

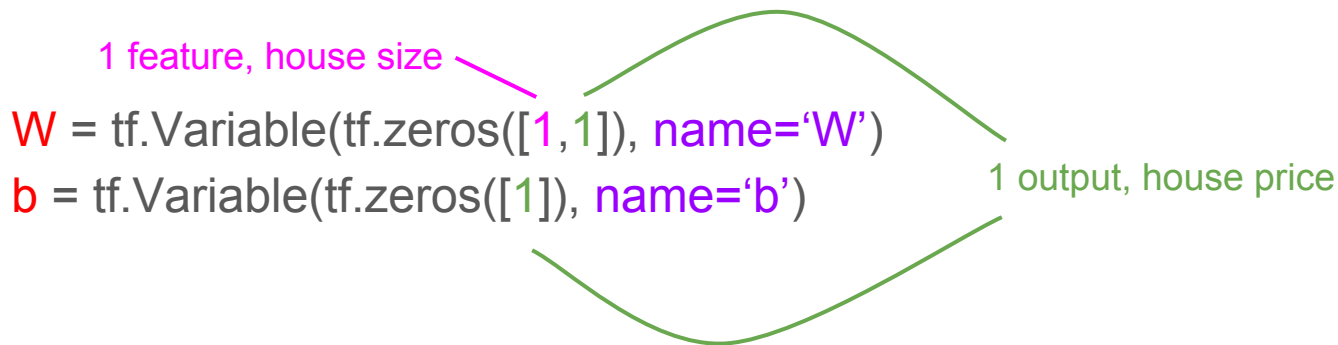


Tensorflow Graph



Tensorflow Graph: tf.Variable (Enhance visualization)

tf.Variable for holding variables we want to learn, W, b



NOTE: Without `name=XXX`, the variables will be assigned names like `Variable_0`, `Variable_1`, etc.

Tensorflow Graph: tf.placeholder (Enhance visualization)

tf.placeholder for actual data, e.g., house size (x), house price (y_)

```
x = tf.placeholder(tf.float32, [None, 1], name='x')  
y_ = tf.placeholder(tf.float32, [None, 1], name='y_')
```

1 feature, house size

1 output, house price

NOTE: Without `name=XXX`, the placeholders will be assigned names like Placeholder_0, Placeholder_1, etc.

Tensorflow Graph: Model (Enhance visualization)

Model linear regression $y = Wx + b$

with tf.name_scope("Wx_b") as scope:

product = tf.matmul(x, W)

y = product + b

NOTE: With `scope`, the all definitions under it, will appear as a black-box that is expandable!

Tensorflow Graph: Cost (Enhance visualization)

Cost function: Least-squared difference

```
with tf.name_scope("cost") as scope:  
    cost = tf.reduce_mean(tf.square(y_g-y))
```


Tensorflow Graph: Train (Enhance visualization)

```
# Train step
```

```
with tf.name_scope("train") as scope:
```

```
    train_step = tf.train.GradientDescentOptimizer(learning_rate).minimize(cost)
```

Tensorboard

Two Main Components

- Summary
 - Mark data we want to collect
 - W , b , $cost$
- Writer
 - Write data collected to event file

Tensorflow Summary: Scalar

tf.scalar_summary

cost is *single-dimension* so

mark it as data to be collected using *scalar_summary*

cost = tf.reduce_mean(tf.square(y_-y))

cost_sum = tf.scalar_summary("cost", **cost**)

Tensorflow Summary: Histogram

`tf.histogram_summary`

W and b can be **multi-dimensional** so

mark them as data to be collected using *histogram_summary*

```
W = tf.Variable(tf.zeros([1,1]), name="W")
```

```
W_hist = tf.histogram_summary("weights", W)
```

```
b = tf.Variable(tf.zeros([1]), name="b")
```

```
b_hist = tf.histogram_summary("biases", b)
```

Tensorflow Summary: Merge

tf.merge_all_summaries

```
merged = tf.merge_all_summaries()
```

tf.merge_summary

```
merged = tf.merge_summary([W_hist, b_hist, cost_sum])
```

Tensorflow Writer

The specified log directory will be:

* Created if it does not exist

* If exist and not empty data is merged

sess.graph_def will enable Tensorflow to draw graph representation

for network

```
writer = tf.train.SummaryWriter("/event_log_dir", sess.graph_def)
```

Tensorflow Writer: When to Record Data

NOTE: Train using #steps

for i in steps:

...

NOTE: Write summary every 10 samples

if i % 10 == 0:

all_feed = { x: all_xs, y_: all_ys }

NOTE:

* **merged** is our merged summary name

* feed is the batch of all x, y data points, to calculate **TOTAL** cost

result = sess.run(**merged**, feed_dict=all_feed)

writer.add_summary(result, i)

...

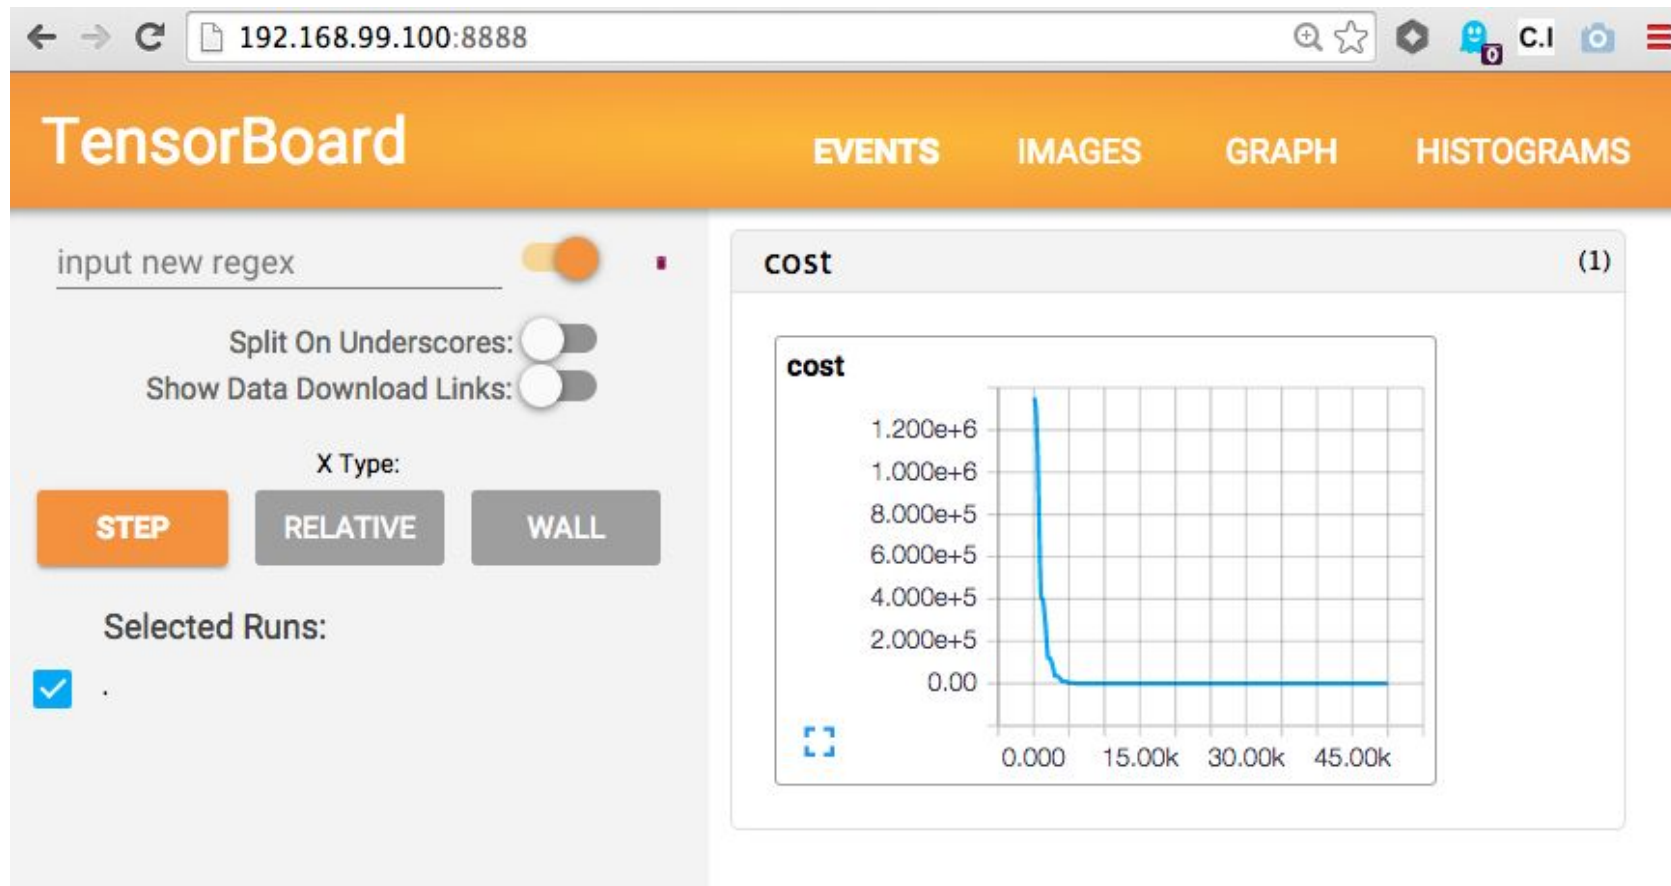
Start Tensorboard Server

NOTE: Single line

```
python /usr/local/lib/python2.7/dist-packages/tensorflow/tensorboard/tensorboard.  
py --logdir /event_log_dir/ --port 8888
```


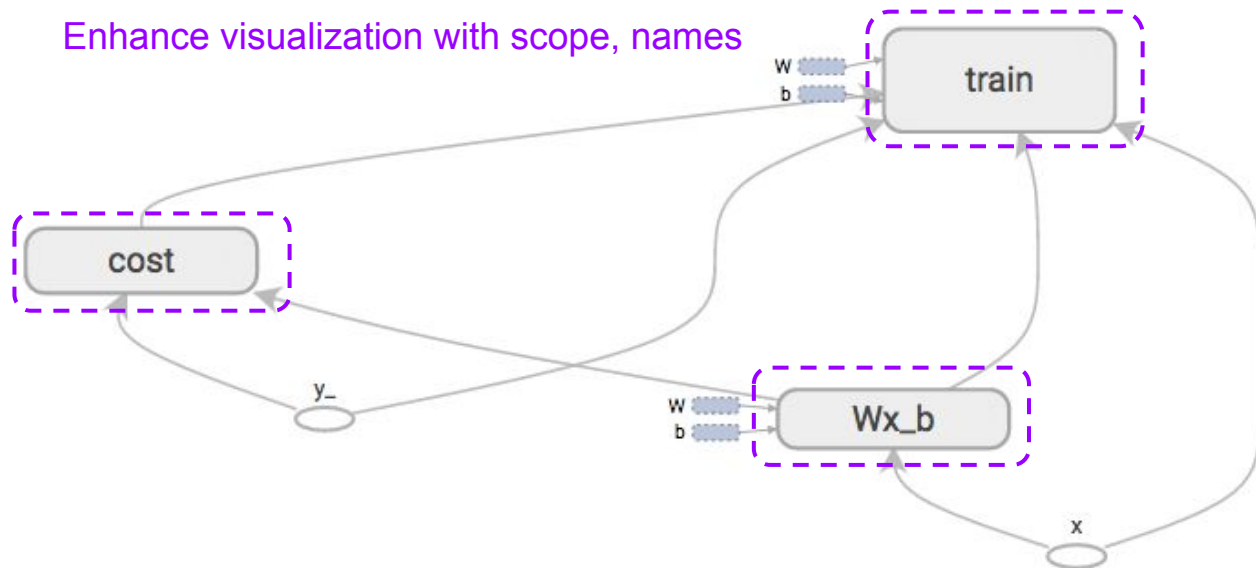
WARNING:tensorflow:Found more than one graph event per run.Overwriting the
graph with the newest event

Access Using Browser




Tensorflow Graph

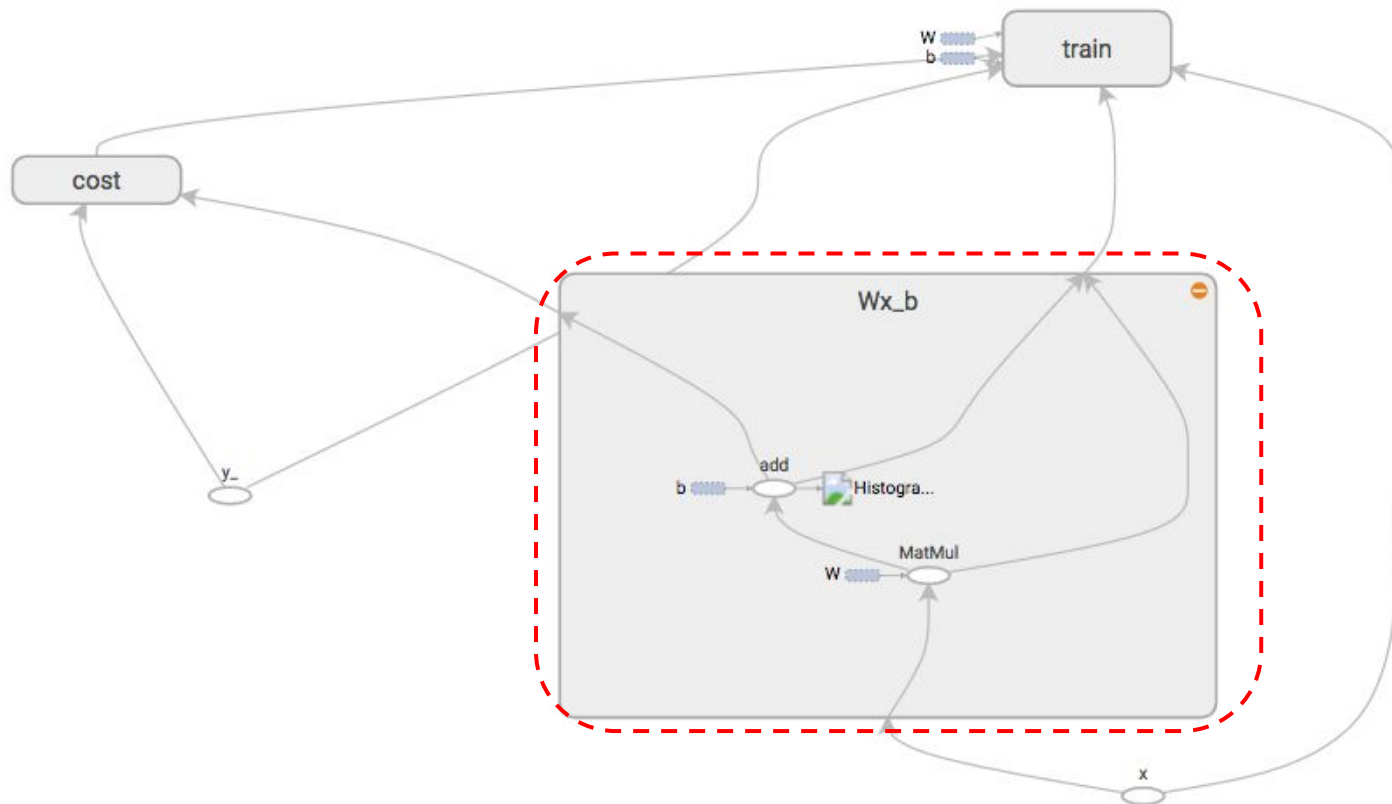
Enhance visualization with scope, names



Tensorflow Graph




Tensorflow Graph: Expanded



Importance of Visualization:
Have we found a good fit?



Linear Regression with 1-Feature

$b \sim 0.0$



Cost reaches zero = good model fit

$W \sim 2.0$



Training Linear Regression with $b > 0$

Generate fake data for $y = 2.x + 10$, i.e., $W = 2$, $b = 10$

```
all_xs = []
```

```
all_ys = []
```

```
for i in range(datapoint_size):
```

```
    # Create fake data for  $y = W.x + b$  where  $W = 2$ ,  $b = 10$ 
```

```
    all_xs.append(i)
```

```
    all_ys.append(2*i+10)
```


all_xs
(house size)

0
1
2
3
⋮
⋮
⋮
⋮
⋮
99


all_ys (house
price)

0
12
14
16
⋮
⋮
⋮
⋮
⋮
208

Linear Regression with 1-Feature ($b > 0$)




Cost appears to reaches zero ...



Linear Regression with 1-Feature ($b > 0$)

$W \sim 2.0$

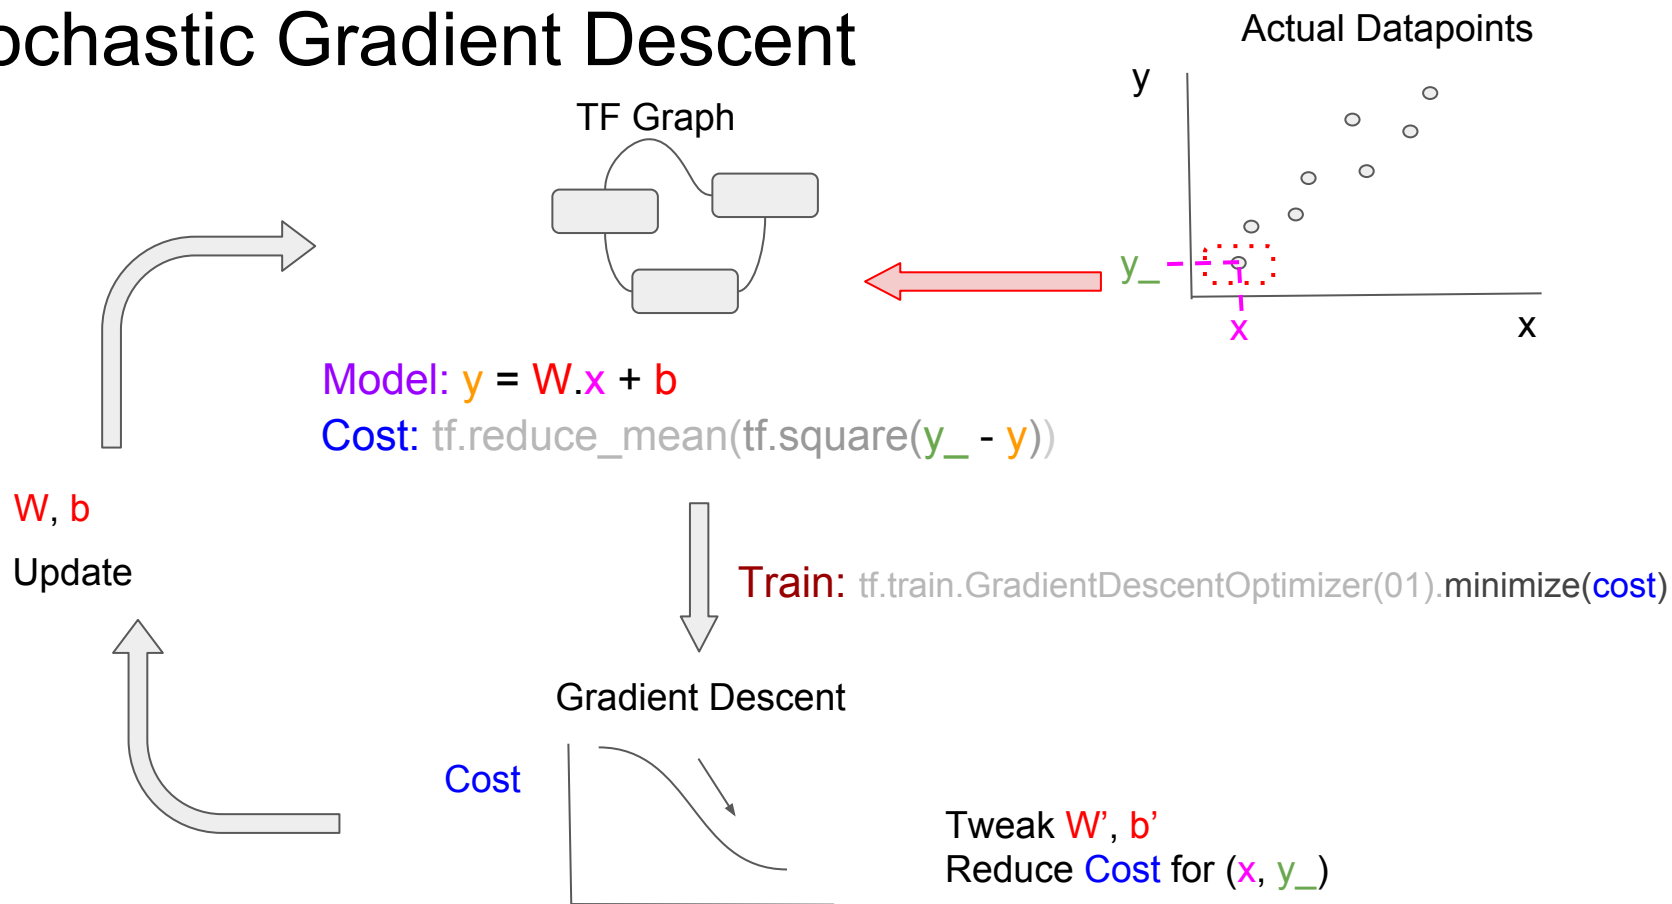


Is there a way to learn faster?



- Faster gradient descent learning rate?
- Different model?
- Different cost function?

Batch Size

Stochastic Gradient Descent



Mini-batch Gradient Descent



Model: $y = W.x + b$

Cost: `tf.reduce_mean(tf.square(y_ - y))`


W, b

Update

Train: `tf.train.GradientDescentOptimizer(01).minimize(cost)`



Gradient Descent

Cost



Tweak W', b'
Reduce Cost for $(x, y_1), (x_2, y_2), (x_3, y_3)$

Batch Gradient Descent



Model: $y = W.x + b$

Cost: `tf.reduce_mean(tf.square(y_ - y))`


W, b

Update

Train: `tf.train.GradientDescentOptimizer(01).minimize(cost)`

Gradient Descent

Cost



Tweak W' , b'
Reduce Cost for $(x, y_1), (x_2, y_2), \dots (x_m, y_m)$

Stochastic/Mini-batch/Batch Gradient Descent

```
x = tf.placeholder(tf.float32, [None, 1])  
y_ = tf.placeholder(tf.float32, [None, 1])  
...  
train_step = tf.train.GradientDescentOptimizer(0.1).minimize(cost)
```

One feature, house size

One output, house price

xs =

ys =

ONLY this part is different!

```
feed = { x: xs, y_: ys }
```

```
sess.run(train_step, feed_dict=feed)
```

Why Feed `train_step`, with `feed_dict = x, y_`

```
train_step = tf.train.GradientDescentOptimizer(learning_rate).minimize(cost)
```

```
cost = tf.reduce_mean(tf.square(y_-y))
```

```
y = tf.matmul(x, W) + b
```

```
y_ = tf.placeholder(tf.float32, [None, 1])
```

```
x = tf.placeholder(tf.float32, [None, 1])
```

`train_step` has dependency on `x`, `y_`

Stochastic Gradient Descent

Stochastic Gradient Descent

```
x = tf.placeholder(tf.float32, [None, 1])
```

One feature, house size

```
y_ = tf.placeholder(tf.float32, [None, 1])
```

One output, house price

```
for i in range(steps):
```

```
# Create fake data for  $y = W.x + b$  where  $W = 2$ ,  $b = 0$ 
```

```
xs = np.array([[i]])
```

Provide datapoint for feature, house size each step

```
ys = np.array([[2*i]])
```

Stochastic Gradient Descent


Provide datapoint for output, house price, each step

```
feed = { x: xs, y_: ys }
```


```
sess.run(train_step, feed_dict=feed)
```

Mini-batch Gradient Descent


Mini-batch Gradient Descent: x , y



Mini-batch Gradient Descent: x , y



Mini-batch Gradient Descent: x, y_



Mini-batch Gradient Descent: x, y_

all_xs


0
1
2
3
4
5
6
7
8
9
.
.
.
.
.
999

all_ys

0
2
4
6
8
10
12
14
16
18
.
.
.
.
.
.
1998

Train step #35

Mini-batch Gradient Descent: x , y



Mini-batch Gradient Descent

```
x = tf.placeholder(tf.float32, [None, 1])
```

One feature, house size

```
y_ = tf.placeholder(tf.float32, [None, 1])
```

One output, house price

```
# Prepare all actual house size (x), house price (y_) before hand
```

```
all_xs = ...
```

```
all_ys = ...
```

```
for i in range(steps):
```

```
    xs = np.array(all_xs[batch_start_idx:batch_end_idx])
```

```
    ys = np.array(all_ys[batch_start_idx:batch_end_idx])
```

```
    feed = { x: xs, y_: ys }
```

```
    sess.run(train_step, feed_dict=feed)
```

Code to do data batching

Mini-batch Gradient Descent

```
x = tf.placeholder(tf.float32, [None, 1], name="x-input")
```

One feature, house size

```
y_ = tf.placeholder(tf.float32, [None, 1], name="y-input")
```

One output, house price

```
# Prepare all actual house size (x), house price (y_) before hand
```

```
all_xs = ...
```

```
all_ys = ...
```

'None' in placeholder means determined at run-time

```
for i in range(steps):
```

```
    xs = np.array(all_xs[batch_start_idx:batch_end_idx])
```

```
    ys = np.array(all_ys[batch_start_idx:batch_end_idx])
```


```
    feed = { x: xs, y_: ys }
```

```
    sess.run(train_step, feed_dict=feed)
```


Code to do data batching

Batch Gradient Descent


Batch Gradient Descent: x, y_



Batch Gradient Descent: $x, y_{\text{...}}$



Batch Gradient Descent: $x, y_{\text{...}}$



Batch Gradient Descent

```
x = tf.placeholder(tf.float32, [None, 1], name="x-input")
```

One feature, house size

```
y_ = tf.placeholder(tf.float32, [None, 1], name="y-input")
```

One output, house price

```
# Prepare all actual house size (x), house price (y_) before hand
```

```
all_xs = ...
```

```
all_ys = ...
```

```
for i in range(steps):
```

```
    xs = all_xs
```

```
    ys = all_ys
```

```
    feed = { x: xs, y_: ys }
```

```
    sess.run(train_step, feed_dict=feed)
```

What is the Significance of Batch Size?

Find steepest gradient descent (**train_step**) from the viewpoint of the points in the batch

Resource to train each step


Less

More

[Batch Size 1] [Batch Size 100] [Full Batch]


Outcome: With same Gradient Descent learning rate ...

Batch Size = 1, Train Steps = 1000




Cost still decreasing towards 0


Still learning ...



Still learning ...





Batch Size = 100, Train Steps = 1000




Cost close to 0 = good model fit

b ~ 0.0

W ~ 2.0





Batch Size = 1000 (Full), Train Steps = 1000



Cost close to 0 = good model fit

b ~ 0.0

W ~ 2.0



Resource to train each step

Less

More

[Batch Size 1] [Batch Size 100] [Full Batch]

Many

Few

Steps to reach good W , b values

References

Code:

- Using Tensorboard for visualization:
 - https://github.com/nethsix/gentle_tensorflow/blob/master/code/linear_regression_one_feature_with_tensorboard.py
- Performing stochastic/mini-batch/batch Gradient Descent using Tensorflow
 - https://github.com/nethsix/gentle_tensorflow/blob/master/code/linear_regression_one_feature_using_mini_batch_with_tensorboard.py