

Big Data Visualization

using

Apache Spark and Zeppelin

Prajod Vettiyattil, Software Architect, Wipro

<https://in.linkedin.com/in/prajod>

@prajods

saltmarch
MEDIA

GREAT INDIAN
DEVELOPER
SUMMIT



Agenda

- Big Data and Ecosystem tools
- Apache Spark
- Apache Zeppelin
- Data Visualization
- Combining Spark and Zeppelin

BIG DATA AND ECOSYSTEM TOOLS

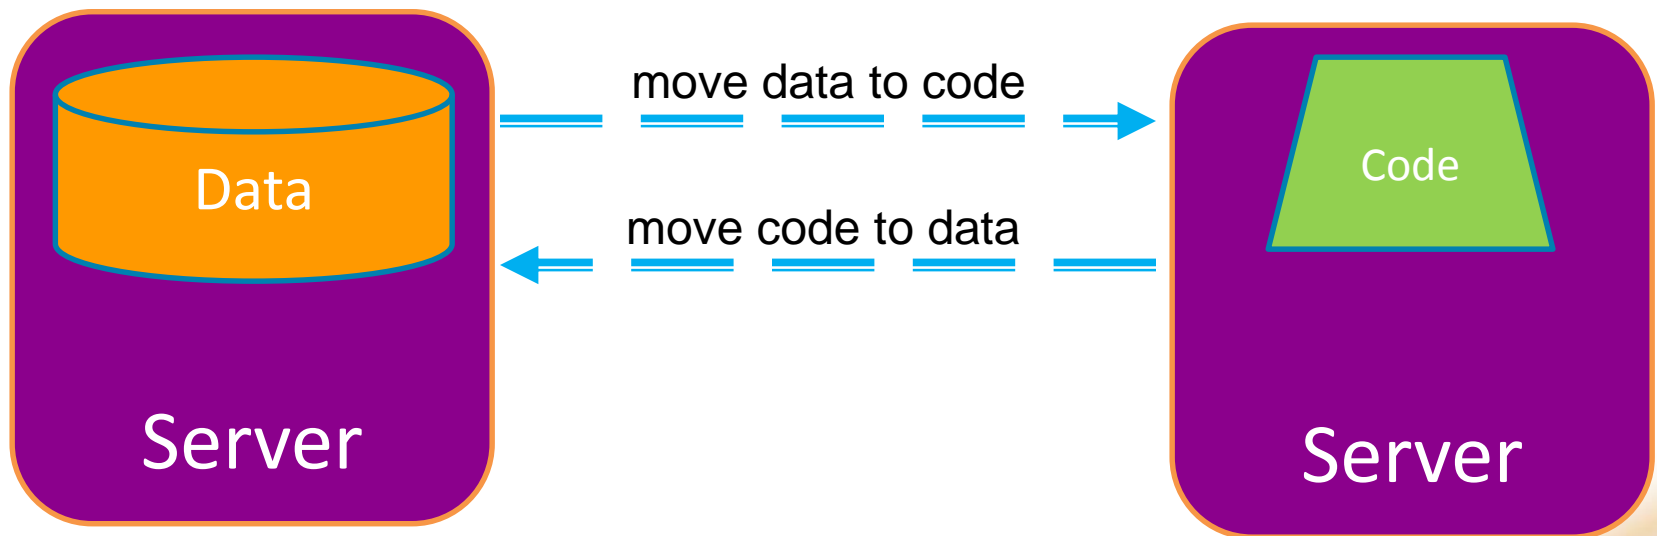


Big Data

- Data size beyond systems capability
 - Terabyte, Petabyte, Exabyte
- Storage
 - Commodity servers, RAID, SAN
- Processing
 - In reasonable response time
 - A challenge here

Tradition processing tools

- Move what ?
 - the data to the code or
 - the code to the data



Traditional processing tools

- Traditional tools
 - RDBMS, DWH, BI
 - High cost
 - Difficult to scale beyond certain data size
 - price/performance skew
 - data variety not supported

Map-Reduce and NoSQL

- Hadoop toolset
 - Free and open source
 - Commodity hardware
 - Scales to exabytes(10^{18}), maybe even more
- Not only SQL
 - Storage and query processing only
 - Complements Hadoop toolset
 - Volume, velocity and variety


All is well ?

- Hadoop was designed for batch processing
- Disk based processing: slow
- Many tools to enhance Hadoop's capabilities
 - Distributed cache, Haloop, Hive, HBase
- Not for interactive and iterative

TOWARDS SINGULARITY



What is singularity ?




Technological singularity

- When AI capability exceeds Human capacity
- AI or non-AI singularity
- 2045: http://en.wikipedia.org/wiki/Ray_Kurzweil
 - The predicted year

APACHE SPARK



History of Spark



Contributors in Spark

- Yahoo
- Intel
- UC Berkeley
- ...
- 50+ organizations

Hadoop and Spark

- Spark complements the Hadoop ecosystem
- Replaces: Hadoop MR
- Spark integrates with
 - HDFS
 - Hive
 - HBase
 - YARN

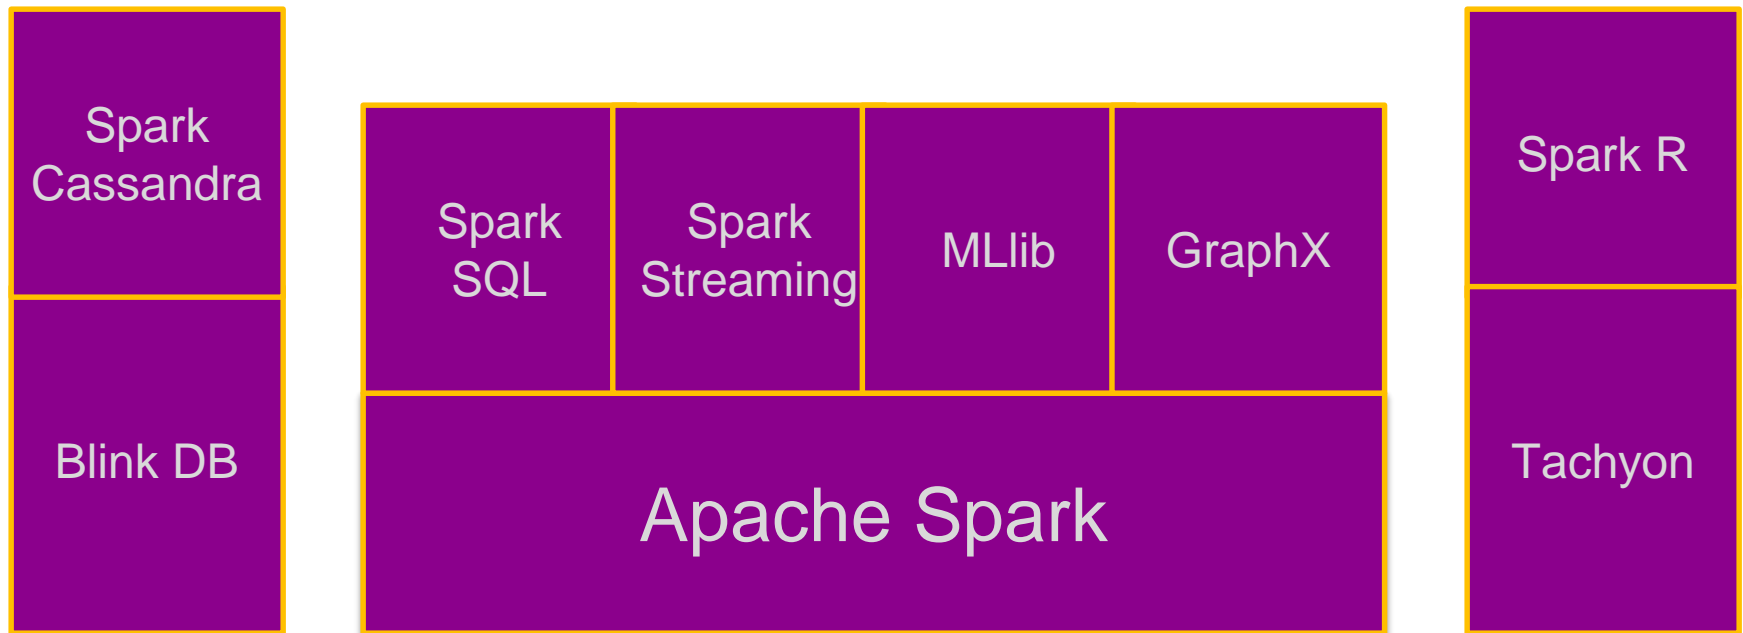
Other big data tools

- Spark also integrates with
 - Kafka
 - ZeroMQ
 - Cassandra
 - Mesos

Programming Spark


- Java
- Scala
- Python
- R

Spark toolset



What is Spark for ?

Batch



Interactive

Streaming

The main difference: speed

- RAM access vs Disk access
 - RAM access is 100,000 times faster !

Latency numbers every programmer should know


L1 cache reference	0.5 ns	
Branch mispredict	5 ns	
L2 cache reference	7 ns	
Mutex lock/unlock	25 ns	
Main memory reference	100 ns	
Compress 1K bytes with Zippy	3,000 ns	= 3 µs
Send 2K bytes over 1 Gbps network	20,000 ns	= 20 µs
SSD random read	150,000 ns	= 150 µs
Read 1 MB sequentially from memory	250,000 ns	= 250 µs
Round trip within same datacenter	500,000 ns	= 0.5 ms
Read 1 MB sequentially from SSD*	1,000,000 ns	= 1 ms
Disk seek	10,000,000 ns	= 10 ms
Read 1 MB sequentially from disk	20,000,000 ns	= 20 ms
Send packet CA->Netherlands->CA	150,000,000 ns	= 150 ms

Assuming ~1GB/sec SSD

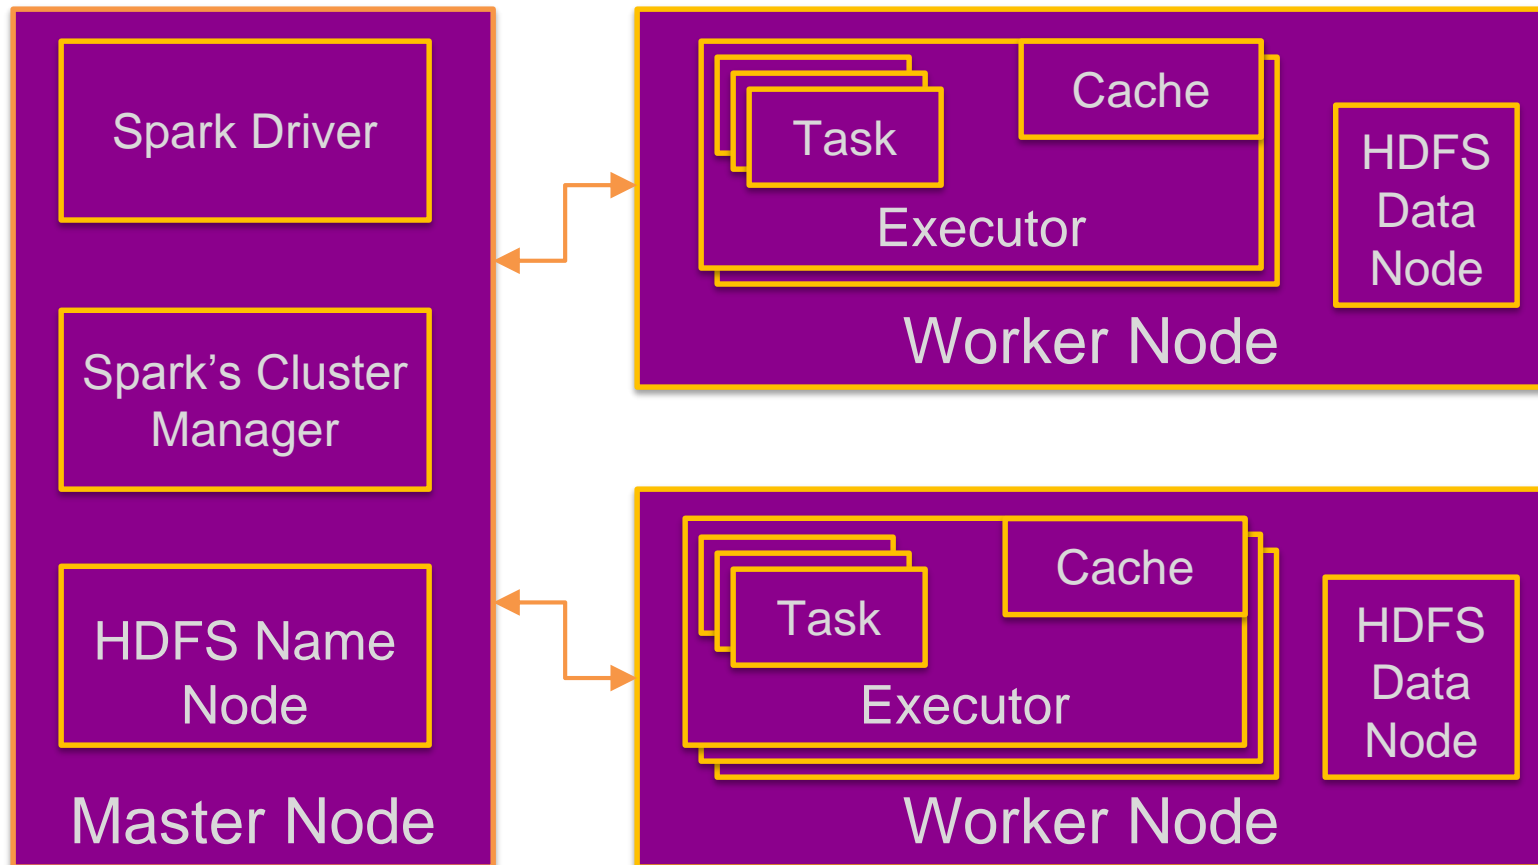
<https://gist.github.com/hellerbarde/2843375>

Lambda Architecture pattern

- Used for Lambda architecture implementation
 - Batch layer
 - Speed layer
 - Serving layer



Deployment Architecture



Core features of Spark

- Rich API
- RDD: Resilient Distributed Datasets
- DAG based execution
- Data caching
- Strong ecosystem tool support

Sample code in scala

- Find the top 100,000 Wikipedia pages by page views
- Log file format: code, title, num_hits
- ```
enPages.map(l => l.split(" "))
 .map(l => (l(1), l(2).toInt))
 .reduceByKey(_+_, 200)
 .filter(x => x._2 > 100000)
 .map(x => (x._2, x._1))
 .collect
 .foreach(println)
```

# APACHE ZEPPELIN

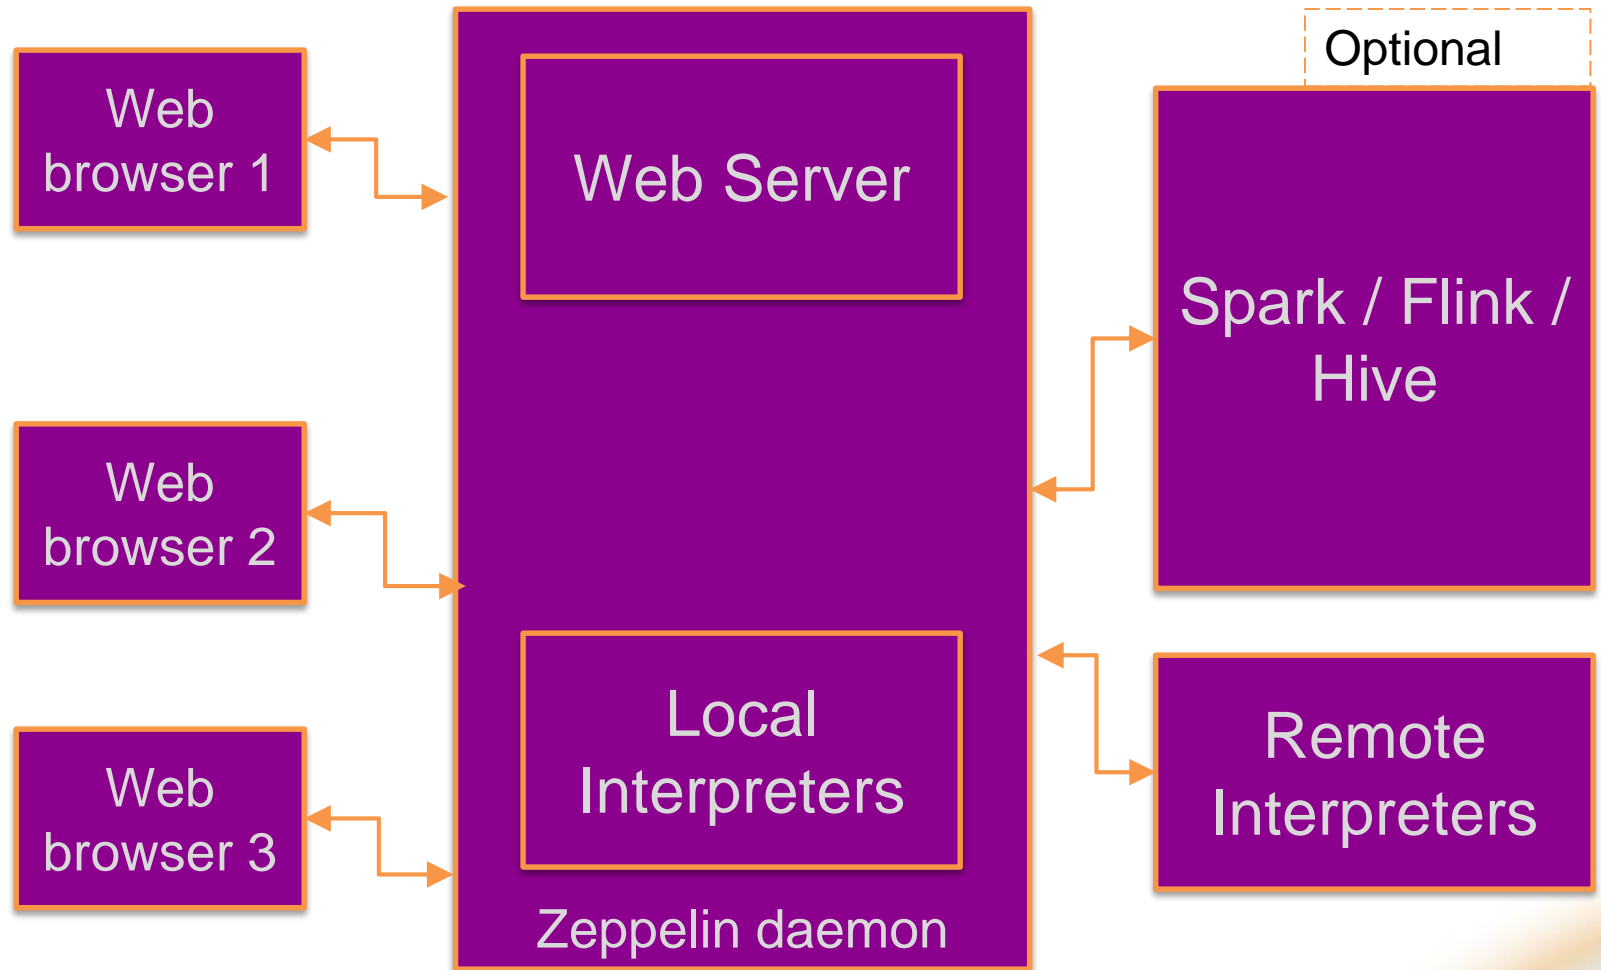


# Interactive data analytics

- For Spark and Flink
- Web front end
- At the back end, it connects to
  - SQL systems(Eg: Hive)
  - Spark
  - Flink



# Deployment Architecture



# Notebook

- Is where you do your data analysis
- Web UI REPL with pluggable interpreters
- Interpreters
  - Scala, Python, Angular, SparkSQL, Markdown and Shell

# Notebook:view

The screenshot shows the Zeppelin Notebook interface in a web browser. The URL is `10.234.10.7:8080/#/notebook/2AGFAVZEP`. The notebook is titled "Untitled" and is in "FINISHED" state. The SQL query being executed is:

```
%sql select age, count(1) from bank where marital="{marital=single,single|divorced|married}" group by age order by age
```

Below the query, there is a dropdown menu for the "marital" variable, currently set to "married". The results are displayed in a table with two columns: "age" and "c1".

| age | c1  |
|-----|-----|
| 20  | 3   |
| 21  | 5   |
| 22  | 9   |
| 23  | 27  |
| 24  | 53  |
| 25  | 98  |
| 26  | 170 |
| 27  | 233 |
| 28  | 325 |

The interface also shows a status bar at the bottom indicating "READY".

# User Interface features

- Markdown
- Dynamic HTML generation
- Dynamic chart generation
- Screen sharing via websockets

# SQL Interpreter

- SQL shell
  - Query spark data using SQL queries
  - Return normal text, HTML or chart type results

# Scala interpreter for Spark


- Similar to the Spark shell
- Upload your data into Spark
- Query the data sets(RDDs) in your Spark server
- Execute map-reduce tasks
- Actions on RDD
- Transformations on RDD



# DATA VISUALIZATION



# Visualization tools



## Arbor.js

A library of force-directed layout algorithms plus abstractions for graph organization and refresh handling.



## CartoDB

A web service for mapping, analyzing and building applications with data.




## Chroma.js

Interactive color space explorer that allows to preview a set of linear interpolated equidistant colors.




## Circos

A software package for visualizing data in a circular layout.




## Cola.js

A library for arranging networks using constraint-based optimization techniques.




## ColorBrewer

A web tool for selecting colors for maps.




## Cubism.js

A library for creating interactive time series and horizon graphs based on D3.js




## Cytoscape

An application for visualizing complex networks and integrating these with any type of attribute data.



## D3.js

A small, flexible and efficient library to create and manipulate interactive documents based on data.


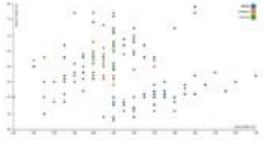

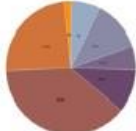

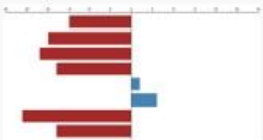
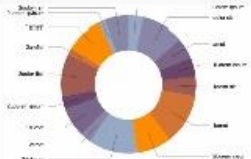
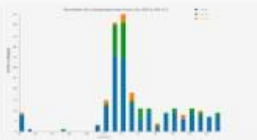
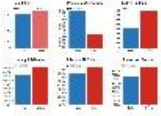





## Dance.js

A simple data-driven visualization framework based on Data.js and Underscore.js

Source: <http://selection.datavisualization.ch/>


# D3 Visualizations

|                                                                                                               |                                                                                                                                |                                                                                                                            |                                                                                                                             |
|---------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------|
| <p>Grouped Bar Chart</p>     | <p>Scatterplot</p>                            | <p>Donut Chart</p>                       | <p>Pie Chart</p>                         |
| <p>Donut Multiples</p>       | <p>Bar Chart with Negative Values</p>         | <p>Animated Donut Chart with Labels</p>  | <p>Stacked Bar Charts on time scale</p>  |
| <p>Bar Chart Multiples</p>  | <p>d3pie - pie chart generator and lib</p>  | <p>3D Donut</p>                         | <p>Gradient Pie</p>                     |

Source: <https://github.com/mbostock/d3/wiki/Gallery>

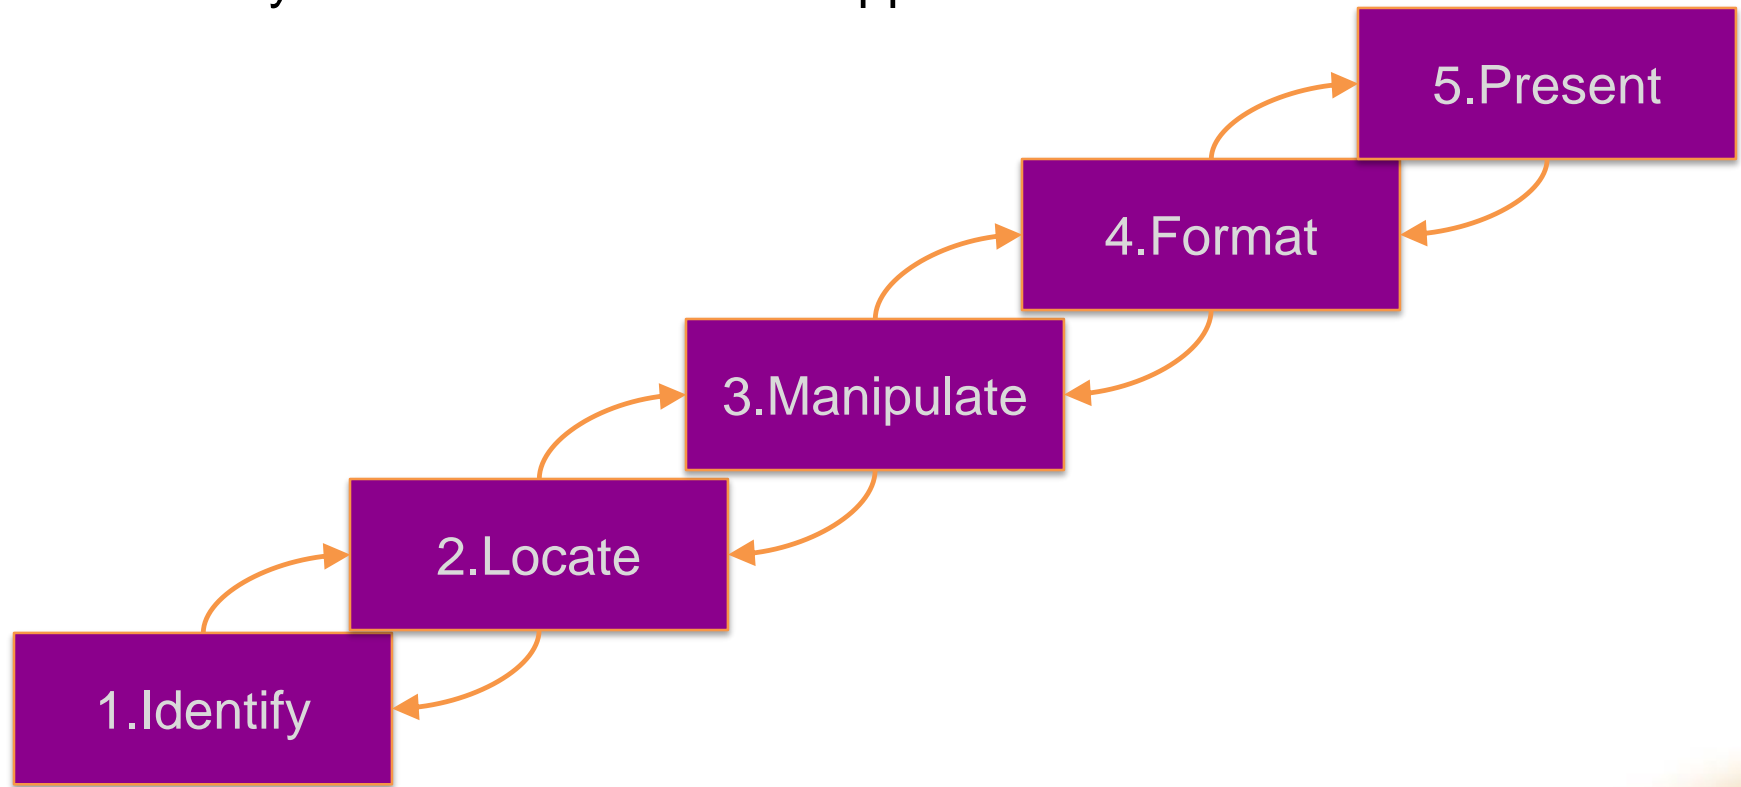


# The need for visualization



# Tools for Data Presentation Architecture

A data analysis tool/toolset would support:

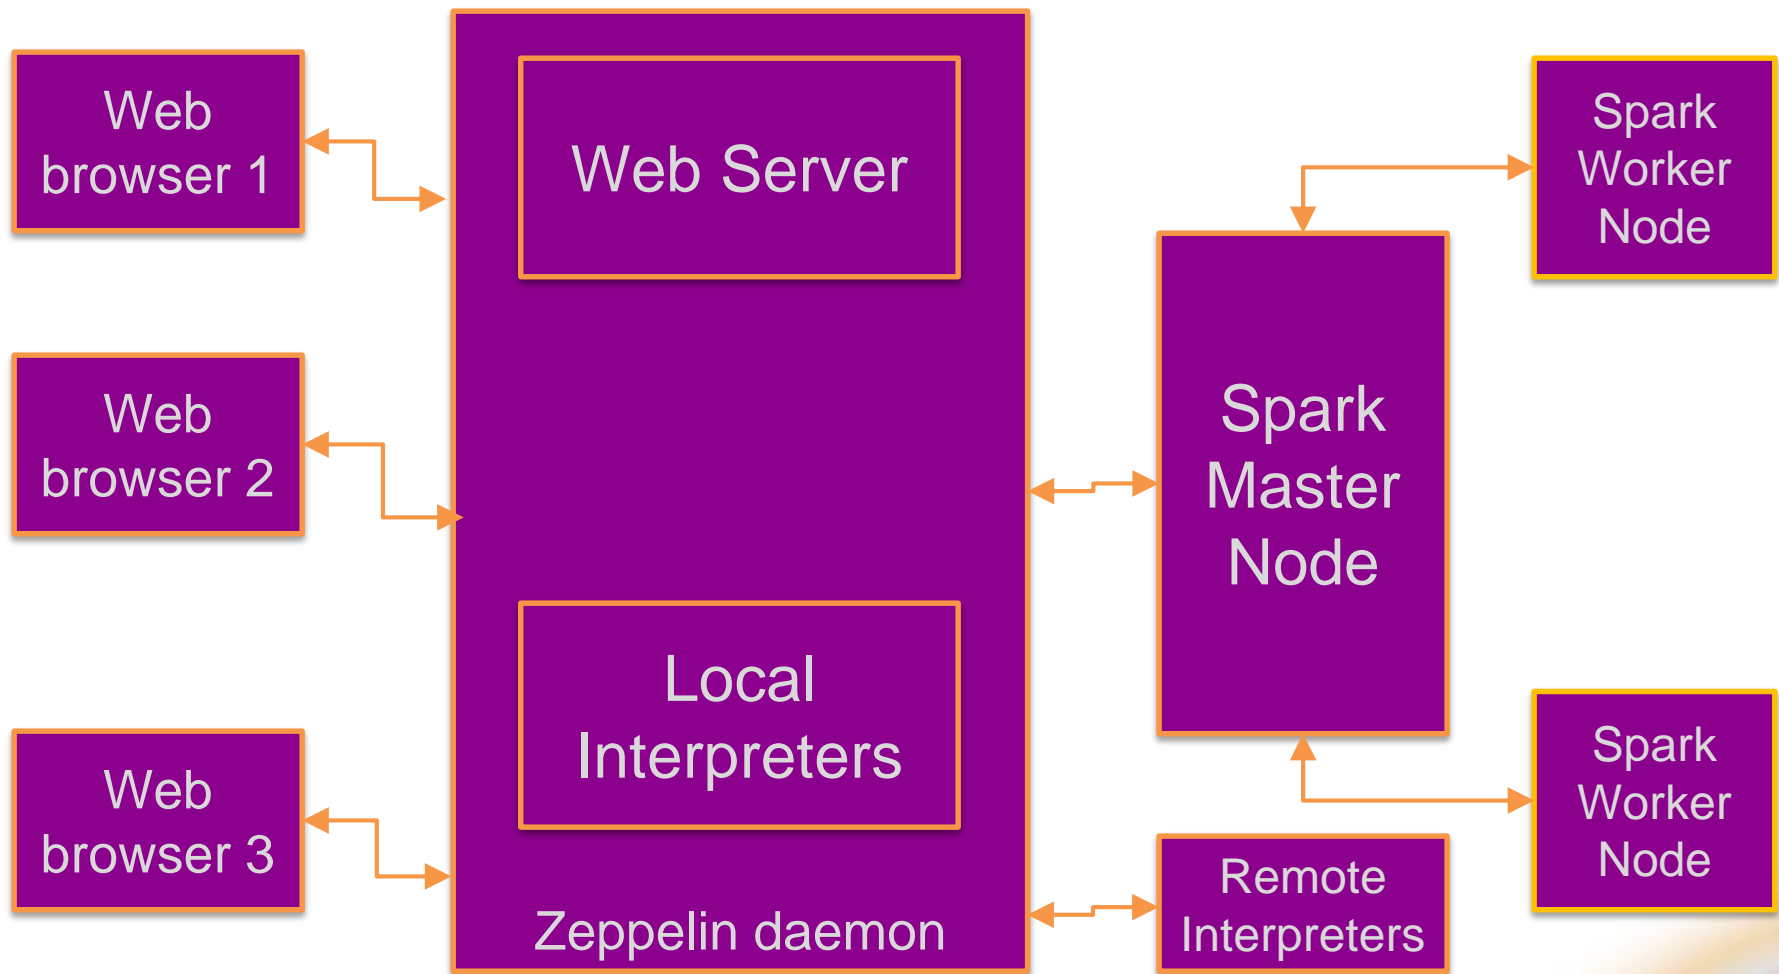


# COMBINING SPARK AND ZEPPELIN

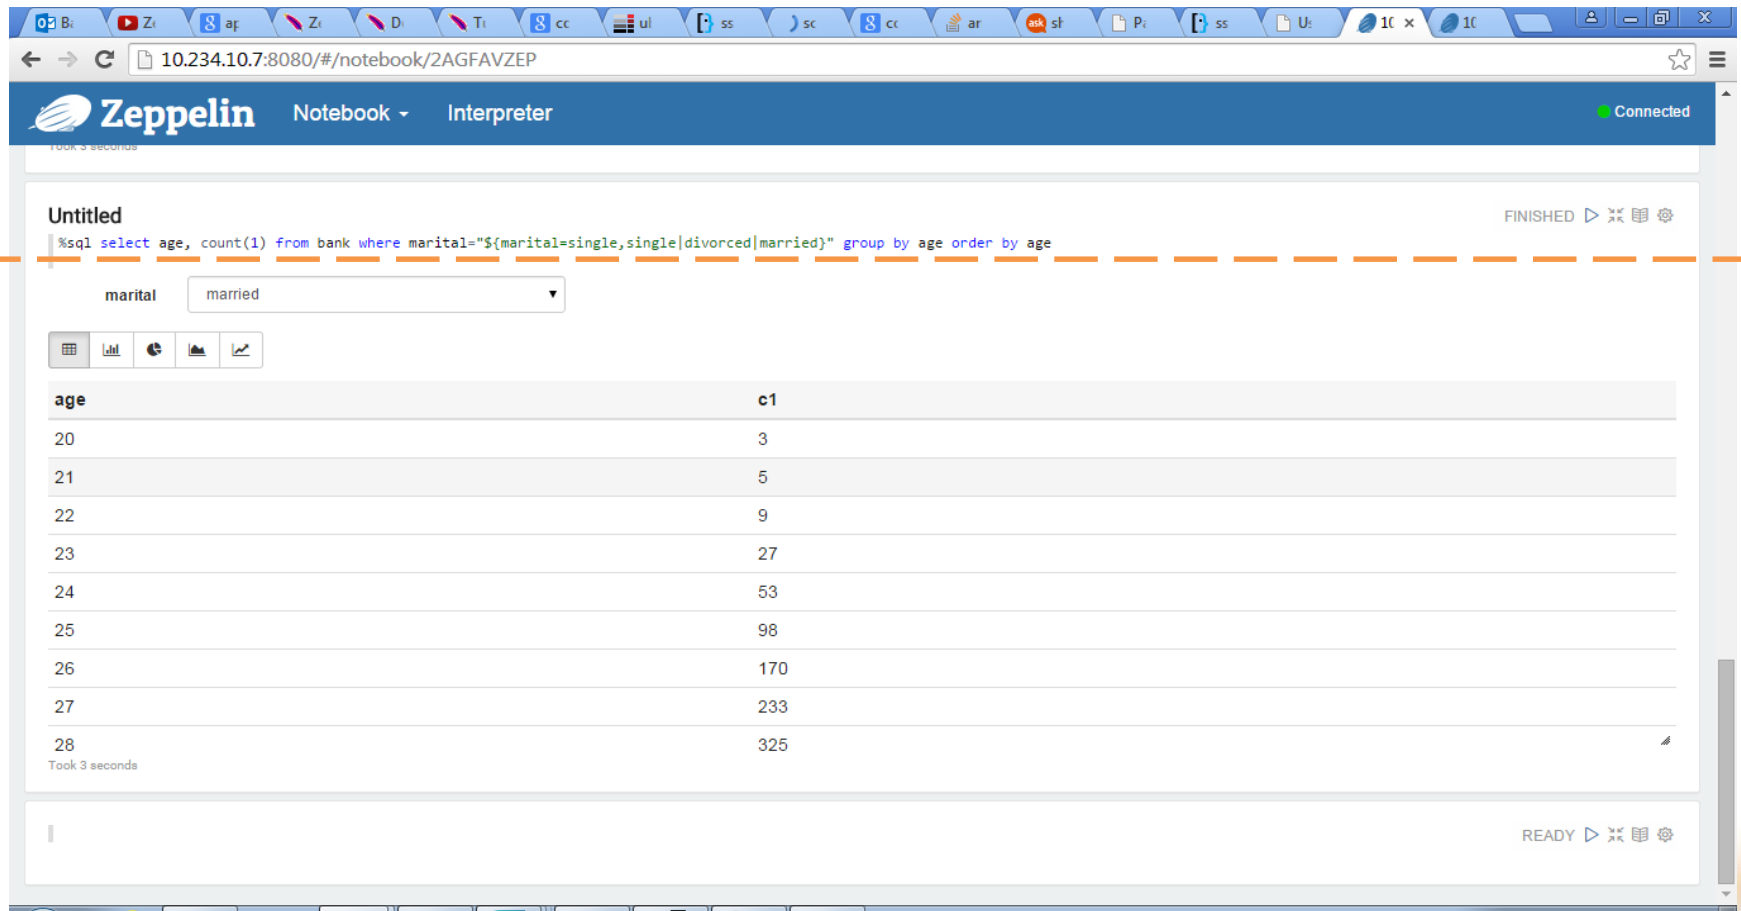




# Spark and Zeppelin



# Zeppelin views: Table from SQL



The screenshot shows the Zeppelin Notebook interface. The top bar includes the Zeppelin logo, 'Notebook' and 'Interpreter' tabs, and a 'Connected' status indicator. The main area displays an 'Untitled' notebook with a SQL query: `%sql select age, count(1) from bank where marital='${marital=single,single|divorced|married}' group by age order by age`. Below the query, a dropdown menu for 'marital' is set to 'married'. The results are displayed in a table view with two columns: 'age' and 'c1'. The table contains 9 rows of data. The status 'FINISHED' is shown in the top right of the notebook area. At the bottom of the notebook area, the status 'READY' is shown.

| age | c1  |
|-----|-----|
| 20  | 3   |
| 21  | 5   |
| 22  | 9   |
| 23  | 27  |
| 24  | 53  |
| 25  | 98  |
| 26  | 170 |
| 27  | 233 |
| 28  | 325 |

# Zeppelin views: Table from SQL

```
%sql select age, count(1) from bank where
marital="${marital=single,single|divorced|married}"
group by age order by age
```

Untitled

```
%sql select age, count(1) from bank where marital="${marital=single,single|divorced|married}" group by age order by age
```

marital

married



| age | c1  |
|-----|-----|
| 20  | 3   |
| 21  | 5   |
| 22  | 9   |
| 23  | 27  |
| 24  | 53  |
| 25  | 98  |
| 26  | 170 |
| 27  | 233 |
| 28  | 325 |

Took 3 seconds

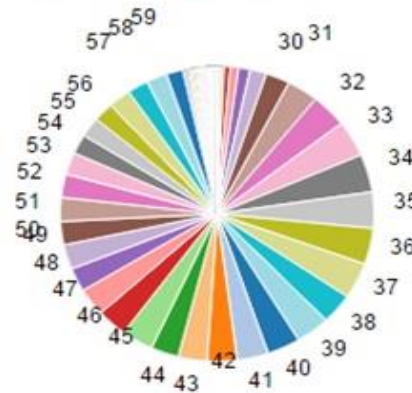

# Zeppelin views: Pie chart from SQL

Untitled

```
%sql select age, count(1) from bank where marital="{marital=single,single|divorced|married}" group by age order by age
```

marital

married



Took 3 seconds


# Zeppelin views: Bar chart from SQL

Untitled

```
%sql select age, count(1) from bank where marital="{marital=single,single|divorced|married}" group by age order by age
```

marital

married



# Zeppelin views: Angular

```
%angular
Write some text in textbox:
<input type="text" ng-model="sometext">

<h1>Hello {{ sometext }}</h1>
```

Write some text in textbox:

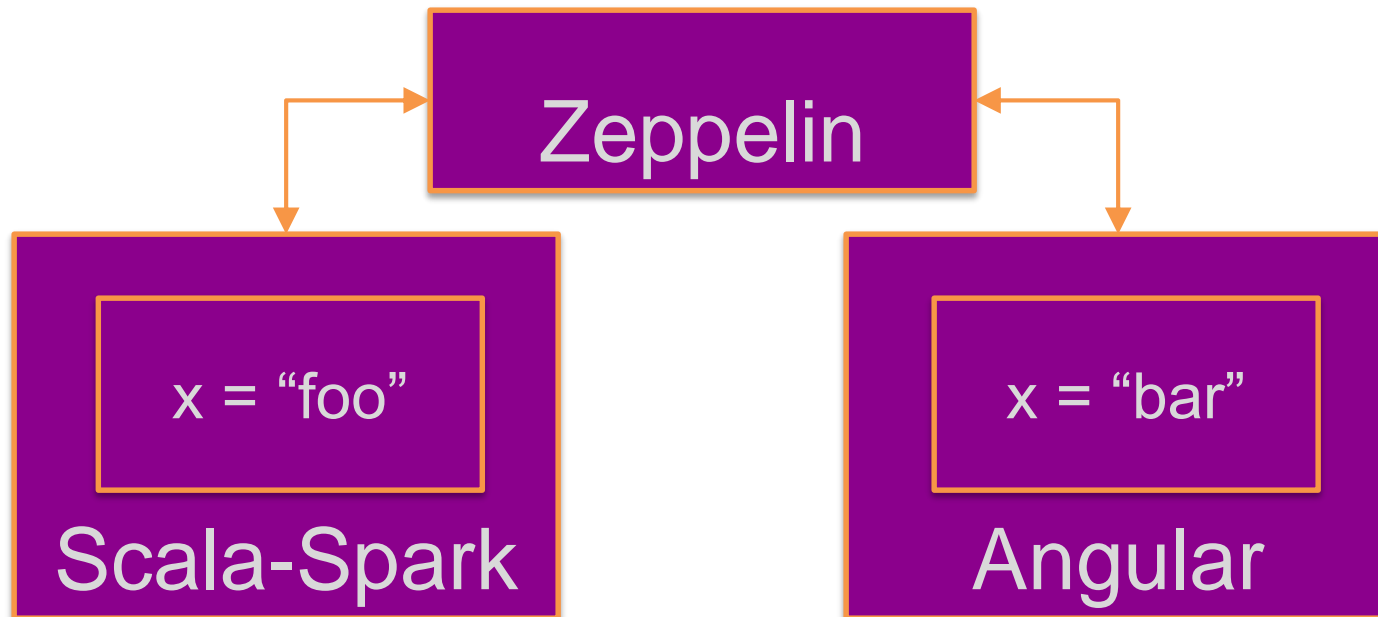
## Hello Big Data

Took 2 seconds



# Share variables: MVVM

- Between Scala/Python/Spark and Angular
- Observe scala variables from angular



# Zeppelin views: Angular-scala binding

```
println("%html <h3>Hello</h3>")
```

FINISHED ▶ ⌵ ⌶ ⚙

Hello

Took 0 seconds

```
println("%angular <h3>Hello</h3>")
```

FINISHED ▶ ⌵ ⌶ ⚙

Hello

Took 1 seconds

```
println("%angular <h3>Hello {{name}}</h3>")
```

FINISHED ▶ ⌵ ⌶ ⚙

Hello Zeppelin!!!!

Took 1 seconds

```
z.angularBind("name", "Zeppelin!!!!")
```

Took 0 seconds

```
println("""%angular <h3>Hello {{name}}</h3>
<button class="btn btn-success"
 ng-click="name='Angular'">click here</button>""")
```

Hello Zeppelin!!!!

click here

Took 1 seconds

# Screen sharing using Zeppelin

- Share your graphical reports
  - Live sharing
  - Get the share URL from zeppelin and share with others
  - Uses websockets
- Embed live reports in web pages

# FUTURE



# Spark and Zeppelin

- Spark
  - Berkeley Data Analytics Stack
  - More source and sinks; SparkSQL
- Zeppelin
  - Notebooks for
    - Machine Learning using Spark
    - GraphX and Mllib
  - Additional interpreters
  - Better graphics, streaming views
  - Report persistence
  - More report templates
  - Better angular integration



# SUMMARY





# Summary

- Spark and tools
- The need for visualization
- The role of Zeppelin
- Zeppelin – Spark integration

# Thank you!

@prajods

saltmarsh  
by

GREAT INDIAN  
**DEVELOPER**  
**SUMMIT**

