

Cost-based query optimization in Apache Hive



Julian Hyde
June 4th, 2014

About me

Julian Hyde

Architect at Hortonworks

Open source:

- Founder & lead, Apache Optiq (query optimization framework)
- Founder & lead, Pentaho Mondrian (analysis engine)
- Committer, Apache Drill
- Contributor, Apache Hive
- Contributor, Cascading Lingual (SQL interface to Cascading)

Past:

- SQLstream (streaming SQL)
- Broadbase (data warehouse)
- Oracle (SQL kernel development)

(Thanks to
John Pullokkaran,
Harish Butani
for presentation content
and actually doing the work.)

Apache Hive

The original “SQL on Hadoop”

Undergoing extensive renovation

- Tez execution engine
- YARN execution environment
- Vectorized data representation
- Column-oriented data storage (ORC)
- ACID transactions
- SQL standards compliance
- SQL authorization model
- **Cost-based query optimization (CBO)**

“Stinger Initiative”




What? Why? How? When?

Incremental cutover to cost-based optimization

Release	Date	Remarks
Apache Hive 0.12	October 2013	<ul style="list-style-type: none">• Rule-based Optimizations• No join reordering• Main optimizations: predicate push-down & partition pruning• Semantic info and operator tree tightly coupled
Apache Hive 0.13	April 2014	<p>“Old-style” JOIN & push-down conditions: <i>... FROM t1, t2 WHERE ...</i></p> <p>CBO just missed the deadline ☹</p>
HDP 2.1	April 2014	<p>Cost-based ordering of joins</p> <ul style="list-style-type: none">• HIVE-6439 “Introduce CBO step in Semantic Analyzer”• HIVE-5775 “Introduce Cost Based Optimizer in Hive”
Apache Hive 0.14	?	<p>CBO patches</p> <p>More rework of internals</p> <p>More cost-based features...</p>

Apache Optiq (incubating)



Apache Optiq

Apache incubator project since May, 2014

Query planning framework

- Extensible
- Usable standalone (JDBC) or embedded

Adoption

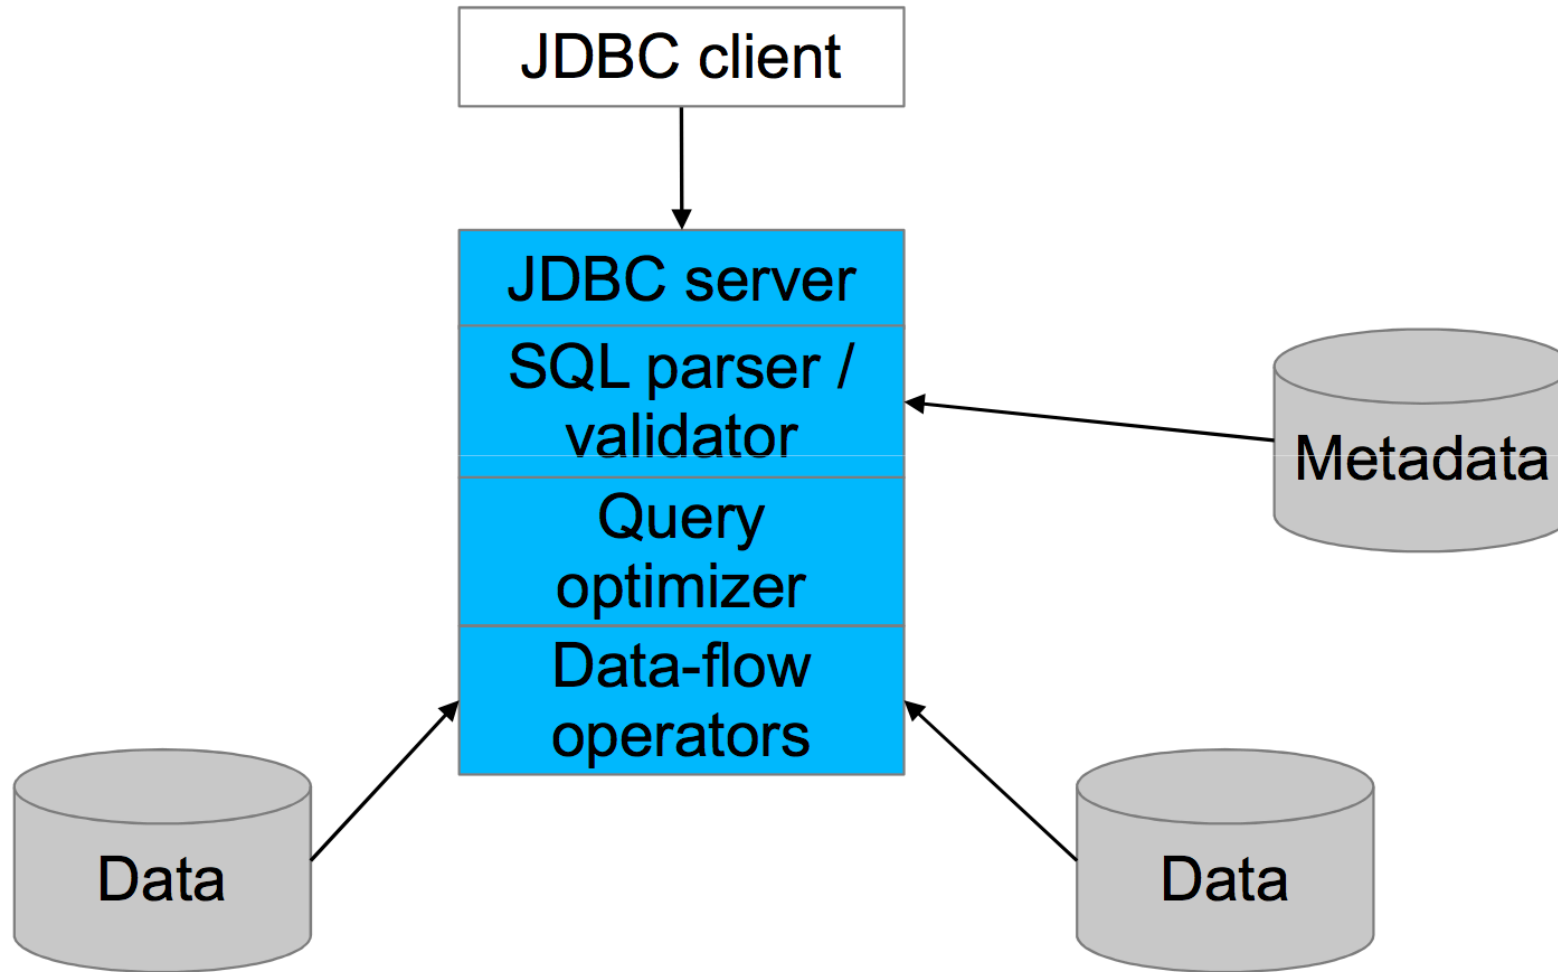
Lingual – SQL interface to Cascading

Apache Drill

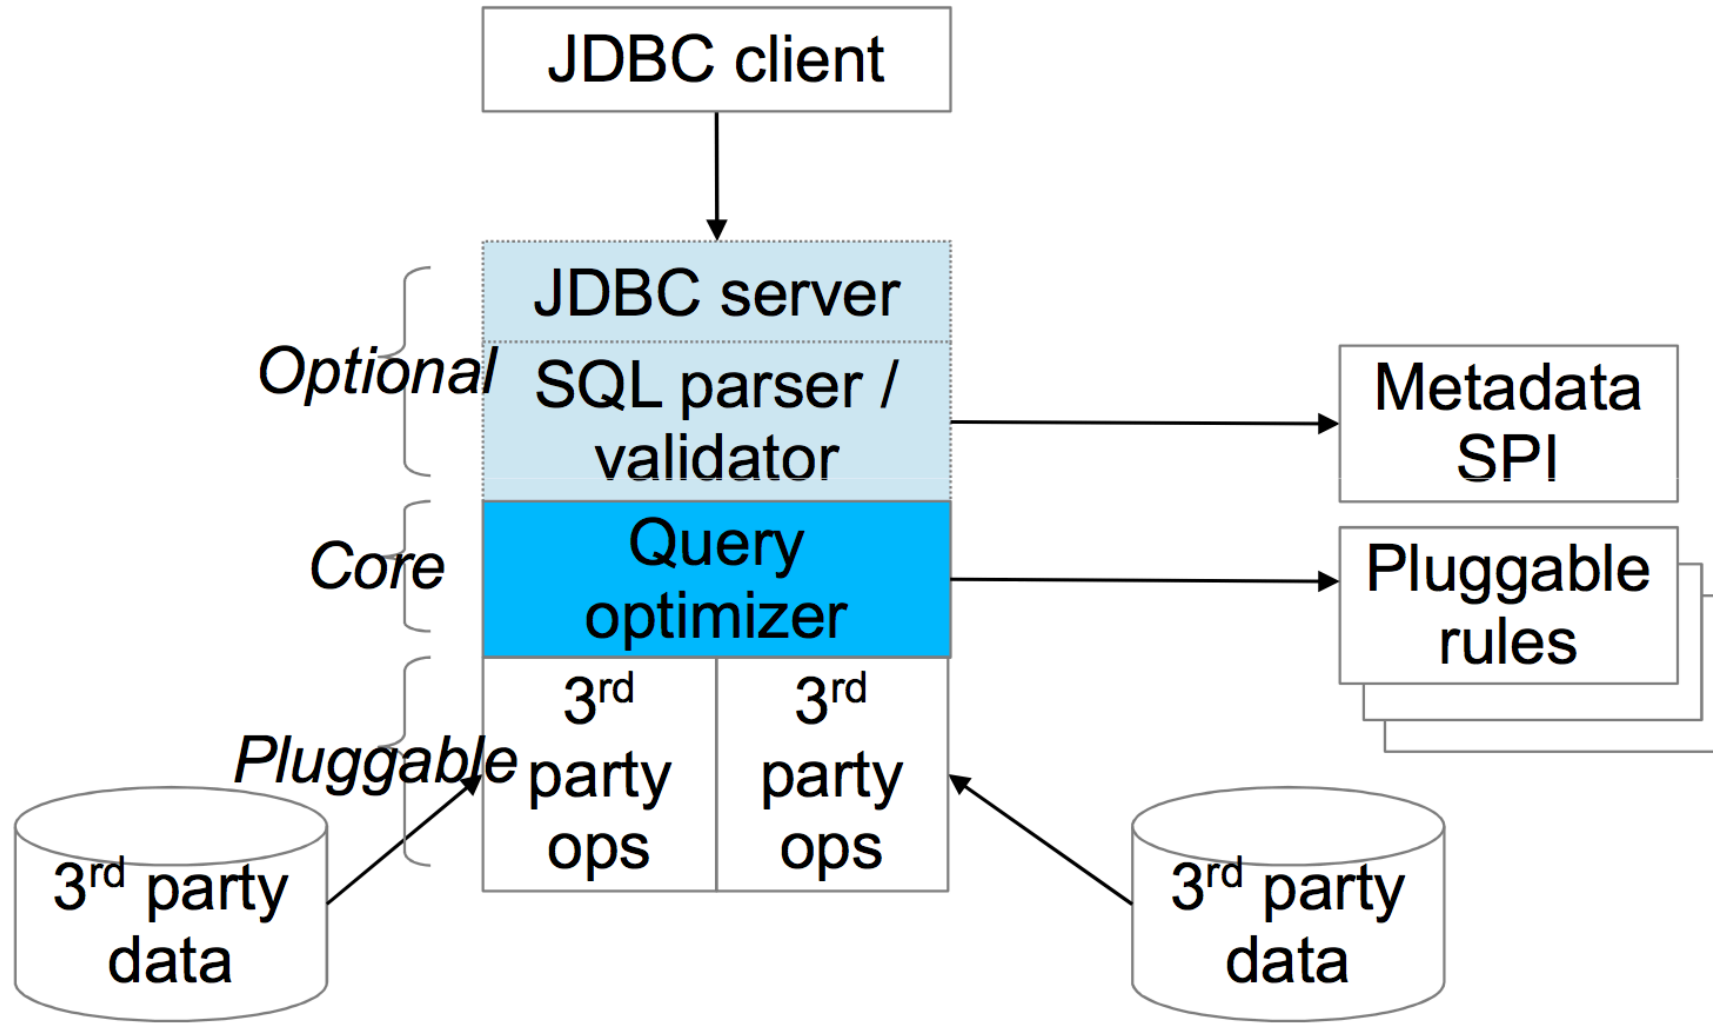
Apache Hive

Adapters: Splunk, Spark, MongoDB, JDBC, CSV, Web tables, In-memory data

Conventional DB architecture



Optiq architecture



Optiq – APIs and SPIs

Relational algebra

RelNode (operator)

- TableScan
- Filter
- Project
- Union
- Aggregate
- ...

RelDataType (type)

RexNode (expression)

RelTrait (physical property)

- RelConvention (calling-convention)
- RelCollation (sortedness)
- TBD (bucketedness/distribution)

Transformation rules

RelOptRule

- MergeFilterRule
- PushAggregateThroughUnionRule
- RemoveCorrelationForScalarProjectRule
- 100+ more

Unification (materialized view)

Column trimming

Cost, statistics

RelOptCost

RelOptCostFactory

RelMetadataProvider

- RelMdColumnUniqueness
- RelMdDistinctRowCount
- RelMdSelectivity

SQL parser

SqlNode

SqlParser

SqlValidator

Metadata

Schema


Table

Function

- TableFunction
- TableMacro

JDBC driver

Now... back to Hive



CBO in Hive

Why cost-based optimization?

Ease of Use – Join Reordering

View Chaining

Ad hoc queries involving multiple views

Enables BI Tools as front ends to Hive

First version

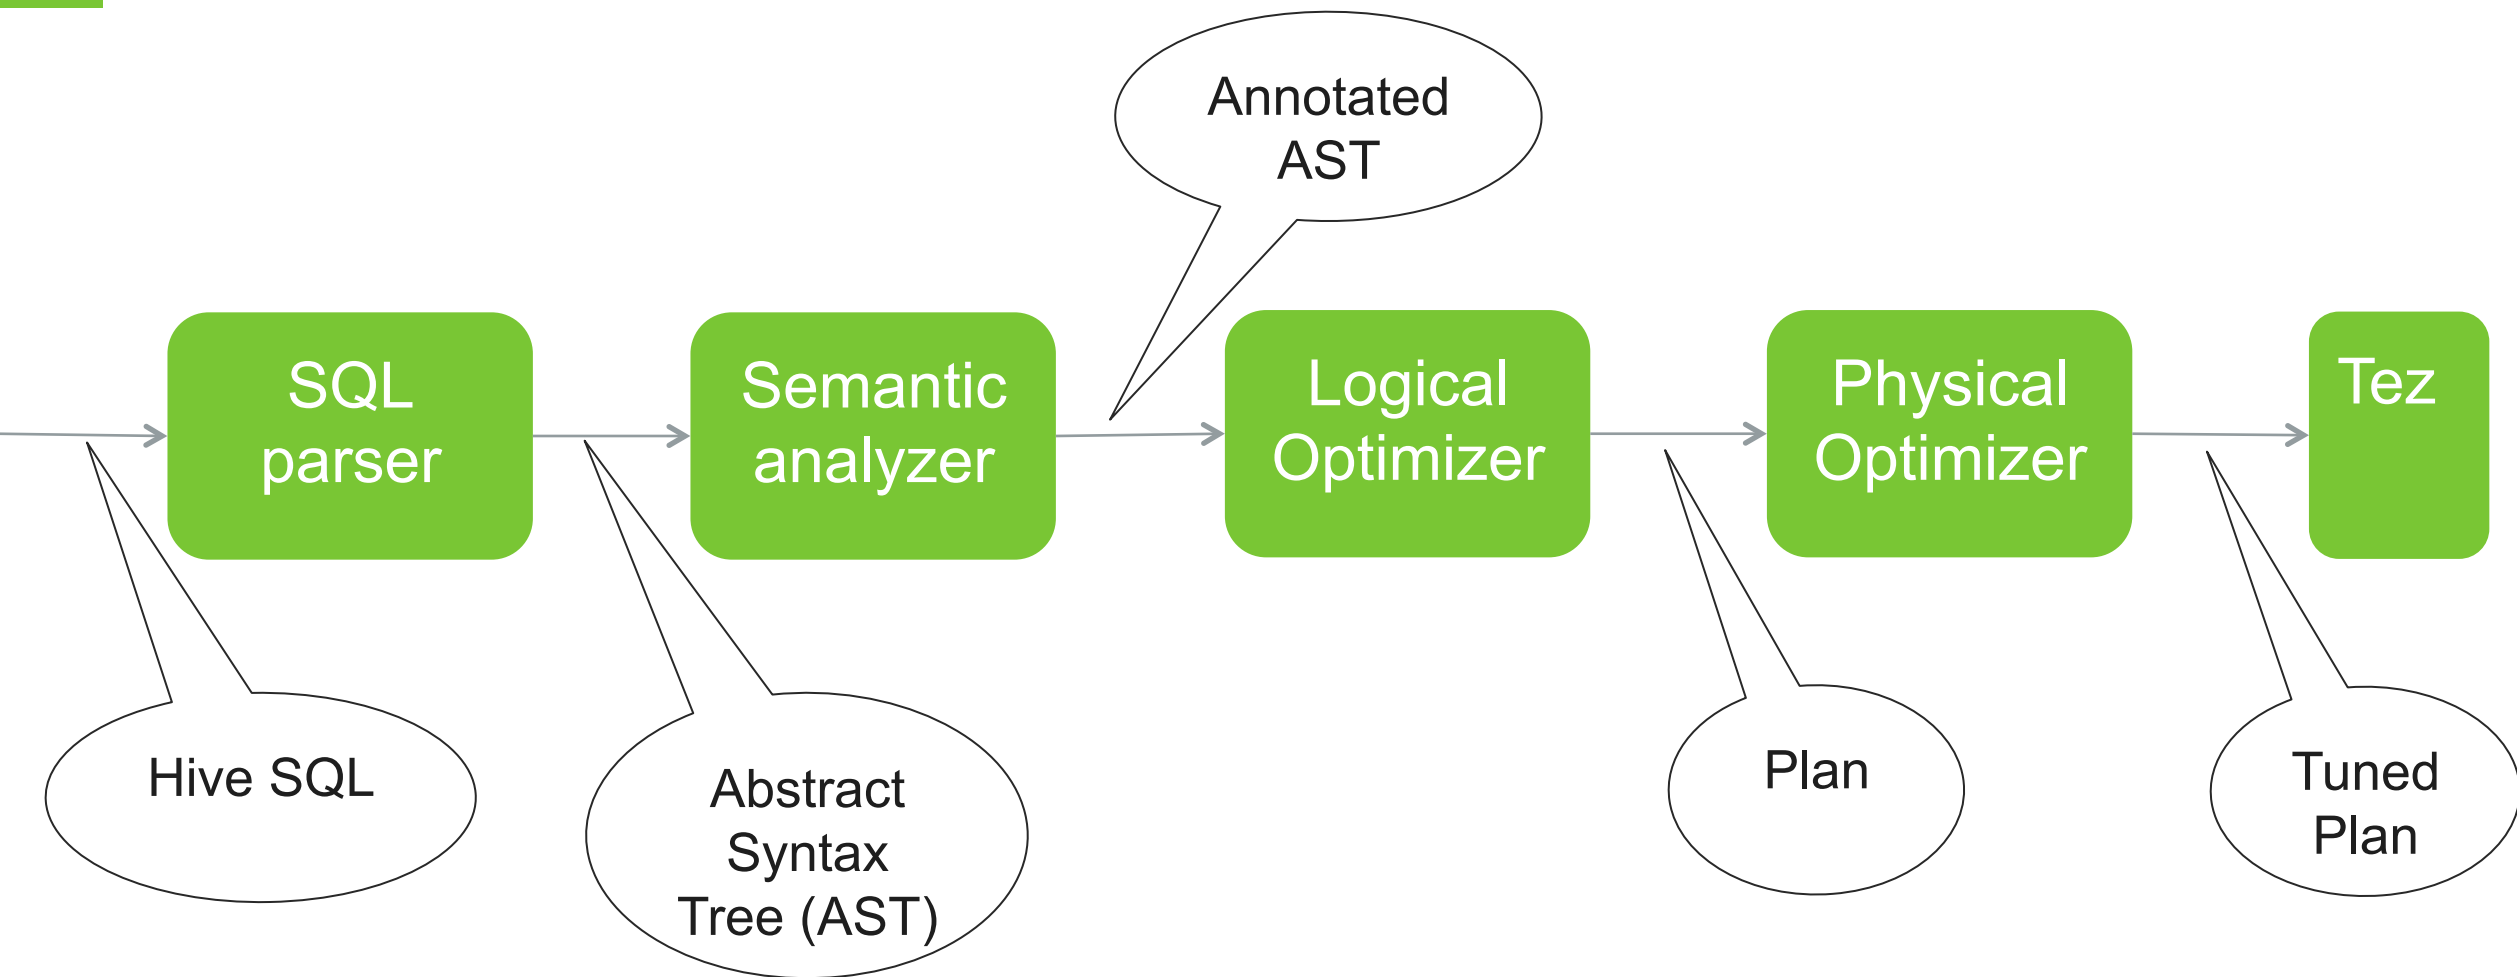
Modest goal

Concrete results


Join re-ordering



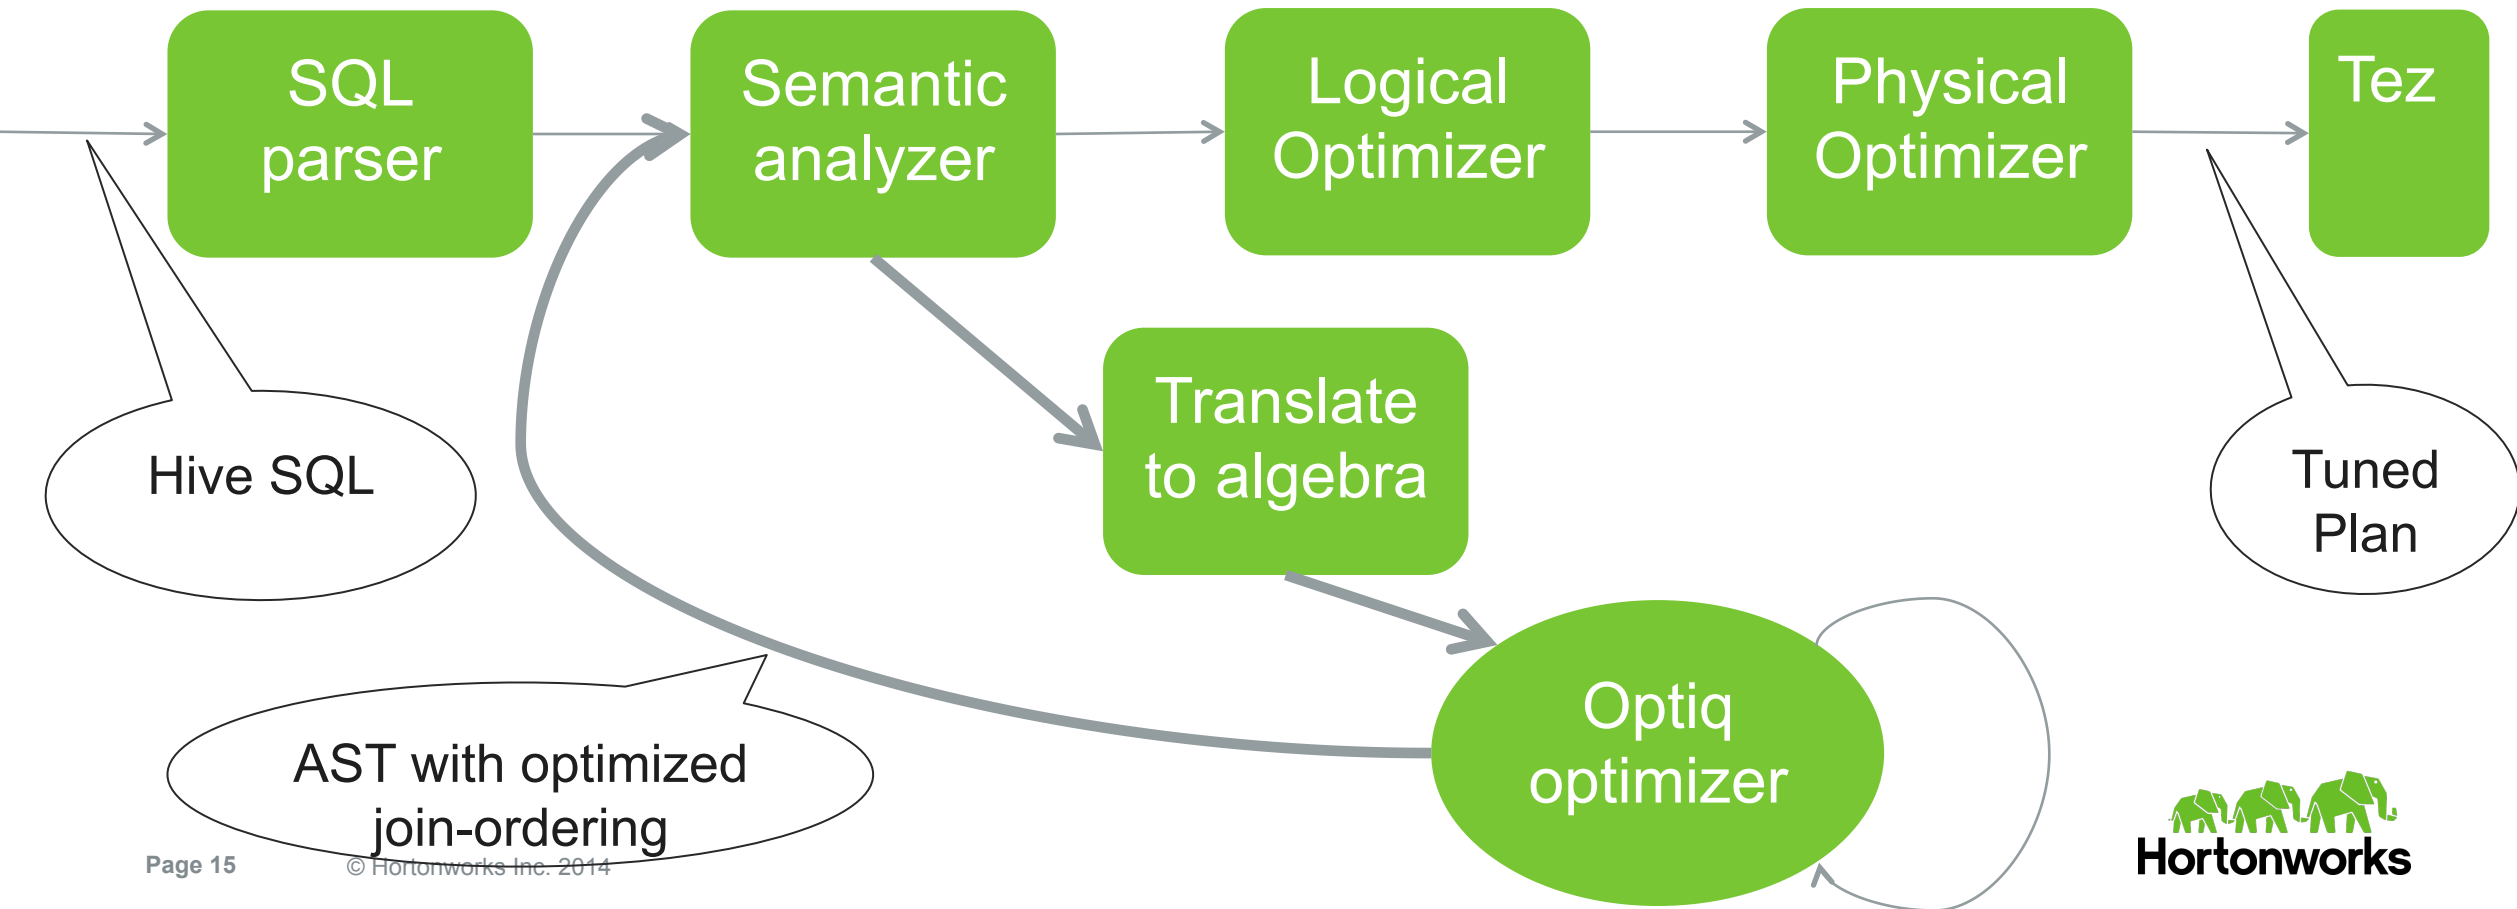
Query preparation – Hive 0.13



Query preparation – full CBO





Query preparation – initial CBO







Query Execution – The basics

```
SELECT R1.x
FROM R1
  JOIN R2 ON R1.x = R2.x
  JOIN R3 on R1.x = R3.x AND R2.x = R3.x
WHERE R1.z > 10;
```

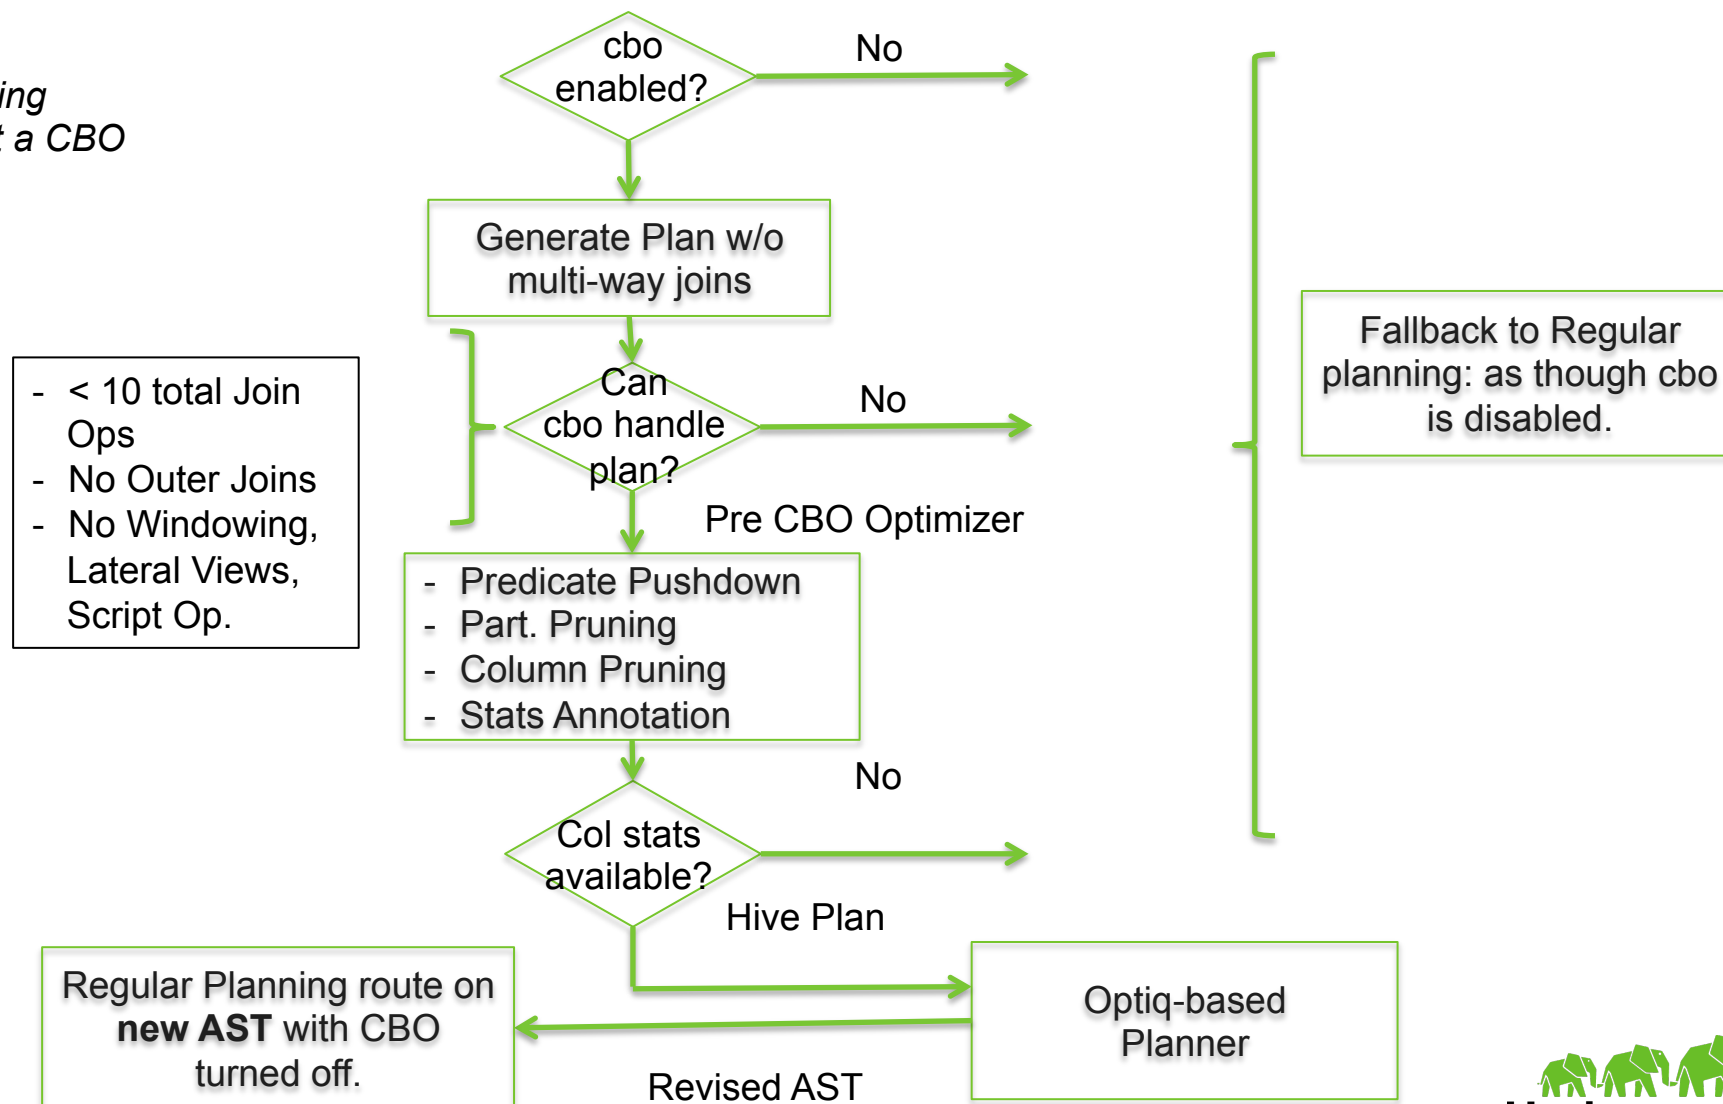


Query Optimization – Rule Based vs. Cost Based

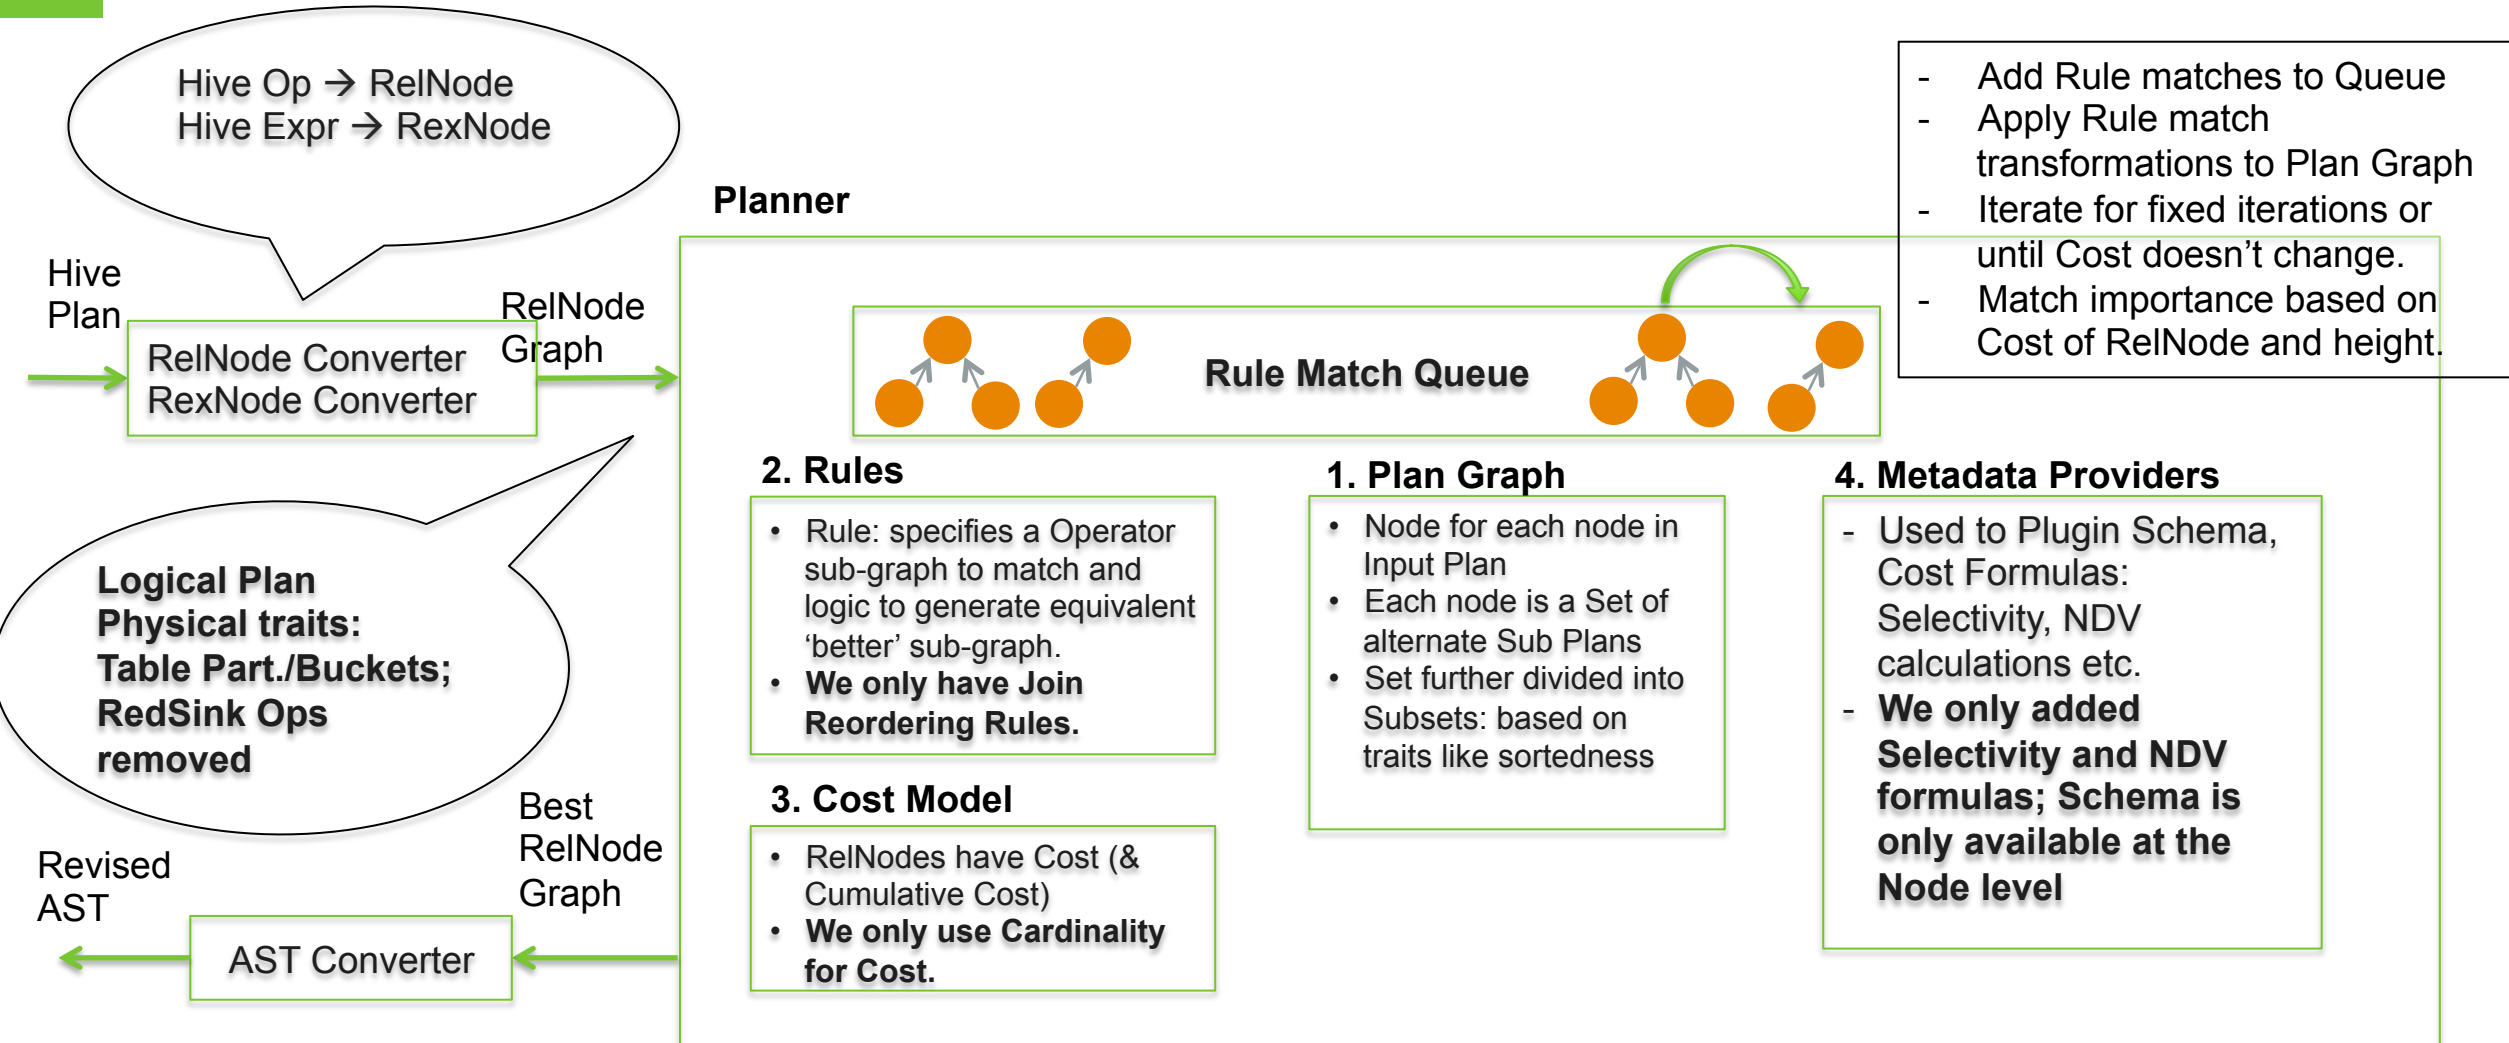


Introduction of CBO into Hive Planning

Series of gating factors to get a CBO Plan.




Optiq Planner Process




Join Reordering Rules

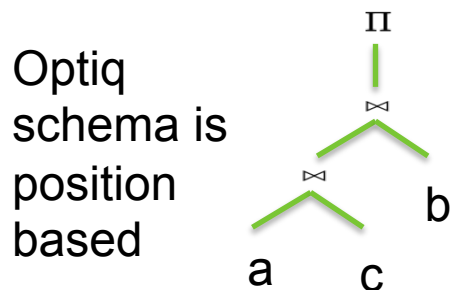
1. Swap Join Rule




2. Push Join Through Join Rule




but is really:



3. So



4. Pull Up Project above Join



5. Merge Projects

Added bonus
Join permutations across sub-query blocks

Summary

Join re-ordering

Join cardinality is used for cost

All other operators are assumed to have tiny cost

Cardinality of filter, join, group-by is based on selectivity

Selectivity is computed based on number-of-distinct-values (NDV)

Table Stats and Column stats are required

Current limitations

Only supports: filter, inner join, group-by, project, order-by, limit

Not all UDFs

Does not attempt all join permutations (e.g. bushy trees; 10-way joins or more)

May not work well for Bucket, SMB & Skew Joins

TPC-DS Query 50


Joins Store Sales, and Store Returns fact tables.

Each of the fact tables are independently restricted by date.

Analysis at Store grain, so this dimension also joined in.



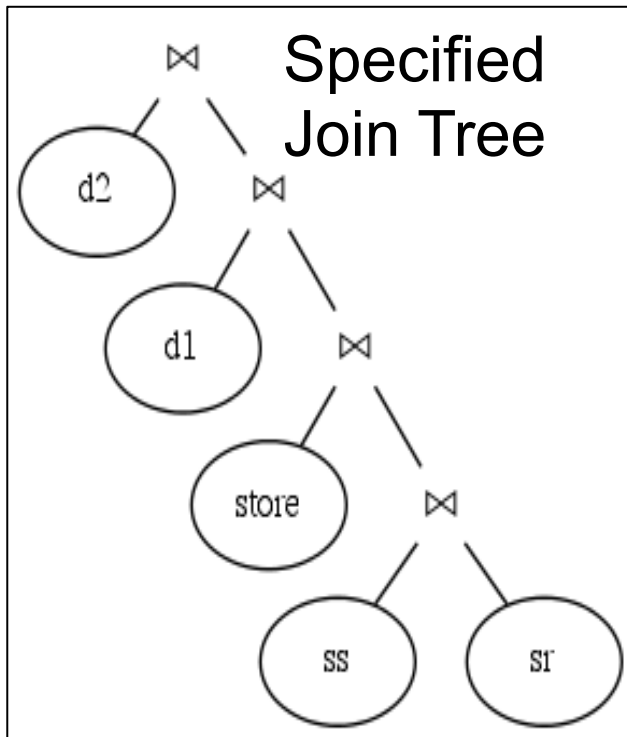
As specified Query starts by joining the 2 Fact tables.

Join Graph



```
select
  s_store_name , .. other store details
, sum(case when (sr_returned_date_sk - ss_sold_date_sk <= 30 ) then 1 else 0 end) as `30 days`, ...
from
  store_sales ss, store_returns sr, store s , date_dim d1 , date_dim d2
where
  d2.d_year = 2000 and d2.d_moy = 9
and ss.ss_ticket_number = sr.sr_ticket_number and ss.ss_item_sk = sr.sr_item_sk
and ss.ss_sold_date_sk = d1.d_date_sk and sr.sr_returned_date_sk = d2.d_date_sk
and ss.ss_customer_sk = sr.sr_customer_sk and ss.ss_store_sk = s.s_store_sk
group by store details
order by store details limit 100;
```

TPC-DS Query 50



TPC-DS Query 50

Facts restricted to 3 months

	Run 1	Run 2
Non CBO	53.1	53.4
CBO	22.5	21.9

- ❖ 1 year test
 - > 10 mins for Non CBO
 - CBO time was about the same
- ❖ Fact tables
 - partitioned by Day,
 - bucketed by Item
- ❖ Bucketing off
 - Bucketing should help CBO plan.
 - SR table much smaller. Better chance of Bucket Join in place of Shuffle Join.

Orderings considered by Planner

Join Ordering	Cost Estimate
['d2', [[['store_sales', 'd1'], 'store_returns'], 'store']]	515074768.659
['d1', [[['store_sales', 'store'], 'store_returns'], 'd2']]	448155.355
...	
['store_returns', 'd2']	9938.93
['store_sales', 'store_returns']	156727295.634
['d1', 'store_sales']	123675664.449

TPC-DS Query 17


Joins Store Sales, Store Returns and Catalog Sales fact tables.

Each of the fact tables are independently restricted by time.

Analysis at Item and Store grain, so these dimensions are also joined in.

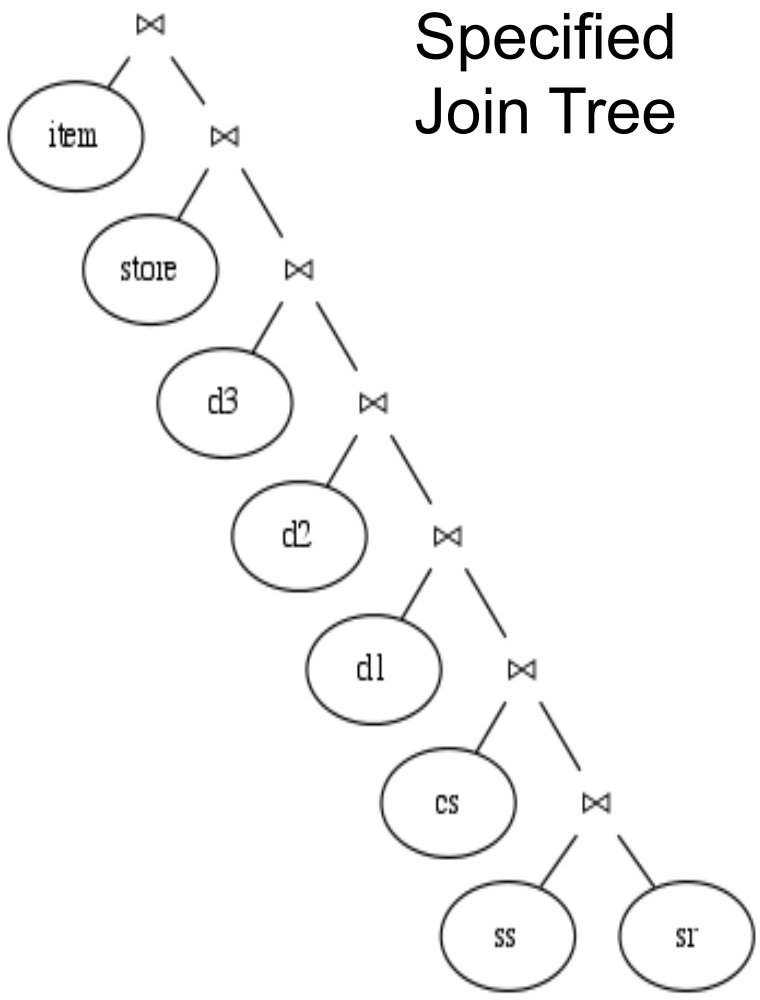
As specified Query starts by joining the 3 Fact tables.

```
select i_item_id
,i_item_desc
,s_state
,count(ss_quantity) as store_sales_quantitycount
,....
from store_sales ss ,store_returns sr, catalog_sales cs,
date_dim d1, date_dim d2, date_dim d3, store s, item i
where d1.d_quarter_name = '2000Q1'
and d1.d_date_sk = ss.ss_sold_date_sk
and i.i_item_sk = ss.ss_item_sk and ...
group by i_item_id ,i_item_desc, ,s_state
order by i_item_id ,i_item_desc, s_state
limit 100;
```

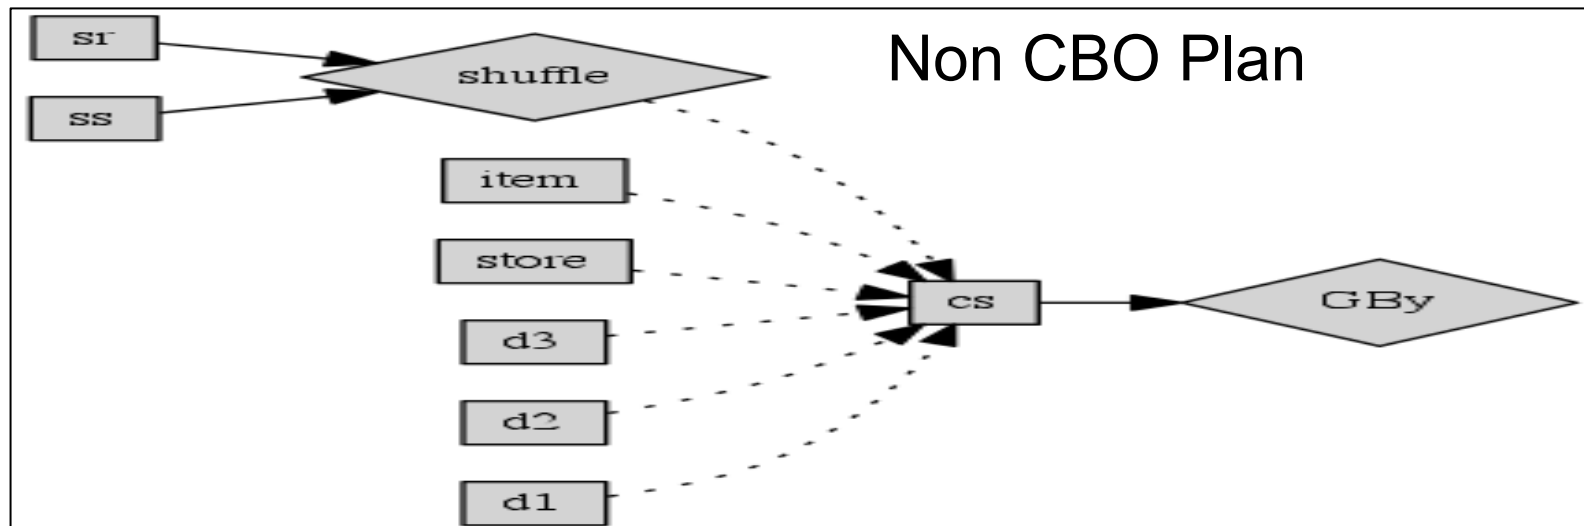


TPC-DS Query 17

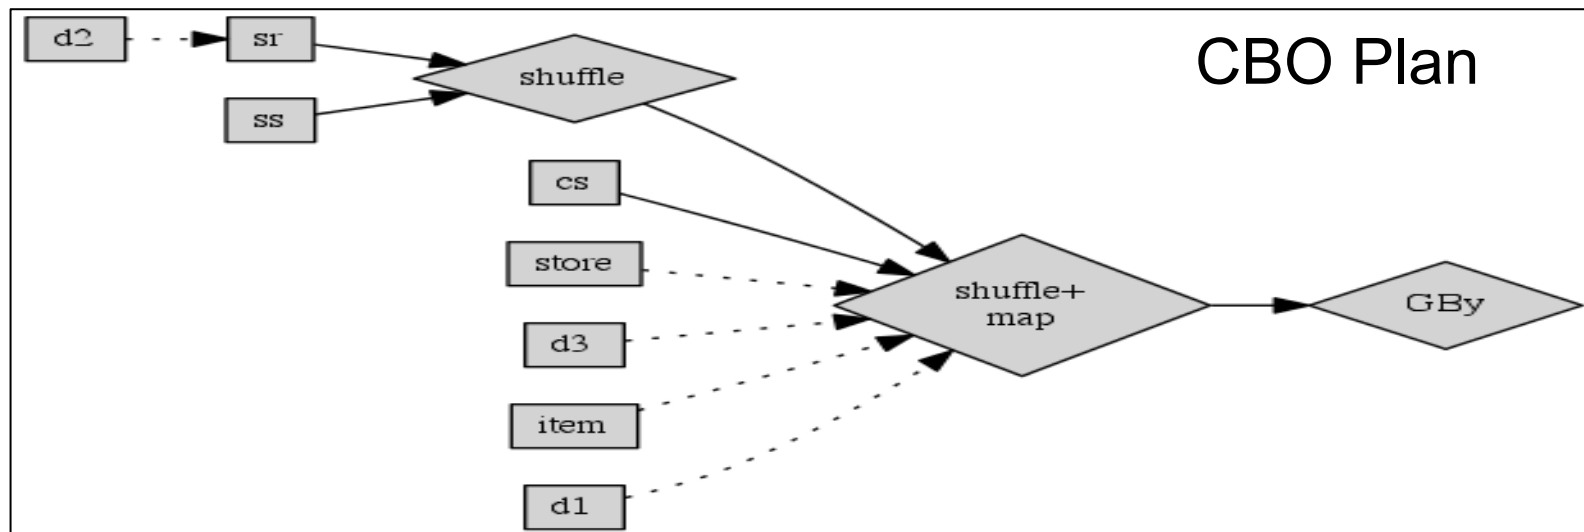
Specified
Join Tree



Non CBO Plan



CBO Plan



TPC-DS Query 17

Facts restricted to 3 months

	Run 1	Run 2
Non CBO	100.71	127.53
CBO	50.9	44.52

- ❖ 1 year test
 - > 10 mins for Non CBO
 - CBO time was about the same
- ❖ Fact tables
 - partitioned by Day,
 - bucketed by Item
- ❖ Bucketing off
 - Bucketing should help CBO plan.
 - SR table much smaller. Better chance of Bucket Join in place of Shuffle Join.

Orderings considered by Planner

Join Ordering	Cost Estimate
['item', [[[[['d2', 'store_returns'], 'store_sales'], 'catalog_sales'], 'd1'], 'd3'], 'store']]	3547898.061
...	
['store_returns', 'd2']	19224.71
['store_sales', 'store_returns']	23057497.991
['d1', 'store_sales']	26142.943

Next?

Outer joins

Scale to larger numbers of joins

Support all expressions (UDFs)

Join algorithm selection

Sortedness & distribution as a trait

Trait propagation

Better cost model

More statistics

Move all pre-planning and logical planning to Optiq

Use Optiq costs/statistics to help physical planning

Constant reduction & tree pruning

Rewrite query to use materialized view

Thank you!



@julianhyde

<http://hive.apache.org/>

<http://incubator.apache.org/projects/optiq.html>