



How Impala Works

Yue Chen

<http://linkedin.com/in/yuechen2>

<http://dataera.wordpress.com>

What' s Impala?



This is Impala...



Goal of Impala

- A general SQL engine for distributed systems, supporting both OLTP and OLAP.
- Interactive (real-time) queries.
- Built on top of HDFS and HBase.
- Engine is written in C++, fast.
- The database execution engine is like that of massively parallel processing (MPP) databases, not using MapReduce.

What's Impala?

- In-memory, distributed SQL query engine (no MapReduce)
- Native backend code (C++)
- Distributed on HDFS data nodes

What's Impala?

- **Interactive SQL**
 - Typically 5-65x faster than Hive (observed up to 100x faster)
 - Responses in seconds instead of minutes (sometimes sub-second)
- **Approx. ANSI-92 standard SQL queries with HiveQL**
 - Compatible SQL interface for existing Hadoop/CDH applications
 - Based on industry standard SQL
- **Natively on Hadoop/HBase storage and metadata**
 - Flexibility, scale, and cost advantages of Hadoop
 - No duplication/synchronization of data and metadata
 - Local processing to avoid network bottlenecks
- **Separate runtime from MapReduce**
 - MapReduce is designed and great for batch
 - Impala is purpose-built for low-latency SQL queries on Hadoop

Why Impala?



FAST!

Why HDFS?

- Low cost
- Reliability
- Easy to scale out

Architecture

Architecture Overview

- `impalad` daemon runs on HDFS nodes
- `statestored` for cluster metadata
- (Hive) `metastore` for database metadata
- Queries run on relevant nodes
- Data streamed to clients

Architecture Overview

- Submit queries via Hue/Beeswax, Thrift API, CLI, ODBC, JDBC
- No fault tolerance (query fails if any query on any node fails)
- Intermediate data never hits disk

statestored

- Acts as a cluster monitor
- Not a single point of failure

Metadata

- Uses Hive metastore
- Daemons cache metadata
- Can create tables in Hive or Impala

Impala Architecture Summary

- **impalad**
 - Runs on every node
 - Handles client requests
 - Handles query planning & execution
- **statestored**
 - Provides name service
 - Metadata distribution
 - Used for finding data
- **catalogd**
 - Relays metadata changes to all impalad' s

Impala: Architecture

Common Hive SQL and interface

SQL application
(Beeswax)

ODBC

Unified metadata store

Hive metastore

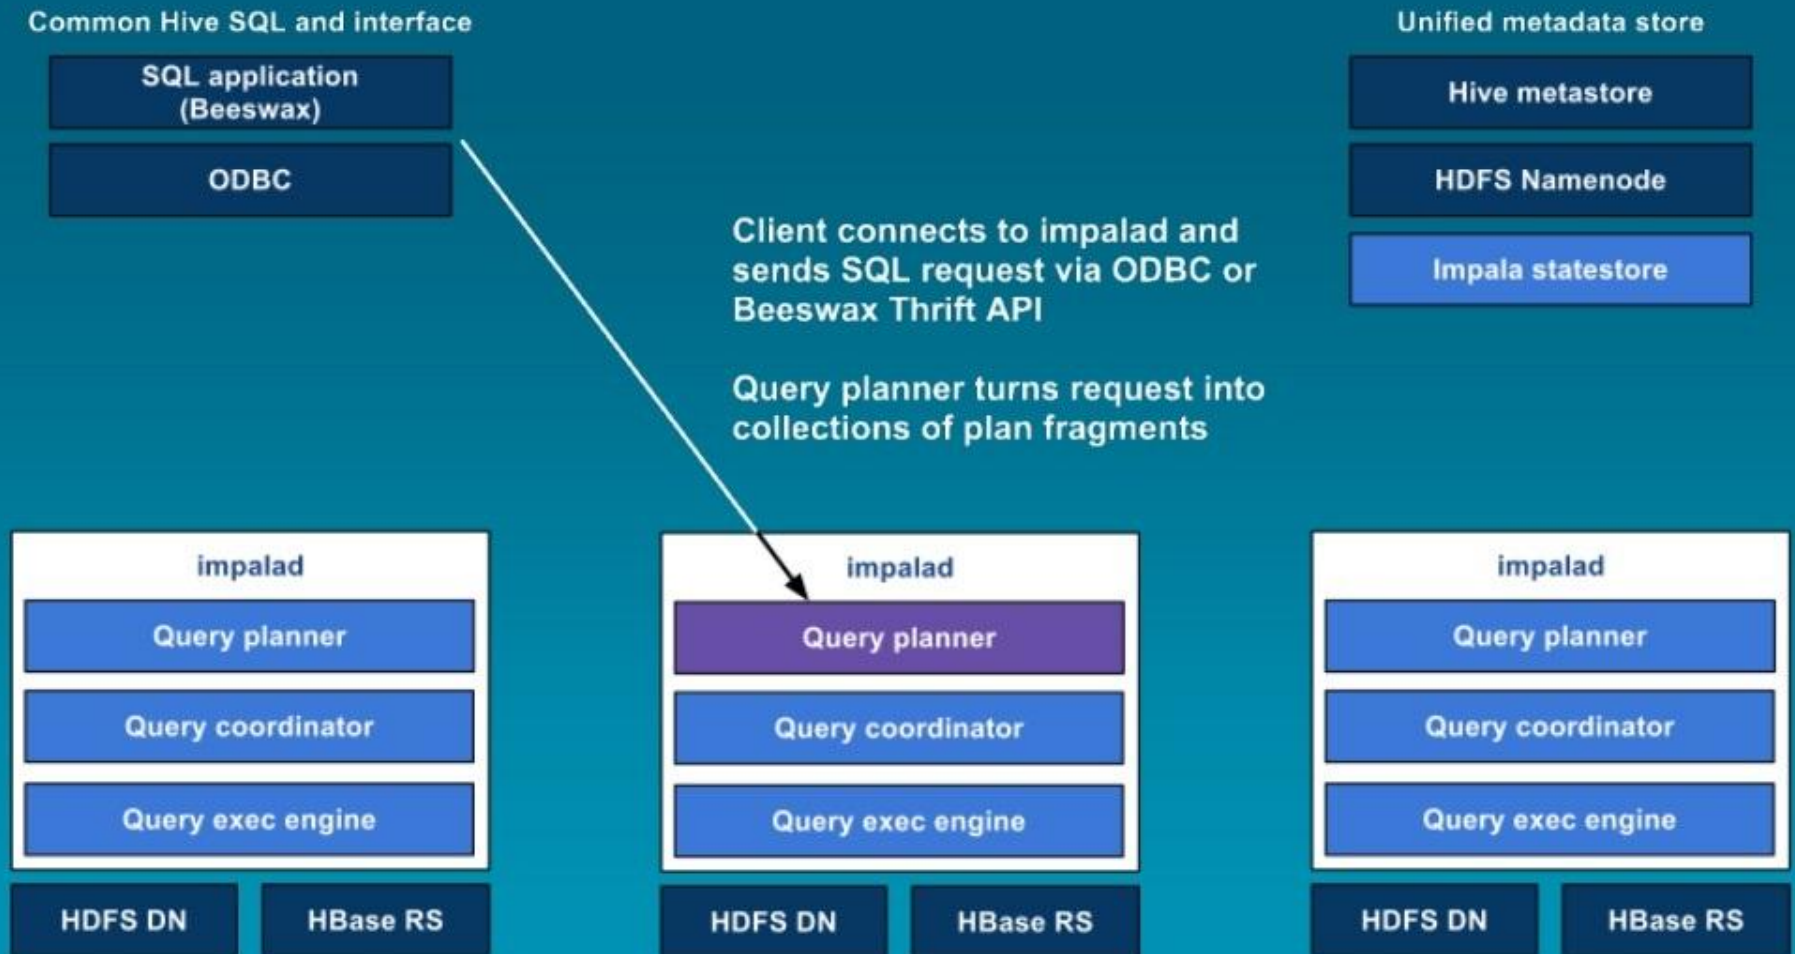
HDFS Namenode

Impala statestore

impalad's continually talk to
statestore to update their state
and to receive metadata to use
for query planning



Impala: Architecture



Impala: Architecture

Common Hive SQL and interface



Unified metadata store



Query coordinator initiates execution on remote impalad's



Impala: Architecture

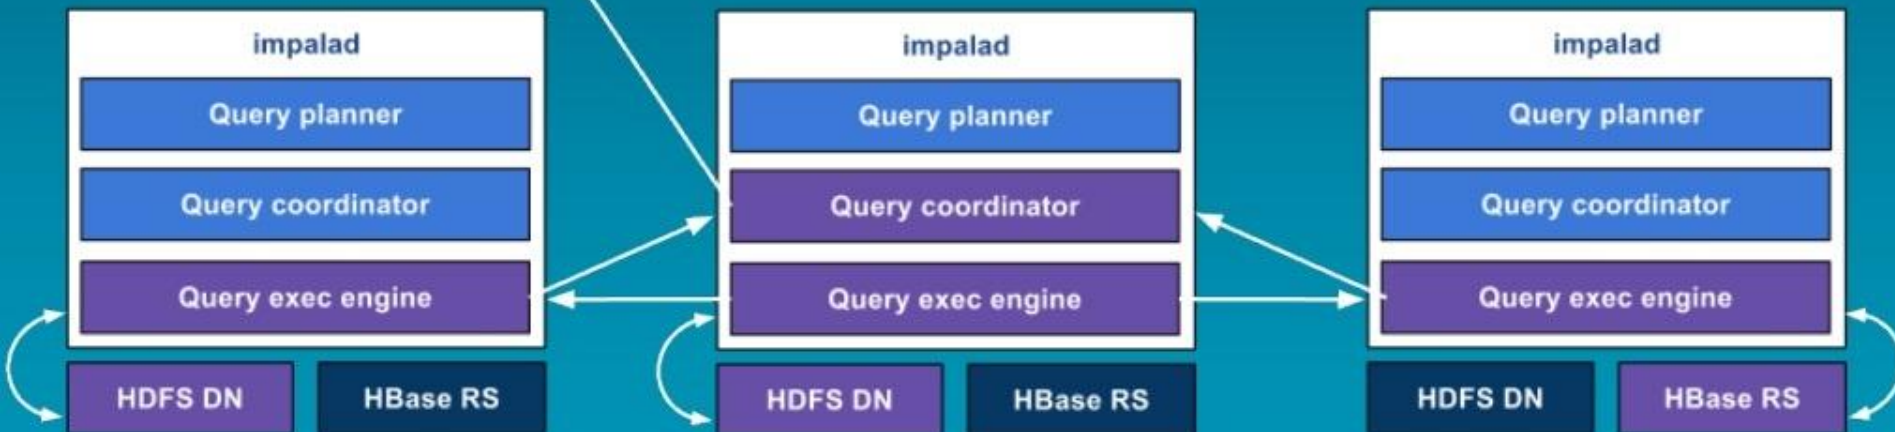
Common Hive SQL and interface



Unified metadata store



Intermediate results are streamed between impalad's, and query results are streamed back to the client



Impala Architecture: Query Execution Phases

- **Client** SQL arrives via ODBC/JDBC/Hue GUI/Shell
- **Planner** turns request into collections of plan fragments
- **Coordinator** initiates execution on impalad's local to data
- During **execution**:
 - intermediate results are streamed between executors
 - query results are streamed back to client
 - subject to limitations imposed to blocking operators (top-n, aggregation)

Query Planning: Overview

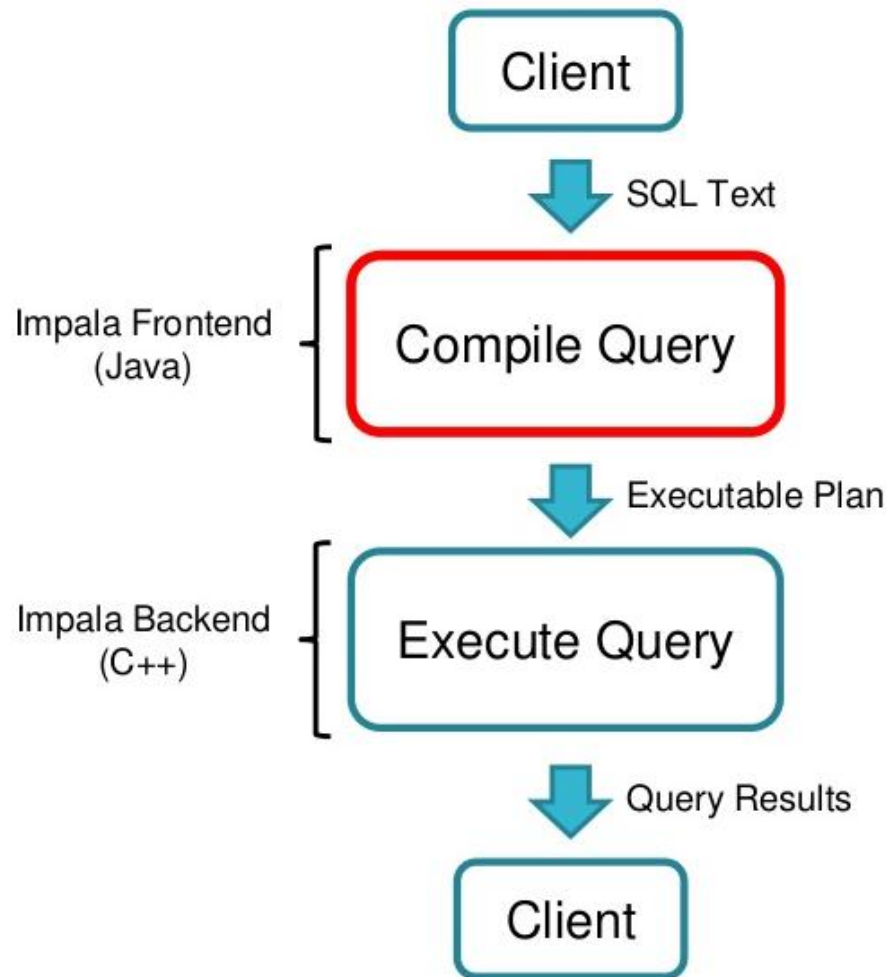
- **Java “front-end”**
- **2-phase planning process:**
 - single-node plan: left-deep tree of plan operators
 - partitioning of operator tree into plan fragments for parallel execution
- **Parallelization of plan operators:**
 - all query operators are fully distributed
 - joins: either broadcast or partitioned: decision is cost based
 - aggregation: fully parallel pre-aggregation and merge aggregation
 - top-n: initial stage done in parallel
- **Join order = FROM clause order** (soon to be cost-based)

Impala Partition

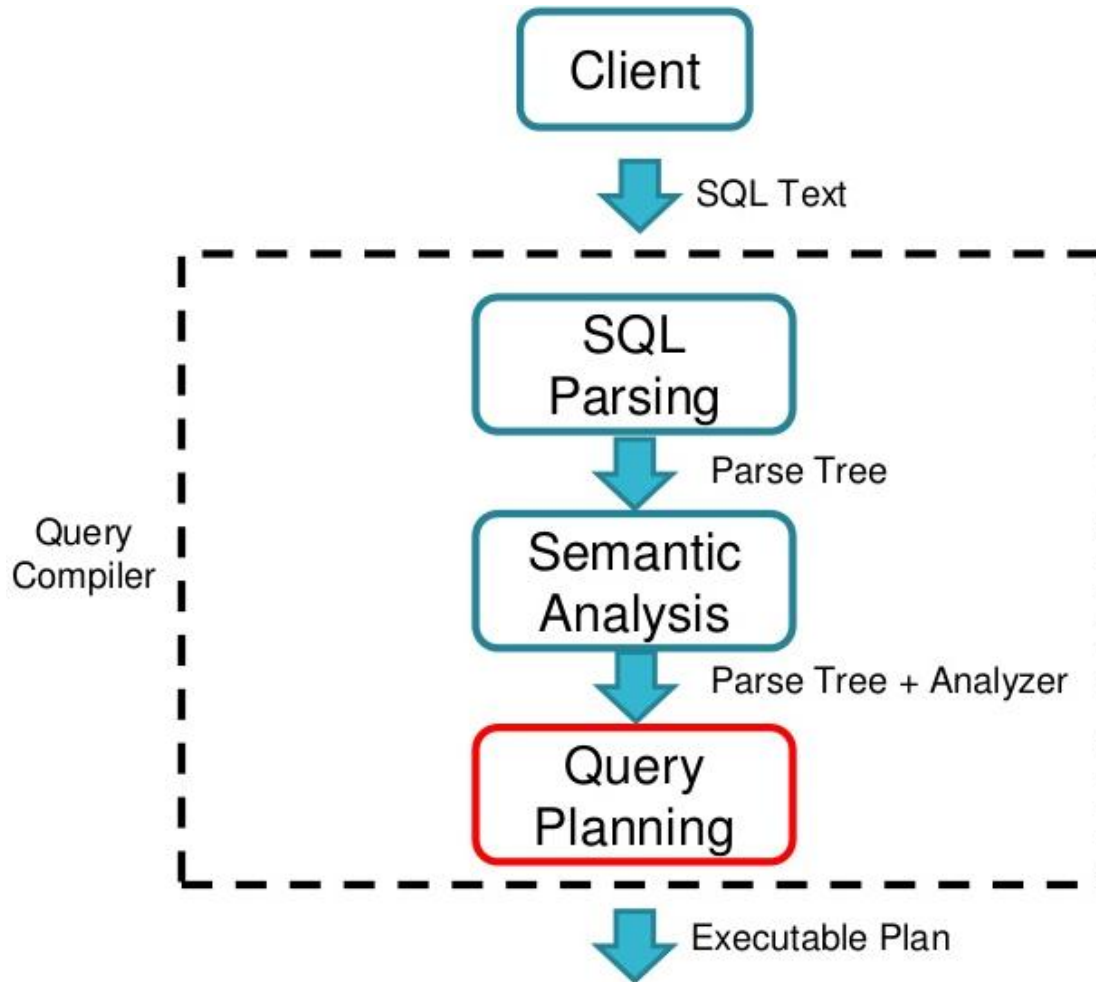
- **Example:**
 - `create table census (name string, census_year int) partitioned by (year int);`
 - `insert into census partition (year=2010) values ('Smith', 2010), ('Jones', 2010);`
- Each partition has its own HDFS directory, and all the data for that partition is stored in a data file in that directory
- To manually define how to partition the table (e.g., `year mod 5 == 0`), we have to create a new column to store the calculation result and then do the partition

Frontend

Flow of a SQL query

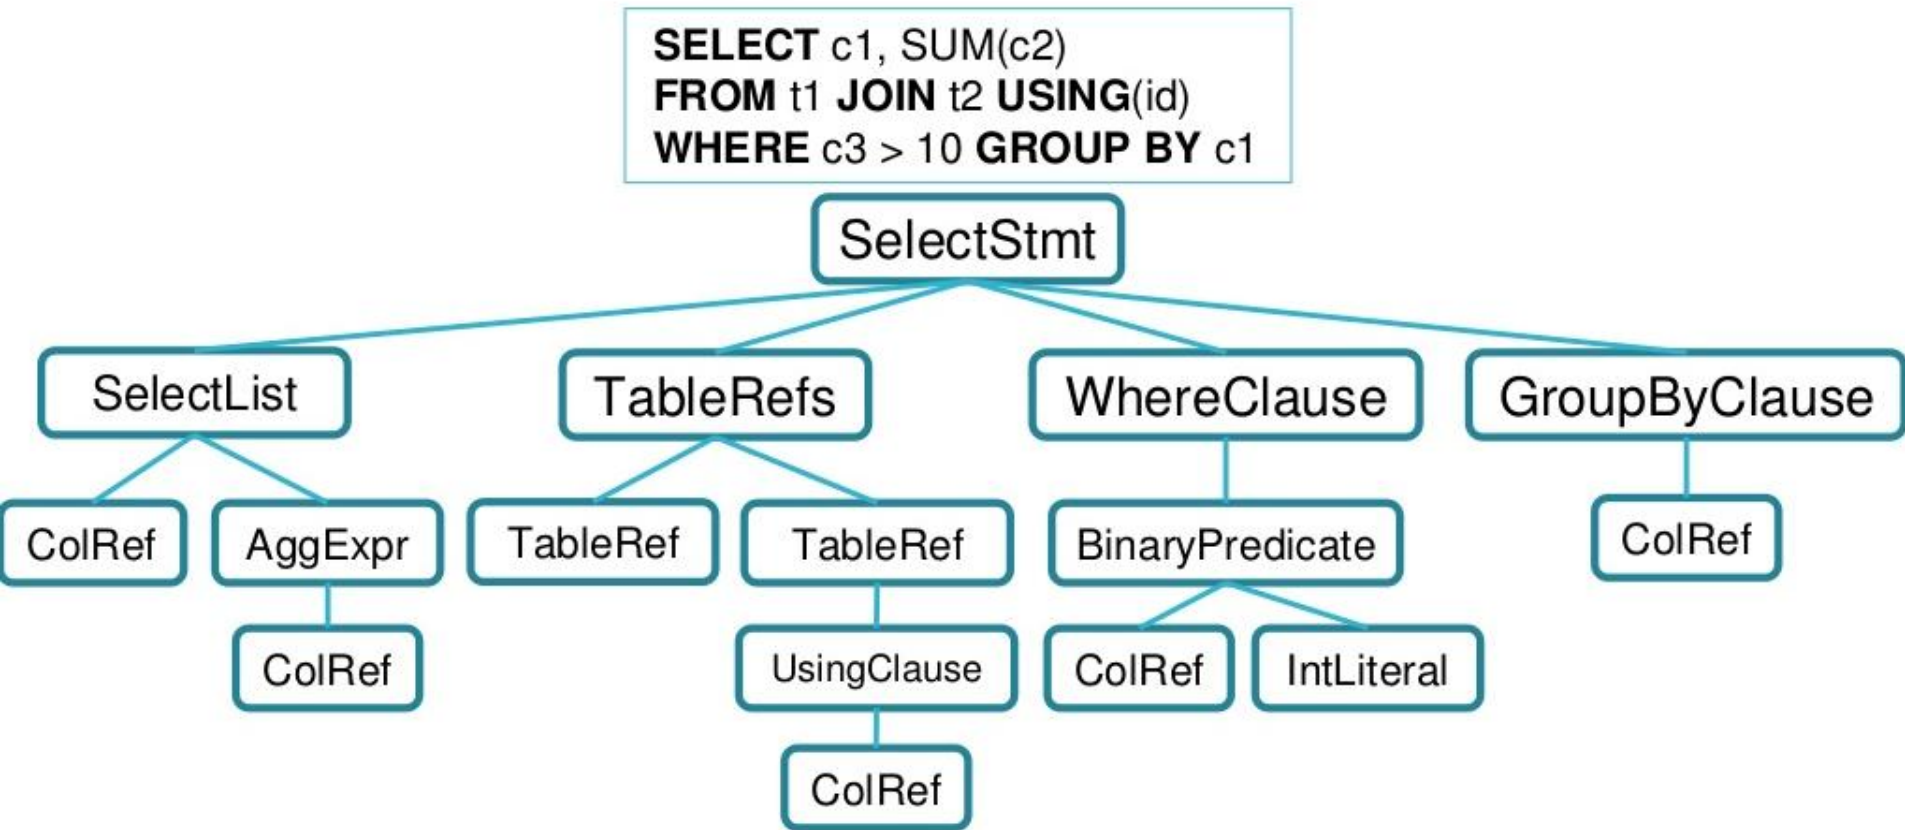


Query Compilation



Query Parsing

- Applies SQL grammar, reports syntax errors
- Produces parse tree capturing syntactic structure of query



Semantic Analysis

```
SELECT c1, SUM(c2)  
FROM t1 JOIN t2 USING(id)  
WHERE c3 > 10 GROUP BY c1
```

- **Precondition: Query is syntactically valid. Analysis operates on parse tree.**
- Consults table metadata
 - Do t1 and t2 exist? Does c1 exist in t1 or t2 (or both → error)? Does id exist in t1 and t2?
 - Does the user have privileges to SELECT from t1?
- Checks type compatibility of expressions, adds implicit casts
 - $c3 > 10 \rightarrow c3 > \text{cast}(10 \text{ as bigint})$
- SQL rules (semantic, not syntactic)
 - Does c1 appear in the GROUP BY clause?

Semantic Analysis

- Expression substitution for views
 - Resolve column references against base tables

```
SELECT c1, SUM(c2)
FROM (SELECT dept AS c1, revenue AS c2,
          month AS c3 FROM t1) AS v
WHERE c3 > 10 GROUP BY c1
```



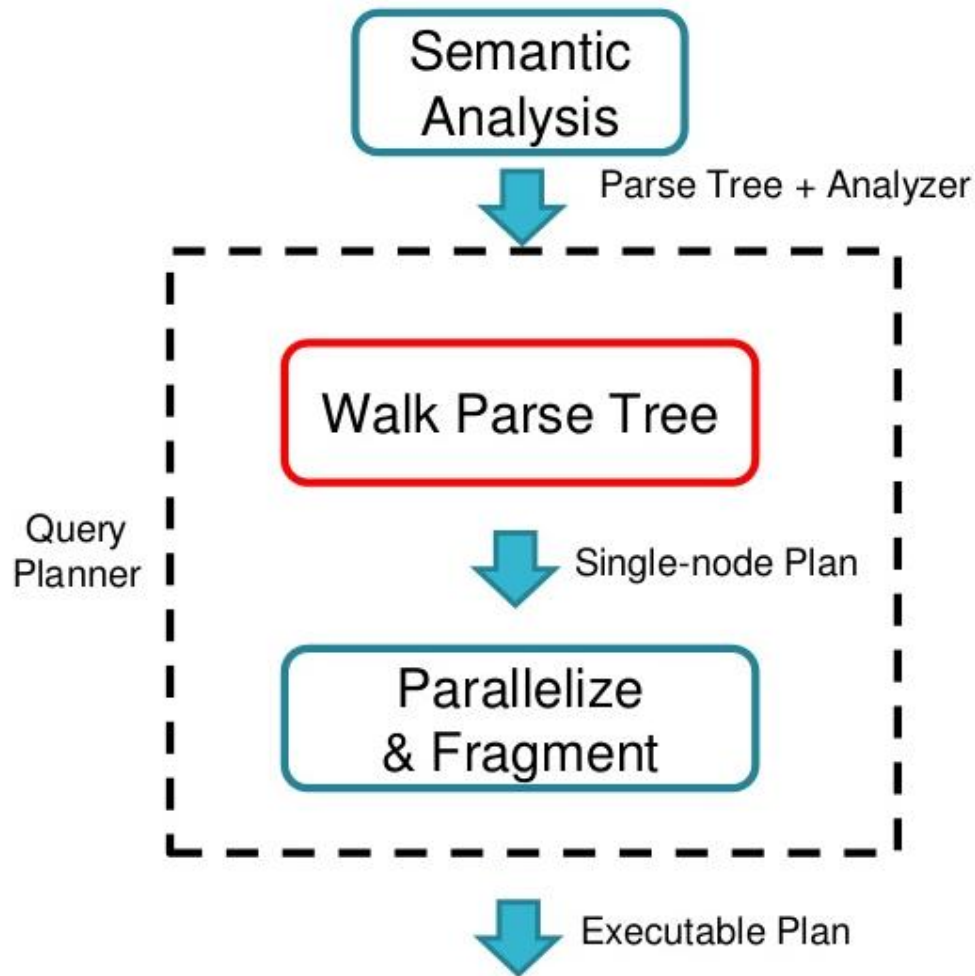
```
SELECT dept, SUM(revenue)
FROM t1
WHERE month > 10
GROUP BY dept
```

- Preparation for Planning
 - Register state in analyzer for correct predicate assignment during planning
 - Register predicates (WHERE, HAVING, ON, USING, etc.)
 - Register outer-joined tables
 - Compute value-transfer graph and equivalence classes for predicate inference
- (...)
- **Postcondition: Query is valid. An executable plan can be produced.**

Query Planning: Goals

- Generates executable plan (“tree” of operators)
 - Maximize scan locality using DataNode block metadata
 - Minimize data movement
- Full distribution of operators
- Query operators
 - Scan, HashJoin, HashAggregation, Union, TopN, Exchange

Query Planning: Overview




Query Planning: Single-Node Plan

- Four major functions:

- ▣ 1. Parse Tree \rightarrow Plan Tree
- ▣ 2. Assigns predicates to lowest plan node
- ▣ Optimizes join order
- ▣ Prunes irrelevant columns

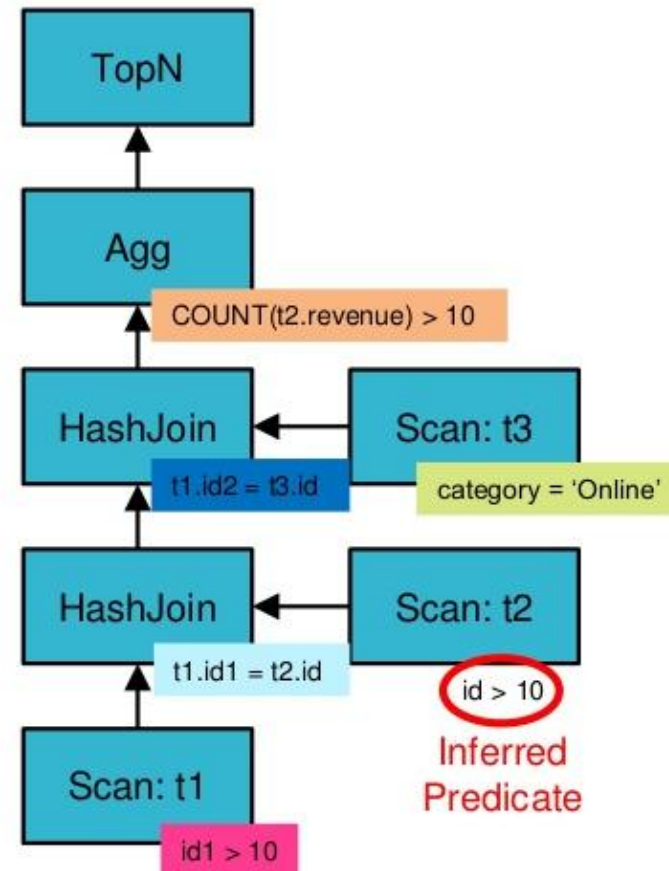
Parse Tree → Single-Node Plan Tree

```
SELECT t1.dept, SUM(t2.revenue)
FROM LargeHdfsTable t1
JOIN HugeHdfsTable t2 ON (t1.id1 = t2.id)
JOIN SmallHbaseTable t3 ON (t1.id2 = t3.id)
WHERE t3.category = 'Online' AND t1.id > 10
GROUP BY t1.dept
HAVING COUNT(t2.revenue) > 10
ORDER BY revenue LIMIT 10
```



Predicate Assignment & Inference

```
SELECT t1.dept, SUM(t2.revenue)
FROM LargeHdfsTable t1
JOIN HugeHdfsTable t2 ON (t1.id1 = t2.id)
JOIN SmallHbaseTable t3 ON (t1.id2 = t3.id)
WHERE t3.category = 'Online' AND t1.id > 10
GROUP BY t1.dept
HAVING COUNT(t2.revenue) > 10
ORDER BY revenue LIMIT 10
```



Join-Order Optimization







- Impala only considers left-deep join trees
- (Right join input is a table, not another join)
- Find cheapest valid join order
- Relies heavily on table and column statistics

Invalid Join Orders

```
SELECT t1.dept, SUM(t2.revenue)
FROM LargeHdfsTable t1
JOIN HugeHdfsTable t2 ON (t1.id1 = t2.id)
JOIN SmallHbaseTable t3 ON (t1.id2 = t3.id)
WHERE t3.category = 'Online' AND t1.id > 10
GROUP BY t1.dept
HAVING COUNT(t2.revenue) > 10
ORDER BY revenue LIMIT 10
```

No explicit or implicit
predicate between t2 and t3

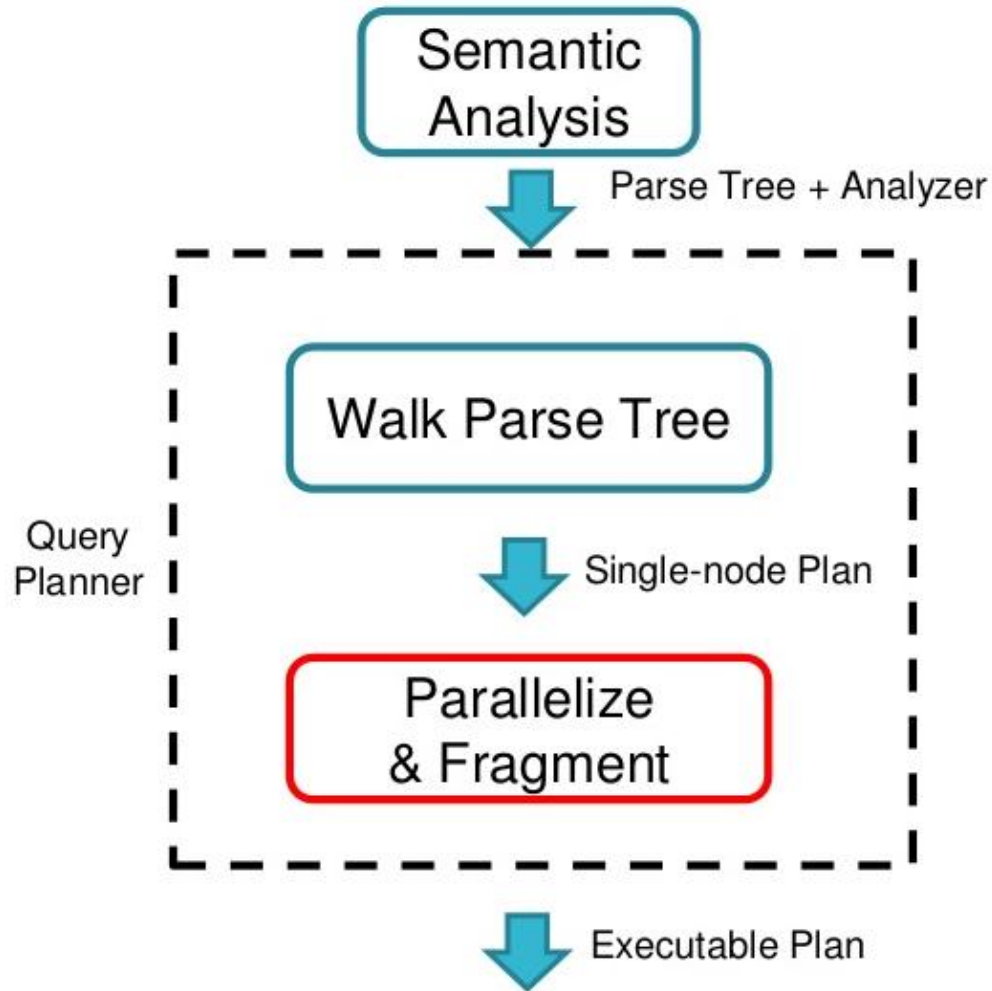
Join-Order Optimization



Join-Order Optimization

- Impala' s Implementation:
 - 1. Heuristic
 - Order tables descending by size
 - Best plan typically has largest table on the left (if valid)
 - 2. Plan enumeration & costing
 - Generates all possible join orders starting from a given left-most table (starting with largest one)
 - Ignore invalid join orders
 - *Estimates intermediate result sizes (key!)*
 - Chooses plan that minimizes intermediate result sizes

Query Planning: Overview



Query Planning: Distributed Plans

- **Goals:**
 - maximize scan locality, minimize data movement
 - full distribution of all query operators
- **Parallel joins:**
 - broadcast join: join is colocated with left input, right-hand side is broadcast to each node executing the join
(preferred for small right-hand side input)
 - partitioned join: both tables are hash-partitioned on join columns
(preferred for large joins)
 - cost-based decision based on column stats and estimated cost of data transfers

Query Planning: Distributed Plans

- **Parallel aggregation:**
 - pre-aggregation where data is first materialized
 - merge aggregation partitioned by grouping columns
- **Parallel top-N:**
 - initial top-N where data is first materialized
 - final top-N in single-node plan fragment

Two Types of Hash Joins

- **Default is BROADCAST (aka Replicated)**
 - Each node ends up with a copy of the right table(s)
 - Left side, read locally and streamed through local hash join
 - Good for one large table and multiple small tables
- **Alternative hash join type is SHUFFLE (aka partitioned)**
 - Right side hashed and shuffled; each node gets $1/N$ of the data
 - Left side hashed and shuffled, then streamed through join
 - Best choice for large_table JOIN large_table

Join Hint

```
select ...  
from large_table  
join [broadcast] small_table
```

```
select ...  
from large_table  
join [shuffle] large_table
```

Determine Join Type from EXPLAIN

```
explain
select
  s_state,
  count(*)
from store_sales
join store on (ss_store_sk = s_store_sk)
group by
  s_state;
```

2:HASH JOIN

```
| join op: INNER JOIN (BROADCAST)
| hash predicates:
|   ss_store_sk = s_store_sk
| tuple ids: 0 1
```

```
| ----4:EXCHANGE
|       tuple ids: 1
```

0:SCAN HDFS

```
table=tpcds.store_sales
tuple ids: 0
```

```
explain
select
  c_preferred_cust_flag,
  count(*)
from store_sales
join customer on (ss_customer_sk = c_customer_sk)
group by
  c_preferred_cust_flag;
```

2:HASH JOIN

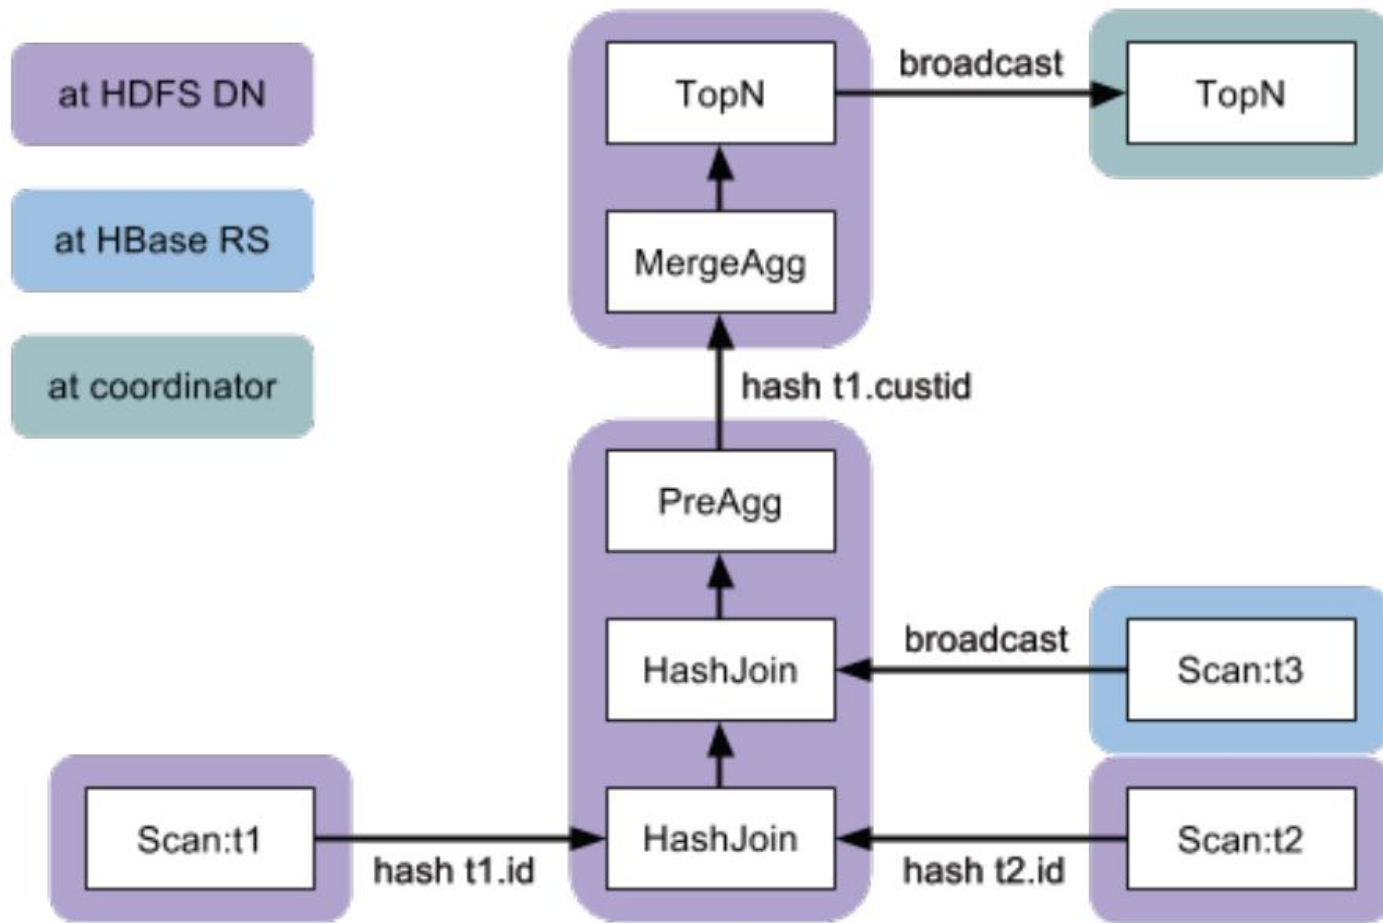
```
| join op: INNER JOIN (PARTITIONED)
| hash predicates:
|   ss_customer_sk = c_customer_sk
| tuple ids: 0 1
```

```
| ----5:EXCHANGE
|       tuple ids: 1
```

4:EXCHANGE

```
tuple ids: 0
```

Query Planning: Distributed Plans



HDFS Improvement Motivated by Impala

- Exposes HDFS block replica disk location information
- Allows for explicitly co-located block replicas across files
- In-memory caching of hot tables/files
- Reduces copies during reading, short-circuit reads

Disk Location of Block Replica

- **Problem:**
 - DataNode knows which DataNode blocks are on, not which disks
 - Only the DNs are aware of block replica->disk mapping
- **Impala wants to make sure that separate plan fragments operate on data on separate disks**
 - Maximizes aggregate available disk throughput

Disk Location of Block Replica

- **Solution:**
 - Adds new RPC call to DataNode to expose which volumes (disks) replicas are stored on
 - During query planning phase, `impalad`...
 - Determines all DNs data for query is stored on
 - Queries these DNs to get volume information
 - During query execution phase, `impalad`...
 - Queues disk reads so that only 1 or 2 reads ever happen to a given disk at a given time
 - With this additional information, Impala is able to ensure disk reads are large, minimizing seeks

Co-located Block Replicas

- **Problem:**

- ▣ When performing a join, a single impalad may have to read from both a local file and a remote file on another DN
- ▣ Ideally all reads should be done on local disks (assuming that local read is faster than remote read)

Co-located Block Replicas

- **Solution:**

- ▣ Adds features to HDFS to specify that a set of files should have their replicas placed on the same set of nodes
- ▣ Gives Impala more control of data
- ▣ Can ensure that tables/files which are joined frequently have their data co-located
- ▣ Additionally, more fine-grained block placement control allows for potential improvement in columnar storage format like Parquet

In-memory Caching

- Problem:
 - Impala queries are IO-bound
- Memory is fast and getting cheaper

In-memory Caching

- **Solution:**

- ▣ Adds facility to HDFS to explicitly read specific HDFS files into memory
- ▣ Allows Impala to read data at full memory bandwidth speed
- ▣ Gives cluster operator control over which files/tables are queried frequently and thus could be kept in memory

Short-circuit Reads

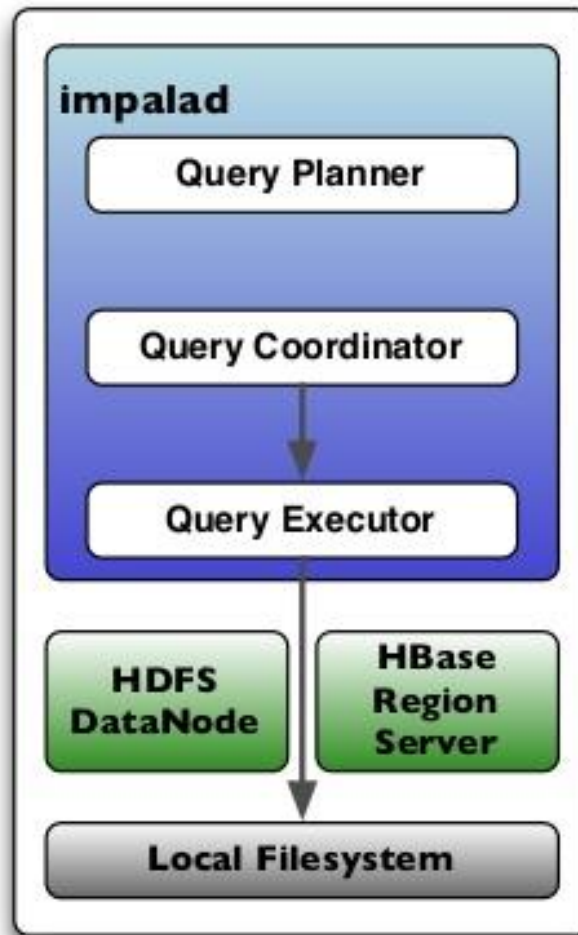
- **Problem:**

- ▣ A typical read in HDFS must be read from disk by DataNode, copied into DN memory, sent over network, copied into client buffers.

Short-circuit Reads

- **Solution:**
 - ▣ Reads are performed directly on local files, using direct buffers
 - ▣ In HDFS, allow for reads to completely bypass DataNode when client is co-located with block replica files, avoiding overhead of HDFS API
 - ▣ Reads data directly from disk to client buffers

Read
directly
from disk



Code Generation

Why Code Generation?

Why Code Generation?

SPEED!

Why Code Generation?

- Code generation (codegen) lets us use query-specific information to do less work
 - ▣ Remove conditionals
 - ▣ Propagate constant offsets, pointers, etc.
 - ▣ Inline virtual function calls

```

void MaterializeTuple(char* tuple) {
    for (int i = 0; i < num_slots_; ++i) {
        char* slot = tuple + offsets_[i];
        switch(types_[i]) {
            case BOOLEAN:
                *slot = ParseBoolean();
                break;
            case INT:
                *slot = ParseInt();
                break;
            case FLOAT: ...
            case STRING: ...
            // etc.
        }
    }
}

```

interpreted

```

void MaterializeTuple(char* tuple) {
    *(tuple + 0) = ParseInt();      // i = 0
    *(tuple + 4) = ParseBoolean();  // i = 1
    *(tuple + 5) = ParseInt();      // i = 2
}

```

codegen'd

```

void MaterializeTuple(char* tuple) {
    for (int i = 0; i < num_slots_; ++i) {
        char* slot = tuple + offsets_[i];
        switch(types_[i]) {
            case BOOLEAN:
                *slot = ParseBoolean();
                break;
            case INT:
                *slot = ParseInt();
                break;
            case FLOAT: ...
            case STRING: ...
            // etc.
        }
    }
}

```

interpreted

```

void MaterializeTuple(char* tuple) {
    *(tuple + 0) = ParseInt();      // i = 0
    *(tuple + 4) = ParseBoolean(); // i = 1
    *(tuple + 5) = ParseInt();      // i = 2
}

```

codegen'd

```

void MaterializeTuple(char* tuple) {
    for (int i = 0; i < num_slots_; ++i) {
        char* slot = tuple + offsets_[i];
        switch(types_[i]) {
            case BOOLEAN:
                *slot = ParseBoolean();
                break;
            case INT:
                *slot = ParseInt();
                break;
            case FLOAT: ...
            case STRING: ...
            // etc.
        }
    }
}

```

interpreted

```

void MaterializeTuple(char* tuple) {
    *(tuple + 0) = ParseInt();      // i = 0
    *(tuple + 4) = ParseBoolean();  // i = 1
    *(tuple + 5) = ParseInt();      // i = 2
}

```

codegen'd

```

void MaterializeTuple(char* tuple) {
    for (int i = 0; i < num_slots_; ++i) {
        char* slot = tuple + offsets_[i];
        switch(types_[i]) {
            case BOOLEAN:
                *slot = ParseBoolean();
                break;
            case INT:
                *slot = ParseInt();
                break;
            case FLOAT: ...
            case STRING: ...
            // etc.
        }
    }
}

```

interpreted

```

void MaterializeTuple(char* tuple) {
    *(tuple + 0) = ParseInt();    // i = 0
    *(tuple + 4) = ParseBoolean(); // i = 1
    *(tuple + 5) = ParseInt();    // i = 2
}

```

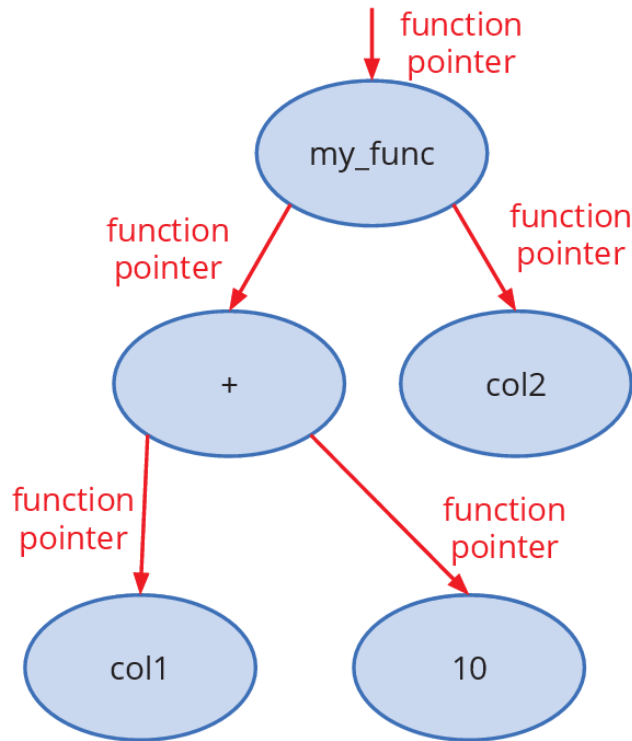
codegen'd

User-Defined Functions (UDFs)

- Allows users to extend Impala's functionality by writing their own functions
- e.g. `select my_func(c1) from table;`
- Defined as C++ functions
- UDFs can be compiled to LLVM IR with Clang \Rightarrow inline UDFs
- IR can be just-in-time compiled (JIT'd) and replace the interpreted functions

```
IntVal my_func(const IntVal& v1, const IntVal& v2) {  
    return IntVal(v1.val * 7 / v2.val);  
}
```

SELECT my_func(col1 + 10, col2) FROM ...



interpreted


$(col1 + 10) * 7 / col2$

codegen'd

Performance (Jan 2014)

- 3TB (TPC-DS scale factor 3,000) across five typical Hadoop DataNodes (dual-socket, 8-core, 16-thread CPU; 96GB memory; 1Gbps Ethernet; 12 x 2TB disk drives)

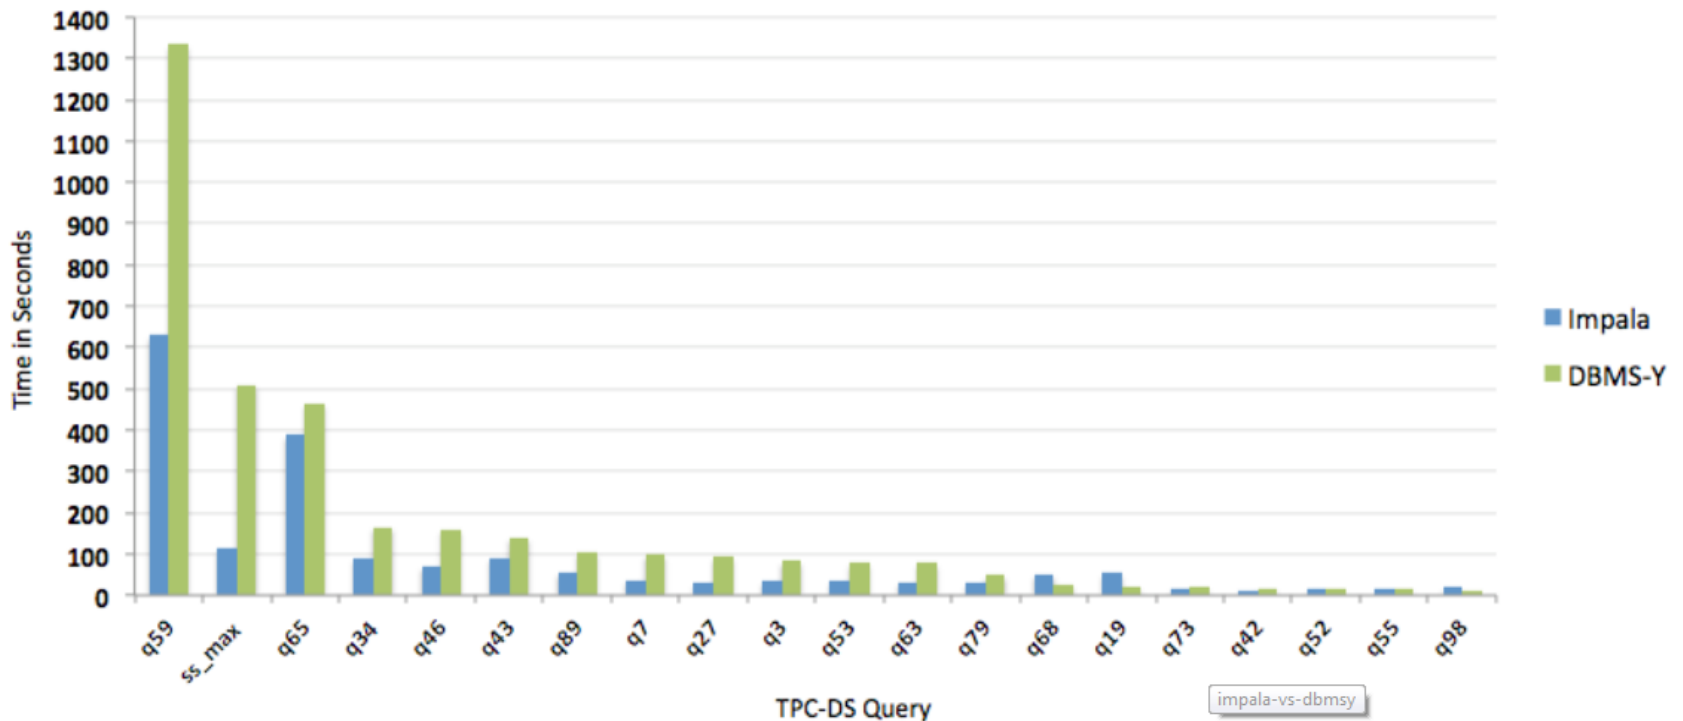
Impala versus Hive 0.12/Stinger (Lower bars are better)



Performance (Jan 2014)

- 30TB set of TPC-DS data (scale factor 30,000), 20 nodes with 96GB memory per node

Impala versus DBMS-Y
(Lower bars are better)



Weaknesses and Limitations

- Data is immutable, no updating
- Response time is not microsecond
- Do not support some operations, like update and delete
- No beyond SQL and advanced data structures (buckets, samples, transforms, arrays, structs, maps, xpath, json)
- When broadcast join, smaller table has to fit in aggregate memory of all executing nodes
- No custom storage format
- LIMIT required when using ORDER BY
- High memory usage

References

- Cloudera Impala official documentation and slides