BUILT FOR THE SPEED OF BUSINESS

Massively Parallel Processing with Procedural Python

How do we use the PyData stack in data science engagements at Pivotal?

Ian Huston, @ianhuston
Data Scientist, Pivotal

Pivotal

Some Links for this talk

- Simple code examples: https://github.com/ihuston/plpython_examples
- IPython notebook rendered with nbviewer: http://tinyurl.com/ih-plpython
- More info (written for PL/R but applies to PL/Python): http://gopivotal.github.io/gp-r/
- Traffic Disruption demo (if we have time) http://ds-demo-transport.cfapps.io

About Pivotal

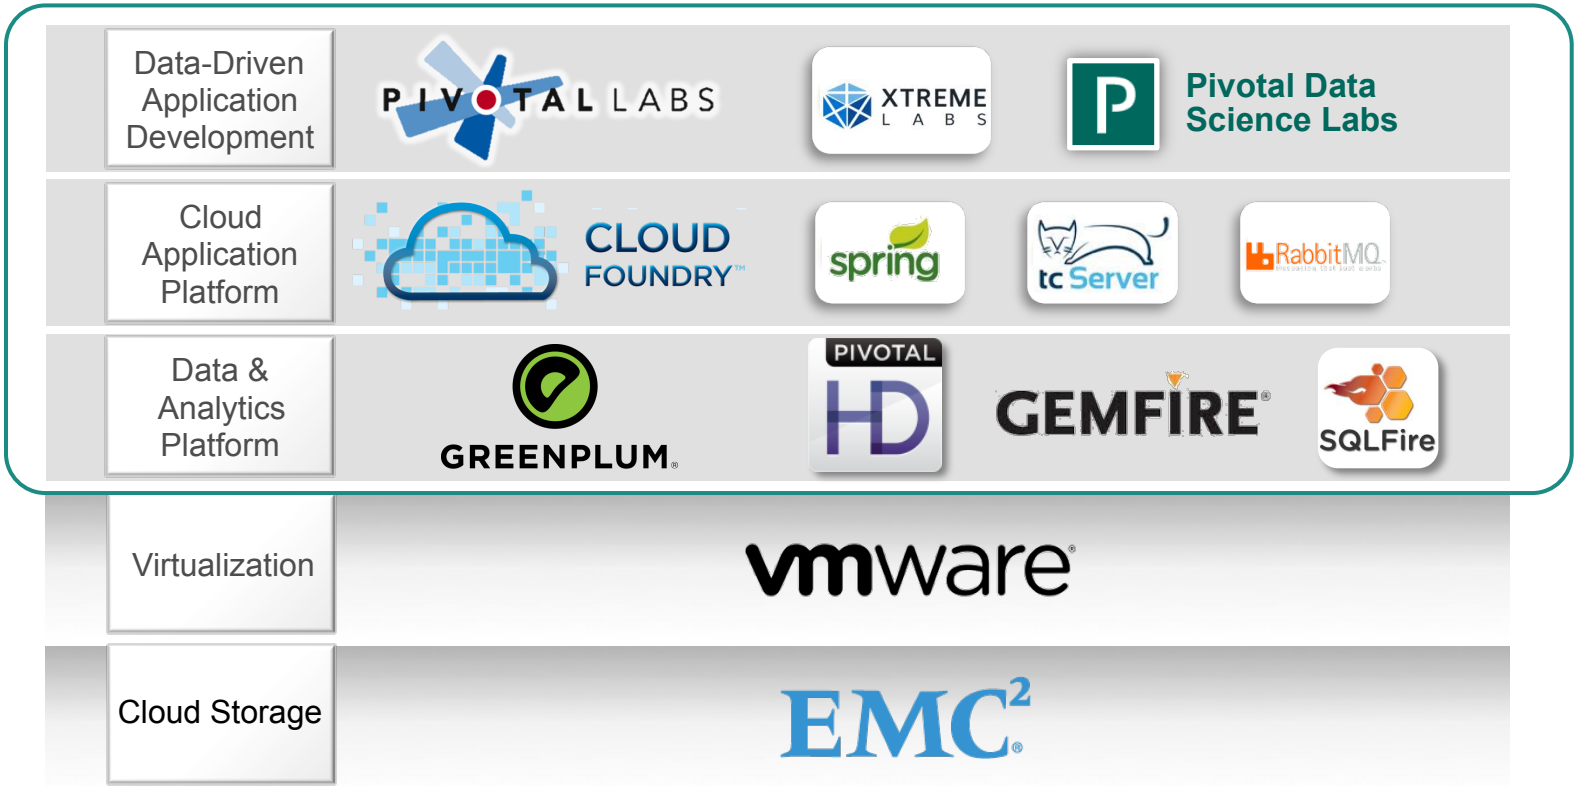

What do our customers look like?

- Large enterprises with lots of data collected
  - Work with 10s of TBs to PBs of data, structured & unstructured
- Not able to get what they want out of their data
  - Old Legacy systems with high cost and no flexibility
  - Response times are too slow for interactive data analysis
  - Can only deal with small samples of data locally
- They want to transform into data driven enterprises

Open Source is Pivotal

Apache Tomcat

Open Chorus
Project

Pivotal's Open Source Contributions

Lots more interesting small projects:

- PyMADlib – Python Wrapper for MADlib
  https://github.com/gopivotal/pymadlib
- PivotalR – R wrapper for MADlib
  http://github.com/madlib-internal/PivotalR
- Part-of-speech tagger for Twitter via SQL
  http://vatsan.github.io/gp-ark-tweet-nlp/
- Pandas via psql
  (interactive PostgreSQL terminal)
  https://github.com/vatsan/pandas_via_psql

Typical Engagement Tech Setup

- Platform:
  - Greenplum Analytics Database (GPDB)
  - Pivotal HD Hadoop Distribution + HAWQ (SQL DB on Hadoop)
- Open Source Options (http://gopivotal.com):
  - Greenplum Community Edition
  - Pivotal HD Community Edition (HAWQ not included)
  - MADlib in-database machine learning library (http://madlib.net)
- Where Python fits in:
  - **PL/Python running in-database**, with nltk, scikit-learn etc
  - IPython for exploratory analysis
  - Pandas, Matplotlib etc.

PIVOTAL DATA SCIENCE TOOLKIT

1 Find Data

Platforms
- Greenplum DB
- Pivotal HD
- Hadoop (other)
- SAS HPA
- AWS

2 Write Code

Editing Tools
- Vi/Vim
- Emacs
- Smultron
- TextWrangler
- Eclipse
- Notepad++
- IPython
- Sublime

Languages
- SQL
- Bash scripting
- C
- C++
- C#
- Java
- Python
- R

3 Run Code

Interfaces
- pgAdminIII
- psql
- psycopg2
- Terminal
- Cygwin
- Putty
- Winscp

4 Write Code for Big Data

In-Database
- SQL
- PL/Python
- PL/Java
- PL/R
- PL/pgSQL

Hadoop
- HAWQ
- Pig
- Hive
- Java

5 Implement Algorithms

Libraries
- MADlib

Java
- Mahout

R
- (Too many to list!)

Text
- OpenNLP
- NLTK
- GPText

C++
- opencv

Python
- NumPy
- SciPy
- scikit-learn
- Pandas

Programs
- Alpine Miner
- Rstudio
- MATLAB
- SAS
- Stata

6 Show Results

Visualization
- python-matplotlib
- python-networkx
- D3.js
- Tableau
- GraphViz
- Gephi
- R (ggplot2, lattice, shiny)
- Excel

7 Collaborate

Sharing Tools
- Chorus
- Confluence
- Socialcast
- Github
- Google Drive & Hangouts

A large and varied tool box!

PL/Python

Pivotal

MPP Architectural Overview

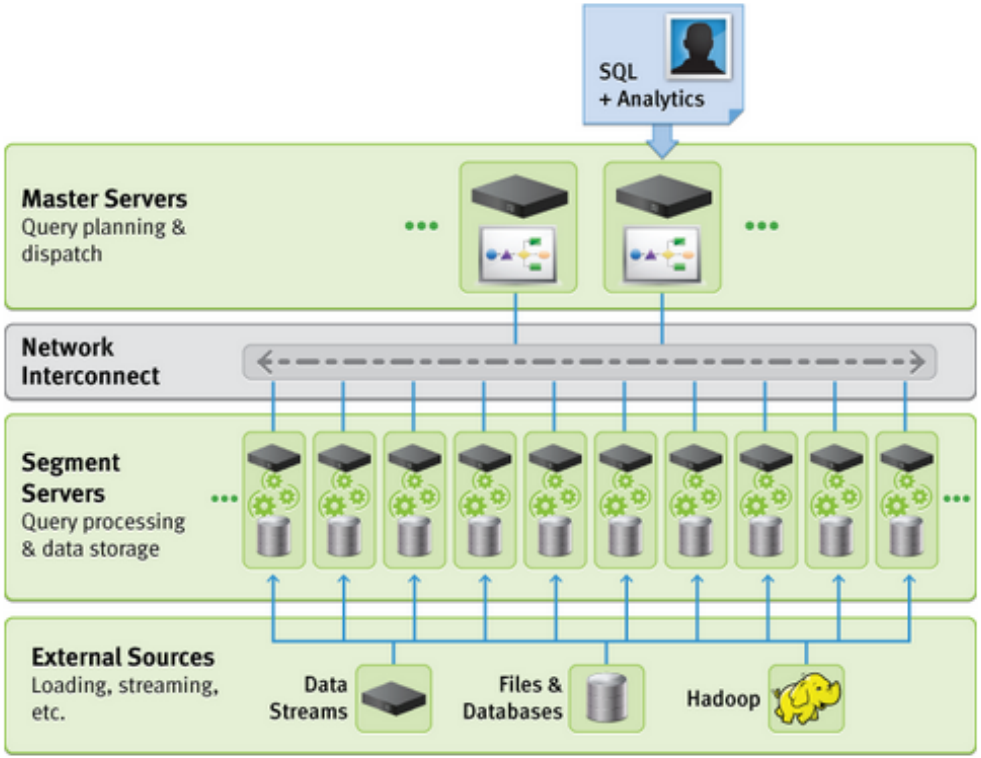

Think of it as multiple PostGreSQL servers

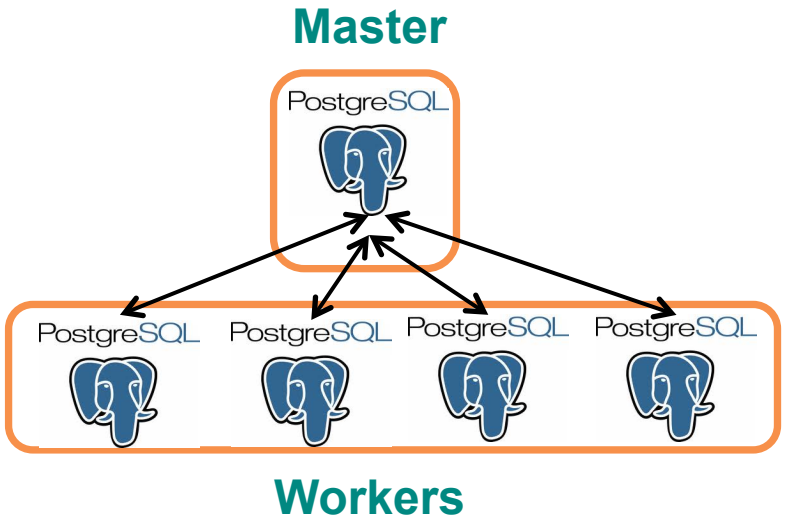

Data Parallelism

- Little or no effort is required to break up the problem into a number of parallel tasks, and there exists no dependency (or communication) between those parallel tasks.
- Examples:
  - Measure the height of each student in a classroom (explicitly parallelizable by student)
  - MapReduce
  - map() function in Python

User-Defined Functions (UDFs)

- PostgreSQL/Greenplum provide lots of flexibility in defining your own functions.
- Simple UDFs are SQL queries with calling arguments and return types.

Definition:

```
CREATE FUNCTION times2(INT)
RETURNS INT
AS $$
    SELECT 2 * $1
$$ LANGUAGE sql;
```

Execution:

```
SELECT times2(1);
 times2
--------
      2
(1 row)
```

PL/X : X in {pgsql, R, Python, Java, Perl, C etc.}

- Allows users to write Greenplum/ PostgreSQL functions in the R/Python/ Java, Perl, pgsql or C languages
- The interpreter/VM of the language 'X' is installed on each node of the Greenplum Database Cluster
- Data Parallelism:
  - PL/X piggybacks on Greenplum's MPP architecture

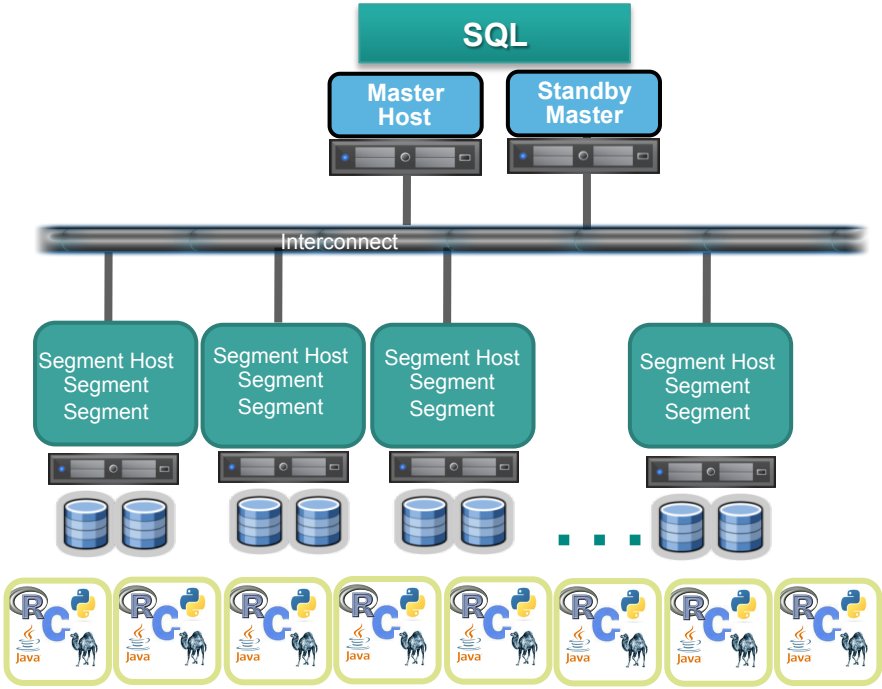

Intro to PL/Python

- Procedural languages need to be installed on each database used.
- Name in SQL is plpythonu, ‘u’ means untrusted so need to be superuser to install.
- Syntax is like normal Python function with function definition line replaced by SQL wrapper. Alternatively like a SQL User Defined Function with Python inside.

SQL wrapper

Normal Python

SQL wrapper

```
CREATE FUNCTION pymax (a integer, b integer)
  RETURNS integer
AS $$
  if a > b:
    return a
  return b
$$ LANGUAGE plpythonu;
```

Examples

Pivotal

Returning Results

- Postgres primitive types (int, bigint, text, float8, double precision, date, NULL etc.)
- Composite types can be returned by creating a composite type in the database:

```
CREATE TYPE named_value AS (
  name  text,
  value  integer
);
```

- Then you can return a list, tuple or dict (not sets) which reference the same structure as the table:

```
CREATE FUNCTION make_pair (name text, value integer)
  RETURNS named_value
AS $$
  return [ name, value ]
  # or alternatively, as tuple: return ( name, value )
  # or as dict: return { "name": name, "value": value }
  # or as an object with attributes .name and .value
$$ LANGUAGE plpythonu;
```

- For functions which return multiple rows, prefix "`setof`" before the return type

Returning more results

You can return multiple results by wrapping them in a sequence (tuple, list or set), an iterator or a generator:

Sequence

```
CREATE FUNCTION make_pair (name text)
  RETURNS SETOF named_value
AS $$
  return ([ name, 1 ], [ name, 2 ], [ name, 3])
$$ LANGUAGE plpythonu;
```

Generator

```
CREATE FUNCTION make_pair (name text)
  RETURNS SETOF named_value  AS $$
  for i in range(3):
      yield (name, i)
$$ LANGUAGE plpythonu;
```

Accessing Packages

- On Greenplum DB: To be available packages must be installed on the individual segment nodes.
  - Can use “parallel ssh” tool `gpssh` to conda/pip install
  - Currently Greenplum DB ships with Python 2.6 (!)
- Then just import as usual inside function:

```
CREATE FUNCTION make_pair (name text)
  RETURNS named_value
AS $$
  import numpy as np
  return ((name,i) for i in np.arange(3))
$$ LANGUAGE plpythonu;
```

Benefits of PL/Python

- Easy to bring your code to the data.
- When SQL falls short leverage your Python (or R/Java/C) experience quickly.
- Apply Python across terabytes of data with minimal overhead or additional requirements.
- Results are already in the database system, ready for further analysis or storage.

MADlib

Pivotal

Going Beyond Data Parallelism

- Data Parallel computation via PL/Python libraries only allow us to run ‘n’ models in parallel.
- This works great when we are building one model for each value of the group by column, but we need parallelized algorithms to be able to build a single model on all the available data
- For this, we use MADlib – an open source library of parallel in-database machine learning algorithms.

MADlib: The Origin

UrbanDictionary

mad *(adj.): an adjective used to enhance a noun.*
1- dude, you got skills.
2- dude, you got mad skills

- First mention of MAD analytics was at VLDB 2009

  MAD Skills: New Analysis Practices for Big Data
  J. Hellerstein, J. Cohen, B. Dolan, M. Dunlap, C. Welton
  (with help from: Noelle Sio, David Hubbard, James Marca)
  http://db.cs.berkeley.edu/papers/vldb09-madskills.pdf

- MADlib project initiated in late 2010:

  Greenplum Analytics team and Prof. Joe Hellerstein

- Open Source!
  https://github.com/madlib/madlib
- Works on Greenplum DB, PostgreSQL and also HAWQ & Impala
- Active development by Pivotal
  - Latest Release: v1.4 (Nov 2013)
- Downloads and Docs:
  http://madlib.net/

MADlib Executes Algorithms In-Place

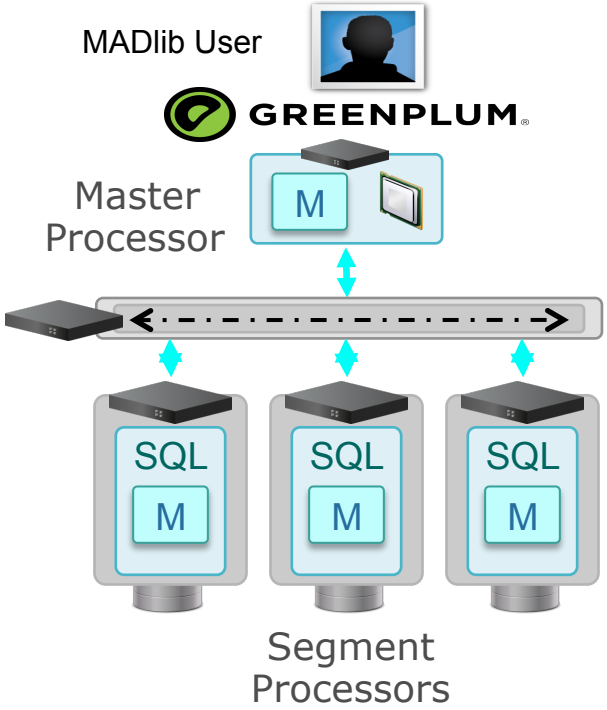

MADlib Advantages

- No Data Movement
- Use MPP architecture's full compute power
- Use MPP architecture's entire memory to process data sets

MADlib In-Database Functions

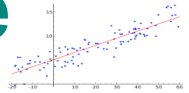
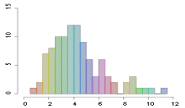
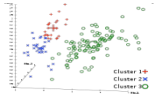

Predictive Modeling Library

Generalized Linear Models
- Linear Regression
- Logistic Regression
- Multinomial Logistic Regression
- Cox Proportional Hazards
- Regression
- Elastic Net Regularization
- Sandwich Estimators (Huber white, clustered, marginal effects)

Machine Learning Algorithms
- Principal Component Analysis (PCA)
- Association Rules (Affinity Analysis, Market Basket)
- Topic Modeling (Parallel LDA)
- Decision Trees
- Ensemble Learners (Random Forests)
- Support Vector Machines
- Conditional Random Field (CRF)
- Clustering (K-means)
- Cross Validation

Matrix Factorization
- Single Value Decomposition (SVD)
- Low-Rank

Linear Systems
- Sparse and Dense Solvers

Descriptive Statistics

Sketch-based Estimators
- CountMin (Cormode-Muthukrishnan)
- FM (Flajolet-Martin)
- MFV (Most Frequent Values)

Correlation
Summary

Support Modules

Array Operations
Sparse Vectors
Random Sampling
Probability Functions

Architecture

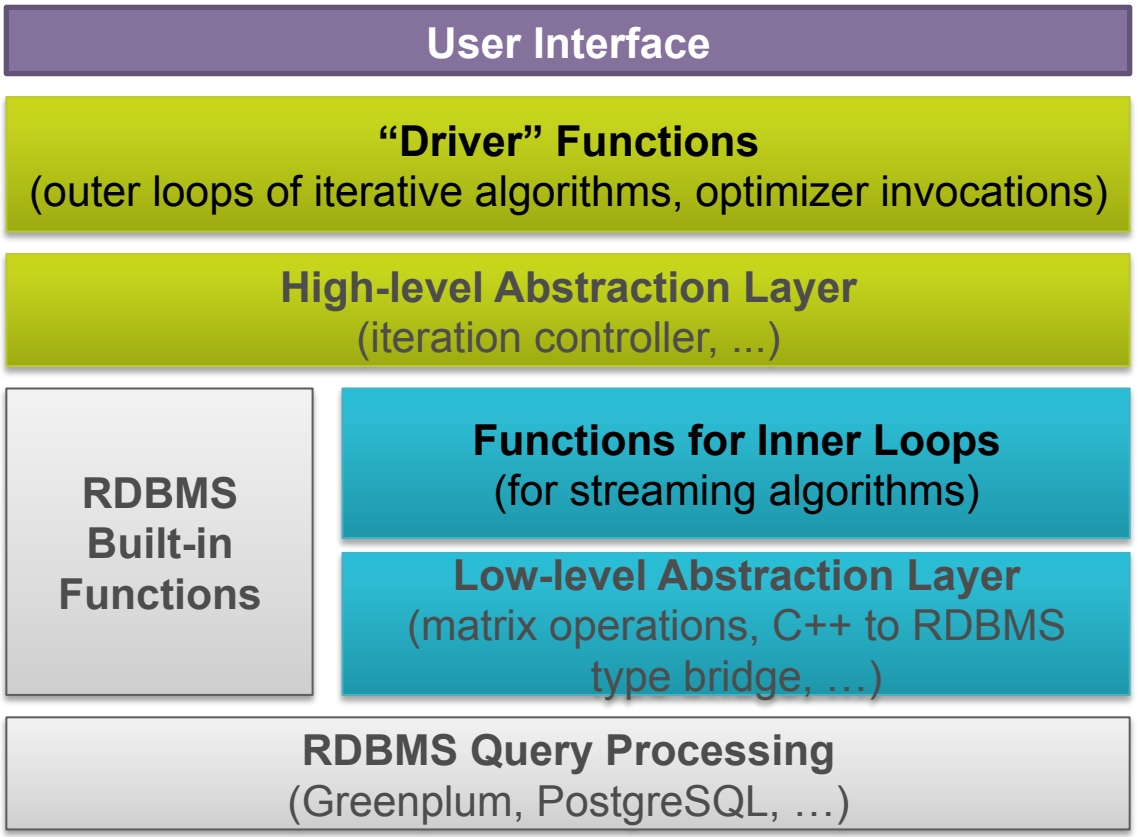

SQL, generated from specification

Python with templated SQL

Python

C++

How does it work ? : A Linear Regression Example

- Finding linear dependencies between variables
  - $y \approx c_0 + c_1 \cdot x_1 + c_2 \cdot x_2$?

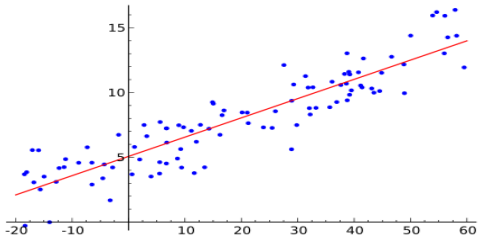

Vector of dependent variables y

```
# select y, x1, x2  from unm limit 6;
```

y	x1	x2
10.14	0	0.3
11.93	0.69	0.6
13.57	1.1	0.9
14.17	1.39	1.2
15.25	1.61	1.5
16.15	1.79	1.8

Design Matrix X

Reminder: Linear-Regression Model

- $E[Y \mid \boldsymbol{x}] = \boldsymbol{x}^T \boldsymbol{c}$
- If residuals i.i.d. Gaussians with standard deviation σ:
  - max likelihood ⇔ min sum of squared residuals

$$f(y \mid \boldsymbol{x}) \propto \exp\left(-\frac{1}{2\sigma^2} \cdot (y - \boldsymbol{x}^T \boldsymbol{c})^2\right)$$

- First-order conditions for the following quadratic objective (in $\boldsymbol{c}$)

$$(\boldsymbol{y} - X\boldsymbol{c})^T (\boldsymbol{y} - X\boldsymbol{c})$$

yield the minimizer

$$\widehat{\boldsymbol{c}} = (X^T X)^{-1} X^T \boldsymbol{y}$$

Linear Regression: Streaming Algorithm

- How to compute with a single table scan?

$$\widehat{\boldsymbol{c}} = (X^T X)^{-1} X^T \boldsymbol{y}$$

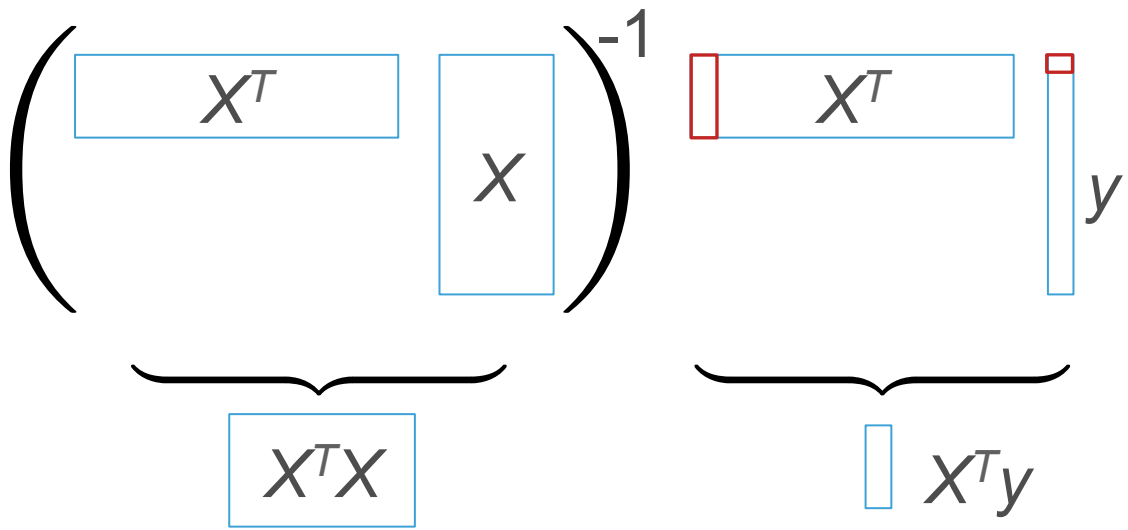

Linear Regression: Parallel Computation

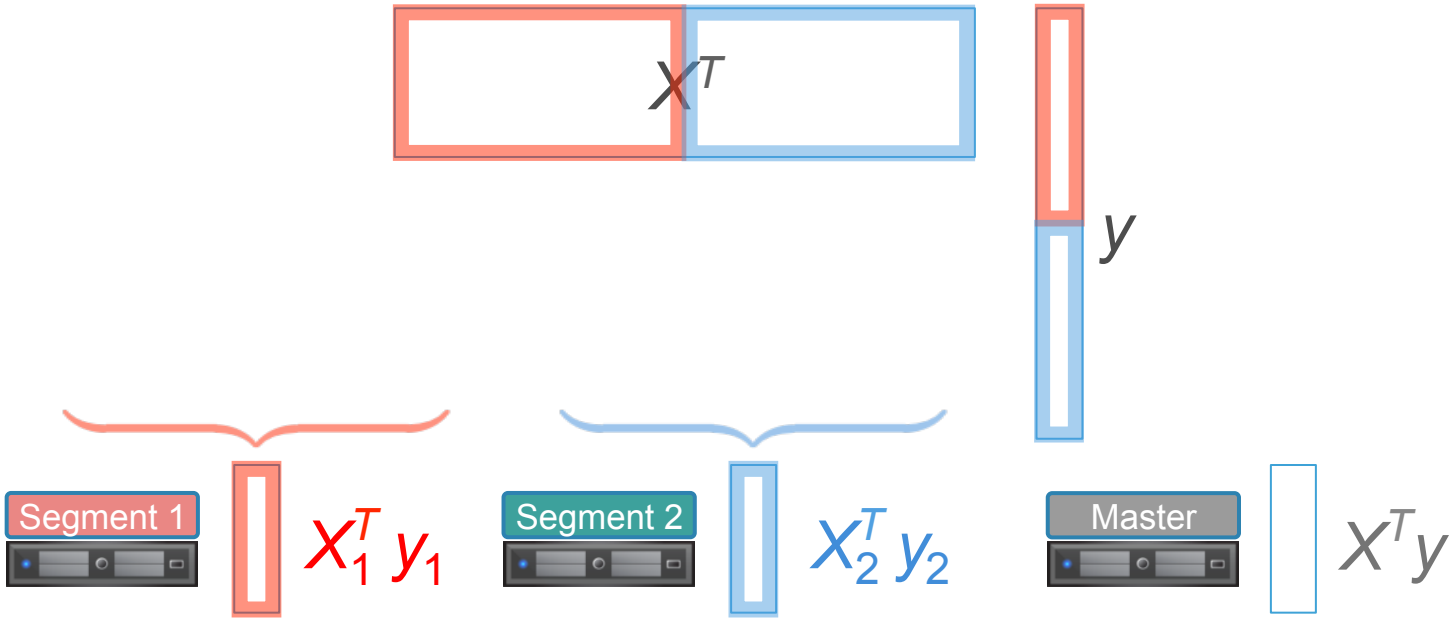

$$X^T \boldsymbol{y} = (X_1^T \; X_2^T) \begin{pmatrix} \boldsymbol{y}_1 \\ \boldsymbol{y}_2 \end{pmatrix} = \sum X_i^T \boldsymbol{y}_i$$

Demos

- We built demos to showcase our technology pipeline, using Python technology.
- Two use cases:
  - Topic and Sentiment Analysis of Tweets
  - London Road Traffic Disruption prediction

Topic and Sentiment Analysis Pipeline

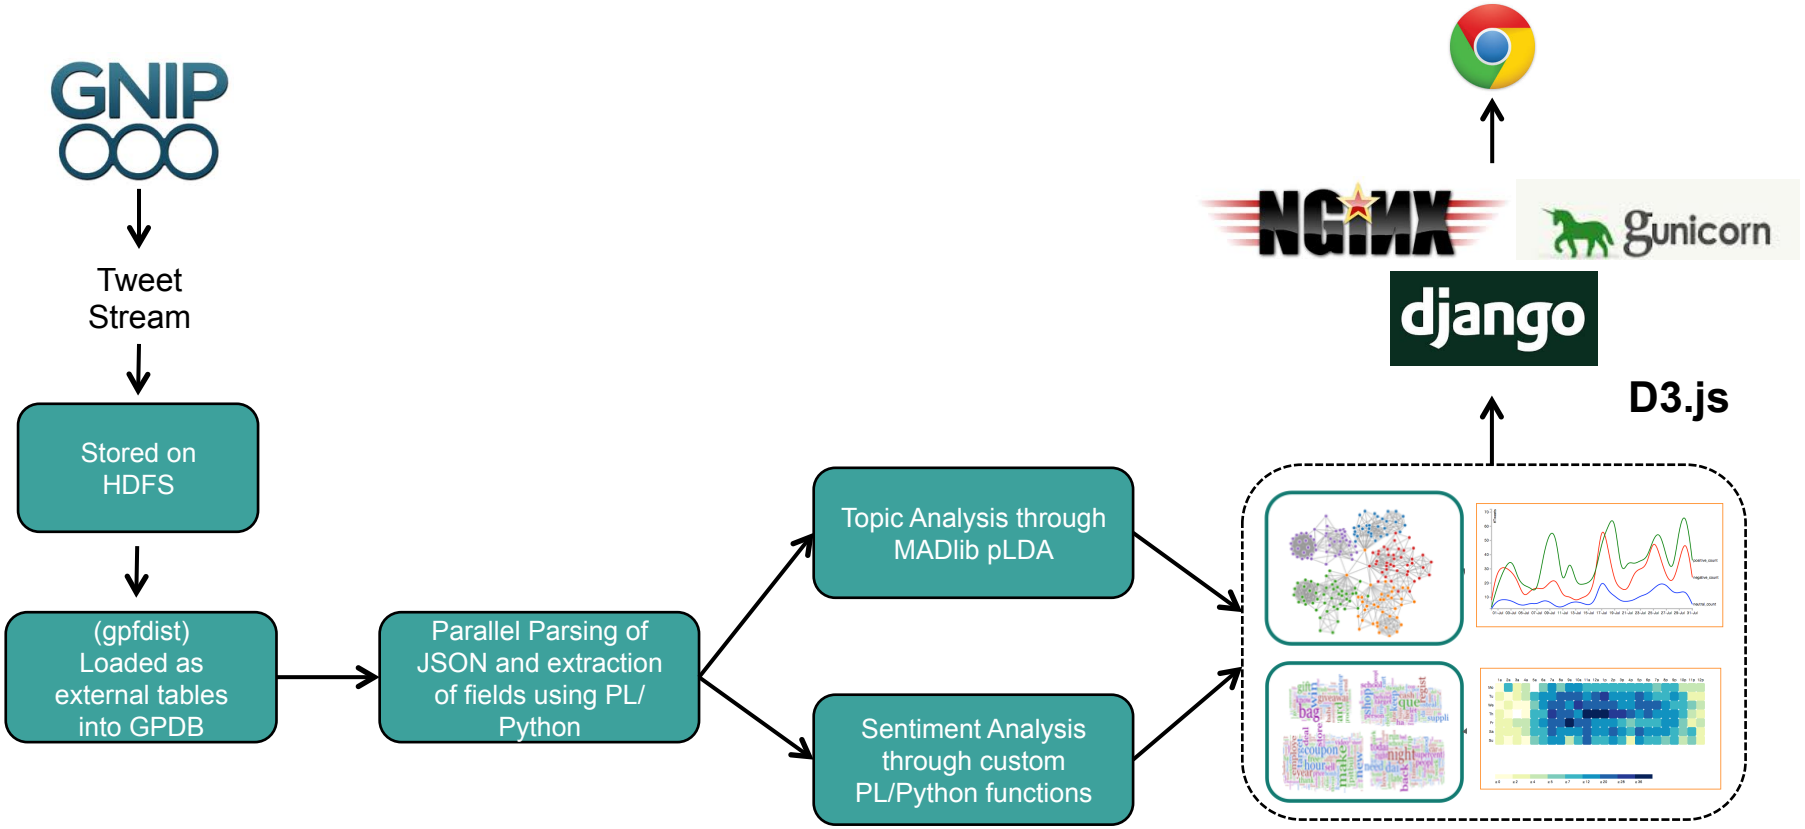

Transport Disruption Prediction Pipeline

Transport for London
Traffic Disruption feed

GREENPLUM®

Pivotal Greenplum
Database

Deduplication

Feature Creation

wunderground.com

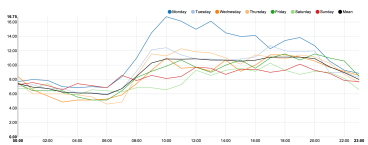

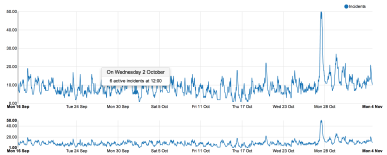

d3.js & NVD3

Interactive SVG figures

CLOUD FOUNDRY™

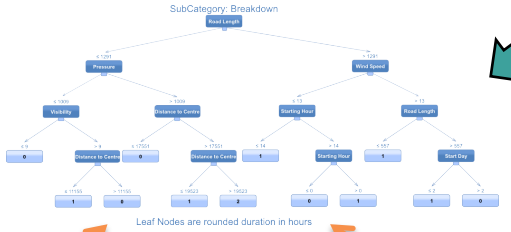

python™

scikit learn

Modelling & Machine Learning

Get in touch

Feel free to contact me about PL/Python, or more generally about Data Science and opportunities available.

@ianhuston

ihuston @ gopivotal.com

http://www.ianhuston.net

BUILT FOR THE SPEED OF BUSINESS