

# In-Memory Computing: Facts & Myths

**Jon Webster**  
**VP of Product**  
**Management**



[www.gridgain.com](http://www.gridgain.com)




[#gridgain](https://twitter.com/gridgain)

# What is In-Memory Computing?

**In-Memory Computing** uses high-performance, integrated, distributed memory systems to compute and transact on large-scale data sets in real-time - orders of magnitude faster than legacy disk-based systems.

# Why Now?



# Paradigm Shift à la 1970s

## 1970s:

- > IBM released “Winchester” IBM 340 disk - **Era of HDD**  
Tapes start to decline
- > SQL  
Era of Structured Data

## 2010s:

- > 64-bit CPUs + DRAM prices drop 30% YoY - **Era of Memory**  
HDDs start to decline
- > NoSQL + SQL  
Era of Unstructured Data

**Gartner** RAM is a new disk, disk is a new tape.

# Memory First vs. Disk First

## > **Memory First Architecture:**

Memory is primary storage, disk for backups

**Reading Record:** API call <-> pointer arithmetic

**Latency:** nanoseconds

## > **Disk First Architecture:**

Disk as primary storage, memory for caching

**Reading Record:** API call <-> OS I/O <-> I/O controller <-> disk

**Latency:** milliseconds

# Bring Computations To Data

- > **Client-Server, J2EE, SMP** 1990s-2005s
  - > Data is moved to application layer for processing:
    - Data not-partitioned and stored in central RDBMS
    - Data sizes are relatively small
    - Technically **impossible** to distributed computations to central RDBMS
- > **In-Memory Computing, Hadoop, MPP** 2005s+
  - > Computations are moved to data:
    - Data is partitioned and stored in distributed systems
    - Data overall sizes are massive
    - Technically **possible** to distribute computations for distributed data

# Myth #1: Too Expensive

## Facts:

- > 2013: 1TB DRAM cluster ~\$25K
- > 2015: 1TB DRAM cluster ~\$10K  
30% reduction YoY
- > Memory Channel Storage (MCS)  
NAND in DRAM form factor, 2x speed of flash, same price as flash
- > Storage Class Memory (SCM)  
~10x slower than DRAM, Flash price, non-volatile

## Myth #2: Not Durable

### Facts:


- > IMC have durable backups and disk storage  
Active or passive replicas, transactional read-through and write-through
- > Mature IMC provide tiered storage  
DRAM - Local Swap - RDBMS/HDFS
- > Operational vs. Historical datasets  
99% of operational datasets < 10TB



## Myth #3: Flash Is Fast Enough

### Facts:

- > Flash on PCI-E is still... **a block device, i.e. disk.** A faster one. Still going through OS I/O, I/O controller, marshaling, buffering.



## Myth #4: Only In-Memory Databases

### Facts:

- > IMC is **not** a product - it is a technology  
It is applied to different products and payloads
- > In-Memory Database is important use case for **today**  
Easiest adoption and a “low-hanging fruit”
- > Streaming is an ideal use case for IMC going forward  
Streaming CAN ONLY be supported on IMC
- > Vertical and PnP products are the future  
Minimal integration, maximum benefit

# Four Use Cases

- > **Real-Time Risk Analytics**  
Able to grow book of business while reducing latencies
- > **Railroad Logistics**  
Able to ingest and process sensor data for instant logistics
- > **Energy Generation**  
Able to decide on trade vs. generate in real-time as demand spikes
- > **Oil & Gas Drilling**  
Able to provide real-time safety monitoring during fracking

# Thank you!



[www.gridgain.com](http://www.gridgain.com)



[#gridgain](https://twitter.com/gridgain)