

# Interactive SQL-on-Hadoop

from Impala to Hive/Tez to  
Spark SQL to JethroData



Ronen Ovadya, Ofir Manor,  
JethroData

# About JethroData

- Founded 2012
- Raised funding from Pitango in 2013
- Engineering in Israel, branch in New York
- Launched beta July 2014
- We're hiring!

# About Me

## Ofir Manor

- Worked in the database industry for 20 years
- Started as developer and DBA
- Worked for Oracle, Greenplum pre-sales roles
- Blogging on Big Data and Hadoop
- Currently Product Manager at JethroData

ofir@jethrodata.com

<http://ofirm.wordpress.com/>

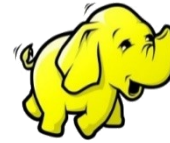
@ofirm

# Agenda

- How we got here?
- How do parallel databases on Hadoop work?
- Impala, Hive/Tez, Spark SQL
- Use case perspective
- JethroData – What? How? Demo

# SQL-on-Hadoop?

SQL?



"Stinger"



Apache Tez?



**IBM** Big SQL  
*(too serious for a cute logo)*



What about NoSQL on Hadoop?



What about Lucene / SOLR?



# Sounds Familiar?

*“Let's bring all the data that our operational systems create into one place, and keep it all forever.*

*When we'll analyze that repository, we will surely uncover critical business insights...”*

## What is this concept called?

- Today - “**Big Data**”
- Last 20 years - “**Enterprise Data Warehouse**” (EDW)

# Big Data vs. Data Warehouse?

Everyone calls themselves “Big Data” now...

But “Big Data” is lead by large web companies with new requirements:

- **Web scale** – orders of magnitude more data
  - Page views and clicks, mobile apps, sensor data etc
- **Cost** – dramatically lower price per node
- **Aversion from vendor lock-in** - open-source preference
- **Methodology** – correctness vs. agility
  - **Classical EDW / Schema-on-write** – data must be cleaned and integrated before business users can access it (*EDW* vs. *ADW*)
  - **Big Data / schema-on-read**– let the data science team handle the unfiltered crap (*in addition to some vetted schema-on-write data sets*)

# Evolution – 90s

- “Big Data” was “Enterprise Data Warehouse”
- You could either build your own using a big Unix machine and enterprise storage
  - “Scale Up”
- Or just buy a Teradata appliance

*World's first production 1TB EDW, 1992*



at





# Evolution – 2000s

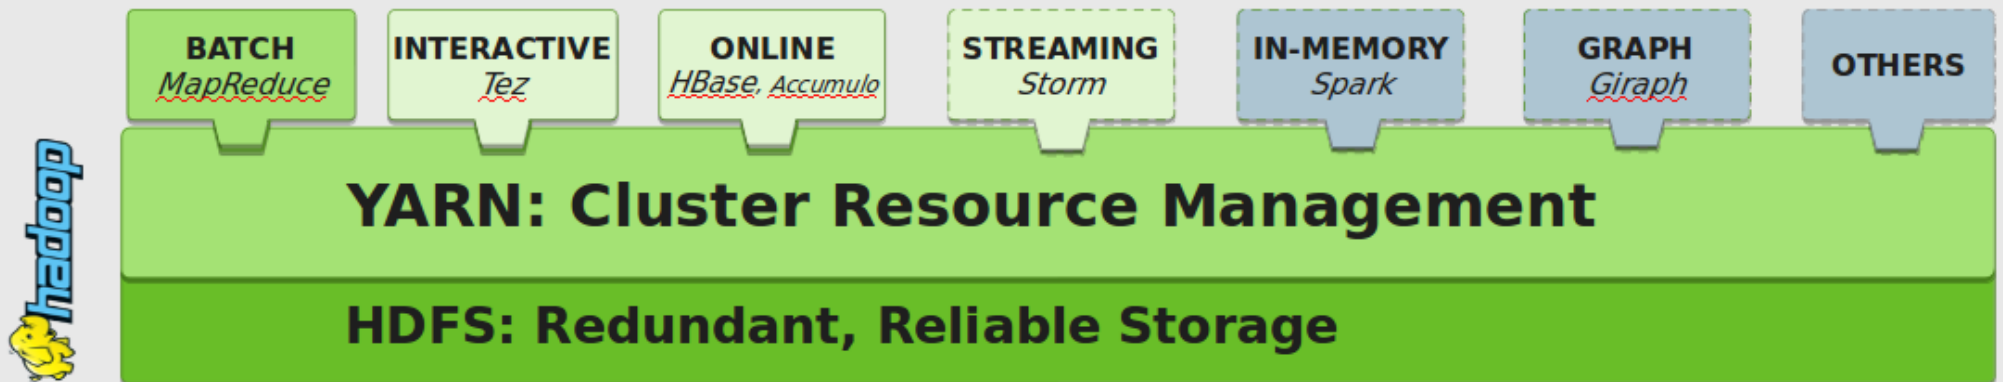
- A new generation of parallel databases -  
***“10x-100x faster, 10x cheaper”***
  - Netezza, Vertica, Greenplum, Aster, Paracel etc
  - All used a cluster of cheap servers without enterprise storage to run parallel analytic SQL
  - *Shared-Nothing Architecture /  
MPP (Massive Parallel Processing) /  
Scale-Out*

# Evolution – Our Days

- It's all about Hadoop
- Hadoop started as a batch platform - **HDFS** and **MapReduce**
- Lately, it became a shared platform for any type of parallel processing framework

*Example – Hortonworks slide*

## Data Processing Engines Run Natively IN Hadoop



# SQL-on-Hadoop

## Underlining Design

- Hadoop uses the same parallel design pattern as the parallel databases from last decade
- Surprisingly, all SQL-on-Hadoop also uses the same design pattern of previous parallel databases!

- HDFS
- MapReduce
- Tez
- Spark

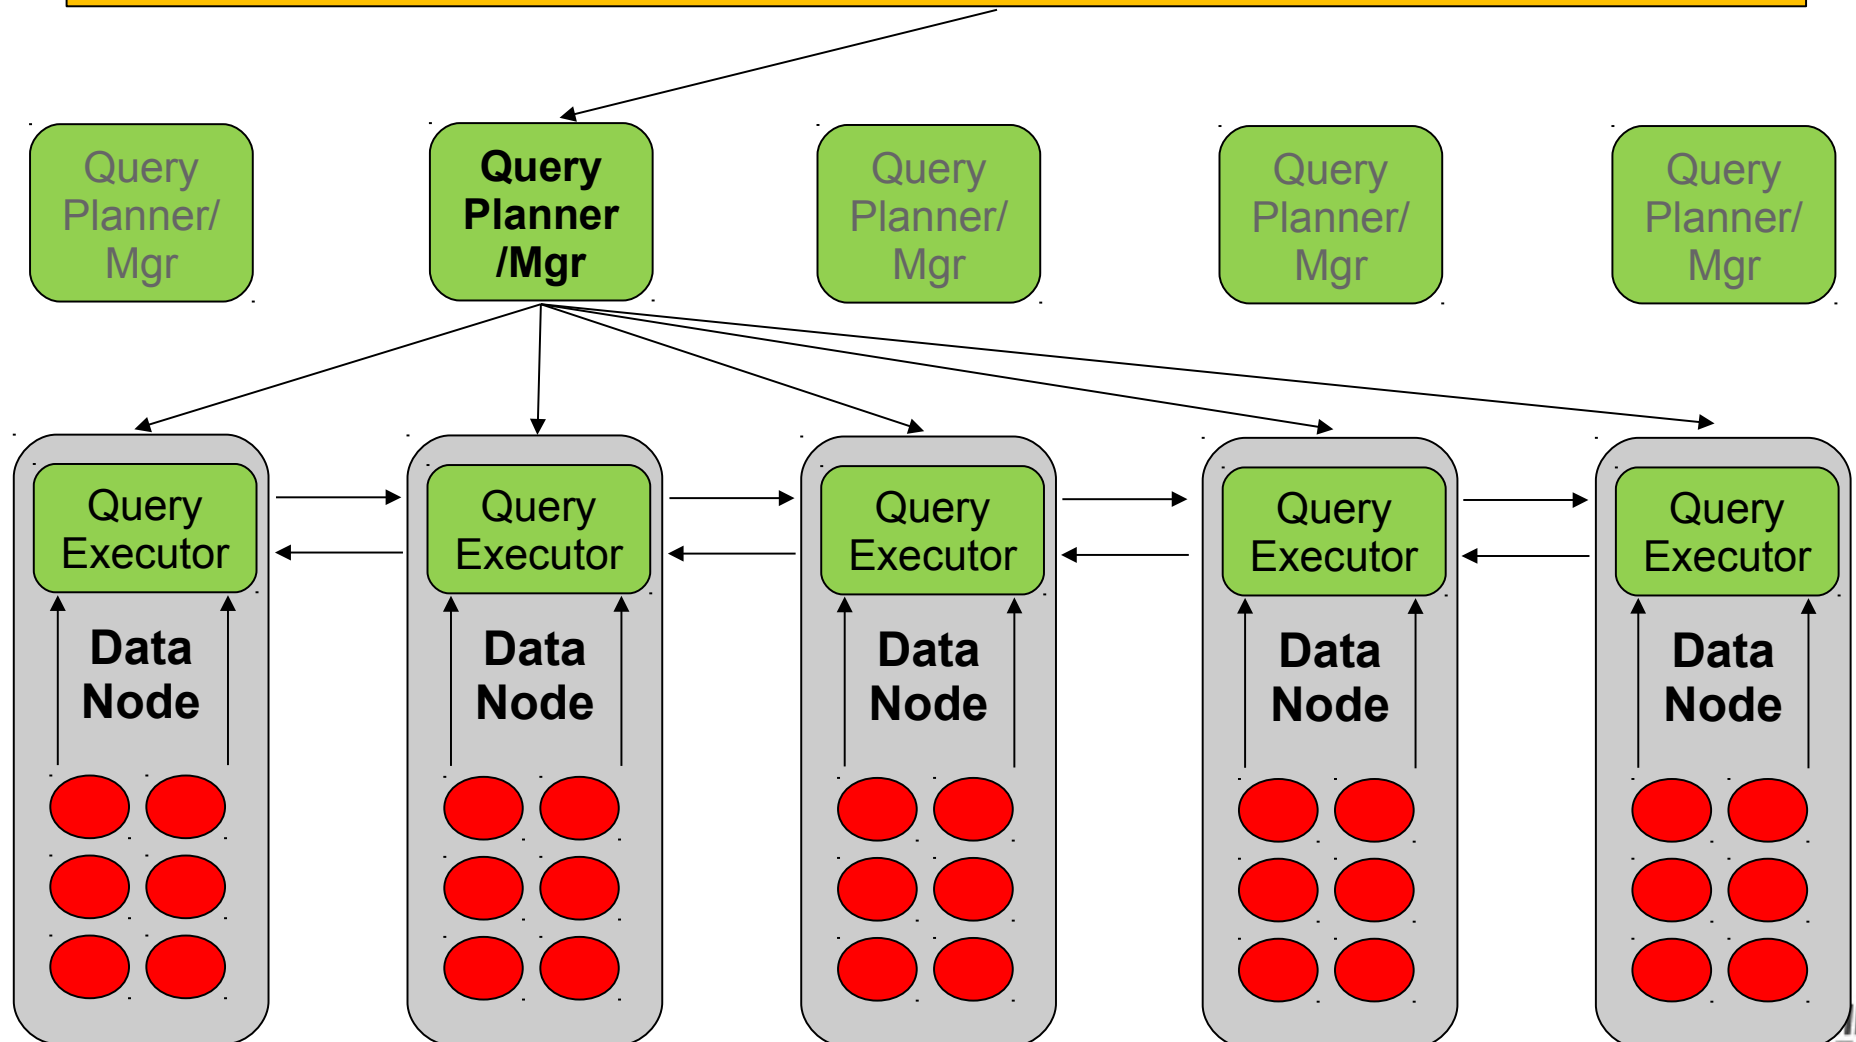
- Hive
- Impala
- Presto
- Tajo
- Drill
- Shark
- Spark SQL

- Pivotal HAWQ
- IBM BigSQL
- Teradata Aster
- Hadapt
- Rainstor
- ...

# The Parallel Design

## (Shared-Nothing MPP)

**Client:** SELECT day, sum(sales) FROM t1 WHERE prod='abc' GROUP BY day



# The Parallel Design Principles

## (Shared-Nothing MPP)

- **Full Scan**
  - Each node reads all its local portion of the table
  - Optimize with large sequential reads
- **Maximize Locality**
  - minimize inter-node work
  - Work should be evenly distributed to avoid bottlenecks
  - Avoid data and processing skew
- **That leads to a hard design-time trade-off**
  - Greatly depends the physical table organization
  - *Choose partition keys, distribution keys and sort keys*

# How to Get Decent Performance

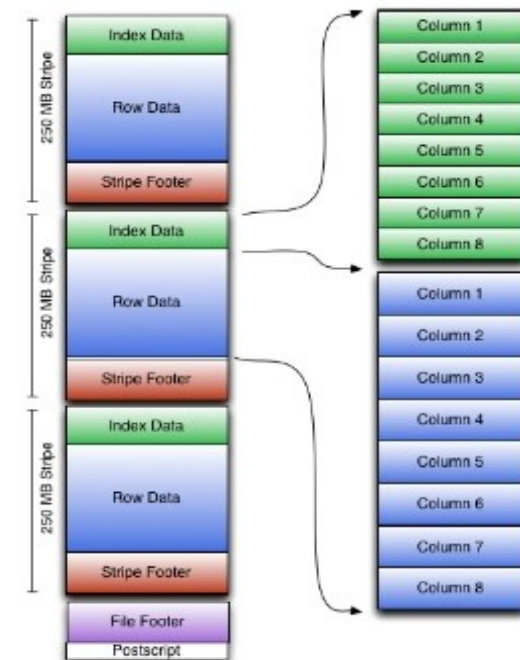
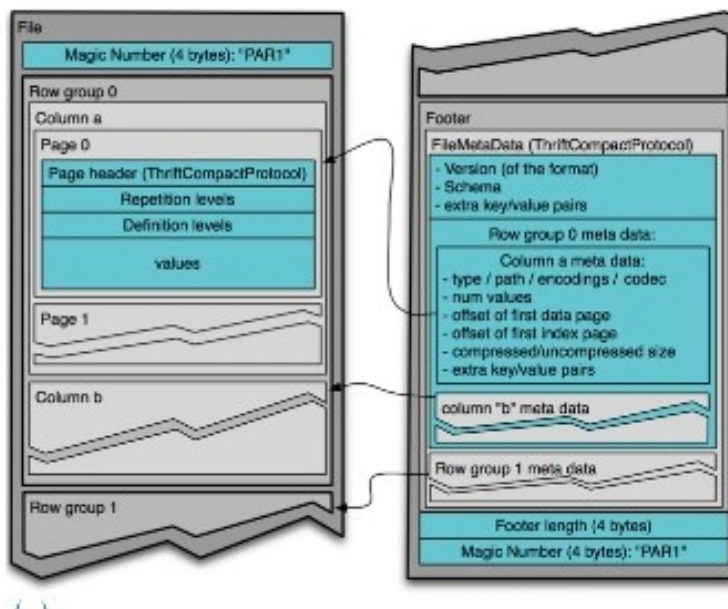
## Hive vs. Impala

1. Minimize Query I/O
  - Skip reading unneeded data
2. Optimize Query Execution
  - Pick best plan
  - Optimize each step
  - Connect the steps efficiently

# 1. Minimize Query I/O

- **Using Partitioning** (Hive, Impala etc)
  - Full scan of less data
  - Manual tuning – too few vs. too many
- **Using columnar, scan-optimized file formats**
  - “Write once, full scan a thousand times”
  - Skip unneeded columns
  - Full scan smaller files - encode and compress per-column
  - Skip blocks (*for sorted data*)
- Big effort in the Hadoop space, mostly done
  - Built two comparable formats - **ORC** and **Parquet**
  - Use the right one – Hive/ORC or Impala/Parquet

# How Parquet and ORC columnar format works?



- Data is divided into blocks - chunks of rows
- In each block, data is physically stored column by column
- Store additional metadata per column in a block, like min / max values



# 2. Optimize Query Execution

1. Pick the best execution plan – **Query Optimizer**
  - Cost-based Optimizations - currently generally weak
2. Optimize each step – **Efficient Processing**
  - Vectorized operations, expression compilation etc
3. Combine the steps efficiently - **Execution Engines**
  - Batch-oriented (*MapReduce*) – focus on recoverability
    - write intermediate results to disk after every step
  - Streaming-oriented (*Tez, Impala, HAWK etc*) – focus on performance
    - Move intermediate results directly between processes
    - Required much more resources at once
  - Hybrid (*Spark*) – enjoy both worlds
    - Stream and re-use intermediate results, optimize for in-memory
    - But can recover / recompute on node failure

# Impala



- ***Impala Highlights***

- Basic partitioning (partition per key)
- Optimized I/O with Parquet
- Built their own streaming execution engine
- Efficient processing
- Basic cost-based query optimizations

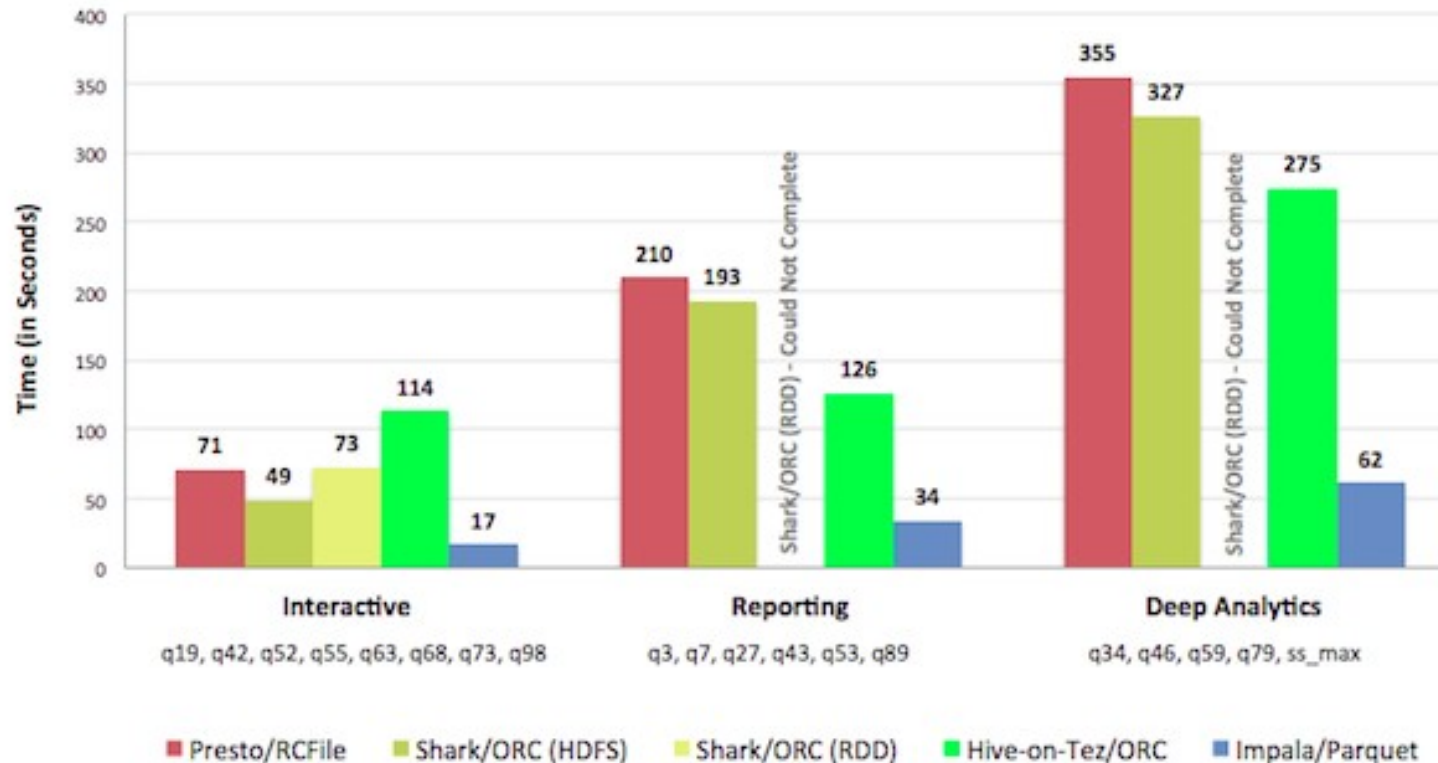
- ***Impala vs. Presto vs. Drill vs. Tajo***

- Generally same high-level design
- Four independent teams implementing it in their own way
- Impala way ahead of the pack
  - Performance, functionality, market share, support

# Impala



Single User Response Time  
(Lower bars are better)



- Sample benchmark results from Cloudera (May 2014)
- 20 nodes running TPCDS data, scale factor 15,000 (GB)
- *Impala 1.3.0 vs. Hive 0.13 on Tez vs. Shark 0.9.2 vs. Presto 0.6.0*

Source: [Cloudera blog](#) - New SQL Choices in the Apache Hadoop Ecosystem: Why Impala Continues to Lead

# Hive



- ***Hive Highlights***

- Rich SQL support
- Basic partitioning (partition per key)
- Optimized I/O with ORC
- Cost-based optimizer coming (Hive 0.14)
- Efficient processing – vectorized query execution
- Reliable execution (**MapReduce**) or fast execution (**Tez**)
- Hive on Spark (**Shark**)
  - Have recently reached “End-of-life”, sort-of

# Spark SQL



- ***Spark SQL Highlights***

- Very early days / alpha state
  - Announced Mar 2014
- A SQL-on-Spark solution written from scratch
- Should support reading / writing from Hadoop, specifically from/to Hive and Parquet
- Could become an interesting player next year
  - Mostly vs. Impala

# Query Use Cases

## Querying Raw Data

### Raw Data



- Ad-hoc “Data Science” investigative work
  - Typically in the minutes to hours range
- Need to make sense of many, ever changing data sources
- Need to make sense of text / semi-structured / dynamic schema sources
- Need to mix SQL, text processing (UDFs) and machine learning algorithm in a manual, multi-step fashion

# Query Use Cases

## Reporting

Raw Data

Reporting



- Internal analyst teams repeatedly run reports on common data sets
- Likely cleaned and vetted, potentially aggregated, shared across many users
- Use latest Hive/ORC/Tez or Impala/Parquet
  - Improves response time from hours to minutes

# Query Use Cases

## Pre-Compute Specific Queries (1)



- Queries embedded in external websites / apps / dashboards
  - Customers expecting page loads of a second or two
- **Must pre-compute results to get response time**
- “Old Style” implementation
  1. Daily massive batch computation (**MapReduce**)
  2. Results are typically pushed to an external OLAP solution (**cubes**) or to a key-value store (**HBase** / Redis etc)
- Great performance but only for a few queries or counters
- “Freshness”: data is typically a day or more behind production
- Complexity: multiple steps and systems, massive daily spike in computation, storage, data movement



# Query Use Cases

## Pre-Compute Specific Queries (2)



### “New Style” implementation – **Stream Processing**

- Continuously update pre-computed results as new events arrive
  - Events pile up in some queue (**Apache Kafka**)
  - Process events in micro-batches (**Apache Storm**)
  - The new computation results are constantly stored / updated in a key-value store
- Solves the “freshness” problem
- Not suitable for complex pre-computations - best for counters
- Bleeding-edge technology

# Query Use Cases

## Fast Ad-hoc Queries



- **Interactive BI is the “Holy Grail” of SQL on Hadoop**
- Analysts / business users want to interact with select data sets from their BI tool
  - drag columns in the BI tool, “slice n' dice”, drill-downs
  - Response time of seconds to tens of seconds from their BI tool
- Existing solutions are too slow – users are stuck with reporting
- Very hard to achieve with existing Hadoop technologies
- Need a different solution – maybe something that have worked for databases in the last 30 years?
- It's time to introduce **JethroData**...

# JethroData

- **Index-based, columnar SQL-on-Hadoop**
  - Delivers interactive BI on Hadoop
  - Focused on ad-hoc queries from BI tools
  - The more you drill down, the faster you go
- **SQL-on-Hadoop**
  - Column data and indexes are stored on HDFS
  - Compute is on dedicated edge nodes
    - Non-intrusive, nothing installed on the Hadoop cluster
    - Our own SQL engine, not using MapReduce / Tez
  - Supports standard ANSI SQL, ODBC/JDBC

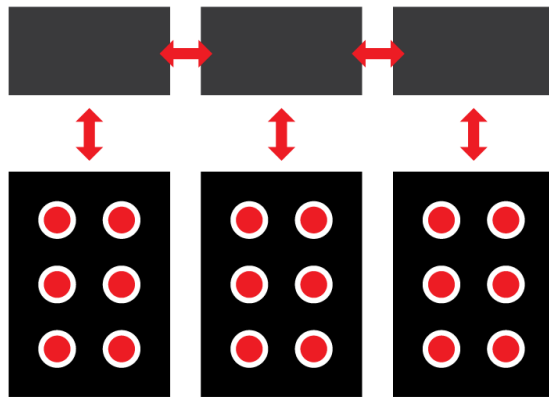
# JethroData

## Working Differently

### Full Scan / Brute Force

(All SQL-on-Hadoop Solutions)

1. Read entire dataset. Every time.

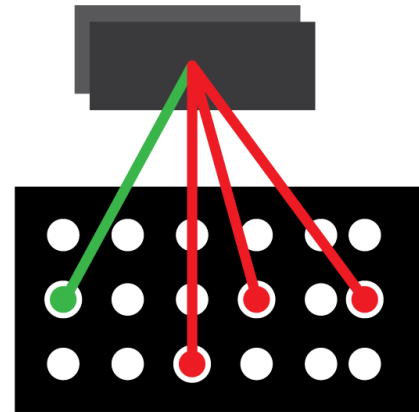


- *Massive # of unnecessary I/Os*
- *Increasing demands on CPU and memory*

### Index

(JethroData)

1. Analyze index
2. Fetch only needed data

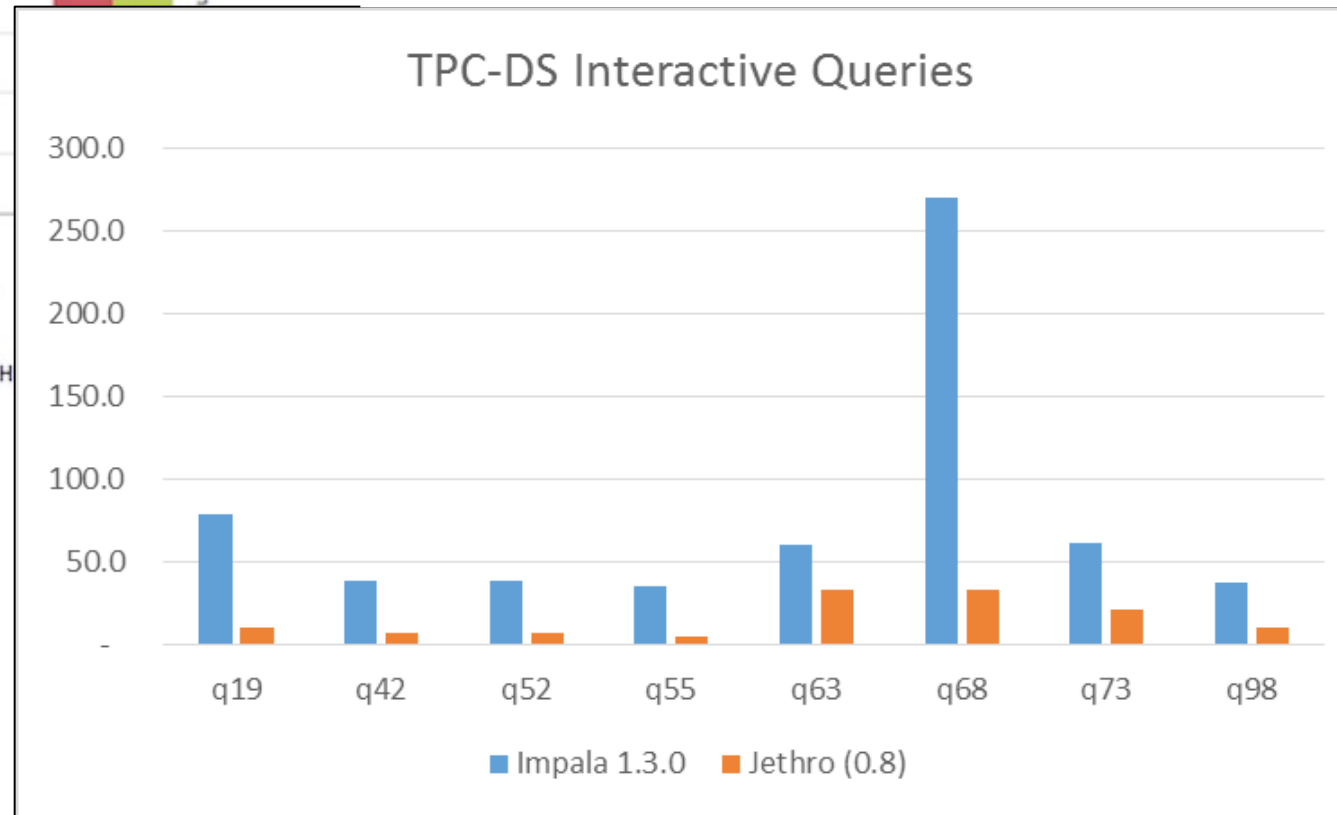
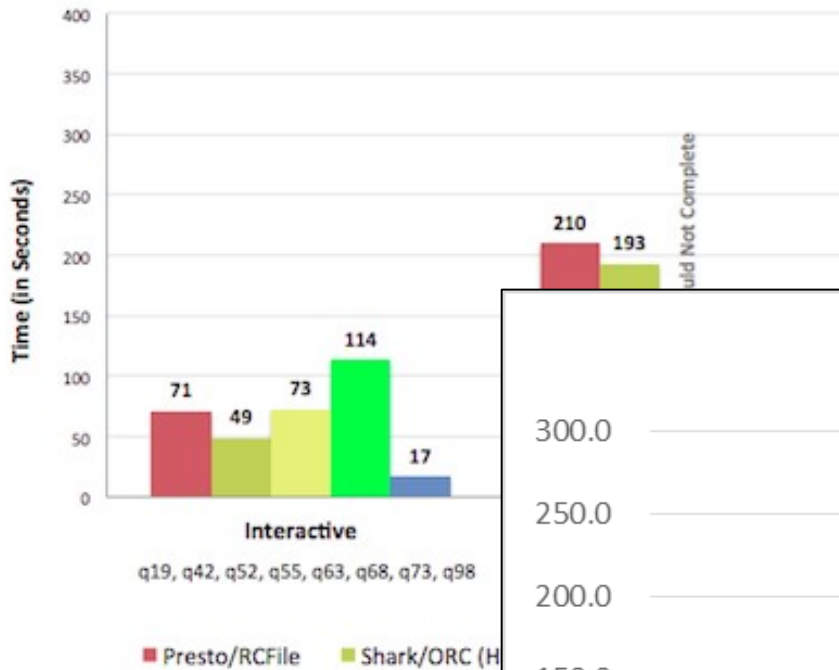


- *Drastically lower cluster load*
- *Low I/O CPU and memory usage*

# Comparing to Impala

Single User Response Time  
(Lower bars are better)

- Using Cloudera's benchmark (data and SQLs)
- Scale Factor 1000(GB)
- Running on AWS



# Demo 1

*Access JethroData from a remote JDBC client - SQL Workbench*

*Small CDH5 cluster on AWS*

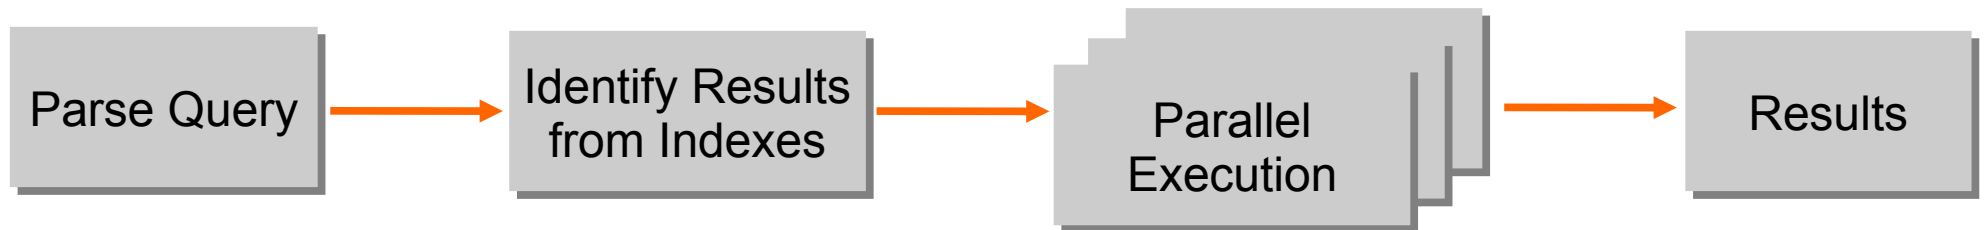
# Jethro Indexes

- Indexes map each column value to a set of rows
- Jethro stores indexes as hierarchical compressed bitmaps

Value	Rows
FR	rows 5,9,10,11,14
IL	rows 1,3,7,12,13
US	rows 2,4,6,8,15

- Very fast query operations – `AND` / `OR` / `NOT`
  - Can process the entire `WHERE` clause to a final list of rows.
  - Fully indexed – all columns are automatically indexed
- **INSERT Performance**
  - Jethro Indexes are append-only
  - If needed, a new, repeating entries are allowed
  - `INSERT` is very fast – files are appended, no random read/write
  - Compatible with HDFS
  - Periodic background merge (non-blocking)

# Jethro Execution Logic



## Index Phase

Use all relevant indexes to compute which rows are needed for query results

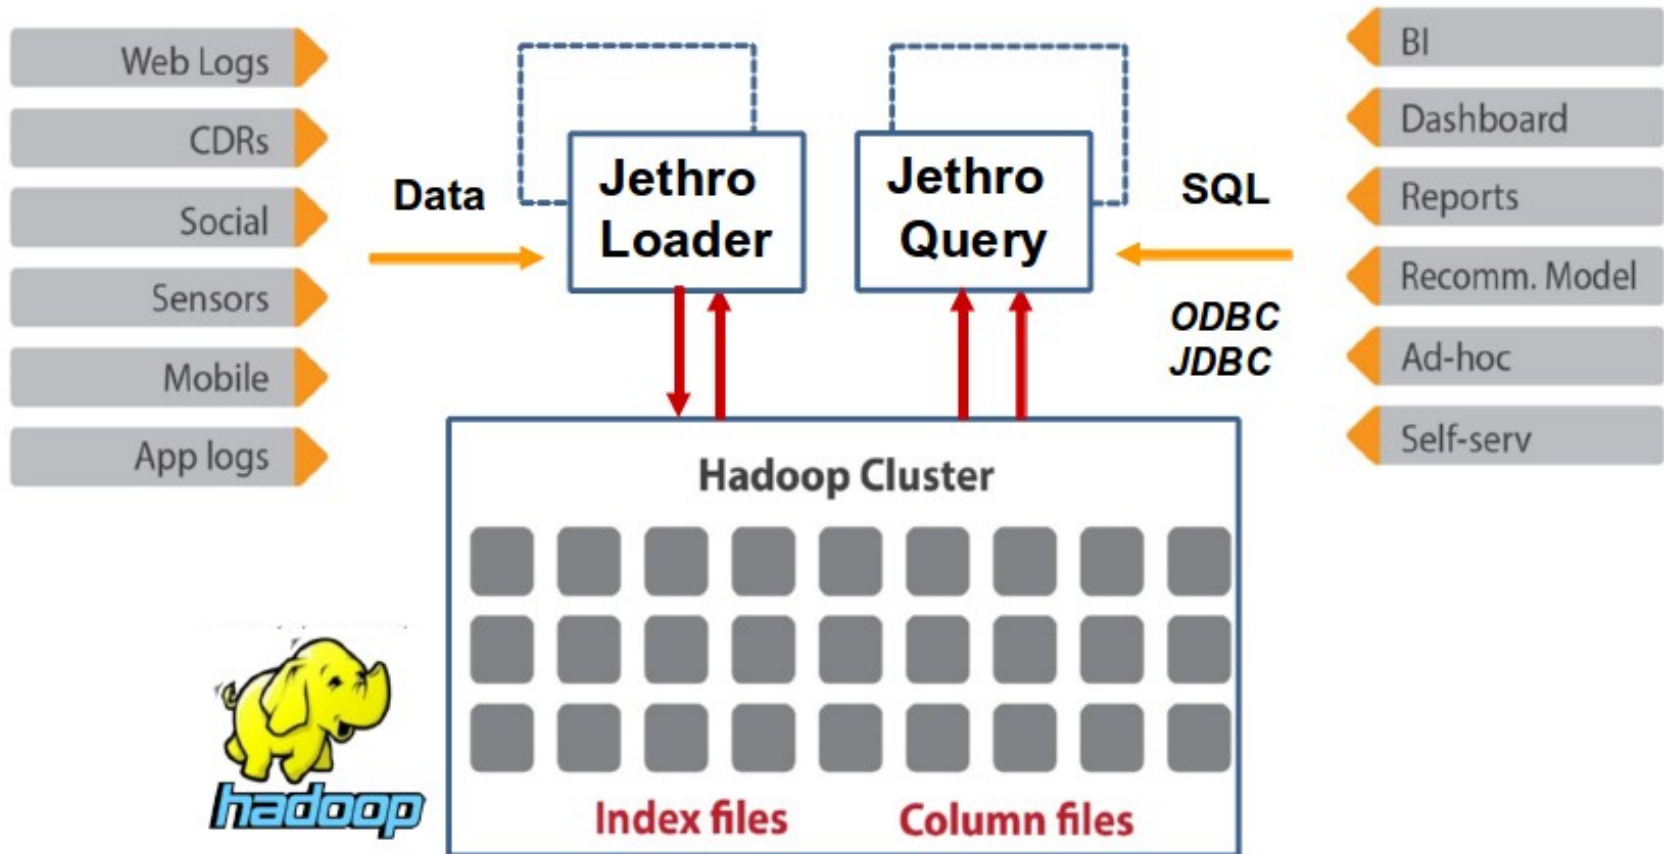
Divide the work to multiple parallel units of work.

## Parallel Execution Phase

- Parallelize both fetch and compute
- Works with or without partitions
- Currently multi-threaded in a single server, designed for multi-node execution



# Technical Architecture



# Minimize Query I/O

## part 1

- Automatically **combine all relevant indexes** for query execution – dramatically reduce fetch I/O
  - Generate a bitmap of all relevant rows after the entire WHERE clause
  - Skip indexes when their value is minimal (`WHERE a>0`) or when indexes are not applicable (`WHERE function(col)=5`)
- Use **Jethro columnar file format** to:
  - Skip columns
  - Encode and compress data (smaller files)
  - Skip blocks (*when a column must be scanned*)
- **HDFS I/O optimizations**
  - Optimized fetch size for bulk fetches
  - Using skip scans to convert random reads to single I/O, if possible

# Minimize Query I/O

## part 2

- **Automatic Local Cache**
  - Simple to set up
  - Metadata is automatically cached locally by priority, in the background
  - Can ask to cache specific columns or tables
  - Improves latency and reduces HDFS roundtrips
- **Use partitioning?**
  - Only small effect – we only read relevant rows, with or without partitioning
  - At the high-end, helps operating on smaller bitmaps
  - Mostly for rolling window operations

# Optimize Query Execution

- **Query Optimizer** - pick the best plan (set of steps)
  - Using up-to-date detailed statistics from the indexes
  - For example, star transformation
- **Efficient Processing** – optimize each step
  - Efficient bitmap operations
- **Execution Engines** – combine the steps efficiently
  - Multi-threaded, streaming, parallel execution
  - Parallelize everything:  
Fetch, Filter, Join, Aggregate, Sort etc

# Scalability

## scale to 100s of billions of rows

- Scale out **HDFS** (Hadoop Cluster), if needed
  - Provide extra storage or extra I/O performance (rare)
- Scale out **Jethro nodes**
  - Jethro query nodes are stateless
  - Add nodes to support additional concurrent queries
- Leverage **partitioning**
  - Support rolling window maintenance at scale
  - Works best with a few billion rows per partition – usually partition for maintenance, not performance

# High-Availability

- Leverage all **HDFS HA** goodies
  - DataNode Replication
  - NameNode HA
  - HDFS Snapshots
  - Any HDFS DR solution
- All Jethro nodes are **stateless**
  - Service auto-start on restart
  - Can start and stop them on demand

# Demo 2

*Access JethroData from Tableau Server over ODBC*

*Small CDH5 cluster on AWS*

# Questions?

Talk to us:

[ronen@jethrodata.com](mailto:ronen@jethrodata.com)

[ofir@jethrodata.com](mailto:ofir@jethrodata.com)

Join our beta!

[www.jethrodata.com](http://www.jethrodata.com)