

# CSE 344

## Lectures 9: Relational Algebra

# Announcements

- Homework 2 due tonight!
- Homework 3 is posted, due next Thursday!
- Webquiz 3 due tomorrow night!
- Reminder about Discussion Board
  - Post your question on HW/WQ here
  - Feel free to answer your friends' questions and discuss concepts (no solution please 😊)
  - You can set alert to get email notification
- Today's lecture: 2.4 and 5.1-5.2

# Where We Are

- Motivation for using a DBMS for managing data
- SQL, SQL, SQL
  - Declaring the schema for our data (CREATE TABLE)
  - Inserting data one row at a time or in bulk (INSERT/.import)
  - Modifying the schema and updating the data (ALTER/UPDATE)
  - Querying the data (SELECT)
  - Tuning queries (CREATE INDEX)
- Next step: More knowledge of how DBMSs work
  - Relational algebra and query execution
  - Client-server architecture

# Relational Algebra

# Sets v.s. Bags

- Sets:  $\{a,b,c\}$ ,  $\{a,d,e,f\}$ ,  $\{ \}$ , . . .
- Bags:  $\{a, a, b, c\}$ ,  $\{b, b, b, b, b\}$ , . . .

Relational Algebra has two semantics:

- Set semantics = standard Relational Algebra
- Bag semantics = extended Relational Algebra

# Relational Algebra Operators

- Union  $\cup$ , intersection  $\cap$ , difference  $-$
- Selection  $\sigma$
- Projection  $\Pi$
- Cartesian product  $\times$ , join  $\bowtie$
- Rename  $\rho$
- Duplicate elimination  $\delta$
- Grouping and aggregation  $\gamma$
- Sorting  $\tau$

RA

Extended RA

# Why learn RA?

- SQL incorporates RA at its center
- When DBMS processes a query, it is translated into an RA expression internally and is used by the query optimizer
- Why Algebra?

# Why learn RA?

- Why Algebra?
  - Has both Operators and Atomic Operands
  - $(x+y) * (z - 3)$
  - Similarly,  
 $\pi_{\text{zip}} (\sigma_{\text{disease}=\text{'heart'}}(\text{Patient}))$



# Union and Difference

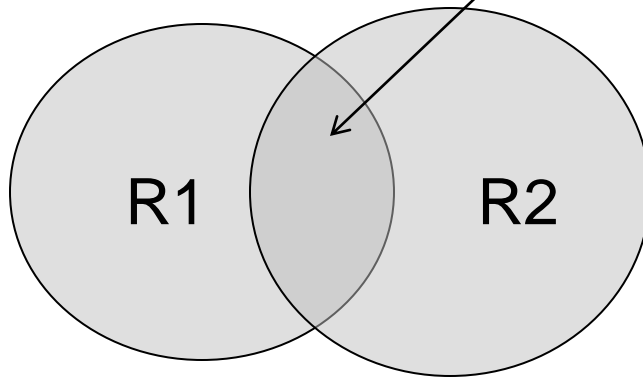
$$R1 \cup R2$$
$$R1 - R2$$

For set operations, R1 and R2

- must have identical schemas
- their attributes must have the same order, i.e. R1(A, B) and R2(B, A) is not allowed

# What about Intersection ?

- Can you derive  $R1 \cap R2$  using union/minus?



$$R2 - (R2 - R1)$$

$$R1 - (R1 - R2)$$

$$(R1 \cup R2) - ((R1 \cup R2) - R1) - ((R1 \cup R2) - R2)$$

# What about Intersection ?

- Derived operator using minus

$$R1 \cap R2 = R1 - (R1 - R2)$$

- Derived using join (will explain later)

$$R1 \cap R2 = R1 \bowtie R2$$

# Union, Difference, Intersection over Bags

$R1 \cup R2$

$R1 - R2$

$R1 \cap R2$

What do they mean over bags ?

R1

A
1
1
1

R2

A
1
1

How many 1's in

- $R1 \cup R2$ : 5
- $R1 - R2$ : 1
- $R2 - R1$ : 0
- $R1 \cap R2$ : 2

# Selection

- Returns all tuples which satisfy a condition

$$\sigma_c(R)$$

What does Selection  
Correspond to in SQL?  
Ans: WHERE clause

- Examples
  - $\sigma_{\text{Salary} > 40000}(\text{Employee})$
  - $\sigma_{\text{name} = \text{"Smith"}}(\text{Employee})$
- The condition c can be =, <, ≤, >, ≥, <>

Employee

SSN	Name	Salary
1234545	John	200000
5423341	Smith	600000
4352342	Fred	500000

$\sigma_{\text{Salary} > 40000}$  (Employee)

SSN	Name	Salary
5423341	Smith	600000
4352342	Fred	500000

# Projection

- Eliminates columns

$$\Pi_{A_1, \dots, A_n} (R)$$

What does Projection  
Correspond to in SQL?  
Ans: SELECT clause

- Example: project social-security number and names:
  - $\Pi_{SSN, Name} (Employee)$
  - Answer(SSN, Name)

Different semantics over sets or bags! Why?

Employee

SSN	Name	Salary
1234545	John	20000
5423341	John	60000
4352342	John	20000

$\Pi_{\text{Name,Salary}}(\text{Employee})$

Name	Salary
John	20000
John	60000
John	20000

Bag semantics

Name	Salary
John	20000
John	60000

Set semantics

Which is more efficient? Ans: Bag

Checking and removing duplicates is expensive



# Composing RA Operators

Patient

no	name	zip	disease
1	p1	98125	flu
2	p2	98125	heart
3	p3	98120	lung
4	p4	98120	heart

$\pi_{\text{zip,disease}}(\text{Patient})$

zip	disease
98125	flu
98125	heart
98120	lung
98120	heart

$\sigma_{\text{disease='heart'}}(\text{Patient})$

no	name	zip	disease
2	p2	98125	heart
4	p4	98120	heart

$\pi_{\text{zip}}(\sigma_{\text{disease='heart'}}(\text{Patient}))$

zip
98120
98125

# Cartesian/Cross Product

- Each tuple in R1 with each tuple in R2

$$R1 \times R2$$

- Rare in practice; mainly used to express joins

# Cross-Product Example

**Employee**

Name	SSN
John	9999999999
Tony	7777777777

**Dependent**

EmpSSN	DepName
9999999999	Emily
7777777777	Joe

**Employee  $\bowtie$  Dependent**

Name	SSN	EmpSSN	DepName
John	9999999999	9999999999	Emily
John	9999999999	7777777777	Joe
Tony	7777777777	9999999999	Emily
Tony	7777777777	7777777777	Joe

Disambiguate  
attributes if necessary  
**Employee.EmpSSN**  
**Dependent.EmpSSN**

# Renaming

- Changes the schema, not the instance

$$\rho_{B1, \dots, Bn}(R)$$

- Example:

- $\rho_{N, S}(\text{Employee}) \rightarrow \text{Answer}(N, S)$
- Given  $R(A, B)$ 
  - $\rho_S(R)$  : Renamed relation  $S(A, B)$
  - $\rho_{S(X, Y)}(R)$  or  $S = \rho_{X, Y}(R)$ : Renamed relation  $S(X, Y)$
  - Sometimes written as  $S = \rho_{A \rightarrow X, B \rightarrow Y}(R)$

Not really used by systems, but needed on paper

# Natural Join

$$R1 \bowtie R2$$

- Meaning:  $R1 \bowtie R2 = \Pi_A(\sigma(R1 \times R2))$
- Where:
  - Selection  $\sigma$  checks equality of all common attributes
  - Projection eliminates duplicate common attributes

# Natural Join Example

**R**

A	B
X	Y
X	Z
Y	Z
Z	V

**S**

B	C
Z	U
V	W
Z	V

**R** ⋈ **S** =

$\Pi_{ABC}(\sigma_{R.B=S.B}(R \times S))$

A	B	C
X	Z	U
X	Z	V
Y	Z	U
Y	Z	V
Z	V	W

# Natural Join Example 2

AnonPatient P

age	zip	disease
54	98125	heart
20	98120	flu

Voters V

name	age	zip
p1	54	98125
p2	20	98120

$P \bowtie V$

age	zip	disease	name
54	98125	heart	p1
20	98120	flu	p2

# Natural Join

- Given schemas  $R(A, B, C, D)$ ,  $S(A, C, E)$ , what is the schema of  $T = R \bowtie S$  ?
- Ans:  $T(A, B, C, D, E)$
- Given  $R(A, B, C)$ ,  $S(D, E)$ , what is  $R \bowtie S$  ?
- Ans:  $R \times S$
- Given  $R(A, B)$ ,  $S(A, B)$ , what is  $R \bowtie S$  ?
- Ans:  $R \cap S$



# Theta Join

- A join that involves a predicate

$$R1 \bowtie_{\theta} R2 = \sigma_{\theta} (R1 \times R2)$$

- Here  $\theta$  can be any condition
- For our voters/disease example:

$$P \bowtie_{P.zip = V.zip \text{ and } P.age < V.age + 5 \text{ and } P.age > V.age - 5} V$$

# Equijoin

- A theta join where  $\theta$  is an equality

$$R1 \bowtie_{A=B} R2 = \sigma_{A=B} (R1 \times R2)$$

- This is by far the most used variant of join in practice

# Equijoin Example

AnonPatient P

age	zip	disease
54	98125	heart
20	98120	flu

Voters V

name	age	zip
p1	54	98125
p2	20	98120

$P \bowtie_{P.age=V.age} V$

age	P.zip	disease	name	V.zip
54	98125	heart	p1	98125
20	98120	flu	p2	98120

# Join Summary

- **Theta-join:**  $R \bowtie_{\theta} S = \sigma_{\theta}(R \times S)$ 
  - Join of R and S with a join condition  $\theta$
  - Cross-product followed by selection  $\theta$
- **Equijoin:**  $R \bowtie_{\theta} S = \pi_A (\sigma_{\theta}(R \times S))$ 
  - Join condition  $\theta$  consists only of equalities
  - Projection  $\pi_A$  drops all redundant attributes
- **Natural join:**  $R \bowtie S = \pi_A (\sigma_{\theta}(R \times S))$ 
  - Equijoin
  - Equality on **all** fields with same name in R and in S

# So Which Join Is It ?

- When we write  $R \bowtie S$  we usually mean an equijoin, but we often omit the equality predicate when it is clear from the context

# More Joins

- **Outer join**
  - Include tuples with no matches in the output
  - Use NULL values for missing attributes
- Variants
  - Left outer join
  - Right outer join
  - Full outer join

# Outer Join Example

AnonPatient P

age	zip	disease
54	98125	heart
20	98120	flu
33	98120	lung

AnnonJob J

job	age	zip
lawyer	54	98125
cashier	20	98120

P  $\bowtie$  V

age	zip	disease	job
54	98125	heart	lawyer
20	98120	flu	cashier
33	98120	lung	null

# Some Examples

Supplier(sno,sname,scity,sstate)

Part(pno,pname,psize,pcolor)

Supply(sno,pno,qty,price)

Q2: Name of supplier of parts with size greater than 10

$\pi_{\text{sname}}(\text{Supplier} \bowtie \text{Supply} \bowtie (\sigma_{\text{psize} > 10}(\text{Part})))$

Q3: Name of supplier of red parts or parts with size greater than 10

$\pi_{\text{sname}}(\text{Supplier} \bowtie \text{Supply} \bowtie (\sigma_{\text{psize} > 10}(\text{Part}) \cup \sigma_{\text{pcolor} = \text{'red'}}(\text{Part})))$



# From SQL to RA

# From SQL to RA

Product(pid, name, price)

Purchase(pid, cid, store)

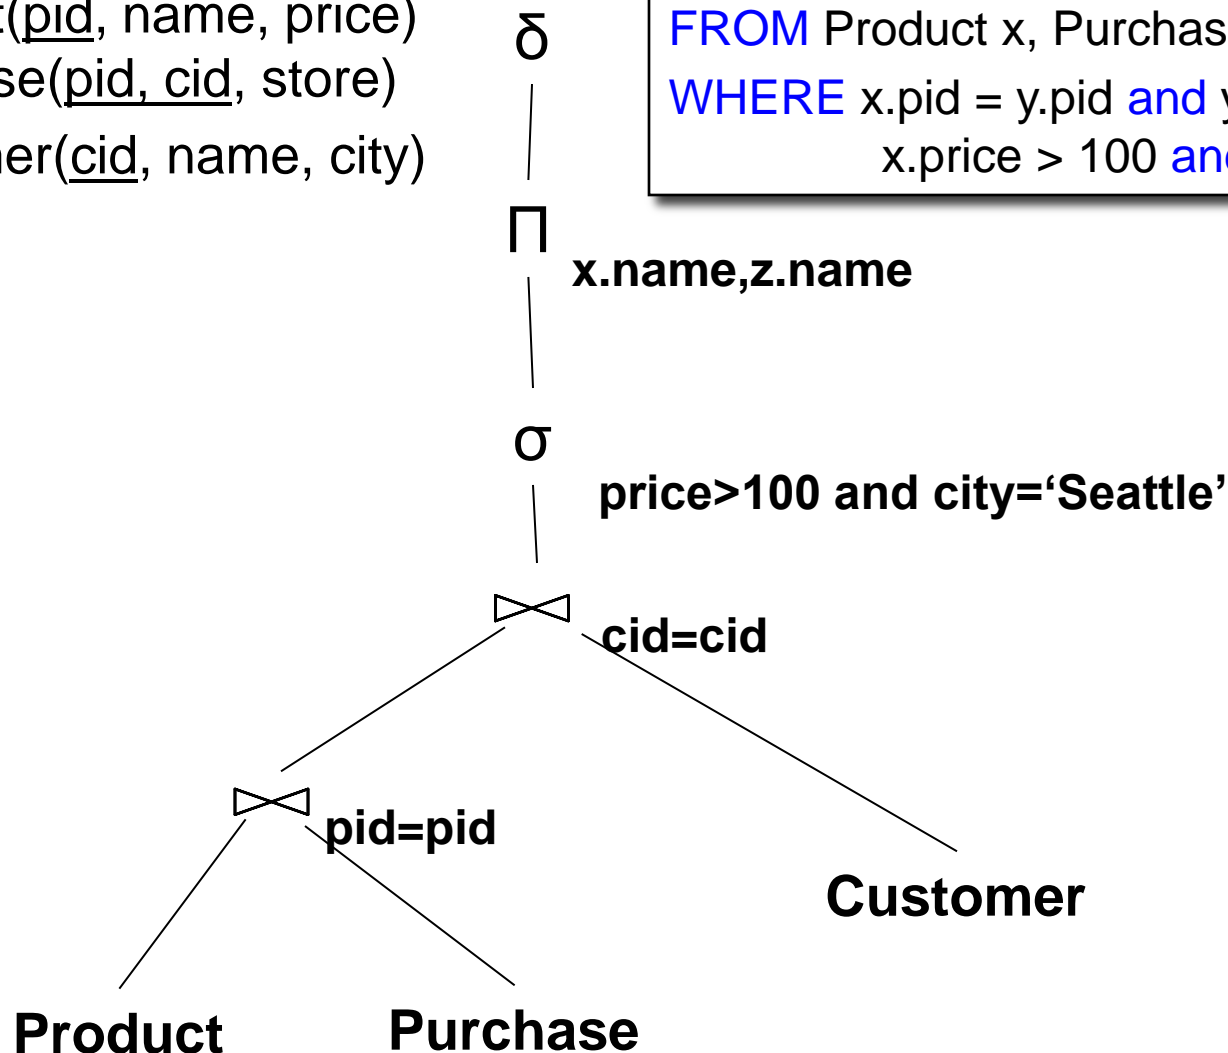
Customer(cid, name, city)

```
SELECT DISTINCT x.name, z.name
FROM Product x, Purchase y, Customer z
WHERE x.pid = y.pid and y.cid = z.cid and
      x.price > 100 and z.city = 'Seattle'
```

# From SQL to RA

Product(pid, name, price)  
Purchase(pid, cid, store)  
Customer(cid, name, city)

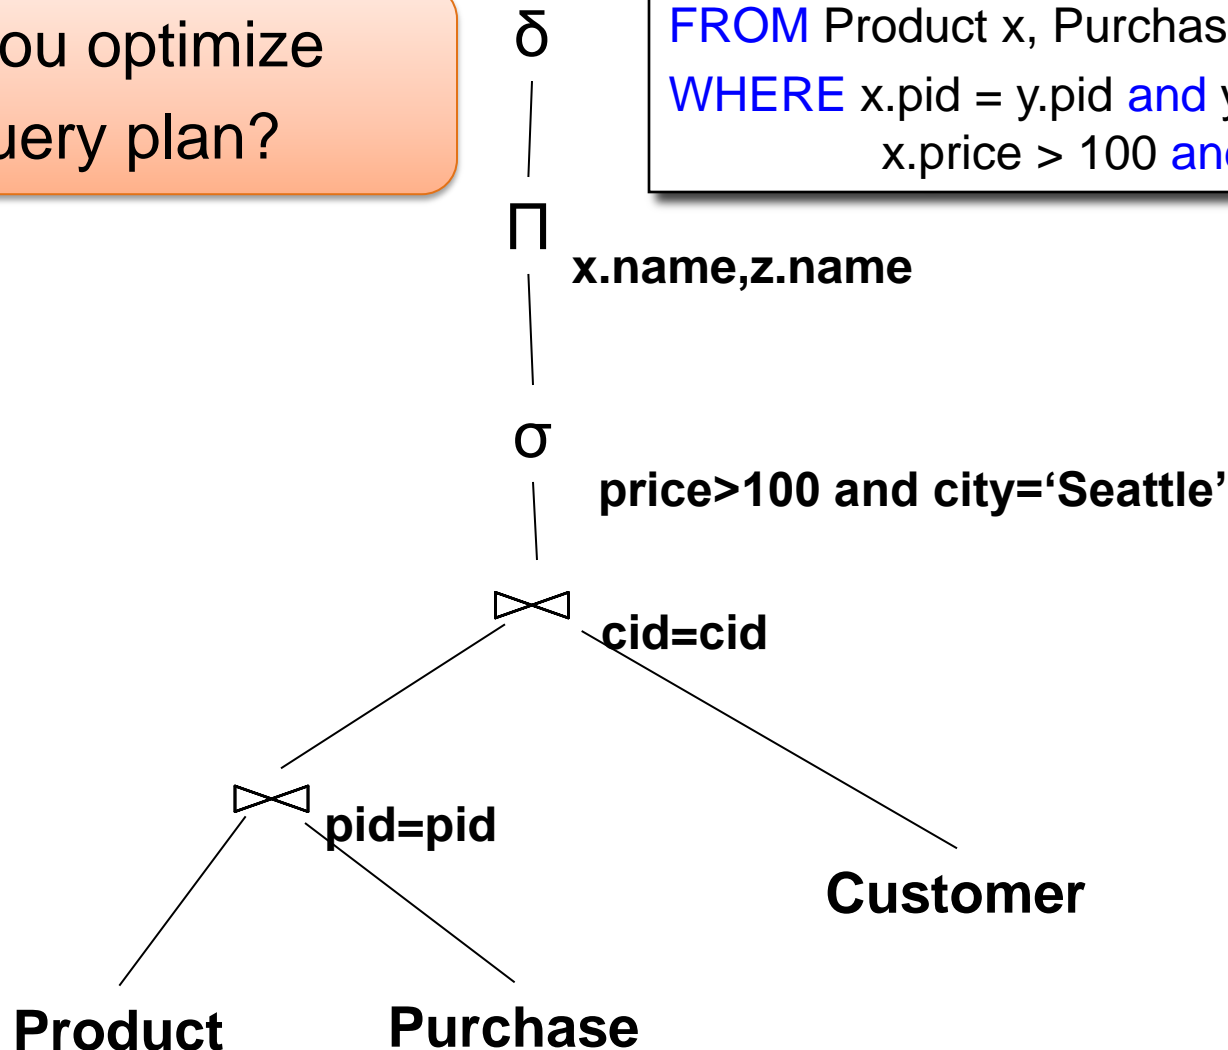
```
SELECT DISTINCT x.name, z.name  
FROM Product x, Purchase y, Customer z  
WHERE x.pid = y.pid and y.cid = z.cid and  
       x.price > 100 and z.city = 'Seattle'
```



# From SQL to RA

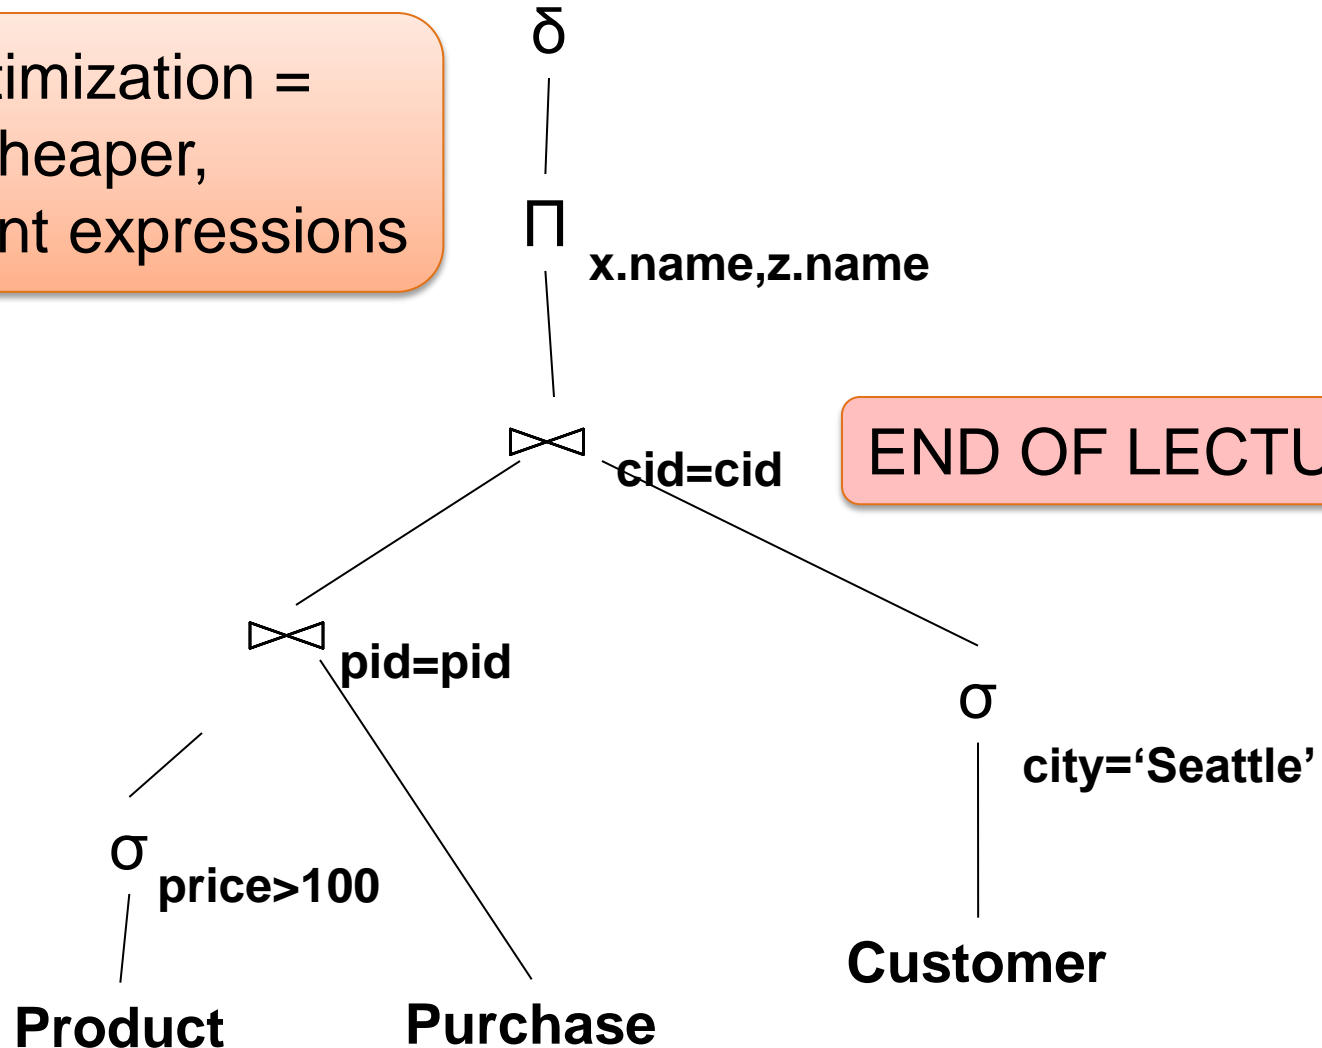
Can you optimize this query plan?

```
SELECT DISTINCT x.name, z.name
FROM Product x, Purchase y, Customer z
WHERE x.pid = y.pid and y.cid = z.cid and
      x.price > 100 and z.city = 'Seattle'
```



# An Equivalent Expression

Query optimization =  
finding cheaper,  
equivalent expressions



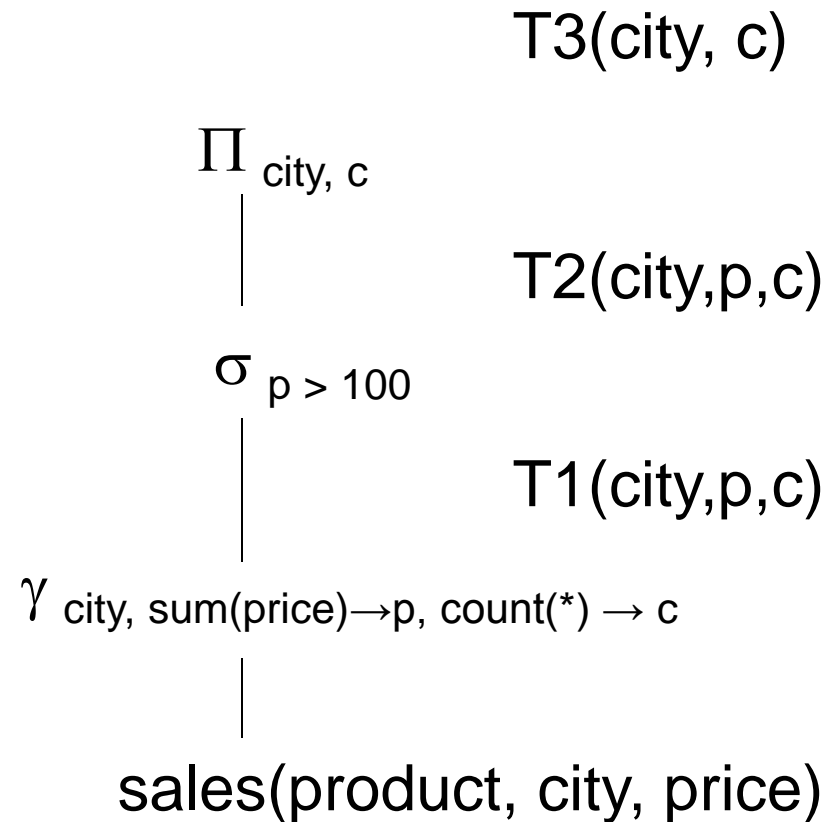
END OF LECTURE09

# Extended RA: Operators on Bags

- Duplicate elimination  $\delta$
- Grouping  $\gamma$
- Sorting  $\tau$

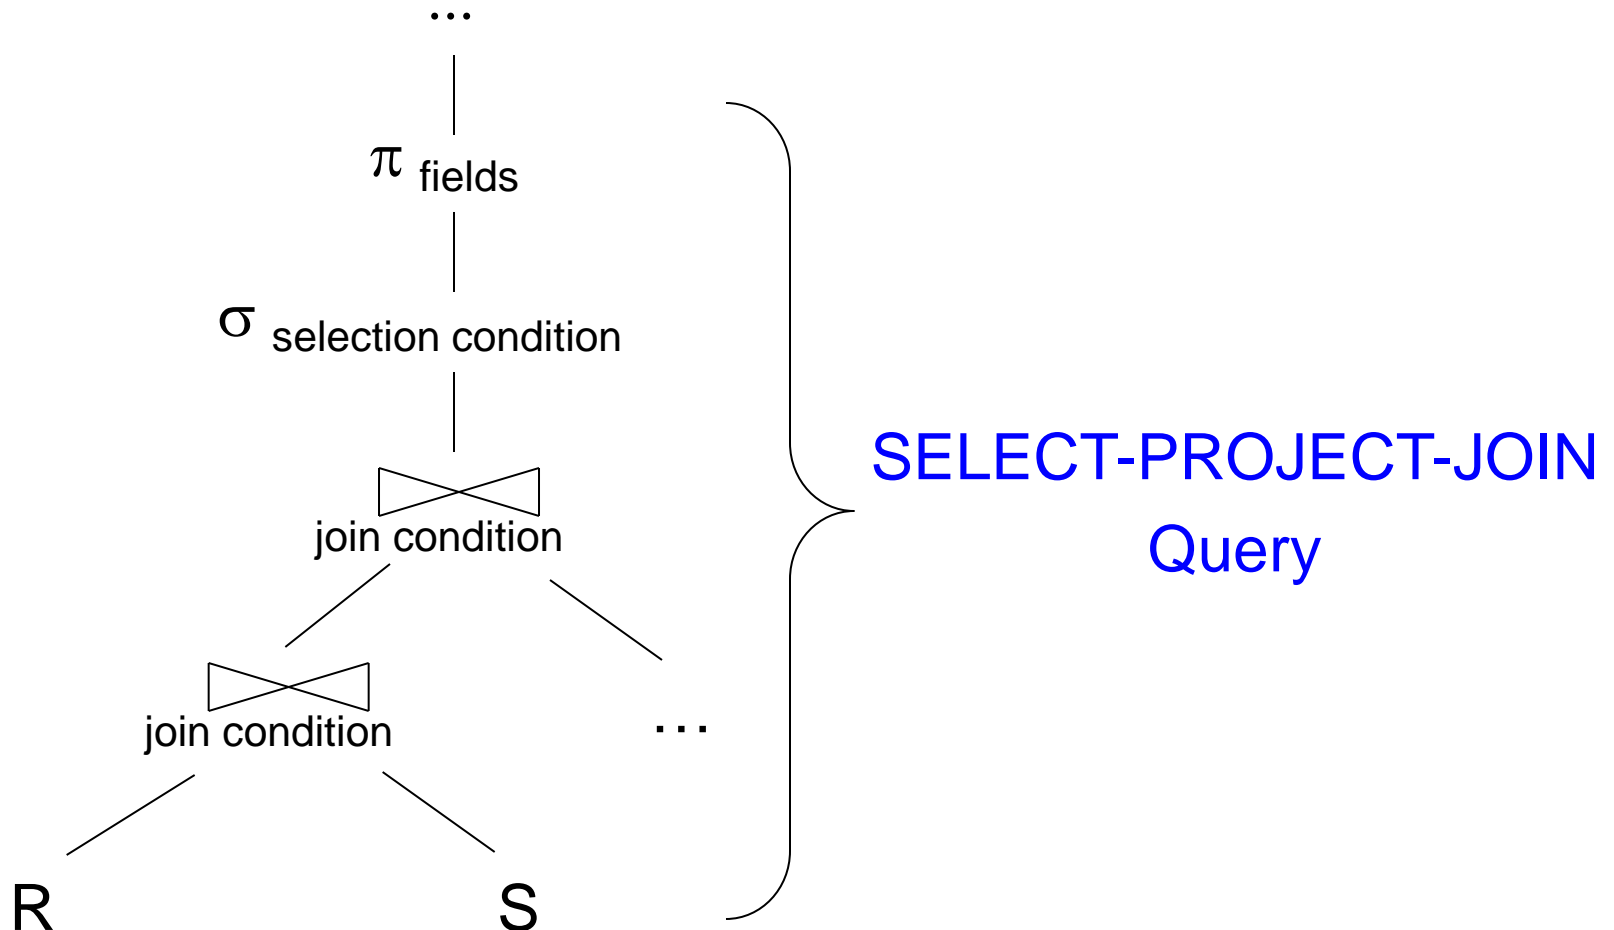
# Logical Query Plan

```
SELECT city, count(*)  
FROM sales  
GROUP BY city  
HAVING sum(price) > 100
```



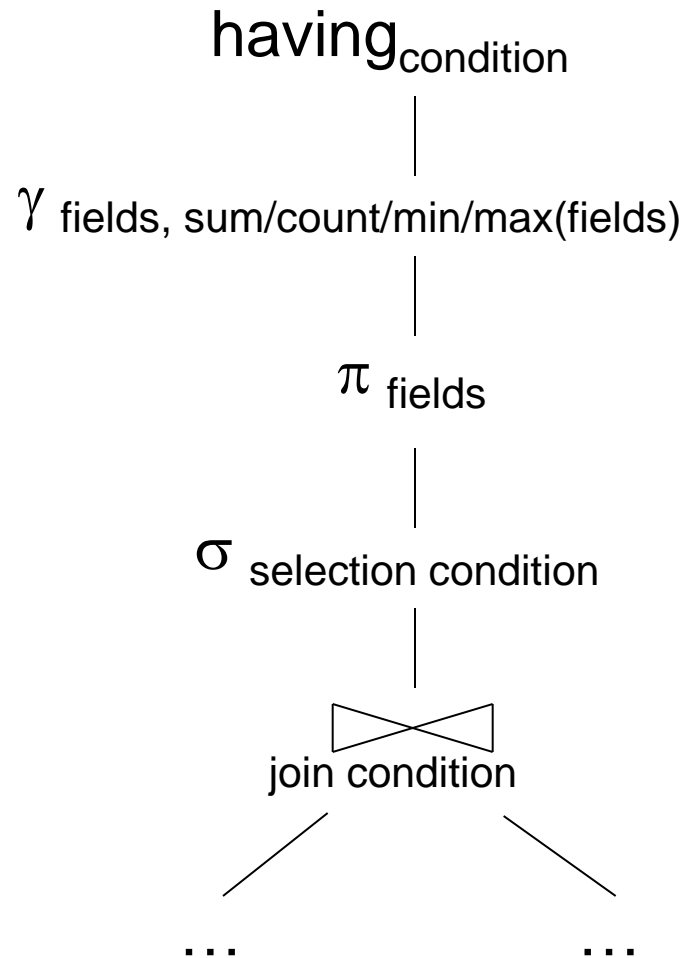
T1, T2, T3 = temporary tables

# Typical Plan for Block (1/2)





# Typical Plan For Block (2/2)



Supplier(sno,sname,scity,sstate)  
Part(pno,pname,psize,pcolor)  
Supply(sno,pno,price)

# How about Subqueries?

```
SELECT Q.sno
FROM Supplier Q
WHERE Q.sstate = 'WA'
      and not exists
      (SELECT *
       FROM Supply P
       WHERE P.sno = Q.sno
              and P.price > 100)
```

Supplier(sno,sname,scity,sstate)  
Part(pno,pname,psize,pcolor)  
Supply(sno,pno,price)

# How about Subqueries?

```
SELECT Q.sno  
FROM Supplier Q  
WHERE Q.sstate = 'WA'  
and not exists  
  (SELECT *  
   FROM Supply P  
   WHERE P.sno = Q.sno  
    and P.price > 100)
```

Correlation !

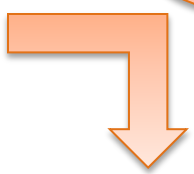


Supplier(sno,sname,scity,sstate)  
Part(pno,pname,psize,pcolor)  
Supply(sno,pno,price)

# How about Subqueries?

```
SELECT Q.sno
FROM Supplier Q
WHERE Q.sstate = 'WA'
      and not exists
      (SELECT *
       FROM Supply P
       WHERE P.sno = Q.sno
              and P.price > 100)
```

De-Correlation



```
SELECT Q.sno
FROM Supplier Q
WHERE Q.sstate = 'WA'
      and Q.sno not in
      (SELECT P.sno
       FROM Supply P
       WHERE P.price > 100)
```

Supplier(sno,sname,scity,sstate)  
Part(pno,pname,psize,pcolor)  
Supply(sno,pno,price)

# How about Subqueries?

Un-nesting

```
(SELECT Q.sno  
FROM Supplier Q  
WHERE Q.sstate = 'WA')  
EXCEPT  
(SELECT P.sno  
FROM Supply P  
WHERE P.price > 100)
```

EXCEPT = set difference

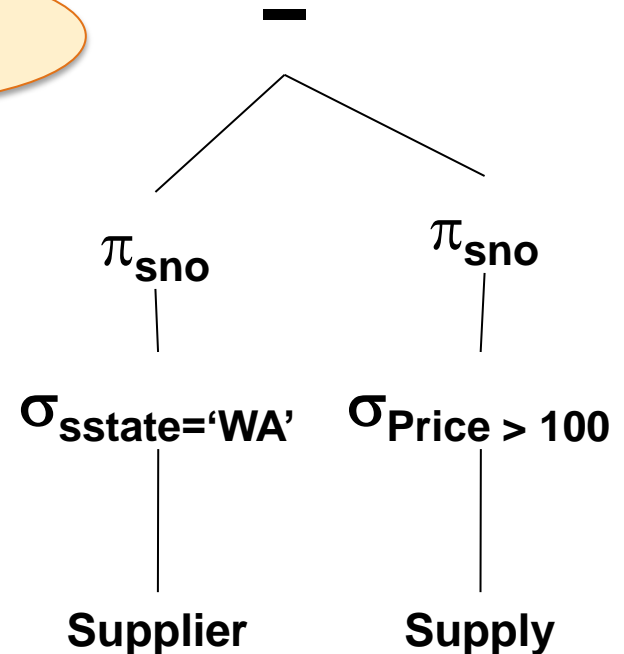
```
SELECT Q.sno  
FROM Supplier Q  
WHERE Q.sstate = 'WA'  
and Q.sno not in  
(SELECT P.sno  
FROM Supply P  
WHERE P.price > 100)
```

Supplier(sno,sname,scity,sstate)  
Part(pno,pname,psize,pcolor)  
Supply(sno,pno,price)

# How about Subqueries?

```
(SELECT Q.sno  
FROM Supplier Q  
WHERE Q.sstate = 'WA')  
EXCEPT  
(SELECT P.sno  
FROM Supply P  
WHERE P.price > 100)
```

Finally...



# From Logical Plans to Physical Plans

Supplier(sid, sname, scity, sstate)

Supply(sid, pno, quantity)

# Example

```
SELECT sname
FROM Supplier x, Supply y
WHERE x.sid = y.sid
      and y.pno = 2
      and x.scity = 'Seattle'
      and x.sstate = 'WA'
```

Give a relational algebra expression for this query



Supplier(sid, sname, scity, sstate)

Supply(sid, pno, quantity)

# Relational Algebra

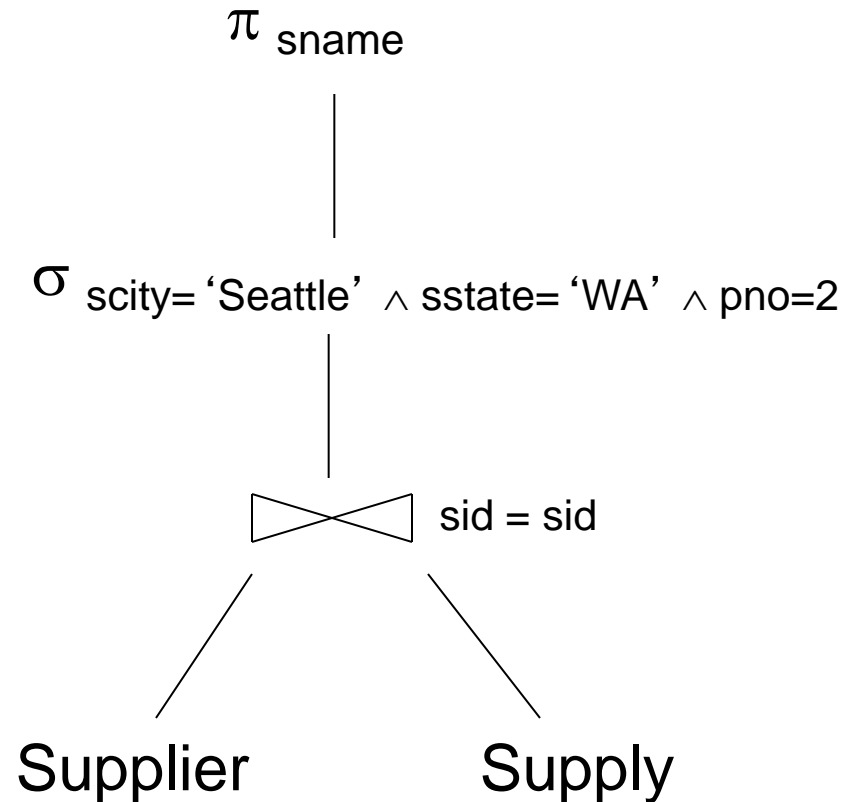
$$\pi_{\text{sname}}(\sigma_{\text{scity} = \text{'Seattle'} \wedge \text{sstate} = \text{'WA'} \wedge \text{pno} = 2} (\text{Supplier} \bowtie_{\text{sid} = \text{sid}} \text{Supply}))$$

Supplier(sid, sname, scity, sstate)

Supply(sid, pno, quantity)

# Relational Algebra

Relational algebra expression is also called the “logical query plan”

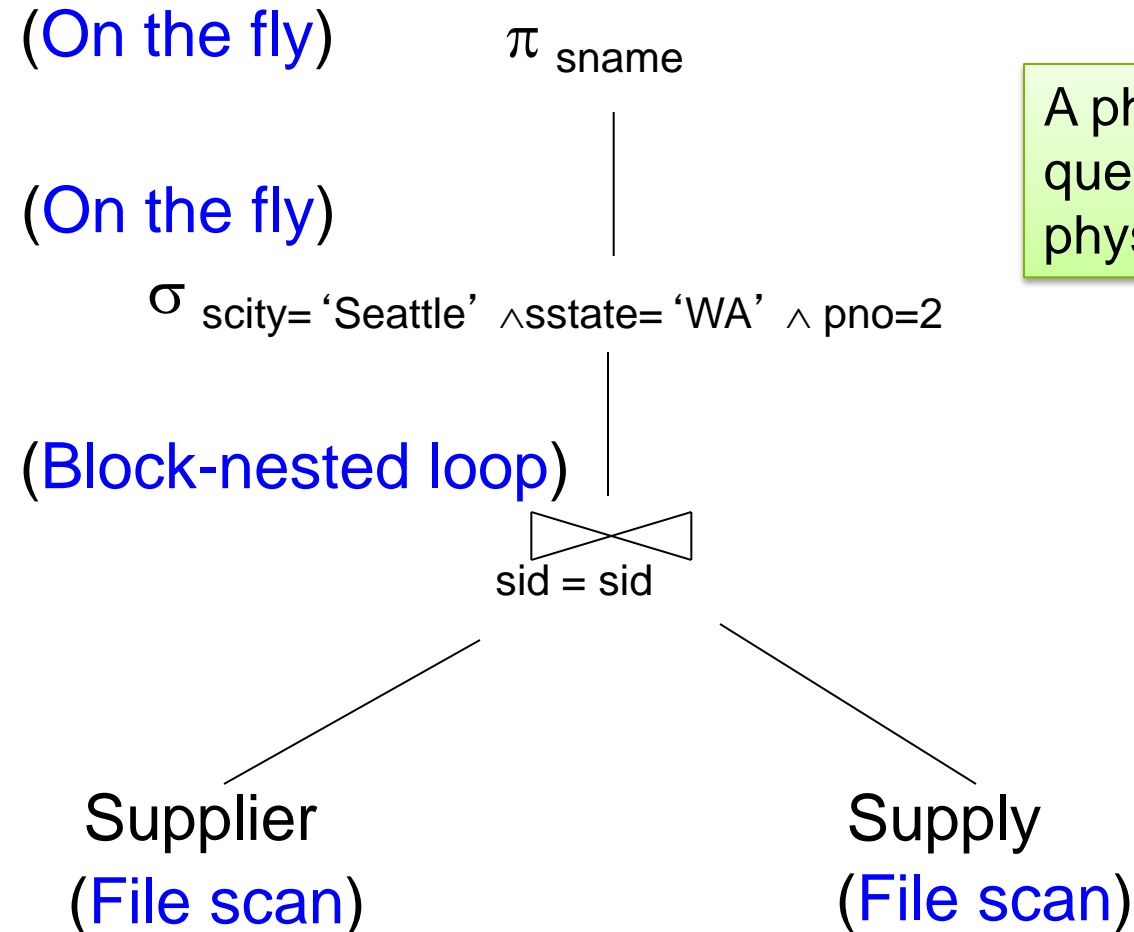


Supplier(sid, sname, scity, sstate)

Supply(sid, pno, quantity)

# Physical Query Plan 1

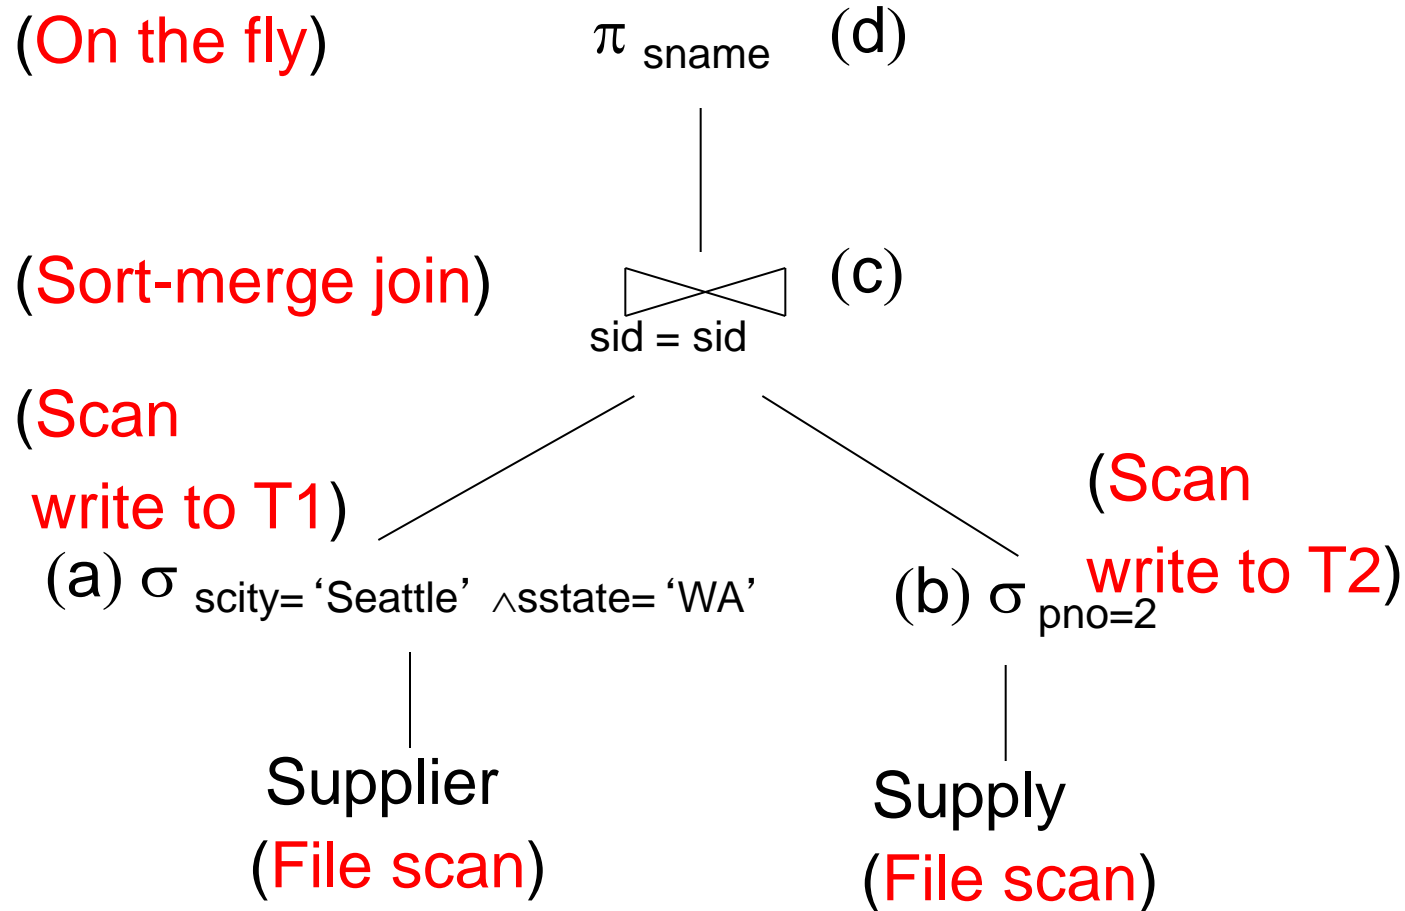
A physical query plan is a logical query plan annotated with physical implementation details



Supplier(sid, sname, scity, sstate)

Supply(sid, pno, quantity)

# Physical Query Plan 2



Supplier(sid, sname, scity, sstate)

Supply(sid, pno, quantity)

# Physical Query Plan 3

(On the fly) (d)  $\pi_{\text{sname}}$

(On the fly)

(c)  $\sigma_{\text{scity}='Seattle' \wedge \text{sstate}='WA'}$

(b)  $\text{sid} = \text{sid}$  (Index nested loop)

(a)  $\sigma_{\text{pno}=2}$

Supply

Supplier

(Index lookup on pno) (Index lookup on sid)

Assume: clustered

CSE 344 - Winter 2014

Doesn't matter if clustered or not

# Physical Data Independence

- Means that applications are insulated from changes in physical storage details
  - E.g., can add/remove indexes without changing apps
  - Can do other physical tunings for performance
- SQL and relational algebra facilitate physical data independence because both languages are “set-at-a-time”: Relations as input and output