

Introduction to Data Management

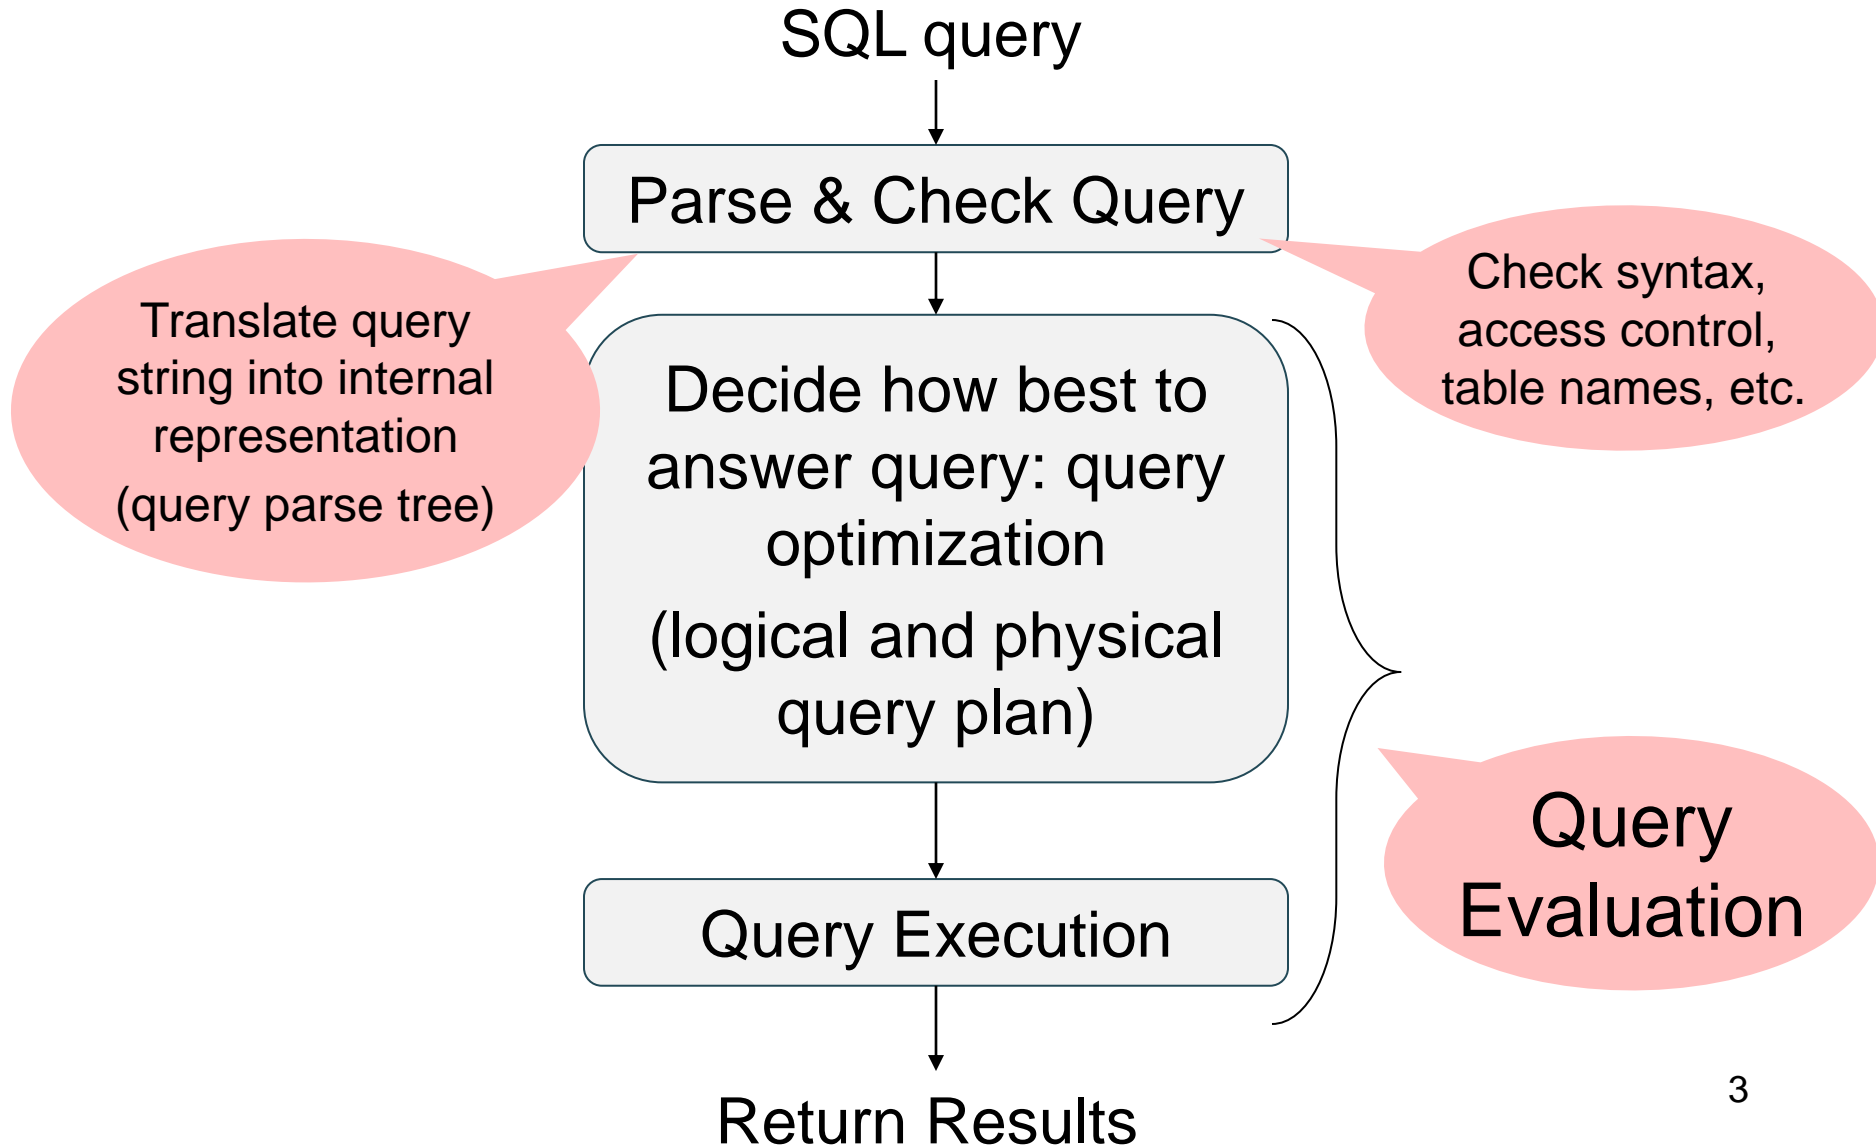
CSE 344

Lecture 10: Relational Algebra Wrap-up Systems Architecture

Announcements

- WQ4 is due next Tuesday
- HW3 is due next Thursday
- Today's lecture: How DBMSs work
 - Relational algebra and query execution
 - 2.4, 5.1-5.2, 16.2-16.3
 - (Optional) Chapter 15, more in CSE 444
 - Client-server architecture
 - 9.1

Query Evaluation Steps



The WHAT and the HOW

- SQL = **WHAT** we want to get from the data
- Relational Algebra = **HOW** to get the data we want
- The passage from **WHAT** to **HOW** is called **query optimization**

Overview: SQL = WHAT

Product(pid, name, price)

Purchase(pid, cid, store)

Customer(cid, name, city)

```
SELECT DISTINCT x.name, z.name
FROM Product x, Purchase y, Customer z
WHERE x.pid = y.pid and y.cid = y.cid and
      x.price > 100 and z.city = 'Seattle'
```

It's clear WHAT we want, unclear HOW to get it

Query Optimizer = HOW

Product(pid, name, price)

Purchase(pid, cid, store)

Customer(cid, name, city)

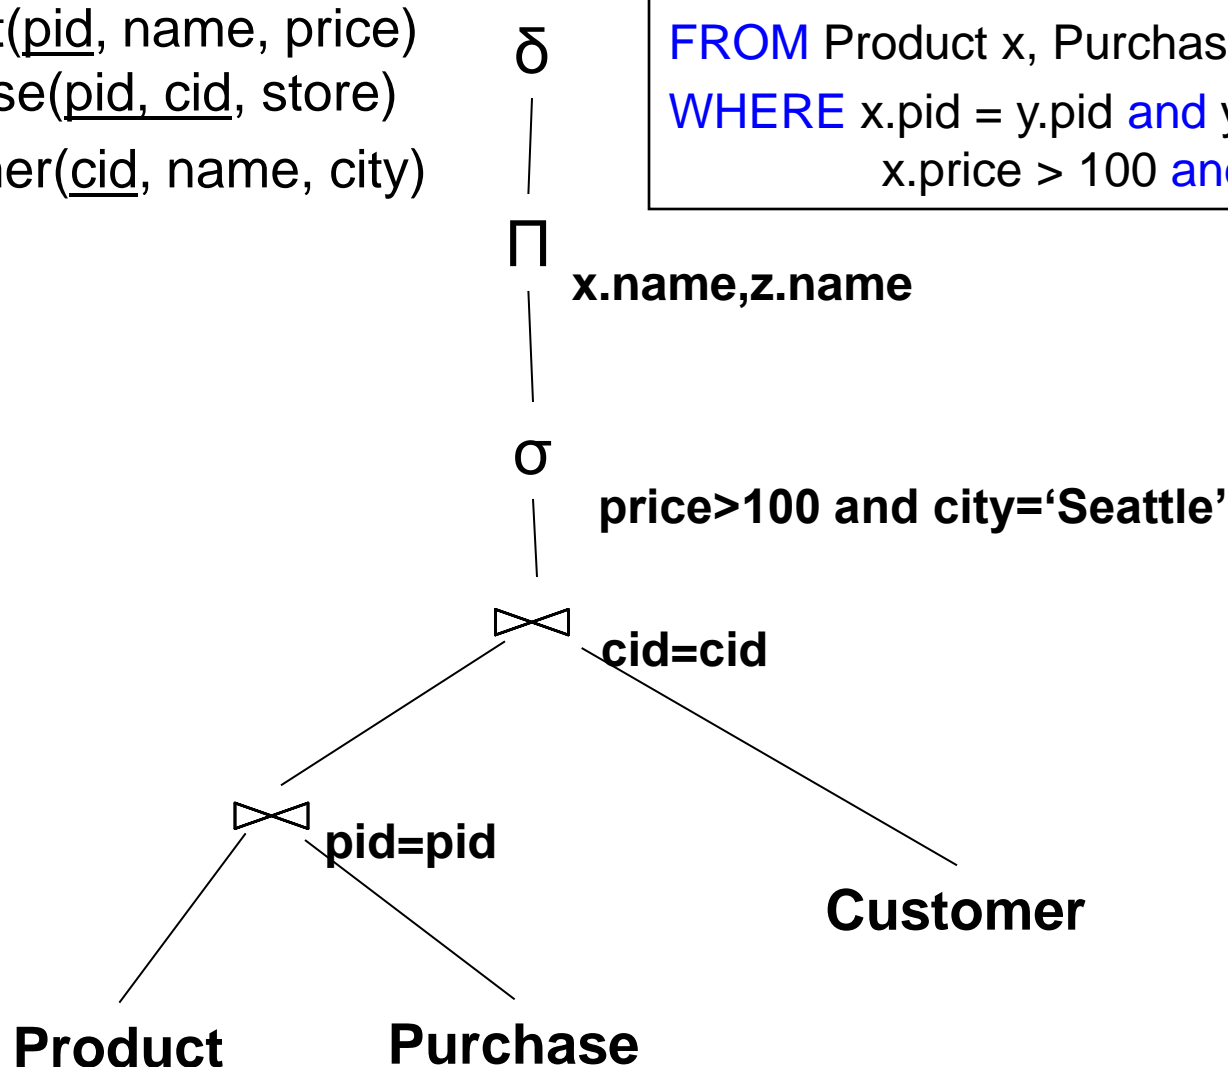
```
SELECT DISTINCT x.name, z.name  
FROM Product x, Purchase y, Customer z  
WHERE x.pid = y.pid and y.cid = z.cid and  
       x.price > 100 and z.city = 'Seattle'
```

1. Which (equivalent) logical plan is the most efficient?
2. Physical plan:
 - How to implement each operation in the plan?
 - How to pass data from one operation to the other?
(pipeline/on-the-fly, main-memory buffer, disk)

From SQL to RA

Product(pid, name, price)
Purchase(pid, cid, store)
Customer(cid, name, city)

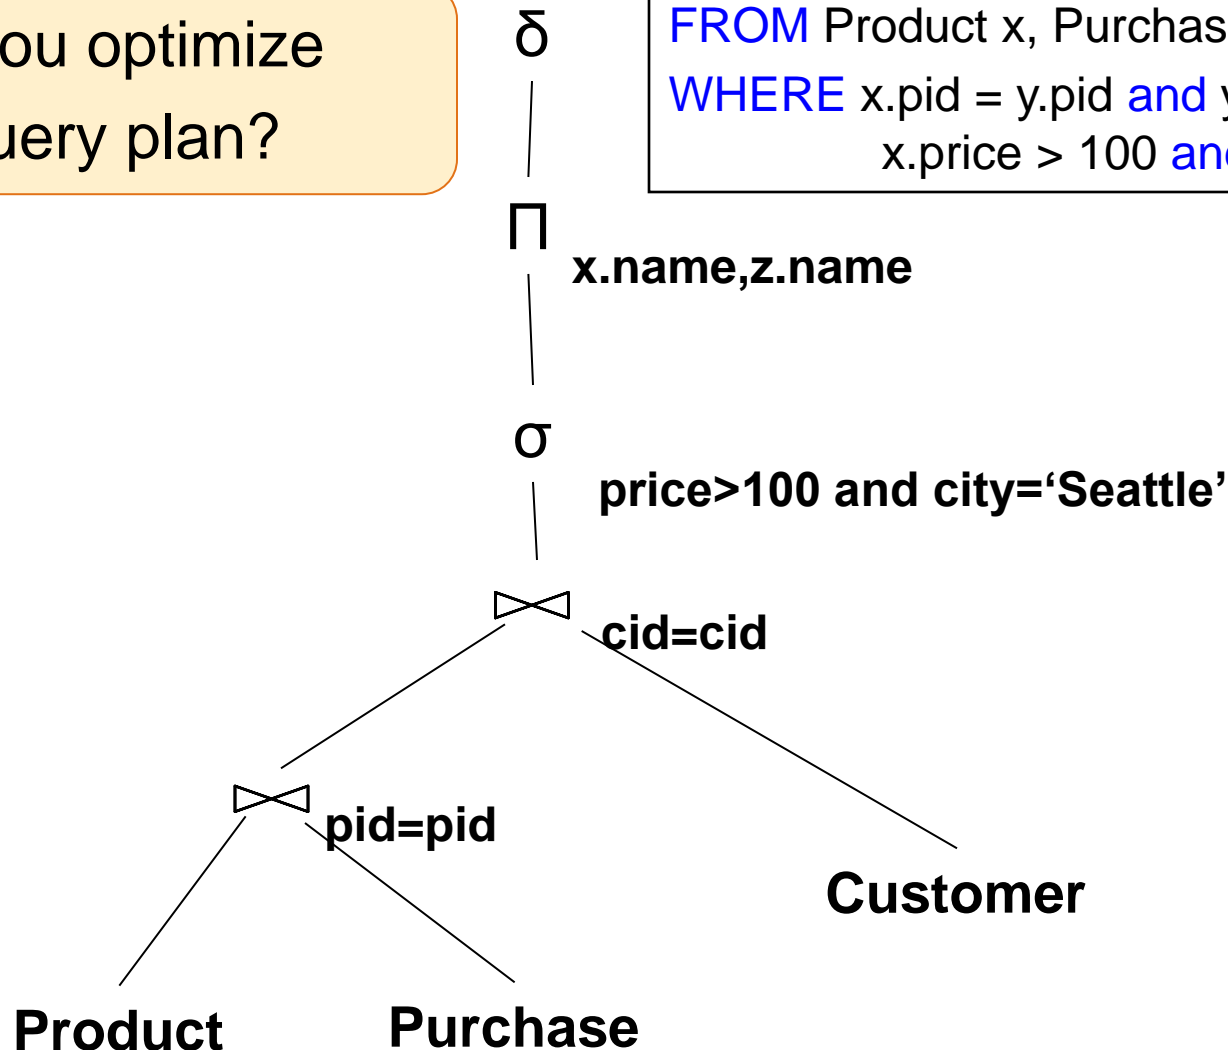
```
SELECT DISTINCT x.name, z.name
FROM Product x, Purchase y, Customer z
WHERE x.pid = y.pid and y.cid = z.cid and
      x.price > 100 and z.city = 'Seattle'
```



From SQL to RA

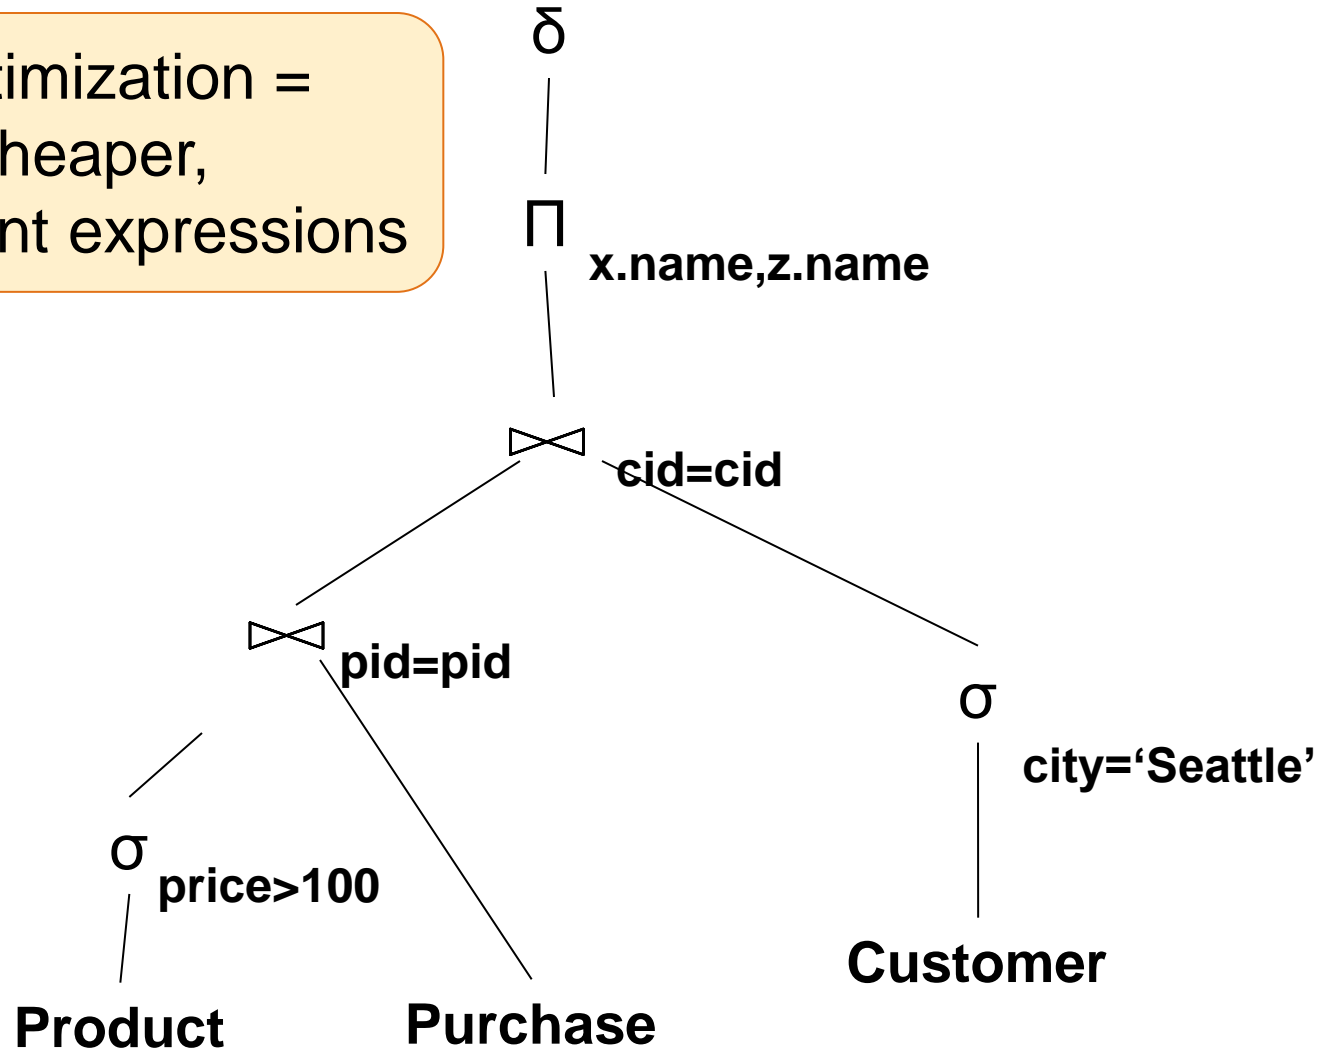
Can you optimize
this query plan?

```
SELECT DISTINCT x.name, z.name
FROM Product x, Purchase y, Customer z
WHERE x.pid = y.pid and y.cid = z.cid and
      x.price > 100 and z.city = 'Seattle'
```



An Equivalent Expression

Query optimization =
finding cheaper,
equivalent expressions

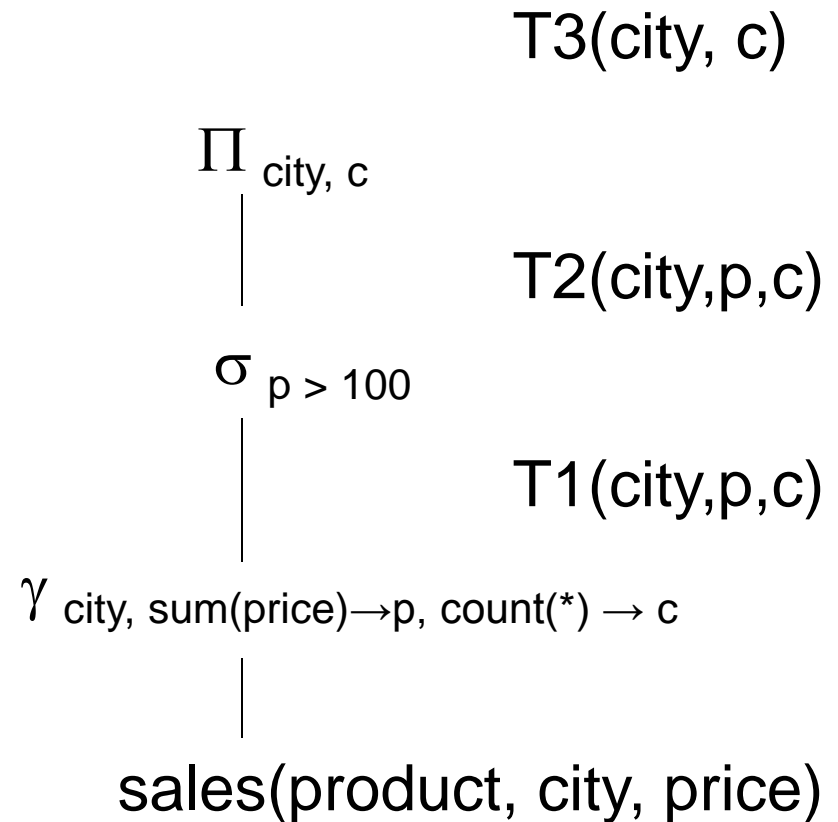


Extended RA: Operators on Bags

- Duplicate elimination δ
- Grouping γ
- Sorting τ

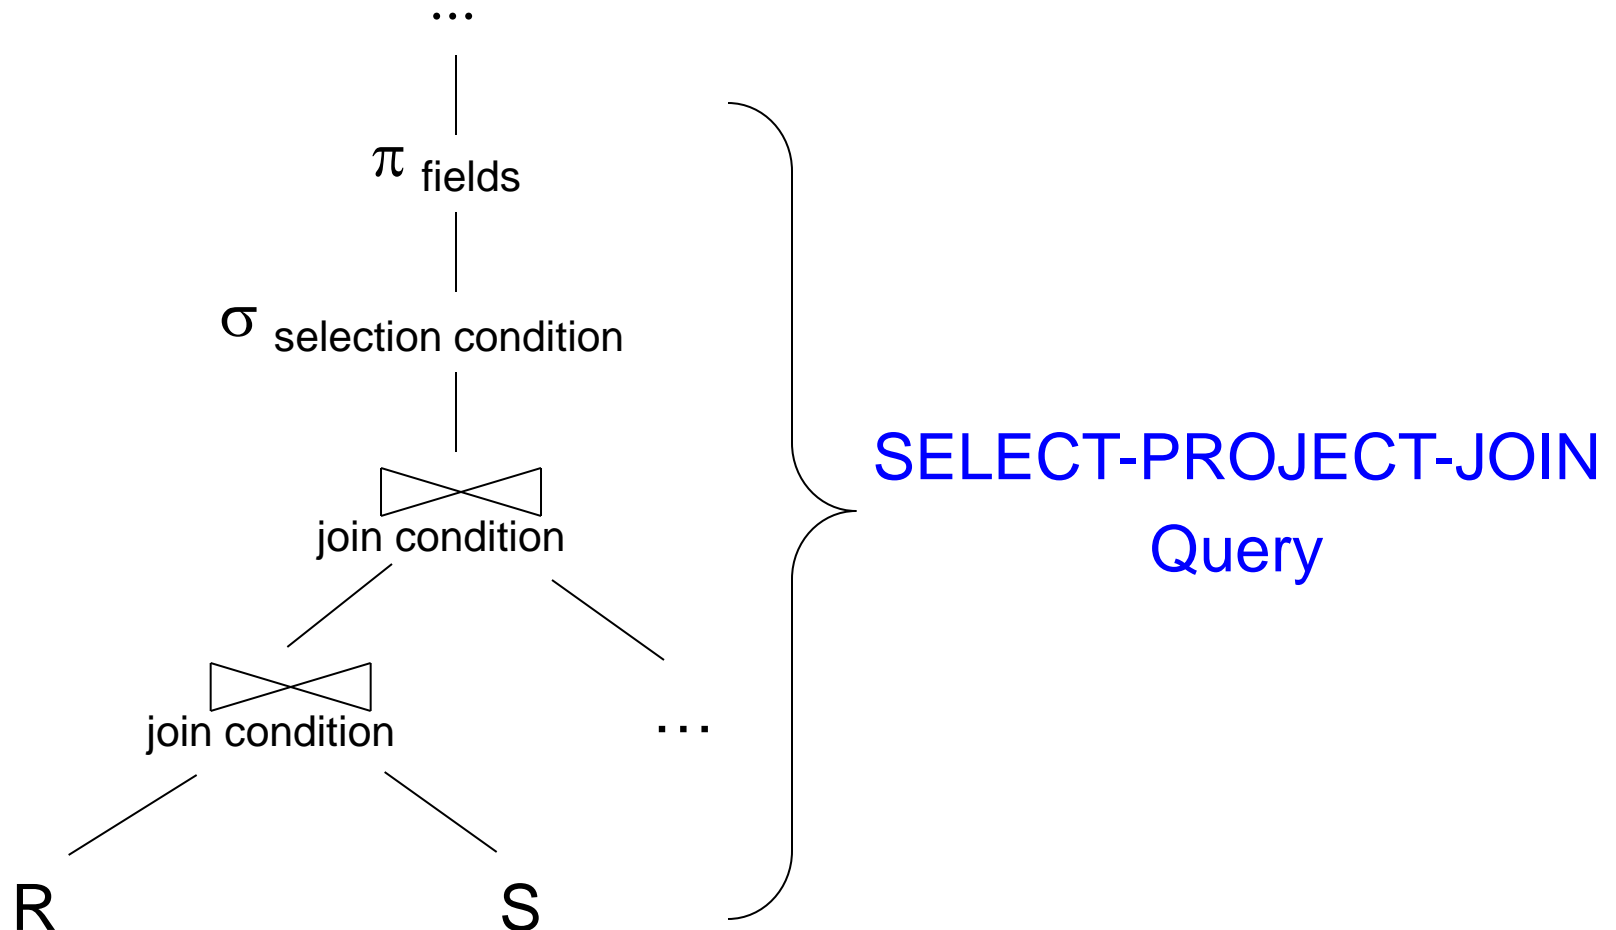
Logical Query Plan

```
SELECT city, count(*)  
FROM sales  
GROUP BY city  
HAVING sum(price) > 100
```

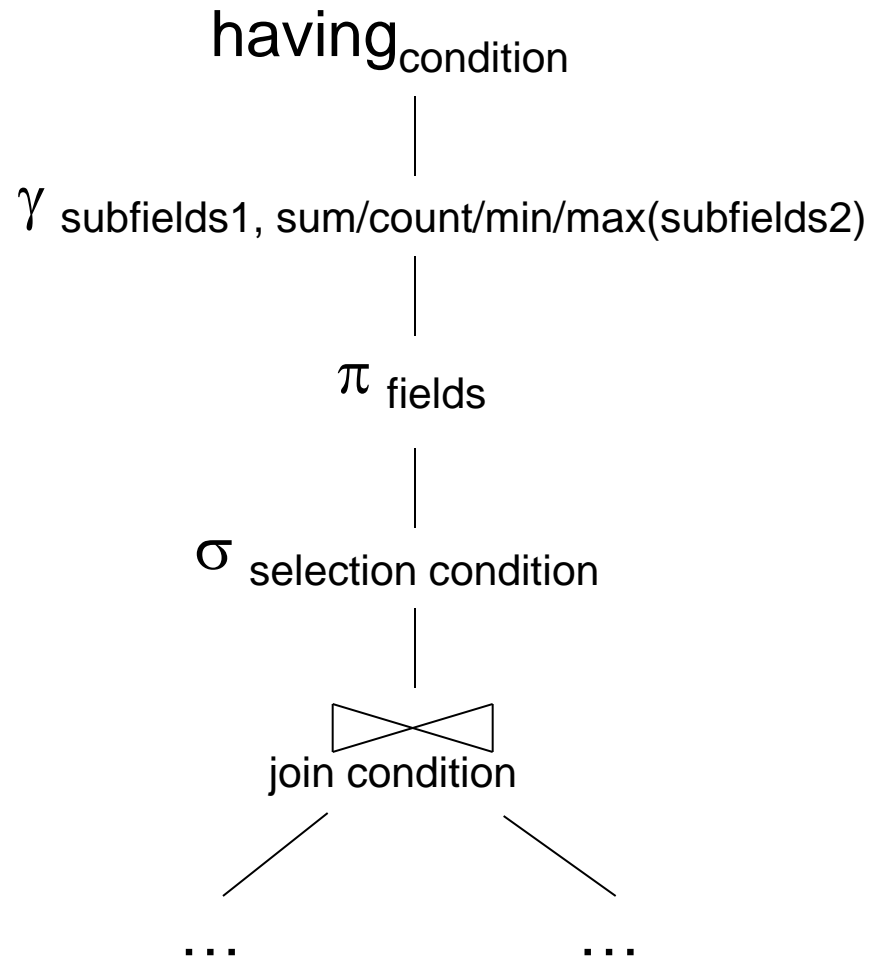


T1, T2, T3 = temporary tables

Typical Plan for Block (1/2)



Typical Plan For Block (2/2)



Supplier(sno,sname,scity,sstate)
Part(pno,pname,psize,pcolor)
Supply(sno,pno,price)

How about Subqueries?

```
SELECT Q.sno
FROM Supplier Q
WHERE Q.sstate = 'WA'
      and not exists
      (SELECT *
       FROM Supply P
       WHERE P.sno = Q.sno
              and P.price > 100)
```

Sno of Suppliers
from WA
who did not
Supply a part
with price > 100

Supplier(sno,sname,scity,sstate)
Part(pno,pname,psize,pcolor)
Supply(sno,pno,price)

How about Subqueries?

```
SELECT Q.sno
FROM Supplier Q
WHERE Q.sstate = 'WA'
and not exists
(SELECT *
FROM Supply P
WHERE P.sno = Q.sno
and P.price > 100)
```


Correlation !

Sno of Suppliers
from WA
who did not
Supply a part
with price > 100

Supplier(sno,sname,scity,sstate)
Part(pno,pname,psize,pcolor)
Supply(sno,pno,price)

Let's try to De-Correlate! (soln-1)

```
SELECT Q.sno
FROM Supplier Q
WHERE Q.sstate = 'WA'
      and not exists
      (SELECT *
       FROM Supply P
       WHERE P.sno = Q.sno
              and P.price > 100)
```



De-Correlation


(Solution from class)

```
SELECT Q.sno
FROM Supplier Q
WHERE Q.sstate = 'WA'
      and not in (SELECT P.sno
                  FROM Supply P
                  WHERE P.sno = Q.sno
                         and P.price > 100)
```


Supplier(sno,sname,scity,sstate)
Part(pno,pname,psize,pcolor)
Supply(sno,pno,price)

Let's try to De-Correlate! (soln-2)

```
SELECT Q.sno
FROM Supplier Q
WHERE Q.sstate = 'WA'
      and not exists
      (SELECT *
       FROM Supply P
       WHERE P.sno = Q.sno
              and P.price > 100)
```



De-Correlation

(Solution from class)

```
SELECT Q.sno
FROM Supplier Q, Supply P
WHERE Q.sstate = 'WA'
      And Q.sno = P.sno
GROUP BY Q.sno
HAVING MAX(P.price) <= 100
```

How about Subqueries? (Decorrelation)

Supplier(sno,sname,scity,sstate)
Part(pno,pname,psize,pcolor)
Supply(sno,pno,price)

```
SELECT Q.sno
FROM Supplier Q
WHERE Q.sstate = 'WA'
      and not exists
      (SELECT *
       FROM Supply P
       WHERE P.sno = Q.sno
              and P.price > 100)
```

How to model “not in” (nested)
Using RA operators?



De-Correlation

```
SELECT Q.sno
FROM Supplier Q
WHERE Q.sstate = 'WA'
      and Q.sno not in
      (SELECT P.sno
       FROM Supply P
       WHERE P.price > 100)
```

How about Subqueries? (Un-nesting)

Supplier(sno,sname,scity,sstate)
Part(pno,pname,psize,pcolor)
Supply(sno,pno,price)

Un-nesting

```
(SELECT Q.sno
FROM Supplier Q
WHERE Q.sstate = 'WA')
EXCEPT
(SELECT P.sno
FROM Supply P
WHERE P.price > 100)
```

EXCEPT = set difference


```
SELECT Q.sno
FROM Supplier Q
WHERE Q.sstate = 'WA'
and Q.sno not in
(SELECT P.sno
FROM Supply P
WHERE P.price > 100)
```

Supplier(sno,sname,scity,sstate)
Part(pno,pname,psize,pcolor)
Supply(sno,pno,price)

Conversion to RA, finally!

```
(SELECT Q.sno  
FROM Supplier Q  
WHERE Q.sstate = 'WA')  
EXCEPT  
(SELECT P.sno  
FROM Supply P  
WHERE P.price > 100)
```

Finally...



From Logical Plans to Physical Plans

Supplier(sid, sname, scity, sstate)

Supply(sid, pno, quantity)

Example

```
SELECT sname
FROM Supplier x, Supply y
WHERE x.sid = y.sid
      and y.pno = 2
      and x.scity = 'Seattle'
      and x.sstate = 'WA'
```

Give a relational algebra expression for this query

Supply(sid, pno, quantity)

Relational Algebra

```
SELECT sname
FROM Supplier x, Supply y
WHERE x.sid = y.sid
      and y.pno = 2
      and x.scity = 'Seattle'
      and x.sstate = 'WA'
```

Give a relational algebra expression for this query

$$\pi_{\text{-----}}(\sigma_{\text{-----}}(\text{Supplier} \bowtie \text{Supply}))$$

Supplier(sid, sname, scity, sstate)

Supply(sid, pno, quantity)

Relational Algebra

```
SELECT sname
FROM Supplier x, Supply y
WHERE x.sid = y.sid
      and y.pno = 2
      and x.scity = 'Seattle'
      and x.sstate = 'WA'
```

Give a relational algebra expression for this query


$$\pi_{\text{sname}}(\sigma_{\text{scity}='Seattle' \wedge \text{sstate}='WA' \wedge \text{pno}=2}(\text{Supplier} \bowtie_{\text{sid}=\text{sid}} \text{Supply}))$$

Supplier(sid, sname, scity, sstate)

Supply(sid, pno, quantity)

Relational Algebra

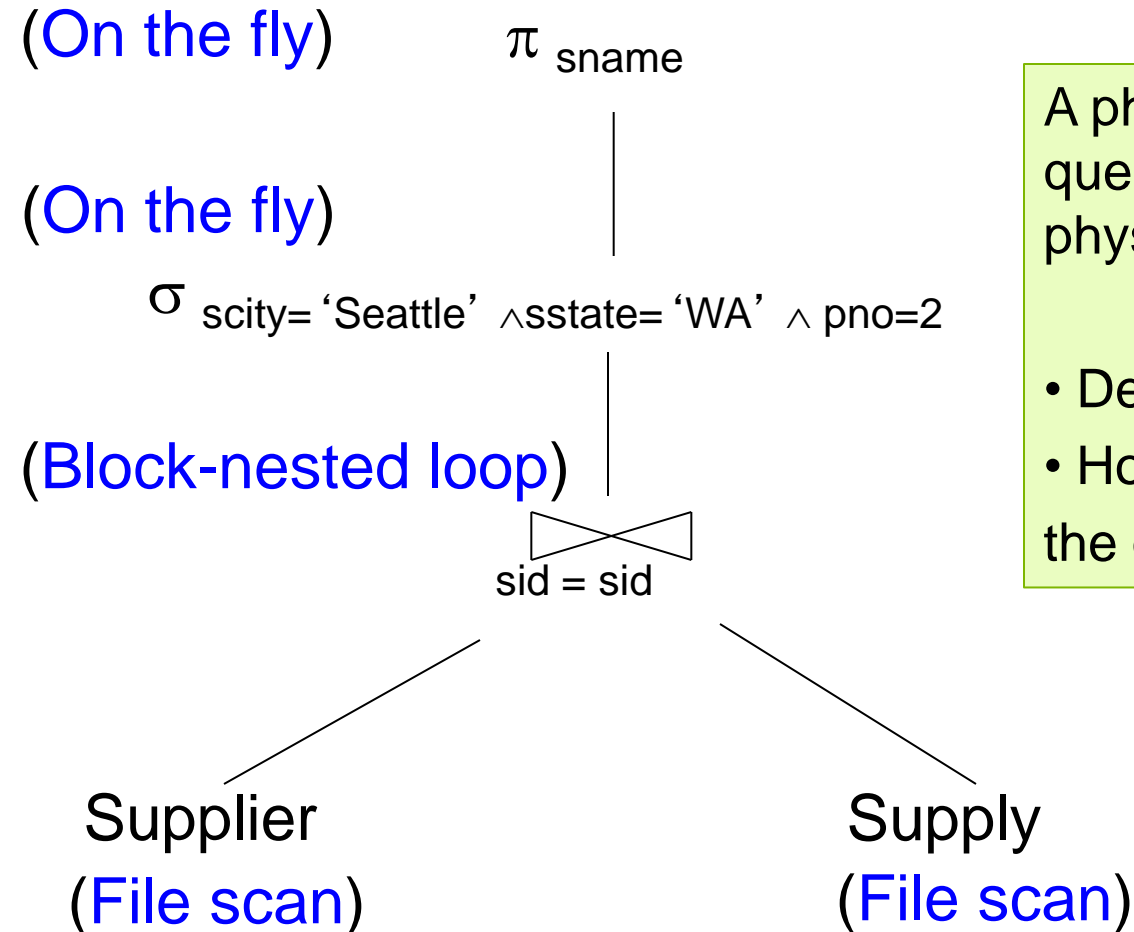
Relational algebra expression is also called the “logical query plan”



Supplier(sid, sname, scity, sstate)

Supply(sid, pno, quantity)

Physical Query Plan 1



A physical query plan is a logical query plan annotated with physical implementation details

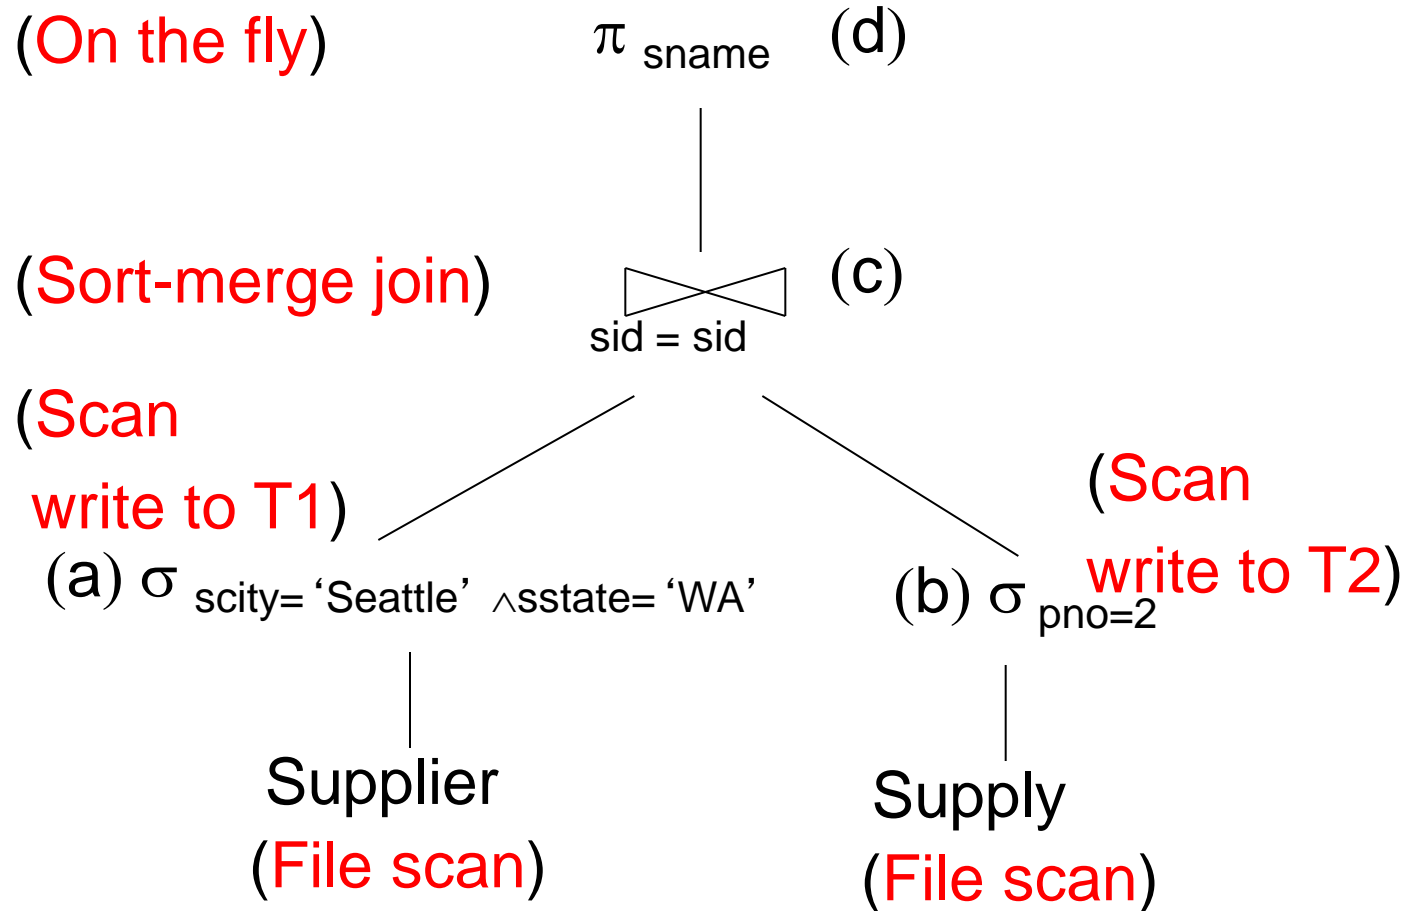
- Details of each operator
- How info is passed between the operators

Can you think of any
Other types of
Join algo? 😊

Supplier(sid, sname, scity, sstate)

Supply(sid, pno, quantity)

Physical Query Plan 2



Supplier(sid, sname, scity, sstate)

Supply(sid, pno, quantity)

Physical Query Plan 3

(On the fly) (d) π_{sname}

(On the fly)

(c) $\sigma_{\text{scity}='Seattle' \wedge \text{sstate}='WA'}$

(b) $\text{sid} = \text{sid}$ (Index nested loop)

(a) $\sigma_{\text{pno}=2}$

Supply


Supplier

(Index lookup on pno) (Index lookup on sid)

Assume: clustered

Doesn't matter if clustered or not

Why?



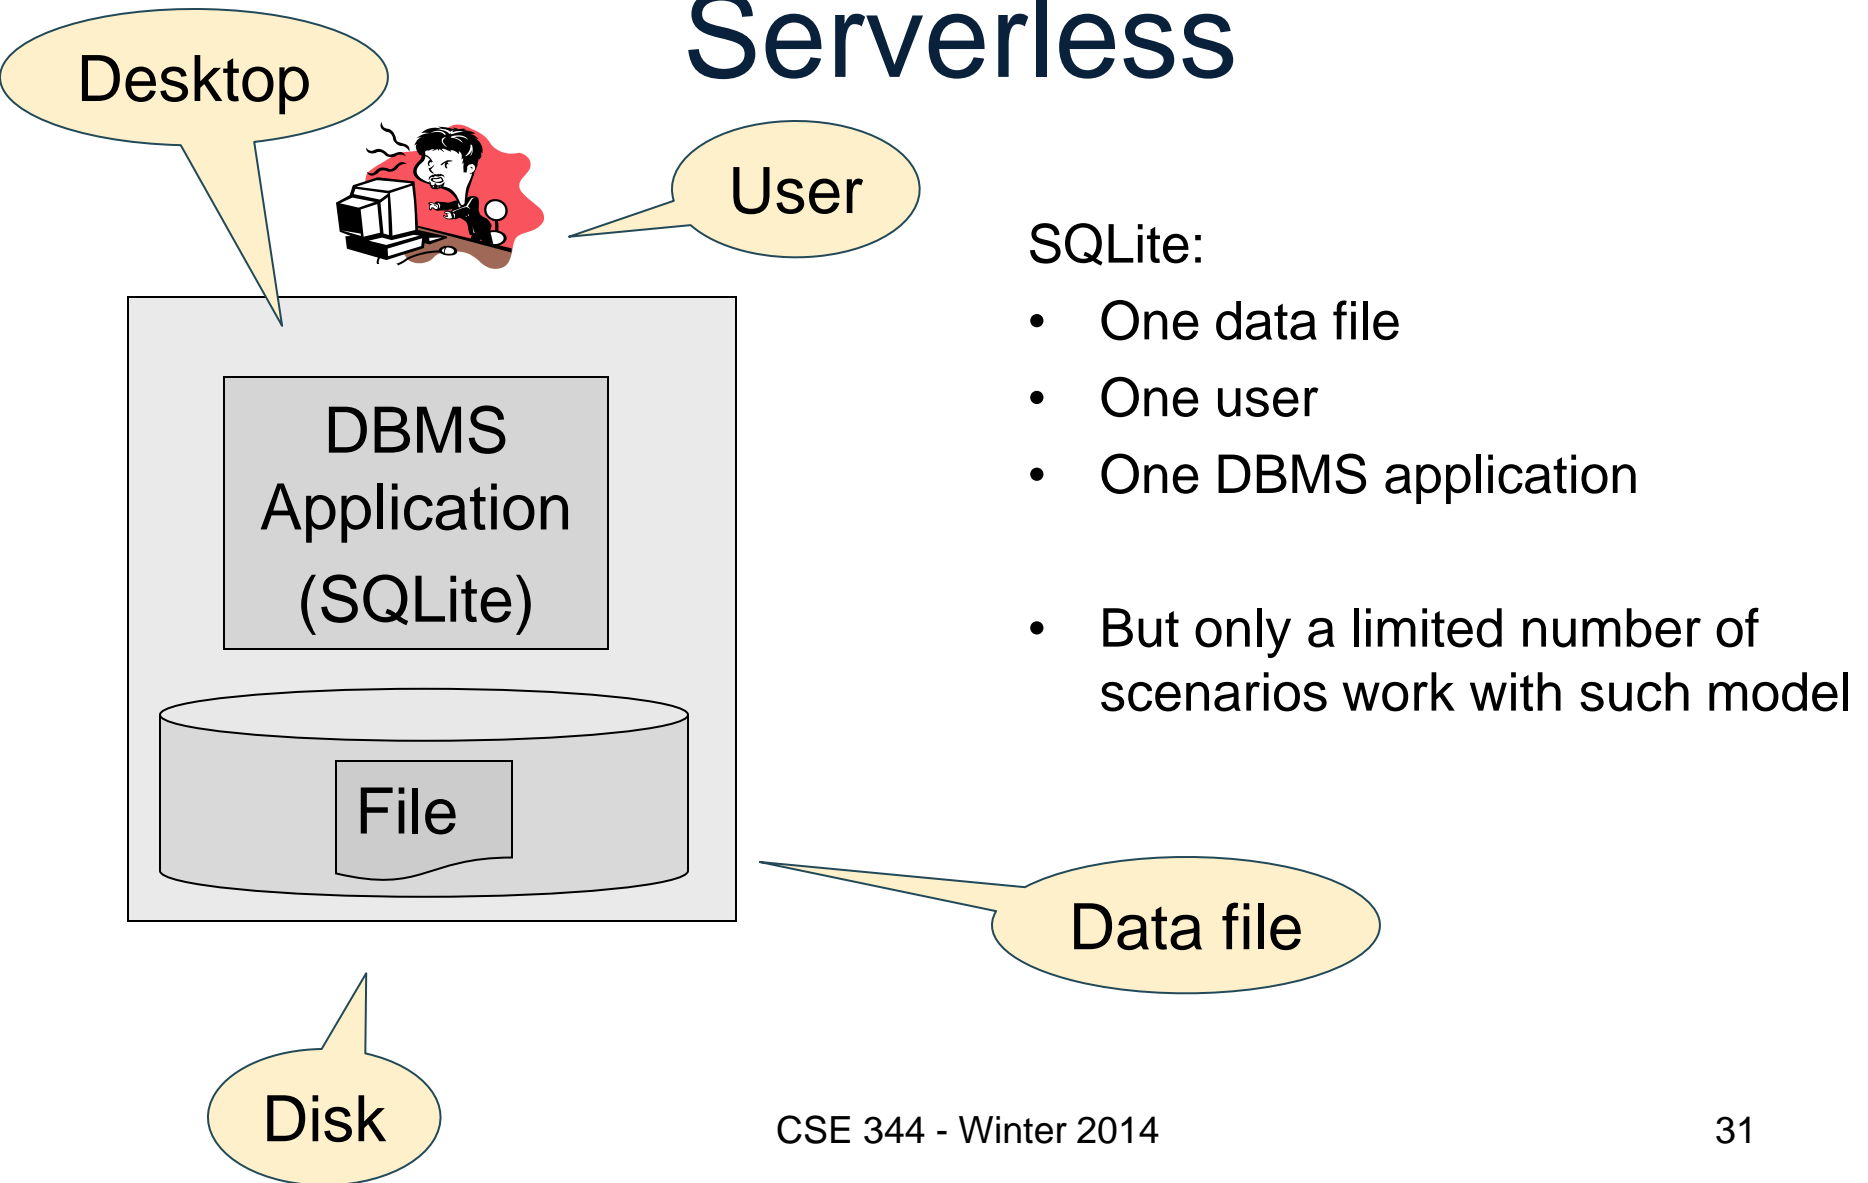
Physical Data Independence

- Means that applications are insulated from changes in physical storage details
 - E.g., can add/remove indexes without changing apps
 - Can do other physical tunings for performance
- SQL and relational algebra facilitate physical data independence because both languages are “set-at-a-time”: Relations as input and output

Architectures

1. Serverless
2. Two tier: client/server
3. Three tier: client/app-server/db-server

Serverless




Client-Server

Supports many apps and many users simultaneously

Client Applications

Server Machine



Connection (JDBC, ODBC)




- One server running the database
- Many clients, connecting via the ODBC or JDBC (Java Database Connectivity) protocol

Client-Server


- One *server* that runs the DBMS (or RDBMS):
 - Your own desktop, or
 - Some beefy system, or
 - A cloud service (SQL Azure)
- Many *clients* run apps and connect to DBMS
 - Microsoft's Management Studio (for SQL Server), or
 - psql (for postgres)
 - Some Java program (HW5) or some C++ program
- Clients “talk” to server using JDBC/ODBC protocol

3-Tiers DBMS Deployment

Web-based applications



Connection
(e.g., JDBC)



Browser



HTTP/SSL



DBS: 6. Executes queries

AppS: Runs “business logic”

e.g. converts price from USD to EUR

5. Forms and asks queries to DBS,

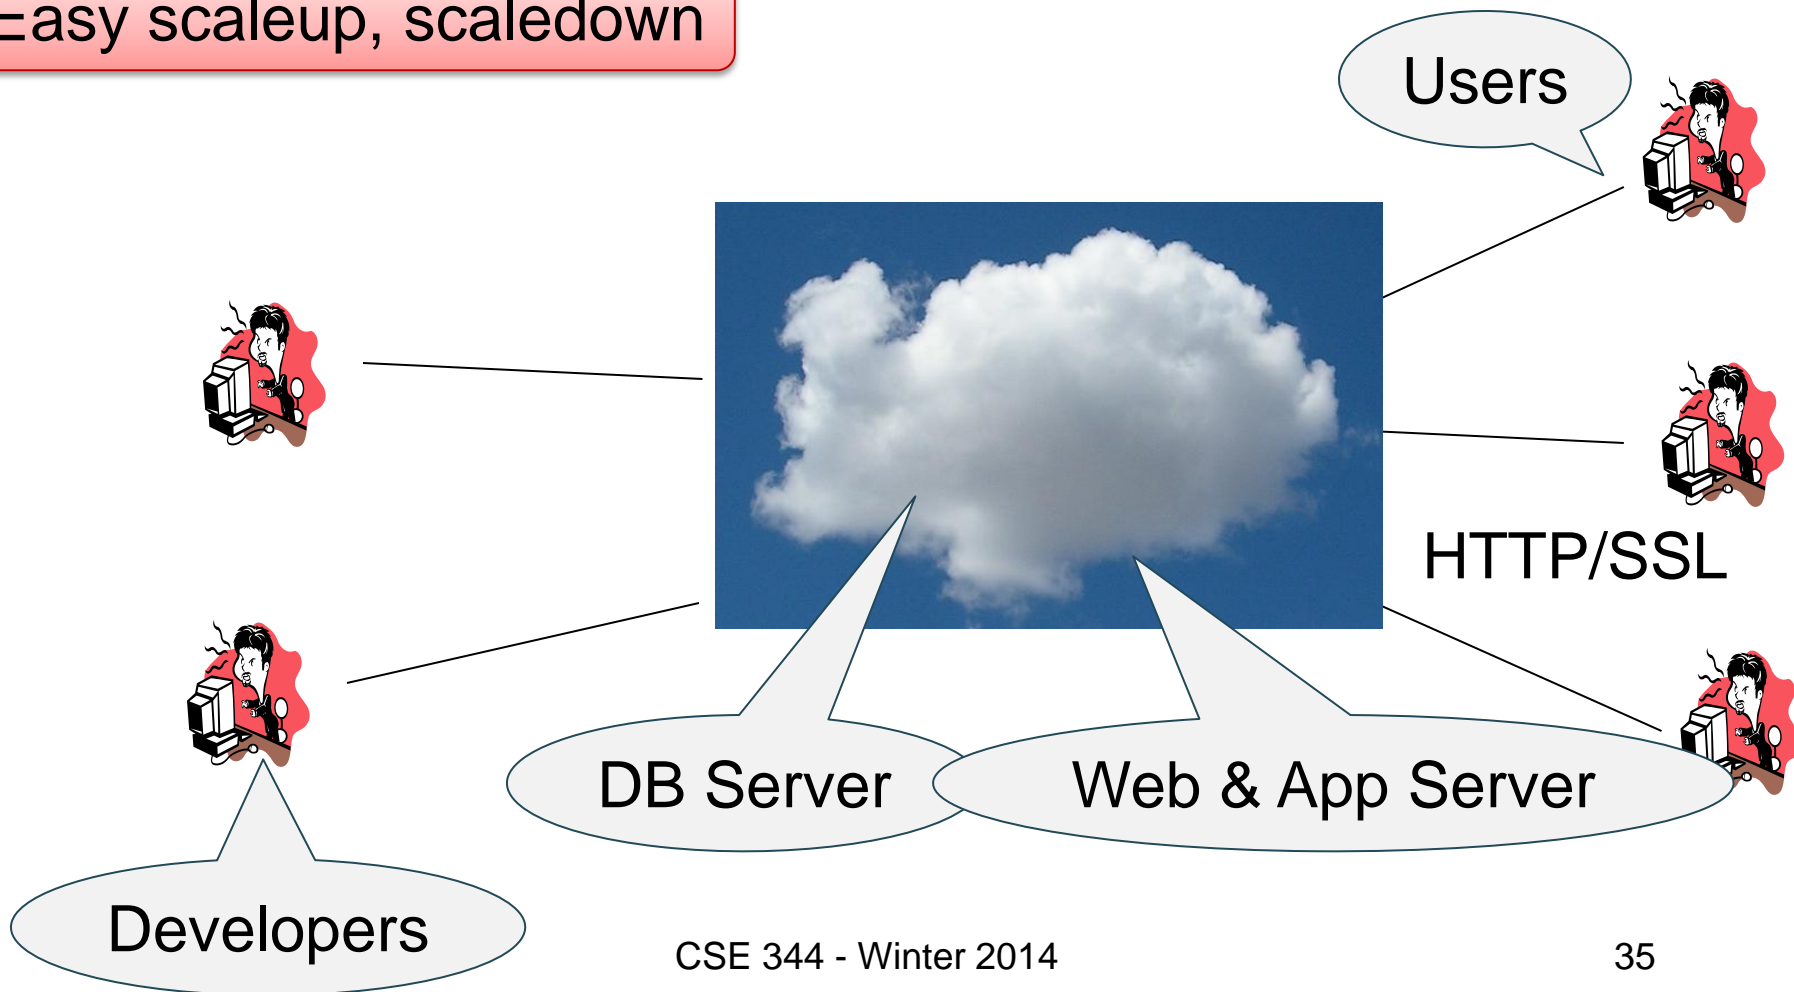
7. Returns results to WS

WS: 2. Provides the page
WS: 4. Sends info to AppS
8. Returns answers from AppS to users' browsers

1. Open amazon.com
3. Search for books
9. Gets the results 😊

DBMS Deployment: Cloud

Easy scaleup, scaledown



Using a DBMS Server

1. Client application establishes connection to server
2. Client must authenticate self
3. Client submits SQL commands to server
4. Server executes commands and returns results

