

Mesos for Spark Users


mesos.apache.org

[@ApacheMesos](https://twitter.com/ApacheMesos)

Benjamin Hindman – [@benh](https://twitter.com/benh)




agenda

- ① Mesos 
- ② Spark on Mesos
- ③ why Mesos?
 - ① multi-tenancy
 - ② fine-grained sharing
 - ③ why not?
- ④ long-lived services and other frameworks


a little history

Mesos started as a research project at Berkeley in early 2009 by Benjamin Hindman, Andy Konwinski, Matei Zaharia, Ali Ghodsi, Anthony D. Joseph, Randy Katz, Scott Shenker, Ion Stoica


our motivation



our intuition



static partitioning considered harmful



datacenter

static partitioning considered harmful

hadoop




Spark 





static partitioning considered harmful


hadoop




Spark




static partitioning considered harmful




static partitioning considered harmful



static partitioning considered harmful




our intuition




***"Map/Reduce is a big hammer,
but not everything is a nail!"***



anatomy of Spark



anatomy of a framework




framework


\approx

distributed system

anatomy of a framework



anatomy of a framework




execution coordination

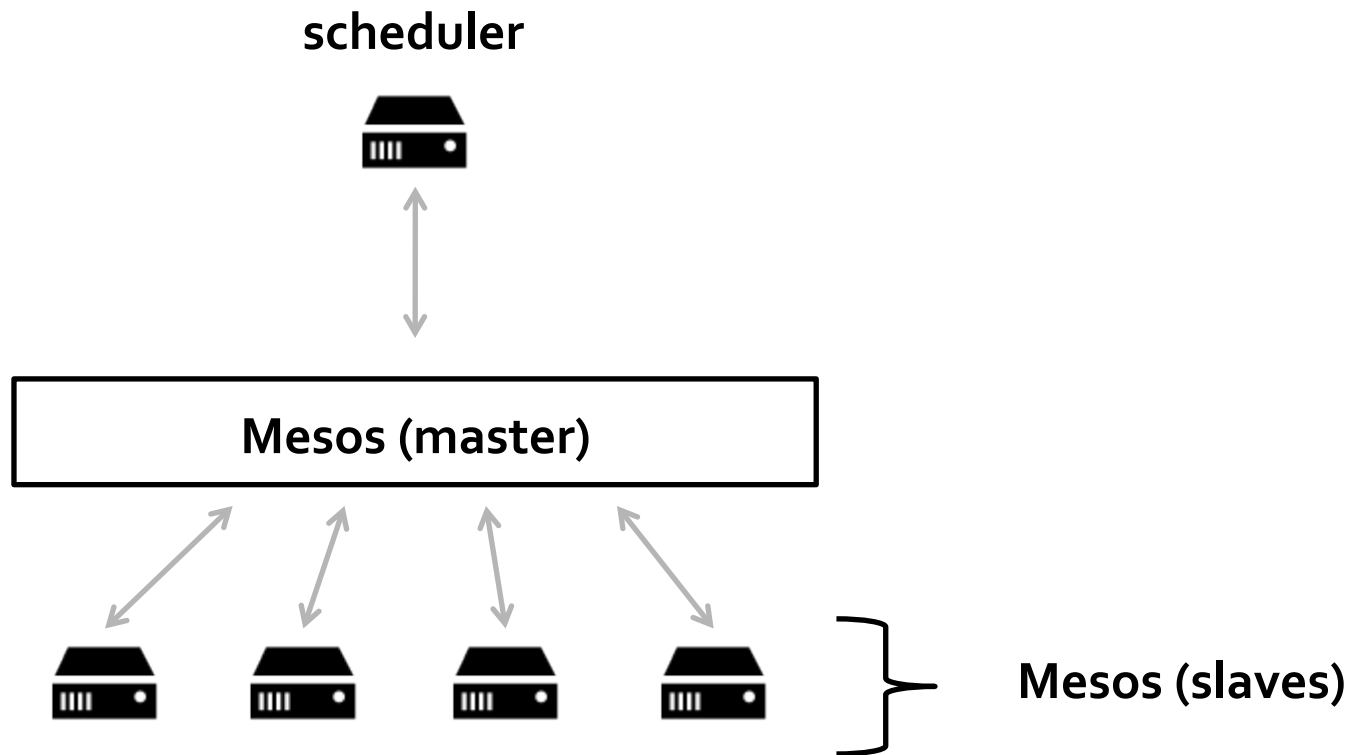
==

scheduling

Mesos: level of indirection



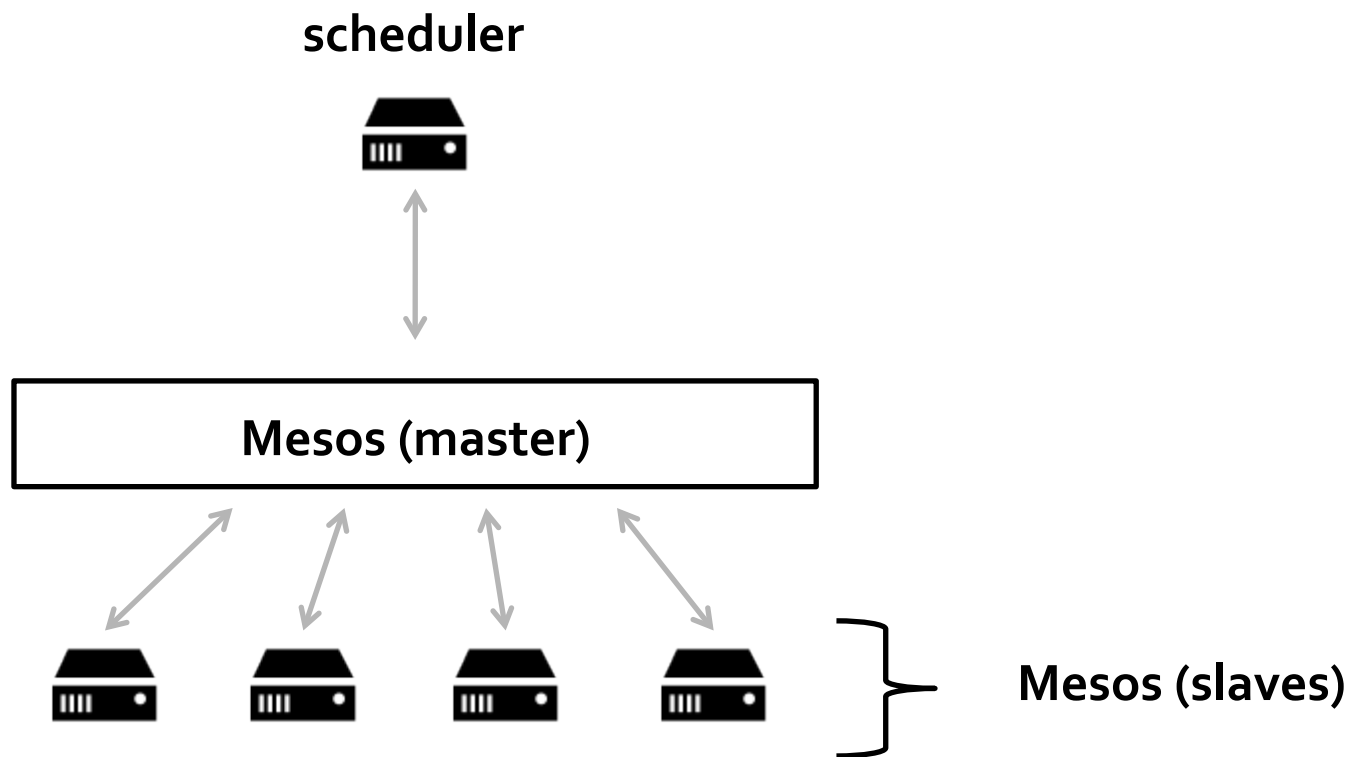
Mesos: level of indirection




Mesos: a level of indirection

+ provide common functionality every new distributed system *re-implements* like failure detection, task distribution, task starting, task monitoring, task killing, task cleanup!


Mesos: level of abstraction



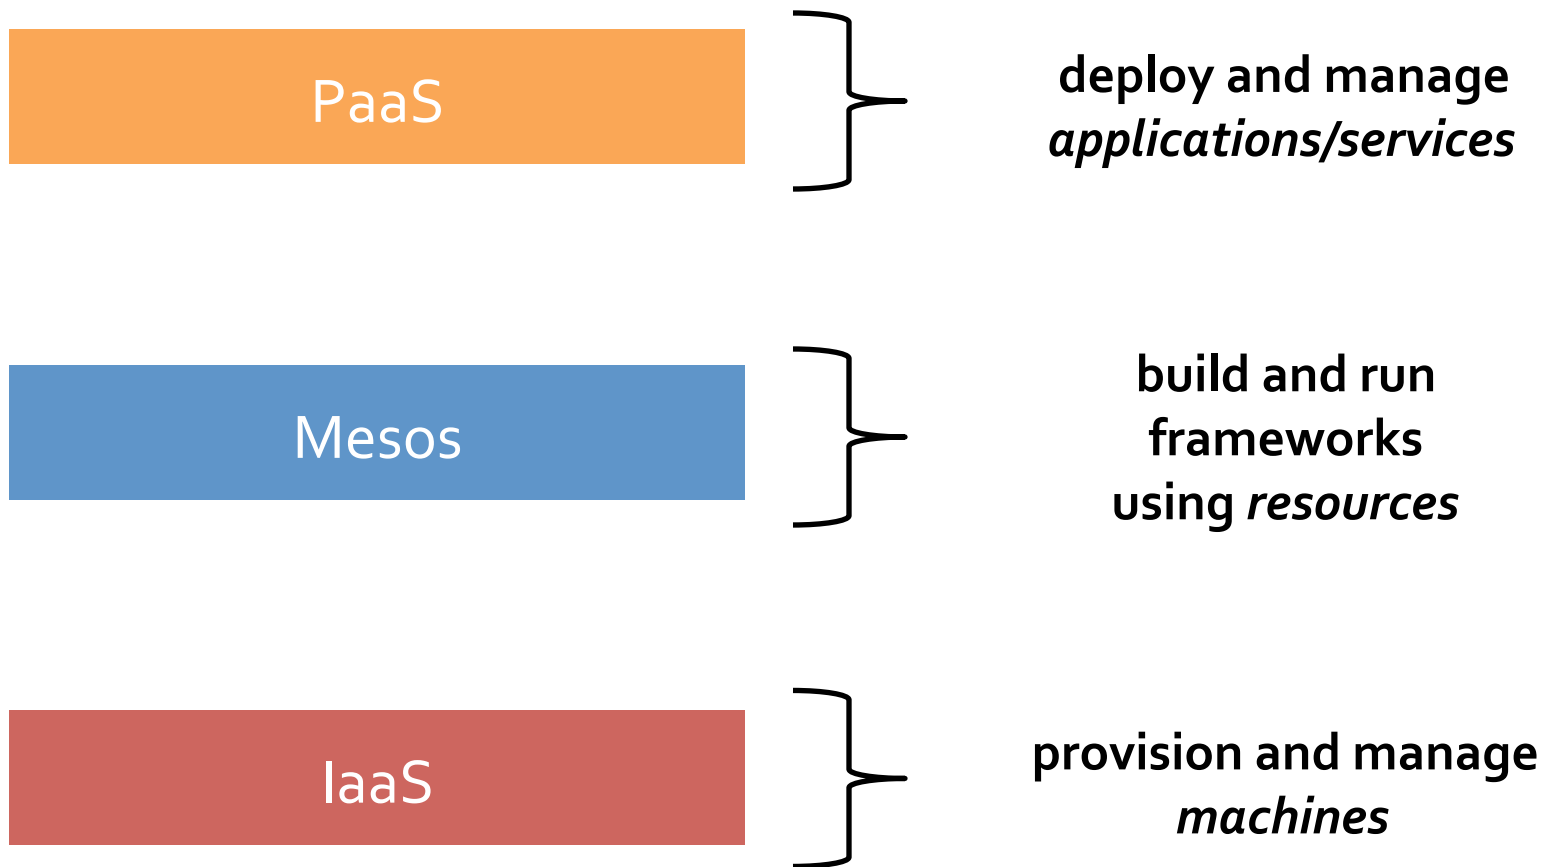
Mesos: level of abstraction




Mesos: level of abstraction




Mesos: level of abstraction




PaaS on Mesos




Mesos on IaaS



Mesos on IaaS++




Mesos: datacenter kernel




Apache Mesos is a
distributed system
for running and building
other distributed systems

Mesos is a cluster manager


agenda

- ① Mesos
- ② Spark on Mesos 
- ③ why Mesos?
 - ① multi-tenancy
 - ② fine-grained sharing
 - ③ why not?
- ④ long-lived services and other frameworks

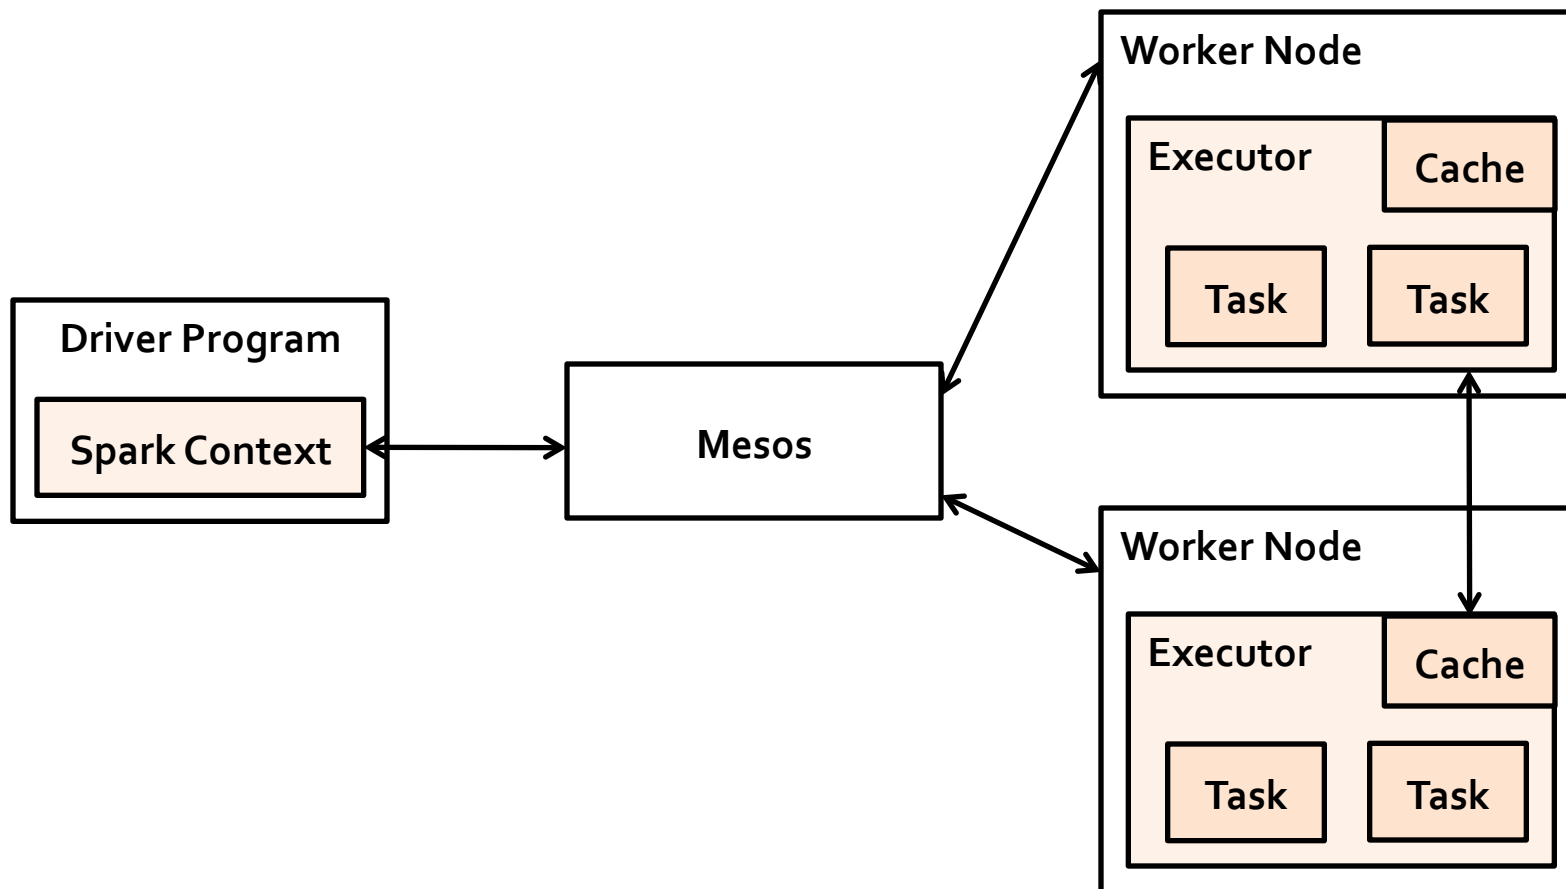
anatomy of Spark




anatomy of Spark



anatomy of Spark




Mesos is a distributed system with a master/slave architecture



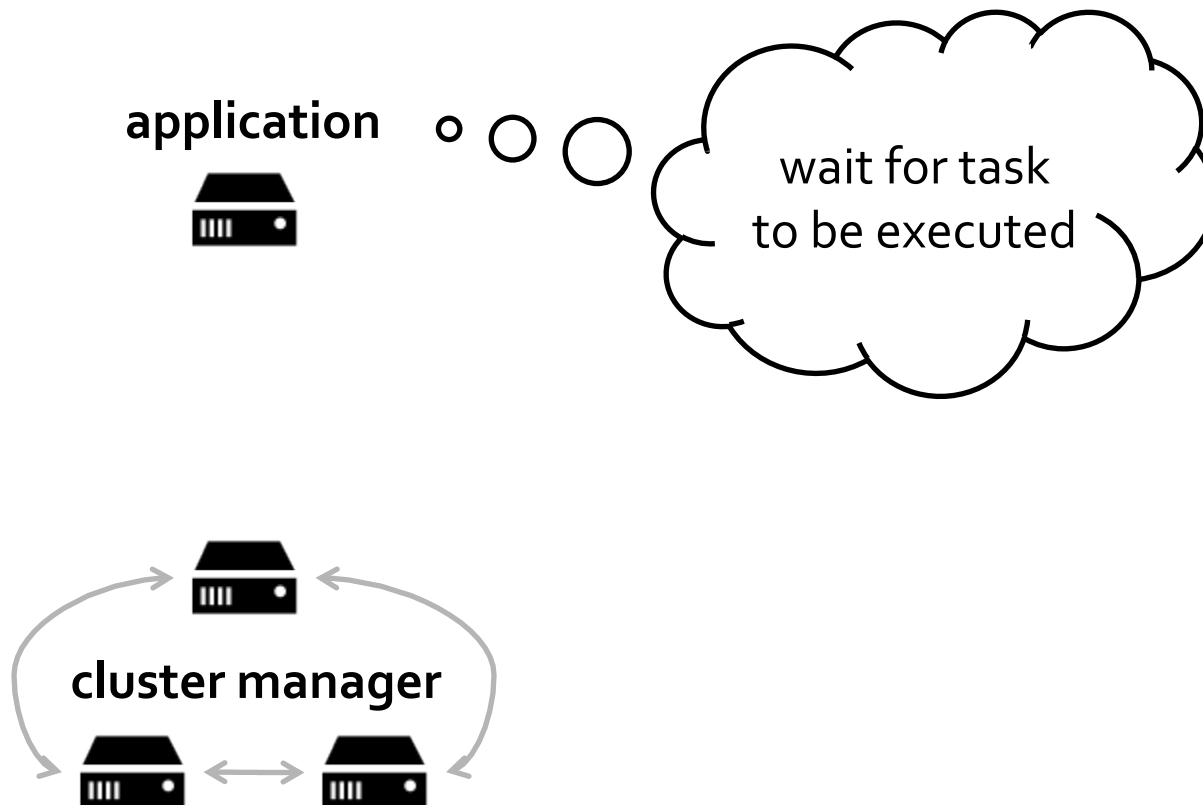
**Mesos challenged
the status quo
of cluster managers**

cluster manager status quo




the specification includes as much information as possible to assist the cluster manager in scheduling and execution

cluster manager status quo




cluster manager status quo



problems with specifications

- ① hard to specify certain desires or constraints
- ② hard to update specifications dynamically as tasks executed and finished/failed

an alternative model



a request is purposely simplified subset of a specification, mainly including the required resources

**question: what should you do
if you can't satisfy a request?**

**question: what should you do
if you can't satisfy a request?**

① wait until you can ...

**question: what should you do
if you can't satisfy a request?**

① wait until you can ...


② *offer* best you can immediately

**question: what should you do
if you can't satisfy a request?**


① wait until you can ...

② *offer* best you can immediately

Mesos model



Mesos model



an analogue: non-blocking sockets

application



kernel

```
write(s, buffer, size);
```

an analogue: non-blocking sockets

application




kernel

42 of 100 bytes written!

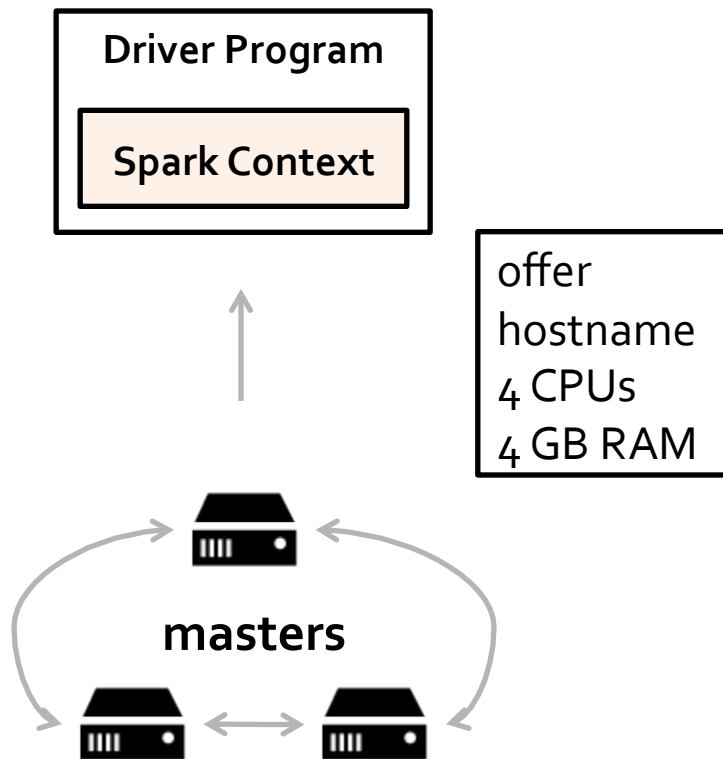
**offers represent the
current *snapshot*
of available resources
a framework can use**

**(requests are complimentary,
but not necessary; see
Google's Omega)**


Spark on Mesos



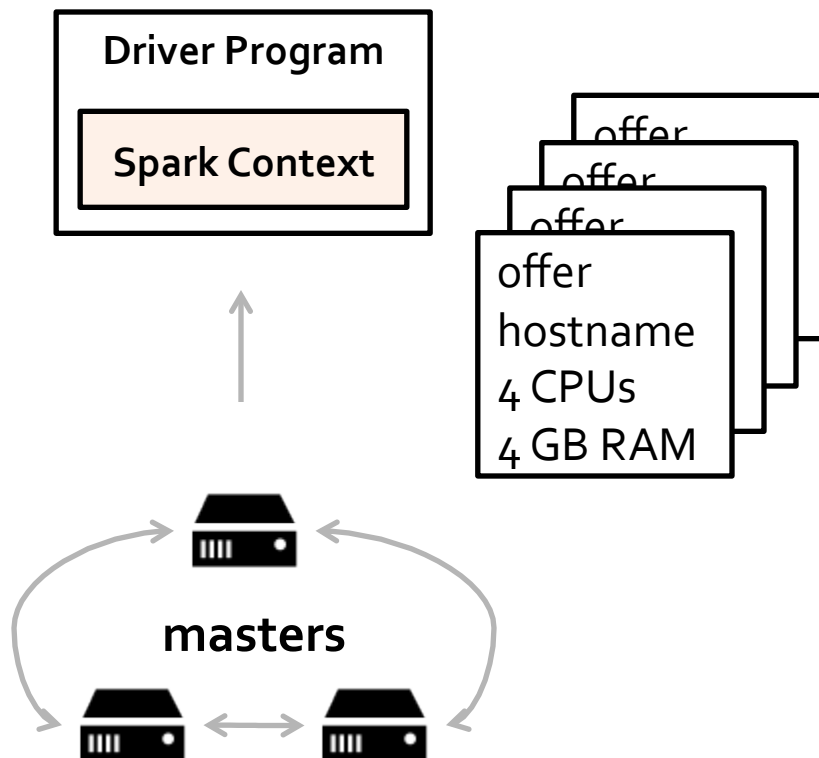
Spark on Mesos



Spark on Mesos

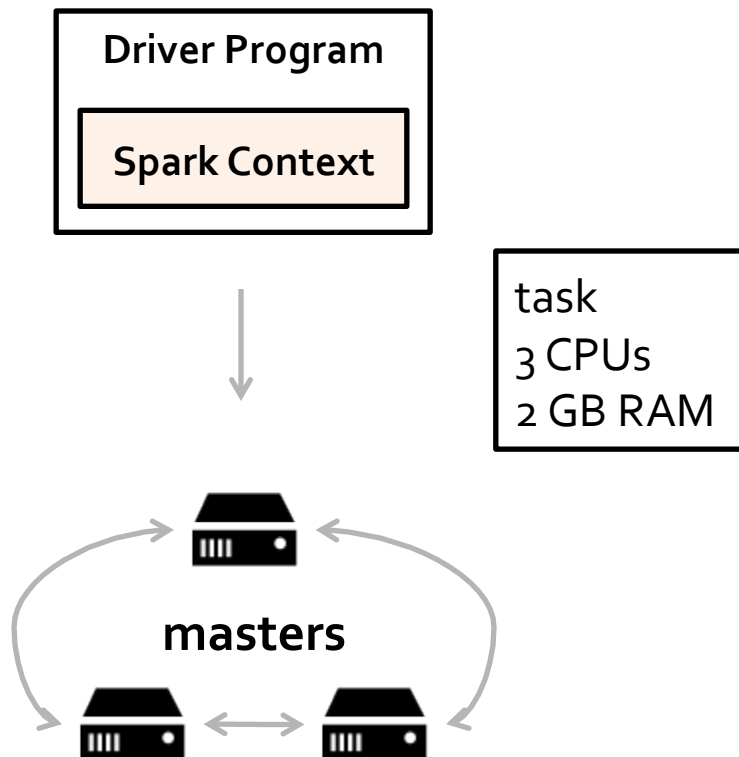


Spark on Mesos




Spark uses the offers to perform its own scheduling

Spark on Mesos



Spark uses the offers to perform its own scheduling

Spark on Mesos



Spark uses the offers to perform its own scheduling


“two-level scheduling”

“two-level scheduling”

Mesos: controls resource *allocations* to Spark


Spark: makes decisions about what tasks to run
given available resources

execution



frameworks launch fine-grained tasks for execution


execution




frameworks launch fine-grained tasks for execution

if necessary, a framework can provide an *executor* to handle the execution of a task


a task with a *command*




a task with a *command*




a task with a *command*




a task with an *executor*




a task with an *executor*




a task with an *executor*




a task with an *executor*



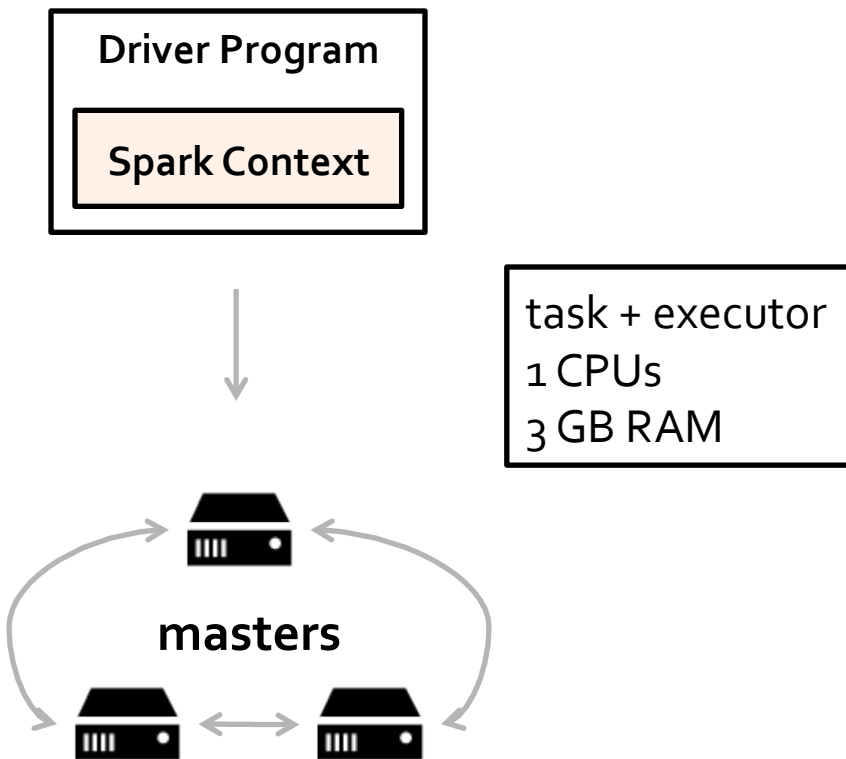
a task with an *executor*




a task with an *executor*




Spark execution




Spark execution




Spark execution




Spark execution



Spark execution



Spark execution



resource isolation

Mesos has containerization support on Linux
(built-in usage of cgroups and namespaces)

islator modules:


CPU (upper and lower bounds)

memory


network I/O (in development)

filesystem (using LVM, planned)


resource isolation




resource isolation



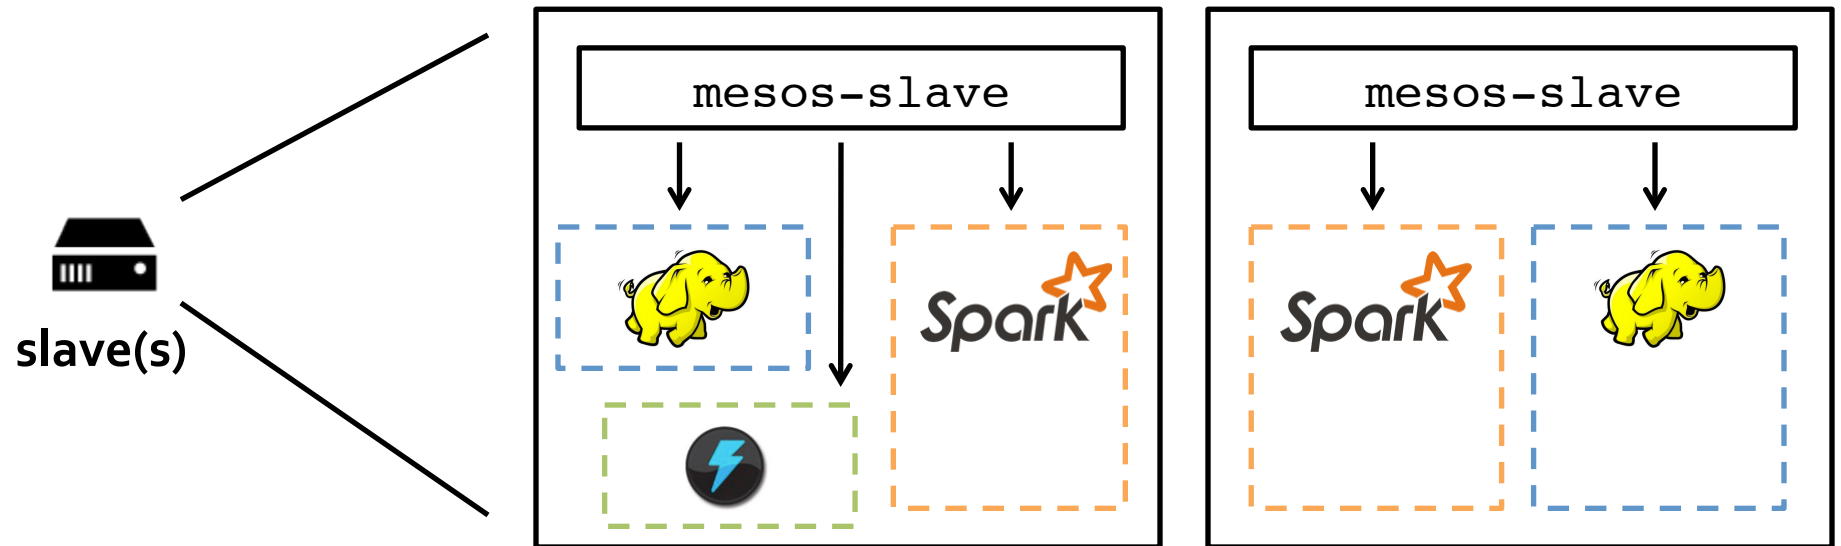
resource isolation



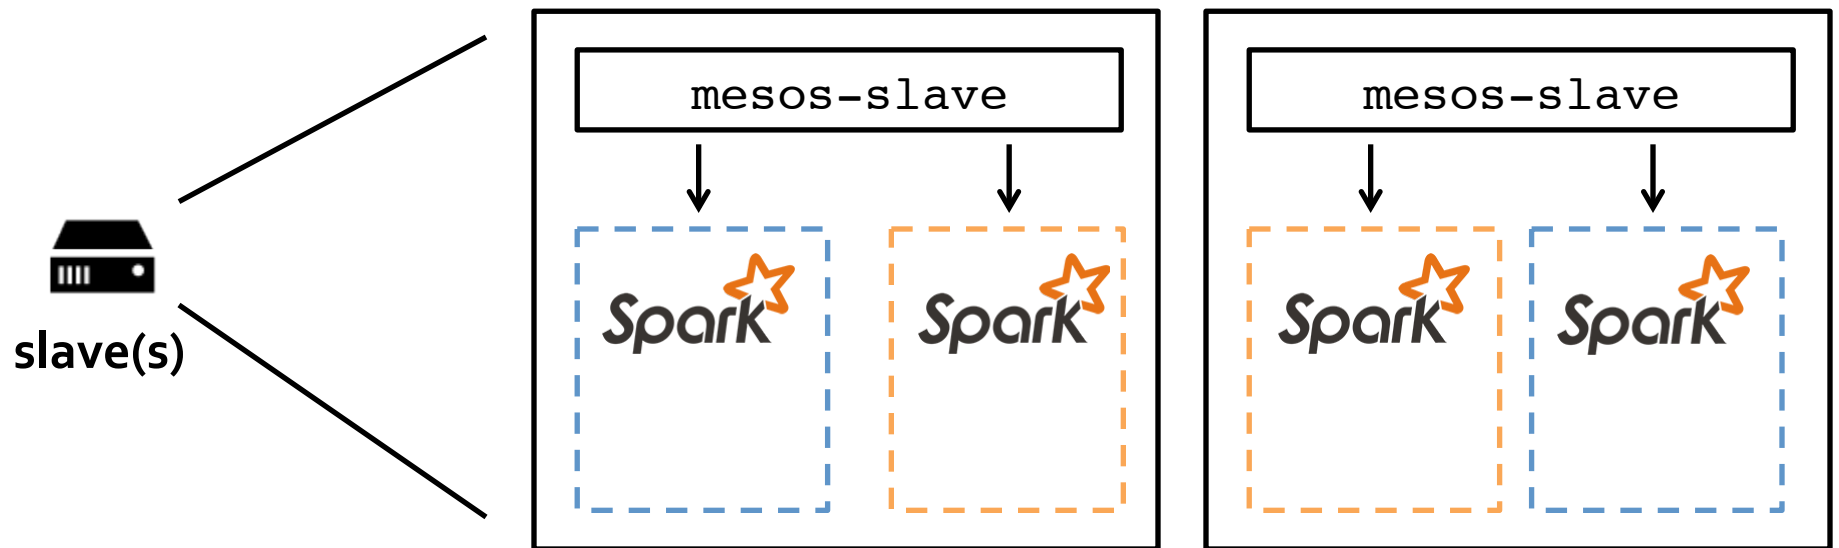
agenda

- ① Mesos
 - ② Spark on Mesos
 - ③ why Mesos?
 - ① multi-tenancy
 - ② fine-grained sharing
 - ③ why not?
 - ④ long-lived services and other frameworks
- 

multi-tenancy

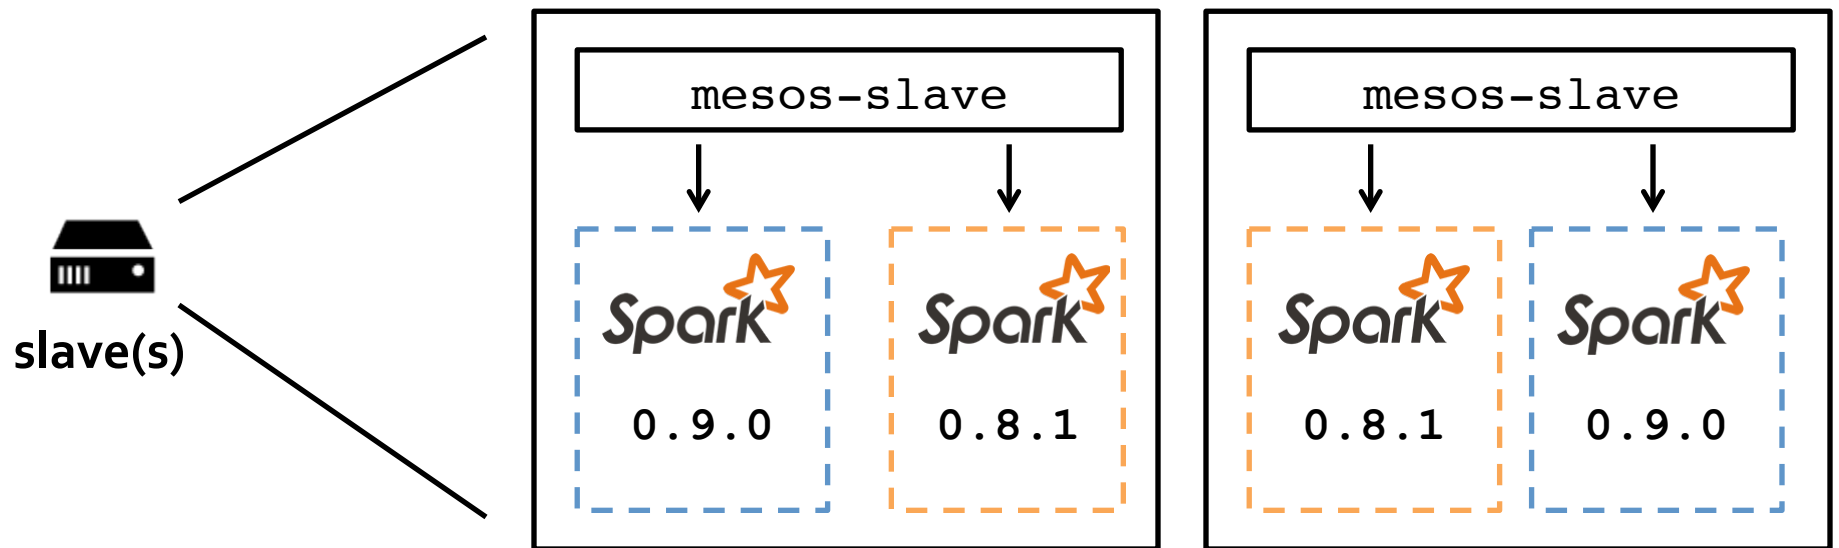


multi-tenancy (only Spark)




(can approximate w/ standalone mode by setting max # cores per application, otherwise get FIFO execution)

multi-tenancy (only Spark)




(run the tried and true and test out the new at the same time!)

agenda


- ① Mesos
- ② Spark on Mesos
- ③ why Mesos?
 - ① multi-tenancy
 - ② fine-grained sharing 
 - ③ why not?
- ④ long-lived services and other frameworks

fine-grained sharing




Spark executors only consume memory, can share CPU between

fine-grained sharing




Spark executors only consume memory, can share CPU between

fine-grained sharing




Spark executors only consume memory, can share CPU between

fine-grained sharing




Spark executors only consume memory, can share CPU between

fine-grained sharing




Spark executors only consume memory, can share CPU between

fine-grained sharing




Spark executors only consume memory, can share CPU between

fine-grained sharing




Spark executors only consume memory, can share CPU between

fine-grained sharing




Spark executors only consume memory, can share CPU between

fine-grained sharing




Spark executors only consume memory, can share CPU between

fine-grained sharing




Spark executors only consume memory, can share CPU between

fine-grained sharing



Spark executors only consume memory, can share CPU between


agenda

- ① Mesos
- ② Spark on Mesos
- ③ why Mesos?
 - ① multi-tenancy
 - ② fine-grained sharing
 - ③ why not? 
- ④ long-lived services and other frameworks

why not?

more moving pieces means more things to learn
and more things that can fail ...


master failover




after a new master is elected all
frameworks and slaves connect to
the new master

*all tasks keep running across master
failover!*


slave failover




slave failover




slave failover




slave failover




slave failover



slave failover @twitter




framework failover




framework re-registers with
master and resumes operation

*all tasks keep running across
framework failover!*

agenda


- ① Mesos
- ② Spark on Mesos
- ③ why Mesos?
 - ① multi-tenancy
 - ② fine-grained sharing
 - ③ why not?
- ④ long-lived services and other frameworks 

Apache Aurora (incubating)




Aurora is a Mesos framework that makes it easy to launch services written in Ruby, Java, Scala, Python, Go, etc!

Marathon (from Mesosphere)



Marathon



Marathon is a Mesos framework that makes it easy to launch services written in Ruby, Java, Scala, Python, Go, etc!



Jenkins on Mesos

ebay™ tech blog Where e-commerce meets world-class technology



Delivering eBay's CI Solution with Apache Mesos – Part I

by THE EBAY PAAS TEAM on 04/04/2014


in [CLOUD, DATA INFRASTRUCTURE AND SERVICES](#), [SOFTWARE ENGINEERING](#)

LINKS

[eBay Careers](#)

[eBay Developers Program](#)

(<http://bit.ly/1frLrLf>)



**Apache Mesos and Jenkins -
elastic build slaves**

(<http://bit.ly/1nHwM3r>)





Elastic Mesos: elastic.mesosphere.io



Launch an Apache Mesos Cluster in [3](#) [2](#) [1](#)

Packages:

Apache Mesos 0.14.2 [Release announcement](#)

04 Nov 2013

Apache Mesos 0.14.2 for [Ubuntu 13.04 \(AMD 64\)](#) and [Instructions](#)

[SHA 256](#)

Apache Mesos 0.14.2 for [Ubuntu 12.10 \(AMD 64\)](#) and [Instructions](#)

[SHA 256](#)

Apache Mesos 0.14.2 for [Ubuntu 12.04 \(AMD 64\)](#) and [Instructions](#)

[SHA 256](#)

Apache Mesos 0.14.2 for [Debian 7 \(AMD 64\)](#) and [Instructions](#)

[SHA 256](#)

Apache Mesos 0.14.2 for [CentOS 6 \(x86_64\)](#) and [Instructions](#)

[SHA 256](#)

Apache Mesos 0.14.2 for [Red Hat 6 \(x86_64\)](#) and [Instructions](#)

[SHA 256](#)

Thank You!

mesos.apache.org

mesos.apache.org/blog

[@ApacheMesos](https://twitter.com/ApacheMesos)

