

ONOS

Open Network Operating System

Architecture Update

ONRC October 2014

Thomas Vachuska

tom@onlab.us

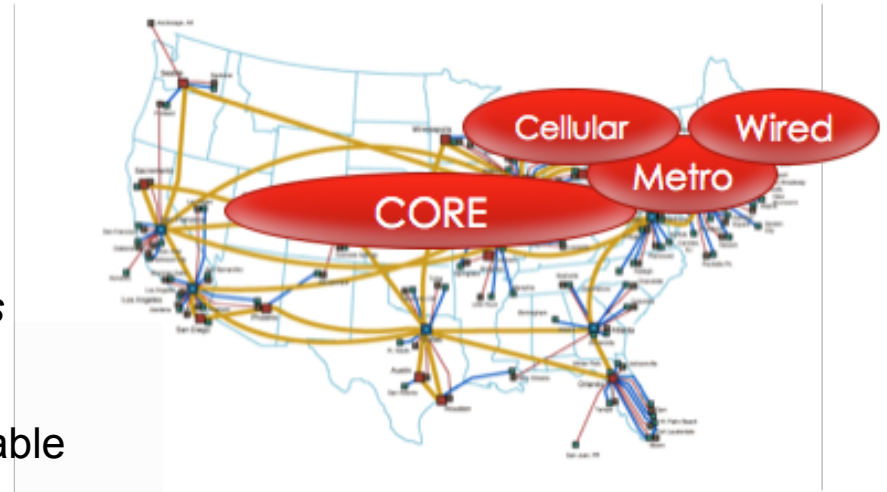
ONOS:

SDN OS for Service Provider Networks

- Scalability, High Availability & Performance
- Northbound & Southbound Abstractions
- Modularity

Service Provider Networks

- WAN core backbone
 - Multi-Protocol Label Switching (MPLS) with Traffic Engineering (TE)
 - *200-500 routers, 5-10K ports*
- Metro Networks
 - Metro cores for access networks
 - *10-50K routers, 2-3M ports*
- Cellular Access Networks
 - LTE for a metro area
 - *20-100K devices, 100K-100M ports*
- Wired access / aggregation
 - Access network for homes; DSL/Cable
 - *10-50K devices, 100K-1M ports*

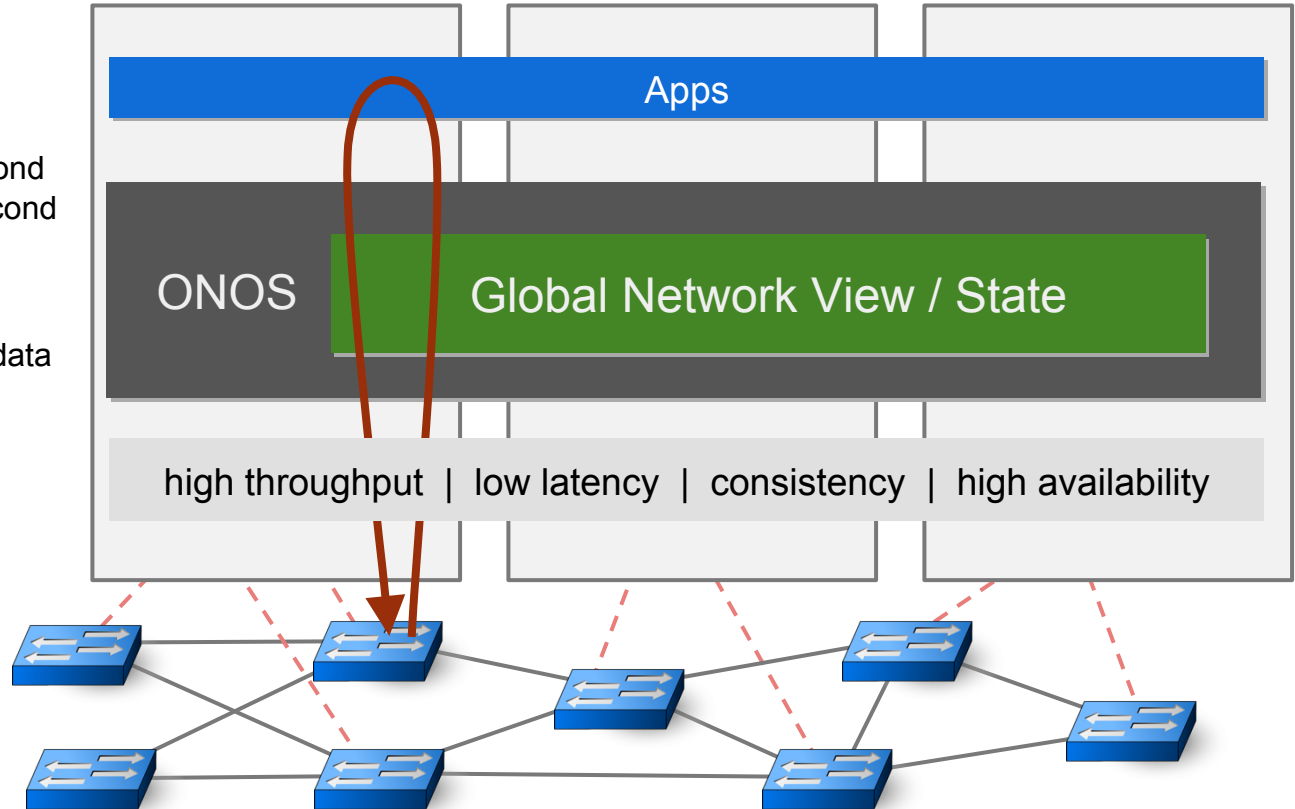


Key Performance Requirements

High Throughput:
~500K-1M paths setups / second
~3-6M network state ops / second

High Volume:
~500GB-1TB of network state data

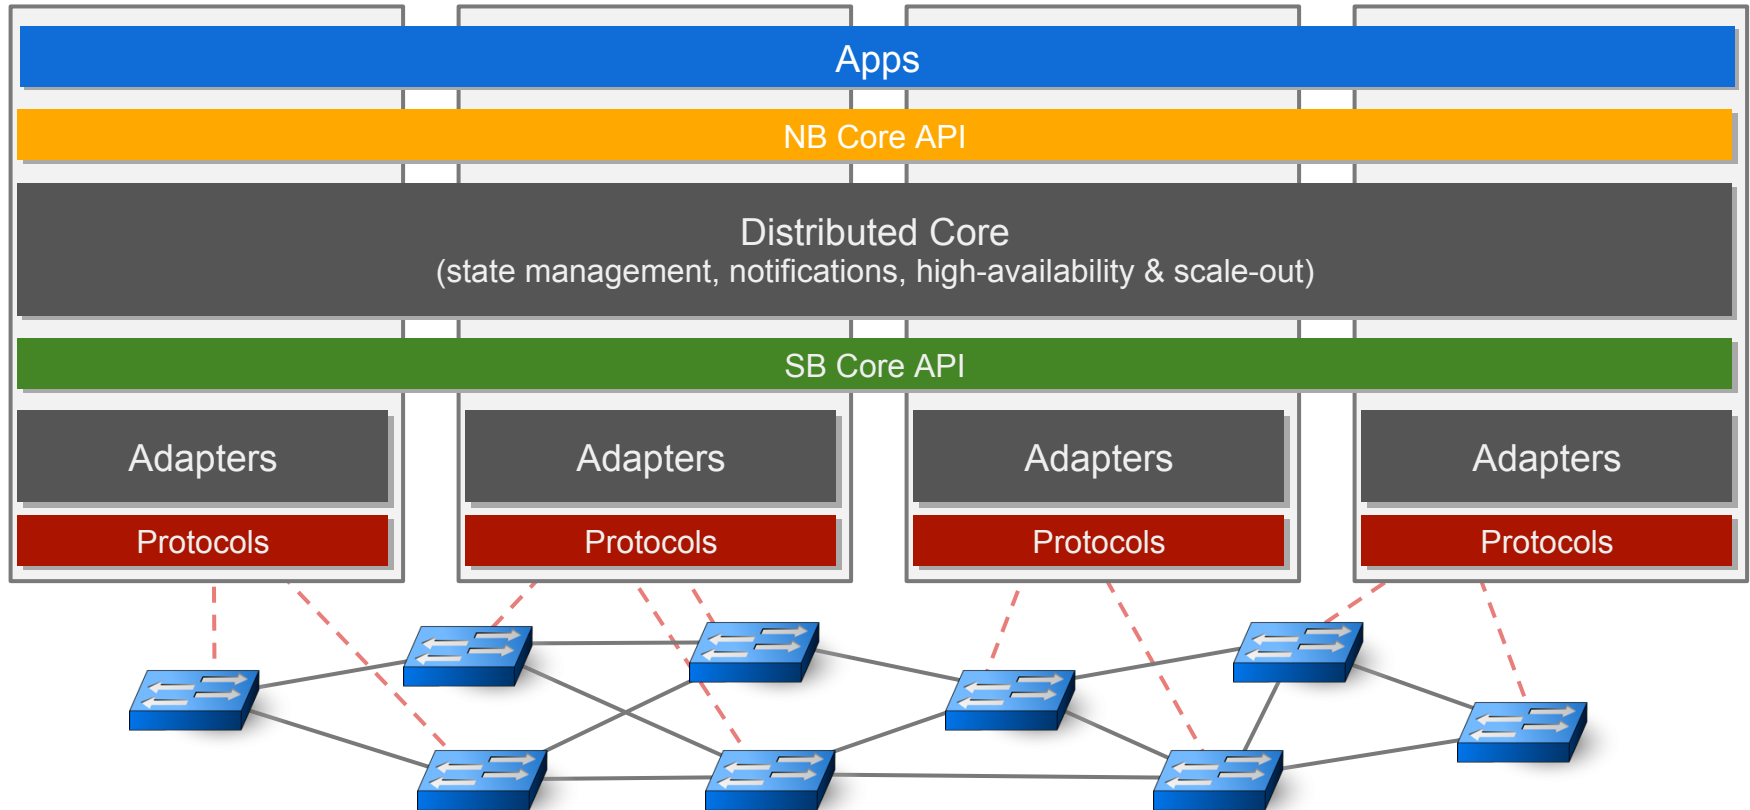
Difficult challenge!



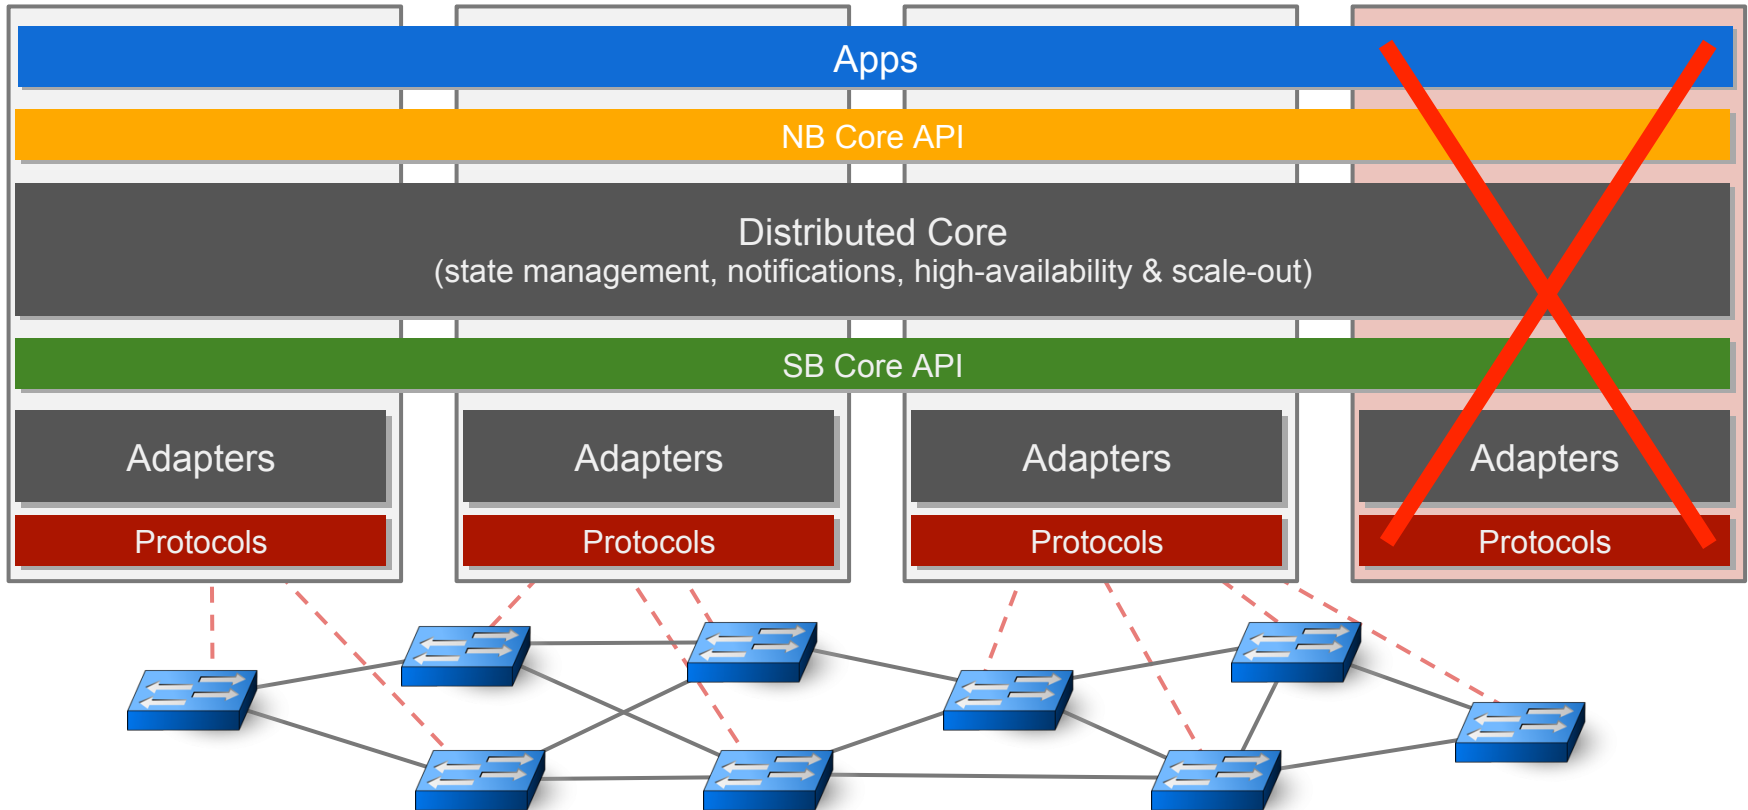
Why Operating System?

- Provides useful services to applications
 - e.g. maintains connection persistence
- Provides framework for driving devices via arbitrary protocols
- Arbitrates shared network resources
- Provides abstractions to simplify resource sharing
 - application intent, network graph & device abstractions
- Isolates and protects resources, tenants & users
 - resource virtualization
- Comes with an SDK
 - APIs & docs, debugging, emulation, monitoring

Distributed Architecture



Distributed Architecture



ONOS Evolution

- Written in Java
- First prototype
 - basic functionality, OpenFlow 1.0
 - scale-out, high-availability, northbound graph abstraction
- Second prototype
 - performance, scale improvements over first generation
- Both
 - prototype quality code
 - OpenFlow as the only southbound protocol
 - relied heavily on open-source off-the-shelf components

ONOS November Release

- Many improvements to distributed core
 - revamped NB & SB interfaces
 - revamped distributed state management
- New abstraction & API
 - application intents
- New & pluggable southbound
 - OpenFlow 1.3 support
 - plugin architecture for legacy protocols
- Improved GUI & CLI
- Modularity
 - revamped code-base for modularity
 - built atop OSGi container - Apache Karaf

ONOS November Release

Northbound Abstraction:

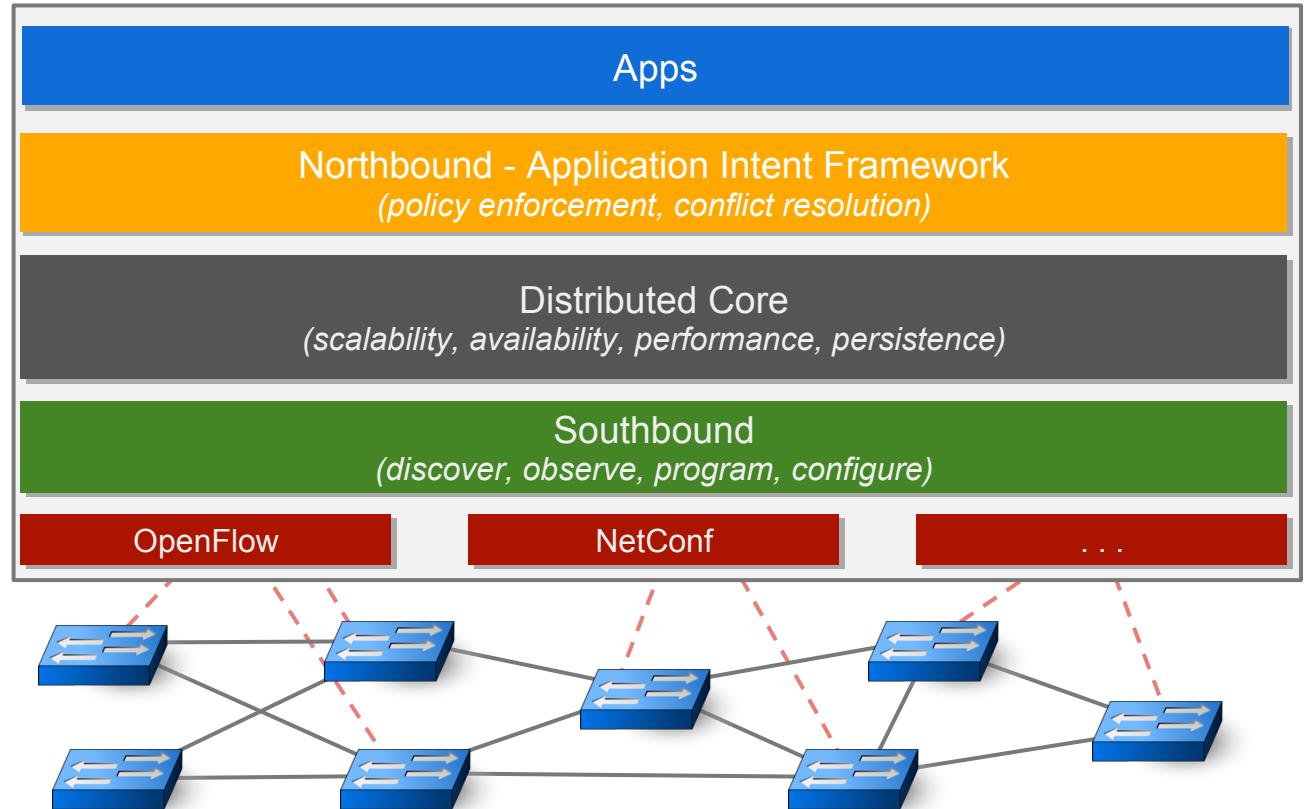
- network graph
- application intents

Core:

- distributed
- protocol independent

Southbound Abstraction:

- generalized OpenFlow
- pluggable & extensible



Application Intent Framework

- Application specifies high-level intents; not low-level rules
 - focus on *what* should be done, rather than *how* it should be done
- Intents are compiled into actionable objectives which are installed into the environment
 - e.g. *HostToHostIntent* compiles into two *PathIntents*
- Resources required by objectives are then monitored
 - e.g. link vanishes, capacity or lambda becomes available
- Intent subsystem reacts by recompiling intent and re-installing revised objectives

Distributed Core

- Distributed state management framework
 - built for high-availability and scale-out
- Different types of state require different types of synchronization
 - fully replicated
 - master / slave replicated
 - partitioned / distributed
- Novel topology replication technique
 - *logical* clock in each instance timestamps events observed in underlying network
 - *logical* timestamps ensure state evolves in consistent and *ordered* fashion
 - allows rapid convergence without complex coordination
 - applications receive notifications about topology changes

Distributed Core

Application Intents




- immutable
- durable & replicated



- ## 3-way replication
- H/A execution via distributed queues

Global Network View

- eventually consistent
- fully replicated






Optimistic Replication

- gossip based
- anti-entropy
- partial ordering

Flow Table Entries

- strongly consistent
- partitioned



Master/Backup Replication

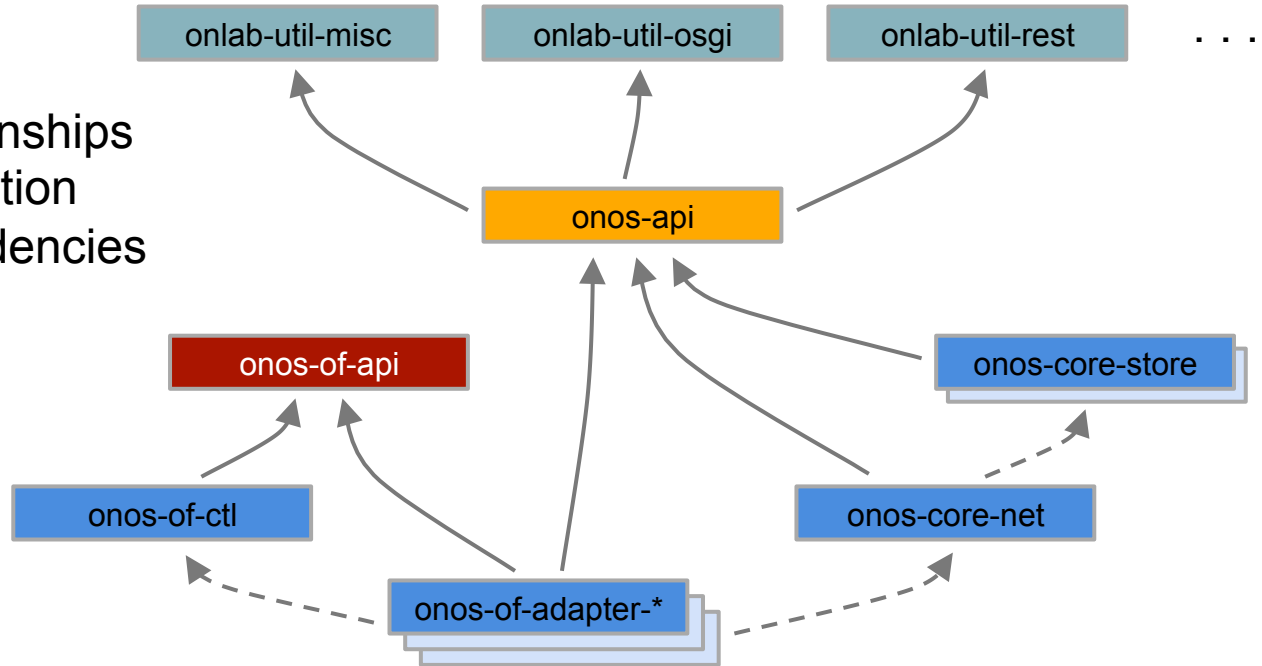
- Distribution & replication methods optimized for the type of state
- Based on size and read/write access patterns

Modularity Objectives

- Increase architectural coherence, testability and maintainability
 - establish tiers with crisply defined boundaries and responsibilities
 - setup code-base to follow and enforce the tiers
- Facilitate extensibility and customization by partners and users
 - unit of replacement is a module
- Avoid speciation of the ONOS code-base
 - APIs setup to encourage extensibility and community code contributions
- Preempt code entanglement, i.e. cyclic dependencies
 - reasonably small modules serve as firewalls against cycles
- Facilitate pluggable southbound

ONOS Modules

- Well-defined relationships
- Basis for customization
- Avoid cyclic dependencies



What's coming on December 5th?

- ONOS with all its key features
 - scalability, high-availability, performance
 - northbound abstractions (application intents)
 - southbound abstractions (OpenFlow adapters)
 - modular code-base
- Open source
 - ONOS code-base on GitHub
 - documentation & infrastructure processes to engage the community
- Use-case demonstrations
 - SDN-IP, Packet-Optical
- Sample applications
 - reactive forwarding, mobility, proxy arp

Check out posters -- learn details from ONOS architects and developers