



ONOS

Open Network Operating System

An Experimental Open-Source Distributed SDN OS

Umesh Krishnaswamy, Pankaj Berde, Jonathan Hart, Masayoshi Kobayashi, Pavlin Radoslavov, Tim Lindberg, Rachel Sverdlov, Suibin Zhang, William Snow, Guru Parulkar

Logically Centralized NOS – Key Questions


Routing

TE

Mobility

How to realize
global network view?

Global Network View



Network OS

Openflow

Scale-out?

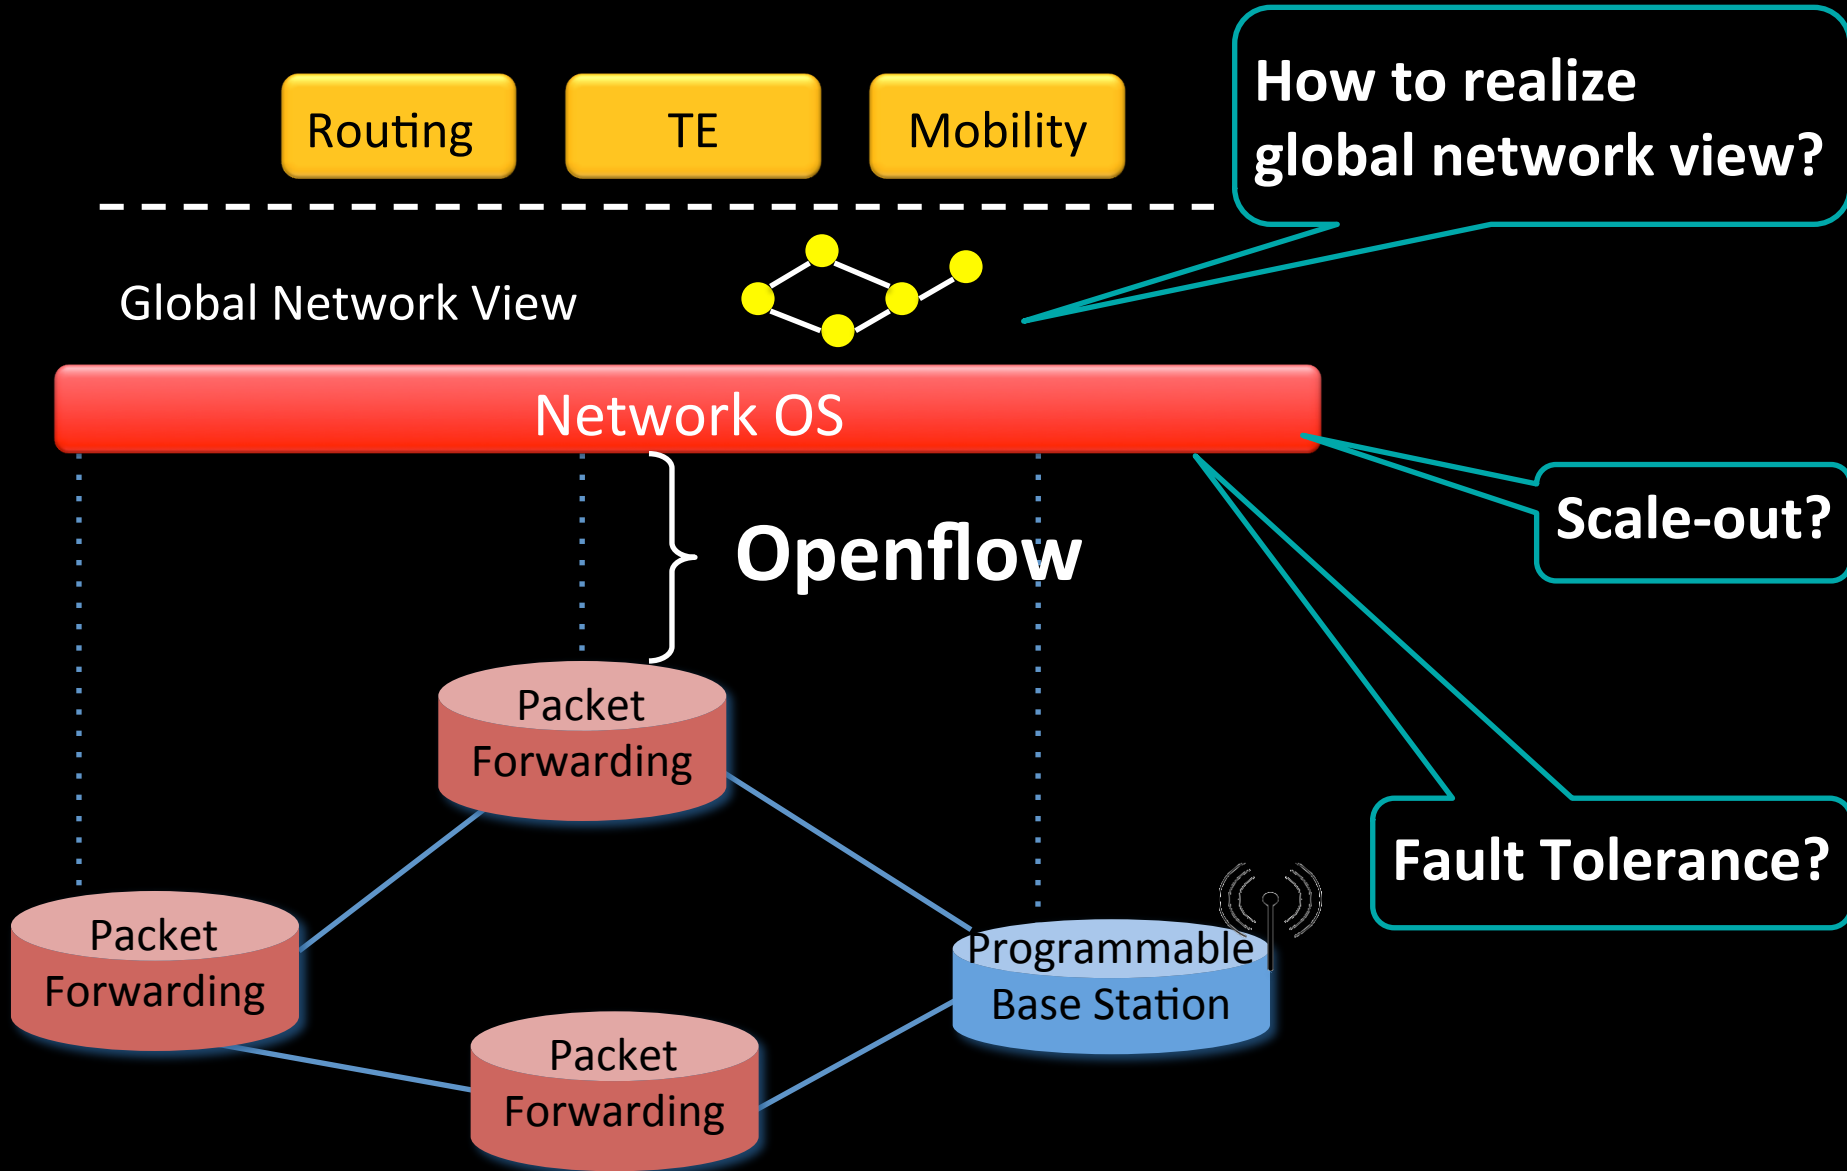
Fault Tolerance?

Packet
Forwarding

Packet
Forwarding

Packet
Forwarding

Programmable
Base Station



Prior Work

ONIX

Distributed control platform for large-scale networks

Focus on reliability, scalability, and generality

State distribution primitives, global network view, ONIX API

Other Work

Helios (NEC), Midonet (Midokura), Hyperflow, Maestro, Kandoo

NOX, POX, Beacon, Floodlight, Trema controllers

Community needs an open source distributed SDN OS

ONOS High Level Architecture

Network Graph
Eventually consistent



Titan Graph DB

Cassandra In-Memory DHT

Distributed Registry
Strongly Consistent



Zookeeper

Instance 1

OF Controller



Floodlight

Instance 2

OF Controller




Floodlight

Instance 3

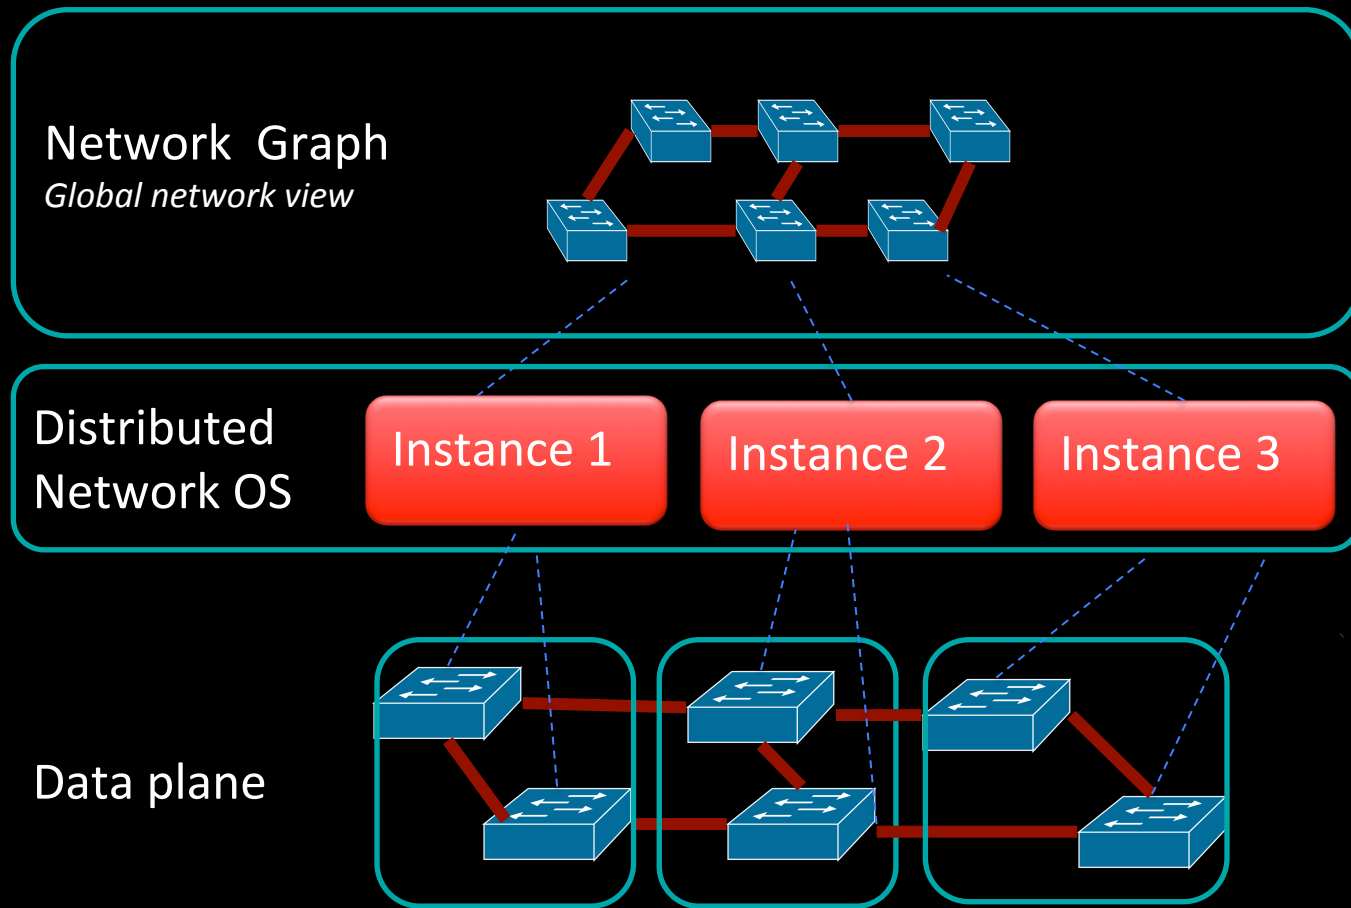
OF Controller



Floodlight



ONOS Scale-Out



An instance is responsible for maintaining a part of network graph
Control capacity can grow with network size or application need

ONOS Control Plane Failover

Distributed
Registry

Master
Switch A = ONOS2
Candidates = ONOS3

Master
Switch A = ONOS2
Candidates = ONOS3


Master
Switch A = ONOS2
Candidates = ONOS3

Distributed
Network OS

Instance 1

Instance 2

Instance 3

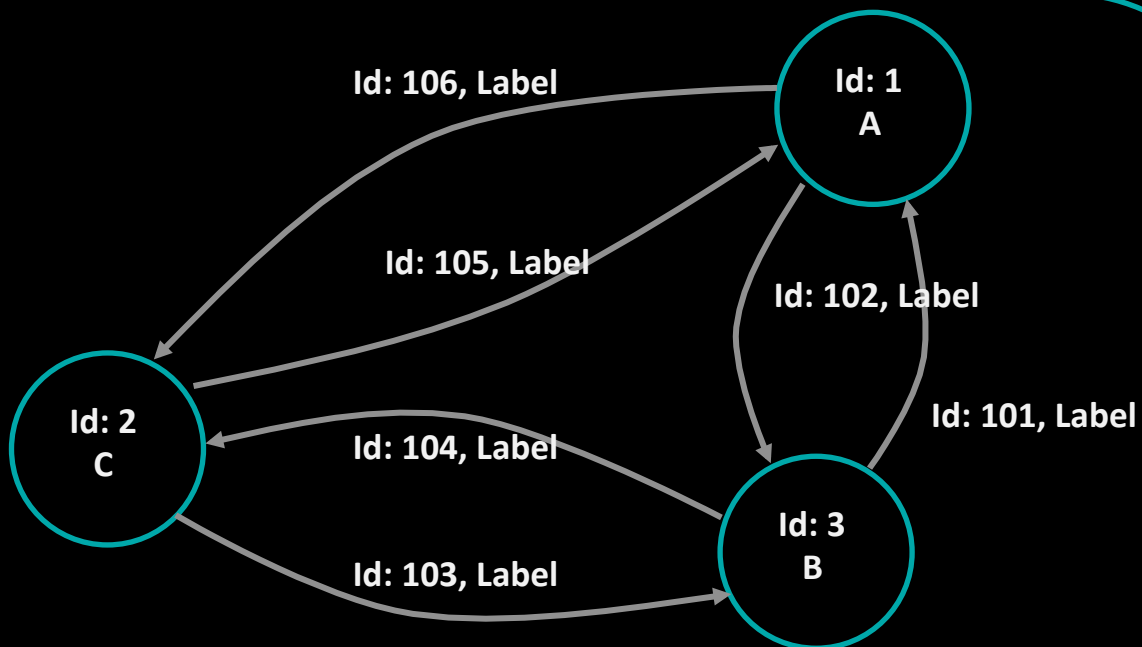


ONOS Network Graph Abstraction


Network Graph




Titan Graph DB



Cassandra
In-memory DHT




Network Graph



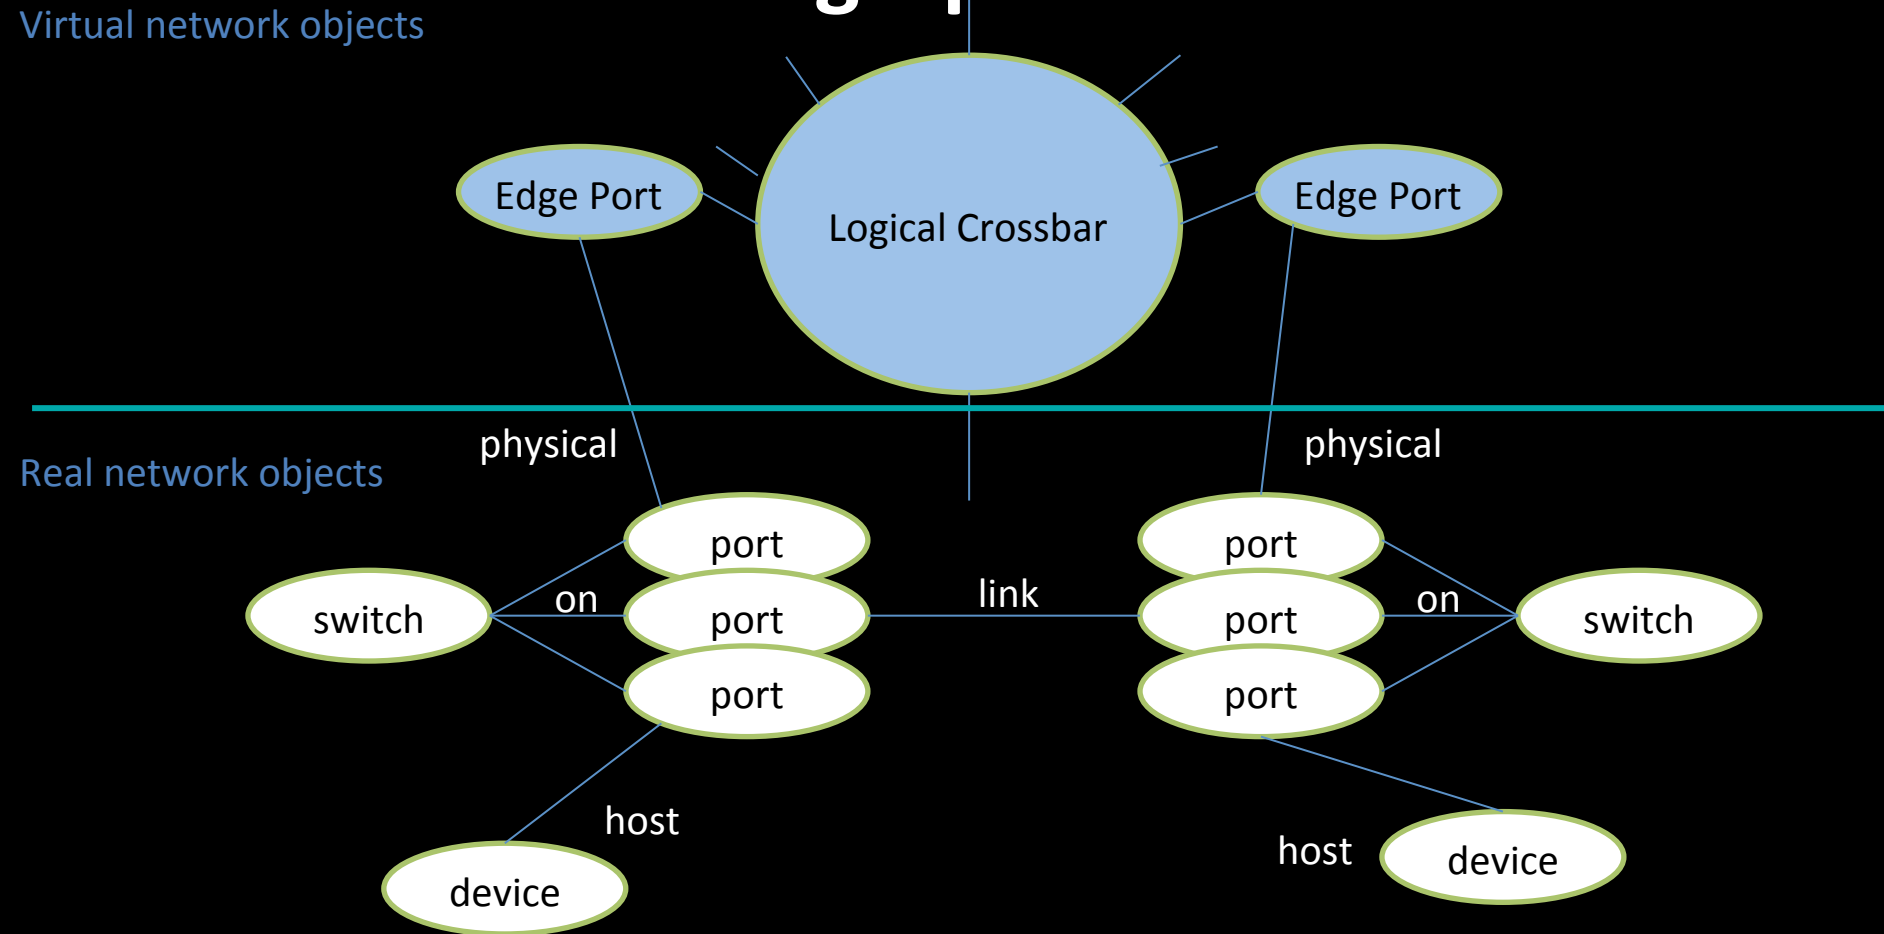
- Network state is naturally represented as a graph
- Graph has basic network objects like switch, port, device and links
- Application writes to this graph & programs the data plane

Example: Path Computation App on Network Graph



- Application computes path by traversing the links from source to destination
 - Application writes each flow entry for the path
- Thus path computation app does not need to worry about topology maintenance


Example: A simpler abstraction on network graph?



- App or service on top of ONOS
- Maintains mapping from simpler to complex

Thus makes applications even simpler and enables new abstractions

Network Graph in Titan and Cassandra



Vertex
represented as
Cassandra row

Property (e.g. dpid)	Property (e.g. state)	Property ...	Edge	Edge
----------------------------	-----------------------------	--------------	------	------


Row indices for fast vertex centric queries

Edge represented
as Cassandra
column

Column				Value	
Label id + direction	Primary key	Edge id	Vertex id	Signature properties	Other properties

Network Graph and Switches


Network Graph: Switches



Switch Manager



Switch Manager

Switch Manager





Network Graph and Link Discovery

Network Graph: Links



Devices and Network Graph

Network Graph: Devices




Path Computation with Network Graph

Path Computation

Path Computation

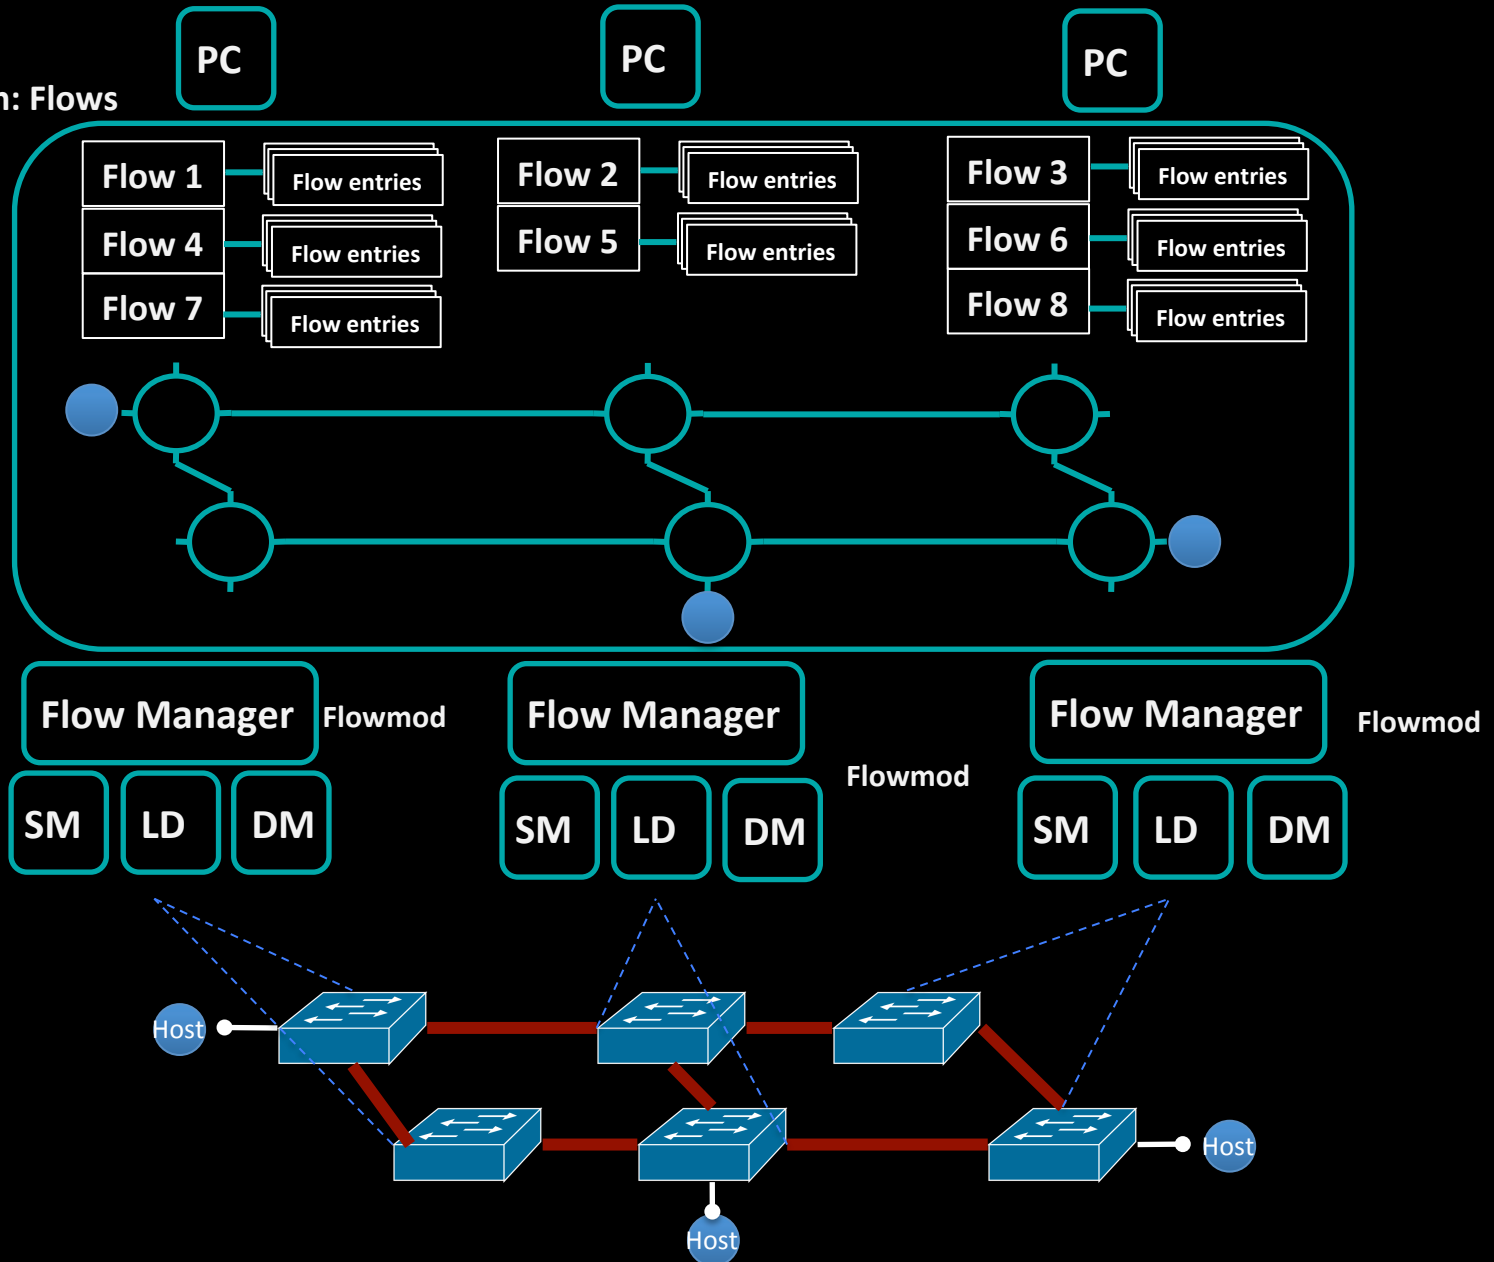
Path Computation

Network Graph: Flow Paths



Network Graph and Flow Manager

Network Graph: Flows




Consistency Definition

- **Strong Consistency:** Upon an update to the network state by an instance, all subsequent reads by any instance returns the last updated value.
- **Strong consistency adds complexity and latency to distributed data management.**
- **Eventual consistency is slight relaxation – allowing readers to be behind for a short period of time.**

Strong Consistency using Registry

Network Graph



Distributed Network OS

Instance 1

Instance 2

Instance 3

Registry

Switch A
Master = ONOS 1

Switch A
Master = ONOS 1

Switch A
Master = ONOS 1

All instances
Switch A Master = NONE

All instances
Switch A Master = NONE

Instance 1 Switch A Master = ONOS 1
Instance 2 Switch A Master = ONOS 1
Instance 3 Switch A Master = ONOS 1

Timeline

Master elected for switch A


Cost of Locking

Why Strong Consistency is needed for Master Election

- Weaker consistency might mean Master election on instance 1 will not be available on other instances.
- That can lead to having multiple masters for a switch.
- Multiple Masters will break our semantic of control isolation.
- Strong locking semantic is needed for Master Election

Eventual Consistency in Network Graph

Network Graph



Distributed
Network OS

Instance 1

Instance 2


Instance 3

DHT

Switch A
State = ACTIVE

Switch A
State = ACTIVE

Switch A
STATE = ACTIVE



Cost of Eventual Consistency

- Short delay will mean the switch A state is not ACTIVE on some ONOS instances in previous example.
- Applications on one instance will compute flow through the switch A while other instances will not use the switch A for path computation.
- Eventual consistency becomes more visible during control plane network congestion.

Why is Eventual Consistency good enough for Network State?

- Physical network state changes asynchronously
 - Strong consistency across data and control plane is too hard
 - Control apps know how to deal with eventual consistency
- In the current distributed control plane, each router makes its own decision based on old info from other parts of the network and it works fine
- Strong Consistency is more likely to lead to inaccuracy of network state as network congestions are real.

What is Next for ONOS



ONOS Core

Performance tuning

Events, callbacks and publish/subscribe API

Expand graph abstraction for more types of network state



ONOS Apps

Intelligent control plane: SDN-IP

Demonstrate two new use cases by end of the year



Community

Limited availability release in Q3

Build and assist developer community outside ON.LAB

Support deployments in R&E networks

ON.LAB

www.onlab.us