

**Lecture 6:**

# **Performance Optimization Part 1: Work Distribution**

---

**Parallel Computer Architecture and Programming  
CMU 15-418/15-618, Spring 2015**

# Tunes

## The Heavy

Colleen

**(Great Vengeance and Furious Fire)**

*“Colleen? Ha, that wasn’t about a girl. We wrote that one about the dangers of premature program optimization. It burns everyone, and it’s certainly burned me.”*

*- Kelvin Swaby*


# Today

- **Solver example in the message passing model**
- **Begin discussing techniques for optimizing parallel programs**
- **CUDA and GPU programming quick-review**

# **The grid solver in a message passing programming model**

# One more time... recall the grid-based solver example

Previously expressed using mechanisms from data parallel and SPMD programming models



**Update all red cells in parallel**

**When done updating red cells ,  
update all black cells in parallel  
(respect dependency on red cells)**

**Repeat until convergence**

# Recall: data-parallel solver implementation

## ■ Synchronization:

- forall loop iterations are independent (can be parallelized)
- Implicit barrier at end of outer forall loop body

## ■ Communication

- Implicit in loads and stores (like shared address space)
- Special built-in primitives: e.g., reduce

```
int n;  
float* A = allocate(n+2, n+2);  
  
void solve(float* A) {  
    bool done = false;  
    float diff = 0.f;  
    while (!done) {  
        for_all (red cells (i,j)) {  
            float prev = A[i,j];  
            A[i,j] = 0.2f * (A[i-1,j] + A[i,j-1] + A[i,j] +  
                           A[i+1,j] + A[i,j+1]);  
            reduceAdd(diff, abs(A[i,j] - prev));  
        }  
        if (diff/(n*n) < TOLERANCE)  
            done = true;  
    }  
}
```

# Recall: shared address space implementation with explicit synchronization (locks and barriers)

```
int n; // grid size
bool done = false;
float diff = 0.0;
LOCK myLock;
BARRIER myBarrier;
```

```
// allocate grid
float* A = allocate(n+2, n+2);
void solve(float* A) {
    float myDiff;
    int threadId = getThreadId();
    int myMin = 1 + (threadId * n / NUM_PROCESSORS);
    int myMax = myMin + (n / NUM_PROCESSORS)

    while (!done) {
        float myDiff = 0.f;
        diff = 0.f;
        barrier(myBarrier, NUM_PROCESSORS);
        for (j=myMin to myMax) {
            for (i = red cells in this row) {
                float prev = A[i,j];
                A[i,j] = 0.2f * (A[i-1,j] + A[i,j-1] + A[i,j] +
                               A[i+1,j], A[i,j+1]);
                myDiff += abs(A[i,j] - prev));
            }
            lock(myLock);
            diff += myDiff;
            unlock(myLock);
            barrier(myBarrier, NUM_PROCESSORS);
            if (diff/(n*n) < TOLERANCE)
                done = true;
            barrier(myBarrier, NUM_PROCESSORS);
        }
    }
}
```

I asked the class:

Could you do better than three barriers?

// check convergence, all threads get same answer

# Shared address space solver: one barrier

## Idea:

Remove dependencies by using different `diff` variables in successive loop iterations

Trade off footprint for removing dependencies!

Three variables instead of one.

But now one barrier instead of three.

(a common parallel programming technique)

```
int n;                // grid size
bool done = false;
LOCK    myLock;
BARRIER myBarrier;
float diff[3]; // global diff, but now 3 copies

float *A = allocate(n+2, n+2);

void solve(float* A) {
    float myDiff; // thread local variable
    int index = 0; // thread local variable

    diff[0] = 0.0f;
    barrier(myBarrier, NUM_PROCESSORS); // one-time only: just for init

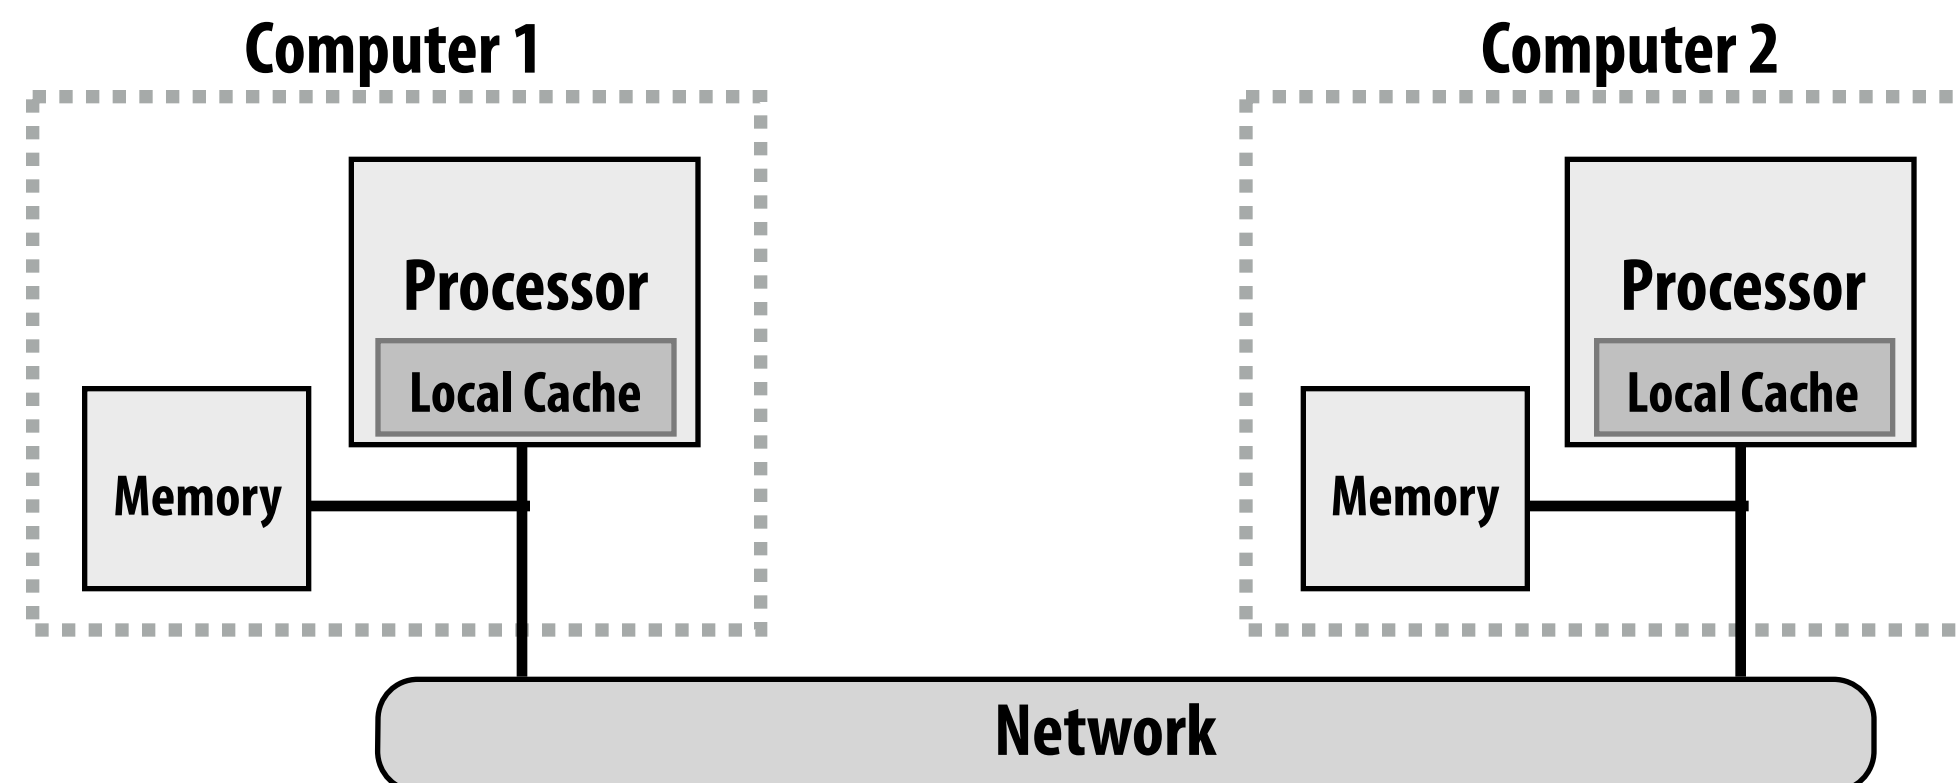
    while (!done) {
        myDiff = 0.0f;
        //
        // perform computation (accumulate locally into myDiff)
        //
        lock(myLock);
        diff[index] += myDiff; // atomically update global diff
        unlock(myLock);
        diff[(index+1) % 3] = 0.0f;
        barrier(myBarrier, NUM_PROCESSORS);
        if (diff[index]/(n*n) < TOLERANCE)
            break;
        index = (index + 1) % 3;
    }
}
```



# Let's think about expressing the grid solver in a message passing model

- No shared address space abstraction (i.e., no shared variables)
- Each thread has its own address space
- Threads communicate & synchronize by sending/receiving messages


**One possible message passing machine configuration: a cluster of two workstations (you could make yourself such a machine in the GHC labs)**







# Review: assignment in a shared address space

- Grid data resided in a single array in shared address space
  - This array was accessible for access by all threads

```
float* A = allocate(n+2, n+2);
```
- Each thread manipulated the region of array it was assigned to process
  - Different assignments may yield different amounts of communication, impacting performance




# Message passing model: each thread operates in its own address space



**This figure: four threads**

**So the grid data is partitioned into allocations residing in each of the four unique address spaces (four per-thread private arrays)**

# Data replication is now required to correctly execute the program



## Example:

After red cell processing is complete, thread 1 and thread 3 send row of data to thread 2 (otherwise thread 2 does not have up-to-date red cell information needed in the subsequent phase)

Commonly used term: "ghost cells"

"Ghost cells" are grid cells replicated from a remote address space. It's common to say that information in ghost cells is "owned" by other threads.

Thread 2 logic:

```
float local_data[(N+2)*rows_per_thread];  
float ghost_row_top[N+2]; // ghost row storage  
float ghost_row_bot[N+2]; // ghost row storage
```

```
int tid = get_thread_id();  
int bytes = sizeof(float) * (N+2);  
recv(ghost_row_top, bytes, tid-1, TOP_MSG_ID);  
recv(ghost_row_bot, bytes, tid+1, BOT_MSG_ID);
```

```
// Thread 2 now has data necessary to perform  
// computation
```

# Message passing solver

Similar structure to shared address space solver, but now communication is explicit in message sends and receives

Send and receive ghost rows to “neighbor threads”

Perform computation

All threads send local my\_diff to thread 0

Thread 0 computes global diff, evaluates termination predicate and sends result back to all other threads

```
int n;
int tid = get_thread_id();
int rows_per_thread = n / get_num_threads();

float* localA = allocate(sizeof(float) * (rows_per_thread+2) * (n+2));

// assume localA is initialized with starting values
// assume MSG_ID_ROW, MSG_ID_DONE, MSG_ID_DIFF are all constants

////////////////////////////////////

void solve() {
    bool done = false;
    while (!done) {

        float my_diff = 0.0f;

        if (tid != 0)
            send(&localA[1,0], sizeof(float)*(N+2), tid-1, MSG_ID_ROW); // send row 0
        if (tid != get_num_threads()-1)
            send(&localA[rows_per_thread-2,0], sizeof(float)*(N+2), tid+1, MSG_ID_ROW);

        if (tid != 0)
            recv(&localA[0,0], sizeof(float)*(N+2), tid-1, MSG_ID_ROW);
        if (tid != get_num_threads()-1)
            recv(&localA[rows_per_thread-1,0], sizeof(float)*(N+2), tid+1, MSG_ID_ROW);

        for (int i=1; i<rows_per_thread-1; i++) {
            for (int j=1; j<n+1; j++) {
                float prev = localA[i,j];
                localA[i,j] = 0.2 * (localA[i-1,j] + localA[i,j] + localA[i+1,j] +
                                   localA[i,j-1] + localA[i,j+1]);
                my_diff += fabs(localA[i,j] - prev);
            }
        }

        if (tid != 0) {
            send(&mydiff, sizeof(float), 0, MSG_ID_DIFF);
            recv(&done, sizeof(bool), 0, MSG_ID_DONE);
        } else {
            float remote_diff;
            for (int i=1; i<get_num_threads()-1; i++) {
                recv(&remote_diff, sizeof(float), i, MSG_ID_DIFF);
                my_diff += remote_diff;
            }
            if (my_diff/(n*n) < TOLERANCE)
                done = true;
            if (int i=1; i<gen_num_threads()-1; i++)
                send(&done, sizeof(bool), i, MSD_ID_DONE);
        }
    }
}
```

# Notes on message passing example

## ■ Computation

- Array indexing is relative to local address space (not global grid coordinates)

## ■ Communication:

- Performed by sending and receiving messages
- Communicate entire rows at a time (not individual elements)


## ■ Synchronization:

- Performed by sending and receiving messages
- Think of how to implement mutual exclusion, barriers, flags using messages

## ■ For convenience: message passing libraries often include higher-level primitives (implemented using send and receive)

```
reduce_add(0, &my_diff, sizeof(float));           // add up all my_diffs, result provided to thread 0
if (pid == 0 && my_diff/(n*n) < TOLERANCE)
    done = true;
broadcast(0, &done, sizeof(bool), MSG_DONE);      // thread 0 sends done to all threads
```

# Variants of send and receive messages



## ■ Synchronous:

- **SEND:** call returns when sender receives acknowledgement that message data resides in address space of receiver
- **RECV:** call returns when data from received message is copied into address space of receiver and acknowledgement sent back to sender

Sender:

Receiver:

Call SEND(foo)

Call RECV(bar)

Copy data from sender's address space buffer 'foo' into network buffer

Send message

Receive message

Copy data into receiver's address space buffer 'bar'

Receive ack

Send ack

SEND() returns

RECV() returns

**As implemented on the prior slide, if our message passing solver uses blocking send/rcv it would deadlock!**

**Why?**

**How can we fix it?**

**(while still using blocking send/rcv)**



# Message passing solver

Similar structure to shared address space solver, but now communication is explicit in message sends and receives

Send and receive ghost rows to “neighbor threads”

Perform computation

All threads send local my\_diff to thread 0

Thread 0 computes global diff, evaluates termination predicate and sends result back to all other threads

```
int n;
int tid = get_thread_id();
int rows_per_thread = n / get_num_threads();

float* localA = allocate(sizeof(float) * (rows_per_thread+2) * (n+2));

// assume localA is initialized with starting values
// assume MSG_ID_ROW, MSG_ID_DONE, MSG_ID_DIFF are all constants

////////////////////////////////////

void solve() {
    bool done = false;
    while (!done) {

        float my_diff = 0.0f;


        if (tid != 0)
            send(&localA[1,0], sizeof(float)*(N+2), tid-1, MSG_ID_ROW);    // send row 0
        if (tid != get_num_threads()-1)
            send(&localA[rows_per_thread-2,0], sizeof(float)*(N+2), tid+1, MSG_ID_ROW);

        if (tid != 0)
            recv(&localA[0,0], sizeof(float)*(N+2), tid-1, MSG_ID_ROW);
        if (tid != get_num_threads()-1)
            recv(&localA[rows_per_thread-1,0], sizeof(float)*(N+2), tid+1, MSG_ID_ROW);

        for (int i=1; i<rows_per_thread-1; i++) {
            for (int j=1; j<n+1; j++) {
                float prev = localA[i,j];
                localA[i,j] = 0.2 * (localA[i-1,j] + localA[i,j] + localA[i+1,j] +
                                   localA[i,j-1] + localA[i,j+1]);
                my_diff += fabs(localA[i,j] - prev);
            }
        }

        if (tid != 0) {
            send(&mydiff, sizeof(float), 0, MSG_ID_DIFF);
            recv(&done, sizeof(bool), 0, MSG_ID_DONE);
        } else {
            float remote_diff;
            for (int i=1; i<get_num_threads()-1; i++) {
                recv(&remote_diff, sizeof(float), i, MSG_ID_DIFF);
                my_diff += remote_diff;
            }
            if (my_diff/(n*n) < TOLERANCE)
                done = true;
            if (int i=1; i<gen_num_threads()-1; i++)
                send(&done, sizeof(bool), i, MSD_ID_DONE);
        }
    }
}
```

# Variants of send and receive messages



## ■ Blocking async:

- **SEND:** call copies data from address space into system buffers, then returns
  - Does not guarantee message has been received (or even sent)
- **RECV:** call returns when data copied into address space, but no ack sent

Sender:

Receiver:

Call SEND(foo)

Call RECV(bar)

Copy data from sender's address space buffer 'foo' into network buffer

SEND(foo) returns, calling thread continues execution

**Send message**


Receive message

Copy data into receiver's address space buffer

RECV(bar) returns

**RED TEXT = executes concurrently with sender's application thread**

# Variants of send and receive messages



## ■ Non-blocking asynchronous: (“non-blocking”)

- **SEND:** call returns immediately. Buffer provided to SEND cannot be modified by calling thread since message processing occurs concurrently with thread execution
- **RECV:** call posts intent to receive, returns immediately.
- Use **SENDPROBE**, **RECVPROBE** to determine actual send/receipt status

Sender:

Receiver:

Call SEND(foo)

SEND(foo) returns handle h1

Call RECV(bar)

RECV(bar) returns handle h2

**Copy data from 'foo' into network buffer**

**Send message**

**Receive message**

**Messaging library copies data into 'bar'**

Call SENDPROBE(h1) // if message sent, now safe for thread to modify 'foo'


Call RECVPROBE(h2)

// if received, now safe for thread

// to access 'bar'

**RED TEXT = executes concurrently with application thread**

# Variants of send and receive messages



**The variants of send/recv provide different levels of programming complexity / opportunity to optimize performance**

# Solver implementation in THREE programming models

## 1. Data-parallel model

- Synchronization:
  - `forall` loop iterations are independent (can be parallelized)
  - Implicit barrier at end of outer `forall` loop body
- Communication
  - Implicit in loads and stores (like shared address space)
  - Special built-in primitives: e.g., `reduce`

## 2. Shared address space model

- Synchronization:
  - Locks used to ensure mutual exclusion
  - Barriers used to express coarse dependencies (e.g., between phases of computation)
- Communication
  - Implicit in loads/stores to shared variables

## 3. Message passing model

- Synchronization:
  - Implemented via messages
  - Mutual exclusion exists by default: no shared data structures
- Communication:
  - Explicit communication via `send/recv` needed for parallel program correctness
  - Bulk communication for efficiency: e.g., communicate entire rows, not single elements
  - Several variants of `send/recv`, each has different semantics

# **Optimizing parallel program performance**

**( how to be l33t )**

# Programming for high performance

- **Optimizing the performance of parallel programs is an iterative process of refining choices for decomposition, assignment, and orchestration...**
- **Key goals (that are at odds with each other)**
  - **Balance workload onto available execution resources**
  - **Reduce communication (to avoid stalls)**
  - **Reduce extra work (overhead) performed to increase parallelism, manage assignment, etc.**
- **We are going to talk about a rich space of techniques**


**TIP #1: Always implement the simplest solution first, then measure/analyze performance to determine if you need to do better.**

**“My solution scales” = your code scales as much as you need it to  
(if you anticipate only running low core count machines, it may be unnecessary to implement a complex approach that creates and hundreds or thousands of pieces of independent work)**



# Balancing the workload

**Ideally: all processors are computing all the time during program execution  
(they are computing simultaneously, and they finish their portion of the work at the same time)**




**Recall Amdahl's Law:  
Only small amount of load imbalance can  
significantly bound maximum speedup**

**P4 does 20% more work → P4 takes 20% longer to complete  
→ 20% of parallel program runtime is  
essentially serial execution**

**(work in serialized section here is about 5% of a sequential  
implementation execution time:  $S=.05$  in Amdahl's law equation)**

# Static assignment


- Assignment of work to threads is pre-determined
  - Not necessarily determined at compile-time (assignment algorithm may depend on runtime parameters such as input data size, number of threads, etc.)
- Recall solver example: assign equal number of grid cells (work) to each thread (worker)
  - We discussed blocked and interleaved static assignments of work to workers



- Good properties of static assignment: simple, essentially zero runtime overhead (in this example: extra work to implement assignment is a little bit of indexing math)

# Static assignment

- When is static assignment applicable?
- When the cost (execution time) of work and the amount of work is predictable (so the programmer can work out assignment in advance)
- Simplest example: it is known up front that all work has the same cost



In the example above:


There are 12 tasks, and it is known each have the same cost.

Statically assign three tasks to each of the four processors.

# Static assignment

## ■ When is static assignment applicable?


- Example 2: predictable, but not all jobs have same cost (see example below)
- Example 3: Statistics about execution time are known (e.g., same cost on average)



Jobs have unequal, but known cost: assign to processors to ensure overall good load balance

# “Semi-static” assignment

- Cost of work is predictable for near-term future
  - Recent past good predictor of near future
- Periodically profile application and re-adjust assignment
  - Assignment is static during interval between re-adjustment



**Particle simulation:**

**Redistribute particles as they move over course of simulation (if motion is slow, redistribution need not occur often)**




Image credit: <http://typhon.sourceforge.net/spip/spip.php?article22>

**Adaptive mesh:**

**Mesh is changed as object moves or flow over object changes, but changes occur slowly (color indicates assignment of parts of mesh to processors)**

# Dynamic assignment

- Assignment is determined at runtime to ensure a well distributed load.  
(The execution time of tasks, or the total number of tasks, is unpredictable.)

## Sequential program (independent loop iterations)

```
int N = 1024;
int* x = new int[N];
bool* prime = new bool[N];

// initialize elements of x

for (int i=0; i<N; i++)
{
    // unknown execution time
    is_prime[i] = test_primalty(x[i]);
}
```


## Parallel program (SPMD execution of multiple threads, shared address space model)

```
LOCK counter_lock;
int counter = 0;    // shared variable (assume
                  // initialization to 0)

int N = 1024;
int* x = new int[N];
bool* is_prime = new bool[N];

// initialize elements of x


while (1) {
    int i;
    lock(counter_lock);
    i = counter++;
    unlock(counter_lock);
    if (i >= N)
        break;
    is_prime[i] = test_primalty(x[i]);
}
```




atomic\_incr(counter);

# Dynamic assignment using work queues


**Sub-problems**  
(a.k.a. “tasks”, “work”)



**Shared work queue: a list of work to do**  
(for now, let's assume each piece of work is independent)



**Worker threads:**  
Pull data from shared work queue  
Push new work to queue as it is created





# What constitutes a piece of work?

## ■ What is a potential problem with this implementation?

```
LOCK counter_lock;
int counter = 0;    // shared variable (assume
                   // initialization to 0)

const int N = 1024;
float* x = new float[N];
bool* prime = new bool[N];

// initialize elements of x

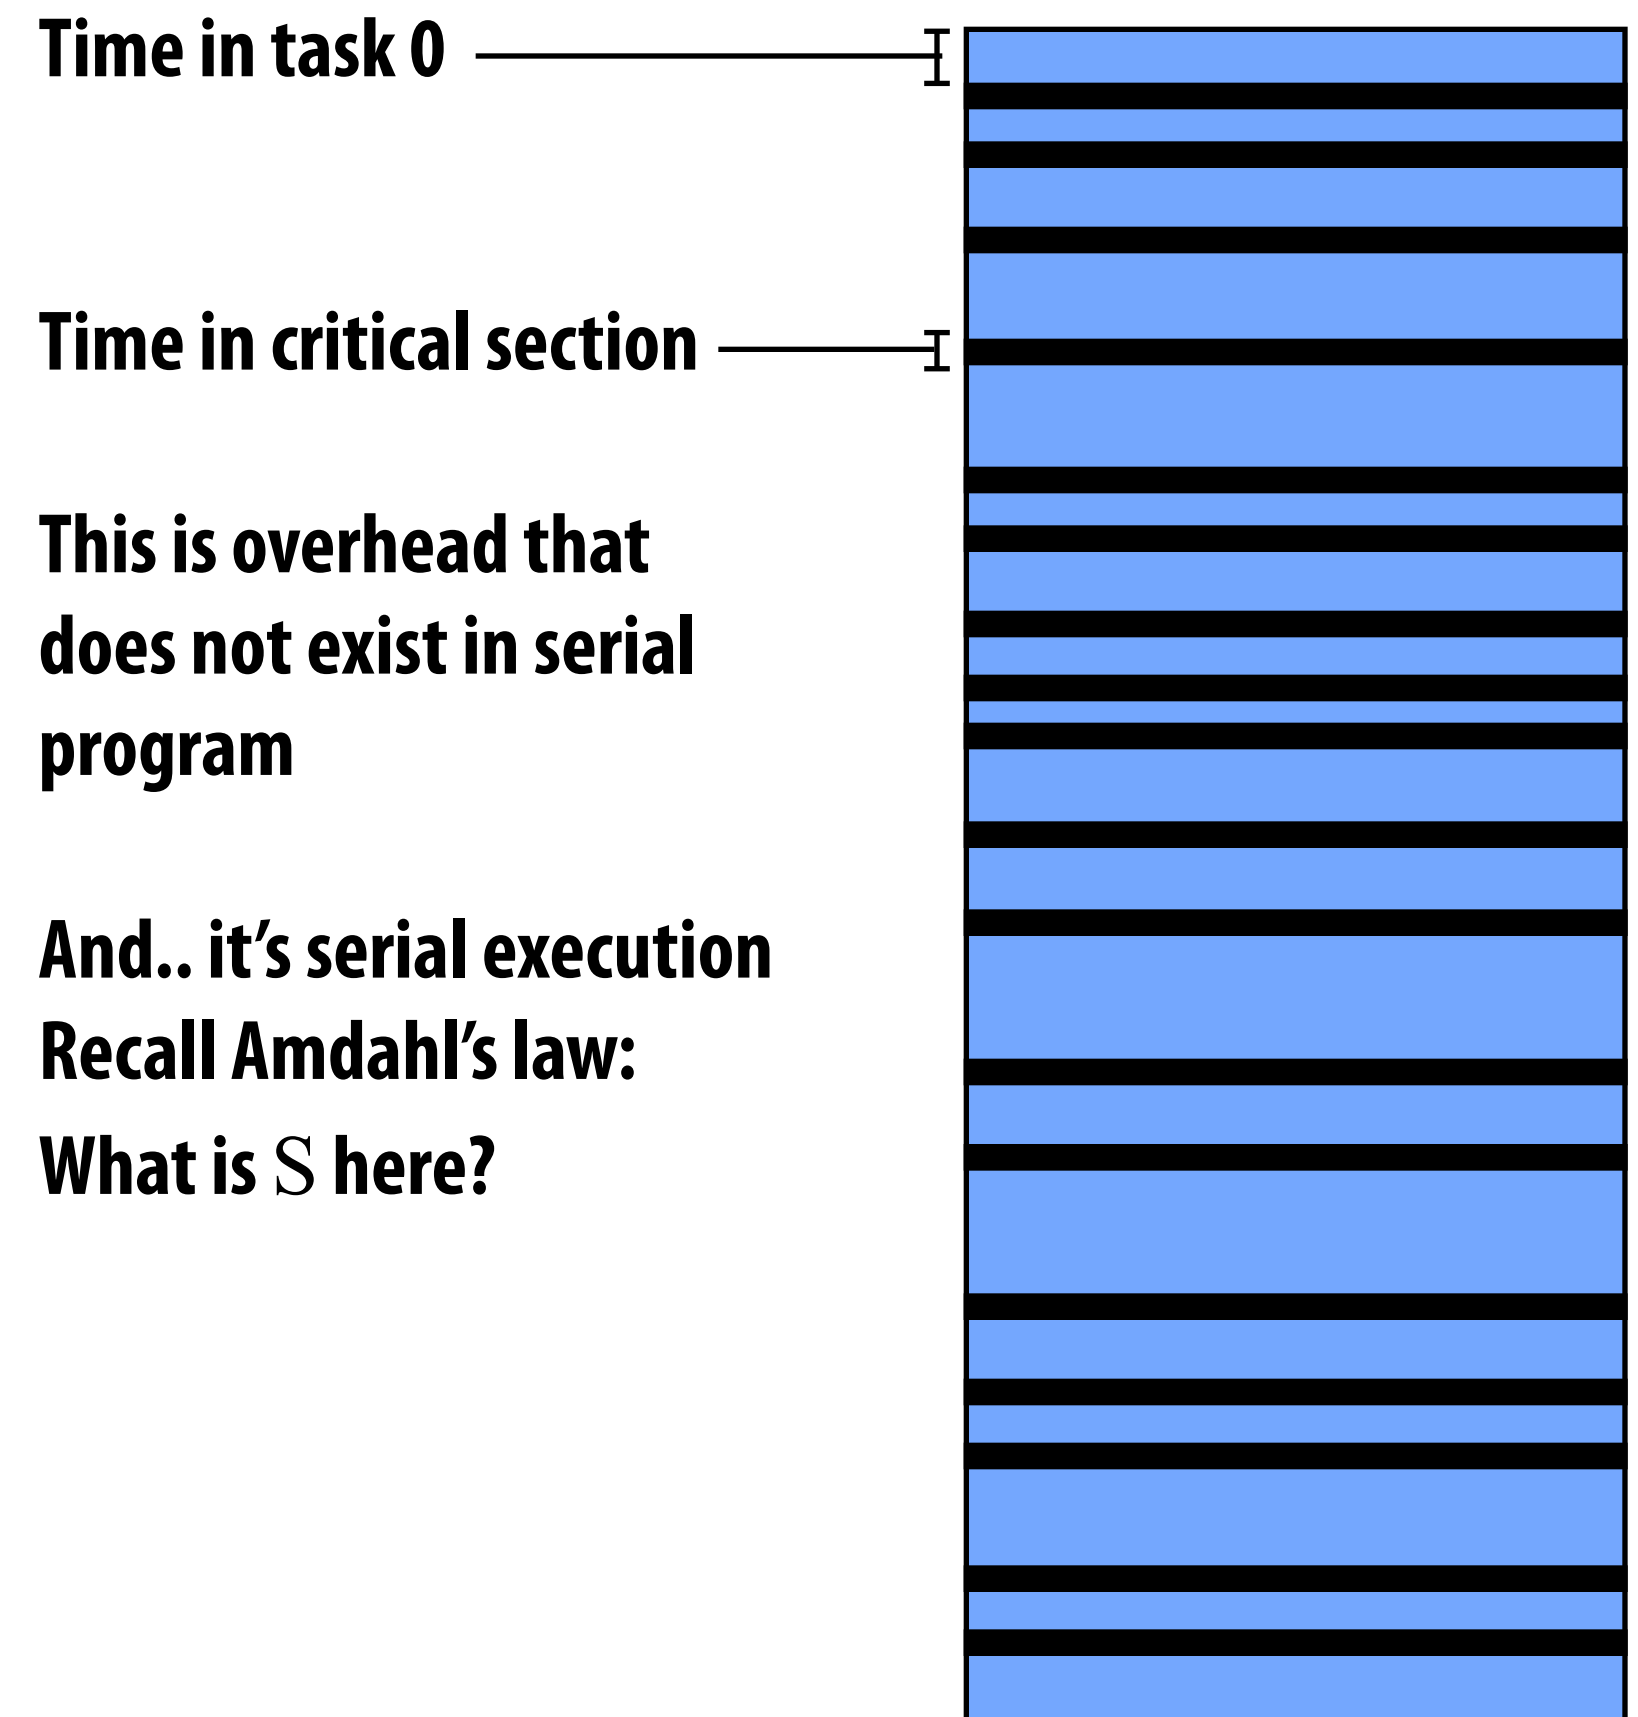
while (1) {
    int i;
    lock(counter_lock);
    i = counter++;
    unlock(counter_lock);
    if (i >= N)
        break;
    is_prime[i] = test_primalty(x[i]);
}
```

**Fine granularity partitioning:**

**Here: 1 “task” = 1 element**

**Likely good workload balance (many small tasks)**

**Potential for high synchronization cost  
(serialization at critical section)**



## So... IS this a problem?




# Increasing task granularity

```
LOCK counter_lock;
int counter = 0;    // shared variable (assume
                   // initialization to 0)

const int N = 1024;
const int GRANULARITY = 10;
float* x = new float[N];
bool* prime = new bool[N];

// initialize elements of x

while (1) {
    int i;
    lock(counter_lock);
    i = counter;
    counter += GRANULARITY;
    unlock(counter_lock);
    if (i >= N)
        break;
    int end = min(i + GRANULARITY, N);
    for (int j=i; j<end; j++)
        is_prime[i] = test_primalty(x[i]);
}
```



**Coarse granularity partitioning:**

**1 "task" = 10 elements**

**Decreased synchronization cost**  
**(Critical section entered 10 times less)**

## So... have we done better?

# Rule of thumb


- **Useful to have many more tasks\* than processors**  
(many small tasks enables good workload balance via dynamic assignment)
  - Motivates small granularity tasks
- **But want as few tasks as possible to minimize overhead of managing the assignment**
  - Motivates large granularity tasks
- **Ideal granularity depends on many factors**  
(Common theme in this course: must know your workload, and your machine)

\* I had to pick a term. Here I'm using "task"  
generally: it's a piece of work, a sub-problem, etc.

# Smarter task scheduling


Consider dynamic scheduling via a shared work queue

What happens if the system assigns these tasks to workers in left-to-right order?



# Smarter task scheduling

What happens if scheduler runs the long task last? Potential for load imbalance!




**One possible solution to imbalance problem:**

**Divide work into a larger number of smaller tasks**

- Hopefully “long pole” gets shorter relative to overall execution time
- May increase synchronization overhead
- May not be possible (perhaps long task is fundamentally sequential)

# Smarter task scheduling

Schedule long task first to reduce “slop” at end of computation



**Another solution: smarter scheduling**

**Schedule long tasks first**

- Thread performing long task performs fewer overall tasks, but approximately the same amount of work as the other threads.
- Requires some knowledge of workload (some predictability of cost)

# Decreasing synchronization overhead

## ■ Distributed work queues

- Replicate data to remove synchronization on single work queue

Subproblems

(a.k.a. “tasks”, “work to do”)

Set of work queues

(In general, one per worker thread)

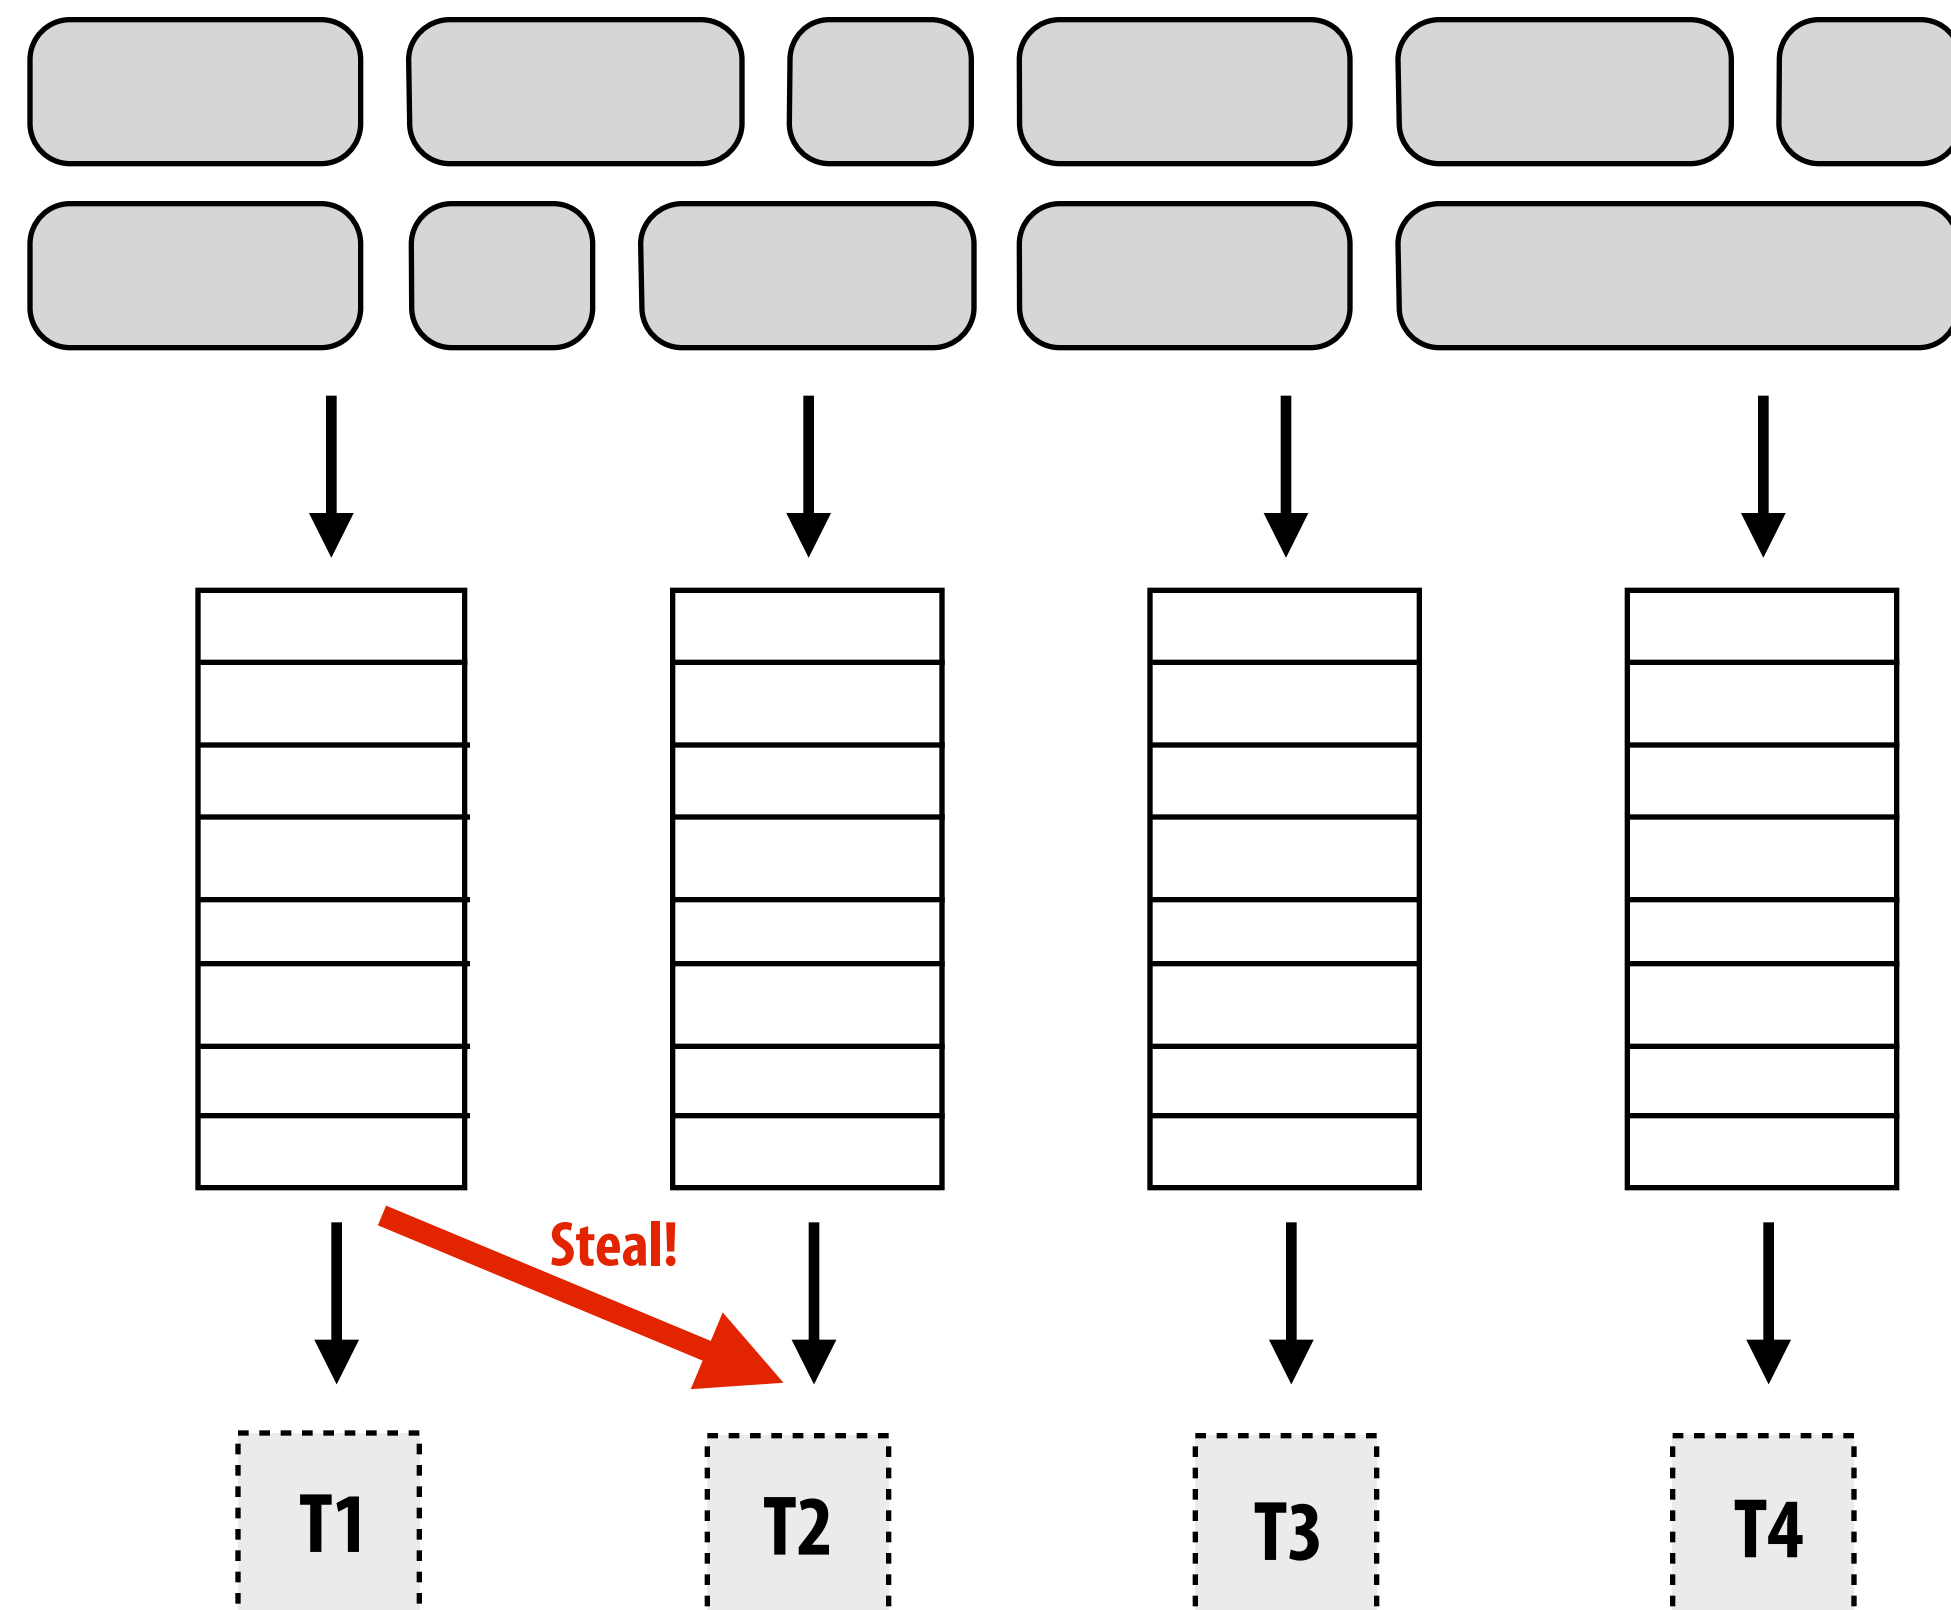
Worker threads:

Pull data from OWN work queue

Push new work to OWN work to queue

When local work queue is empty...

STEAL work from another work queue



# Distributed work queues

## ■ Costly synchronization/communication occurs during stealing


- But not every time a thread takes on new work
- Stealing occurs only when necessary to ensure good load balance

## ■ Leads to increased locality

- Common case: threads work on tasks they create (producer-consumer locality)

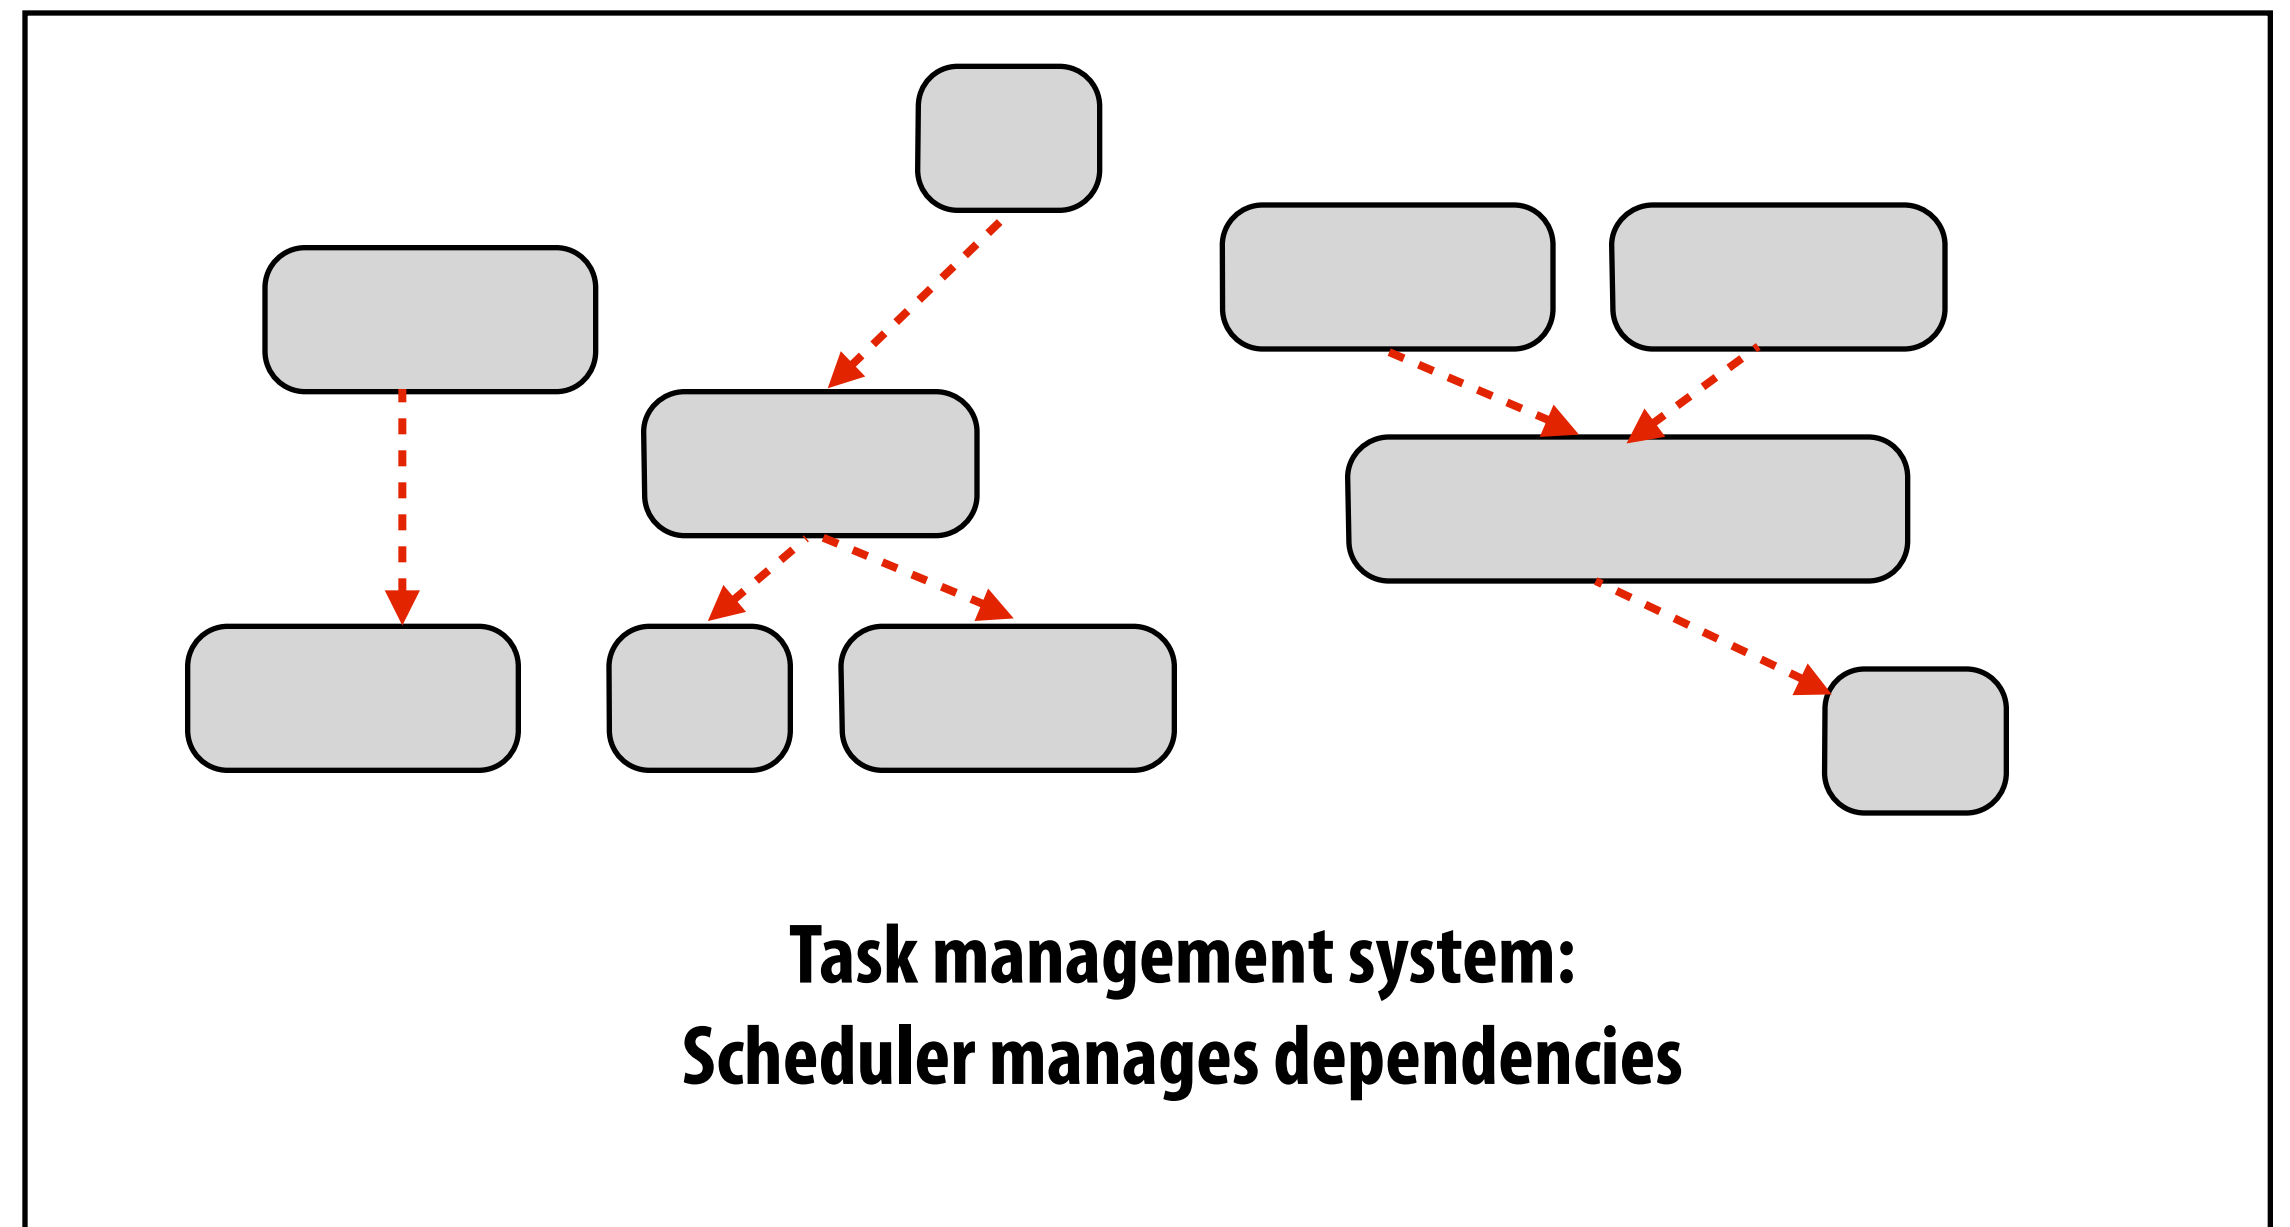
## ■ Implementation challenges

- Who to steal from?
- How much to steal?
- How to detect program termination?
- Ensuring local queue access is fast (while preserving mutual exclusion)




# Work in task queues need not be independent

-----> = application specified dependency



**A task is not removed from queue and assigned to worker thread until all task dependencies are satisfied**

**Workers can submit new tasks (with optional explicit dependencies) to task system**



```
foo_handle = enqueue_task(foo);           // enqueue task foo (independent of all prior tasks, run at any time)
bar_handle = enqueue_task(bar, foo_handle); // enqueue task bar, cannot run until foo is complete.
```



# Summary

- **Challenge: achieving good workload balance**
  - Want all processors working at all the time (otherwise, resources are idle!)
  - But want low cost solution for achieving this balance
    - Minimize computational overhead (e.g., scheduling/assignment logic)
    - Minimize synchronization costs
- **Static assignment vs. dynamic assignment**
  - Really, it is not an either/or decision, there's a continuum of choices
  - Use up-front knowledge about workload as much as possible to reduce load imbalance and task management/synchronization costs (in the limit, if the system knows everything, use fully static assignment)
- **Issues discussed today span decomposition, assignment, and orchestration**

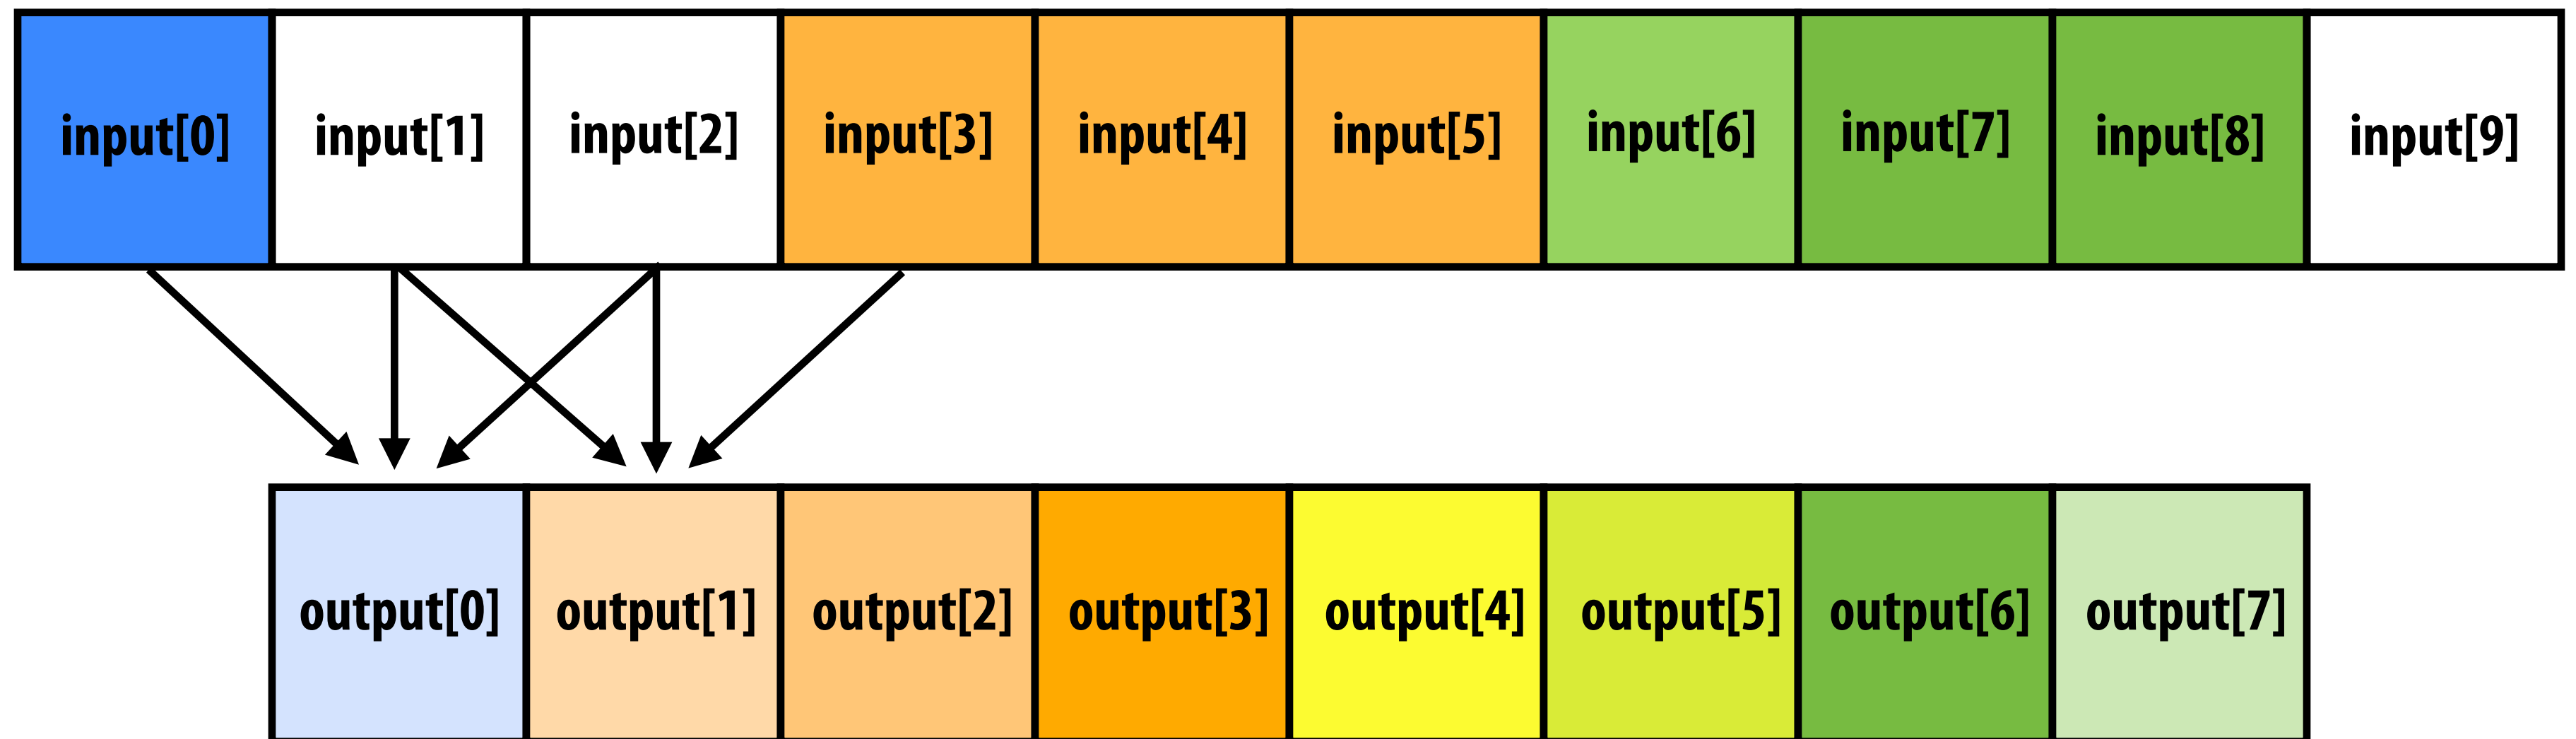
# **CUDA and GPU programming self check**

**(this is also an abstraction vs. implementation self check)**

**(and a work scheduling self check)**


**(and a parallel architecture self check)**

# Recall the 1D convolution example



`output[i] = (input[i] + input[i+1] + input[i+2]) / 3.f;`

# Recalling the 1D convolution in CUDA example



```
__global__ void convolve_1d(int N, float* input, float* output) {  
  
    int index = blockIdx.x * blockDim.x + threadIdx.x; // thread local variable  
  
    float result = 0.0f; // thread-local variable  
    for (int i=0; i<3; i++)  
        result += input[index + i];  
  
    output[index] = result / 3.f;  
}
```

**Simplest possible CUDA kernel for this computation.**


**One CUDA thread per output element.**

**Each thread independently loads input elements it requires.**

**Notice: CUDA threads in thread block do not cooperate.**

**(no logic based on block size in the kernel\*)**

# Implementation using per-block shared memory



```
#define THREADS_PER_BLK 128

__global__ void convolve_1d_shared(int N, float* input, float* output) {

    __shared__ float support[THREADS_PER_BLK+2]; // per block allocation
    int index = blockIdx.x * blockDim.x + threadIdx.x; // thread local variable

    support[threadIdx.x] = input[index];
    if (threadIdx.x < 2) {
        support[THREADS_PER_BLK + threadIdx.x] = input[index+THREADS_PER_BLK];
    }

    __syncthreads();

    float result = 0.0f; // thread-local variable
    for (int i=0; i<3; i++)
        result += support[threadIdx.x + i];

    output[index] = result / 3.f;
}
```

**All threads cooperatively load block's support region from global memory into shared memory**

**total of 130 load instructions instead of 3 \* 128 load instructions**

# Running the kernel

**Kernel's execution requirements:**

**Each thread block must execute 128 CUDA threads**

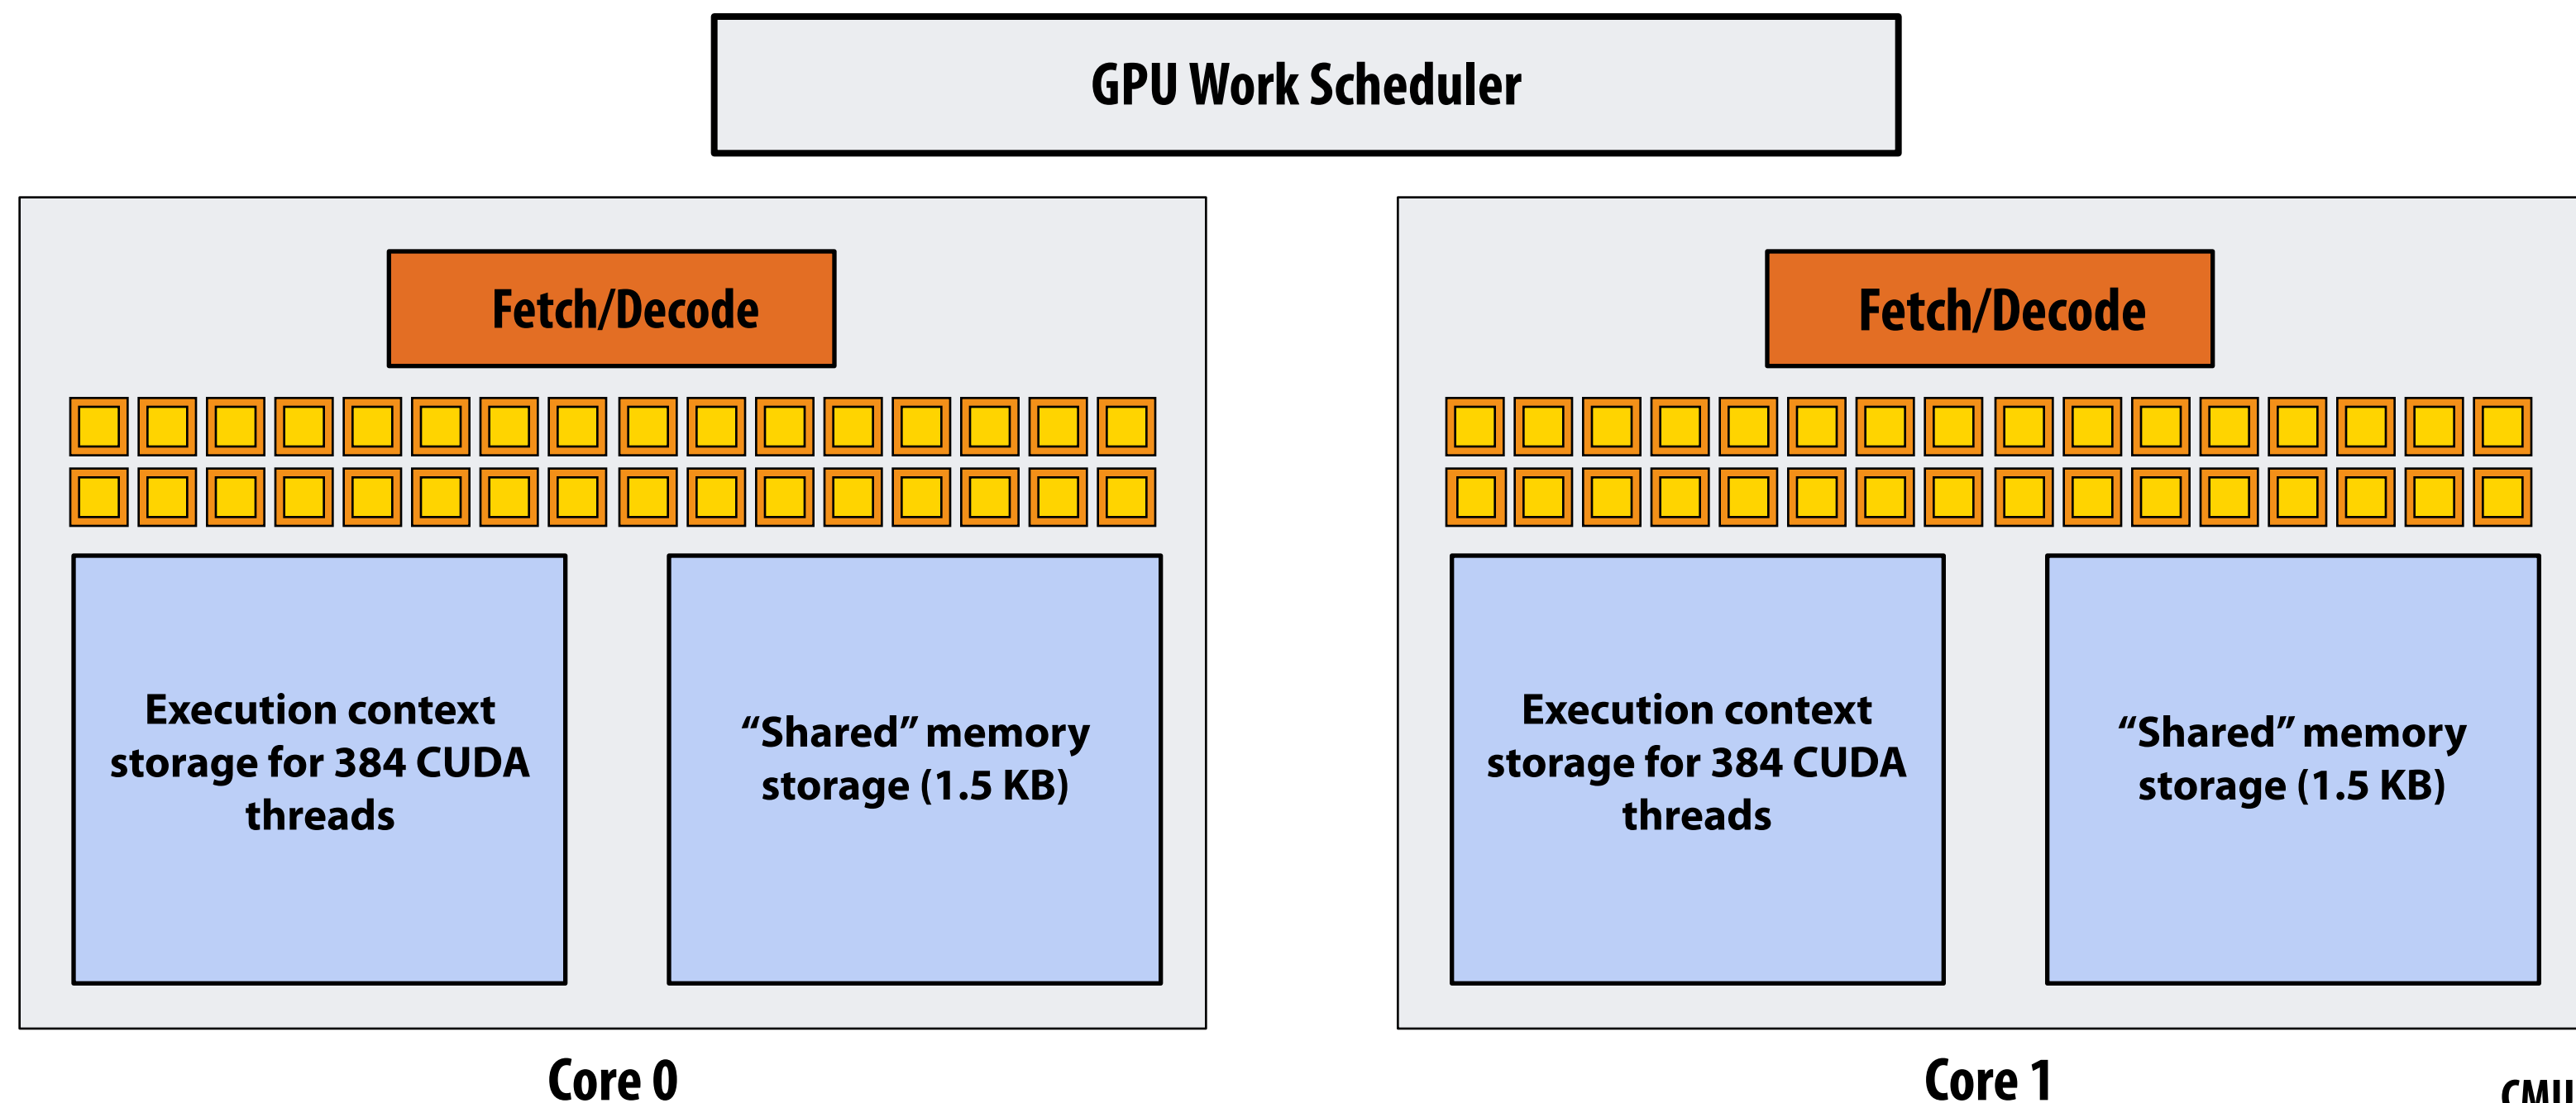
**Each thread block must allocate  $130 * \text{sizeof(float)} = 520$  bytes of shared memory**

**Let's assume array size N is very large, so the host-side kernel launch generates thousands of thread blocks.**

```
#define THREADS_PER_BLK 128
```

```
convolve_1d_shared<<<N/THREADS_PER_BLK, THREADS_PER_BLK>>>(N, input_array, output_array);
```

**Let's run this program on the fictitious two-core GPU below.**



# Running the CUDA kernel

Kernel's execution requirements:

Each thread block must execute 128 CUDA threads

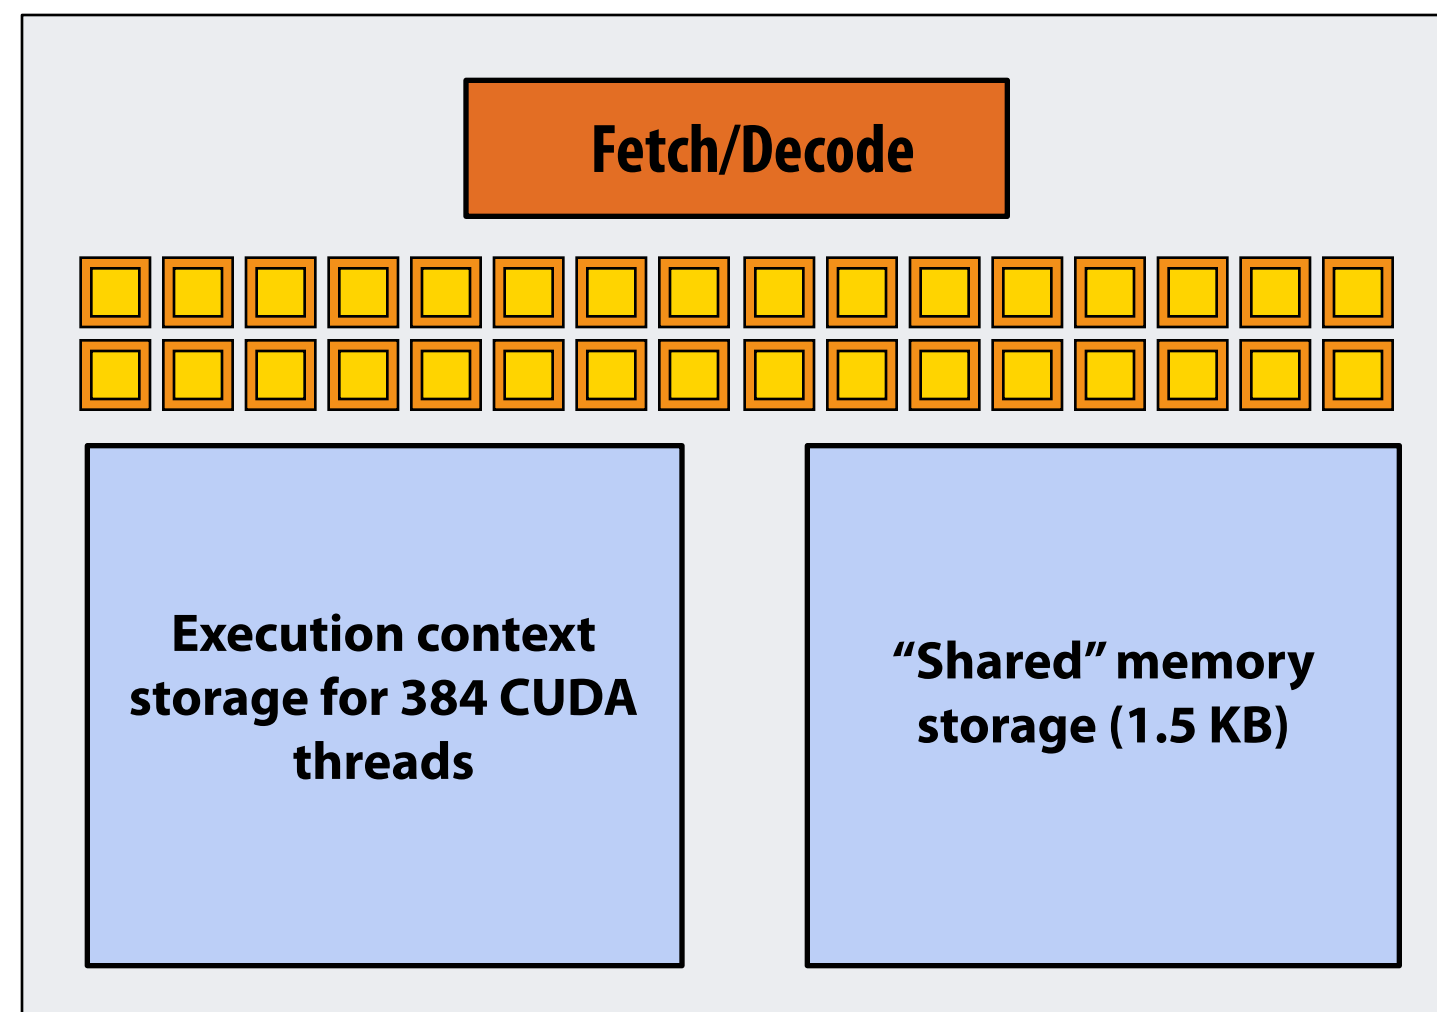
Each thread block must allocate  $130 * \text{sizeof(float)} = 520$  bytes of shared memory

Step 1: host sends CUDA device (GPU) a command ("execute this kernel")

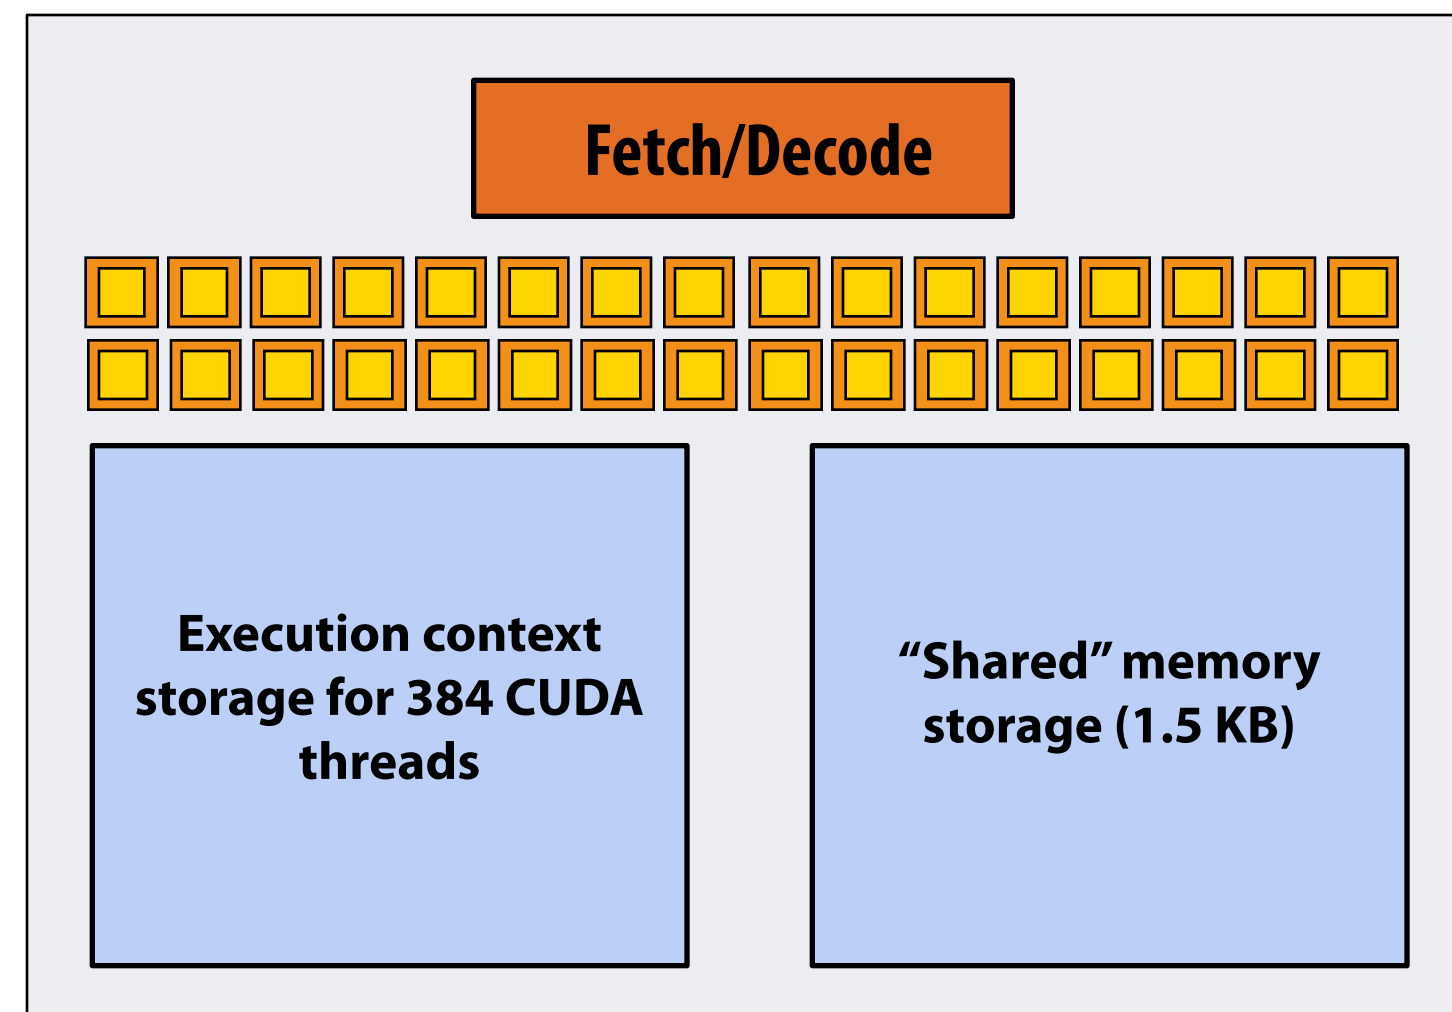


EXECUTE: `convolve_1d_shared`  
ARGS: `N, input_array, output_array`  
NUM\_BLOCKS: `1000`

GPU Work Scheduler



Core 0



Core 1

# Running the CUDA kernel

Kernel's execution requirements:

Each thread block must execute 128 CUDA threads

Each thread block must allocate  $130 * \text{sizeof(float)} = 520$  bytes of shared memory

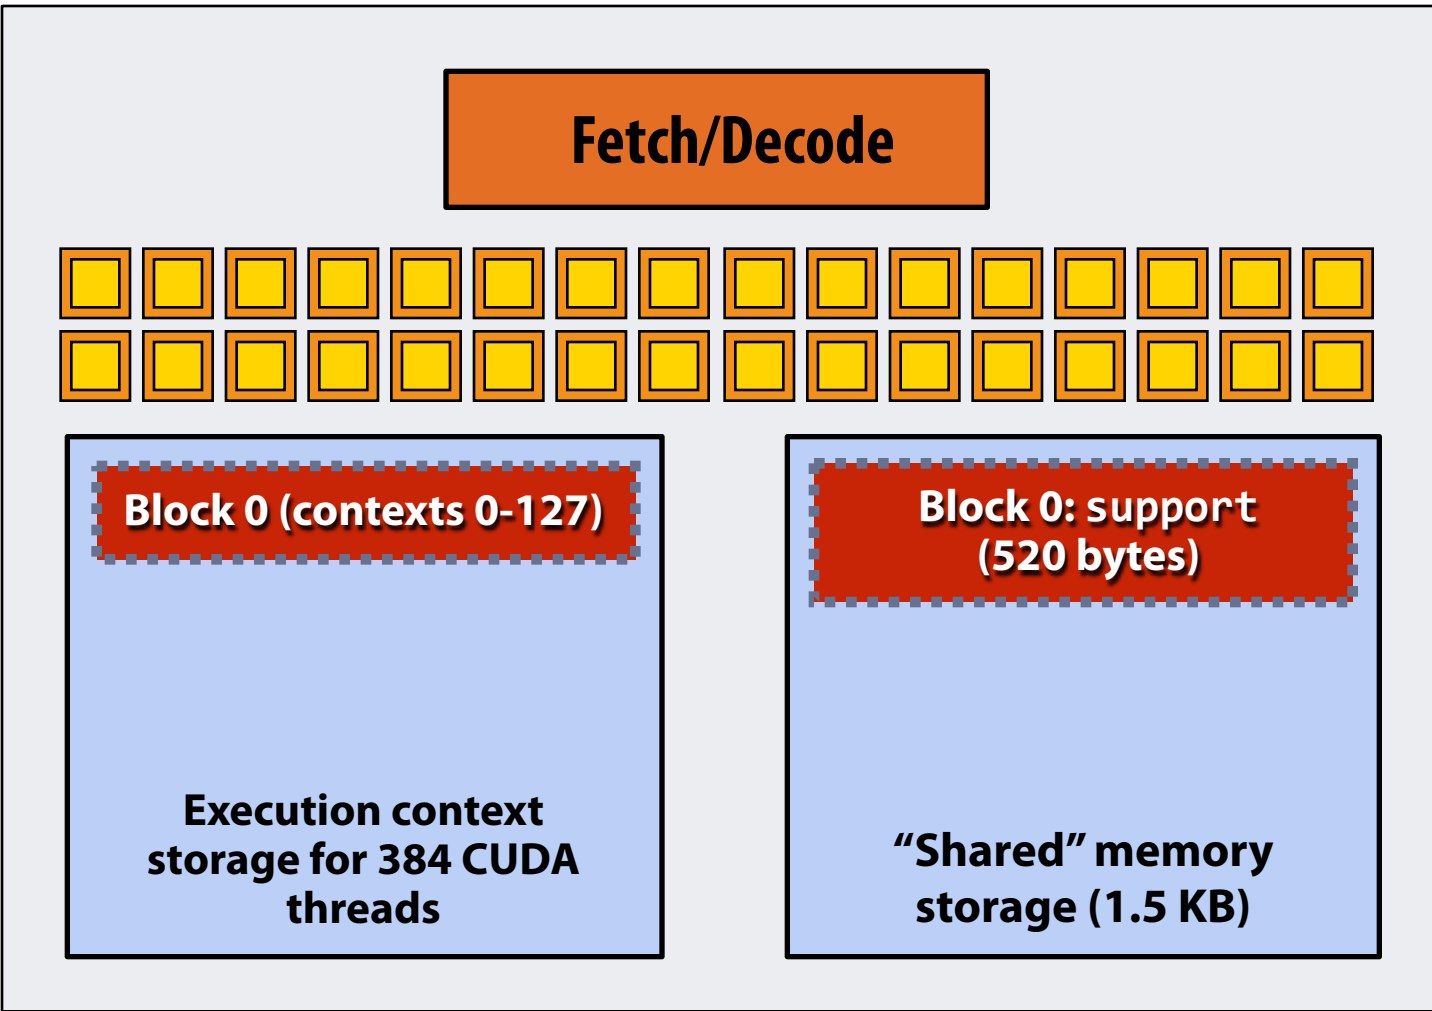
Step 2: scheduler maps block 0 to core 0 (reserves execution contexts and shared storage)



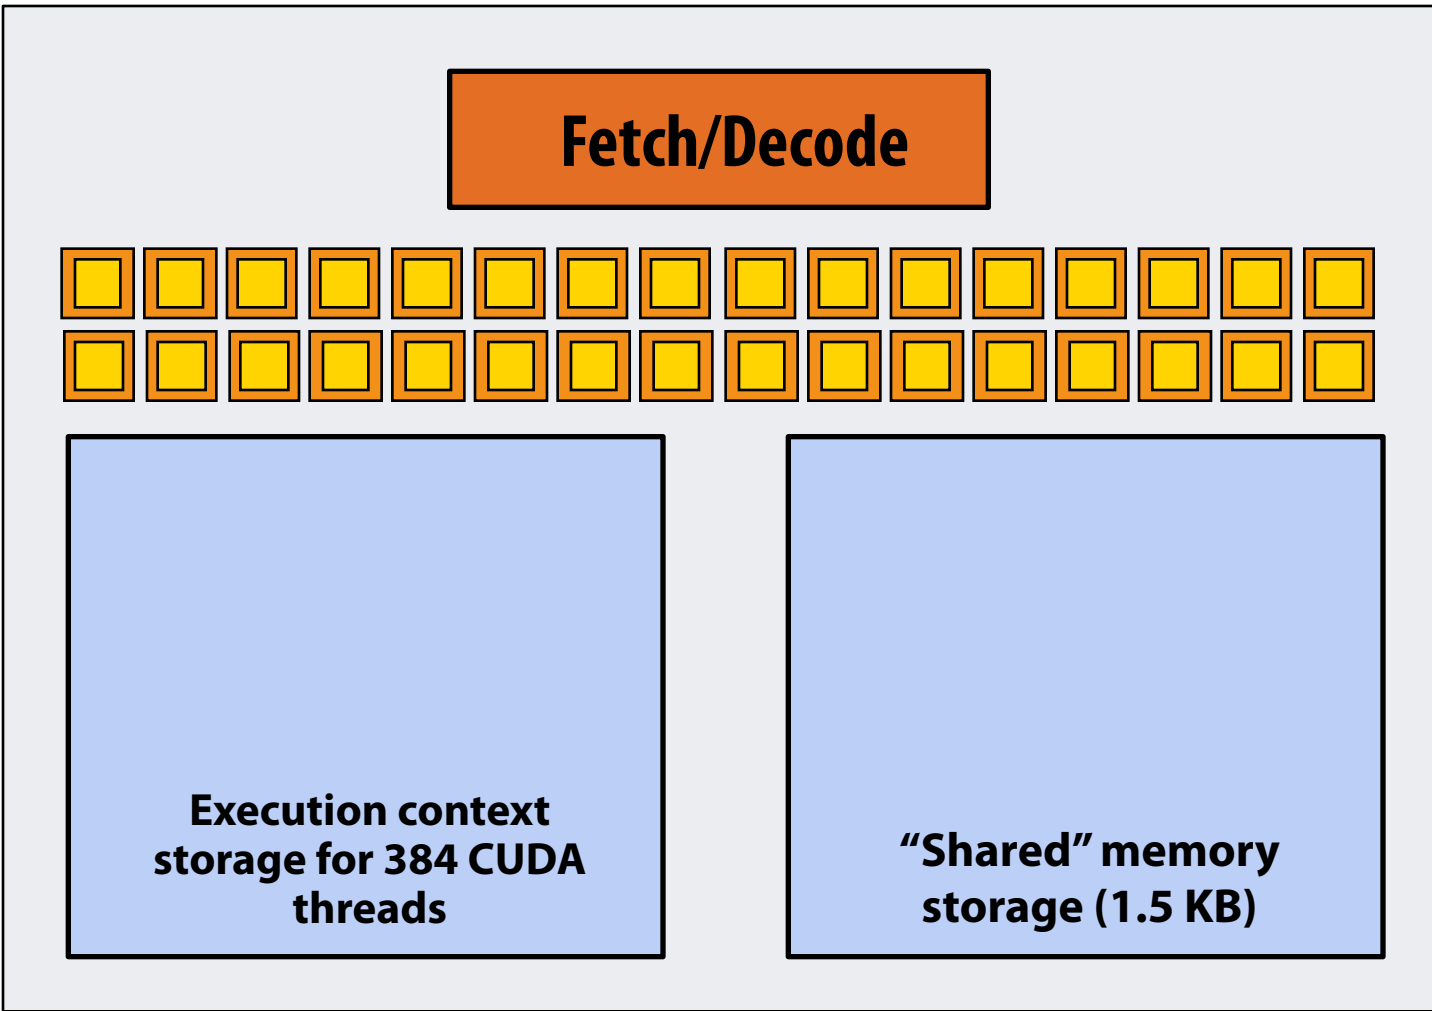
EXECUTE: convolve\_1d\_shared  
ARGS: N, input\_array, output\_array  
NUM\_BLOCKS: 1000

NEXT = 1  
TOTAL = 1000

GPU Work Scheduler



Core 0



Core 1



# Running the CUDA kernel

Kernel's execution requirements:

Each thread block must execute 128 CUDA threads

Each thread block must allocate  $130 * \text{sizeof(float)} = 520$  bytes of shared memory

Step 3: scheduler continues to map blocks to available execution contexts  
(interleaved mapping shown)

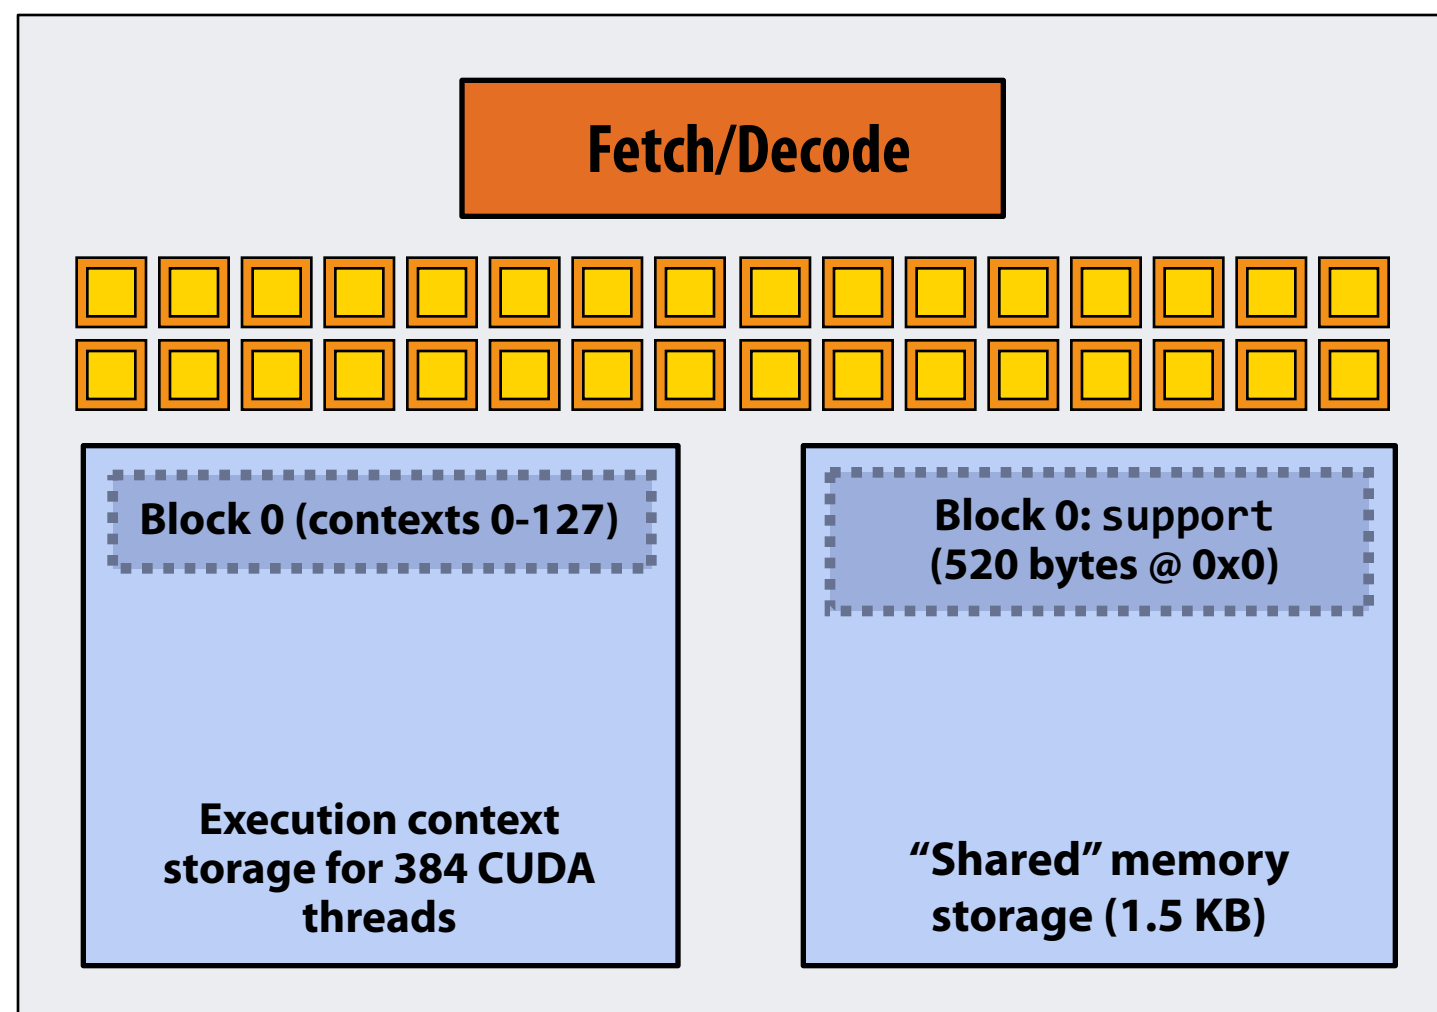


EXECUTE: `convolve_1d_shared`  
ARGS: `N, input_array, output_array`  
NUM\_BLOCKS: 1000

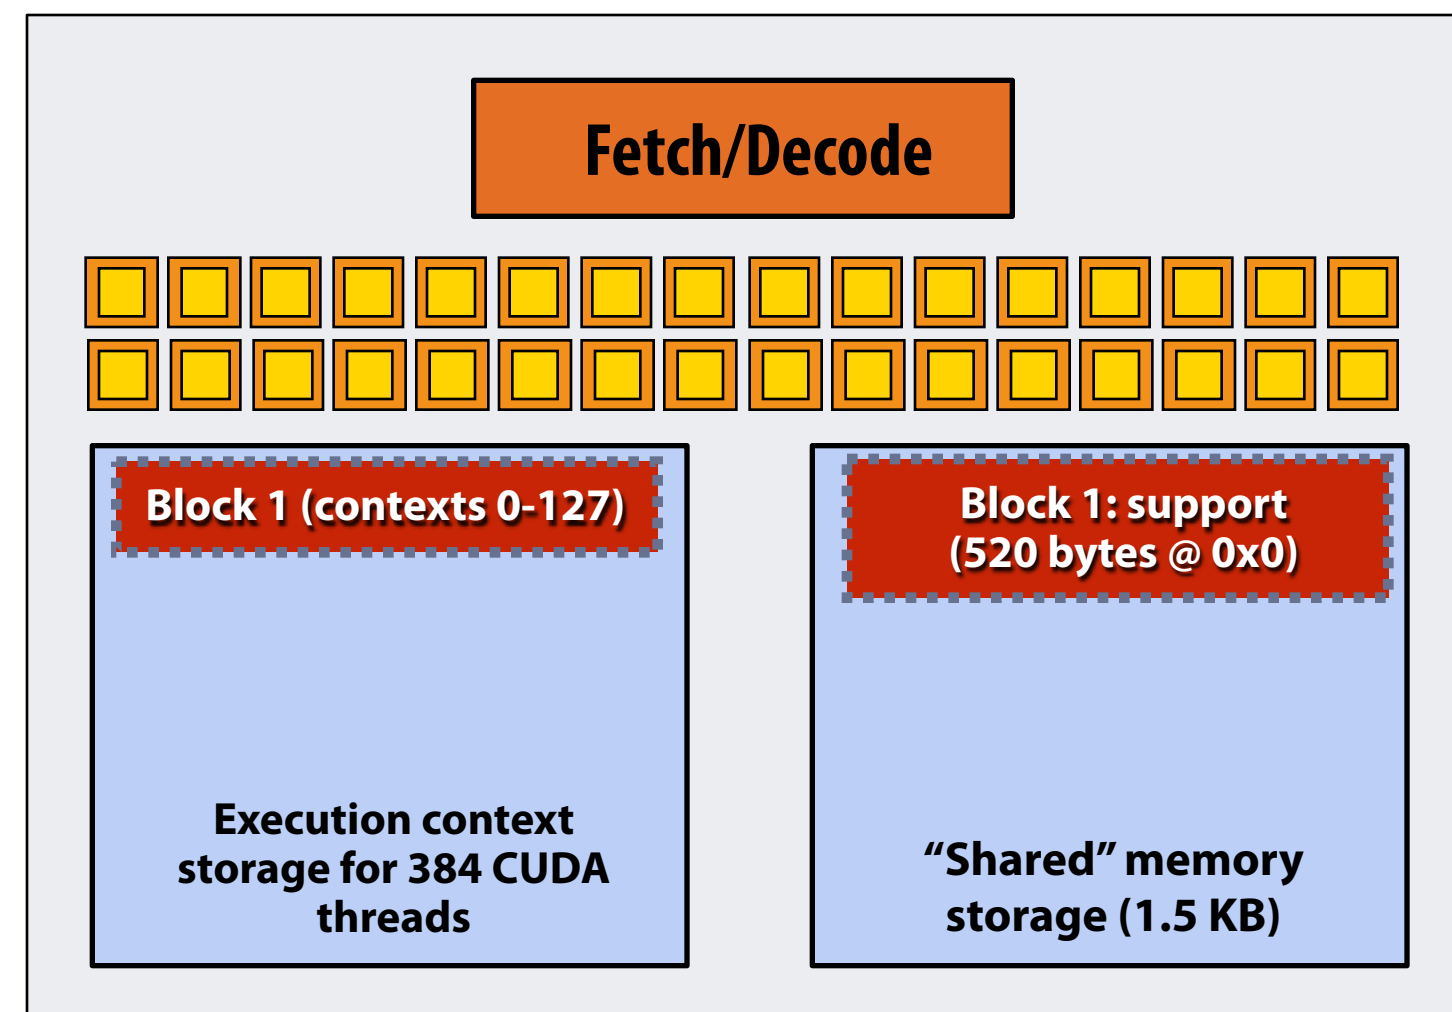
NEXT = 2

GPU Work Scheduler

TOTAL = 1000



Core 0



Core 1

# Running the CUDA kernel

Kernel's execution requirements:

Each thread block must execute 128 CUDA threads

Each thread block must allocate  $130 * \text{sizeof(float)} = 520$  bytes of shared memory

Step 3: scheduler continues to map blocks to available execution contexts  
(interleaved mapping shown)

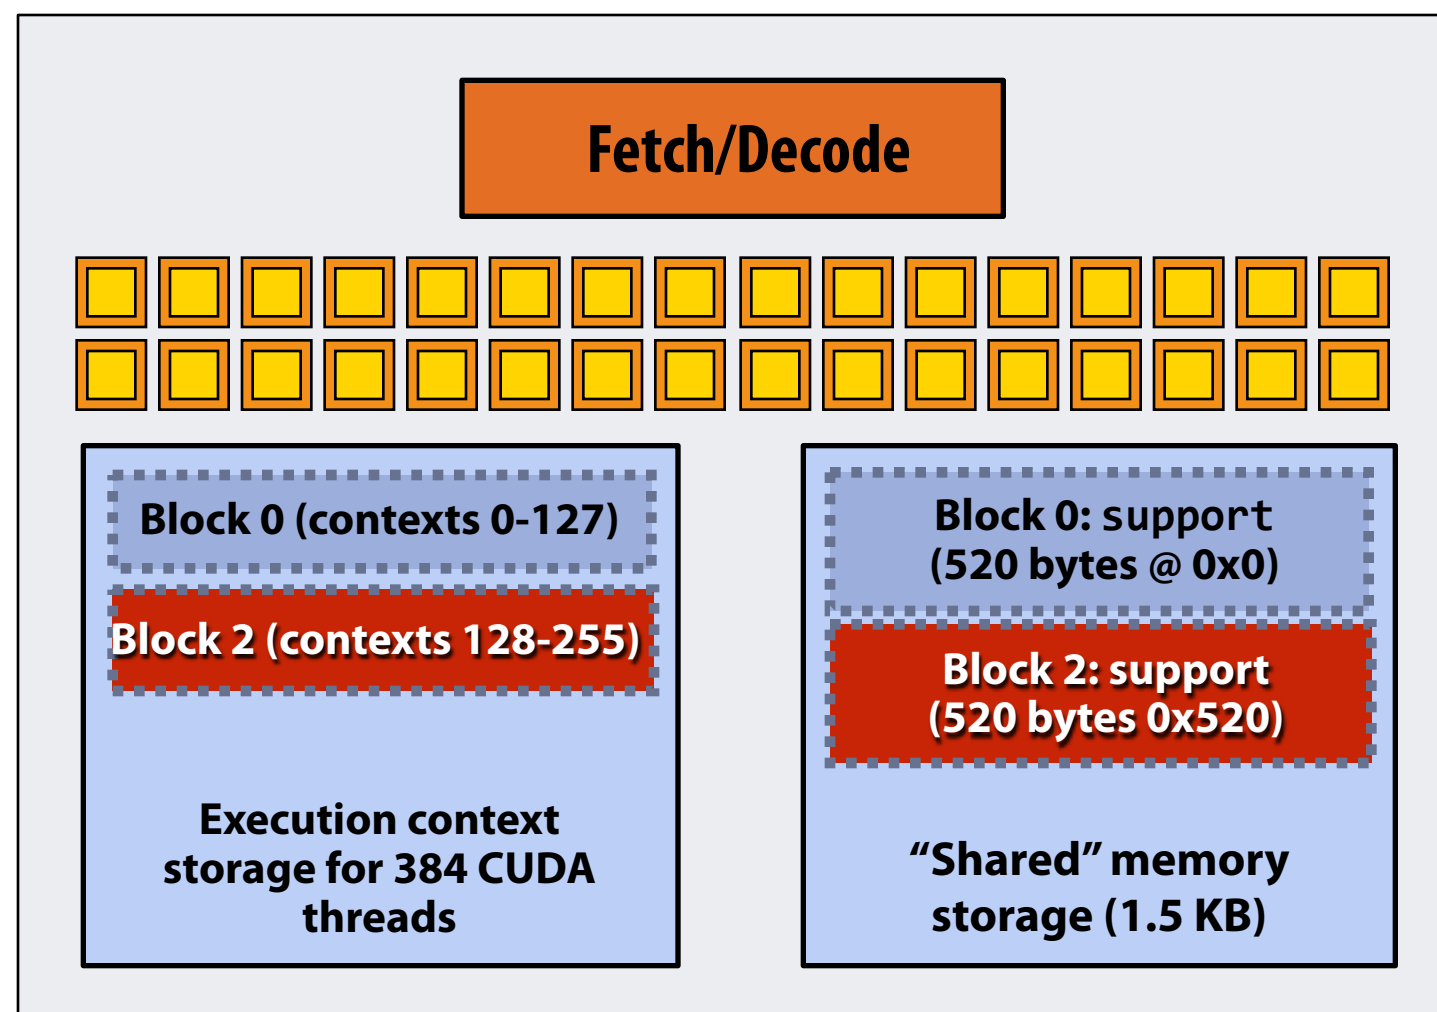


EXECUTE: `convolve_1d_shared`  
ARGS: `N, input_array, output_array`  
NUM\_BLOCKS: 1000

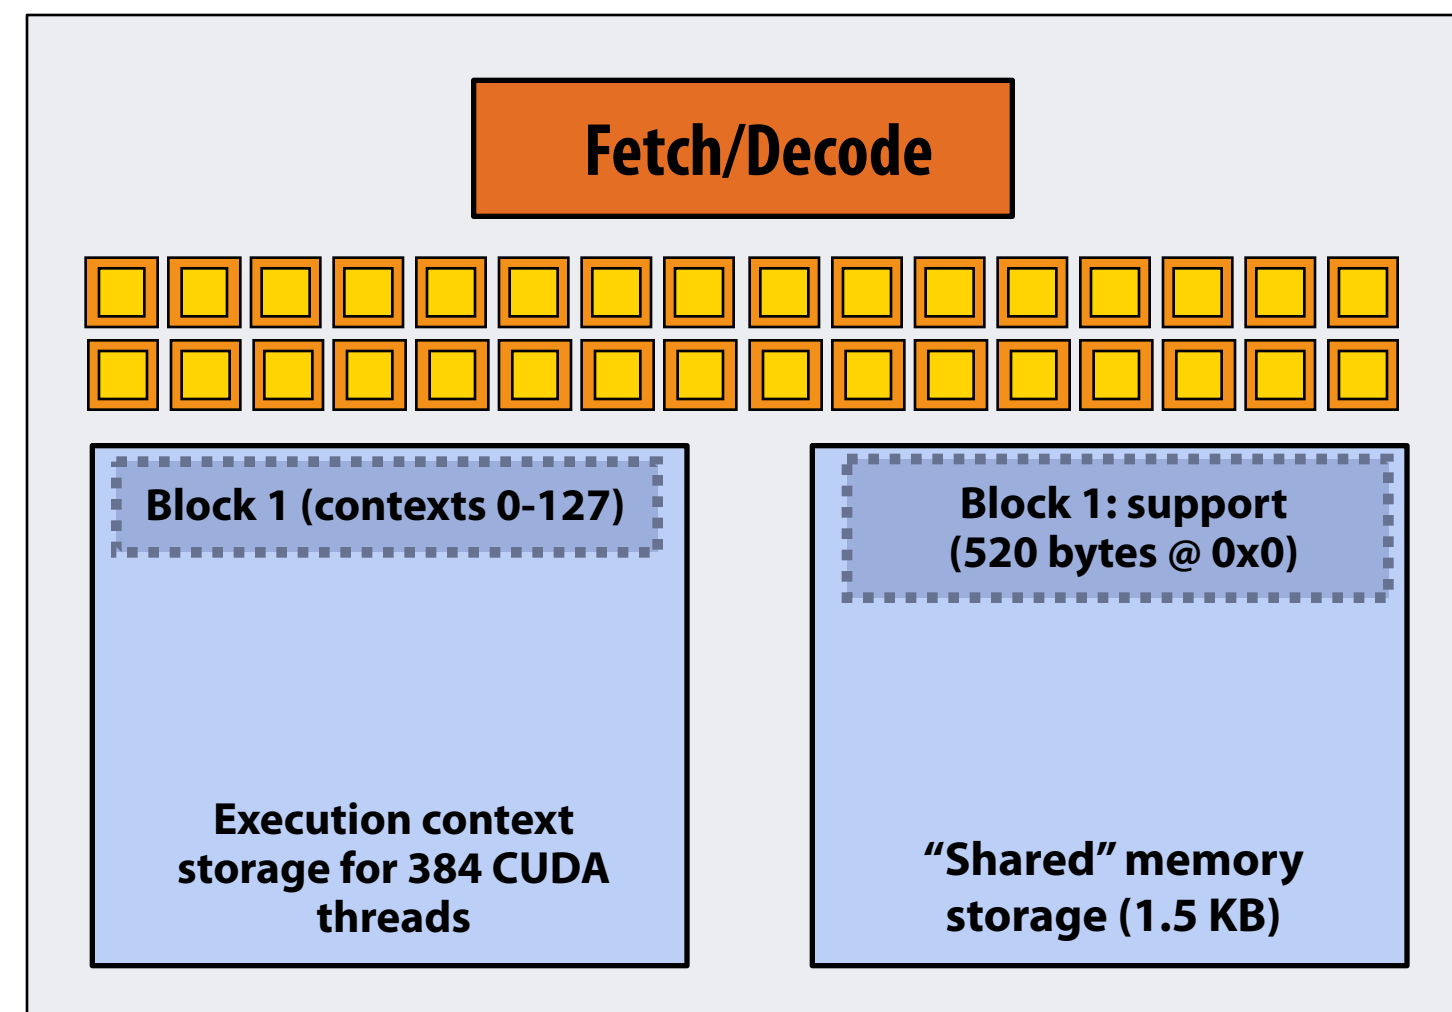
NEXT = 3

GPU Work Scheduler

TOTAL = 1000



Core 0



Core 1

# Running the CUDA kernel

Kernel's execution requirements:

Each thread block must execute 128 CUDA threads

Each thread block must allocate  $130 * \text{sizeof(float)} = 520$  bytes of shared memory

**Step 3: scheduler continues to map blocks to available execution contexts (interleaved mapping shown).  
Only two thread blocks fit on a core (third block won't fit due to insufficient shared storage  $3 * 520 > 1536$ )**

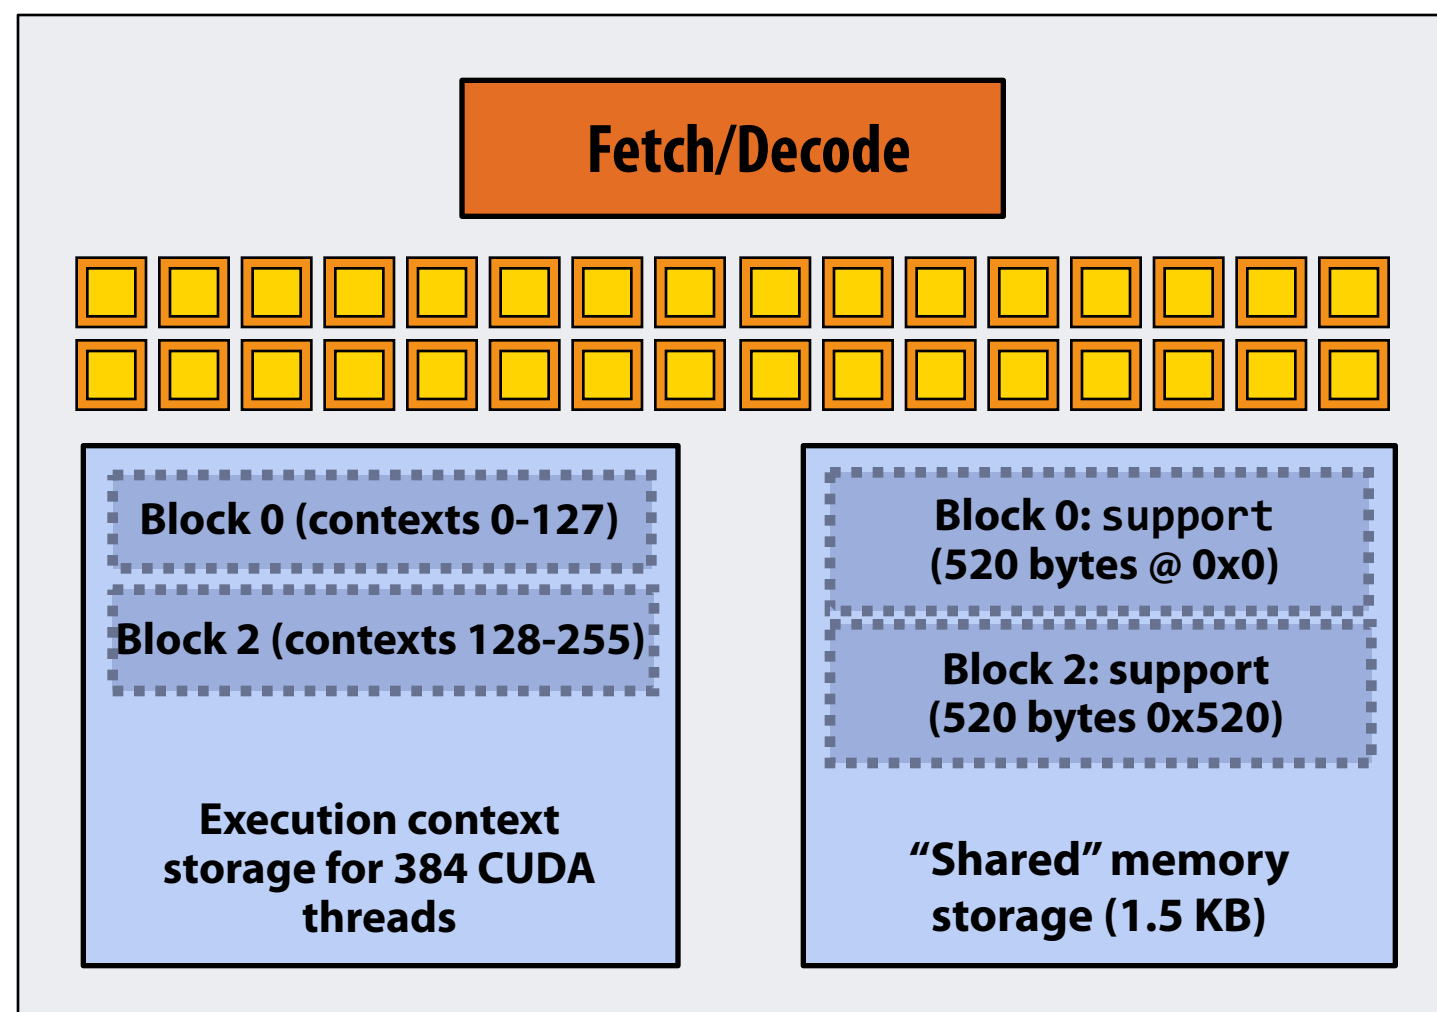


EXECUTE: `convolve_1d_shared`  
ARGS: `N, input_array, output_array`  
NUM\_BLOCKS: 1000

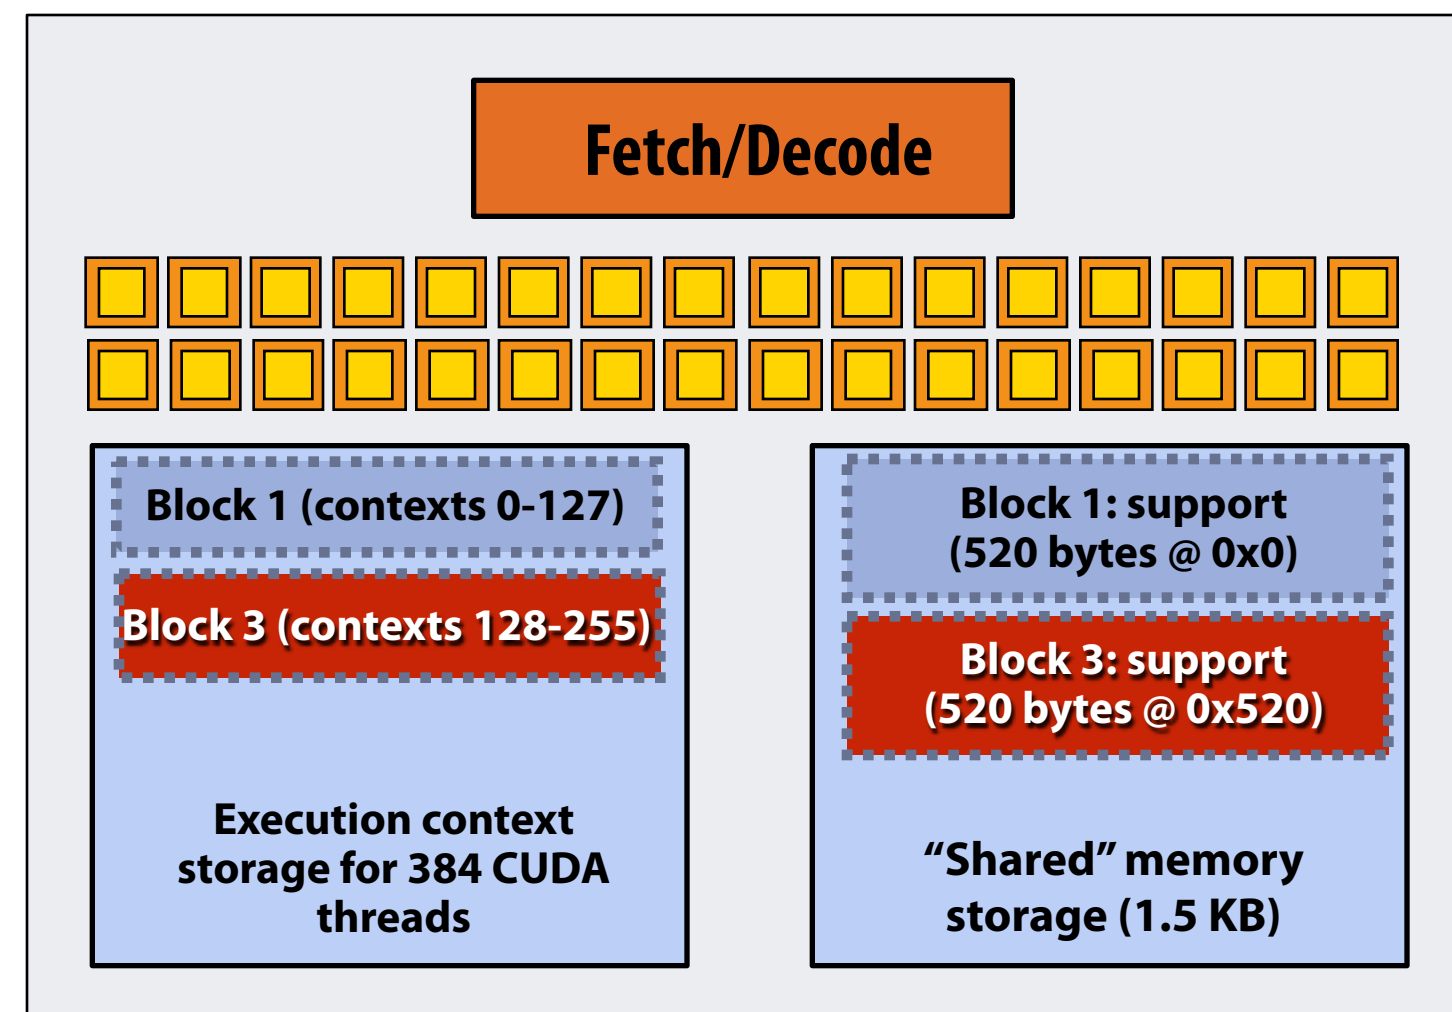
NEXT = 4

GPU Work Scheduler

TOTAL = 1000



Core 0



Core 1

# Running the CUDA kernel

Kernel's execution requirements:

Each thread block must execute 128 CUDA threads

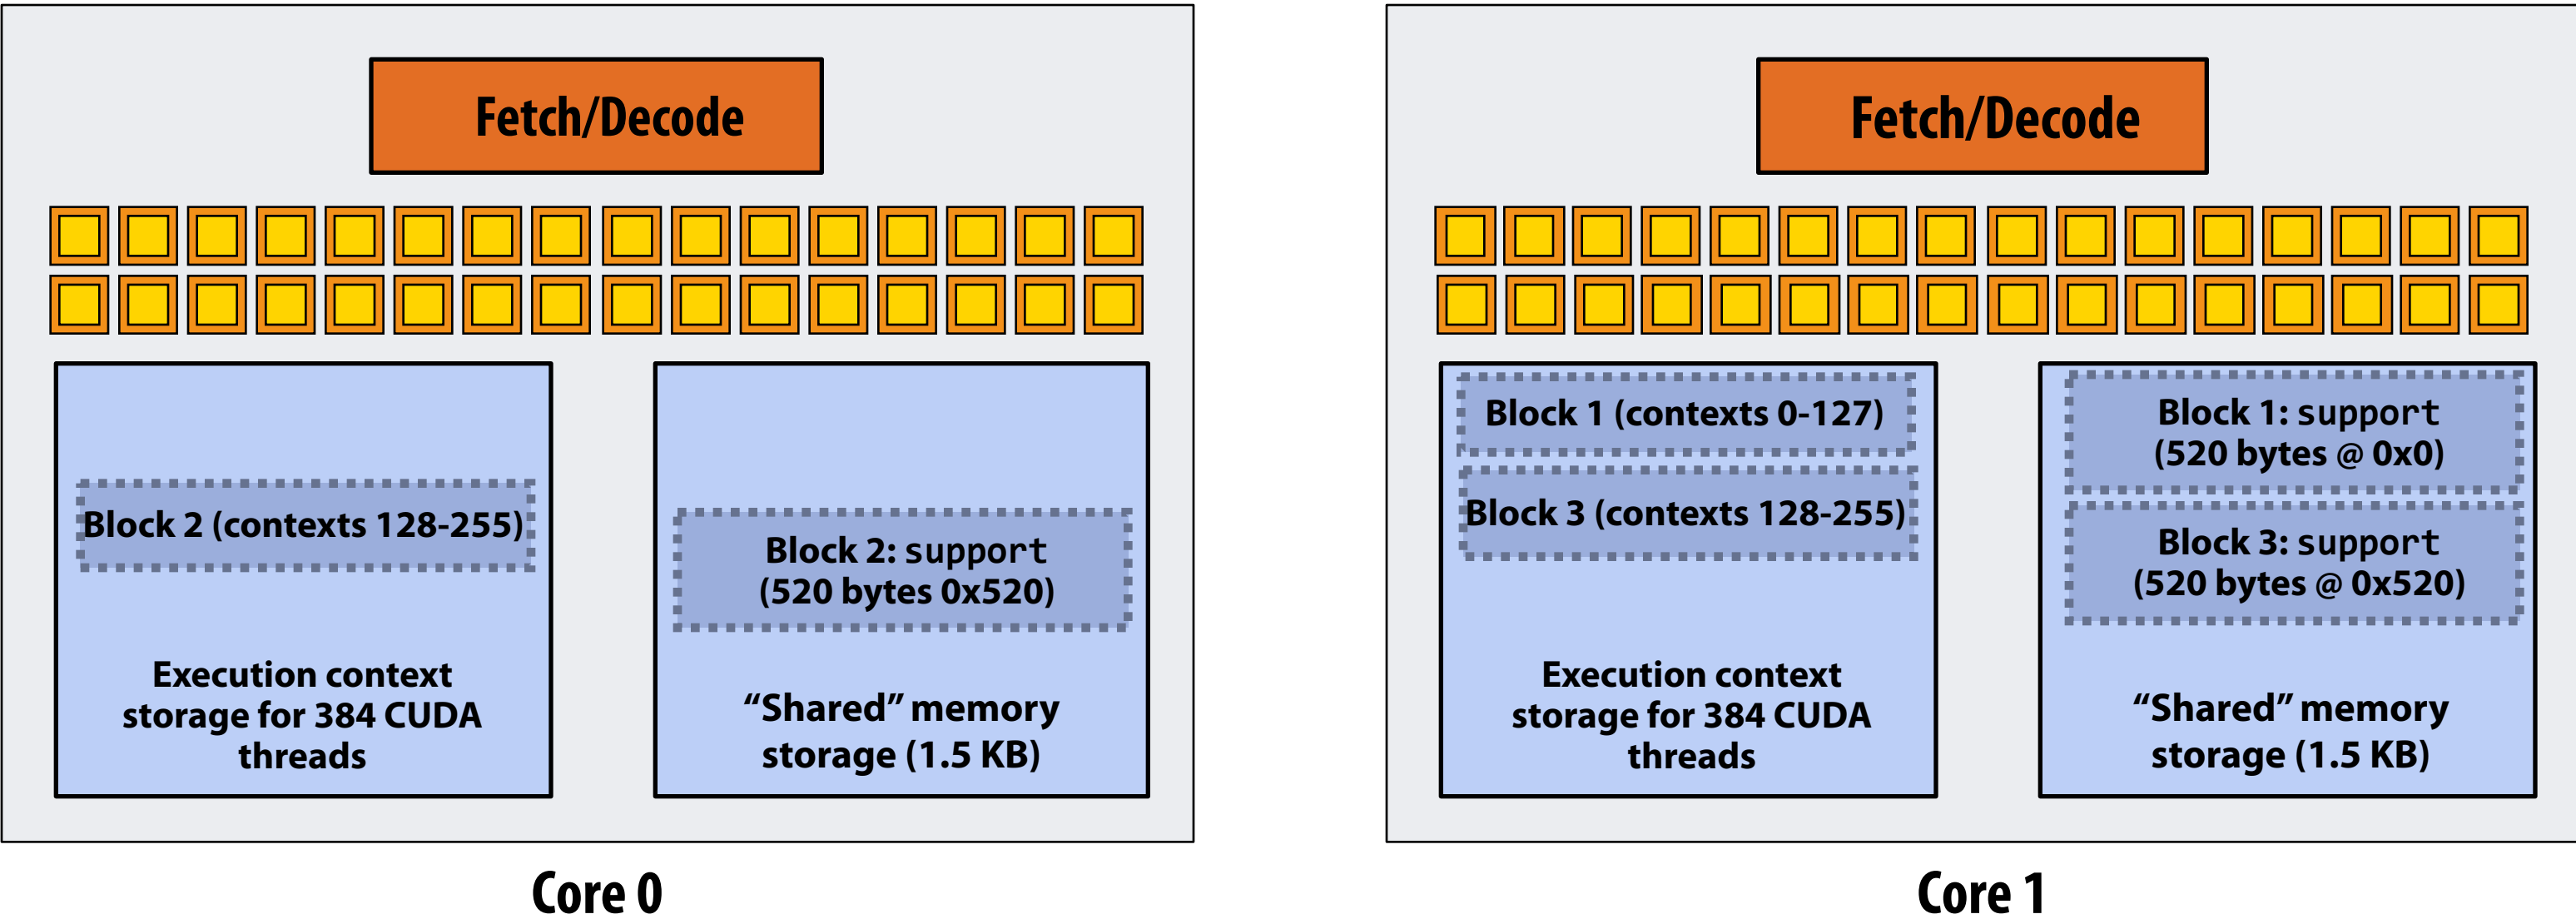
Each thread block must allocate  $130 * \text{sizeof(float)} = 520$  bytes of shared memory

## Step 4: thread block 0 completes on core 0



EXECUTE: `convolve_1d_shared`  
ARGS: `N, input_array, output_array`  
NUM\_BLOCKS: 1000

NEXT = 4      GPU Work Scheduler  
TOTAL = 1000



# Running the CUDA kernel

Kernel's execution requirements:

Each thread block must execute 128 CUDA threads

Each thread block must allocate  $130 * \text{sizeof(float)} = 520$  bytes of shared memory

Step 5: block 4 is scheduled on core 0 (mapped to execution contexts 0-127)

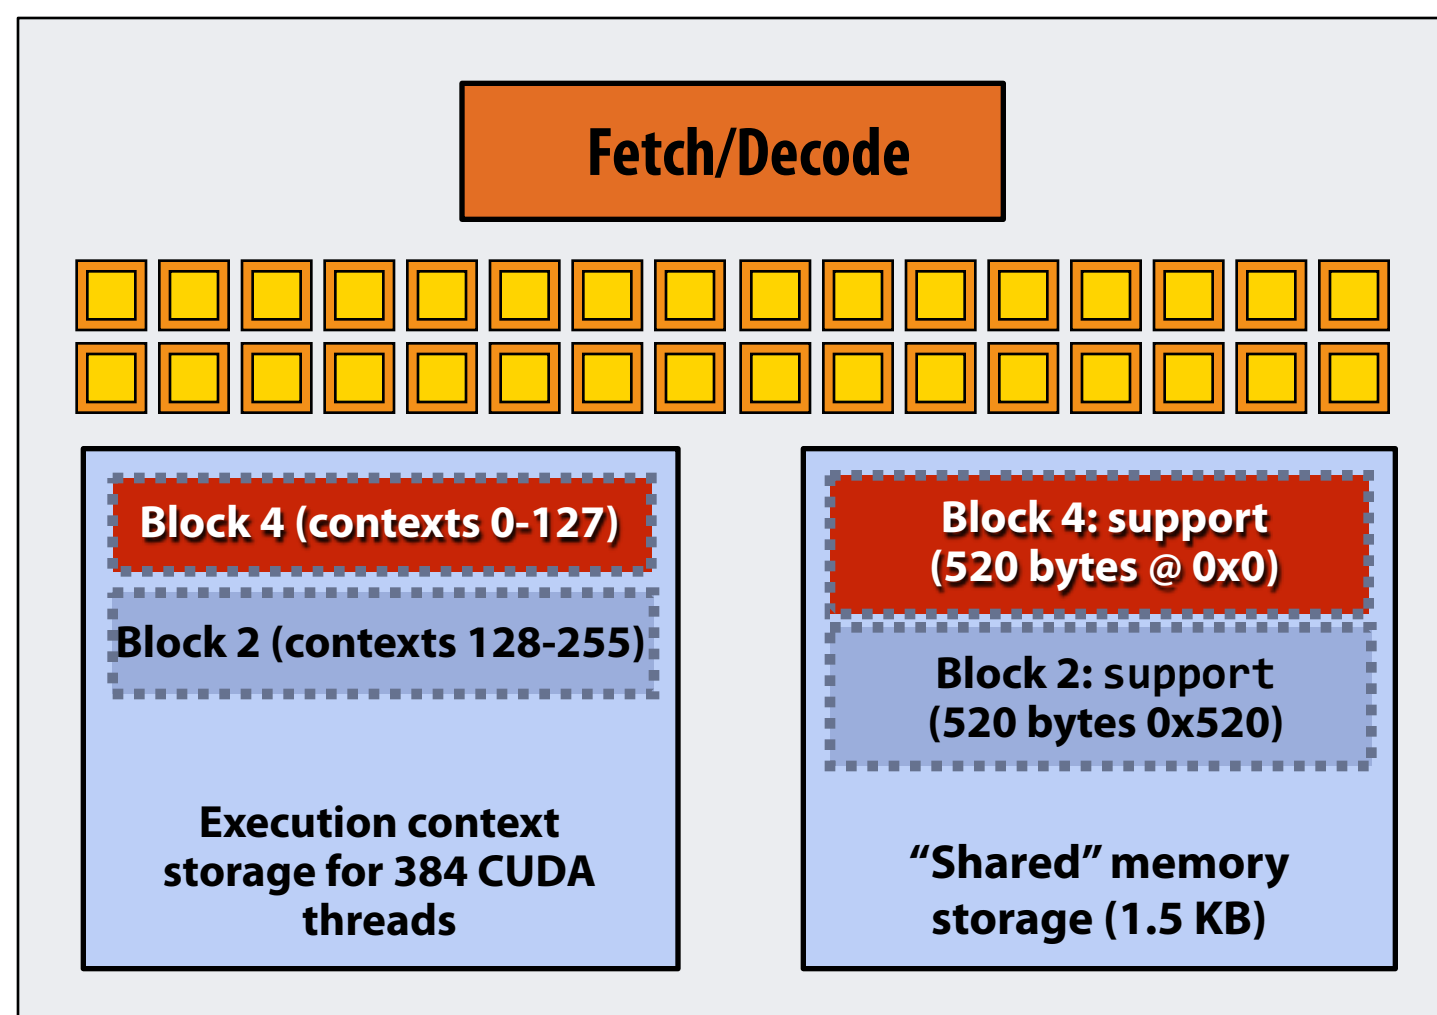


EXECUTE: `convolve_1d_shared`  
ARGS: `N, input_array, output_array`  
NUM\_BLOCKS: 1000

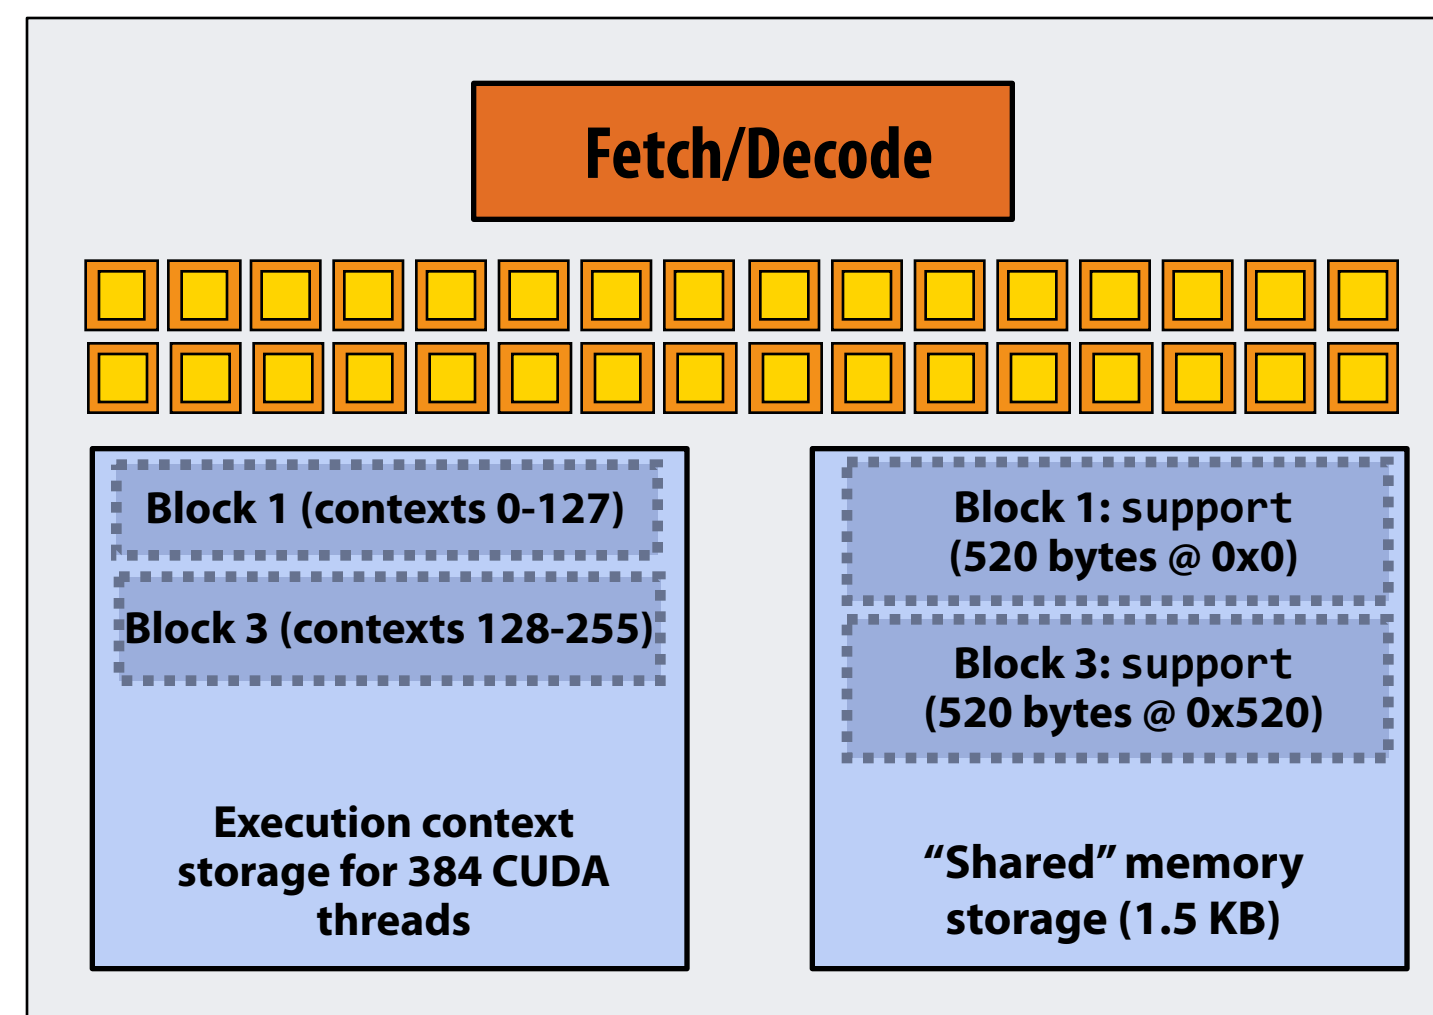
NEXT = 5

GPU Work Scheduler

TOTAL = 1000



Core 0



Core 1

# Running the CUDA kernel

Kernel's execution requirements:

Each thread block must execute 128 CUDA threads

Each thread block must allocate  $130 * \text{sizeof(float)} = 520$  bytes of shared memory

Step 6: thread block 2 completes on core 0

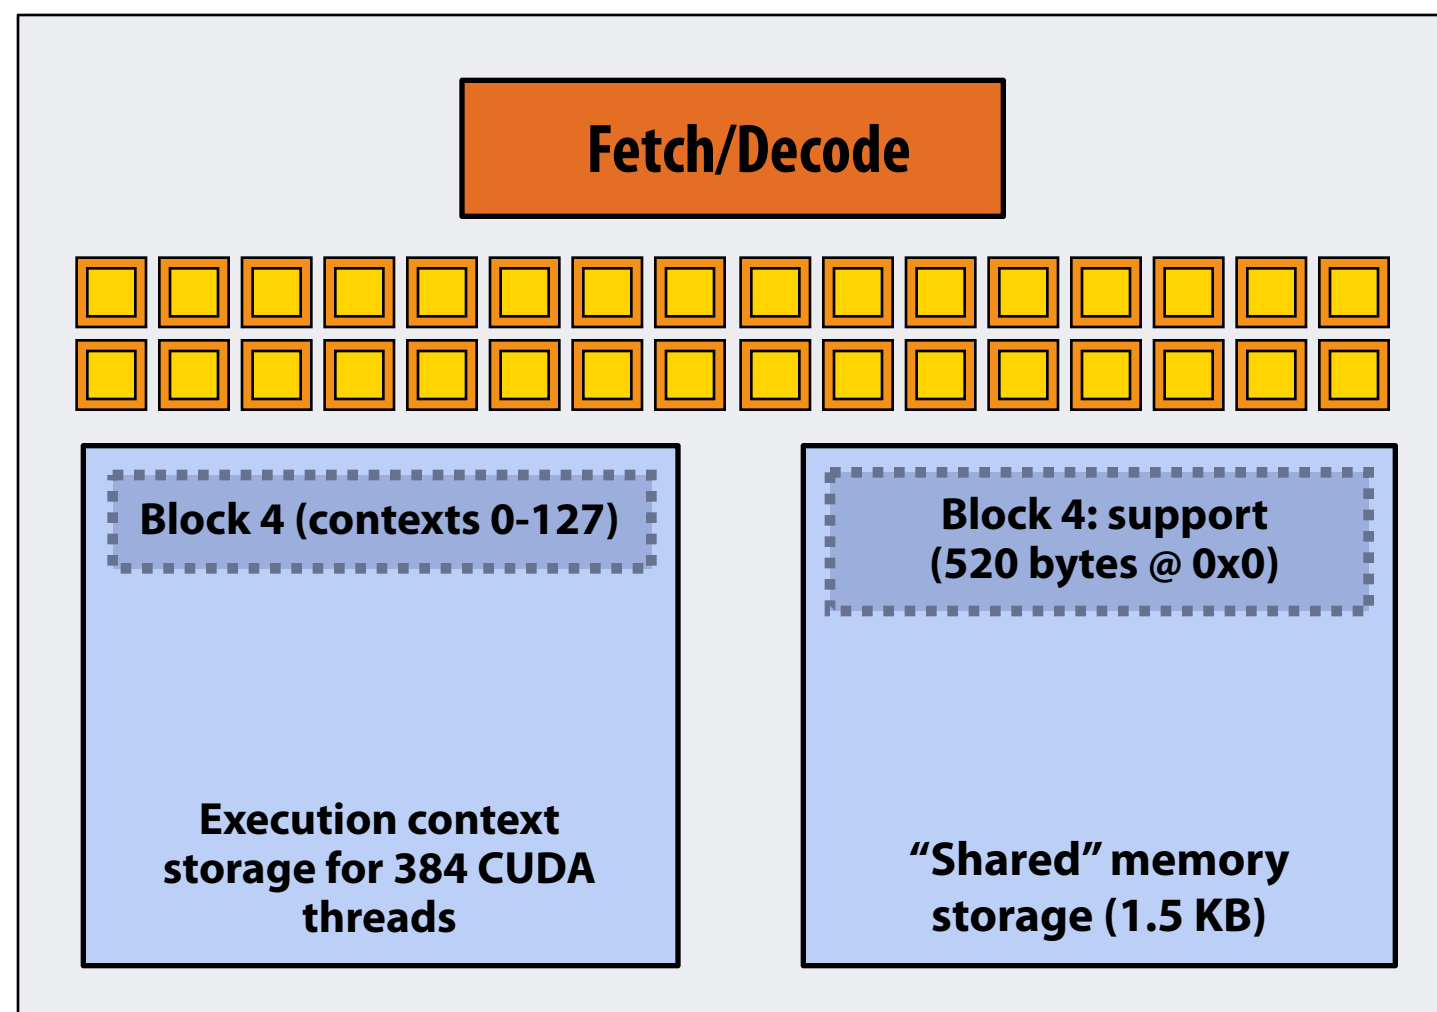


EXECUTE: `convolve_1d_shared`  
ARGS: `N, input_array, output_array`  
NUM\_BLOCKS: 1000

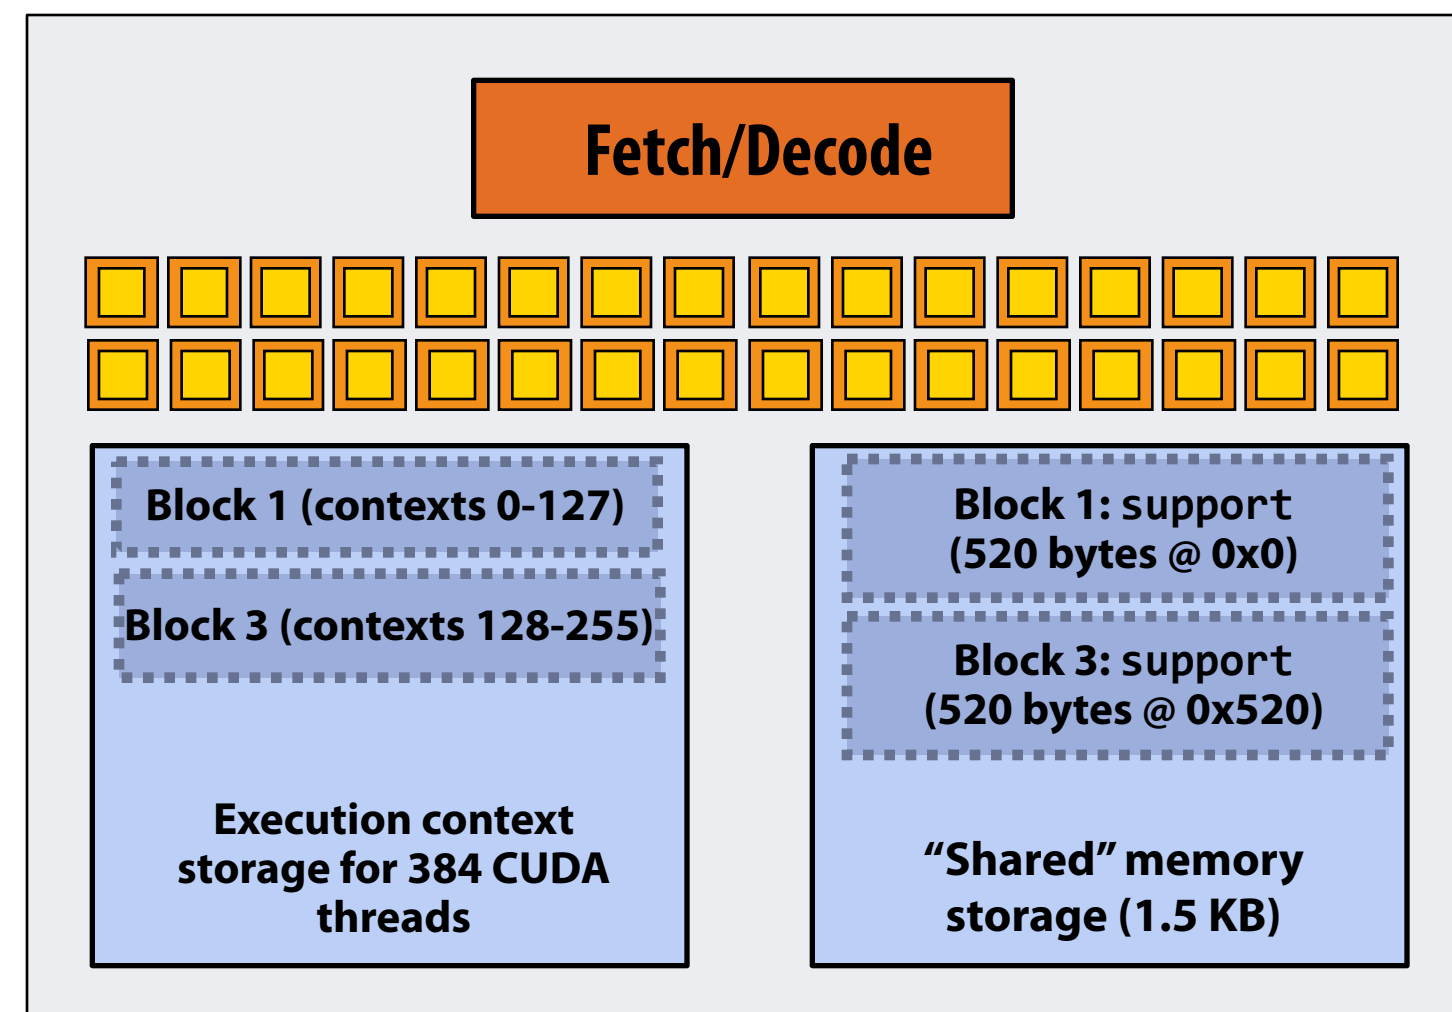
NEXT = 5

GPU Work Scheduler

TOTAL = 1000



Core 0



Core 1

# Running the CUDA kernel

Kernel's execution requirements:

Each thread block must execute 128 CUDA threads

Each thread block must allocate  $130 * \text{sizeof(float)} = 520$  bytes of shared memory

Step 7: thread block 5 is scheduled on core 0 (mapped to execution contexts 128-255)

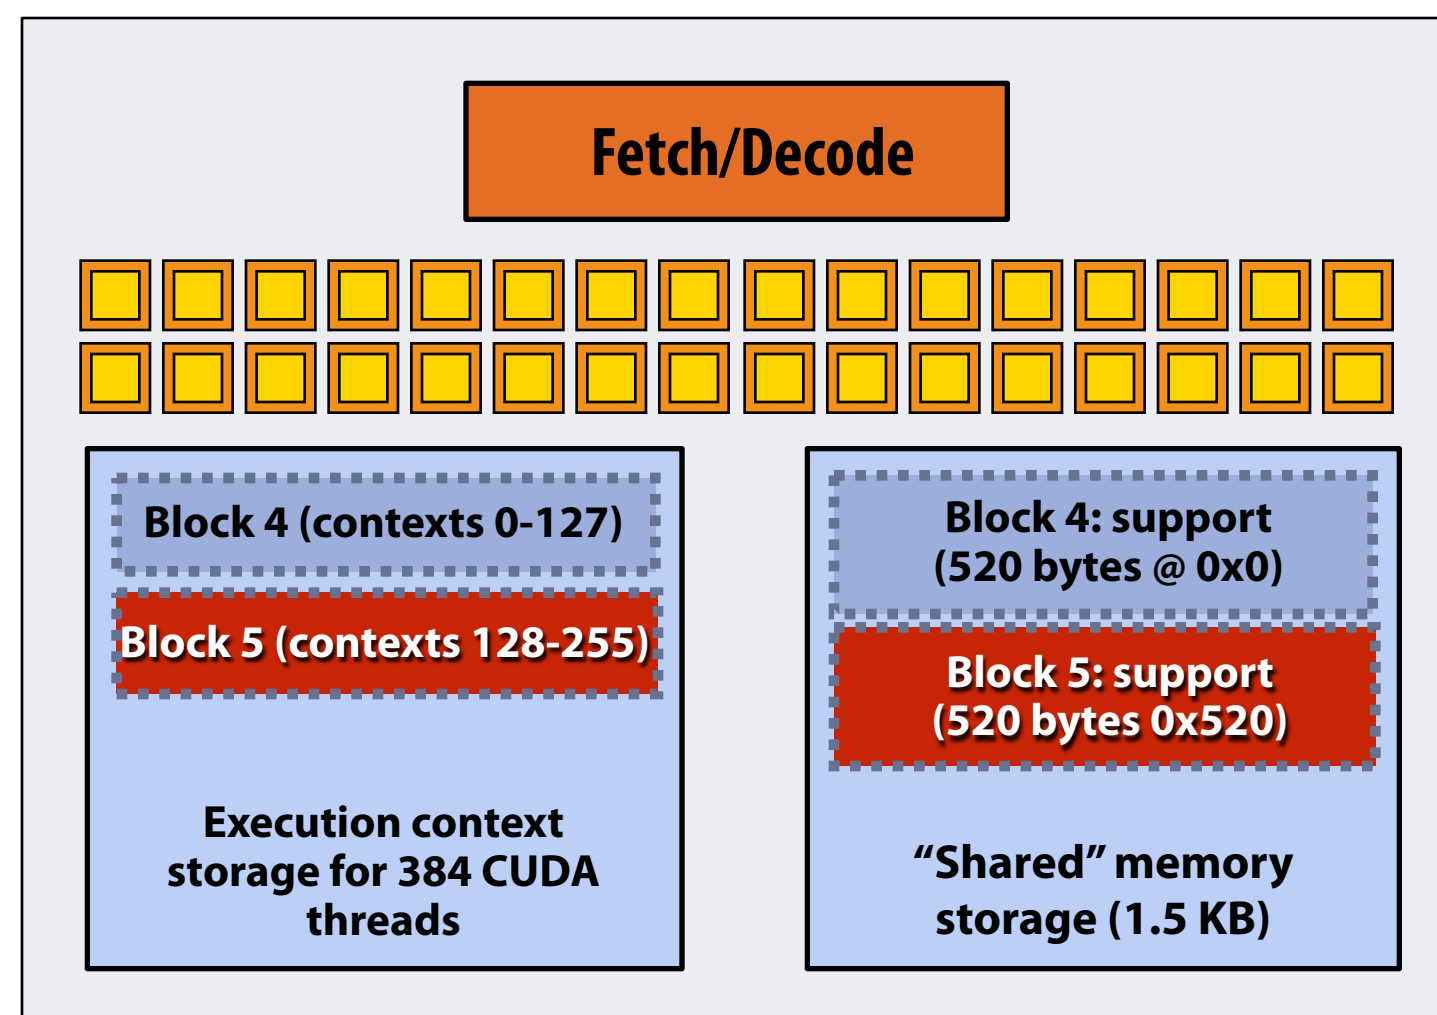


EXECUTE: `convolve_1d_shared`  
ARGS: `N, input_array, output_array`  
NUM\_BLOCKS: 1000

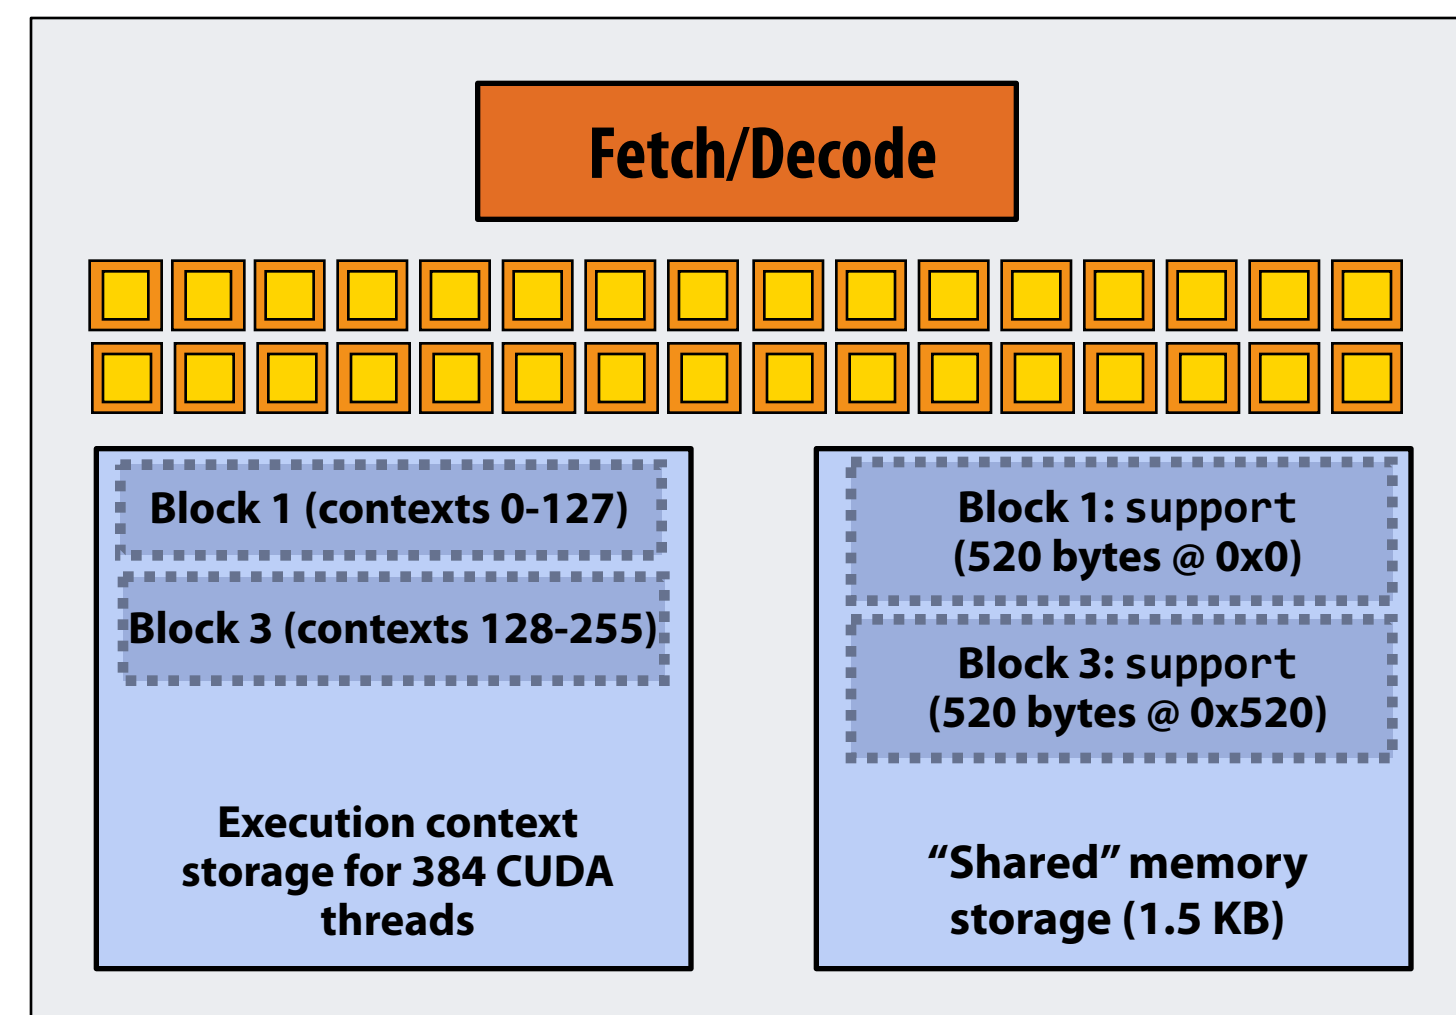
NEXT = 6

GPU Work Scheduler

TOTAL = 1000



Core 0

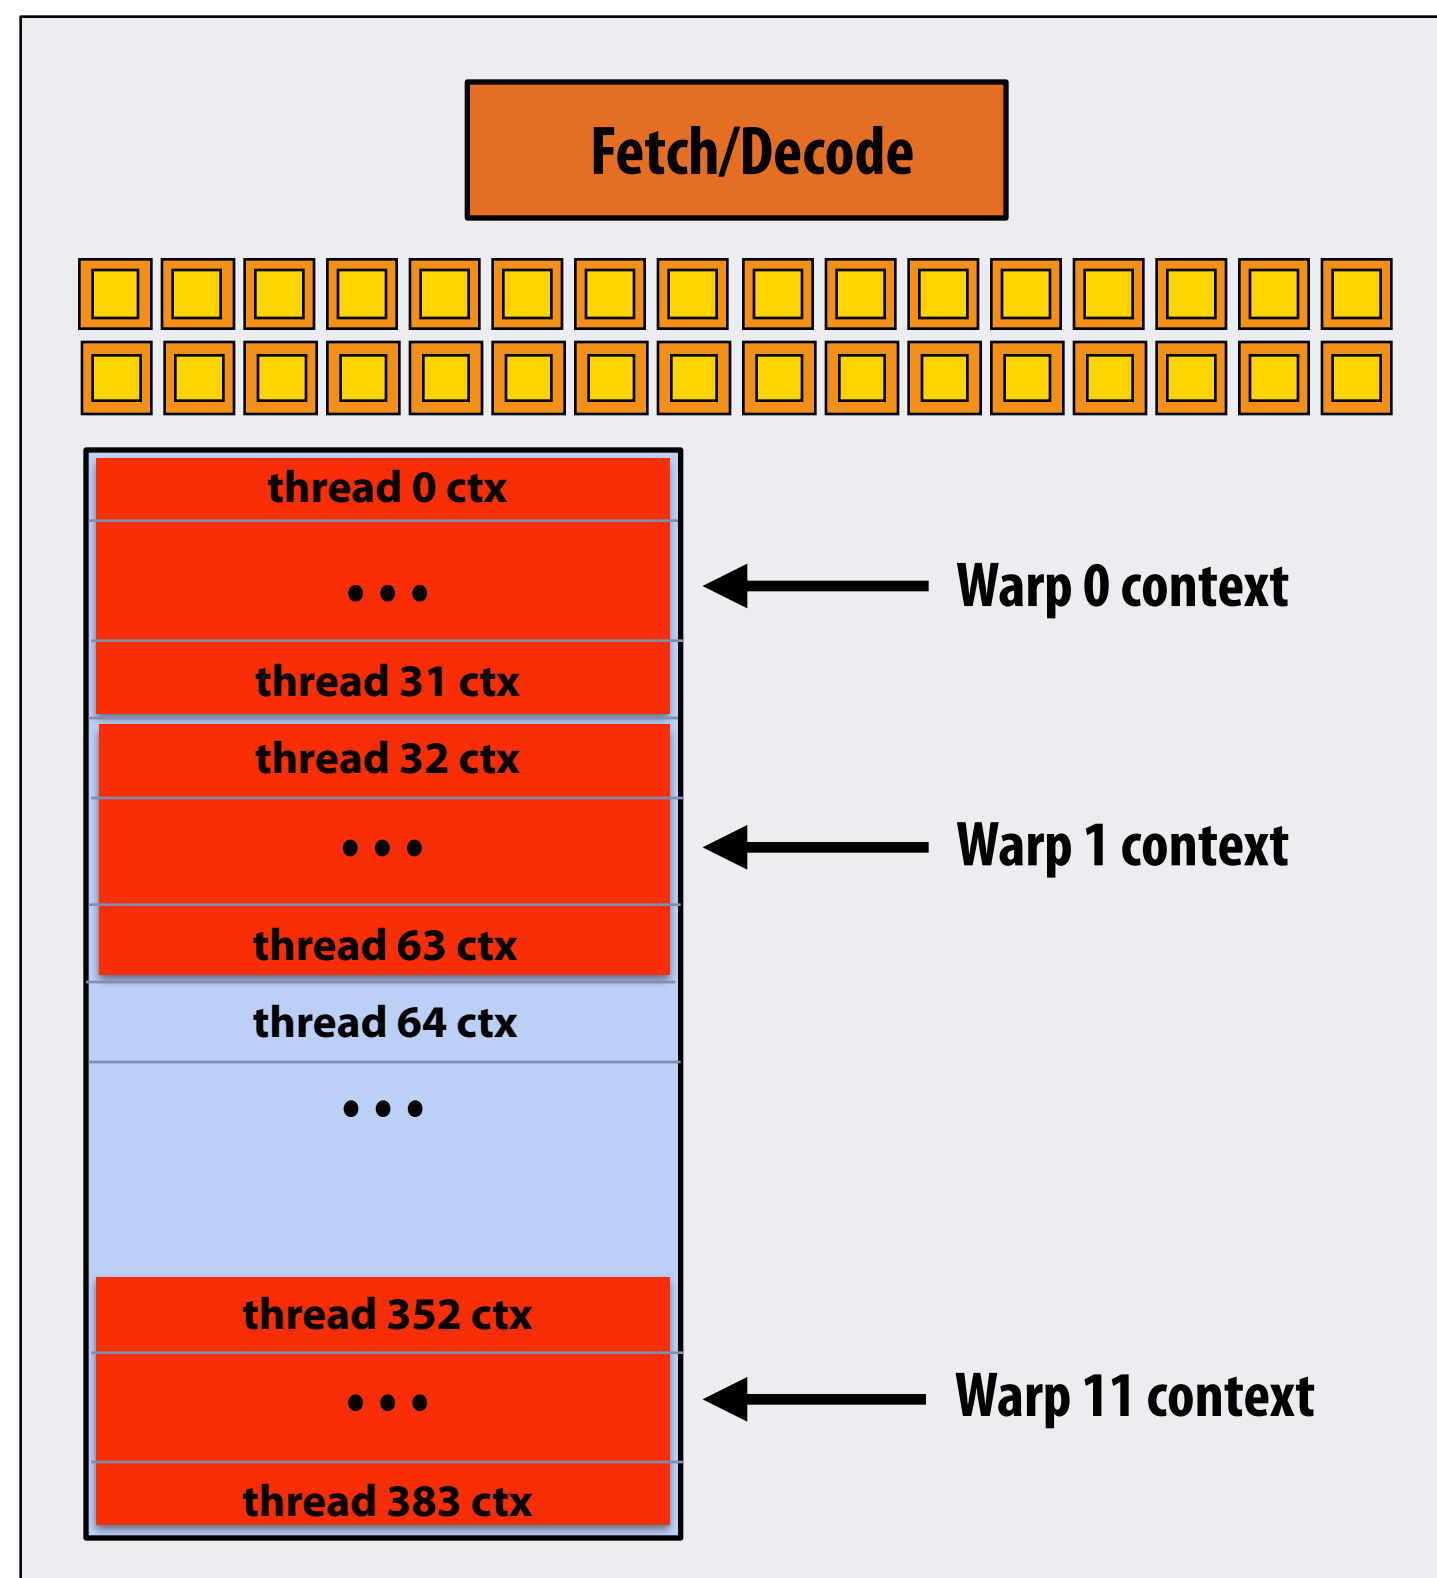


Core 1



# What is a “warp”?

- Before all else: a warp is a CUDA implementation detail on NVIDIA GPUs
- On modern NVIDIA hardware, groups of 32 CUDA threads in a thread block are executed simultaneously using 32-wide SIMD execution.



**In this fictitious NVIDIA GPU example:  
Core maintains contexts for 12 warps  
Selects one warp to run each clock**



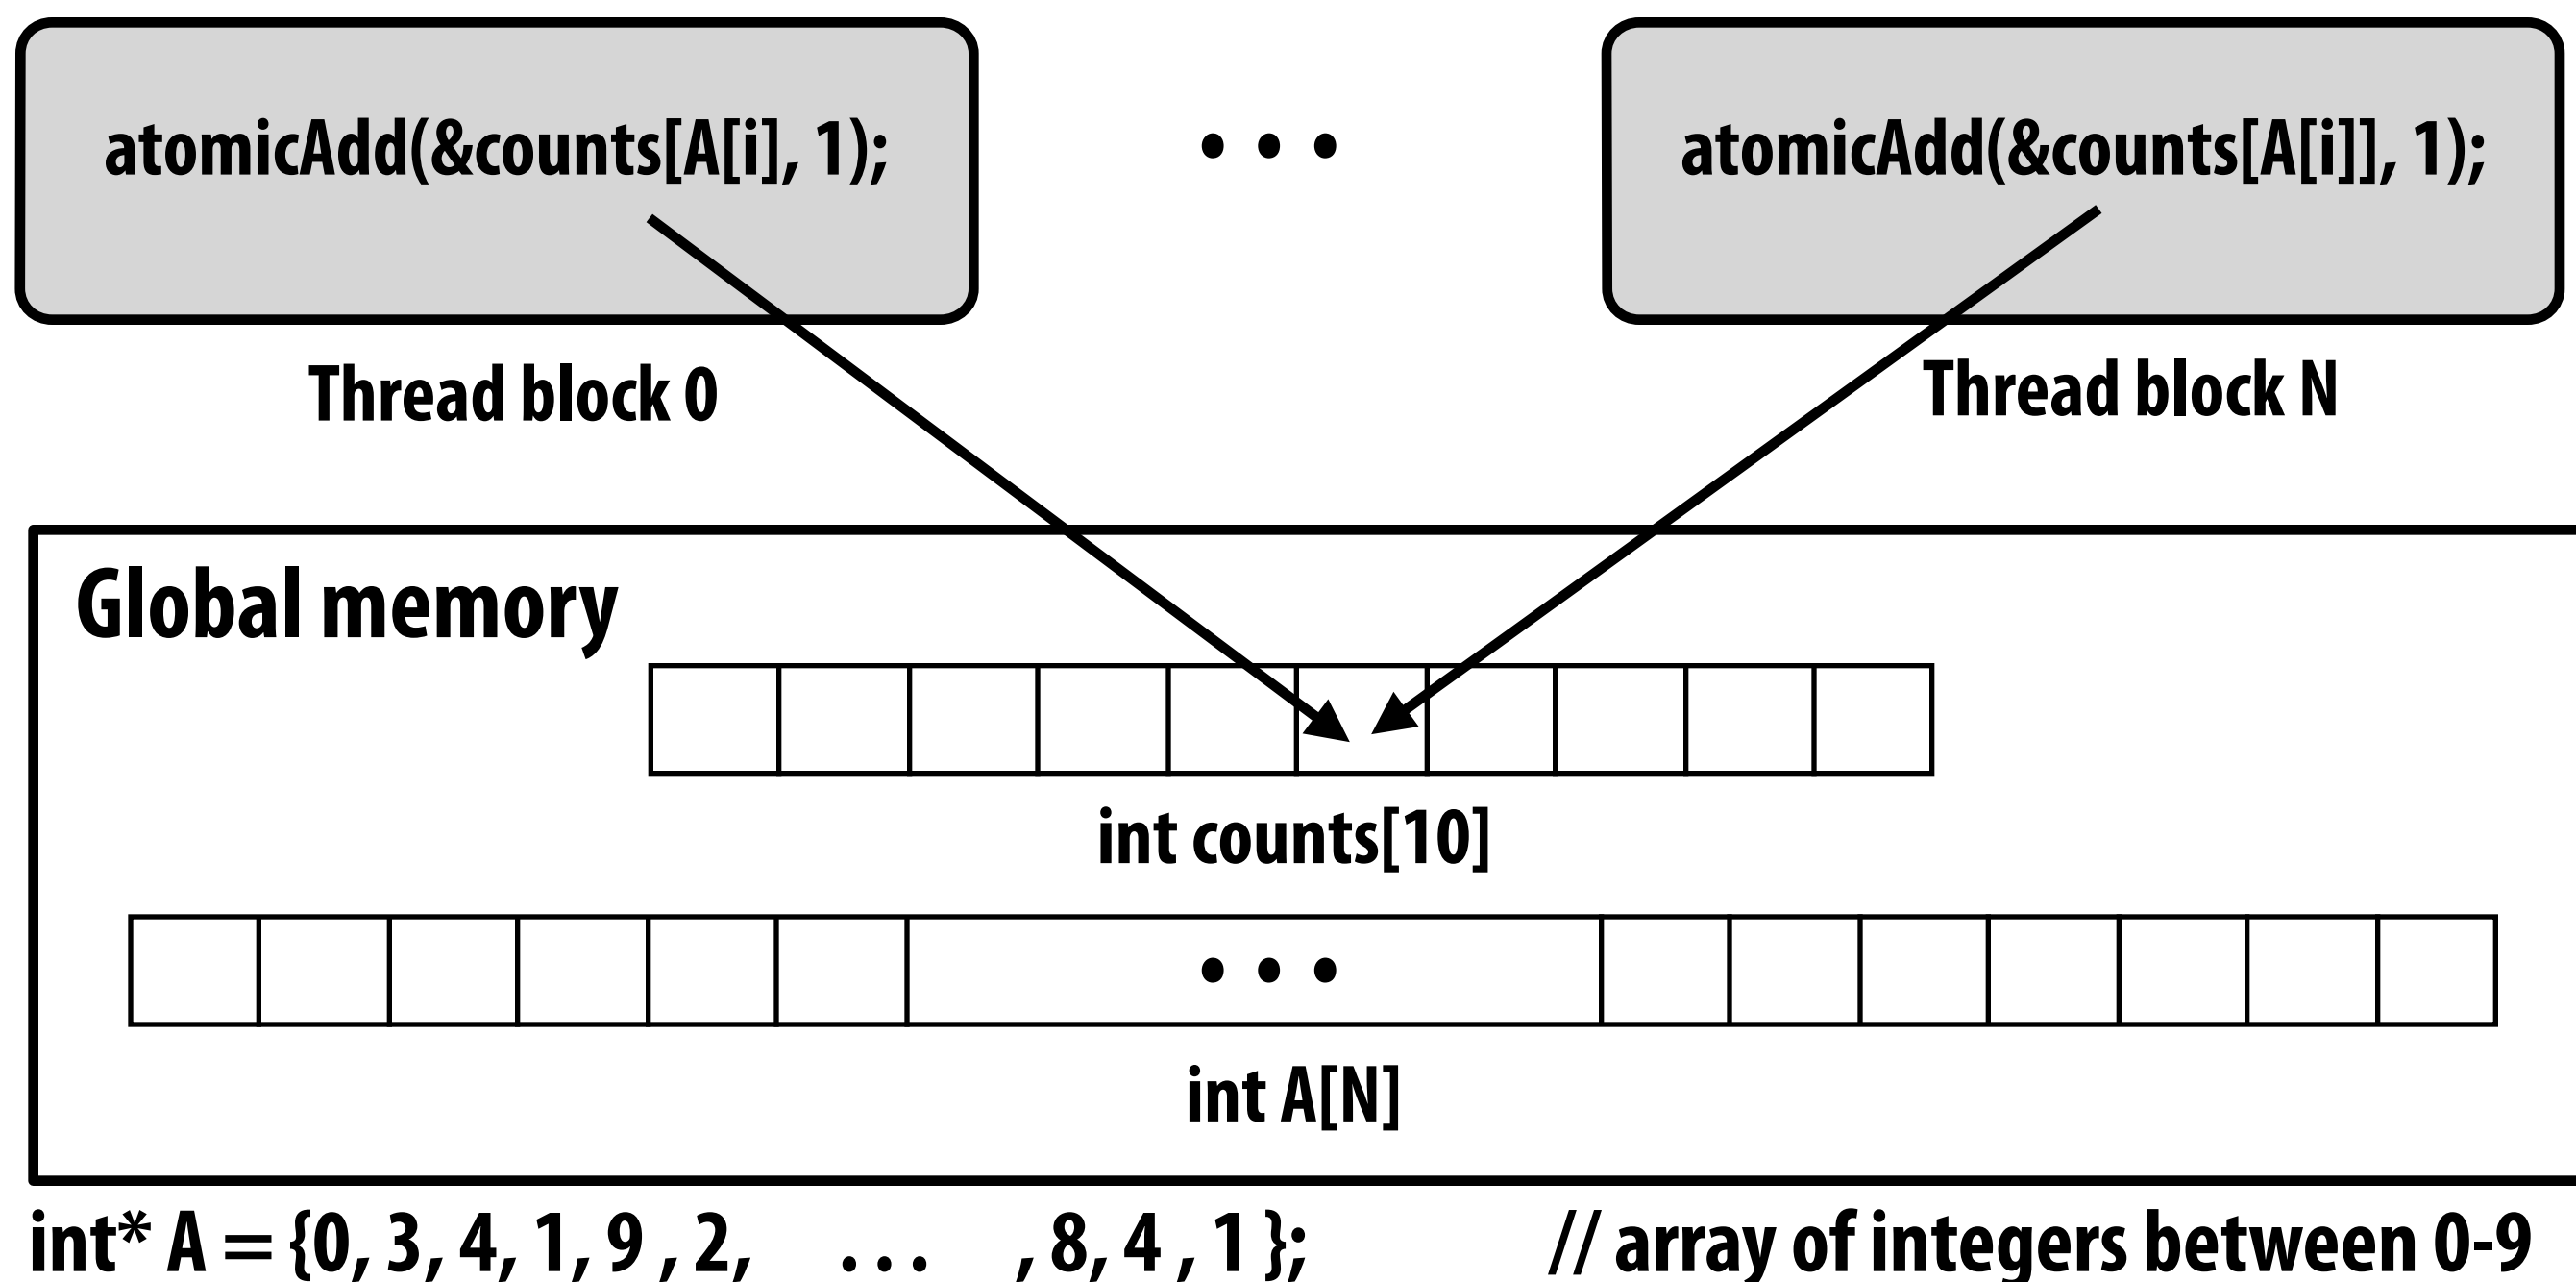
# What is a “warp”?

- Before all else: a warp is a CUDA implementation detail on NVIDIA GPUs
- On modern NVIDIA hardware, groups of 32 CUDA threads in a thread block are executed simultaneously using 32-wide SIMD execution.
  - These 32 logical CUDA threads share an instruction stream and therefore performance can suffer due to divergent execution.
  - This mapping is similar to how ISPC runs program instances in a gang.
- The group of 32 threads sharing an instruction stream is called a warp.
  - In a thread block, threads 0-31 fall into the same warp (so do threads 32-63, etc.)
  - Therefore, a thread block with 256 CUDA threads is mapped to 8 warps.
  - Each “SMX” core in the GTX 680 we discussed last time is capable of scheduling and interleaving execution of up to 64 warps.
  - So a “SMX” core is capable of concurrently executing multiple CUDA thread blocks.

# This program creates a histogram.

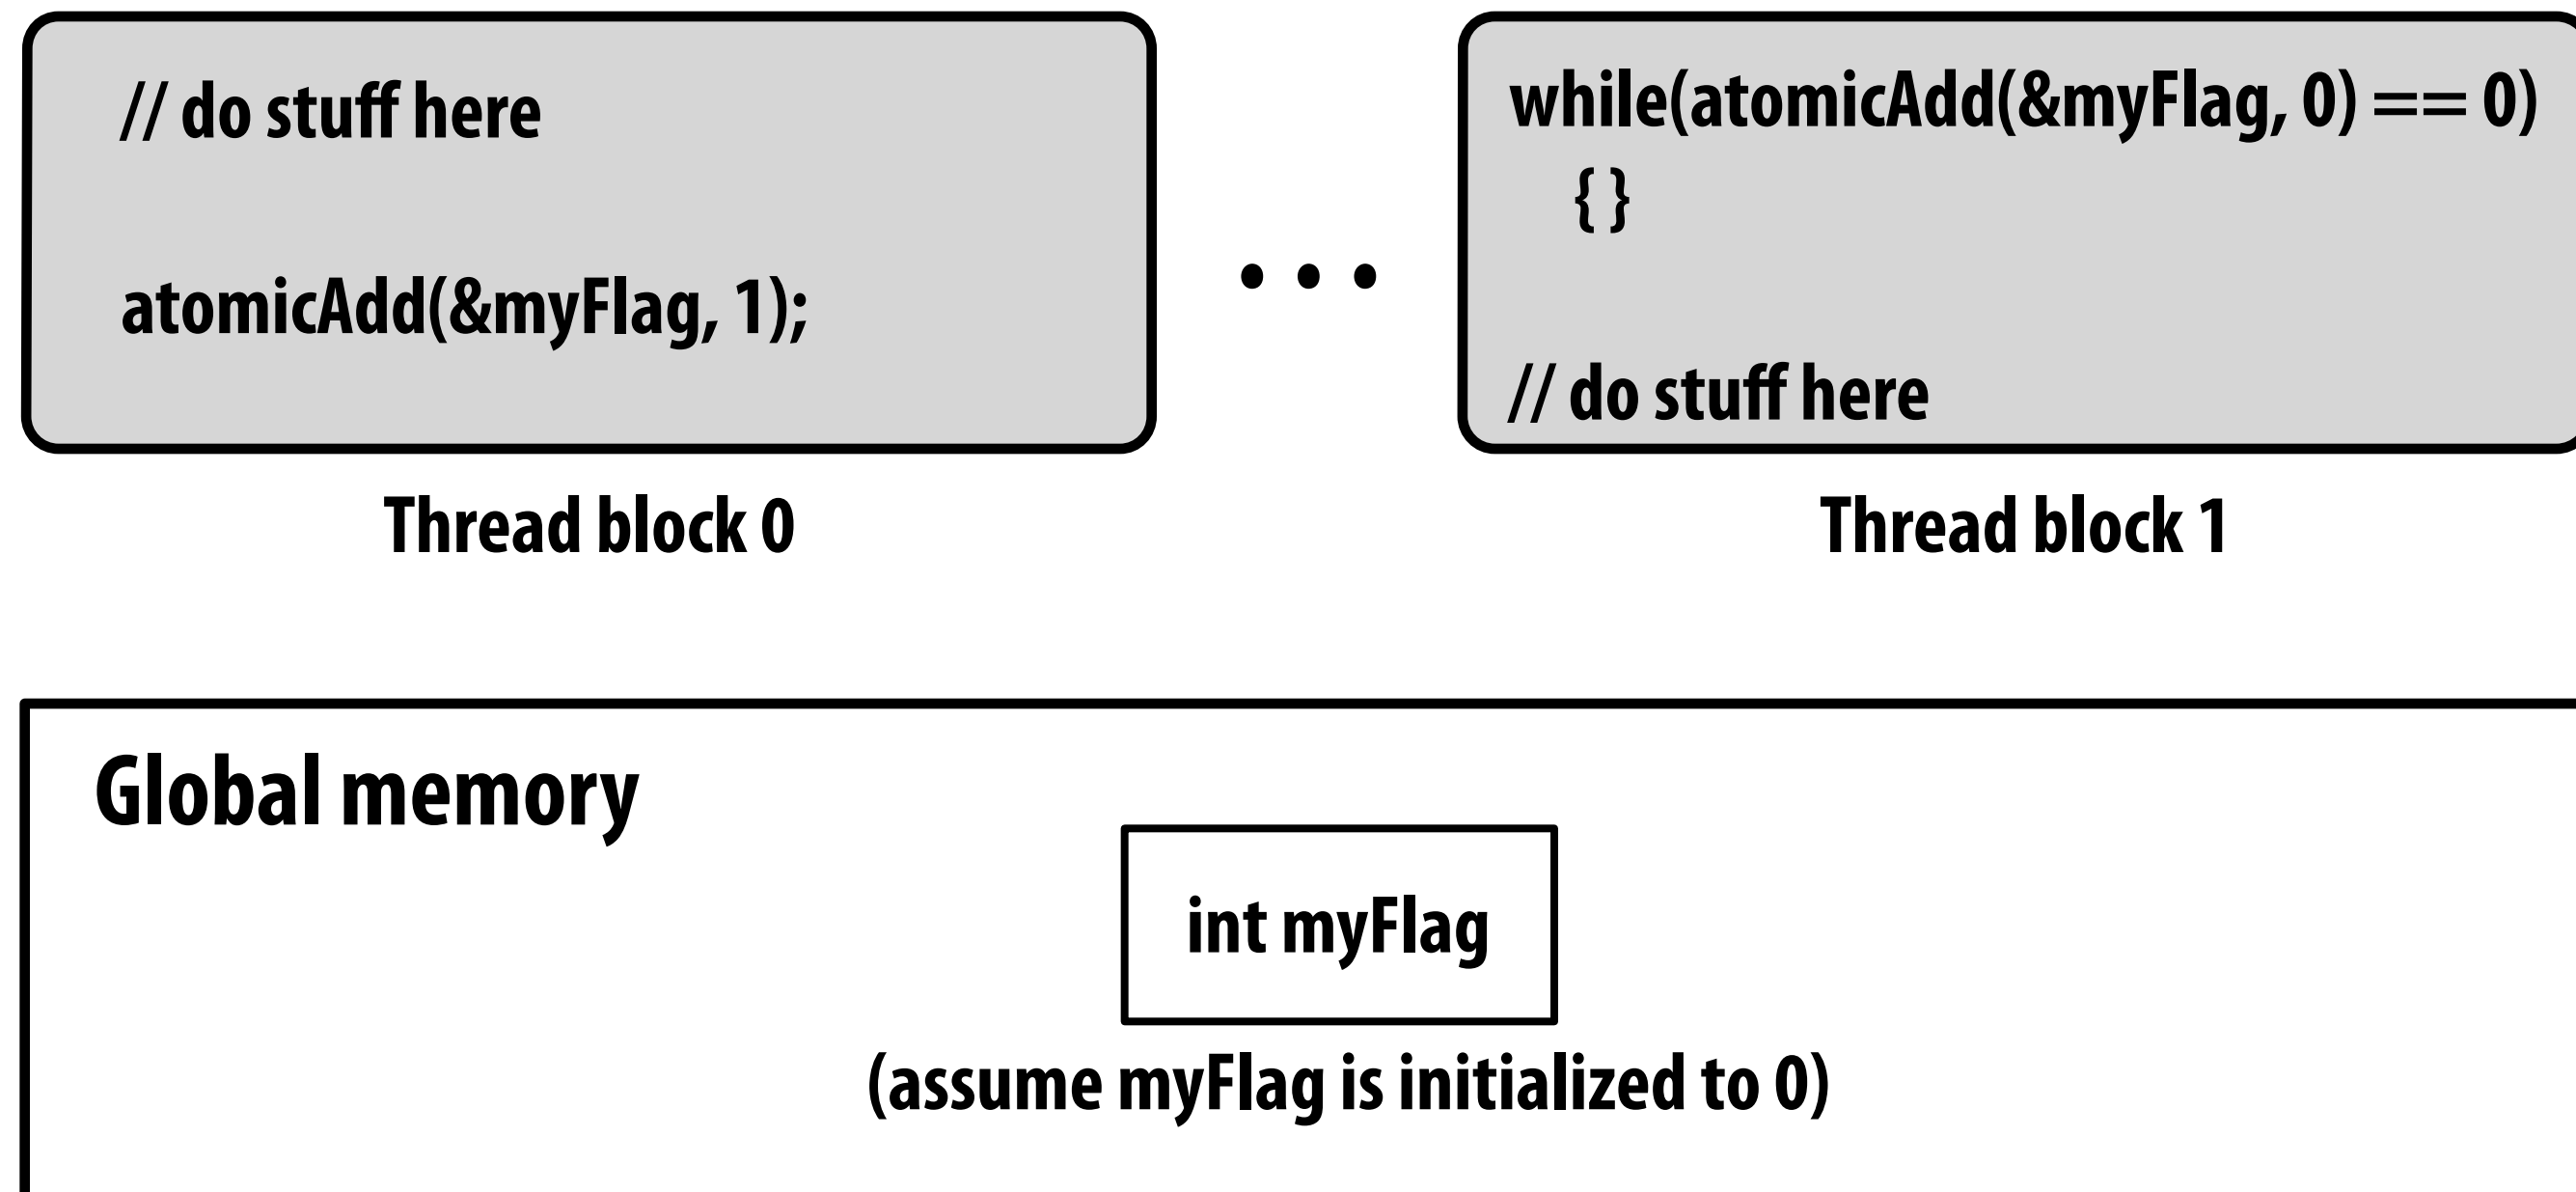
## It is reasonable CUDA code?

- This example: build a histogram of values in an array
  - CUDA threads atomically update shared variables in global memory
- Notice I have never claimed CUDA thread blocks were guaranteed to be independent. I only stated CUDA reserves the right to schedule them in any order.
- This use of atomics does not impact implementation's ability to schedule blocks in any order (atomics used for mutual exclusion, and nothing more)



# But is this reasonable CUDA code?

- Consider implementation of on a single core GPU with resources for one CUDA thread block per core
  - What happens if the CUDA implementation runs block 0 first?
  - What happens if the CUDA implementation runs block 1 first?



# “Persistent thread” CUDA programming style

```
#define THREADS_PER_BLK 128
#define BLOCKS_PER_CHIP 15 * 12    // specific to a certain GTX 480 GPU

__device__ int workCounter = 0;    // global mem variable

__global__ void convolve(int N, float* input, float* output) {
    __shared__ int startingIndex;
    __shared__ float support[THREADS_PER_BLK+2];    // shared across block
    while (1) {

        if (threadIdx.x == 0)
            startingIndex = atomicInc(workCounter, THREADS_PER_BLK);
        __syncthreads();
        if (startingIndex >= N)
            break;

        int index = startingIndex + threadIdx.x;    // thread local
        support[threadIdx.x] = input[index];
        if (threadIdx.x < 2)
            support[THREADS_PER_BLK+threadIdx.x] = input[index+THREADS_PER_BLK];

        __syncthreads();

        float result = 0.0f;    // thread-local variable
        for (int i=0; i<3; i++)
            result += support[threadIdx.x + i];
        output[index] = result;

        __syncthreads();
    }
}

// host code ////////////////////////////////////////
int N = 1024 * 1024;
cudaMalloc(&devInput, N+2);    // allocate array in device memory
cudaMalloc(&devOutput, N);    // allocate array in device memory
// properly initialize contents of devInput here ...

convolve<<<BLOCKS_PER_CHIP, THREADS_PER_BLK>>>(N, devInput, devOutput);
```

**Idea: write CUDA code that requires knowledge of the number of cores and blocks per core that are supported by underlying GPU implementation.**

**Programmer launches exactly as many thread blocks as will fill the GPU**

**(Program makes assumptions about GPU implementation: that GPU will in fact run all blocks concurrently. Ug!)**

**Now, work assignment to blocks is implemented entirely by the application (circumvents GPU thread block scheduler)**

**Now programmer’s mental model is that \*all\* threads are concurrently running on the machine at once.**