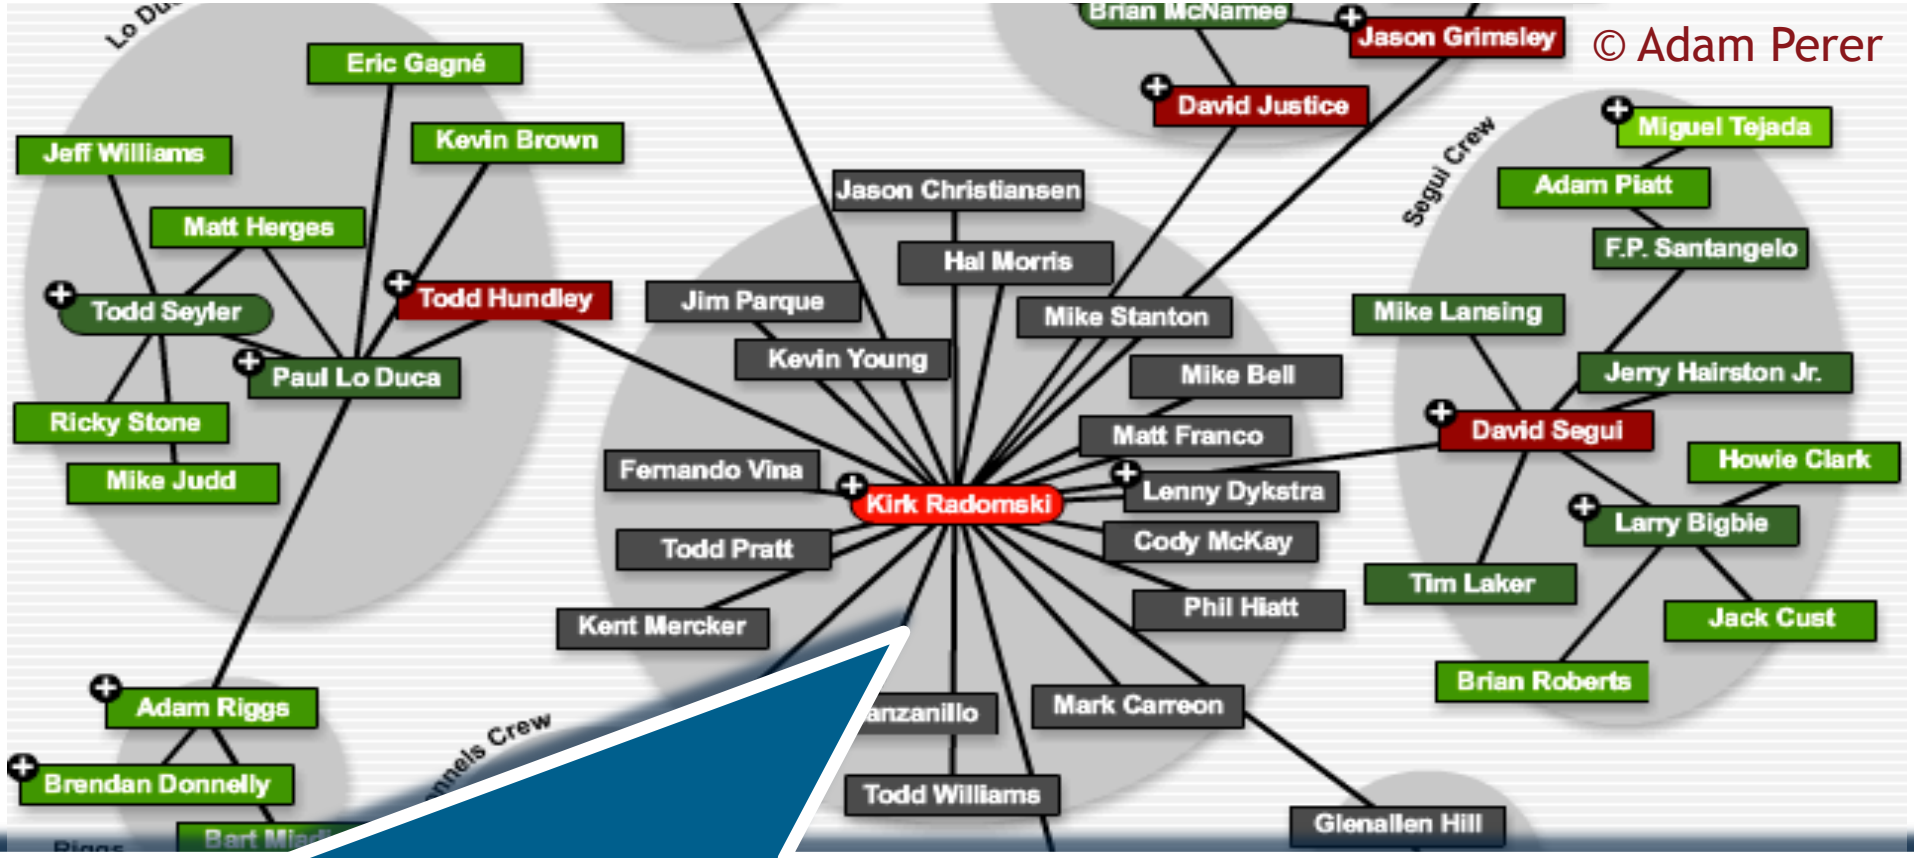

PMatch: Probabilistic Subgraph Matching on Huge Social Networks

Matthias Bröcheler, Andrea Pugliese & V.S. Subrahmanian

UNIVERSITY OF MARYLAND 18 56

UNIVERSITÀ DELLA CALABRIA
Dipartimento di ELETTRONICA, INFORMATICA E SISTEMISTICA

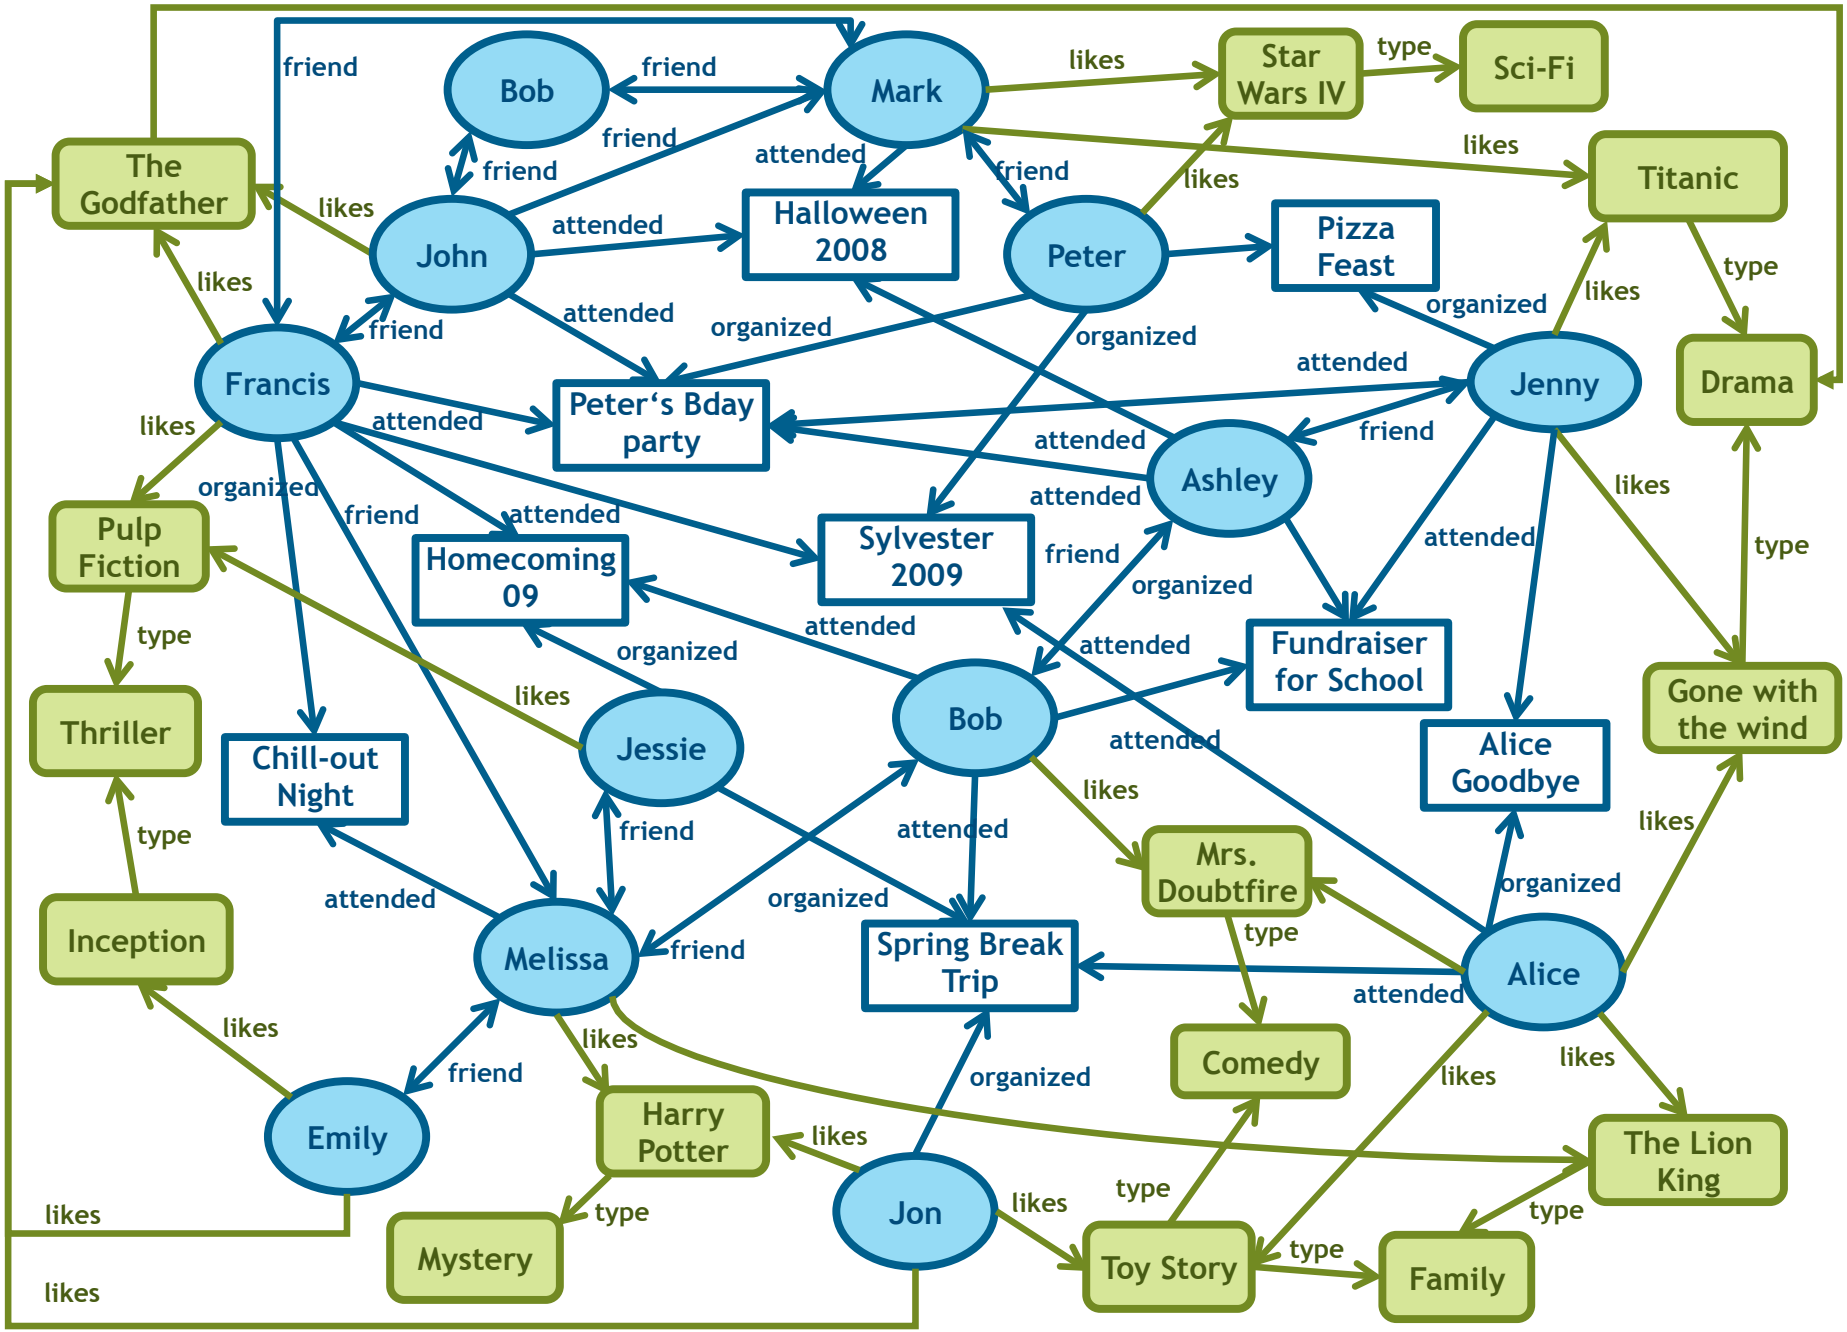
Bob
Mark
Star Wars IV
Sci-Fi
The Godfather
John
Halloween 2008
Peter
Pizza Feast
Titanic
Francis
Peter's Bday party
Jenny
Drama
Ashley
Pulp Fiction
Homecoming 09
Sylvester 2009
Fundraiser for School
Gone with the wind
Thriller
Chill-out Night
Jessie
Bob
Alice Goodbye
Mrs. Doubtfire
Inception
Melissa
Spring Break Trip
Alice
Comedy
Emily
Harry Potter
Jon
The Lion King
Mystery
Toy Story
Family
friend
likes
attended
organized
type

Example Query

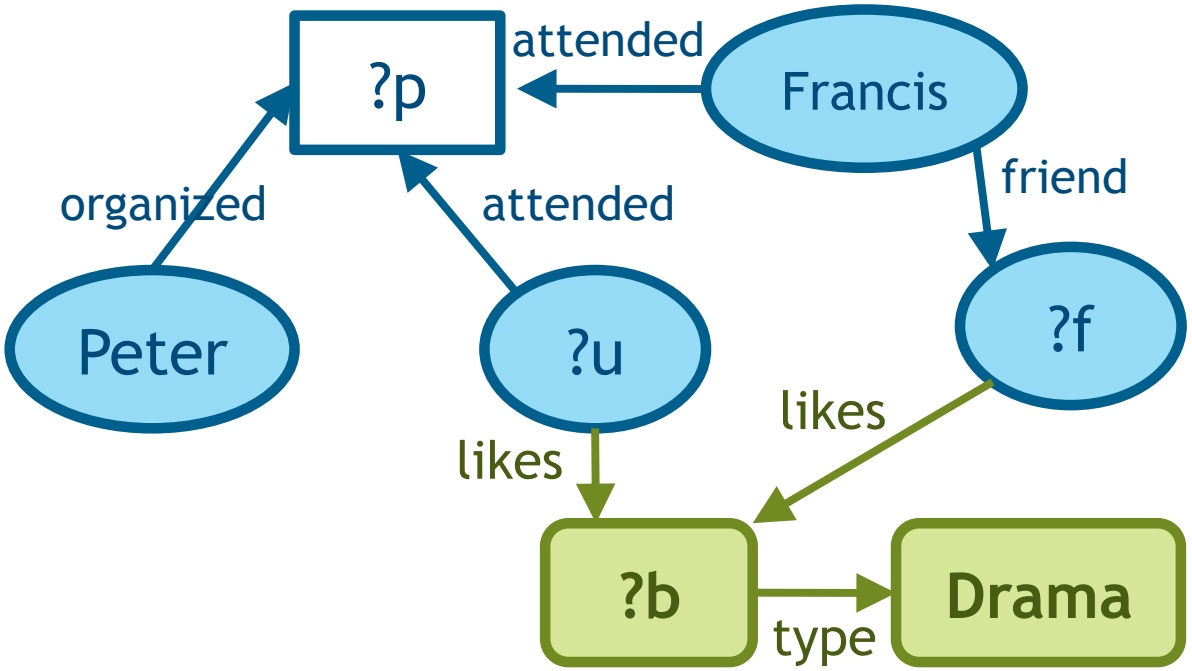

Simple query, yet complex structure and difficult to answer by hand

Query Uncertainty

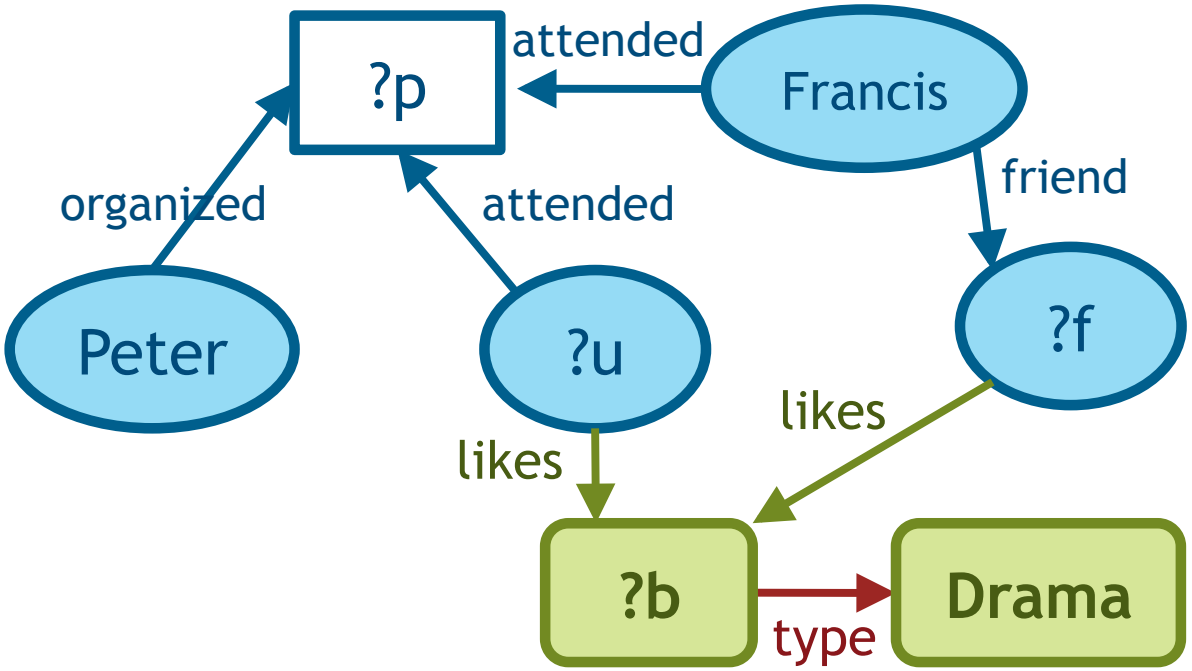

Was it really a drama movie?

Probabilistic Matching Query

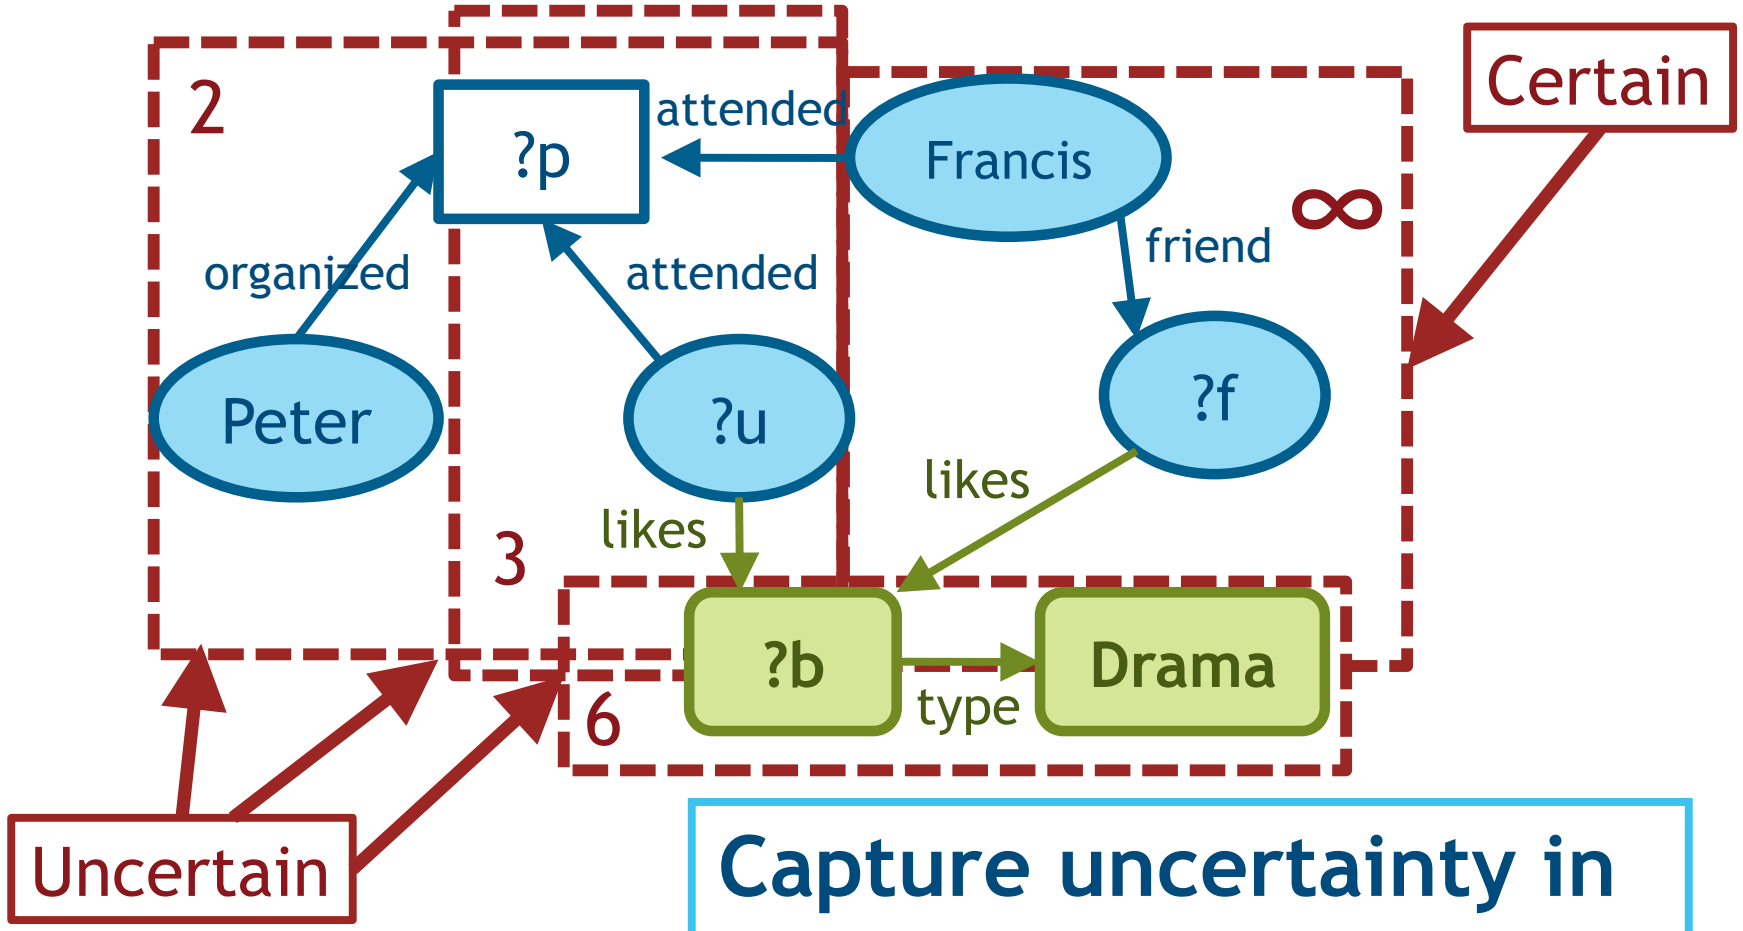

Capture uncertainty in weighted query boxes

Intelligence and Security

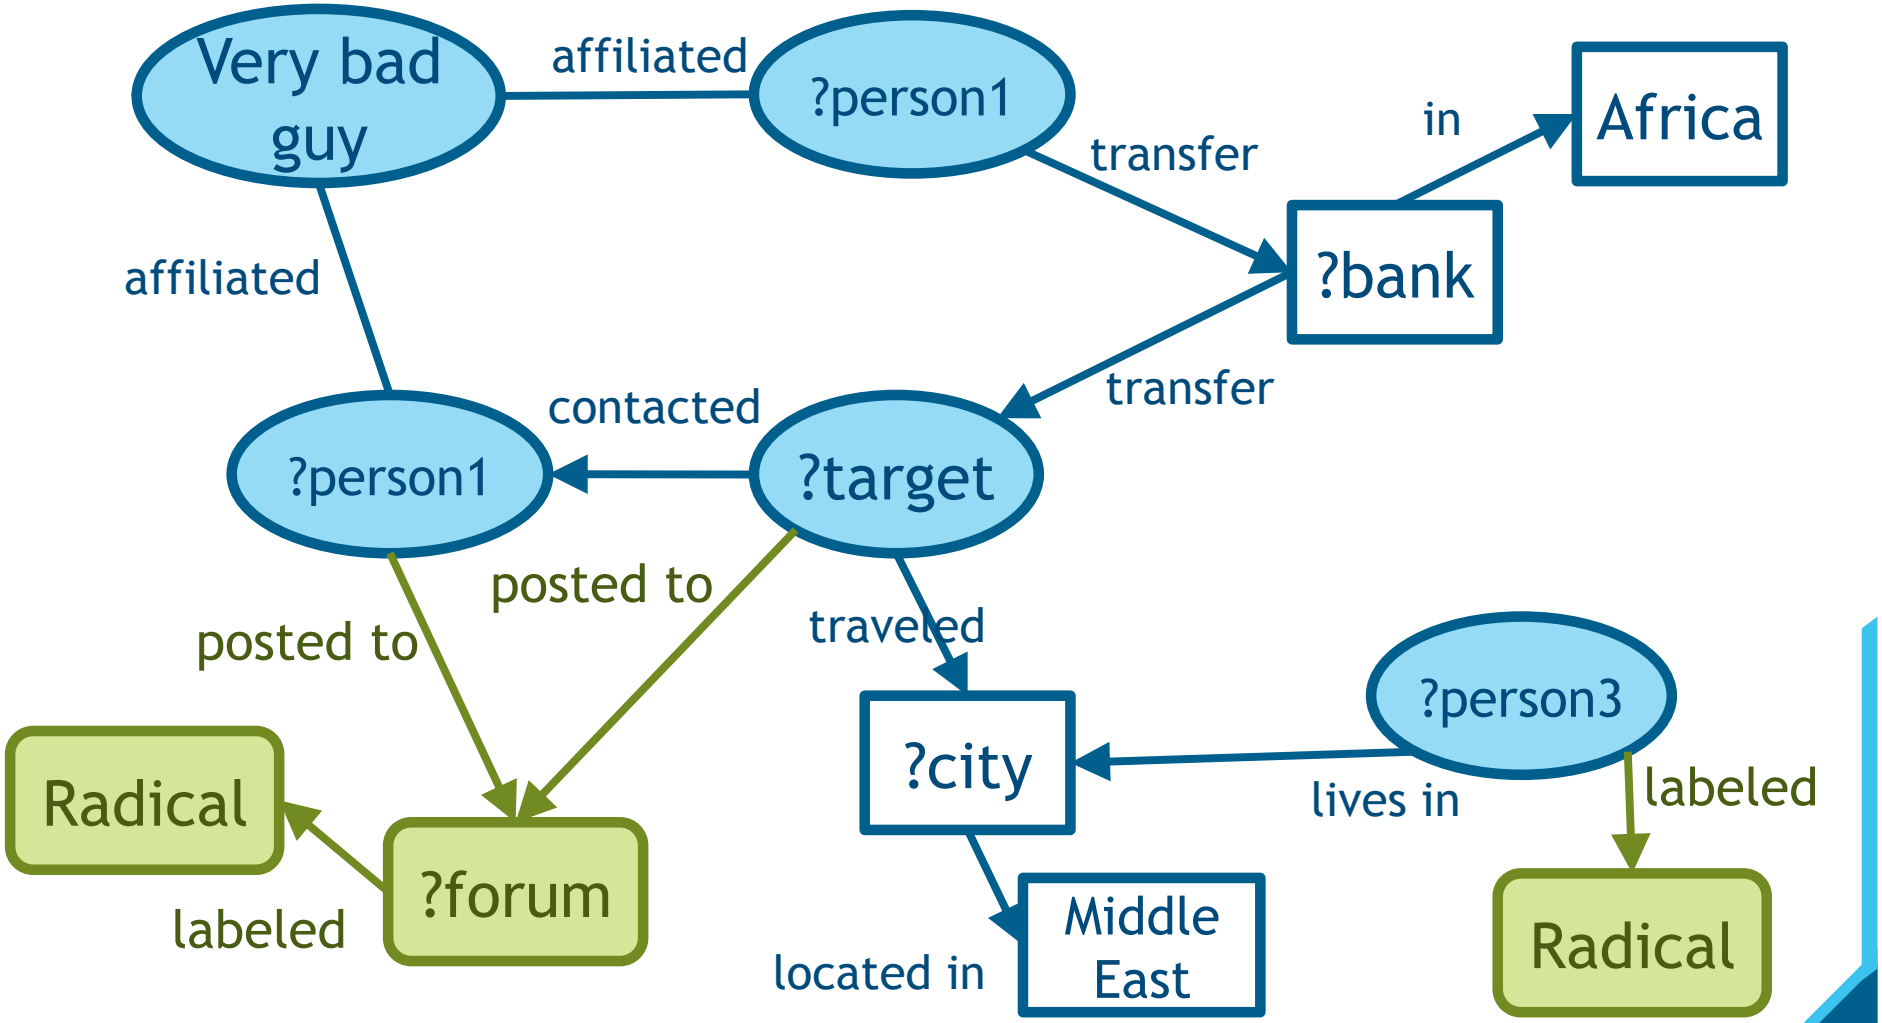

Intelligence and Security

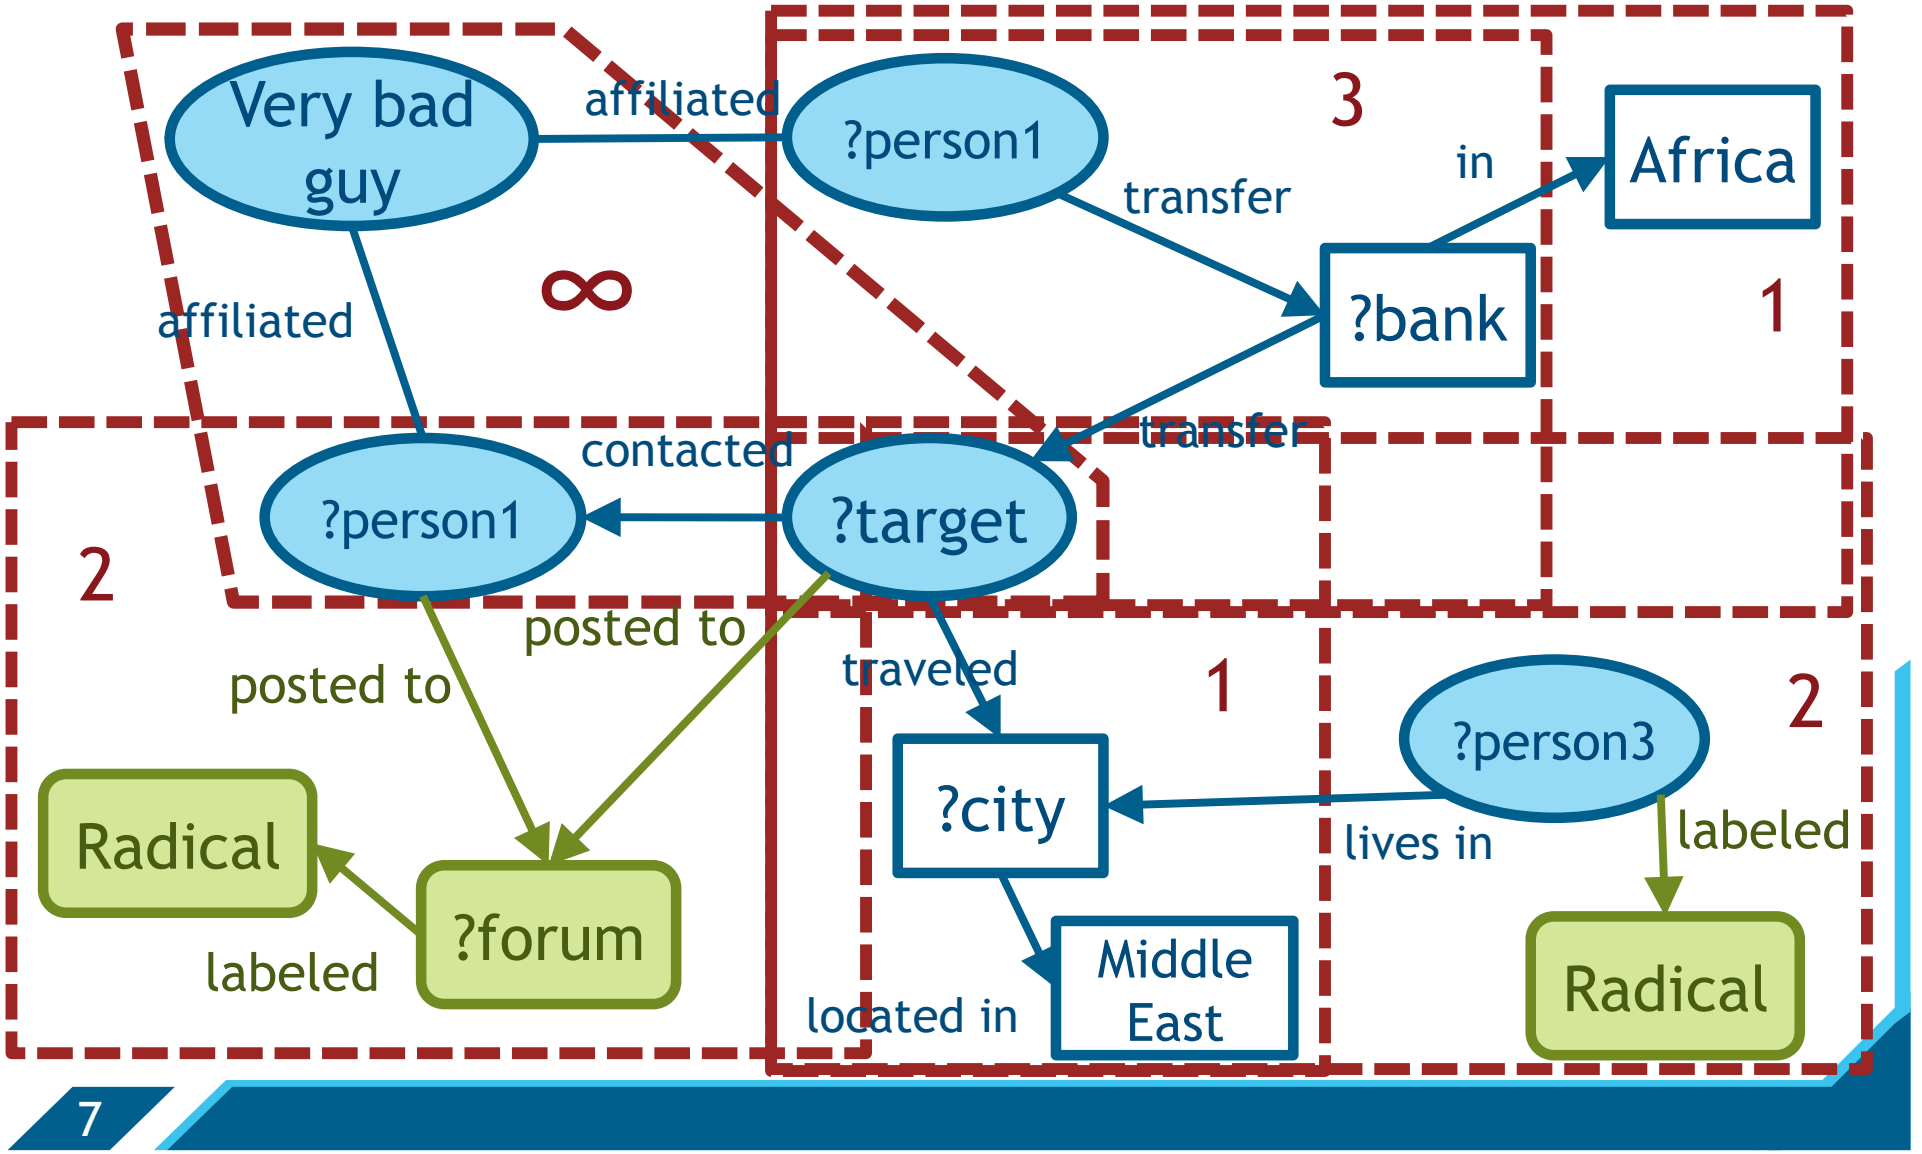

PMatch Applications

- Intelligence
  - Information retrieval on open source data
- Search
  - Facilitate productive use of social network data
- Security
  - Finding suspicious patterns
- Social Network Analysis
  - Querying as a primitive operation for data retrieval

Outline

- Motivation
- PMatch Query Definition
- PMatch Query Answering
- Experiments
- Related Work & Conclusion

PMatch Query

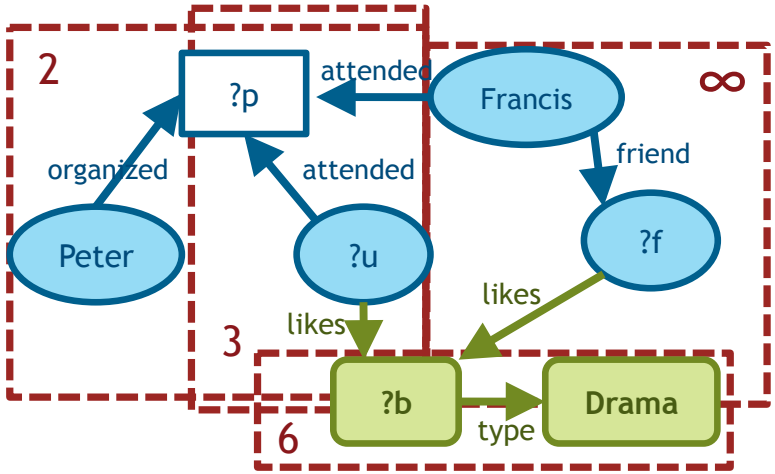

- Set of Boxes
  - Each contains edges
  - Has a weight
    - Infinity → certainty or “core” of query
- Naturally extends exact subgraph matching queries

PMatch Query

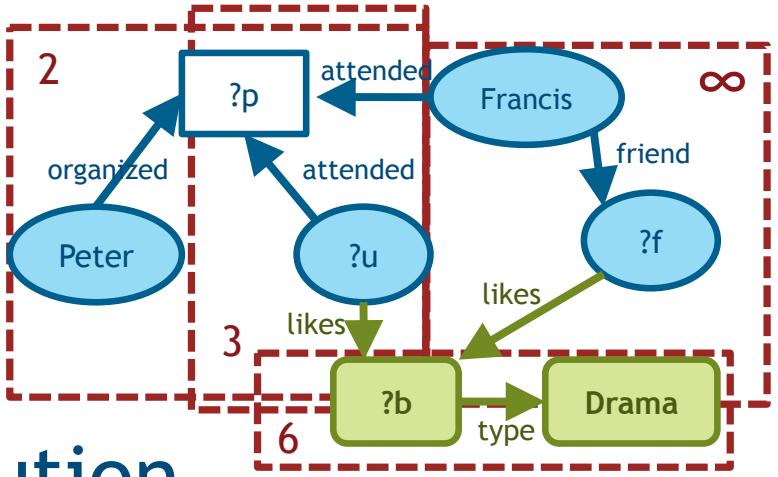

- Matching boxes
  - Let b_i = 1 if box B_i is matched by a substitution else 0 for i=1 to n (= # of boxes)
  - Weights w_i, offset w_0
- Probability of match (user defined)

  =0 if b_0 = 0, core must be matched

  $$= \frac{1}{1 + e^{-(w_0 + \sum_{i=1}^{n} w_i b_i)}}$$

PMatch Query Answer

- All substitutions with probability above user defined threshold τ
- Need to eliminate redundant answers
  - Substitutions that can be extended to higher probability

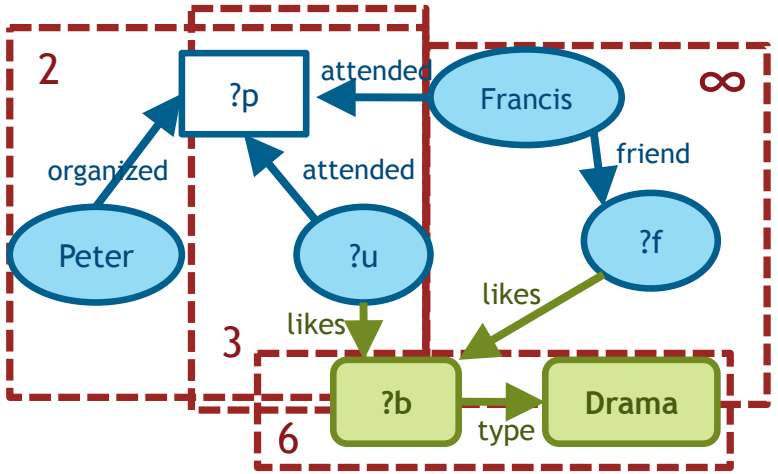

PMatch Answer

- $1/(1+e^{-(-5+6)}) \approx 73\%$
  - $w_0=-5$

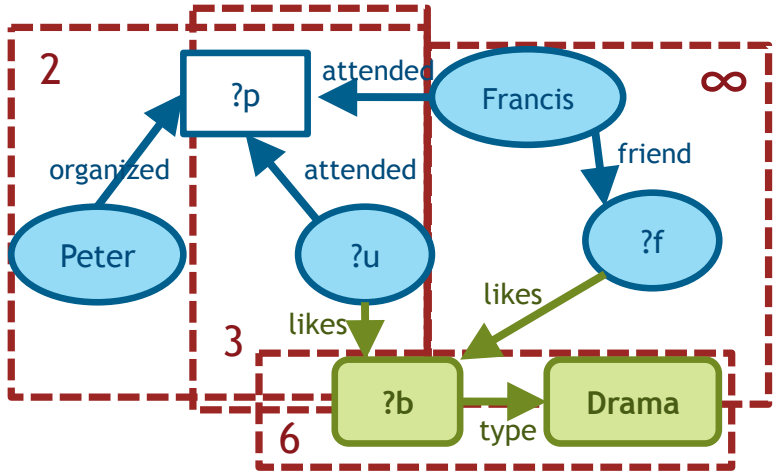

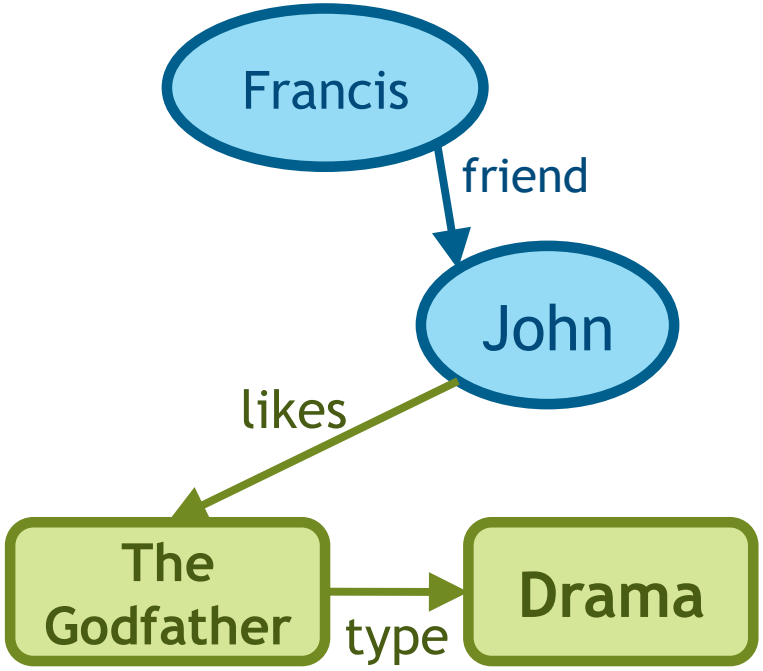

Outline

- Motivation
- PMatch Query Definition
- PMatch Query Answering
- Experiments
- Related Work & Conclusion

PMatch Query Answering

- Try to find all answers at once by exploiting shared substructures and efficiently maintain data structures.
- Depth first-search algorithm
  - Terminates when search space is exhausted or answer probability falls below threshold.
- Provably correct

PMatch Initialization

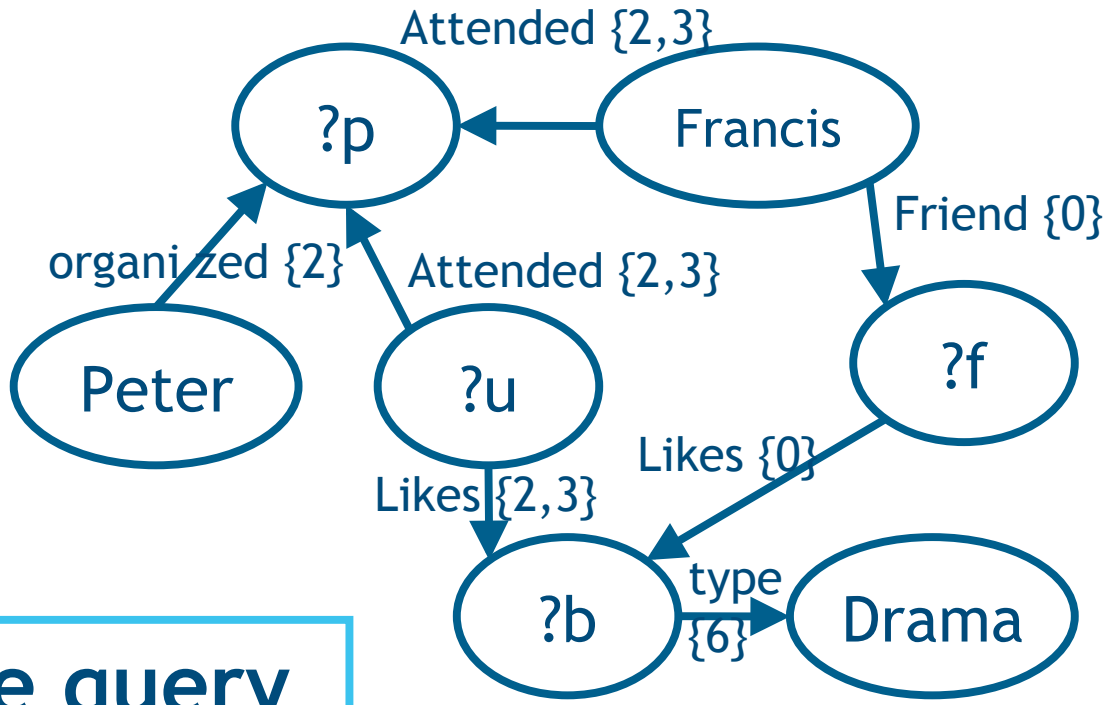

Rewrite query
(0=certain)

PMatch Data Structures

Maintained per vertex:

- R: Set of substitution candidates
  - Pairs of the form <x,H_x>, where x is a potential substitution and H_x the set of box ids it matches.
    H_x maintained as bit vector
- S: Set of box ids for which R has been initialized

PMatch Initialization

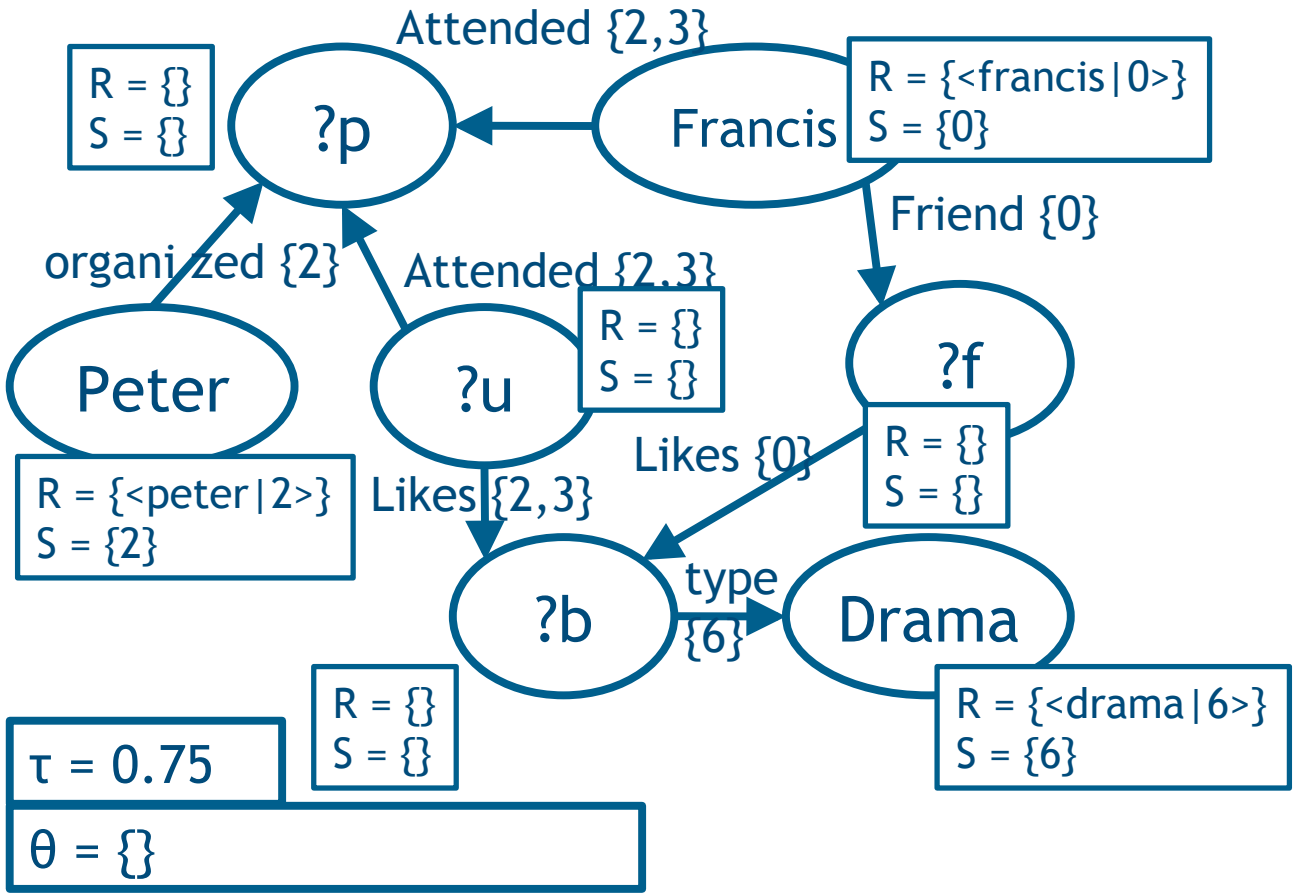

Vertex selection heuristic

- Goodness(v) = $\frac{\text{progress(v)}}{\text{cost(v)}}$
- Progress(v) = cumulative box weight of edges to be processed
  - Normalize weight by box cardinality
- Cost(v) = use standard cost models from exact subgraph matching
  - Cardinality, estimated selectivity

PMatch Algorithm

Initial vertex selection

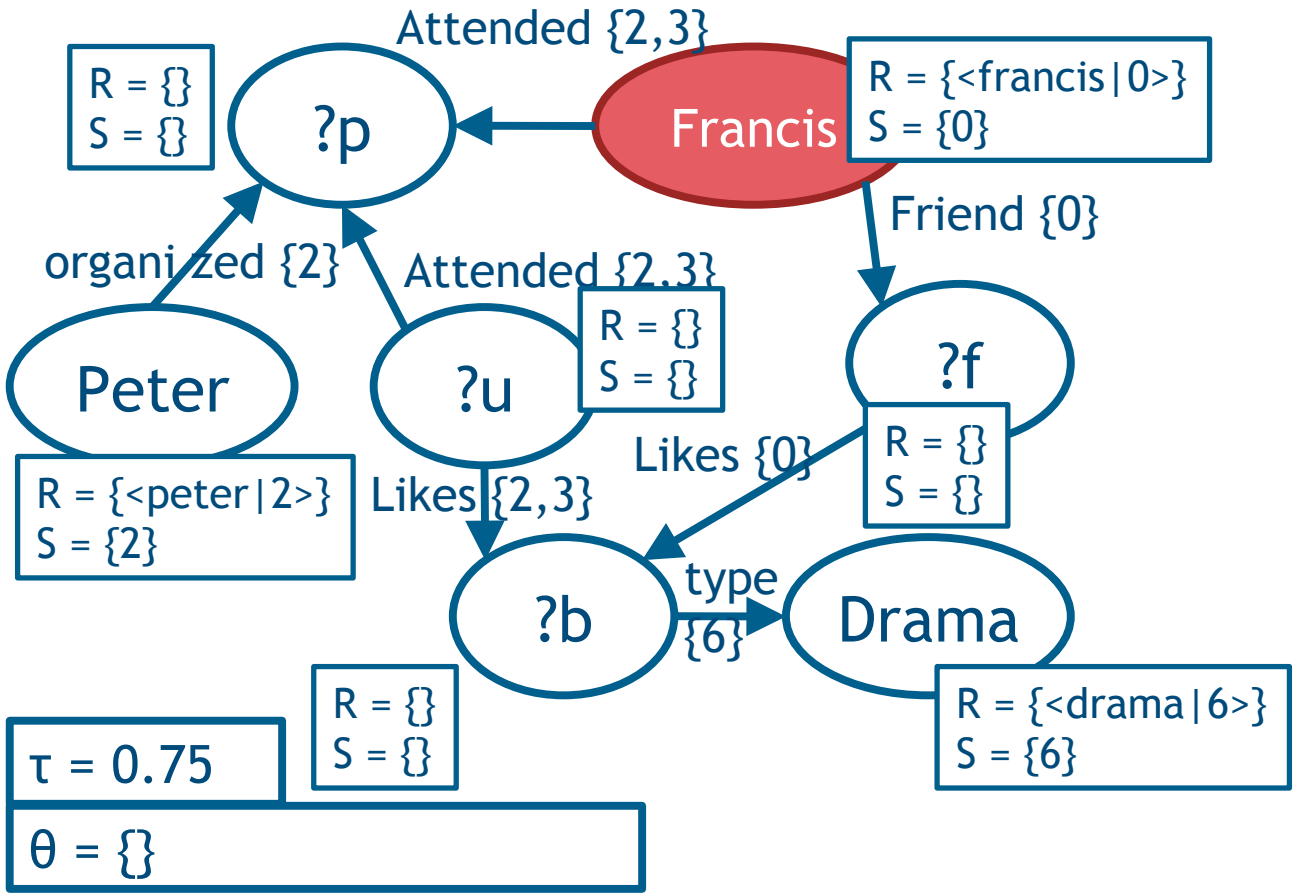

PMatch Algorithm

Processing of „Francis“ vertex

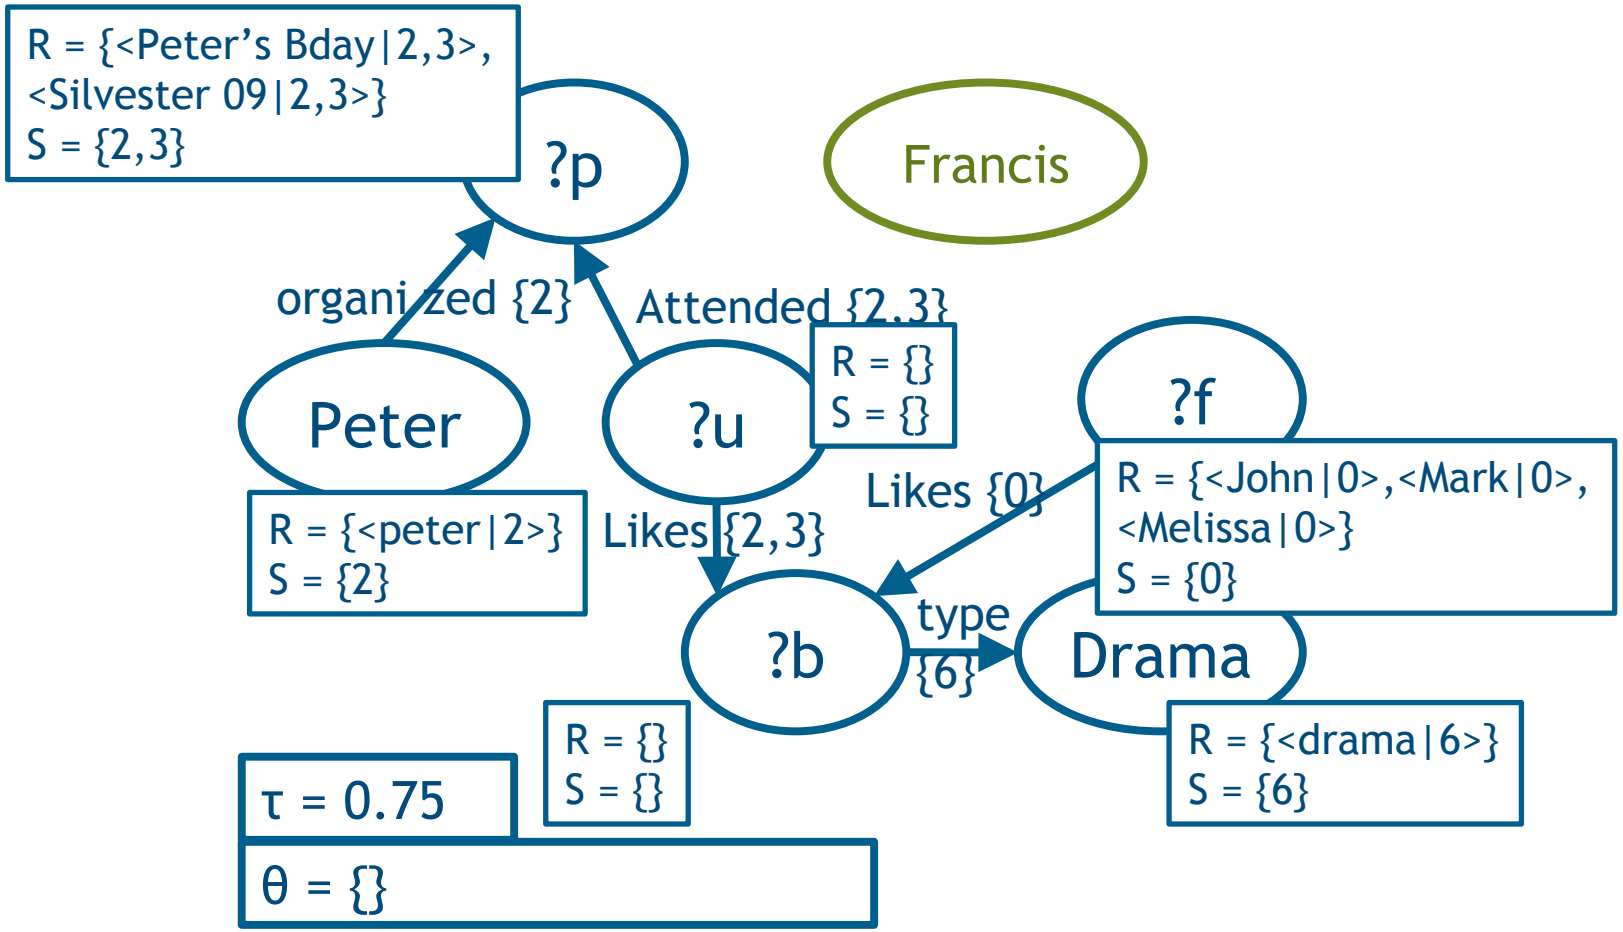

PMatch Algorithm

Threshold Pruning

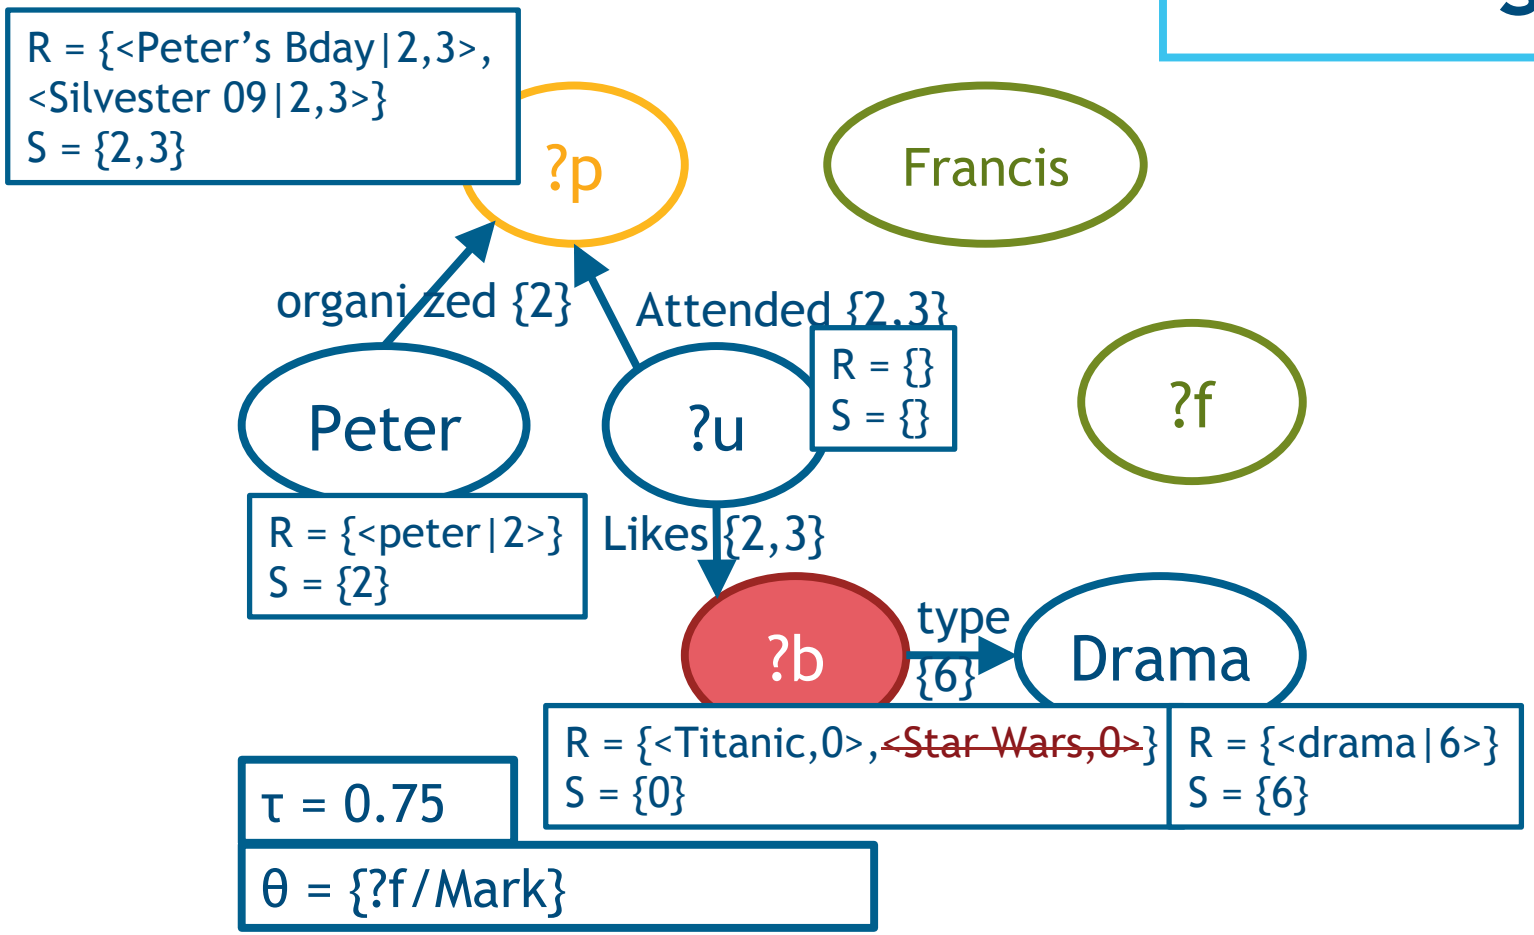

PMatch Algorithm

Processing vertex „?b“. Maintaining box indices

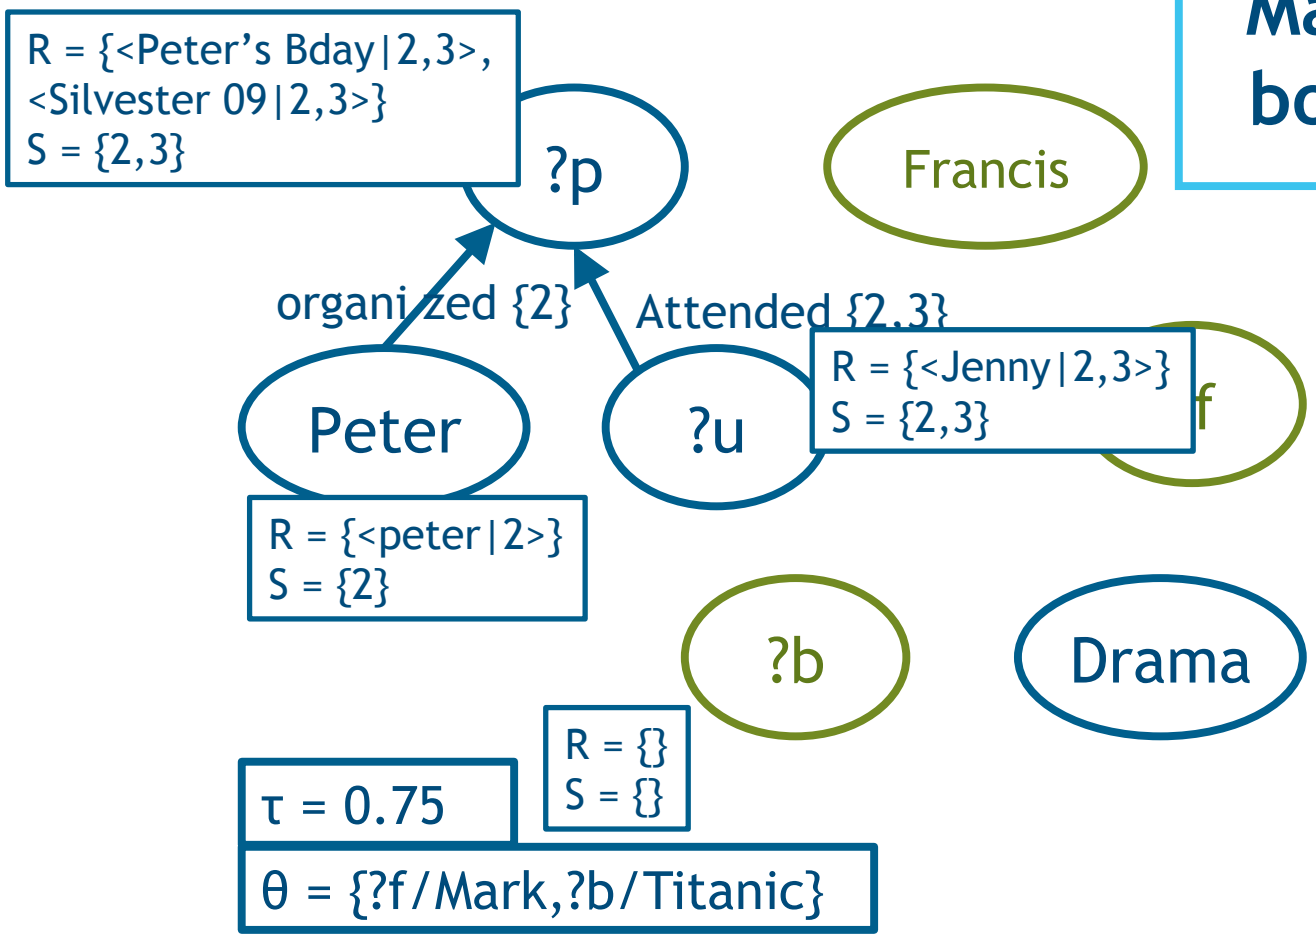

PMatch Algorithm

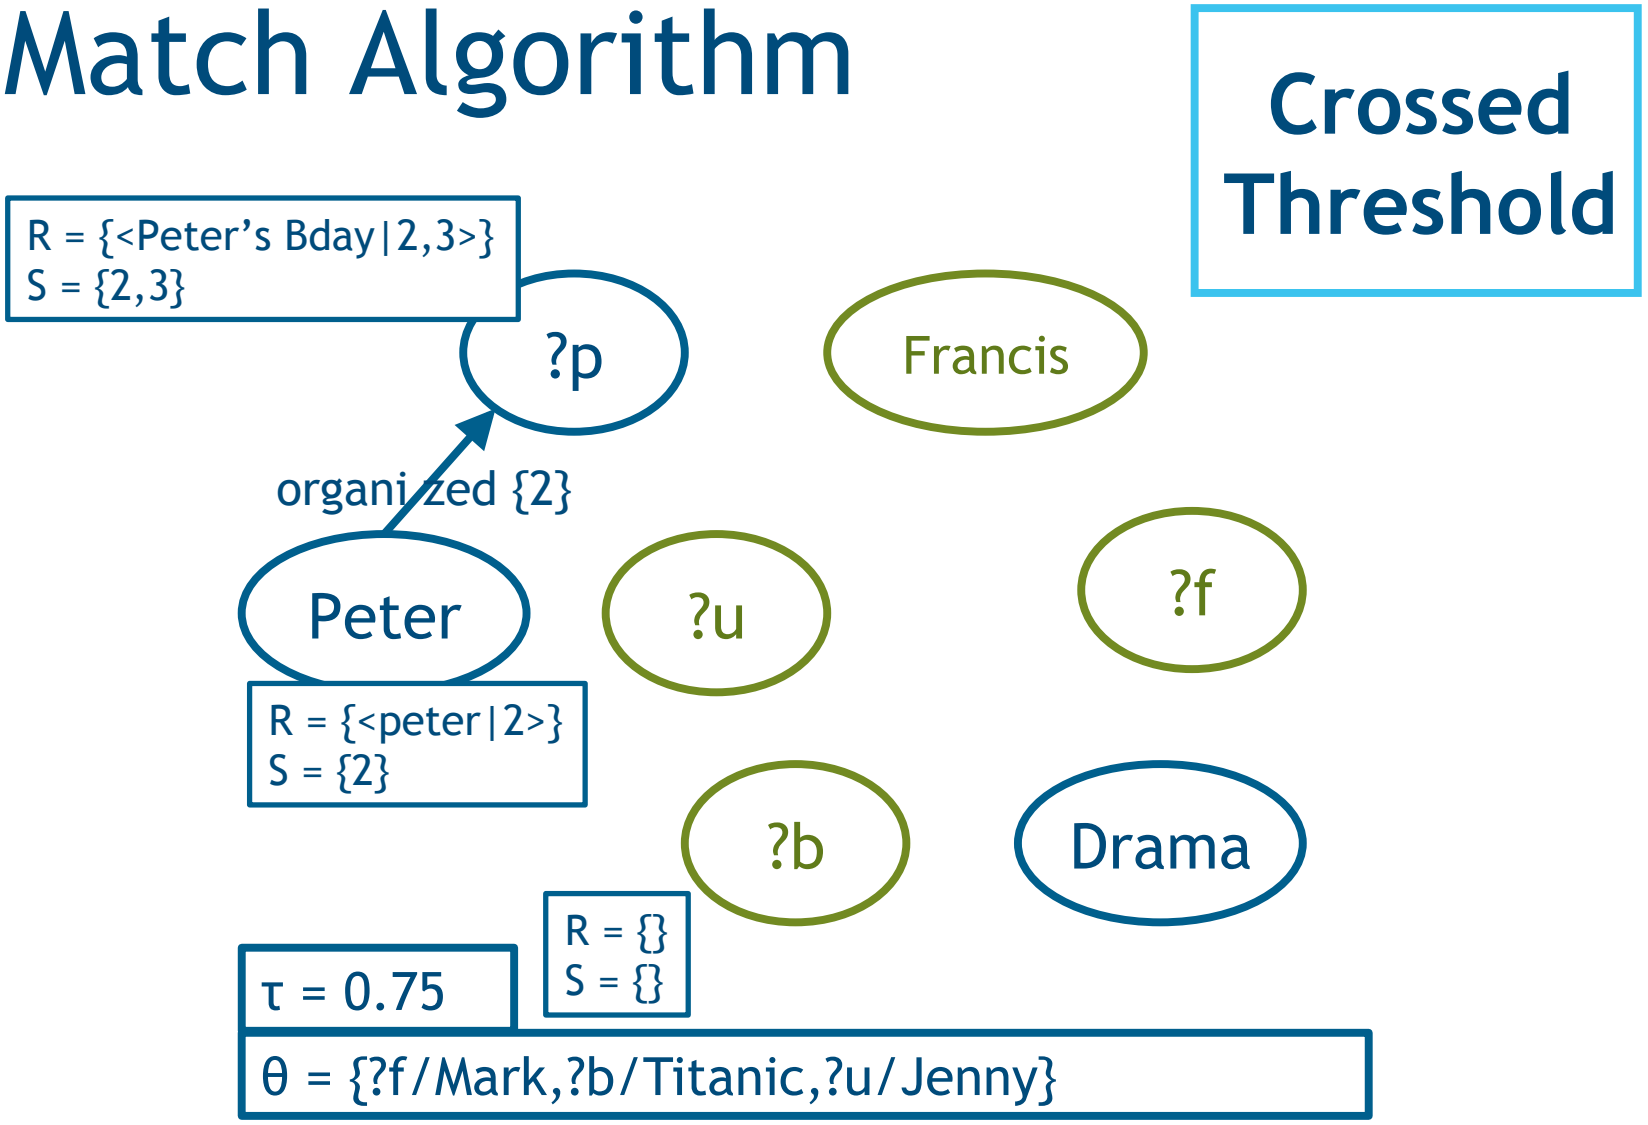

PMatch Algorithm

Found Result

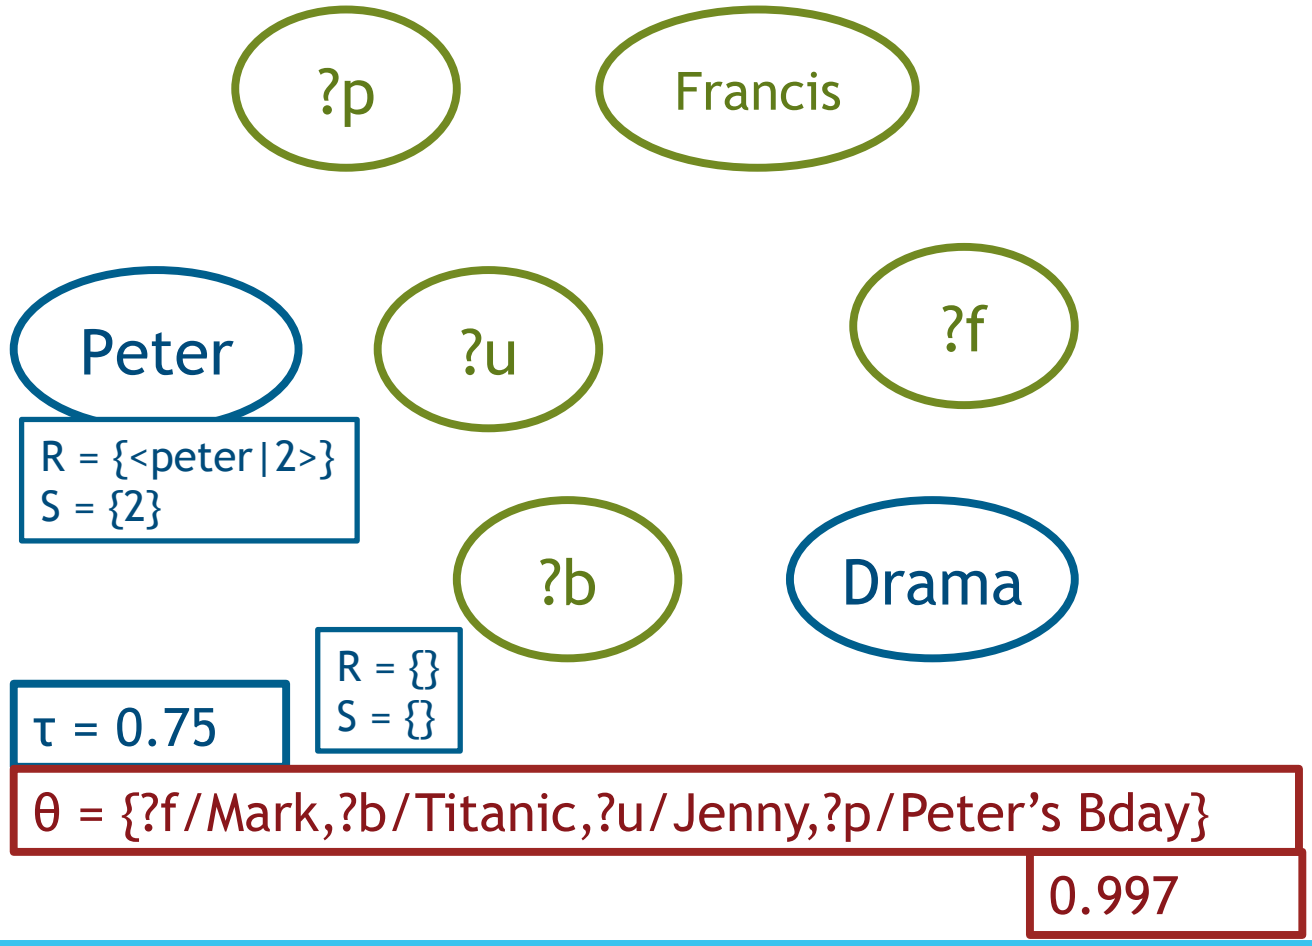

Query Splitting

Selecting „Peter“ first causes query to be split.

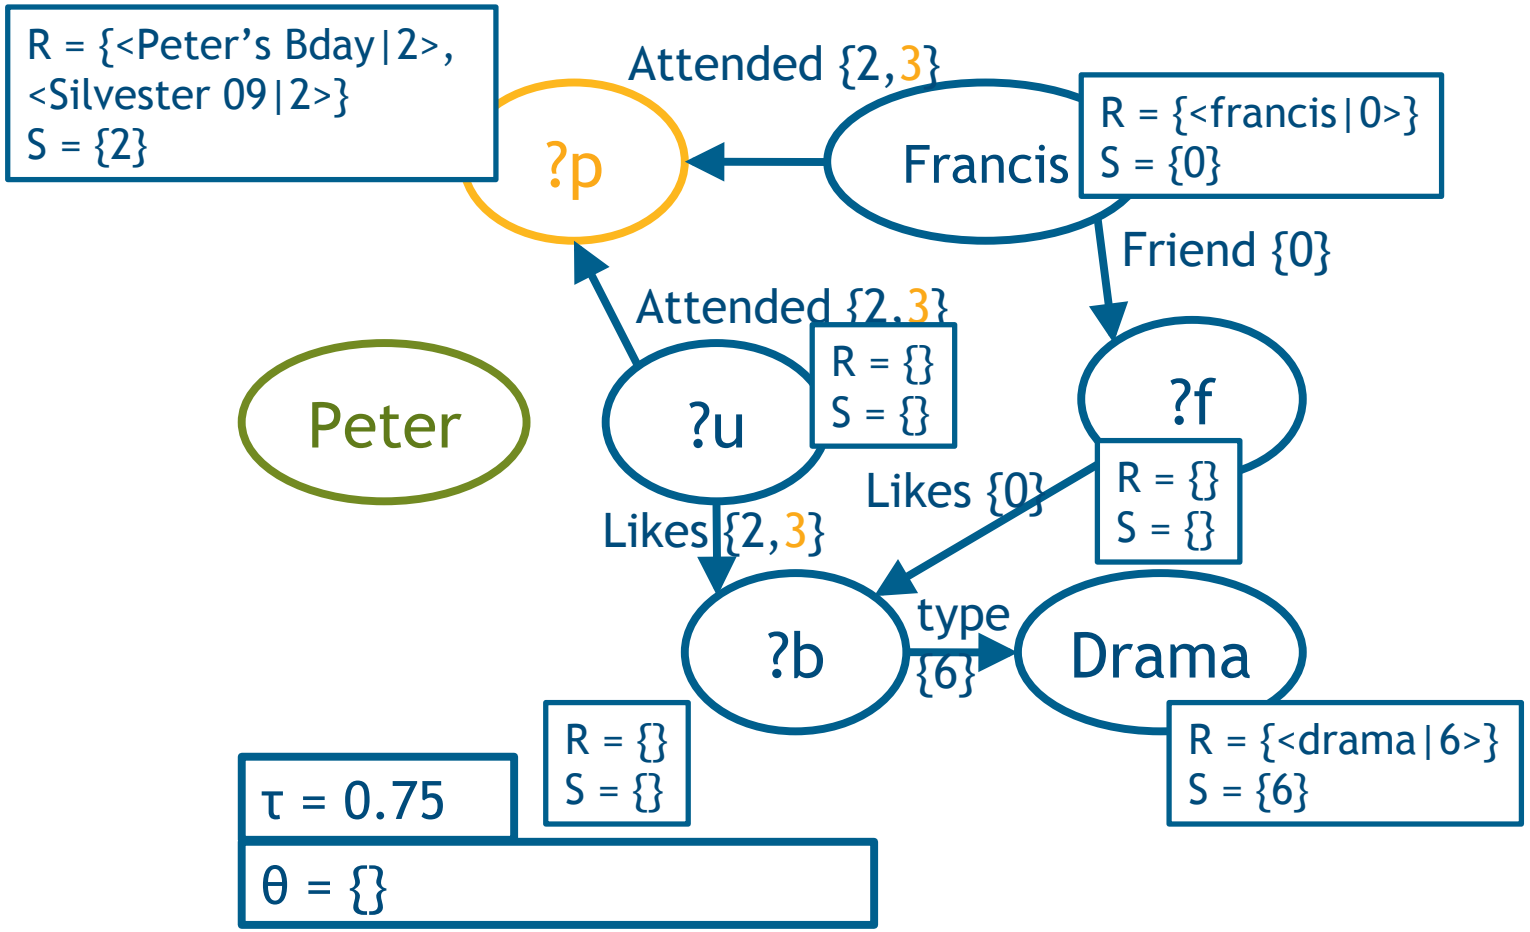

Outline

- Motivation
- PMatch Query Definition
- PMatch Query Answering
- Experiments
- Related Work & Conclusion

Experimental Setup I

- Implementation in Java (approx 5500 loc) on top of Neo4j
- Datasets
  - Friendship Network (Flickr, Orkut, Livejournal, Youtube): 778 million edges
  - Delicious Network: 1.1 billion edges
- Query Benchmark for each dataset: 9 PMatch queries each of increasing complexity

Experimental Setup II

- Compared
  - PMatch-1: box weight
  - PMatch-2: normalized box weight
  - DOGMA: Exact matching for all combination
  - Neo4j: Exact matching for all combinations
    - Too slow – not shown in the charts
- System: AMD Opteron, 15K RPM 300 GB SAS HD, 2GB of heap space

Friendship network (small queries)

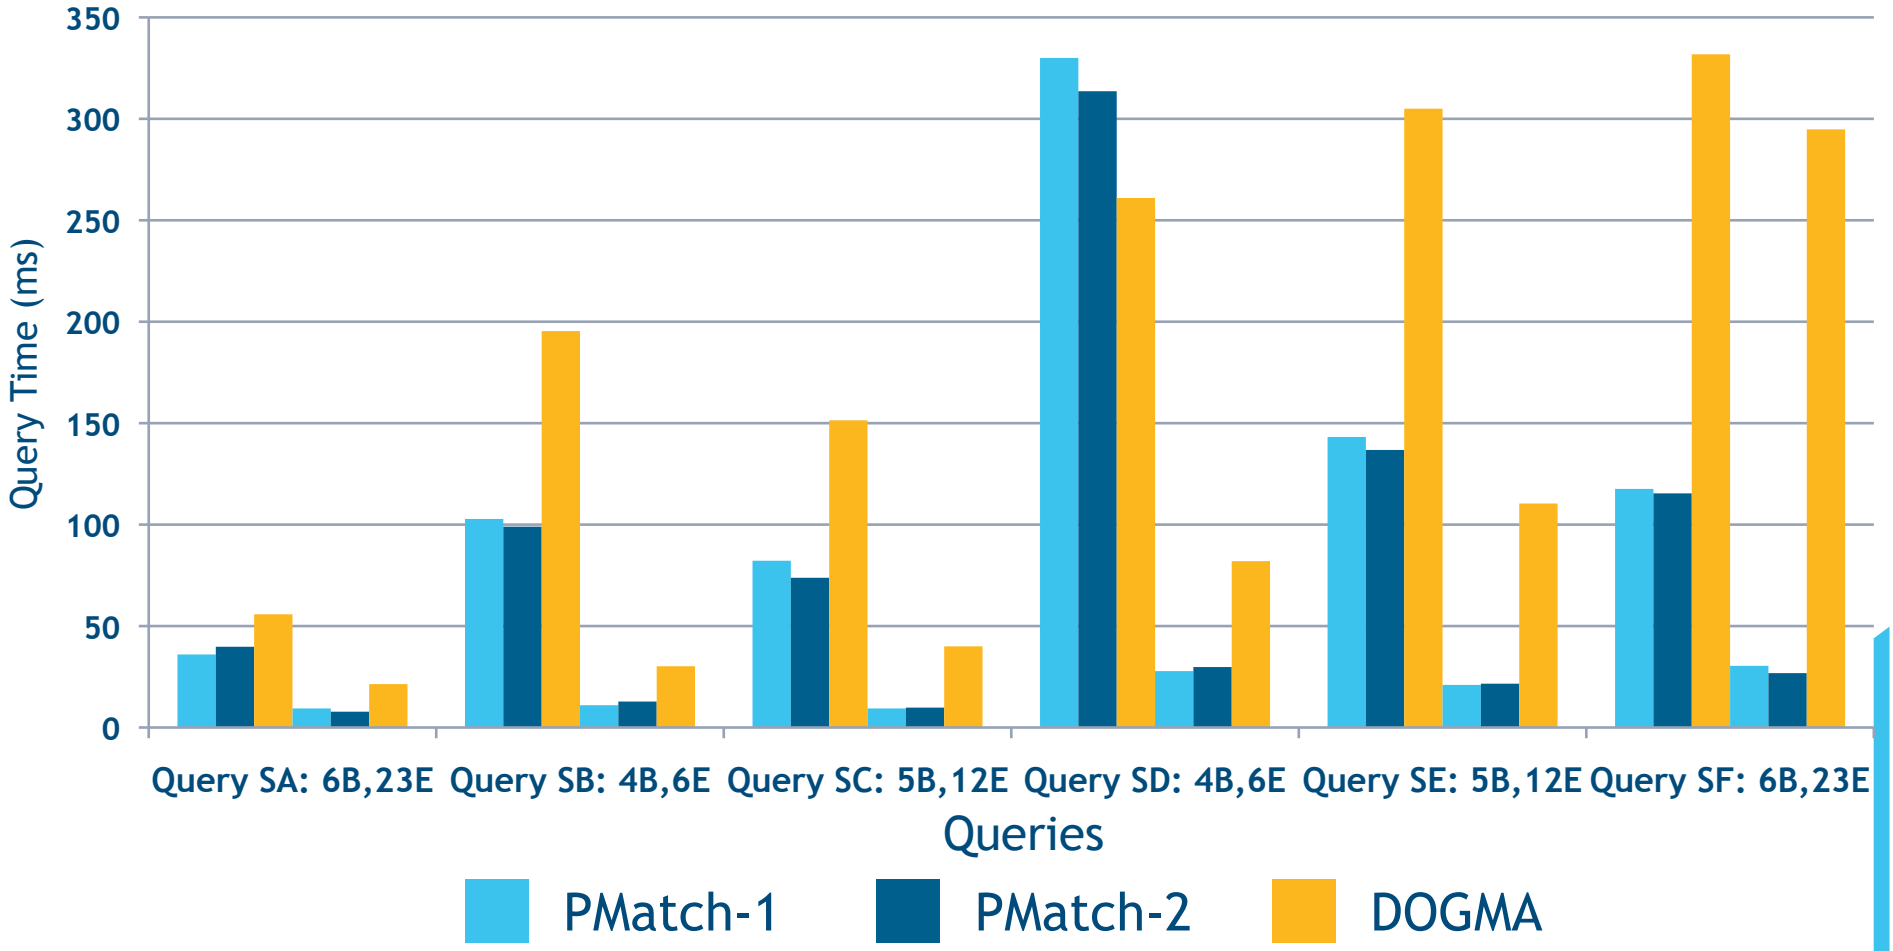

For each query, we report query times with cold (right) and warm (left) caches.

Friendship network (large queries)

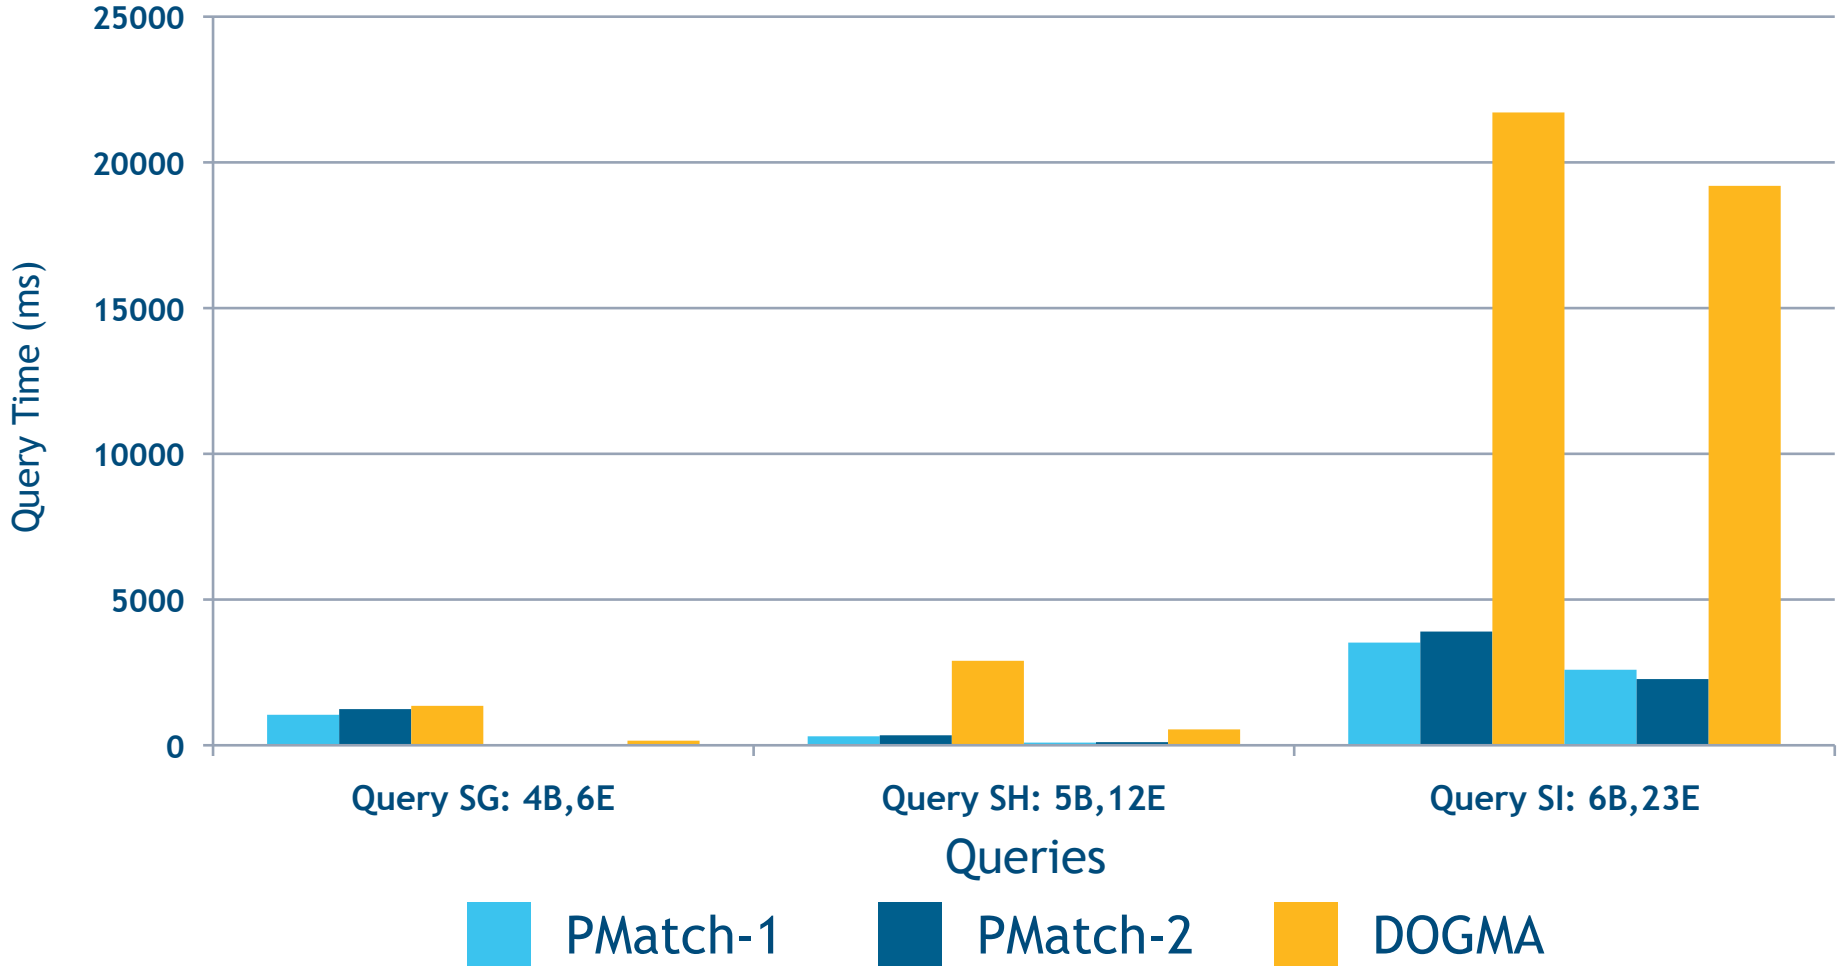

For each query, we report query times with cold (right) and warm (left) caches.

Delicious network (small queries)

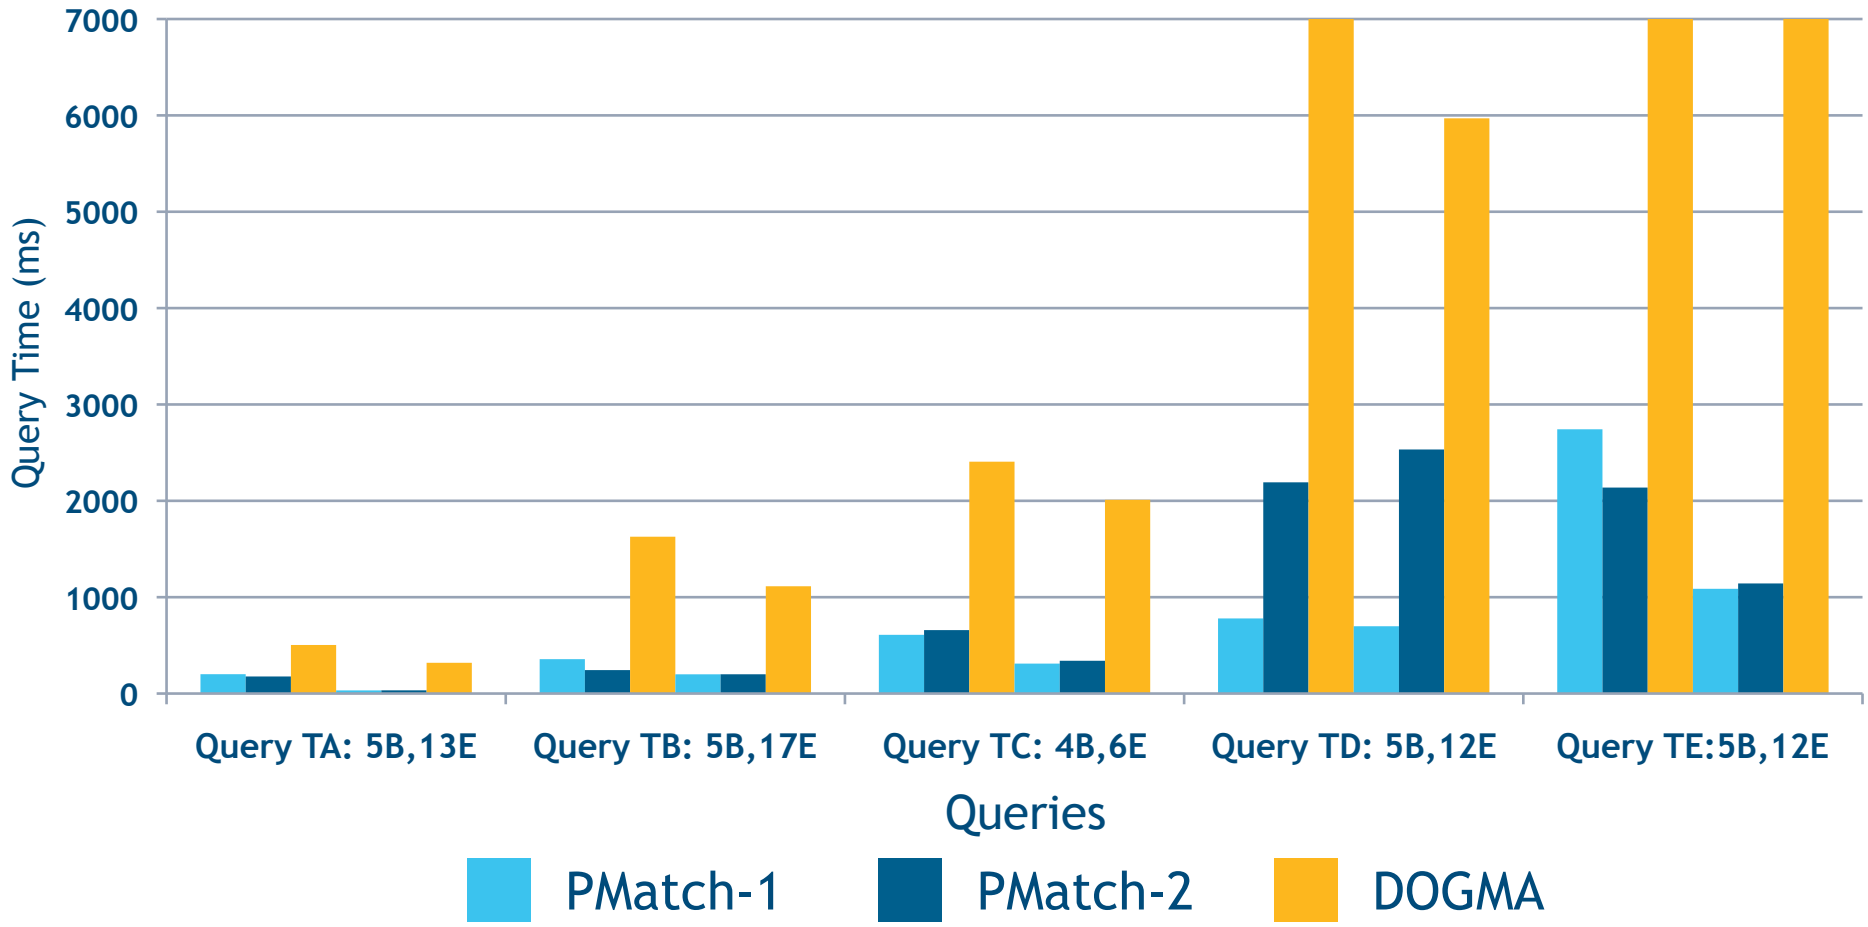

For each query, we report query times with cold (right) and warm (left) caches.

Delicious network (large queries)

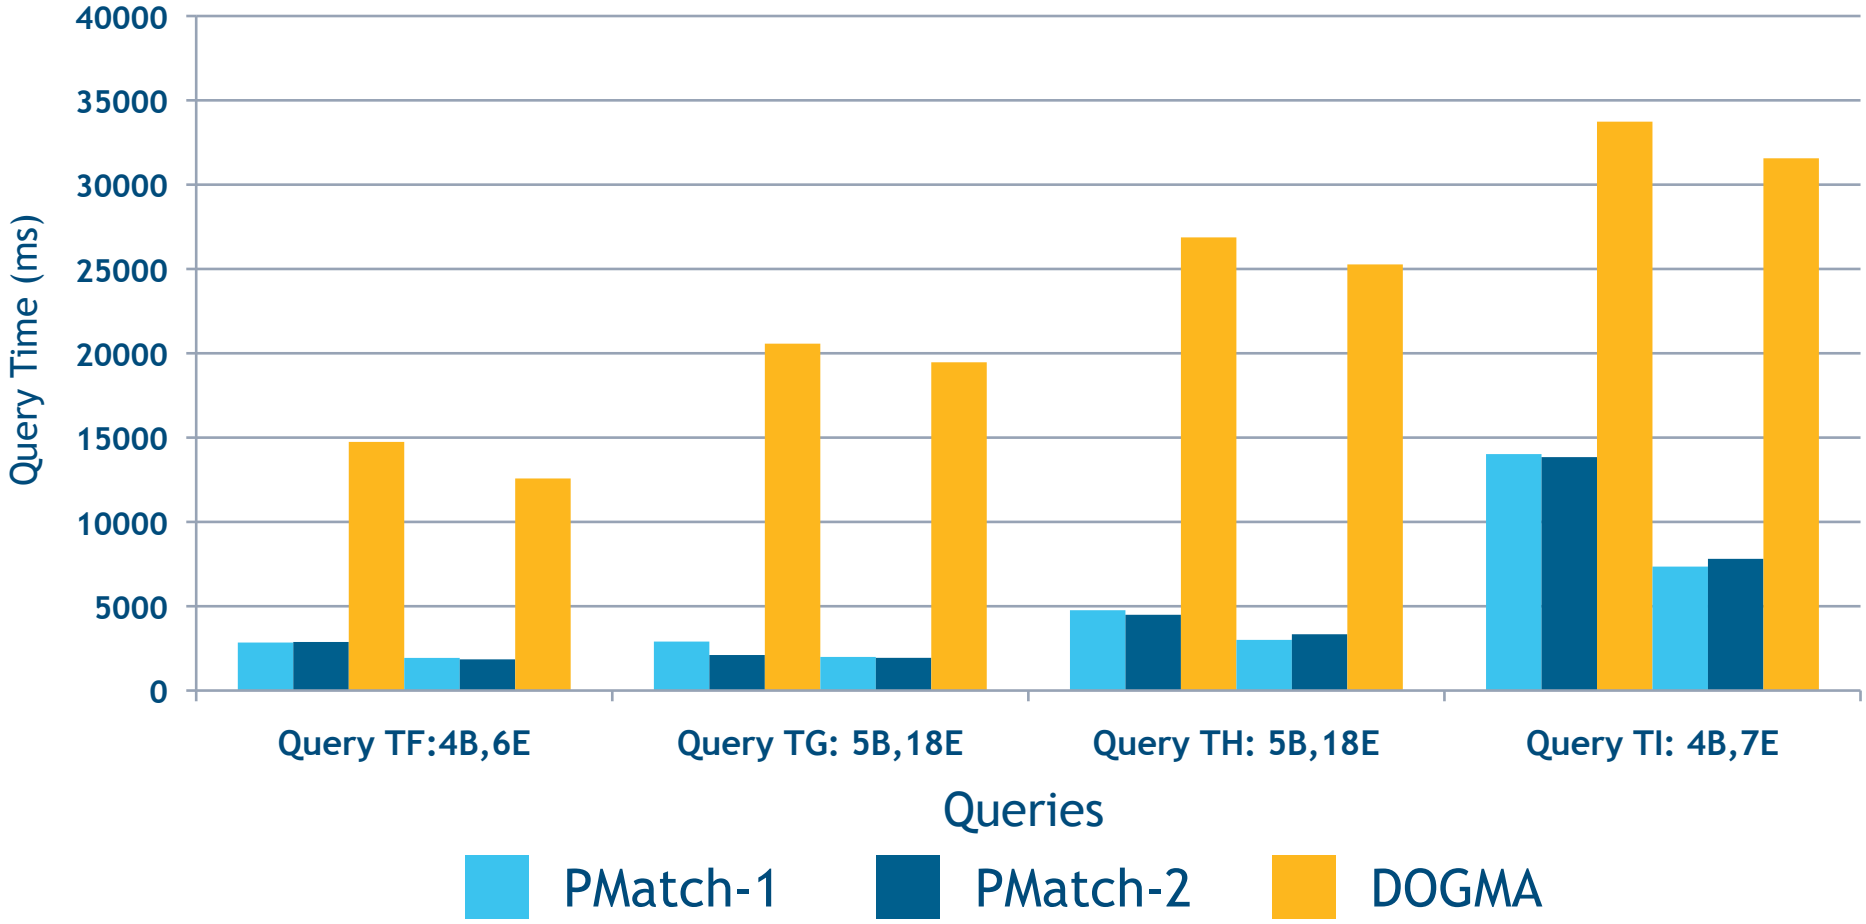

For each query, we report query times with cold (right) and warm (left) caches.

Outline

- Motivation
- PMatch Query Definition
- PMatch Query Answering
- Experiments
- Related Work & Conclusion

Related Work I

- Exact Subgraph Matching
  - Many systems: RDF-3X, OWLIM, YARS (2), Sesame, Jena, DOGMA, COSI etc.
  - Efficient on-disk index structures
  - Query answering algorithms focused on exact queries (e.g. SPARQL).
- No treatment of query uncertainty

Related Work II

- Approximate Subgraph Matching
  - SAGA, TALE, Grafil, etc
  - Database containing many small graphs
  - Do not scale to large social networks
  - Query uncertainty is defined in the system.
- PMatch provides an expressive language to specify probabilistic queries and answers them efficiently.

Conclusion

- Query uncertainty often occurs in social network search, information retrieval and pattern matching.
- PMatch provides a language to concisely express and algorithms to efficiently answer probabilistic subgraph matching queries.

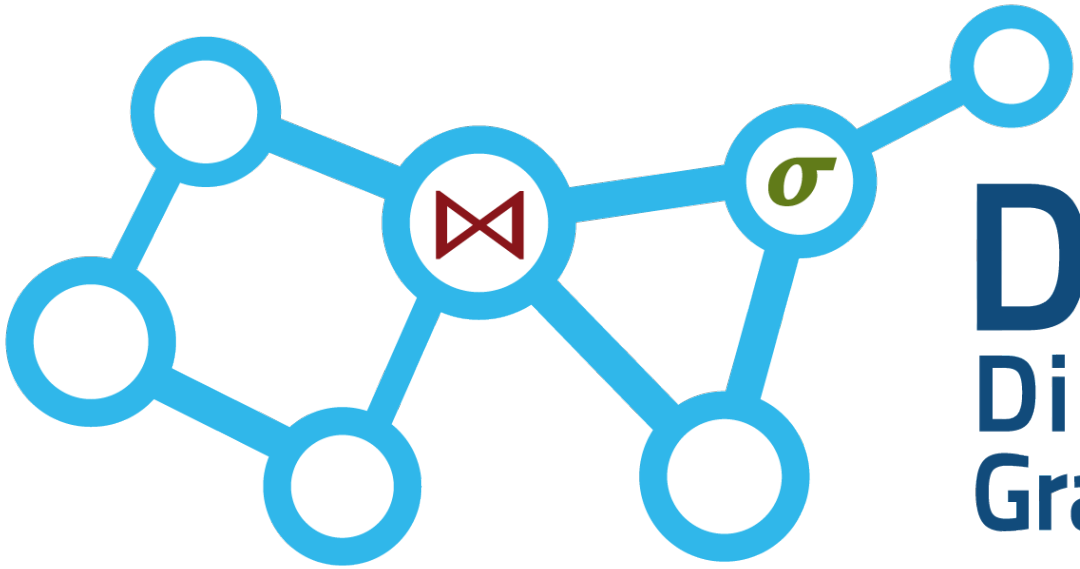

DOGMA

Disk Oriented Graph Matching

dogma.umiacs.umd.edu

Questions?
Comments?

Related Work III

- Data Uncertainty
  - Specified at the tuple or edge level
  - How to answer certain (SQL) queries over uncertain data and what are the answer probabilities?

First Name	Last Name	Income	Probability
John	Smith	80,000	0.6
Alice	Baker	85,000	0.5
Jennifer	Temper	90,000	0.9

Related Work III

- Data Uncertainty
  - How to encode dependencies?
  - Need to make strong independence assumptions to be efficient
  - Where are the probabilities coming form?
- Instead, we assume certain data and model uncertainty in the user specified query.

Query Uncertainty Motivation

- User Uncertainty
  - User does not know what she is looking for exactly
- Lack of schema
  - Lack of schema complicates query design
- Data heterogeneity
  - Can be caused by data integration
- Noisy or Missing Data
  - It’s real world data