



SIGGRAPH2015

Xroads of Discovery





SIGGRAPH2015
Xroads of Discovery

The 42nd International Conference and Exhibition
on Computer Graphics and Interactive Techniques



P I X A R
SOFTWARE R&D

Presto Execution System

An Asynchronous Computation Engine for
Animation

George ElKoura
Pixar


More Info



@multithreadvfx



www.multithreadingandvfx.org



“Multithreading for Visual Effects”



Team

Engineering

Florian Zitzelsberger Florian Sauer John Loy

Thanks

Steve May Guido Quaroni Adam Woodbury
Dirk Van Gelder Alicia Mooty

Pixar Software R&D



Overview

- What is Presto?
- Execution system architecture
- Multithreading strategies
- Switching a large code base to TBB



Presto Execution

- The Presto Execution System:
 - A general purpose execution engine
 - Commonly used for posing points
 - Must be as fast as possible




Phases of Execution

- Execution is broken up into three phases:
 - Compilation
 - Scheduling
 - Evaluation




Compilation

User-Authored Scene Objects




Compilation

Compiled Data Flow Network



Scheduling

Compiled Network




+

Client Request

Request
Body points
Body transform
Head points
Head transform


Scheduling

Schedule




Evaluation

Compiled Network



Schedule



+

Evaluation




Results

Computed Values

```
(1.4, 0.0, 2.4)
(1.4, -0.784, 2.4)
(0.784, -1.4, 2.4)
(0.0, -1.4, 2.4)
(1.3375, 0.0, 2.53125)
(1.3375, -0.749,
2.53125)
(0.749, -1.3375,
2.53125)
(1.4375, 0.0, 2.53125)
(1.4375, -0.805,
2.53125)
(0.805, -1.4375,
2.53125)
(0.0, -1.4375, 2.53125)
(1.5, 0.0, 2.4)
...
```



Exec Engine Architecture



User Extensions



- Plugin support
- Need painless multithreading
- Provide safeguards to help



Execution Callbacks

```
static void  
MyCallback(const Context &context)  
{  
    ...  
}
```



Execution Callbacks

- Iterators are provided for vectorized data access...
- which allows for easier experimentation with multithreading strategies.



User Gotchas

- Avoid global and static variables
- Be careful about singletons
- Using a non-thread-safe 3rd party library requires users to do their own locking
- ... and, of course, Python



Python in Execution



- Global Interpreter Lock (GIL)
- Only one thread at a time
- Do heavy lifting in C++
- Look out for GIL deadlocks



Flexibility to Experiment

- Modular, adaptable design, e.g.:
 - Hardware changes: rewrite executor, possibly scheduler
 - User representation changes: rewrite compilation
- Allows for easy experimentation with different algorithms





Multithreading



- Computed data not in the network
- Stateless computations
- Modular schedulers and executors





Per Node




Per Branch



Per Model



Per Frame



Background Execution

- Evaluates entire frames at different times in parallel and in the background
- Fills in a computed value cache



Interruption

- If values change, we must interrupt, but...
- We can't block to wait for it to interrupt.



Interruption

- Avoid `pthread_cancel` and `pthread_kill`
- Instead, we use a cooperative scheme...
- and throw away results.



TBB



Switching to TBB

- We had our own threading toolkit
- Built on for a decade or two
- Lacked modern features
- Wanted to switch to TBB



libWork

1. Create abstraction
2. Convert clients
 - Don't disrupt shipping software
3. Implement TBB backend
4. Flip the switch



libWork Abstractions

- Parallel For
- Fixed Dispatcher
- Dynamic Dispatcher
- Background Job



Switching to TBB is Easy!

```
#define WORK_USE_TBB 01
```



Unexpected Deadlocks

```
void F()  
{  
    // do something ...  
    A();  
    // do more things ...  
}
```



Unexpected Deadlocks

```
A () {  
    Lock a;  
    if (cache.Lookup())  
        return hit;  
    parallel_for([]{ fill in cache } )  
    return hit;  
}
```




Unexpected Deadlocks

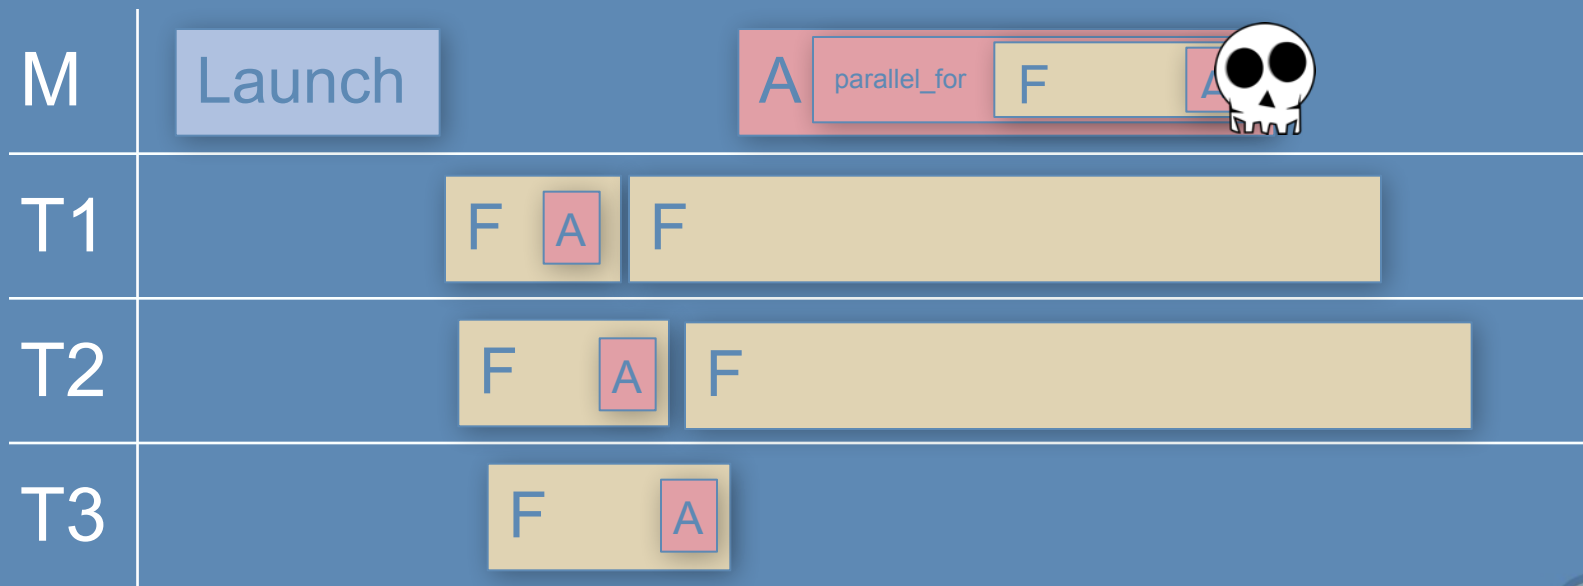
```
int main() {  
    ...  
    Spawn_many_tasks_running_F();  
    A();  
    Wait_for_Fs();  
    ...  
}
```



Without Work Stealing



With Work Stealing



Unexpected Deadlocks

```
A () {  
    Lock a;  
    if (cache.Lookup())  
        return hit;  
    parallel_for([]{ fill in cache })  
    return hit;  
}
```

not smart!



Thread Throttling

- Had strict control
- TBB hungry
- `tbb::task_scheduler_init()`



tbb::task_scheduler_init



- Hard in embedded libs
- Stack object, must go first
- Who goes first?



In Practice

PRESTO

`tbb::task_scheduler_init`

My Fast Library



In Practice



`tbb::task_scheduler_init`

My Fast Library



In Practice



`tbb::task_scheduler_init`

My Fast Lib

`tbb::task_scheduler_init`

Your Fast Lib



In Practice



`tbb::task_scheduler_init`

`tbb::task_scheduler_init`


My Fast Lib

`tbb::task_scheduler_init`

Your Fast Lib



All on one machine...



What we're doing

- Don't oversubscribe the farm
- Centralize decision making
- So far okay, not great



Recap

- Separate your data from your algorithm
- Build modular systems
- TBB is great
- Think in tasks, not threads
- Careful with resource management




THANKS!



@multithreadvfx



www.multithreadingandvfx.org



“Multithreading for Visual Effects”

