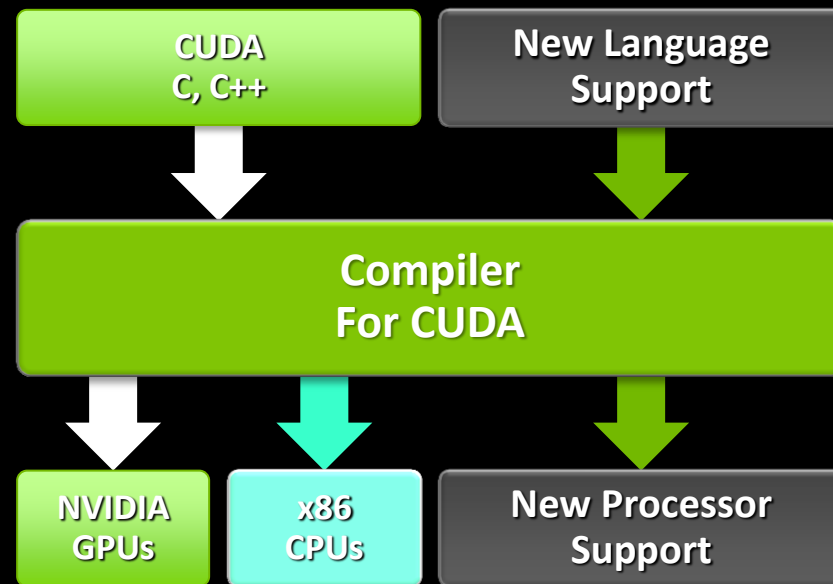


Building GPU Compilers with libNVVM

Yuan Lin

Vision

- Build a platform for GPU computing around foundations of CUDA.
 - Bring other languages to GPUs
 - Enable CUDA for other platforms
- Make that platform available for ISVs, researchers, and hobbyists
 - Create a flourishing eco-system

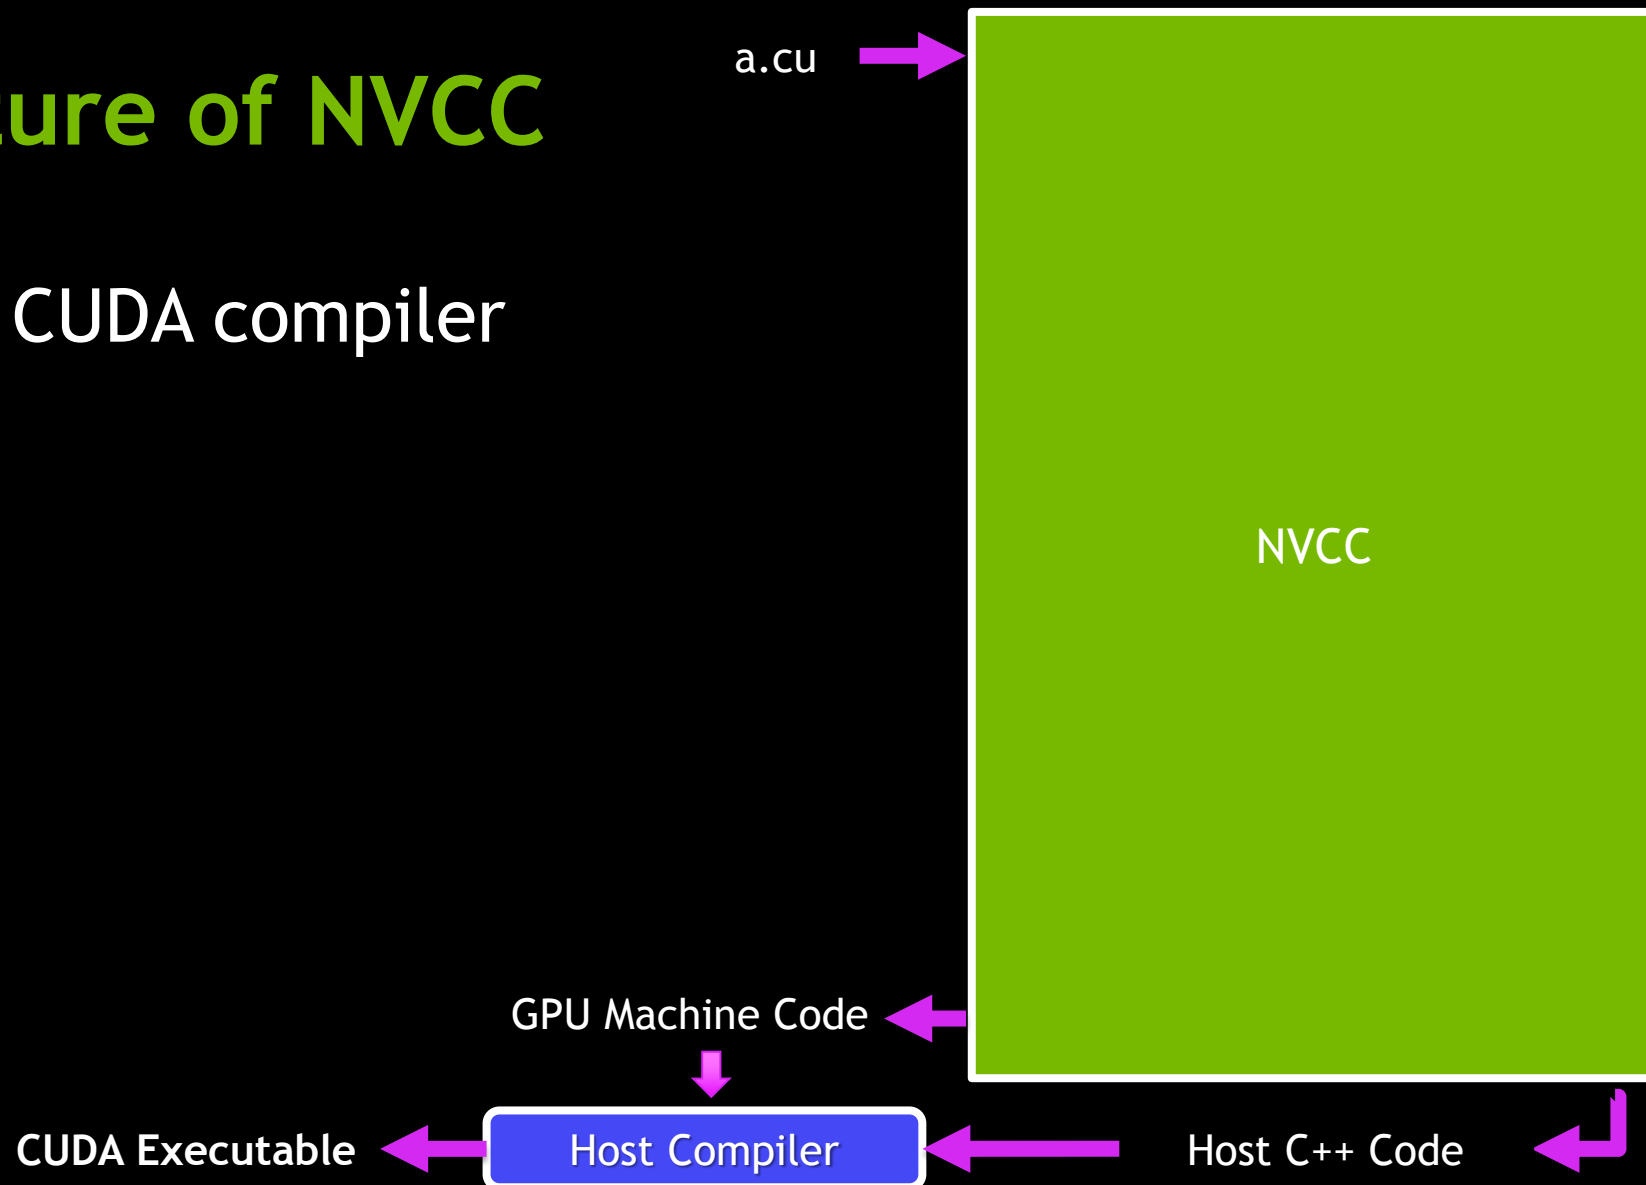


Structure of NVCC

- NVCC: CUDA compiler

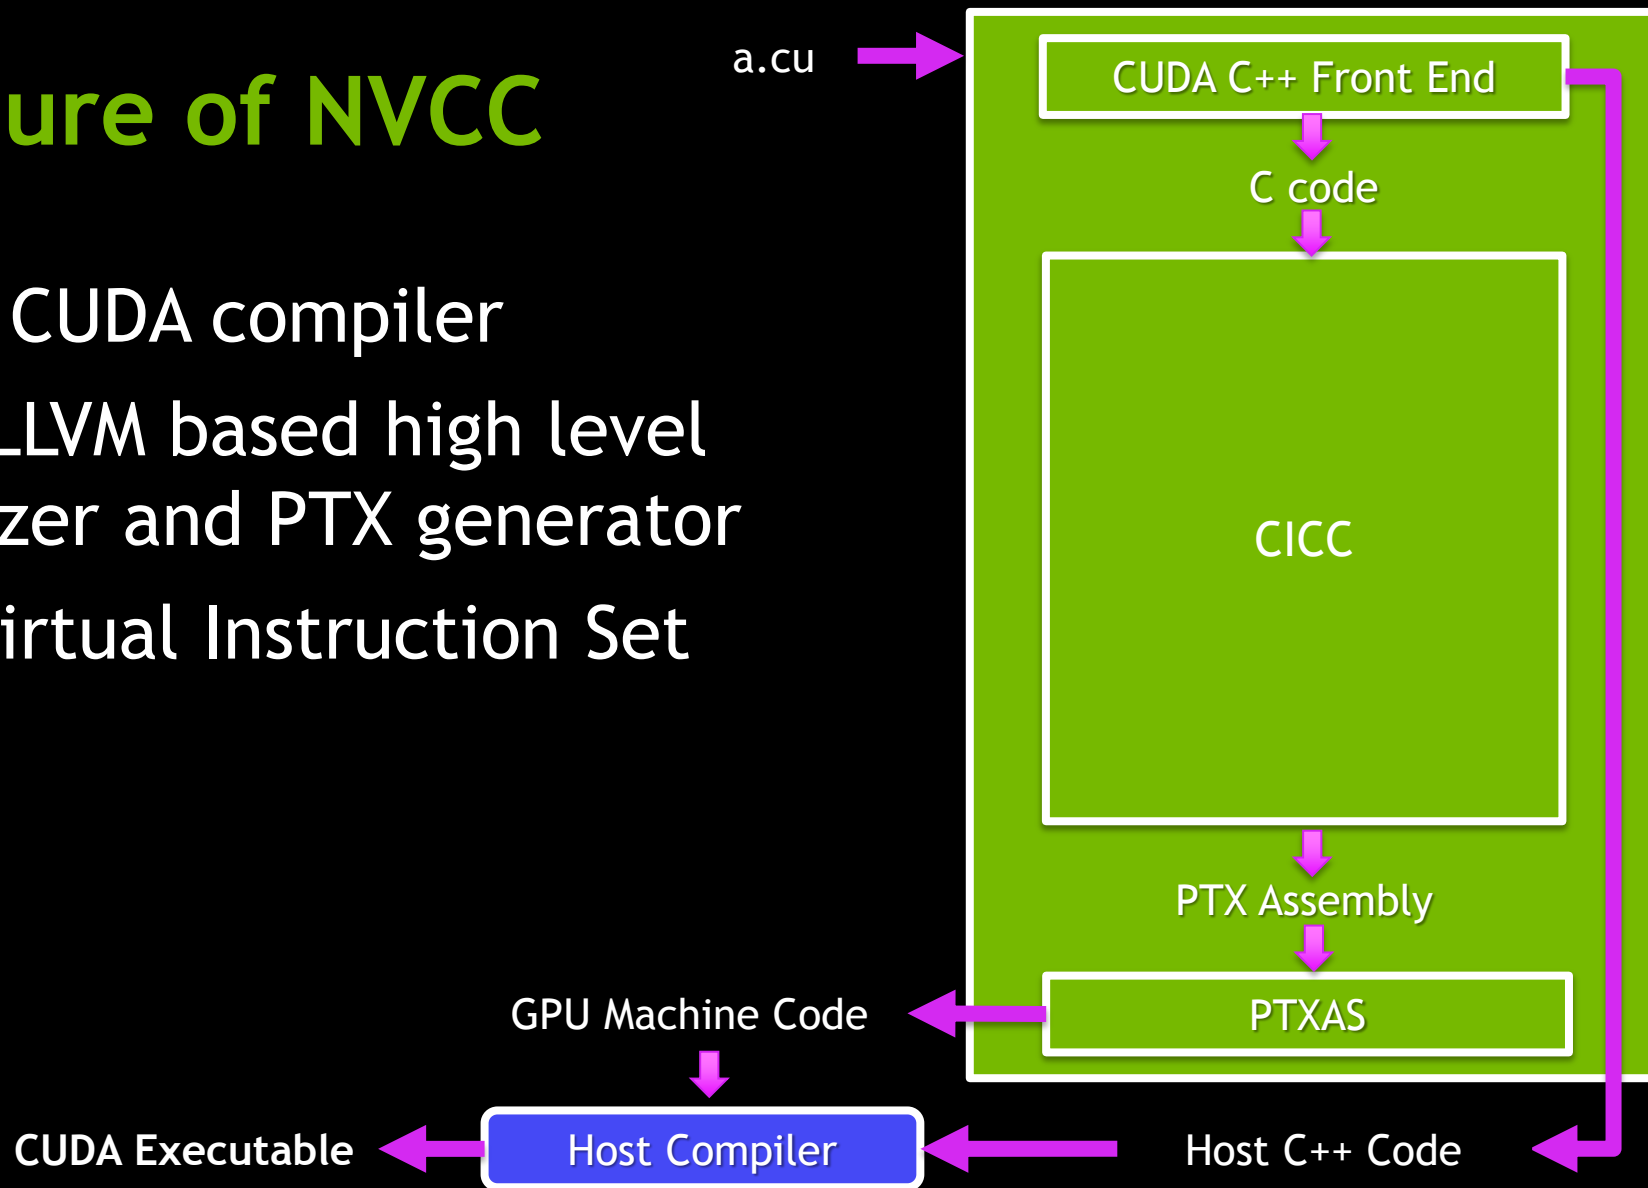
Structure of NVCC

- NVCC: CUDA compiler




Structure of NVCC

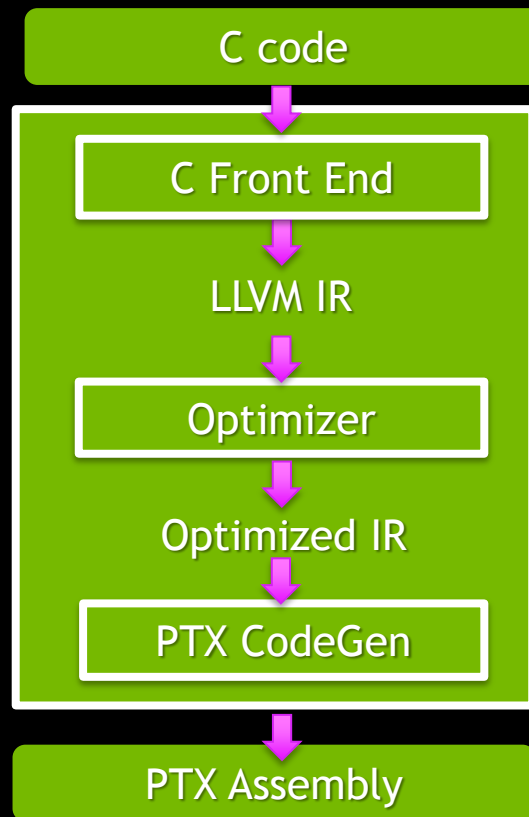
- NVCC: CUDA compiler
- CICC: LLVM based high level optimizer and PTX generator
- PTX: Virtual Instruction Set



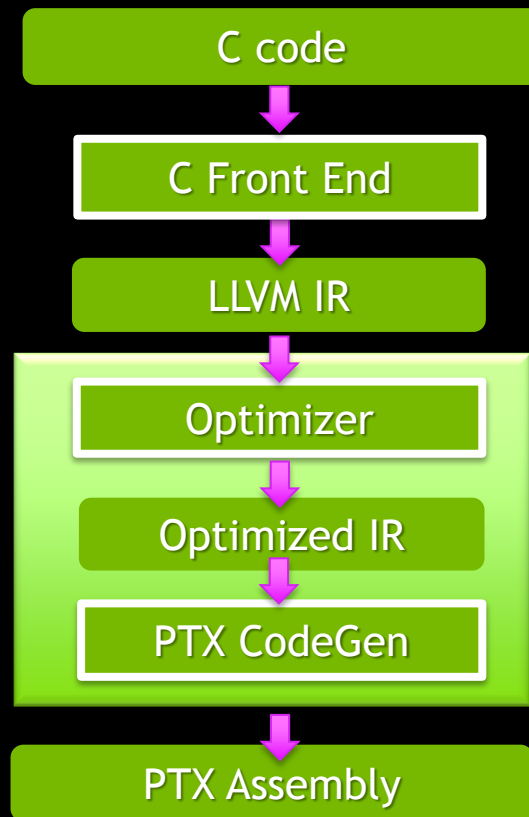
Structure of CICC



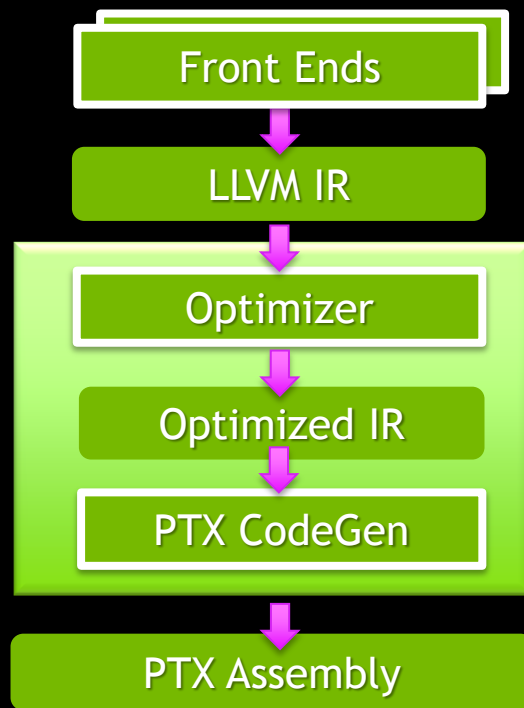
Structure of CICC



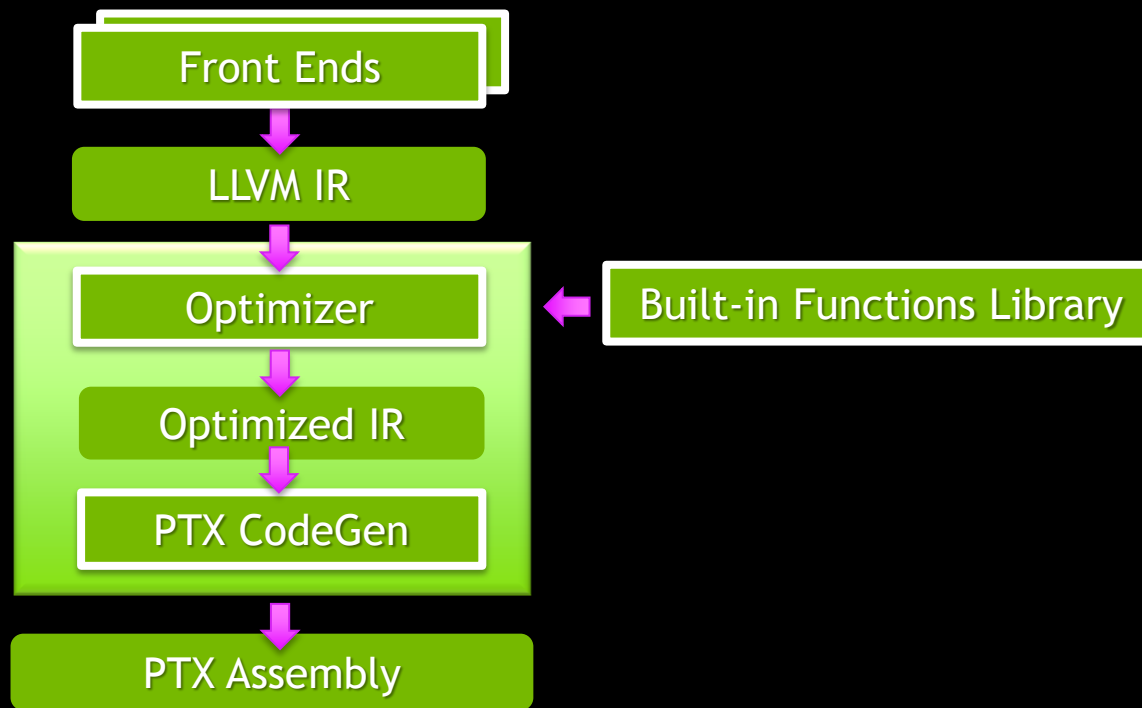
Structure of CICC



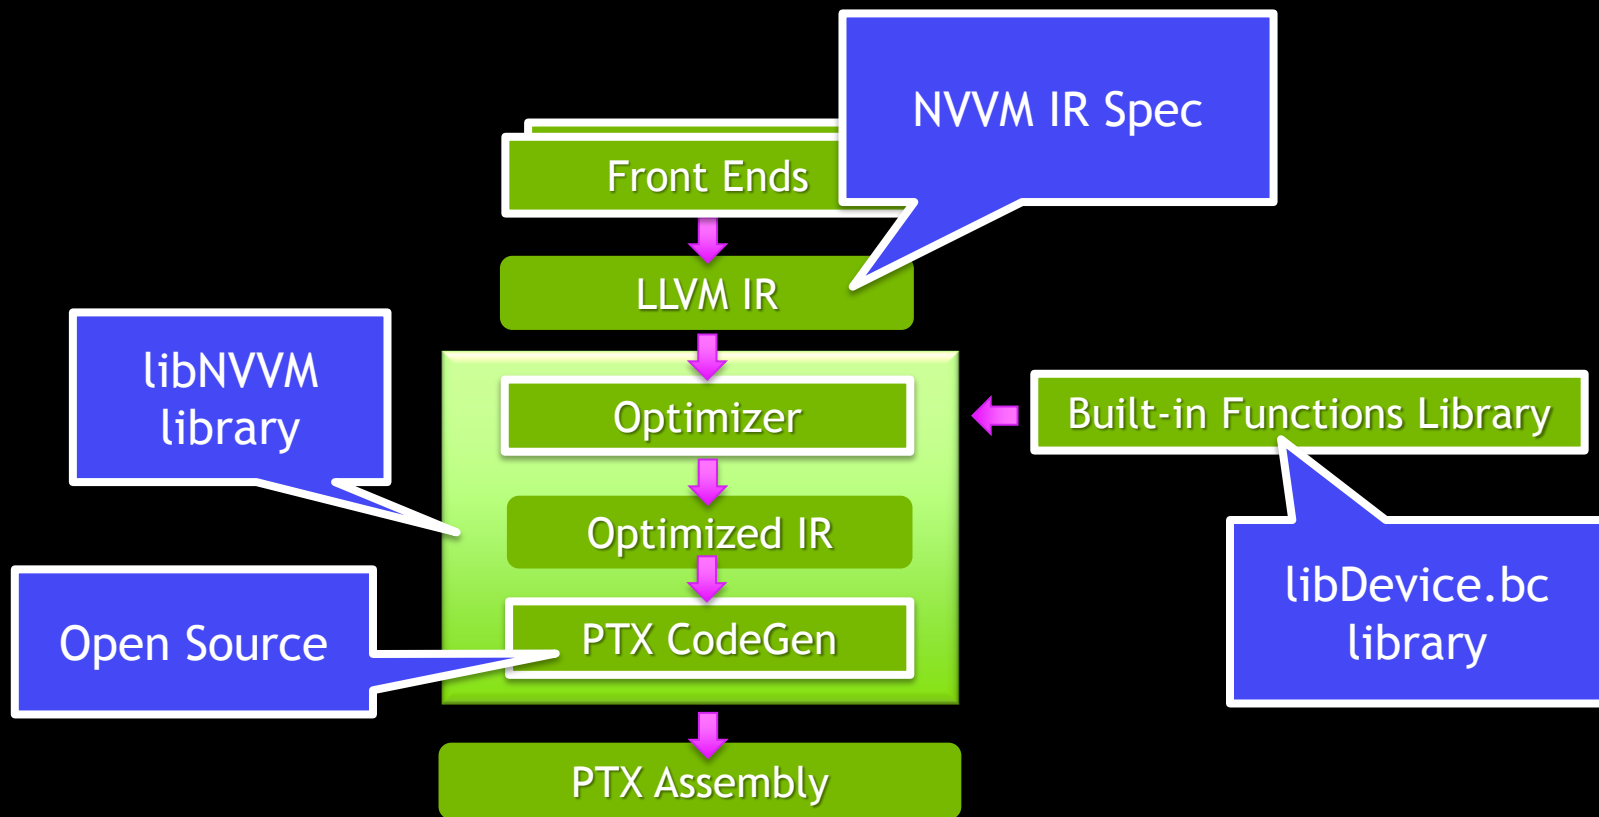
Common Compiler



Common Compiler



Common Compiler



Two-pronged Approach

- Open-source LLVM NVPTX backend
 - Community supported
- NVIDIA Compiler SDK
 - Binary library, header files, documents
 - Supported product

Enabling Open Source GPU Compilers


- Contributed NVPTX code generator sources back to LLVM in summer 2012
- Part of LLVM 3.2 release
- Actively maintained in LLVM trunk
 - by NVIDIA and other LLVM developers
- Standard LLVM License
- Best for
 - Prototyping
 - Developers who work only with LLVM trunk

NVIDIA Compiler SDK

- Preview was released at GTC 2012.
- 1st official release will be included in CUDA 5.5 toolkit.

LLVM NVPTX / libNVVM Users

- Numba Pro: array-oriented compiler for Numpy/Python
- Halide image processing language : MIT (halide-lang.org)
- Jet fluid dynamics DSL : Double Negative
- Alea.CUDA, F# on the GPU : QuantAlea
- Delite parallel EDSL framework : Stanford PPL
- KernelGen: open source compiler project at HPC forge



NVIDIA Compiler SDK

- NVVM IR specification
- libNVVM library and header file
- libDevice
- Code samples
- Developer's guide and API document



libNVVM Library

libNVVM Library


- An optimizing compiler library that generates PTX from NVVM IR
 - Supports LLVM 3.0, 3.1 and 3.2 IR format
 - Supports Fermi, Kepler and later architectures
- Available on 32-bit and 64-bit Windows, Linux and Mac
- Actually used by CUDA 5.5 compiler
 - Dynamically linked by cicc

libNVVM Library

- Analysis and optimizations
 - Address space access optimization
 - Thread variance/convergence analyses
 - Re-materialization
 - Load/store coalescing
 - Sign extension elimination
 - New phi elimination
 - Enhanced alias analysis
 - ...
- Support DWARF generation

libNVVM C APIs

- Create a program
 - Add NVVM IR modules to the program
 - Support NVVM IR level linking
- Verify the input IR
- Compile IR to PTX
- Get result
 - Get back PTX string
 - Get back message log



libDevice

libDevice

- Common device math functions
 - Distributed in LLVM bitcode format
 - Supports Fermi and Kepler
- Supports both ftz and non-ftz mode
- Can be linked with NVVM IR program using libNVVM API
 - Treated as a normal NVVM IR module in libNVVM
- Can be used with the open source LLVM NVPTX backend
- We build our CUDA math functions on top of it in CUDA 5.5.
- Will include more common device functions in the future.



NVVM IR

NVVM IR

- Designed to represent GPU kernel functions and device functions
 - Represents the code executed by each CUDA thread
- NVVM IR Specification 1.0
 - Based on LLVM IR 3.2

NVVM IR and LLVM IR

- NVVM IR
 - Based on LLVM IR
 - With a set of rules and intrinsics
- No new types. No new operators. No new reserved words.
- An NVVM IR program can work with any standard LLVM IR tool
 - ✓ llvm-as ✓ llvm-link ✓ llvm-extract
 - ✓ llvm-dis ✓ llvm-ar ✓ ...
- An NVVM IR program can be built with the standard LLVM distribution.

svn co http://llvm.org/svn/llvm-project/llvm/branches/release_32 llvm

NVVM IR

NVVM IR: Address Spaces



NVVM IR: Address Spaces

- CUDA C++

- Address space is a storage qualifier.
- A pointer is *generic* pointer, which can point to any address space.

```
__global__ int g;  
__shared__ int s;  
__constant__ int c;  
void foo(int a) {  
    int l;  
    int *p ;  
    switch (a) {  
        case 1: p = &g; ...  
        case 2: p = &s; ...  
        case 3: p = &c; ...  
        case 4: p = &l; ...  
    }  
    ...  
}
```


NVVM IR: Address Spaces

■ CUDA C++

- Address space is a storage qualifier.
- A pointer is *generic* pointer, which can point to any address space.

■ OpenCL C

- Address space is part of the type system.
- A pointer type must be qualified with an address space.

```
constant int c;  
  
foo(global int *pg) {  
  
    int l;  
    int *p ;  
    p = &l;  
  
    constant int *pc = &c;  
    ...  
}
```

NVVM IR: Address Spaces

- CUDA C++
 - Address space is a storage qualifier.
 - A pointer is *generic* pointer, which can point to any address space.
- OpenCL C
 - Address space is part of the type system.
 - A pointer type must be qualified with an address space.
- NVVM IR
 - Support both use cases in the same program.

NVVM IR: Address Spaces

- Define address space numbers
- Allow generic pointers and specific pointers
- Provide intrinsics to perform conversions between generic pointers and specific pointers

NVVM IR: Address Spaces

- Allow module scope variables that are in the global address space to have generic address values.
 - Make generating NVVM IR code much easier.

```
// @a is a module scope variable residing in the  
// global address space (1).  
// The address value of @a is a global address value.  
@a = addrspace(1) float 0.000000e+00
```

NVVM IR: Address Spaces

- Allow module scope variables that are in the global address space to have generic address values.
 - Make generating NVVM IR code much easier.

```
// @a is a module scope variable residing in the
// global address space (1).
// The address value of @a is a global address value.
@a = addrspace(1) float 0.000000e+00

// @b is a module scope variable residing in the
// global address space (1).
// The address value of @b is a generic address value.
@b = float 0.000000e+00
```

NVVM IR: GPU Program Properties

- Properties:
 - Maximum expected CTA size from any launch
 - Minimum number of CTAs on an SM
 - Kernel function vs. device function
 - Texture/surface variables
 - more
- Use named metadata

NVVM IR: Intrinsic

- Atomic operations
- Barriers
- Address space conversions
- Special registers read
- Texture/surface access
- more

Programming Guide

Developer's Guide and API Document

- NVVM IR Specification
- libNVVM: Developer's Guide and API Spec
- libDevice: Developer's Guide and API Spec
- Will be available
 - online @ docs.nvidia.com
 - as PDF files

Samples

Samples

- “simple”
 - JIT compile a NVVM IR program using libNVVM
 - launch it using CUDA driver API

libNVVM

CUDA Driver

Samples

- “ptxgen”
 - A simple offline NVVM IR to PTX compiler
 - Links in libDevice

libNVVM

libDevice

Samples

- “cuda-c-linking”
 - Create a NVVM IR program using LLVM IR builder API
 - JIT compile the NVVM IR program using libNVVM
 - Link it with a PTX generated from CUDA C using PTX JIT linking API
 - Launch the final code using CUDA driver API

libNVVM

LLVM

PTX JIT linking

Samples


- “simple”
- “ptxgen”
- “cuda-c-linking”
- Part of CUDA 5.5 toolkit

Samples

- Samples that
 - are relatively big,
 - depend on other open-source package, or
 - need to be updated more frequently

Samples on github

- Samples that
 - are relatively big,
 - depend on other open-source package, or
 - need to be updated more frequently
- Examples
 - Other language bindings for libNVVM: Python, Haskell
 - Kaleidoscope
 - Small utilities
- Open source BSD style license



Example

SAXPY

```
let (saxpy (lambda (a : f32)
                  (x : vf32)
                  (y : vf32) : vf32
                  (map (lambda (xi : f32)
                        (yi : f32) : f32
                        (+ (* a xi) yi) )
                        x y)))
in (saxpy A X Y)
```

SAXPY

```
let (saxpy (lambda (a : f32)
                (x : vf32)
                (y : vf32) : vf32
                (map (lambda (xi : f32)
                        (yi : f32) : f32
                        (+ (* a xi) yi)
                        )
                    x y)))
in (saxpy A X Y)
```

SAXPY

- Execute this on the GPU.
- Use one GPU thread for each vector element.

```
saxpy(float a, float *x, float *y) {  
    y[thread_id] = a * x[thread_id] + y[thread_id];  
}
```

```
let (saxpy (lambda (a : f32)  
                (x : vf32)  
                (y : vf32) : vf32  
                (map (lambda (xi : f32)  
                        (yi : f32) : f32  
                        (+ (* a xi) yi) )  
                    x y)))  
in (saxpy A X Y)
```

@n = internal global i32 0, align 4

@a = internal global float 0.000000e+00, align 4

```
define void @saxpy(float* %x, float* %y) {  
}
```


@n = internal global i32 0, align 4

@a = internal global float 0.000000e+00, align 4

```
define void @saxpy(float* %x, float* %y) {  
}
```

!nvvm.annotations = !{!0}

!0 = metadata !{void (float, float*, float*)* @saxpy, metadata !"kernel", i32 1}

```
target datalayout = "e-p:64:64:64-i1:8:8-i8:8:8-i16:16:16-i32:32:32  
-i64:64:64-f32:32:32-f64:64:64-v16:16:16-v32:32:32-v64:64:64-v128:128:128-n16:32:64"
```

```
@n = internal global i32 0, align 4
```

```
@a = internal global float 0.000000e+00, align 4
```

```
define void @saxpy(float* %x, float* %y) {  
}
```

```
!nvvm.annotations = !{!0}
```

```
!0 = metadata !{void (float, float*, float*)* @saxpy, metadata !"kernel", i32 1}
```

```
define void @saxpy(float %a, float* %x, float* %y) {  
    ; load @n and @a  
    ; int i = blockIdx.x * blockDim.x + threadIdx.x;  
    ; load x[i]  
    ; load y[i]  
    ; a * x + y  
    ; store y[i]  
}
```

```
define void @saxpy(float %a, float* %x, float* %y) {
```

```
    ; load @n and @a
```

```
    %n = load i32 * @n, align 4
```

```
    %a = load float * @a, align4
```

```
    ...
```

```
}
```

```
declare i32 @llvm.nvvm.read.ptx.sreg.ctaid.x()
declare i32 @llvm.nvvm.read.ptx.sreg.ntid.x()
declare i32 @llvm.nvvm.read.ptx.sreg.tid.x()

define void @saxpy(float %a, float* %x, float* %y) {

    ; load @n and @a
    %n = load i32 * @n, align 4
    %a = load float * @a, align4

    ; int i = blockIdx.x * blockDim.x + threadIdx.x;
    %0 = call i32 @llvm.nvvm.read.ptx.sreg.ctaid.x()
    %1 = call i32 @llvm.nvvm.read.ptx.sreg.ntid.x()
    %mul = mul i32 %1, %0
    %2 = call i32 @llvm.nvvm.read.ptx.sreg.tid.x()
    %add = add i32 %mul, %2

    ...
}
```

```
declare i32 @llvm.nvvm.read.ptx.sreg.ctaid.x()
declare i32 @llvm.nvvm.read.ptx.sreg.ntid.x()
declare i32 @llvm.nvvm.read.ptx.sreg.tid.x()

define void @saxpy(float %a, float* %x, float* %y) {

    ; load @n and @a
    %n = load i32 * @n, align 4
    %a = load float * @a, align 4

    ; int i = blockIdx.x * blockDim.x + threadIdx.x;
    %0 = call i32 @llvm.nvvm.read.ptx.sreg.ctaid.x()
    %1 = call i32 @llvm.nvvm.read.ptx.sreg.ntid.x()
    %mul = mul i32 %1, %0
    %2 = call i32 @llvm.nvvm.read.ptx.sreg.tid.x()
    %add = add i32 %mul, %2
```

```
    ; load x[i]
    %3 = sext i32 %add to i64
    %x_ptr = getelementptr float * %x, i64 %3
    %x_value = load float * %x_ptr, align 4
```

```
    ; load y[i]
    %y_ptr = getelementptr float * %y, i64 %3
    %y_value = load float * %y_ptr, align 4
```

```
    ; a * x + y
    %mul11 = fmul float %a, %x_value
    %add16 = fadd float %mul11, %y_value
```

```
    ; store y[i]
    store float %add16, float * %y_ptr, align 4
```

```
    ret void
```

```
}
```

Using libNVVM APIs

```
#include "nvvm.h"
char *compiler ()
{
    nvvmProgram prog;
    nvvmCreateProgram(&prog);
    char *buffer;  size_t size;
    getSAXPYir(&buffer, &size);
    nvvmProgramAddModule(prog, buffer, size);
    getLibDevice(&buffer, &size, "libdevice.compute_" #ARCH "." #MAJOR #MINOR ".bc")
    nvvmProgramAddModule(prog, buffer, size);
    nvvmCompileProgram(prog, 0, NULL);
    size_t ptxSize;
    nvvmGetCompiledResultSize(prog, &ptxSize);
    char *ptx = (char *) malloc(ptxSize);
    nvvmGetCompiledResult(prog, ptx);
    nvvmDestroyProgram(&prog);
    return ptx;
}
```


How to get them?

- Distributed with CUDA 5.5
 - NVVM IR spec
 - libNVVM library and header file
 - libDevice
 - Code samples
 - Developer's guide and API document
- LLVM.org: NVPTX backend
- Github: More samples
- devtalk.nvidia.com: Post your questions/suggestions