

# *LARGE SCALE MACHINE LEARNING WITH THE SIMSQL SYSTEM*

**Chris Jermaine**  
**Rice University**

**Many Current and Past Rice Team Members**  
**Also, Peter J. Haas at IBM Almaden**

# This Talk Is About

- Programming environments/execution platforms for big ML

# First, an Admission...

- I'm a database guy<sup>TM</sup>

# First, an Admission...

- I'm a database guy<sup>TM</sup>
- What does that mean?

# First, an Admission...

- I'm a database guy<sup>TM</sup>
- What does that mean?
- To me, means that I worship at the church of **data independence**
  - Now what in the heck does *that* mean?

# First, an Admission...

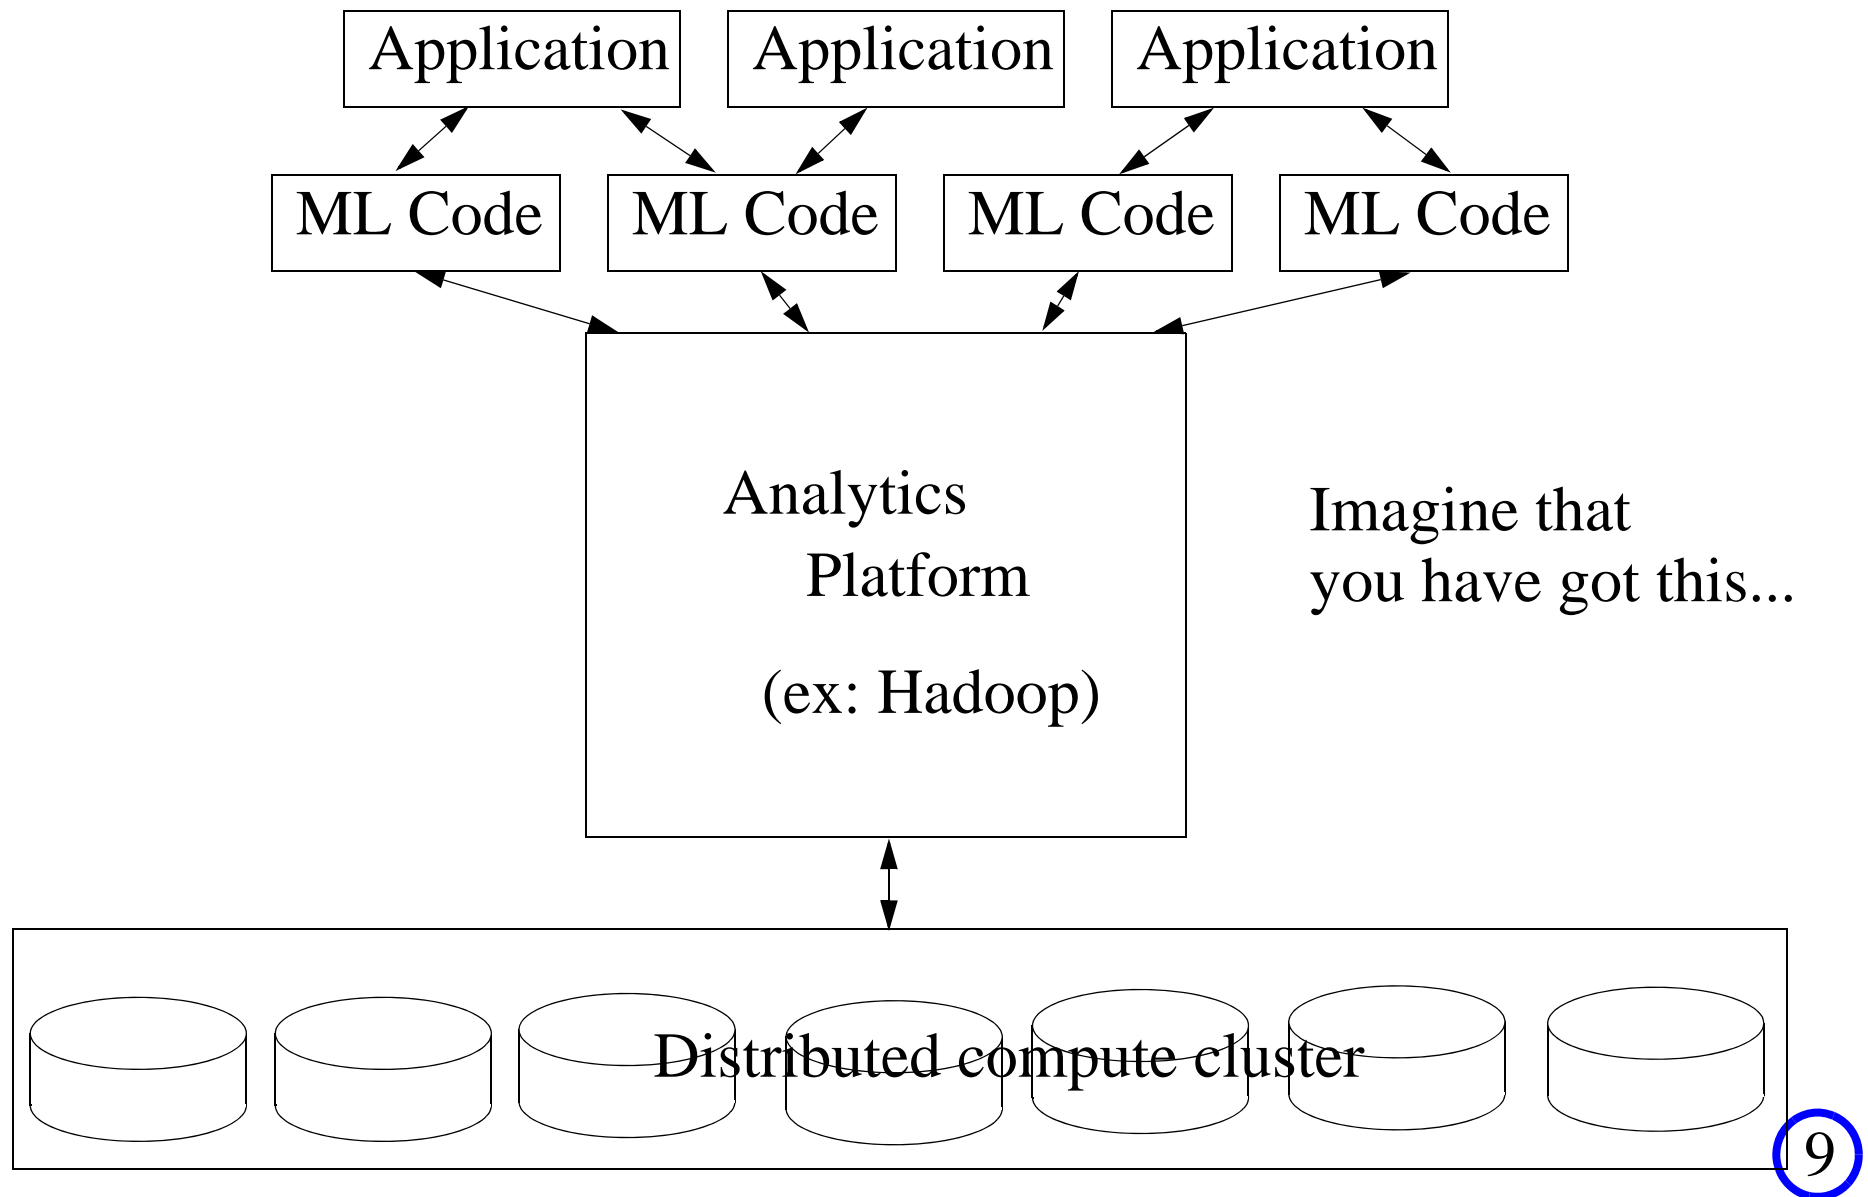
- I'm a database guy<sup>TM</sup>
- What does that mean?
- To me, means that I worship at the church of **data independence**
  - Now what in the heck does *that* mean?
- Means that when one designs a data-processing system...
- It should strive for the following ideal:
  - Coder specifies **what** the computation result should be, not **how** to get there
  - System itself figures out the **how** (the “declarative” paradigm)
  - Means code can be independent of data format, size, schema, processing hardware
  - Same code runs on one box with a GPU and on a 1000-machine cluster

# Why Are Declarative and Data Ind. Good?

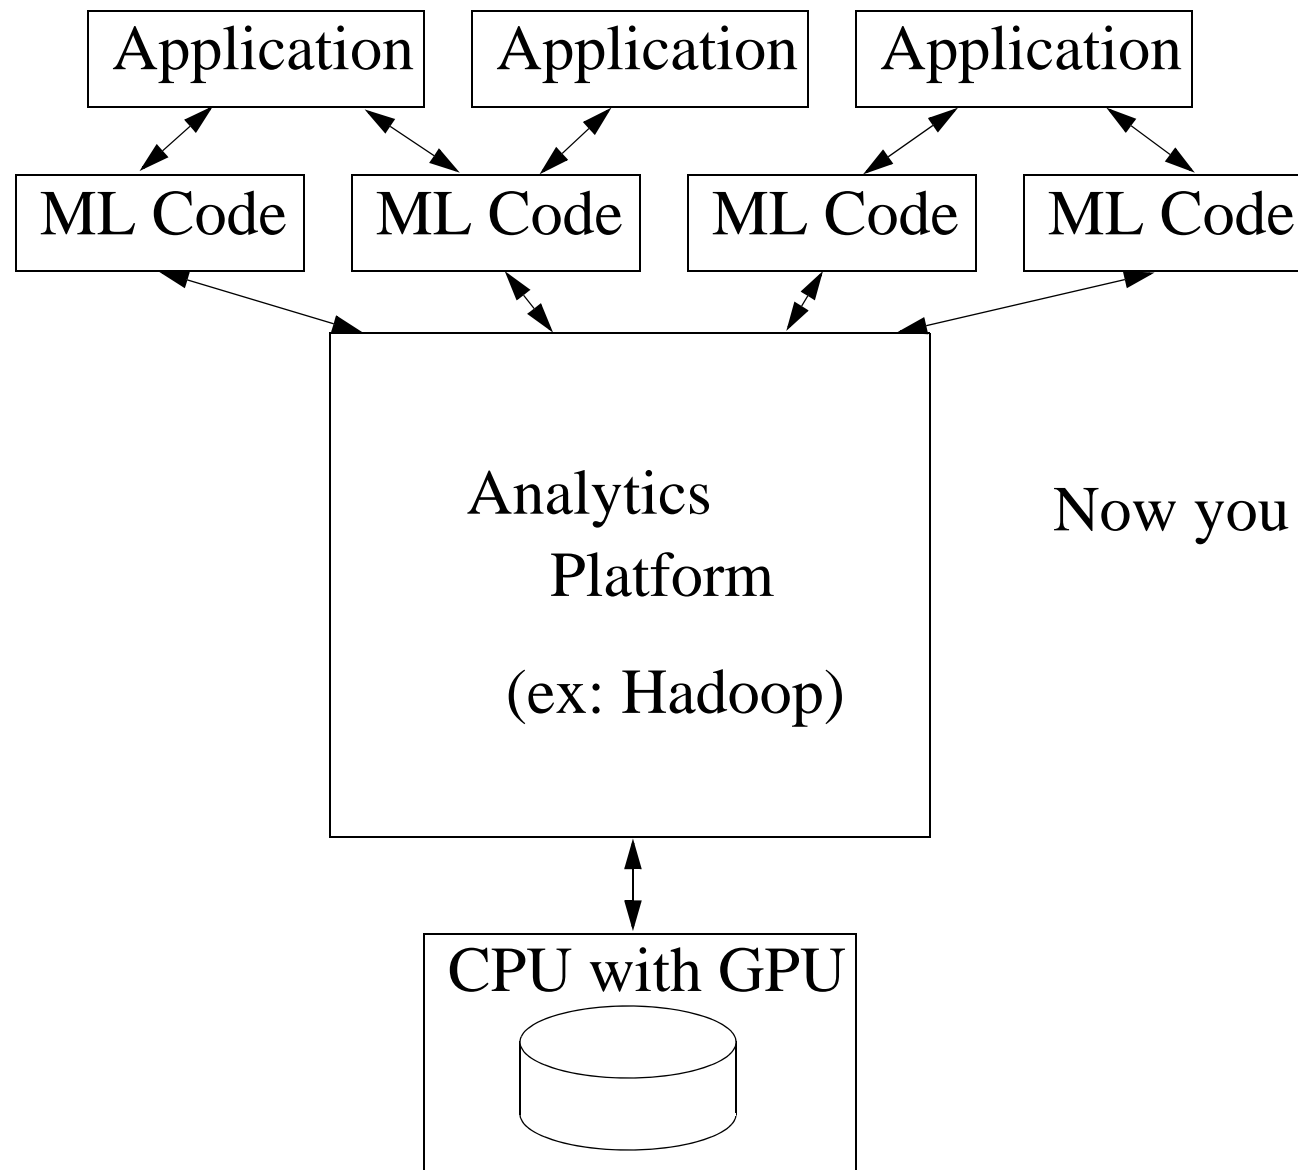
# Declarative Arg 1: One Code, Many Backends



# Declarative Arg 1: One Code, Many Backends

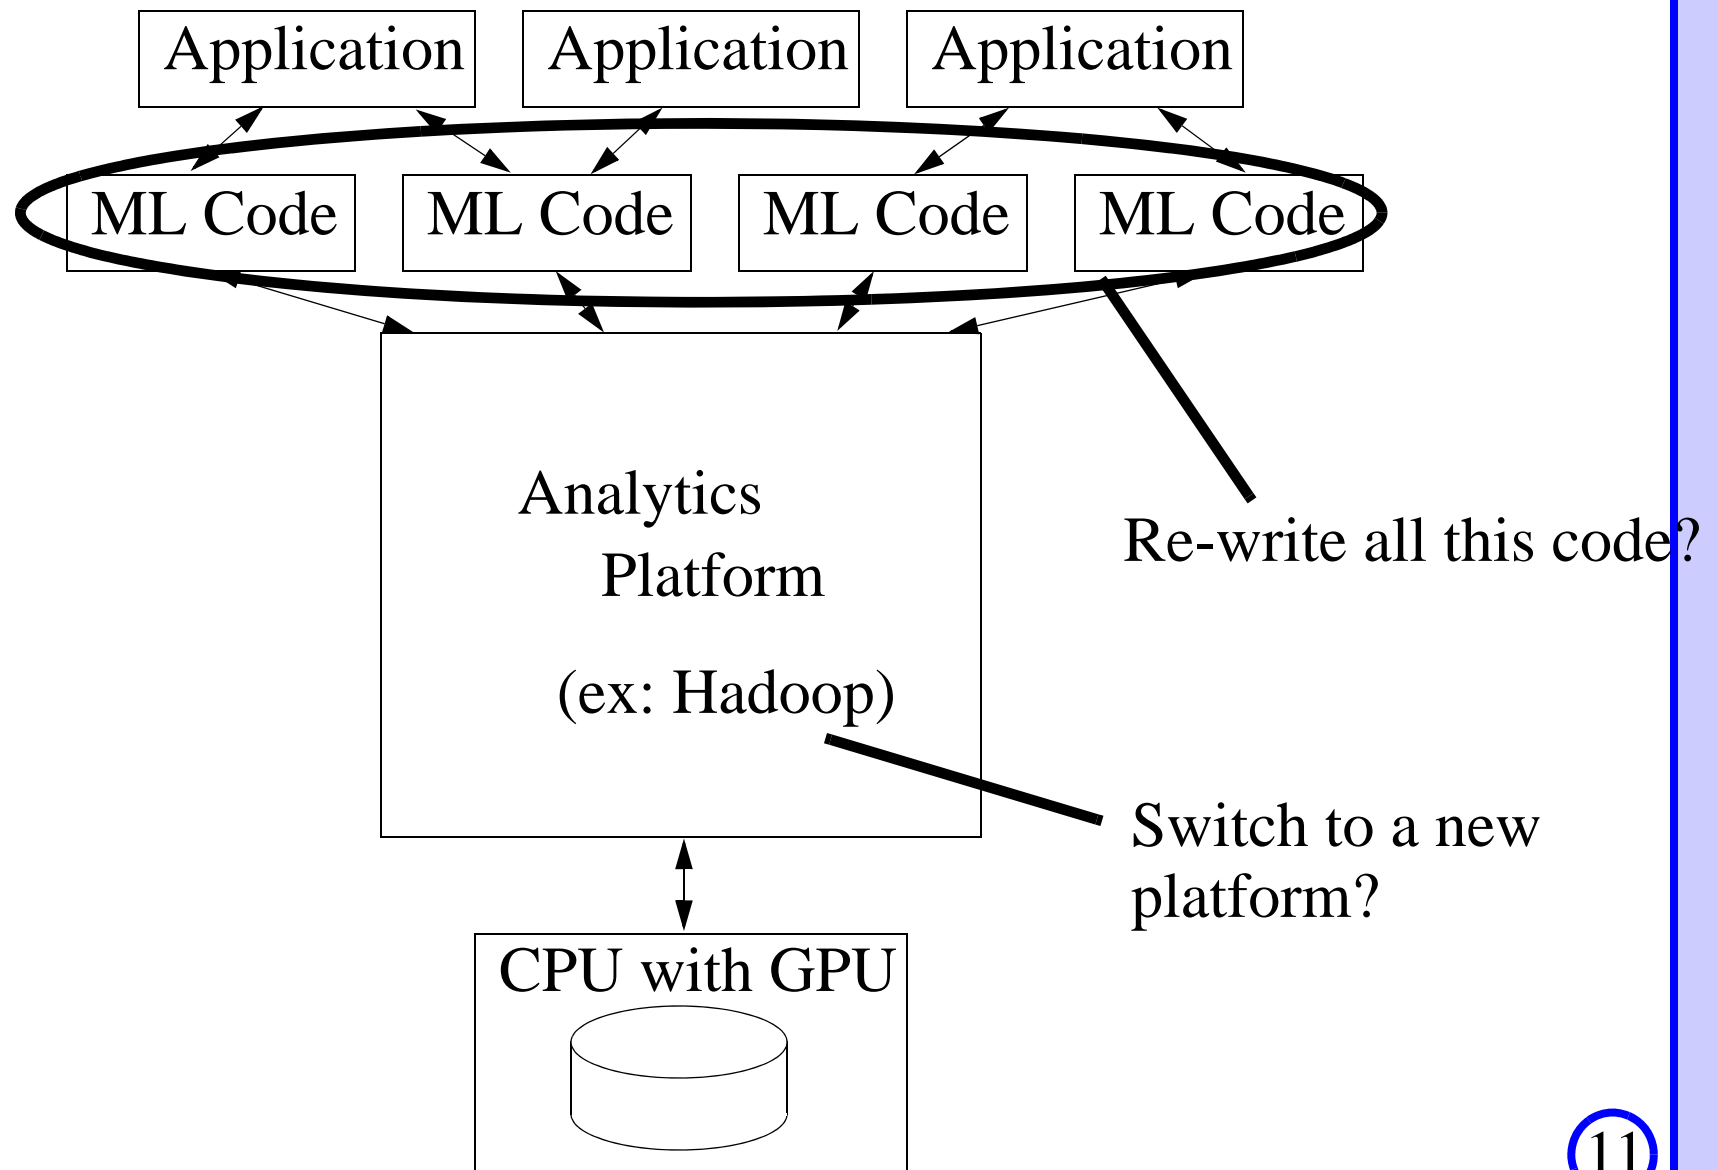


# Declarative Arg 1: One Code, Many Backends

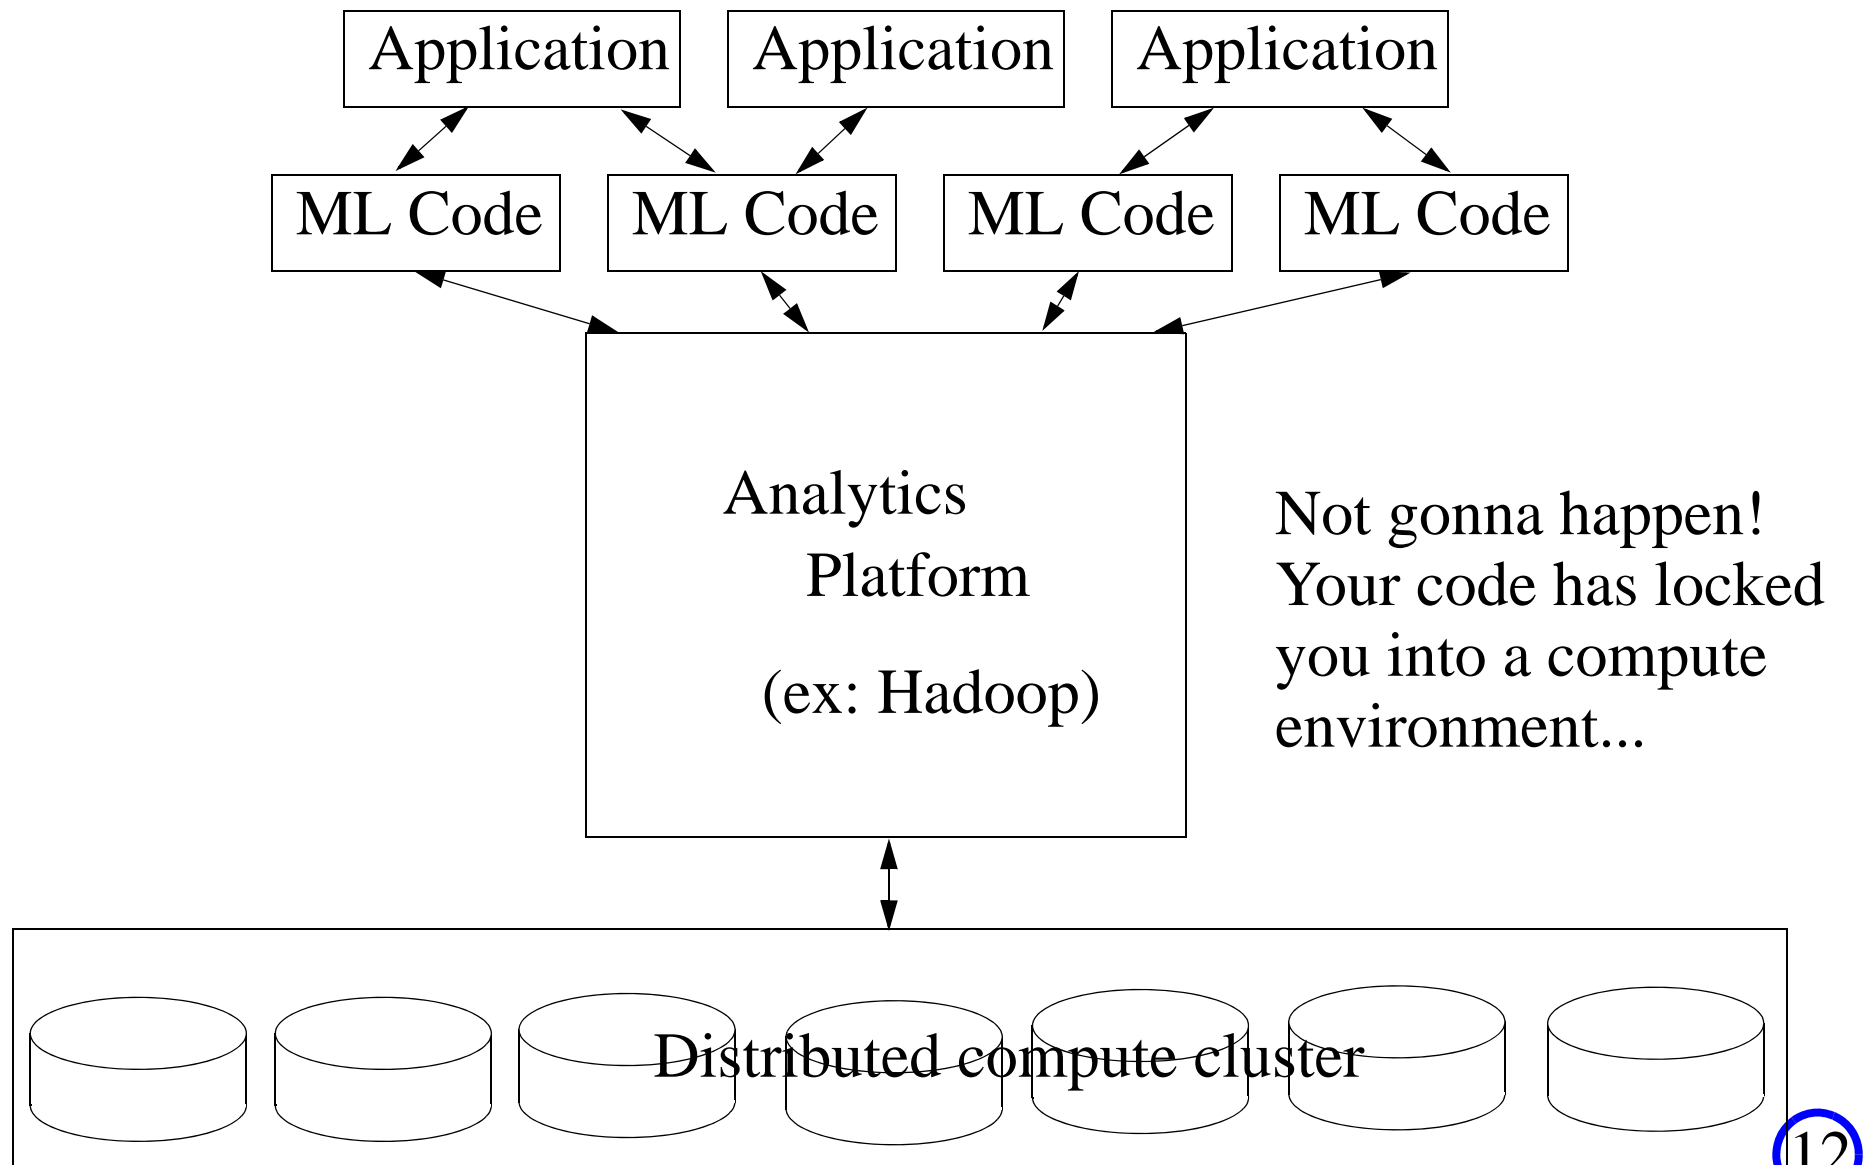


Now you want this...

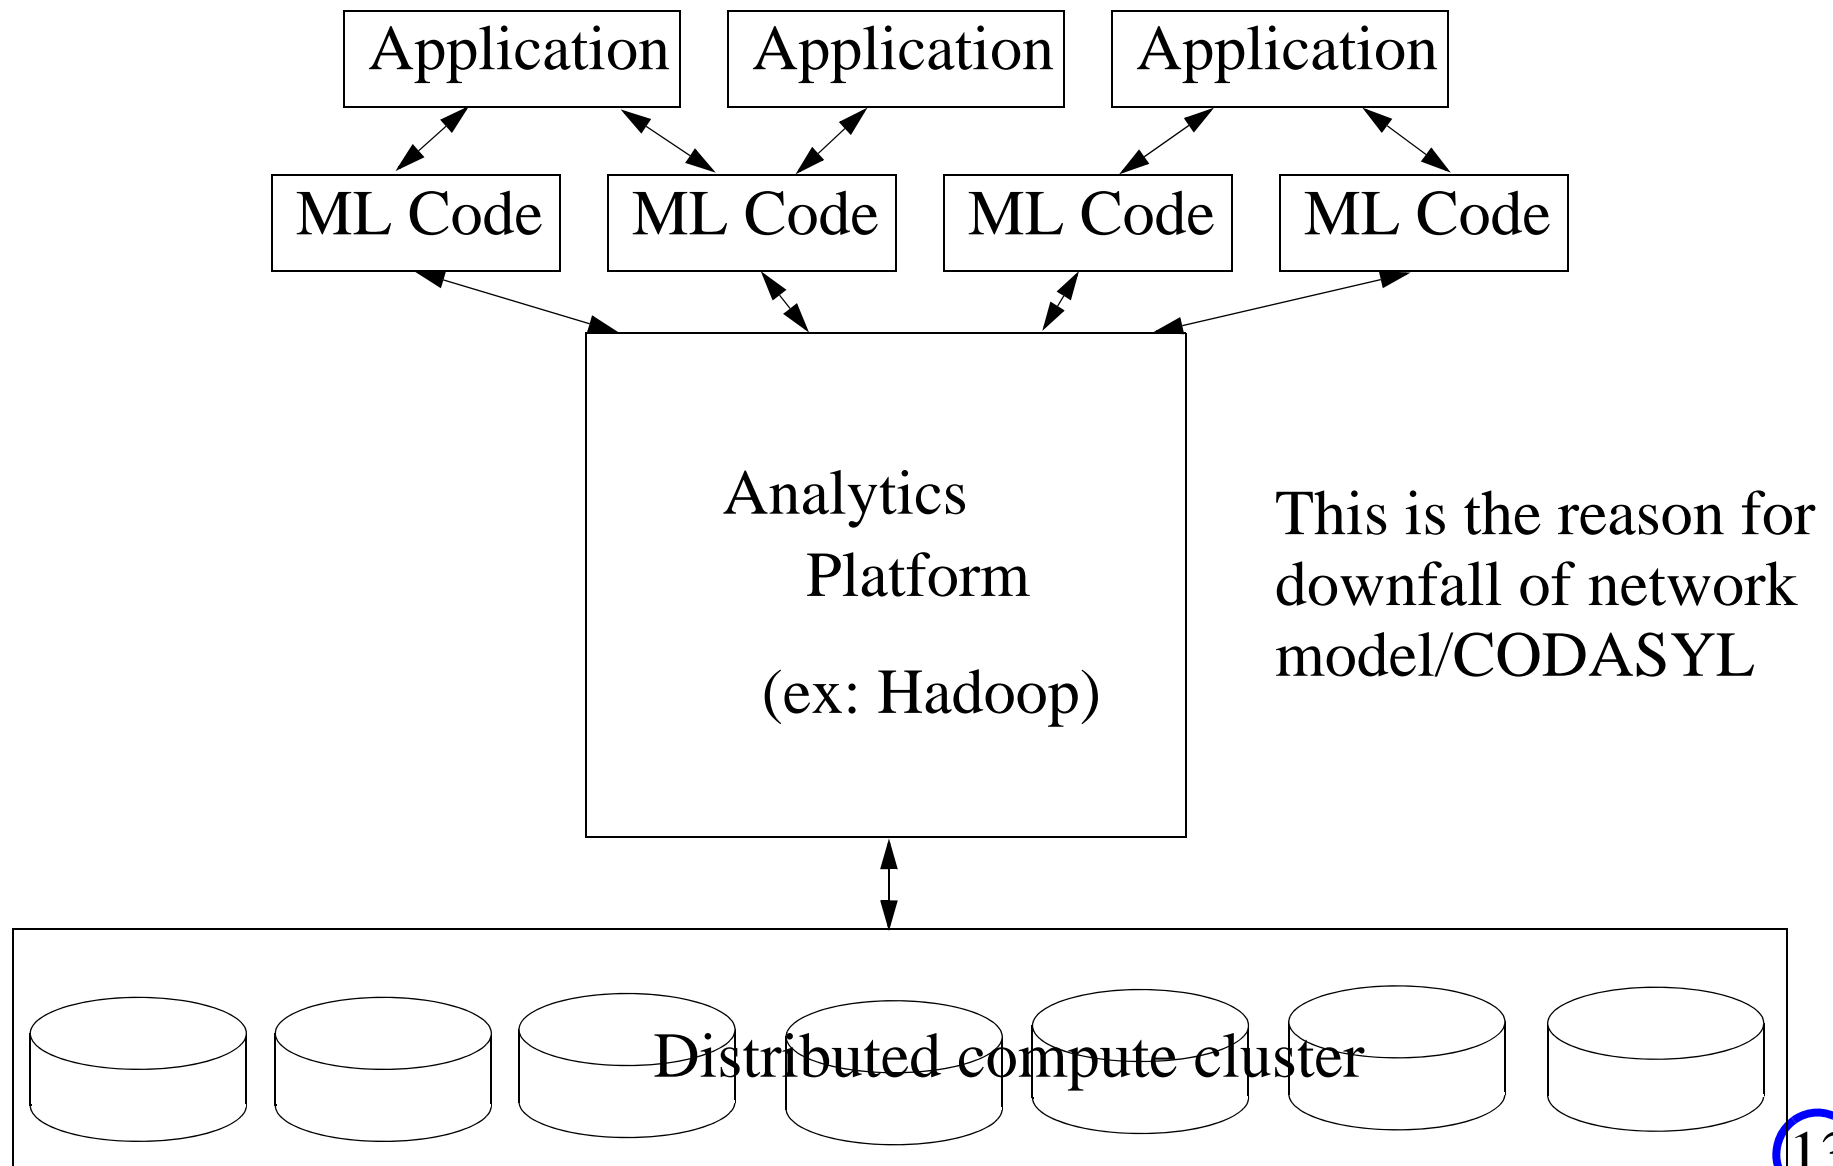
# Declarative Arg 1: One Code, Many Backends



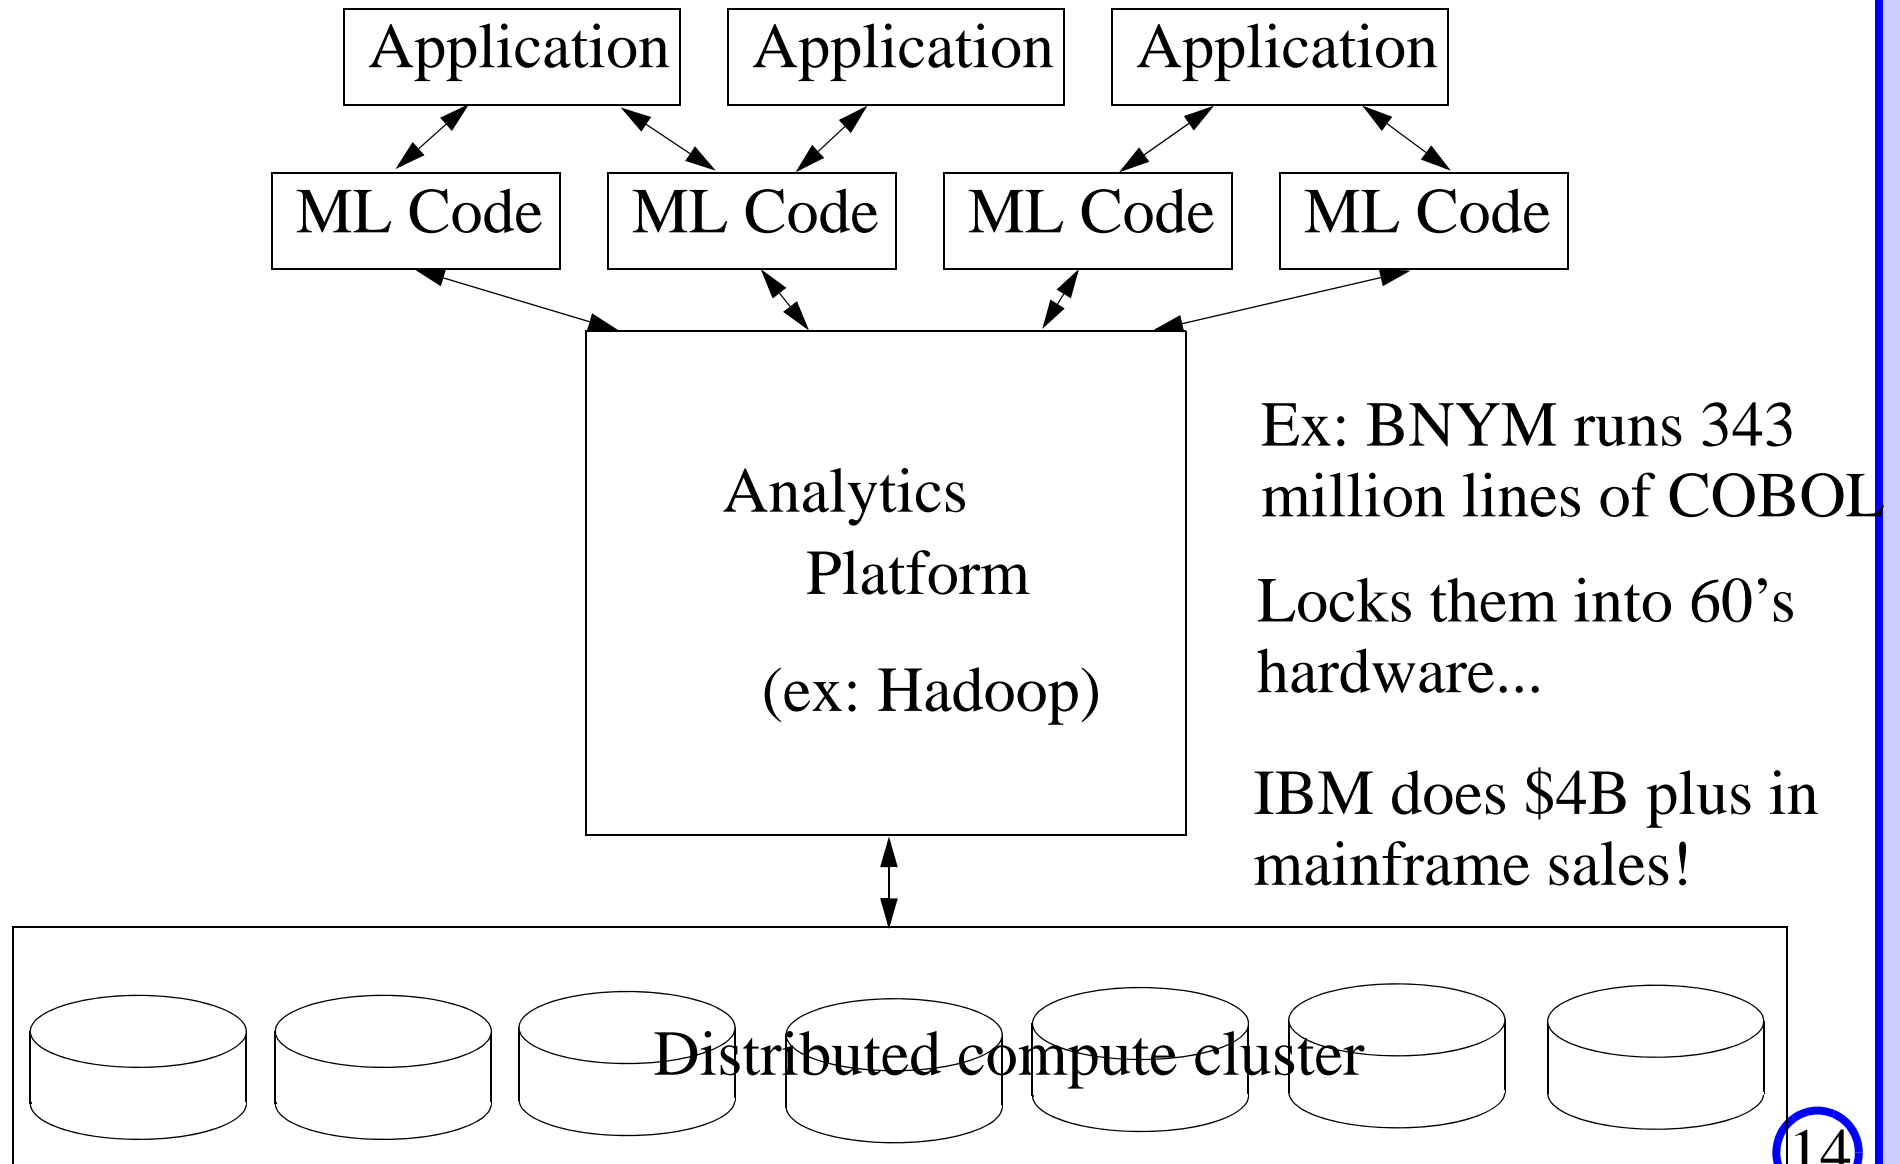
# Declarative Arg 1: One Code, Many Backends



# Declarative Arg 1: One Code, Many Backends

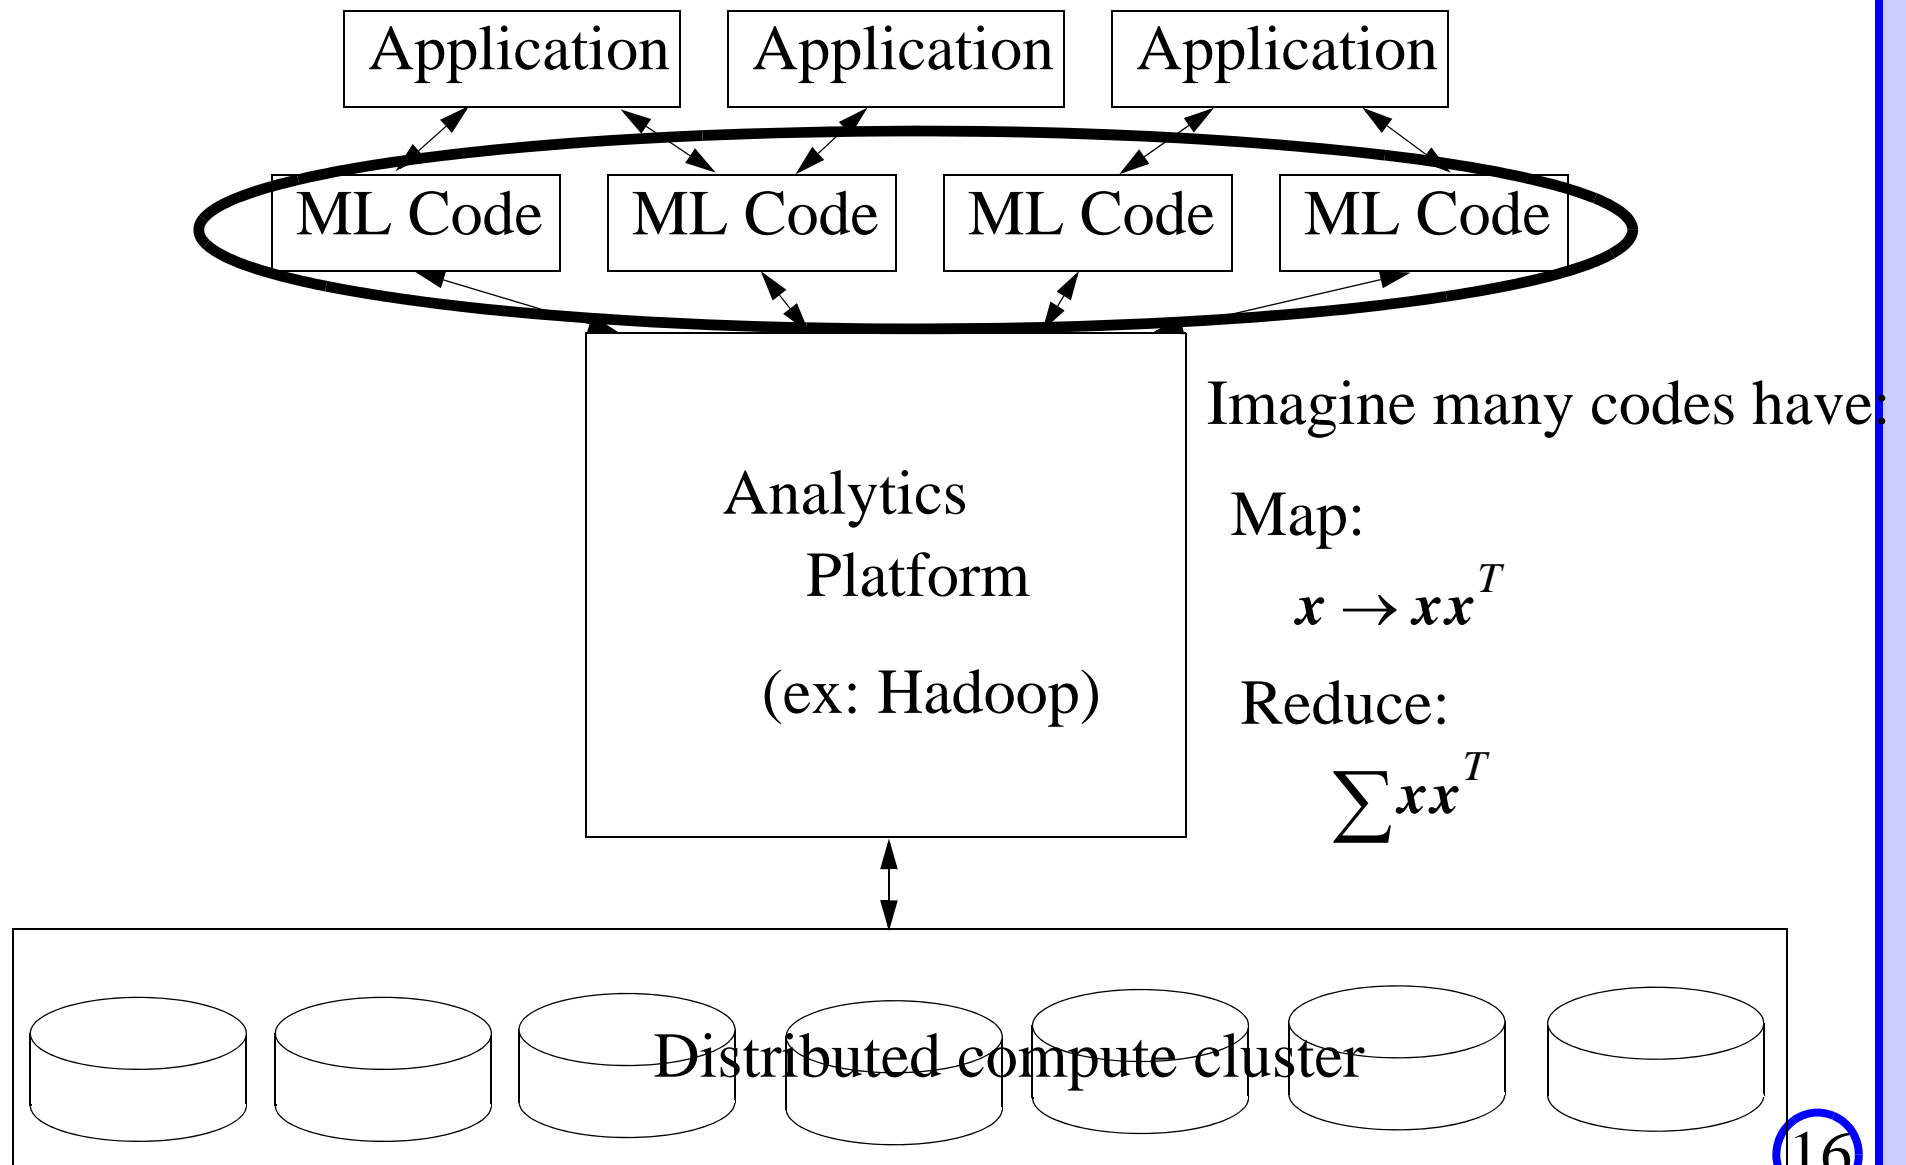


# Declarative Arg 1: One Code, Many Backends



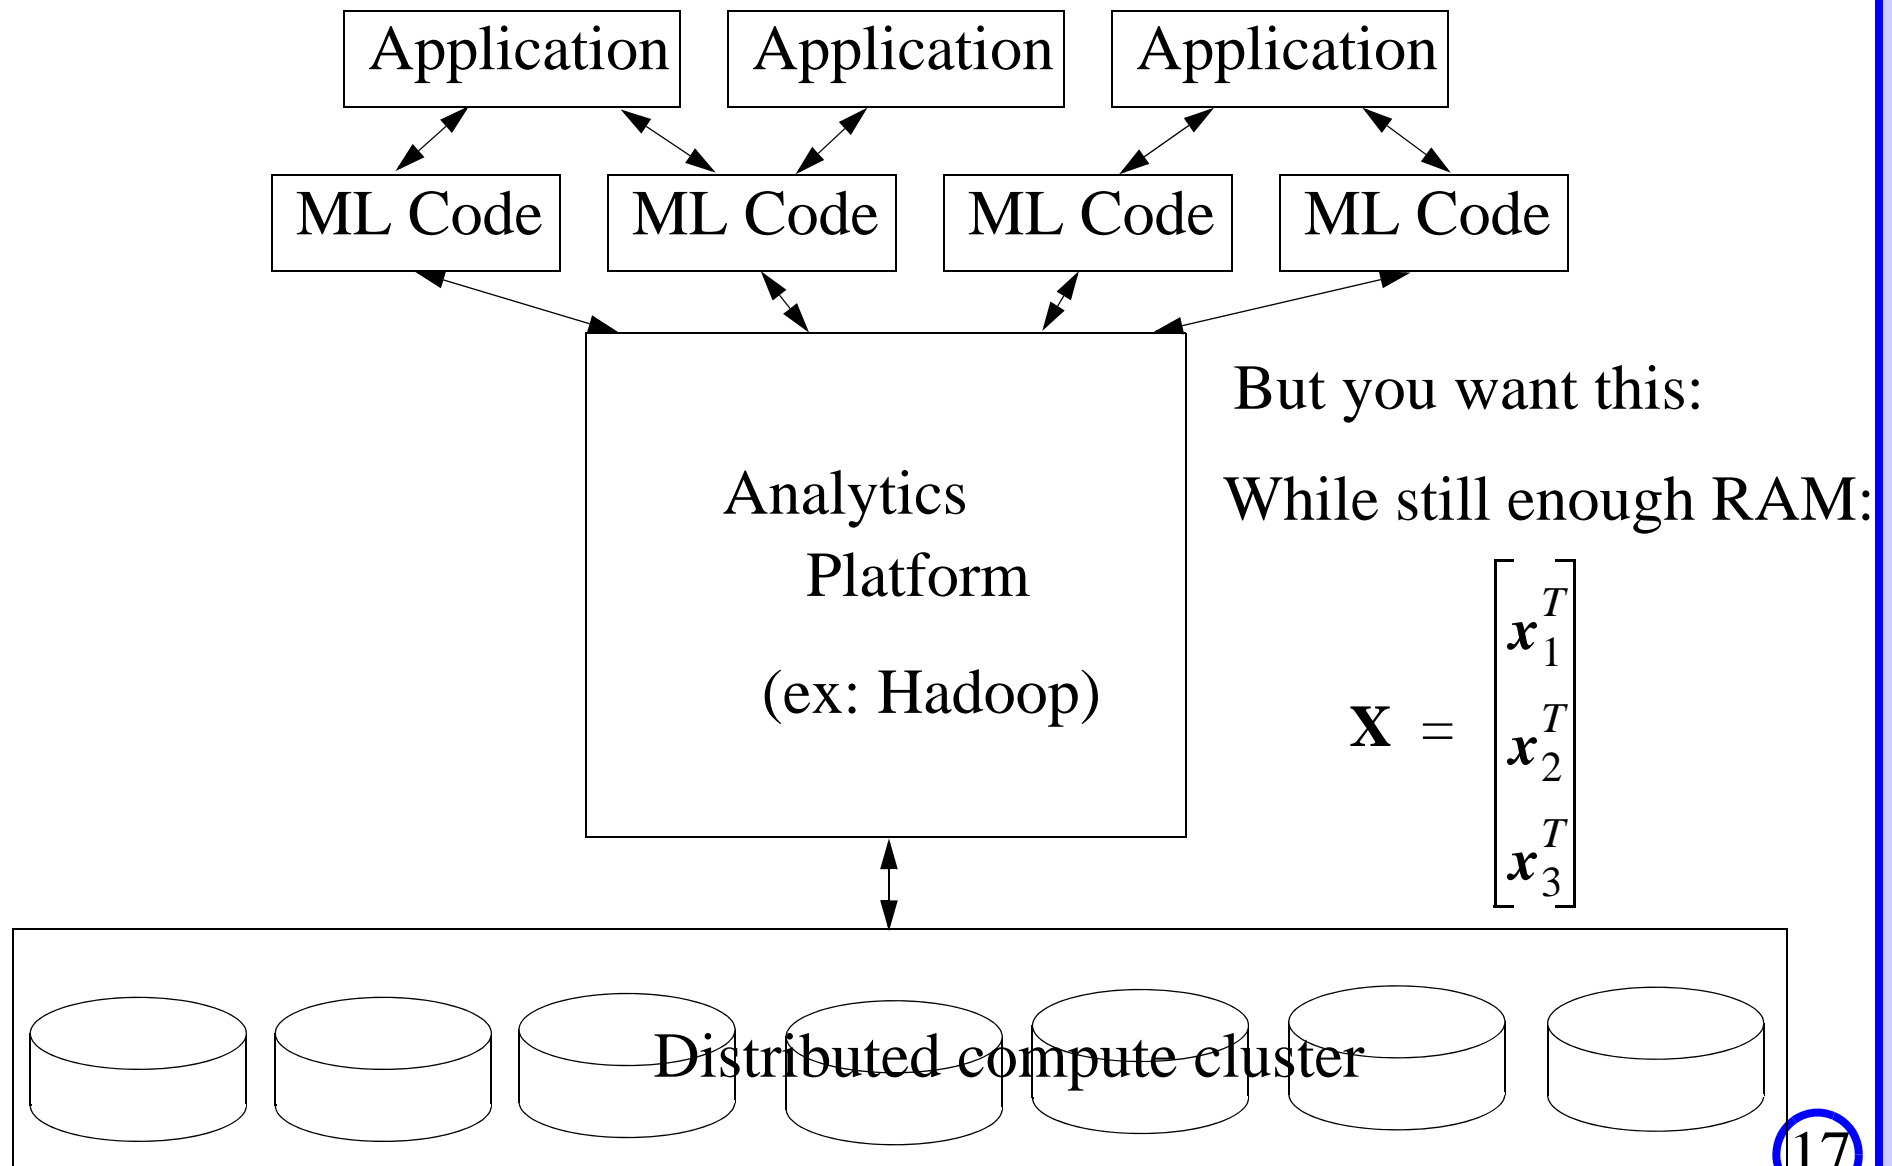
# Declarative Arg 2: Freedom From Algorithms

# Declarative Arg 2: Freedom From Algorithms

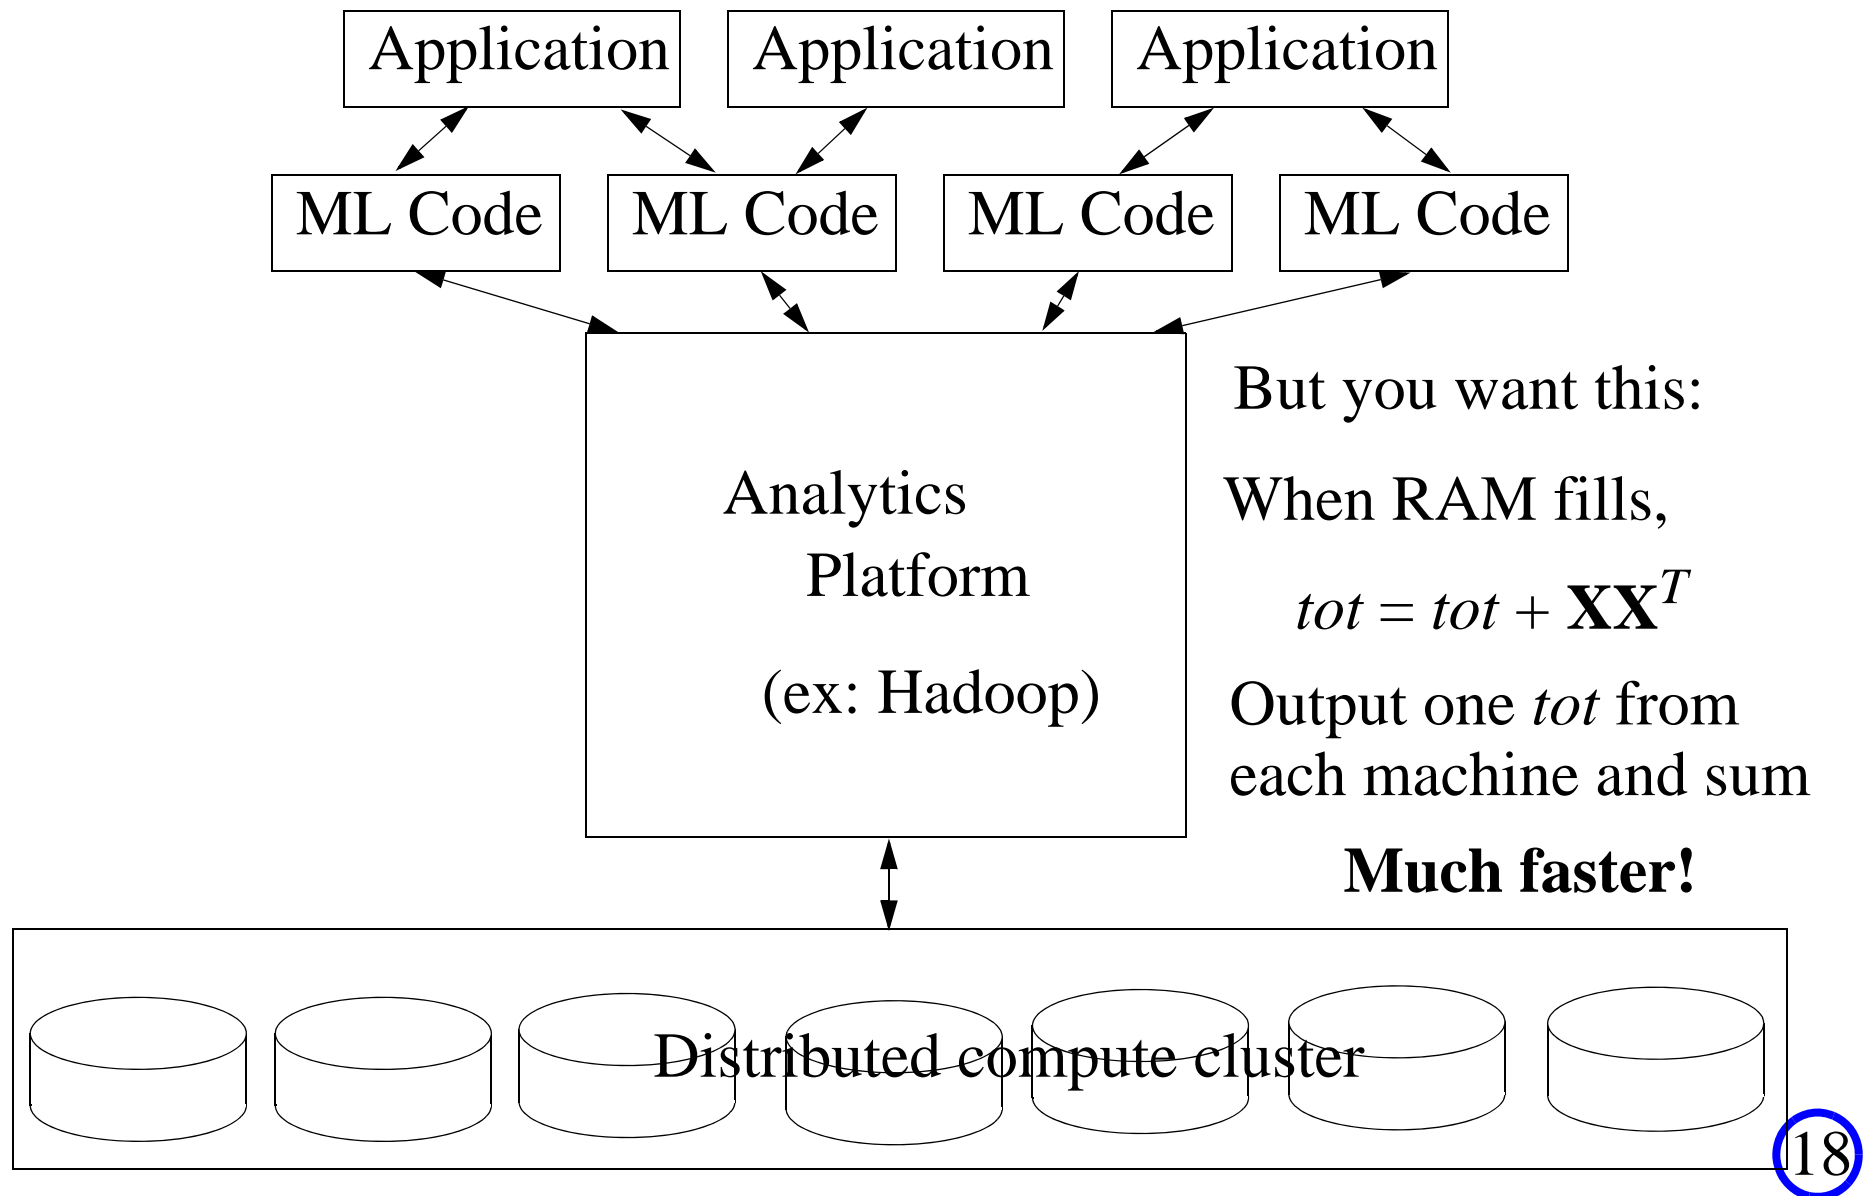




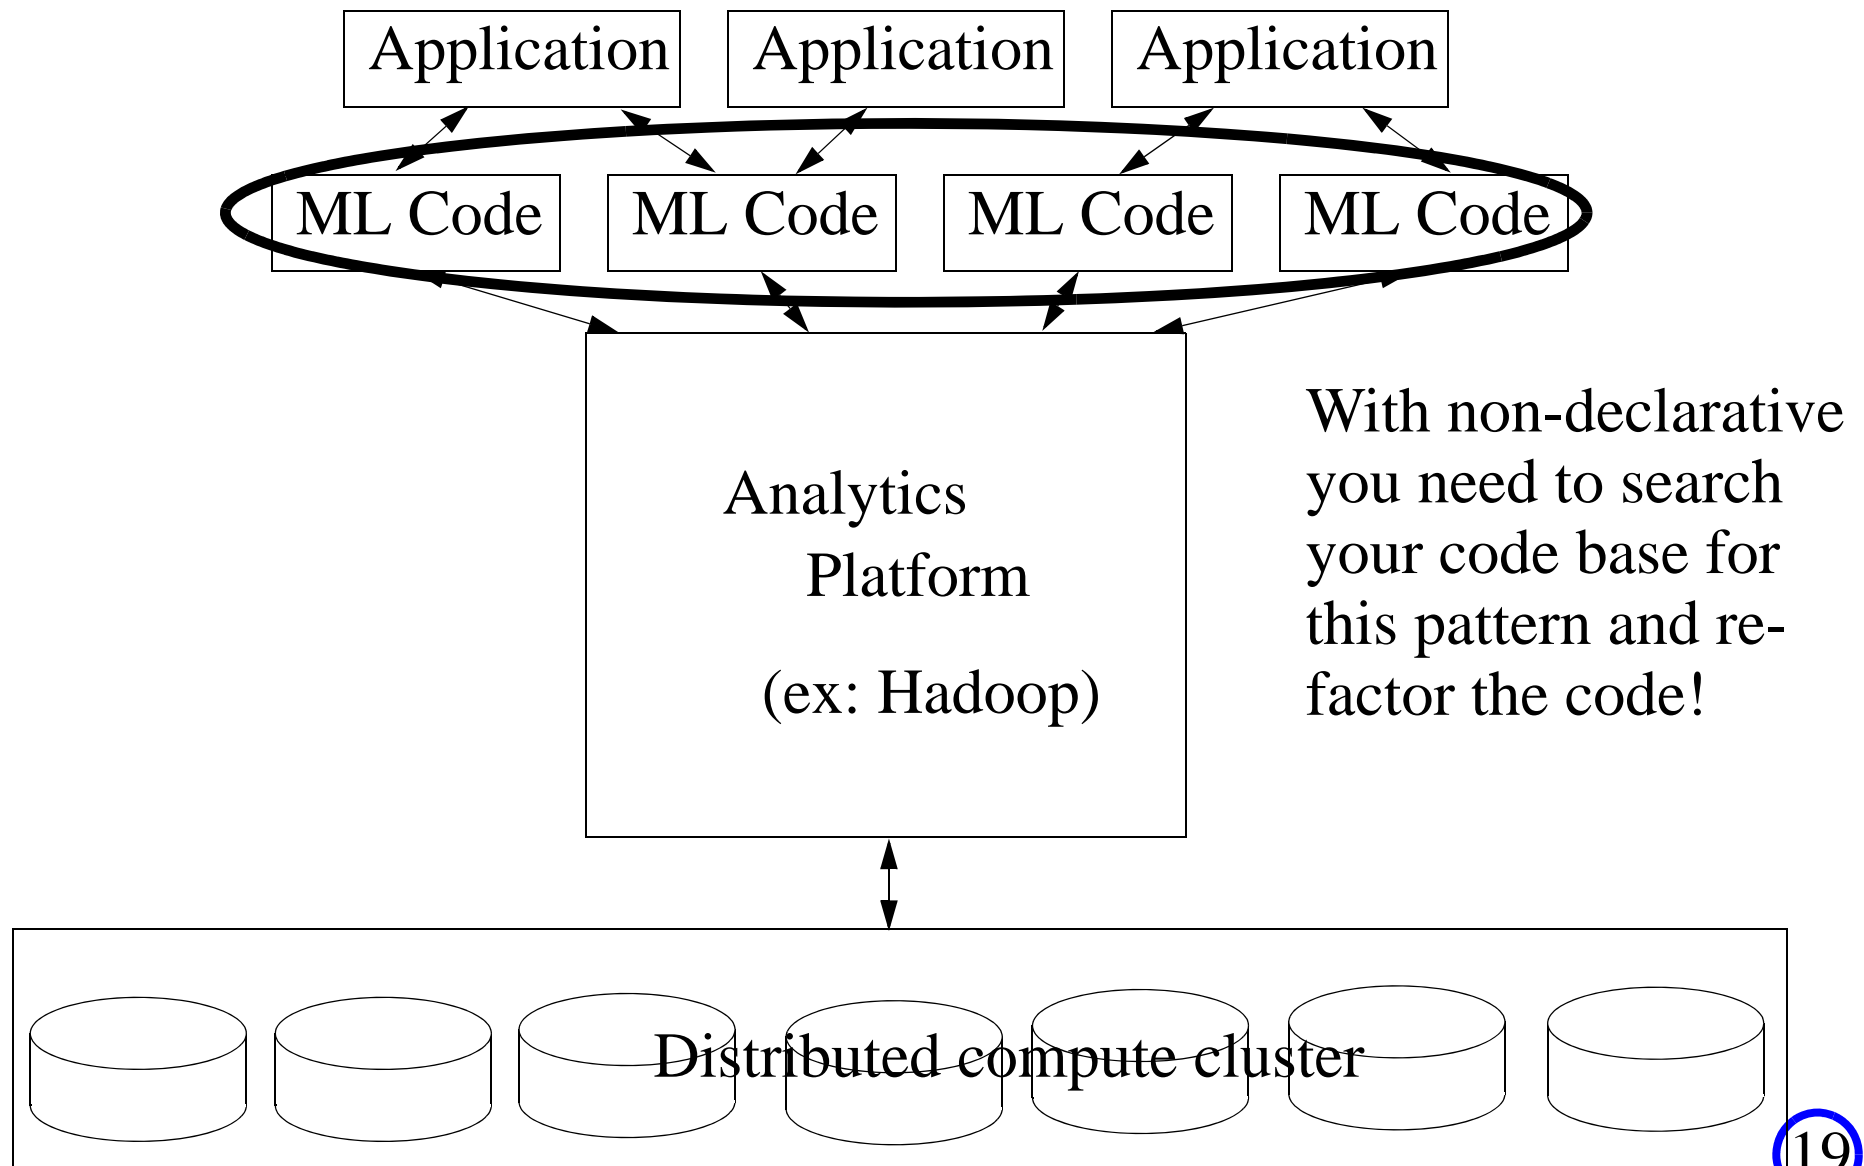
# Declarative Arg 2: Freedom From Algorithms



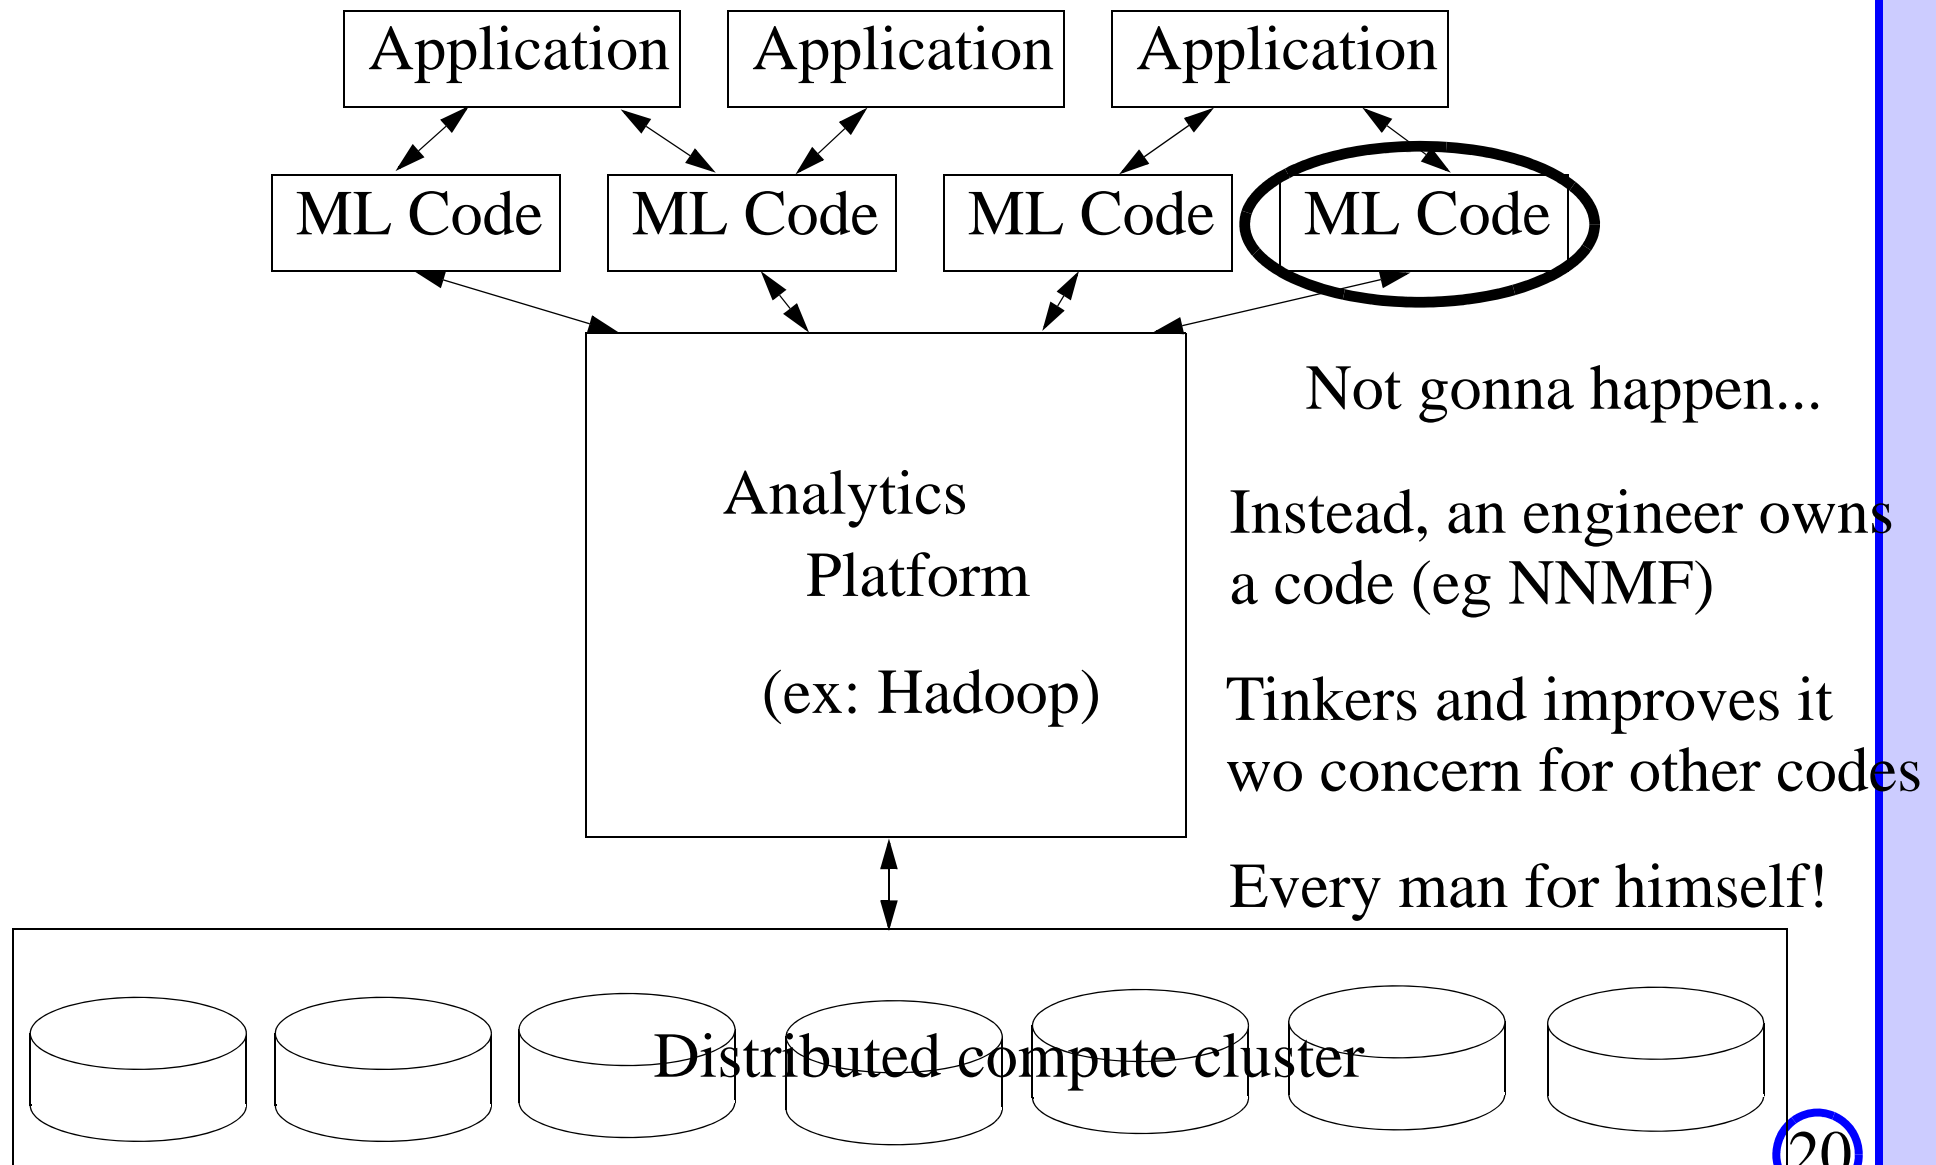
# Declarative Arg 2: Freedom From Algorithms



# Declarative Arg 2: Freedom From Algorithms

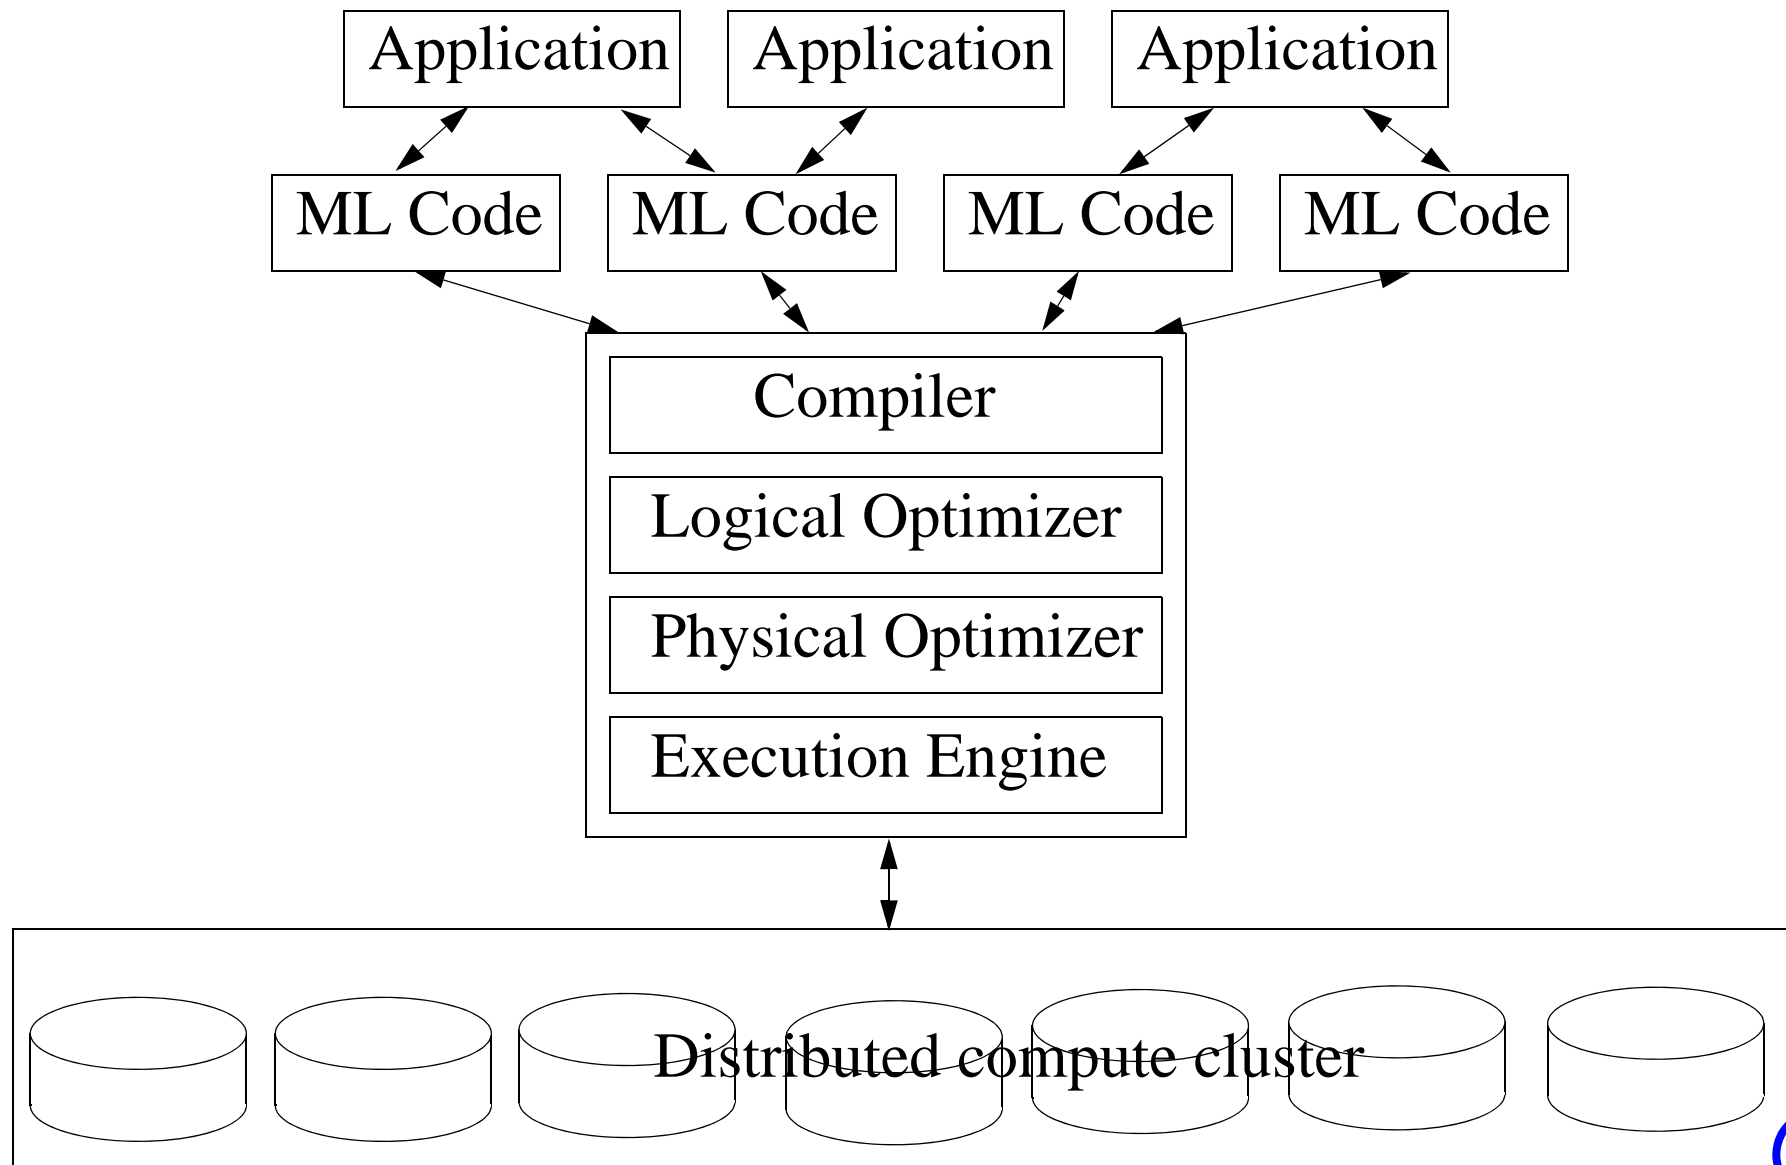


# Declarative Arg 2: Freedom From Algorithms

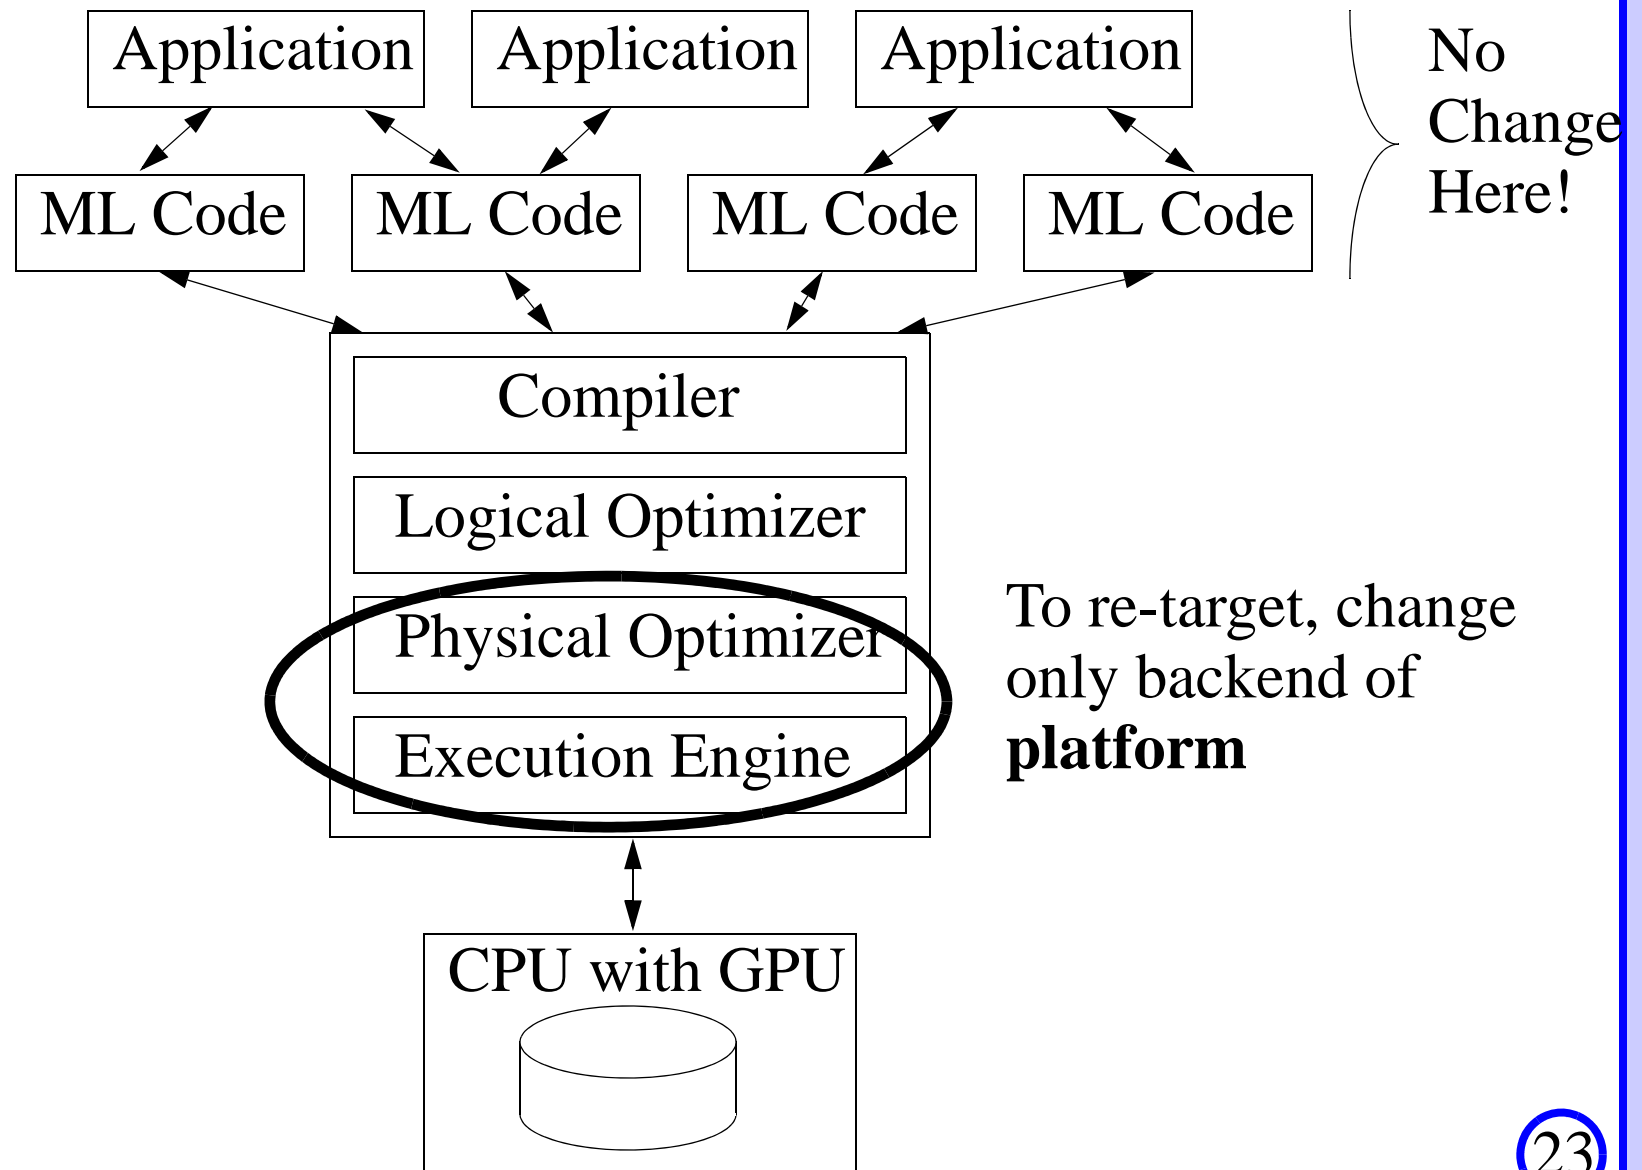


Much Better In System Based On Data Ind.

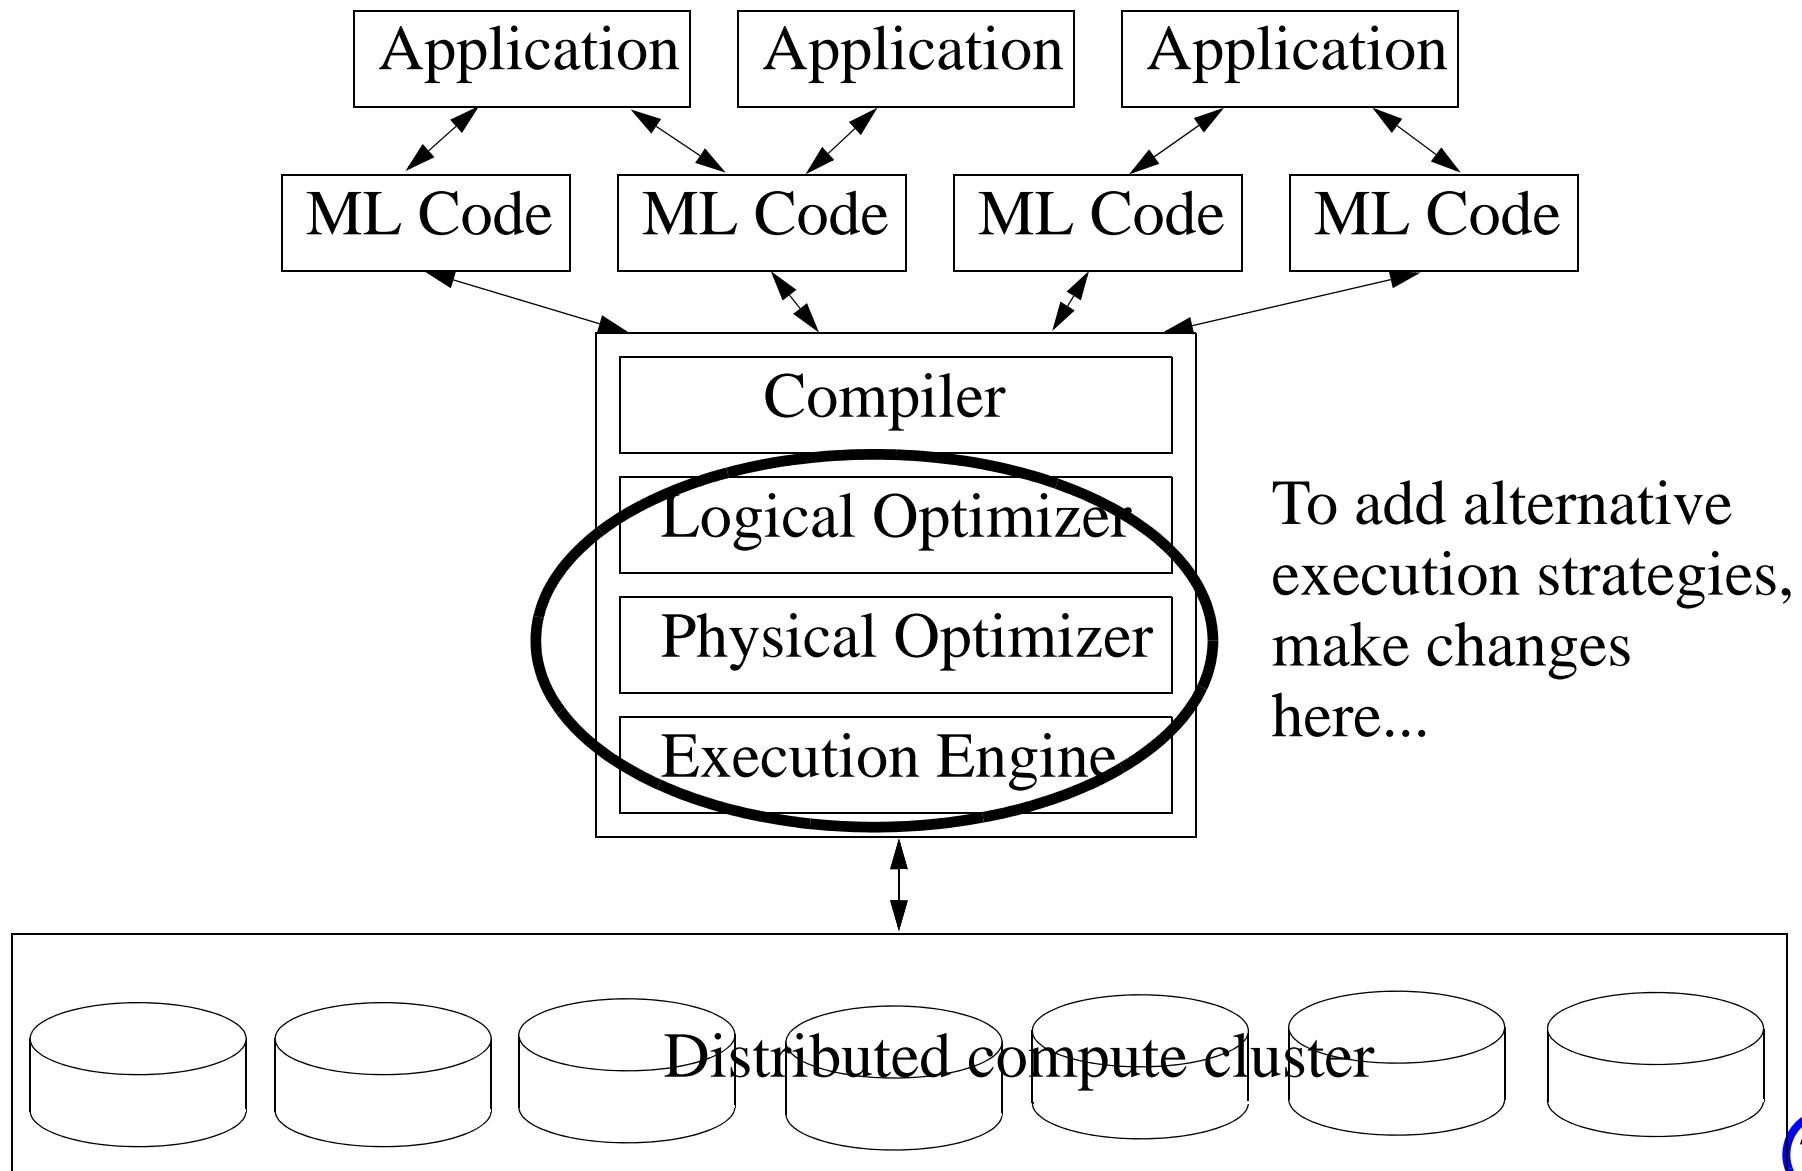
# Much Better In System Based On Data Ind.



# Much Better In System Based On Data Ind.

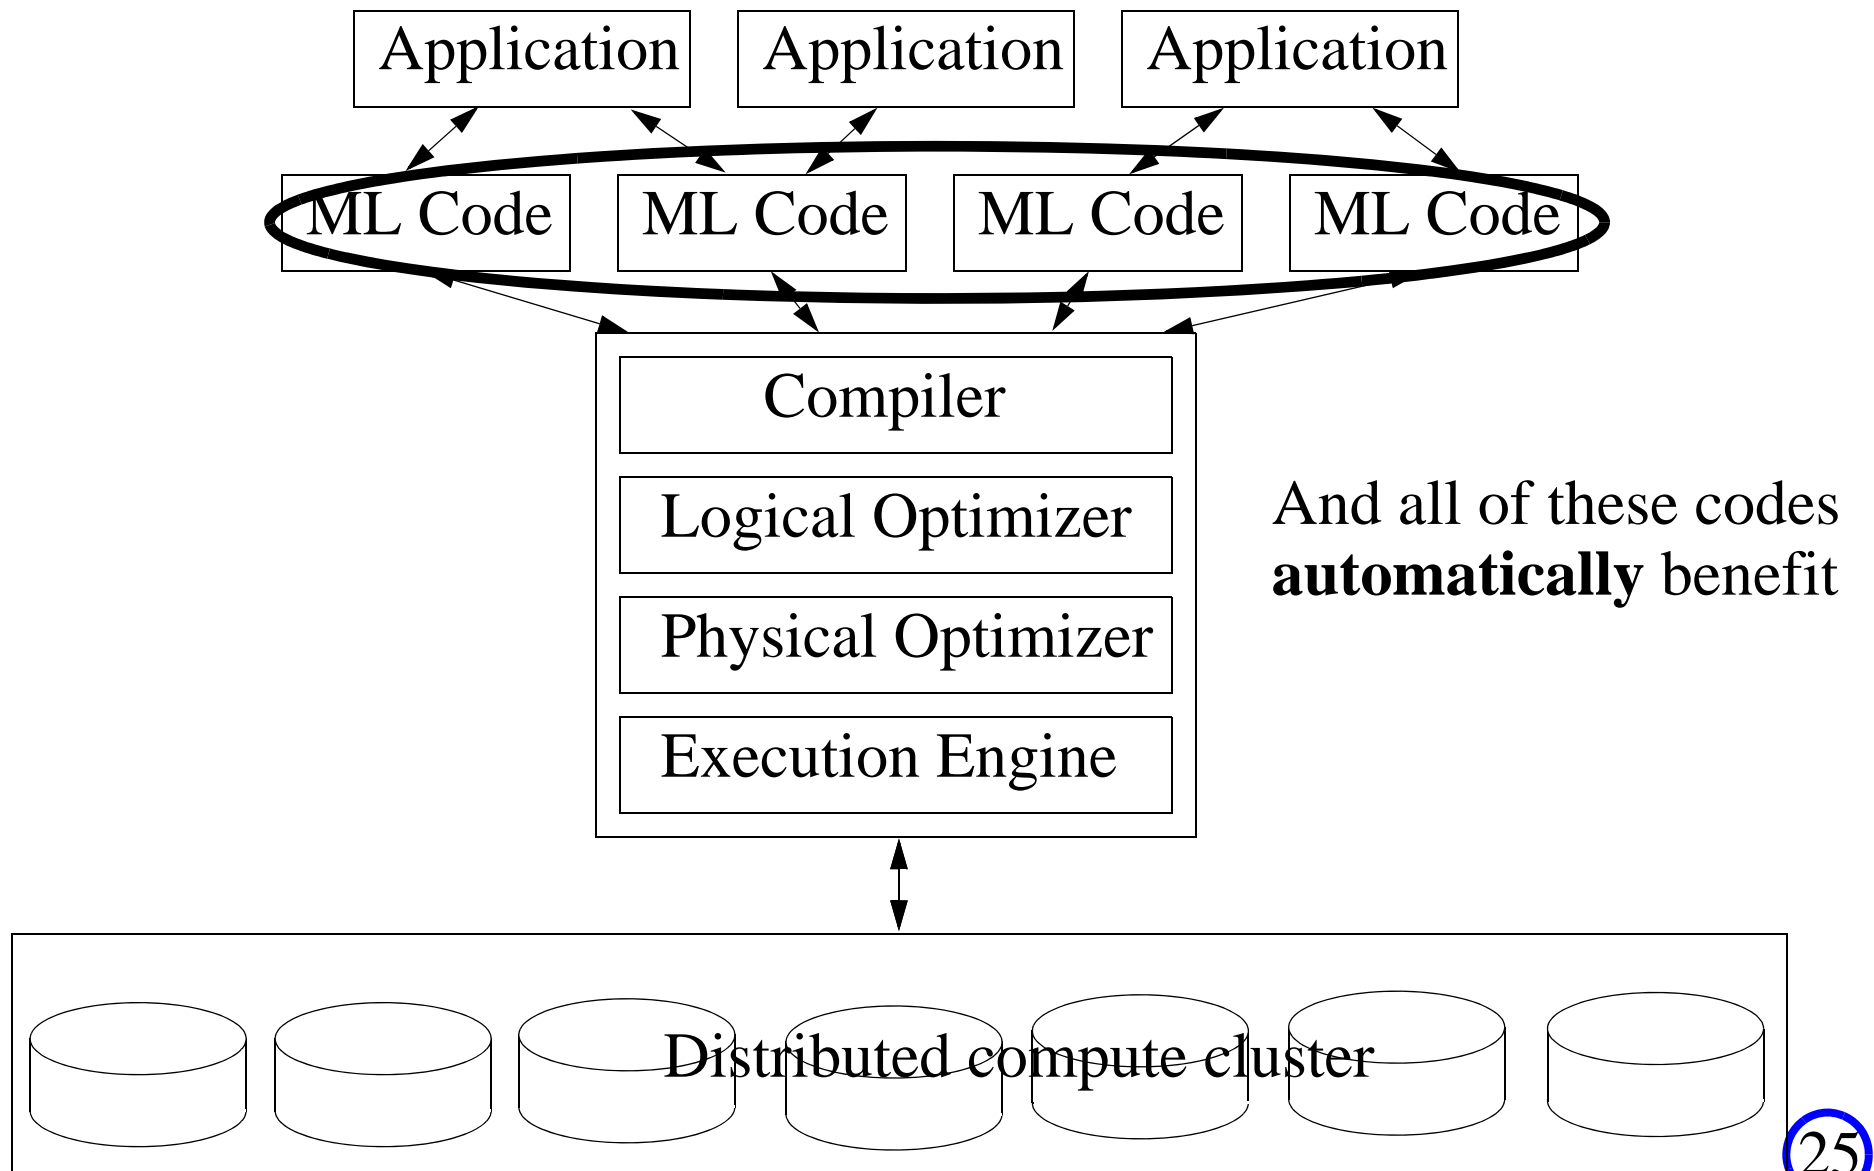


# Much Better In System Based On Data Ind.





# Much Better In System Based On Data Ind.



# Huge Win

- Once written, code at top of stack can remain unchanged
  - #1 selling point of RDBMS tech for 30 years!

# Huge Win

- Once written, code at top of stack can remain unchanged
  - #1 selling point of RDBMS tech for 30 years!
- Unfortunately, not accepted by ML community
  - ML people write great ML-on-GPU papers
  - They design platforms such as GraphLab (higher-level, MPI-like framework)

# Huge Win

- Once written, code at top of stack can remain unchanged
  - #1 selling point of RDBMS tech for 30 years!
- Unfortunately, not accepted by ML community
  - ML people write great ML-on-GPU papers
  - They design platforms such as GraphLab (higher-level, MPI-like framework)
- Some in DB community have looked at declarative dataflow...
  - Spark SQL on Spark
  - Meteor on Stratosphere
  - Asterix from UC Irvine

# Huge Win

- Once written, code at top of stack can remain unchanged
  - #1 selling point of RDBMS tech for 30 years!
- Unfortunately, not accepted by ML community
  - ML people write great ML-on-GPU papers
  - They design platforms such as GraphLab (higher-level, MPI-like framework)
- Some in DB community have looked at declarative dataflow...
  - Spark SQL on Spark
  - Meteor on Stratosphere
  - Asterix from UC Irvine
- But this is far from declarative ML
  - The codes don't look anything like math!

# The Goal

- Start with mathematical spec of learning algorithm...

1.  $r \sim \text{Normal}(\mathbf{A}^{-1} \mathbf{X}^T \tilde{\mathbf{y}}, \sigma^2 \mathbf{A}^{-1})$

2.  $\sigma^2 \sim \text{InvGamma}\left(\frac{(n-1)+p}{2}, \frac{(\tilde{\mathbf{y}} - \mathbf{X}r)^T (\tilde{\mathbf{y}} - \mathbf{X}r) + r^T \mathbf{D}^{-1} r}{2}\right)$

3.  $\tau_j^{-2} \sim \text{InvGaussian}\left(\frac{\lambda \sigma}{r_j}, \lambda^2\right)$

— where  $\mathbf{A} = \mathbf{X}^T \mathbf{X} + \mathbf{D}^{-1}$ ,  $\mathbf{D}^{-1} = \text{diag}(\tau_1^{-2}, \tau_2^{-2}, \dots)$

This is math for the Bayesian Lasso, lifted from original paper

— Bayesian regression model with regularizing prior on regression coefs

# The Goal

- Programmer writes code that looks just like the math...

```
data {  
  n: range (responses); p: range (regressors);  
  X: array[n, p] of real; y: array[n] of real;  
  lam: real  
}  
  
var {  
  sig: real;  
  r, t: array[p] of real; yy, Z: array[n] of real;  
}  
  
A <- inv(X `* X + diag(t));  
yy <- (y[i] - mean(y) | i in 1:n);  
Z <- yy - X * r;  
  
init {  
  sig ~ InvGamma (1, 1);  
  t ~ (InvGauss (1, lam) | j in 1:p);  
}  
  
r ~ Normal (A *' X * yy, sig * A);  
sig ~ InvGamma(((n-1) + p)/2,  
  (Z `* Z + (r * diag(t) `* r)) / 2);  
for (j in 1:p) {  
  t[j] ~ InvGauss (sqrt((lam * sig) / r[j]), lam);  
}
```

We call our  
language “BUDS”

# The Goal

- Write code that looks just like the math...

```
data {
  n: range (responses); p: range (regressors);
  X: array[n, p] of real; y: array[n] of real;
  lam: real
}
```

```
var {
  sig: real;
  r, t: array[p] of real; yy, Z: array[n] of real;
}
```

```
A <- inv(X `*` X + diag(t));  $\mathbf{A} = \mathbf{X}^T \mathbf{X} + \mathbf{D}^{-1}$ 
yy <- (y[i] - mean(y) | i in 1:n);
Z <- yy - X * r;
```

```
init {
  sig ~ InvGamma (1, 1);
  t ~ (InvGauss (1, lam) | j in 1:p);
}
```

```
r ~ Normal (A *' X * yy, sig * A);  $r \sim \text{Normal}(\mathbf{A}^{-1} \mathbf{X}^T \tilde{\mathbf{y}}, \sigma^2 \mathbf{A}^{-1})$ 
sig ~ InvGamma(((n-1) + p)/2,  $\sigma^2 \sim \text{InvGamma}\left(\frac{(n-1) + p}{2}, \frac{(\tilde{\mathbf{y}} - \mathbf{X}r)^T (\tilde{\mathbf{y}} - \mathbf{X}r)^T + r^T \mathbf{D}^{-1} r}{2}\right)$ 
  (Z `*` Z + (r * diag(t) `*` r)) / 2);
for (j in 1:p) {
  t[j] ~ InvGauss (sqrt((lam * sig) / r[j]), lam);  $\tau_j^{-2} \sim \text{InvGaussian}\left(\frac{\lambda \sigma}{r_j}, \lambda^2\right)$ 
}
```



# The Goal

- And the system compiles and executes this for a huge data set
  - On hundreds or thousands of machines...
  - Or on a desktop with a GPU...
  - Or for whatever backend the system can target...

## Also Important

- We don't want to be like everyone and argue for a new DA stack
  - The world has too many dataflow platforms already

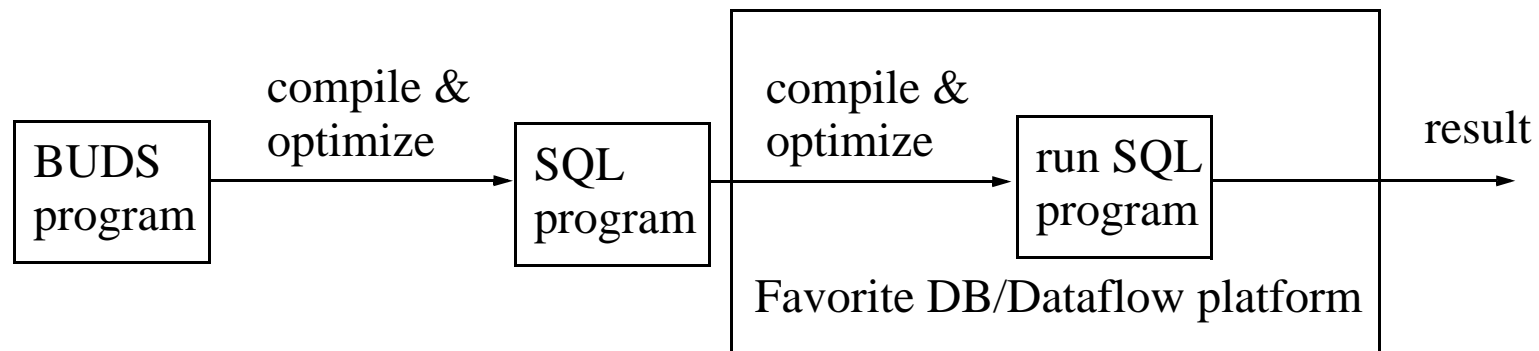
# The Fact Is

- SQL or SQL-like declarative language is the LCD for all systems
  - Including commercial databases
  - Free databases
  - And newfangled dataflow platforms

# The Fact Is

- SQL or SQL-like declarative language is the LCD for all systems
  - Including commercial databases
  - Free databases
  - And newfangled dataflow platforms

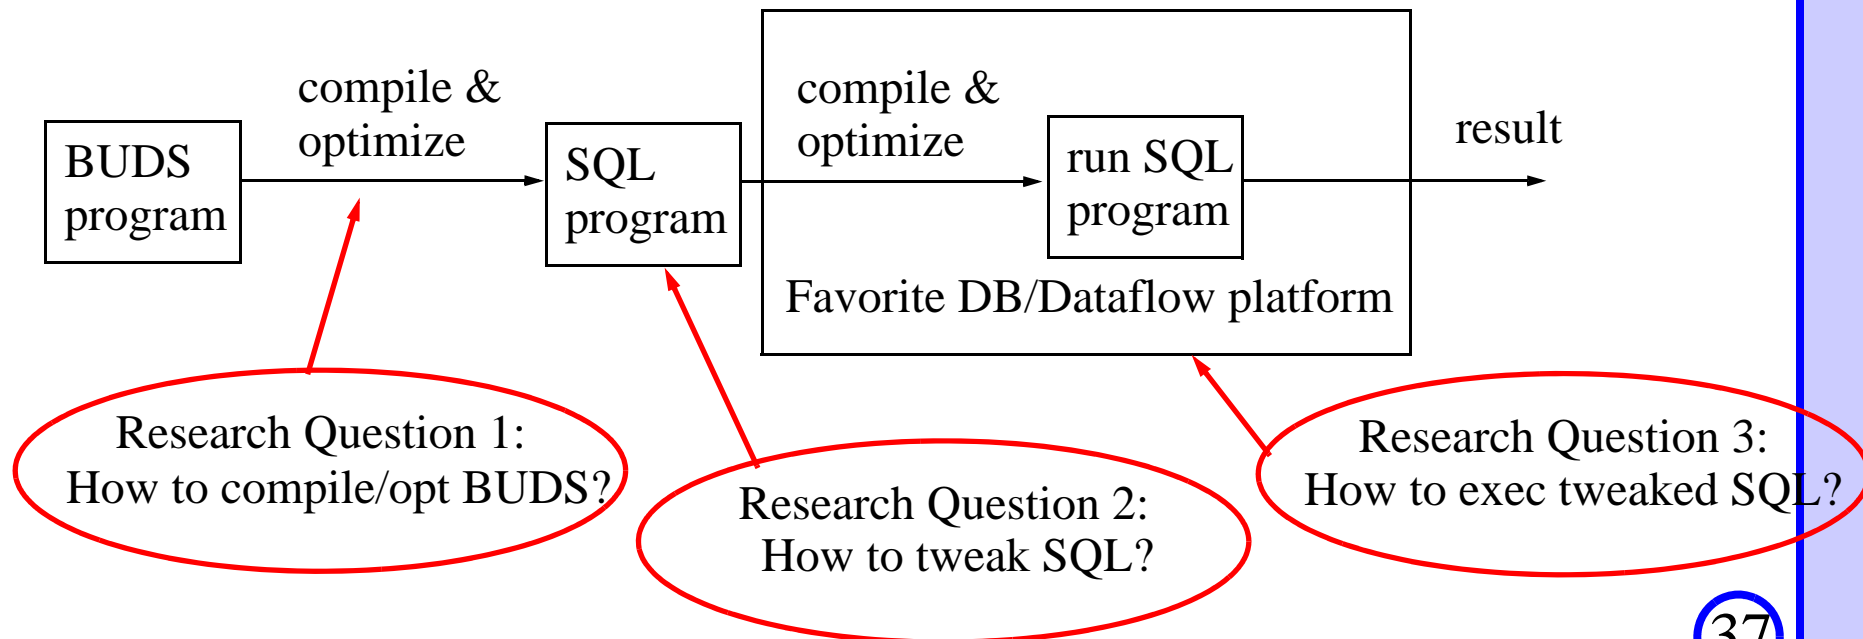
So here's the workflow we envision...



# The Fact Is

- SQL or SQL-like declarative language is the LCD for all systems
  - Including commercial databases
  - Free databases
  - And newfangled dataflow platforms

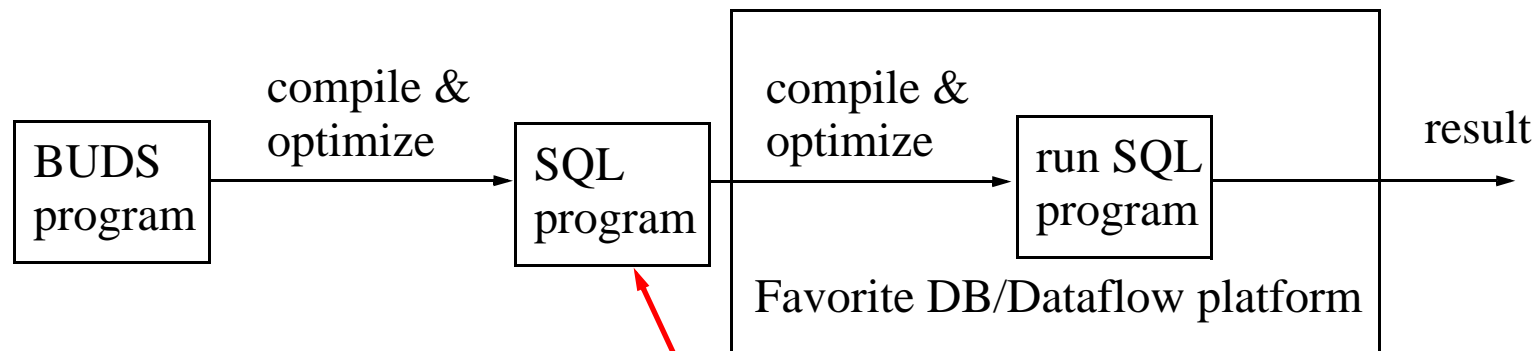
So here's the workflow we envision...



# The Fact Is

- SQL or SQL-like declarative language is the LCD for all systems
  - Including commercial databases
  - Free databases
  - And newfangled dataflow platforms

So here's the workflow we envision...



Focus of rest of talk  
is mostly here...

Research Question 2:  
How to tweak SQL?

Research Question 3:  
How to exec tweaked SQL?

# So, How Must SQL/DA Platform Change?

- More extensive support for recursion
- Fancier table functions (“VG functions”)
- Add native support for vectors/matrices (as att types)
- Support for executing huge “query” plans (1000’s of operations)
- A few new logical/physical operators

# So, How Must SQL/DA Platform Change?

- More extensive support for recursion
- Fancier table functions (“VG functions”)
- Add native support for vectors/matrices (as att types)
- Support for executing huge “query” plans (1000’s of operations)
- A few new logical/physical operators

- We built the SimSQL system to demo how this works
  - Simple shared nothing, parallel DBMS
  - 100K SLOC
  - Java, C++, Prolog
  - Runs queries as Map-only jobs on Hadoop



# So, How Must SQL/DA Platform Change?

- More extensive support for recursion
- Fancier table functions (“VG functions”)
- Add native support for vectors/matrices (as att types)
- Support for executing huge “query” plans (1000’s of operations)
- A few new logical/physical operators

- We built the SimSQL system to demo how this works
  - Simple shared nothing, parallel DBMS
  - 100K SLOC
  - Java, C++, Prolog
  - Runs queries as Map-only jobs on Hadoop

# SimSQL's Specialized for Stochastic Algs

- Due to my own Bayesian bias
  - Though if you can do stochastic, you can do deterministic
- So I'll make a brief foray into MCMC...

# MCMC

- Standard Bayesian ML inference method
- Idea is to simulate a Markov chain
- Whose *stationary distribution* is equal to the target posterior
  - Means that if you run forever then stop, have sample from the target
  - In theory, can be used with virtually any target distribution

# MCMC: Gibbs Sampling


- Many MCMC algorithms; useful example is Gibbs sampling

— Unknown vars/params in  $\theta$ ; state of chain is described by  $\bar{\theta}$

1. Pick subset  $\theta' \subseteq \theta$  (without looking at  $\bar{\theta}$ !)

2. Sample  $\bar{\theta}' \sim f(\theta' | (\bar{\theta} - \bar{\theta}'), X)$

3. Repeat forever!



# MCMC: Gibbs Sampling


- Many MCMC algorithms; useful example is Gibbs sampling

— Unknown vars/params in  $\theta$ ; state of chain is described by  $\bar{\theta}$

1. Pick subset  $\theta' \subseteq \theta$  (without looking at  $\bar{\theta}$ !)

2. Sample  $\bar{\theta}' \sim f(\theta' | (\bar{\theta} - \bar{\theta}'), X)$

3. Repeat forever!



# MCMC: Gibbs Sampling


- Many MCMC algorithms; useful example is Gibbs sampling

— Unknown vars/params in  $\theta$ ; state of chain is described by  $\bar{\theta}$

1. Pick subset  $\theta' \subseteq \theta$  (without looking at  $\bar{\theta}$ !)

2. Sample  $\bar{\theta}' \sim f(\theta' | (\bar{\theta} - \bar{\theta}'), X)$

3. Repeat forever!




# MCMC: Gibbs Sampling

- Many MCMC algorithms; useful example is Gibbs sampling

— Unknown vars/params in  $\theta$ ; state of chain is described by  $\bar{\theta}$

1. Pick subset  $\theta' \subseteq \theta$  (without looking at  $\bar{\theta}$ !)
2. Sample  $\bar{\theta}' \sim f(\theta' | (\bar{\theta} - \bar{\theta}'), X)$
3. Repeat forever!



# MCMC: Gibbs Sampling


- Many MCMC algorithms; useful example is Gibbs sampling

— Unknown vars/params in  $\theta$ ; state of chain is described by  $\bar{\theta}$

1. Pick subset  $\theta' \subseteq \theta$  (without looking at  $\bar{\theta}$ !)

2. Sample  $\bar{\theta}' \sim f(\theta' | (\bar{\theta} - \bar{\theta}'), X)$

3. Repeat forever!





# MCMC: Gibbs Sampling


- Many MCMC algorithms; useful example is Gibbs sampling

— Unknown vars/params in  $\theta$ ; state of chain is described by  $\bar{\theta}$

1. Pick subset  $\theta' \subseteq \theta$  (without looking at  $\bar{\theta}$ !)

2. Sample  $\bar{\theta}' \sim f(\theta' | (\bar{\theta} - \bar{\theta}'), X)$

3. Repeat forever!



# MCMC: Gibbs Sampling


- Many MCMC algorithms; useful example is Gibbs sampling

— Unknown vars/params in  $\theta$ ; state of chain is described by  $\bar{\theta}$

1. Pick subset  $\theta' \subseteq \theta$  (without looking at  $\bar{\theta}$ !)

2. Sample  $\bar{\theta}' \sim f(\theta' | (\bar{\theta} - \bar{\theta}'), X)$

3. Repeat forever!

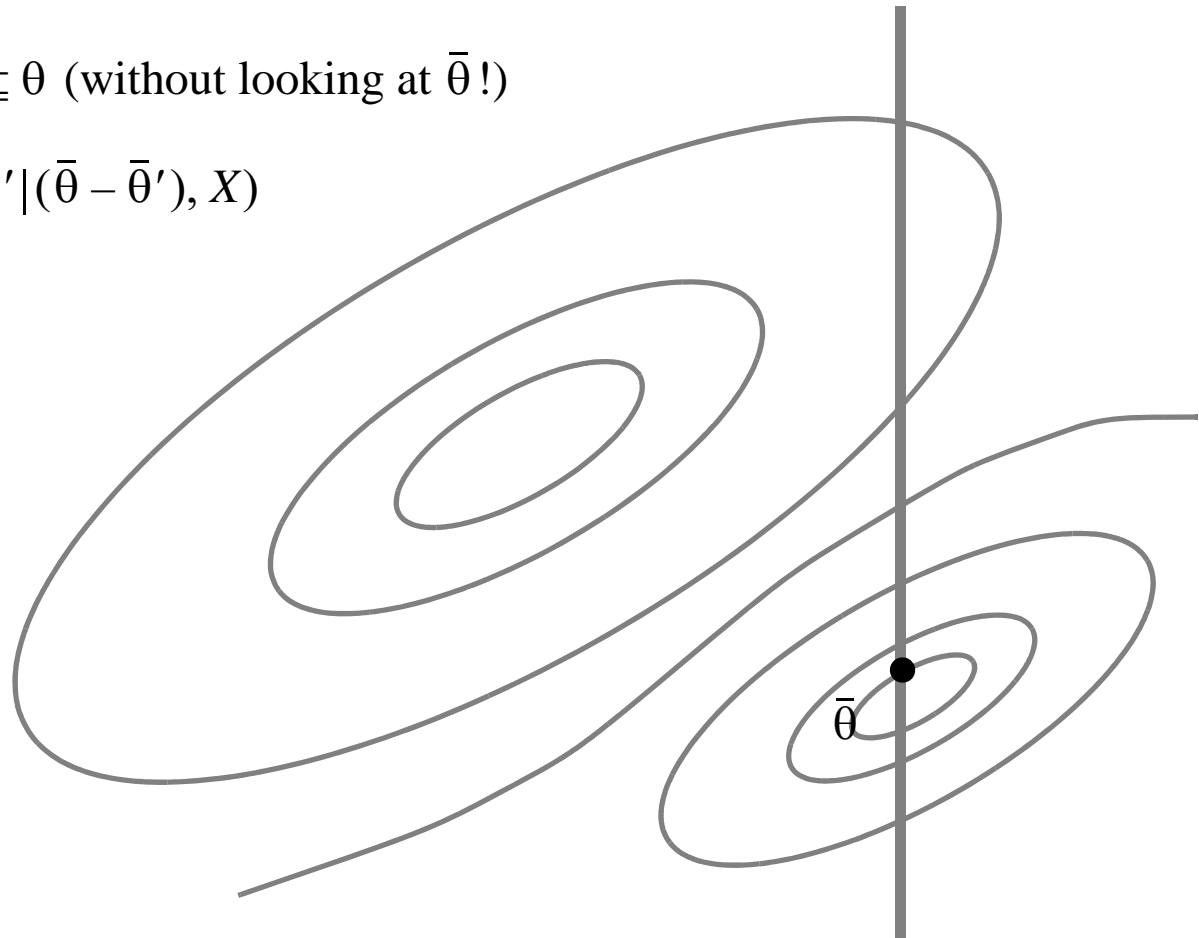


# MCMC: Gibbs Sampling

- Many MCMC algorithms; useful example is Gibbs sampling

— Unknown vars/params in  $\theta$ ; state of chain is described by  $\bar{\theta}$

1. Pick subset  $\theta' \subseteq \theta$  (without looking at  $\bar{\theta}$ !)
2. Sample  $\bar{\theta}' \sim f(\theta' | (\bar{\theta} - \bar{\theta}'), X)$
3. Repeat forever!



# MCMC: Gibbs Sampling

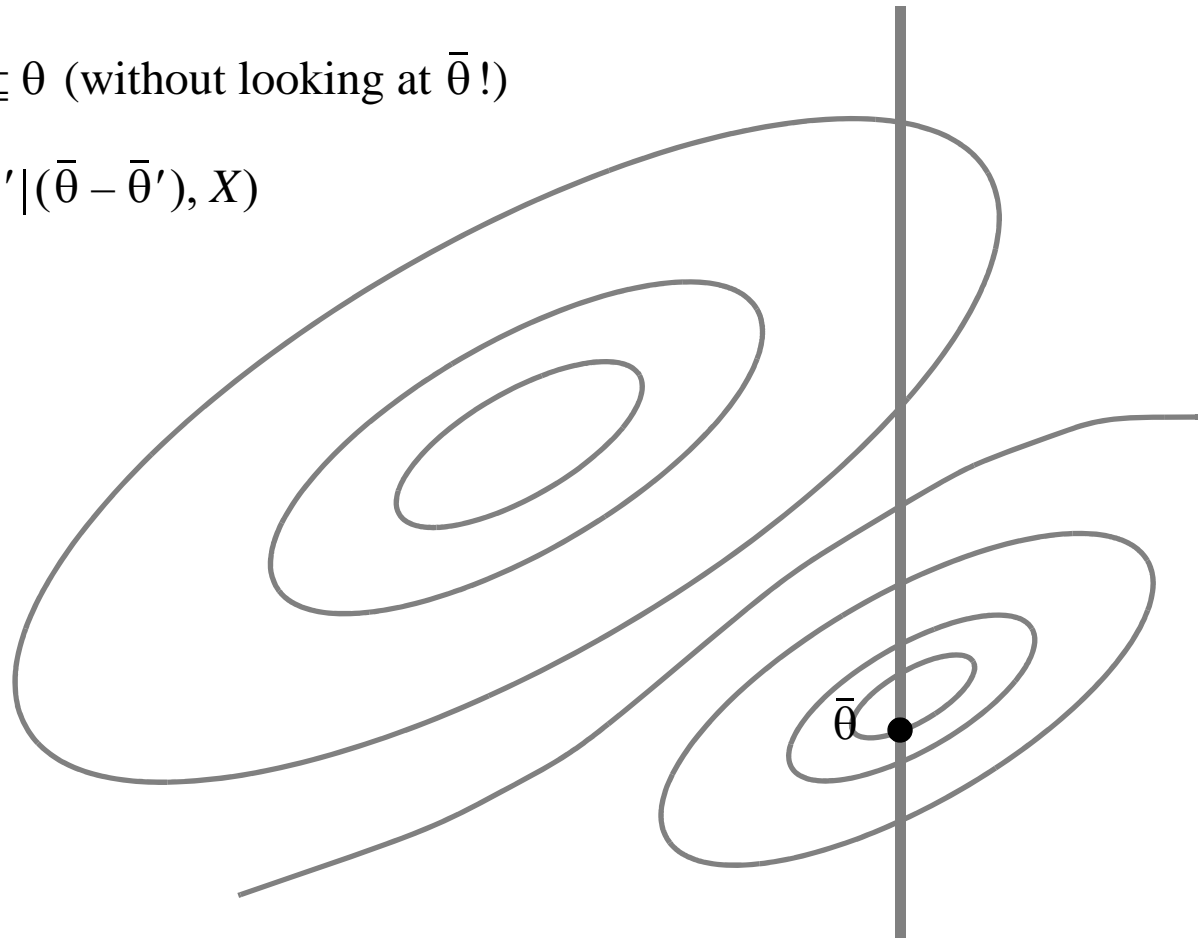
- Many MCMC algorithms; useful example is Gibbs sampling

— Unknown vars/params in  $\theta$ ; state of chain is described by  $\bar{\theta}$

1. Pick subset  $\theta' \subseteq \theta$  (without looking at  $\bar{\theta}$ !)

2. Sample  $\bar{\theta}' \sim f(\theta' | (\bar{\theta} - \bar{\theta}'), X)$

3. Repeat forever!



# MCMC: Gibbs Sampling


- Many MCMC algorithms; useful example is Gibbs sampling

— Unknown vars/params in  $\theta$ ; state of chain is described by  $\bar{\theta}$

1. Pick subset  $\theta' \subseteq \theta$  (without looking at  $\bar{\theta}$ !)

2. Sample  $\bar{\theta}' \sim f(\theta' | (\bar{\theta} - \bar{\theta}'), X)$

3. Repeat forever!



# MCMC: Gibbs Sampling


- Many MCMC algorithms; useful example is Gibbs sampling

— Unknown vars/params in  $\theta$ ; state of chain is described by  $\bar{\theta}$

1. Pick subset  $\theta' \subseteq \theta$  (without looking at  $\bar{\theta}$ !)

2. Sample  $\bar{\theta}' \sim f(\theta' | (\bar{\theta} - \bar{\theta}'), X)$

3. Repeat forever!



# MCMC: Gibbs Sampling


- Many MCMC algorithms; useful example is Gibbs sampling

— Unknown vars/params in  $\theta$ ; state of chain is described by  $\bar{\theta}$

1. Pick subset  $\theta' \subseteq \theta$  (without looking at  $\bar{\theta}$ !)

2. Sample  $\bar{\theta}' \sim f(\theta' | (\bar{\theta} - \bar{\theta}'), X)$

3. Repeat forever!




# MCMC: Gibbs Sampling

- Many MCMC algorithms; useful example is Gibbs sampling

— Unknown vars/params in  $\theta$ ; state of chain is described by  $\bar{\theta}$

1. Pick subset  $\theta' \subseteq \theta$  (without looking at  $\bar{\theta}$ !)
2. Sample  $\bar{\theta}' \sim f(\theta' | (\bar{\theta} - \bar{\theta}'), X)$
3. Repeat forever!





# MCMC: Gibbs Sampling


- Many MCMC algorithms; useful example is Gibbs sampling

— Unknown vars/params in  $\theta$ ; state of chain is described by  $\bar{\theta}$

1. Pick subset  $\theta' \subseteq \theta$  (without looking at  $\bar{\theta}$ !)

2. Sample  $\bar{\theta}' \sim f(\theta' | (\bar{\theta} - \bar{\theta}'), X)$

3. Repeat forever!



# MCMC: Gibbs Sampling


- Many MCMC algorithms; useful example is Gibbs sampling

— Unknown vars/params in  $\theta$ ; state of chain is described by  $\bar{\theta}$

1. Pick subset  $\theta' \subseteq \theta$  (without looking at  $\bar{\theta}$ !)

2. Sample  $\bar{\theta}' \sim f(\theta' | (\bar{\theta} - \bar{\theta}'), X)$

3. Repeat forever!



# How To Do MCMC Inference Over Big Data?

- Easy to spec MCMC simulations in SimSQL SQL

# SimSQL's Version of SQL

- Most fundamental SQL addition is “VG Function” abstraction
- Called via a special, stochastic CREATE TABLE statement
- Example; assuming:

- SBP(MEAN, STD, GENDER)
  - PATIENTS(NAME, GENDER)

- To create a stochastic table, we might have:

```
CREATE TABLE SBP_DATA(NAME, GENDER, SBP) AS
FOR EACH p in PATIENTS
  WITH Res AS Normal (
    SELECT s.MEAN, s.STD
    FROM SPB s WHERE s.GENDER = p.GENDER)
  SELECT p.NAME, p.GENDER, r.VALUE
FROM Res r
```

# How Does This Work?

```
CREATE TABLE SBP_DATA(NAME, GENDER, SBP) AS
FOR EACH p in PATIENTS
  WITH Res AS Normal (
    SELECT s.MEAN, s.STD
    FROM SPB s WHERE s.GENDER = p.GENDER)
  SELECT p.NAME, p.GENDER, r.VALUE
FROM Res r
```

← Loop through PATIENTS

PATIENTS (NAME, GENDER)
( <b>Joe, Male</b> ) "p"
(Tom, Male)
(Jen, Female)
(Sue, Female)
(Jim, Male)

SBP(MEAN, STD, GENDER)
(150, 20, Male)
(130, 25, Female)

# How Does This Work?

```
CREATE TABLE SBP_DATA(NAME, GENDER, SBP) AS
FOR EACH p in PATIENTS
  WITH Res AS Normal (
    SELECT s.MEAN, s.STD
    FROM SPB s WHERE s.GENDER = p.GENDER)
  SELECT p.NAME, p.GENDER, r.VALUE
FROM Res r
```

PATIENTS (NAME, GENDER)
( <b>Joe, Male</b> ) "p"
(Tom, Male)
(Jen, Female)
(Sue, Female)
(Jim, Male)

SBP(MEAN, STD, GENDER)
(150, 20, <b>Male</b> )
(130, 25, Female)

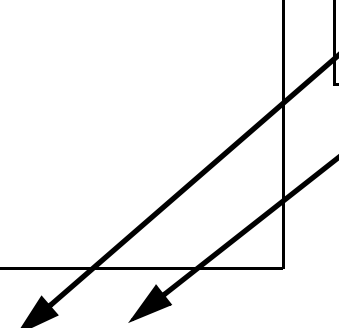
# How Does This Work?

```
CREATE TABLE SBP_DATA(NAME, GENDER, SBP) AS
FOR EACH p in PATIENTS
  WITH Res AS Normal (
    SELECT s.MEAN, s.STD
    FROM SPB s WHERE s.GENDER = p.GENDER)
  SELECT p.NAME, p.GENDER, r.VALUE
FROM Res r
```

PATIENTS (NAME, GENDER)
(Joe, Male) "p"
(Tom, Male)
(Jen, Female)
(Sue, Female)
(Jim, Male)

SBP(MEAN, STD, GENDER)
(150, 20, Male)
(130, 25, Female)

Normal(150,20)



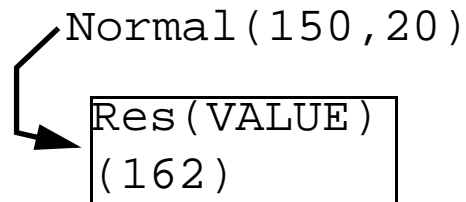
# How Does This Work?

```
CREATE TABLE SBP_DATA(NAME, GENDER, SBP) AS
FOR EACH p in PATIENTS
  WITH Res AS Normal (
    SELECT s.MEAN, s.STD
    FROM SPB s WHERE s.GENDER = p.GENDER)
  SELECT p.NAME, p.GENDER, r.VALUE
FROM Res r
```

PATIENTS (NAME, GENDER)
(Joe, Male) "p"
(Tom, Male)
(Jen, Female)
(Sue, Female)
(Jim, Male)

SBP(MEAN, STD, GENDER)
(150, 20, Male)
(130, 25, Female)

Normal(150,20)



Res(VALUE)
(162)



# How Does This Work?

```
CREATE TABLE SBP_DATA(NAME, GENDER, SBP) AS
FOR EACH p in PATIENTS
  WITH Res AS Normal (
    SELECT s.MEAN, s.STD
    FROM SPB s WHERE s.GENDER = p.GENDER)
  SELECT p.NAME, p.GENDER, r.VALUE
FROM Res r
```

PATIENTS (NAME, GENDER)
(Joe, Male) "p"
(Tom, Male)
(Jen, Female)
(Sue, Female)
(Jim, Male)

SPB(MEAN, STD, GENDER)
(150, 20, Male)
(130, 25, Female)

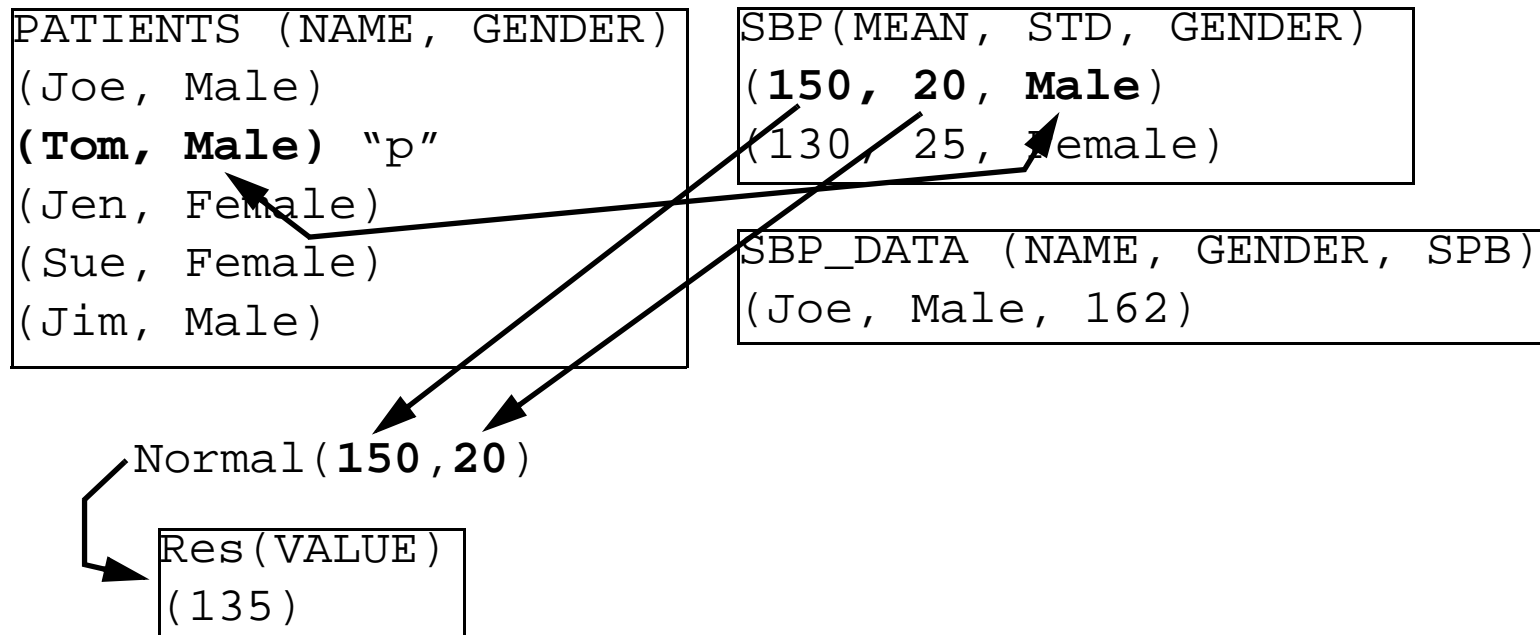
SBP_DATA (NAME, GENDER, SPB)
(Joe, Male, 162)

Normal(150,20)

Res(VALUE)
(162)

# How Does This Work?

```
CREATE TABLE SBP_DATA(NAME, GENDER, SBP) AS
FOR EACH p in PATIENTS
  WITH Res AS Normal (
    SELECT s.MEAN, s.STD
    FROM SPB s WHERE s.GENDER = p.GENDER)
  SELECT p.NAME, p.GENDER, r.VALUE
FROM Res r
```



# How Does This Work?

```
CREATE TABLE SBP_DATA(NAME, GENDER, SBP) AS
FOR EACH p in PATIENTS
  WITH Res AS Normal (
    SELECT s.MEAN, s.STD
    FROM SPB s WHERE s.GENDER = p.GENDER)
  SELECT p.NAME, p.GENDER, r.VALUE
FROM Res r
```

PATIENTS (NAME, GENDER)	
Joe	Male
<b>Tom</b>	<b>Male</b>
Jen	Female
Sue	Female
Jim	Male

SPB(MEAN, STD, GENDER)		
150	20	Male
130	25	Female

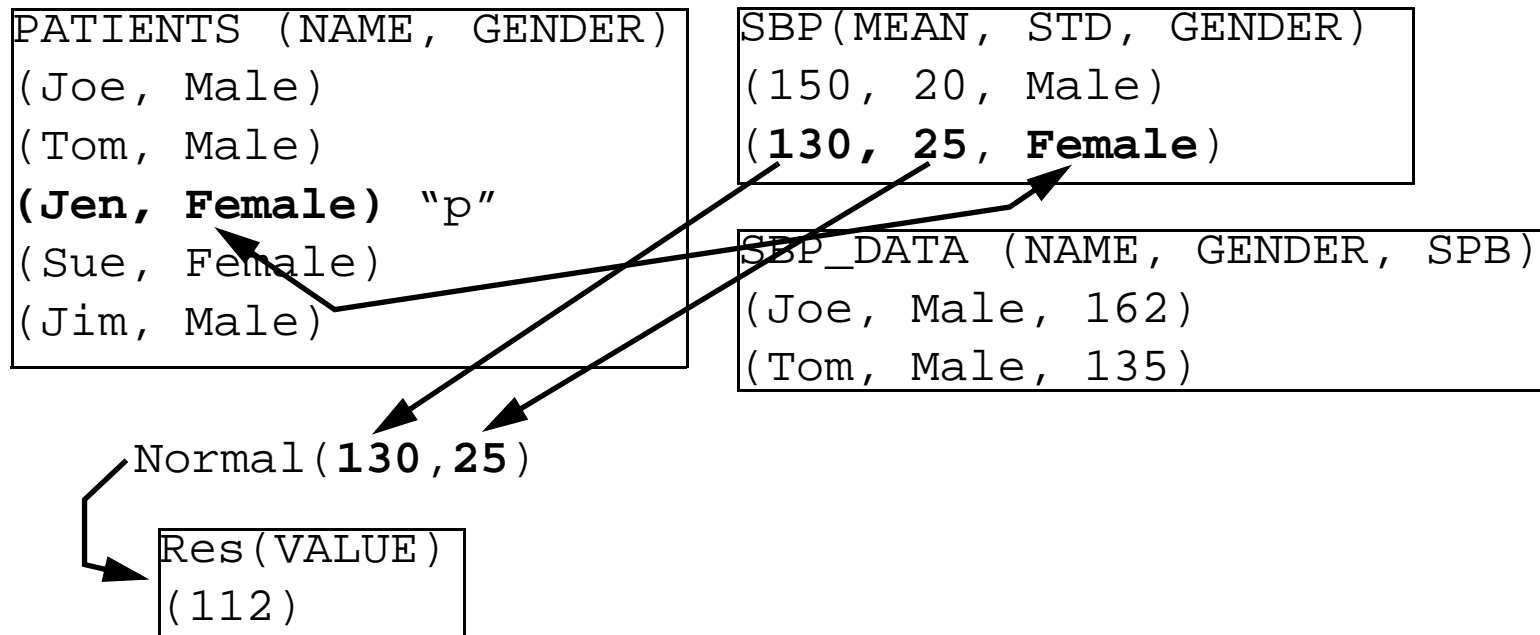
SBP_DATA (NAME, GENDER, SPB)		
Joe	Male	162
<b>Tom</b>	<b>Male</b>	<b>135</b>

Normal(150,20)

Res(VALUE)
<b>(135)</b>

# How Does This Work?

```
CREATE TABLE SBP_DATA(NAME, GENDER, SBP) AS
FOR EACH p in PATIENTS
  WITH Res AS Normal (
    SELECT s.MEAN, s.STD
    FROM SPB s WHERE s.GENDER = p.GENDER)
  SELECT p.NAME, p.GENDER, r.VALUE
FROM Res r
```



# How Does This Work?

```
CREATE TABLE SBP_DATA(NAME, GENDER, SBP) AS
FOR EACH p in PATIENTS
  WITH Res AS Normal (
    SELECT s.MEAN, s.STD
    FROM SPB s WHERE s.GENDER = p.GENDER)
  SELECT p.NAME, p.GENDER, r.VALUE
FROM Res r
```

PATIENTS (NAME, GENDER)	
(Joe, Male)	
(Tom, Male)	
<b>(Jen, Female)</b>	"p"
(Sue, Female)	
(Jim, Male)	

SBP(MEAN, STD, GENDER)		
(150, 20, Male)		
(130, 25, Female)		

Normal(130, 25)

Res(VALUE)
<b>(112)</b>

SBP_DATA (NAME, GENDER, SPB)		
(Joe, Male, 162)		
(Tom, Male, 135)		
<b>(Jen, Female, 112)</b>		

and so on...

# Markov Chain Simulation

- Previous allows for table-valued RVs, not for Markov chains
- But Markov chains are easy in SimSQL
- Here's a silly Markov chain. We have:
  - `PERSON (name)`
  - `LOCTION (name, dim, val)`
  - `MOVEMENT_VAR (name, dim1, dim2, var)`
  - `MOVEMENT_MEAN (name, dim, mean)`
- We want to randomly start each person at a location
- Then move them all randomly around

# Markov Chain Simulation

- To select an initial starting position for each person:

```
CREATE TABLE POSITION[0] (name, dim, val) AS
FOR EACH p IN PERSON
  WITH Pos AS DiscreteChoice (
    SELECT DISTINCT name
    FROM LOCATION)
  SELECT p.name, l.dim, l.val
  FROM Pos, LOCATION l
  WHERE l.name = Pos.val
```

# Markov Chain Simulation

- And then to move them all along:

```
CREATE TABLE POSITION[i] (name, dim, val) AS
FOR EACH p IN PERSON
  WITH Pos AS ConditionalNormal (
    (SELECT pos.dim, pos.val
     FROM POSITION[i - 1] pos
     WHERE pos.dim = i MOD 2 AND pos.name = p.name)
    (SELECT m.dim1, m.dim2, m.var
     FROM MOVEMENT_VAR m
     WHERE m.name = p.name)
    (SELECT m.dim, m.mean
     FROM MOVEMENT_MEAN m
     WHERE m.name = p.name))
  SELECT p.name, Pos.dim, Pos.val
FROM Pos
```

- Now we've fully spec'd a distributed Markov chain simulation!



# Getting This To Run

- Can use a lot of standard parallel DB techniques to implement
- But some problems are quite unique to SimSQL
  - No time to talk about them today!
  - Perhaps informally at end of talk?

# How Well Does All of This Work?

- SimSQL is great in theory...
  - Many will buy the “data independence” argument
  - Will appreciate being able to specify algs at a very high level
- But isn't the declarative approach gonna be slow?

# How Well Does All of This Work?

- SimSQL is great in theory...
  - Many will buy the “data independence” argument
  - Will appreciate being able to specify algs at a very high level
- But isn't the declarative approach gonna be slow?
- Yes, it's slow, compared to C/Fortran + MPI
  - But zero data independence with MPI

# How Well Does All of This Work?


- SimSQL is great in theory...
  - Many will buy the “data independence” argument
  - Will appreciate being able to specify algs at a very high level
- But isn't the declarative approach gonna be slow?
- Yes, it's slow, compared to C/Fortran + MPI
  - But zero data independence with MPI
- But does it compete well with other “Big Data” ML platforms?
  - After all, are many that count ML as the primary (or a motivating) application
  - OptiML, GraphLab, SystemML, MLBase, ScalOps, Pregel, Giraph, Hama, Spark, Ricardo, Nyad, DradLinq
  - How might those compare?

# How Well Does All of This Work?

- We've done a **LOT** of comparisons with other mature platforms
  - Specifically, GraphLab, Giraph, Spark
  - More than 70,000 hours of Amazon EC2 time (\$100,000 @on-demand price)
  - I'd wager that few groups have a better understanding of how well these platforms work in practice!
- Note: point is not to show SimSQL is the fastest (it is not)
  - Only to argue that it can compete well
  - If it competes, it's a strong argument for the declarative approach to ML
- Note: this is hand-coded SimSQL SQL
  - Not SQL compiled from BUDS
  - Will get those results soon!


# Example One: Bayesian GMM

Generative process:



# Example One: Bayesian GMM


Generative process:  
(1) **Pick a cluster**



# Example One: Bayesian GMM

Generative process:

- (1) Pick a cluster
- (2) **Use it to generate point**






# Example One: Bayesian GMM

Generative process:


- (1) **Pick a cluster**
- (2) Use it to generate point



# Example One: Bayesian GMM

Generative process:


- (1) Pick a cluster
- (2) **Use it to generate point**



# Example One: Bayesian GMM

Generative process:


- (1) **Pick a cluster**
- (2) Use it to generate point



# Example One: Bayesian GMM


Generative process:

- (1) Pick a cluster
- (2) **Use it to generate point**




# Example One: Bayesian GMM

Then given **this**



# Example One: Bayesian GMM

Infer **this**



# Example One: Bayesian GMM

- Implemented relevant MCMC simulation on all four platforms
  - SimSQL, GraphLab, Spark, Giraph

# Example One: Bayesian GMM

- Implemented relevant MCMC simulation on all four platforms
  - SimSQL, GraphLab, Spark, Giraph
- Philosophy: be true to the platform
  - Ex: avoid “Hadoop abuse” [Smola & Narayanamurthy, VLDB 2010]



# Example One: Bayesian GMM

- Implemented relevant MCMC simulation on all four platforms
  - SimSQL, GraphLab, Spark, Giraph
- Philosophy: be true to the platform
  - Ex: avoid “Hadoop abuse” [Smola & Narayanamurthy, VLDB 2010]
- Ran on 10 dimensional data, 10 clusters, 10M points per machine
  - Full (non-diagonal) covariance matrix
  - Also on 100 dimensional data, 1M points per machine

# Example One: Bayesian GMM

GMM: Initial Implementations					
		10 dimensions			100 dimensions
	lines of code	5 machines	20 machines	100 machines	5 machines
SimSQL	197	27:55 (13:55)	28:55 (14:38)	35:54 (18:58)	1:51:12 (36:08)
GraphLab	661	Fail	Fail	Fail	Fail
Spark (Python)	236	26:04 (4:10)	37:34 (2:27)	38:09 (2:00)	47:40 (0:52)
Giraph	2131	25:21 (0:18)	30:26 (0:15)	Fail	Fail

- Some notes:

- Times are HH:MM:SS per iteration (time in parens is startup/initialization)
- Amount of data is kept constant per machine in all tests
- “Fail” means that even with much effort and tuning, it crashed

# Example One: Bayesian GMM

GMM: Initial Implementations					
		10 dimensions			100 dimensions
	lines of code	5 machines	20 machines	100 machines	5 machines
SimSQL	197	27:55 (13:55)	28:55 (14:38)	35:54 (18:58)	1:51:12 (36:08)
GraphLab	661	Fail	Fail	Fail	Fail
Spark (Python)	236	26:04 (4:10)	37:34 (2:27)	38:09 (2:00)	47:40 (0:52)
Giraph	2131	25:21 (0:18)	30:26 (0:15)	Fail	Fail

- Not much difference!

- But SimSQL was slower in 100 dims. Why?

- No native support for vectors/matrices at time tests were run
    - Forget array databases, *this* is an important problem!

# Example One: Bayesian GMM

GMM: Initial Implementations					
		10 dimensions			100 dimensions
	lines of code	5 machines	20 machines	100 machines	5 machines
SimSQL	197	27:55 (13:55)	28:55 (14:38)	35:54 (18:58)	1:51:12 (36:08)
GraphLab	661	Fail	Fail	Fail	Fail
Spark (Python)	236	26:04 (4:10)	37:34 (2:27)	38:09 (2:00)	47:40 (0:52)
Giraph	2131	25:21 (0:18)	30:26 (0:15)	Fail	Fail

- Spark is surprisingly slow

— Is Spark slower due to Python vs. Java?

GMM: Alternative Implementations					
		10 dimensions			100 dimensions
	lines of code	5 machines	20 machines	100 machines	5 machines
Spark (Java)	737	12:30 (2:01)	12:25 (2:03)	18:11 (2:26)	6:25:04 (36:08)

# Example One: Bayesian GMM

GMM: Initial Implementations					
		10 dimensions			100 dimensions
	lines of code	5 machines	20 machines	100 machines	5 machines
SimSQL	197	27:55 (13:55)	28:55 (14:38)	35:54 (18:58)	1:51:12 (36:08)
GraphLab	661	Fail	Fail	Fail	Fail
Spark (Python)	236	26:04 (4:10)	37:34 (2:27)	38:09 (2:00)	47:40 (0:52)
Giraph	2131	25:21 (0:18)	30:26 (0:15)	Fail	Fail

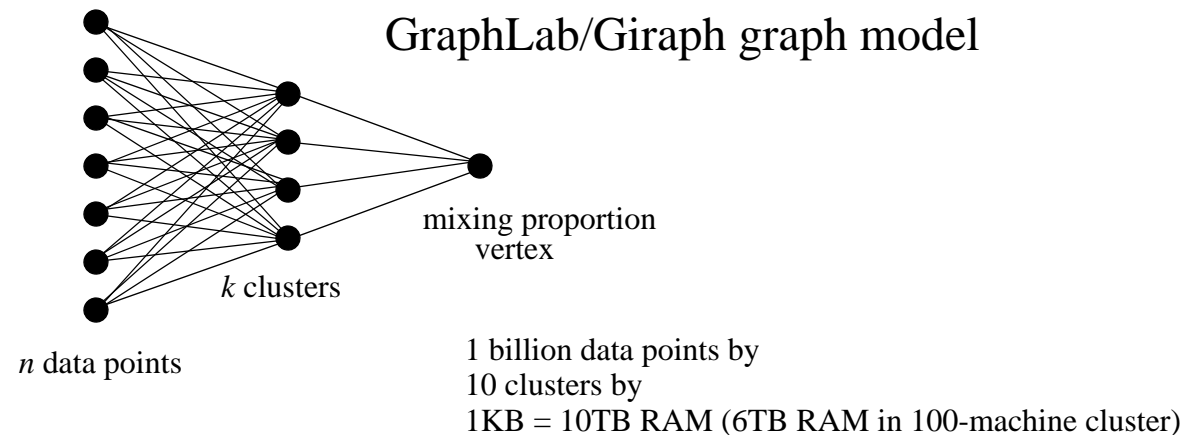
- What about GraphLab?
  - GraphLab failed every time. Why?

# Example One: Bayesian GMM

GMM: Initial Implementations					
		10 dimensions			100 dimensions
	lines of code	5 machines	20 machines	100 machines	5 machines
SimSQL	197	27:55 (13:55)	28:55 (14:38)	35:54 (18:58)	1:51:12 (36:08)
GraphLab	661	Fail	Fail	Fail	Fail
Spark (Python)	236	26:04 (4:10)	37:34 (2:27)	38:09 (2:00)	47:40 (0:52)
Giraph	2131	25:21 (0:18)	30:26 (0:15)	Fail	Fail

- What about GraphLab?

— GraphLab failed every time. Why?



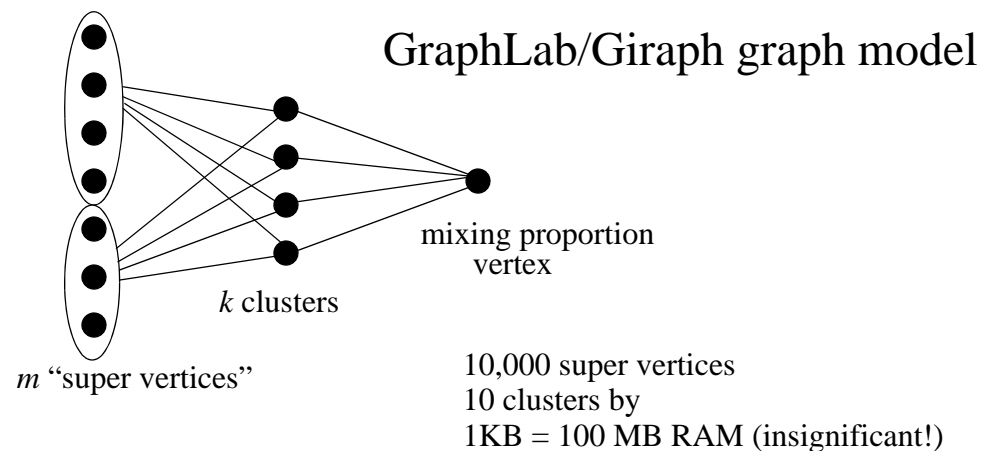
# Example One: Bayesian GMM

GMM: Initial Implementations					
		10 dimensions			100 dimensions
	lines of code	5 machines	20 machines	100 machines	5 machines
SimSQL	197	27:55 (13:55)	28:55 (14:38)	35:54 (18:58)	1:51:12 (36:08)
GraphLab	661	Fail	Fail	Fail	Fail
Spark (Python)	236	26:04 (4:10)	37:34 (2:27)	38:09 (2:00)	47:40 (0:52)
Giraph	2131	25:21 (0:18)	30:26 (0:15)	Fail	Fail

- What about GraphLab?

— GraphLab failed every time. Why?

To Fix...



# Example One: Bayesian GMM

GMM: Alternative Implementations					
		10 dimensions			100 dimensions
	lines of code	5 machines	20 machines	100 machines	5 machines
GraphLab (Super Vertex)	681	6:13 (1:13)	4:36 (2:47)	6:09 (1:21)*	33:32 (0:42)

- Super vertex results
  - GraphLab super vertex screams!



# Example One: Bayesian GMM

GMM: Alternative Implementations					
		10 dimensions			100 dimensions
	lines of code	5 machines	20 machines	100 machines	5 machines
GraphLab (Super Vertex)	681	6:13 (1:13)	4:36 (2:47)	6:09 (1:21)*	33:32 (0:42)

- Super vertex results

- GraphLab super vertex screams!
- But to be fair, others can benefit from super vertices as well...

GMM: Super Vertex Implementations				
	10 dimensions, 5 machines		100 dimensions, 5 machines	
	w/o super vertex	with super vertex	w/o super vertex	with super vertex
SimSQL	27:55 (13:55)	6:20 (12:33)	1:51:12 (36:08)	7:22 (14:07)
GraphLab	Fail	6:13 (1:13)	Fail	33:32 (0:42)
Spark (Python)	26:04 (4:10)	29:12 (4:01)	47:40 (0:52)	47:03 (2:17)
Giraph	25:21 (0:18)	13:48 (0:03)	Fail	6:17:32 (0:03)

# Example Two: Bayesian Lasso

- Experimental setup
  - 1K regressors (dense)
  - 100K points per machine

# Example Two: Bayesian Lasso

- Experimental setup
  - 1K regressors (dense)
  - 100K points per machine

- Results

Bayesian Lasso				
	lines of code	5 machines	20 machines	100 machines
SimSQL	100	7:09 (2:40:06)	8:04 (2:45:28)	12:24 (2:54:45)
GraphLab (Super Vertex)	572	0:36 (0:37)	0:26 (0:35)	0:31 (0:50)
Spark (Python)	168	0:55 (1:26:59)	0:59 (1:33:13)	1:12 (2:06:30)
Giraph	1871	Fail	Fail	Fail
Giraph (Super Vertex)	1953	0:58 (1:14)	1:03 (1:14)	2:08 (6:31)

# Example Two: Bayesian Lasso

- Experimental setup
  - 1K regressors (dense)
  - 100K points per machine

- Results

Bayesian Lasso				
	lines of code	5 machines	20 machines	100 machines
SimSQL	100	7:09 (2:40:06)	8:04 (2:45:28)	12:24 (2:54:45)
GraphLab (Super Vertex)	572	0:36 (0:37)	0:26 (0:35)	0:31 (0:50)
Spark (Python)	168	0:55 (1:26:59)	0:59 (1:33:13)	1:12 (2:06:30)
Giraph	1871	Fail	Fail	Fail
Giraph (Super Vertex)	1953	0:58 (1:14)	1:03 (1:14)	2:08 (6:31)

- Interesting points
  - SimSQL slow (again, lack of support for vectors/matrices is brutal here)...
  - But Spark is almost as slow for startup (computation of Gram matrix)
  - Check out GraphLab: super fast!

## Example Three: LDA

- Sort of a Bayesian variant on PCA (for dimensionality reduction)
- Experimental setup
  - Run over a document database, dictionary size of 10K words
  - 100 “topics” (components) were learned
  - Constant 2.5M documents per machine
- Note: didn’t do collapsed simulation, since hard to parallelize

## Example Three: LDA

- First we considered a “word based” implementation
  - Arguably the most natural
  - One vertex for each word in corpus in graph-based
  - Separate Multinomial call for each word in each doc in SimSQL/Spark
- And a “document based” implementation
  - One vertex for each document in graph-based
  - Update membership for all words at once in SimSQL/Spark (faster ‘cause you broadcast the model, do join with words in doc in user code)

# Example Three: LDA

- Results

LDA: Word-based and document-based implementations				
	Word-based, 5 machines		Document-based, 5 machines	
	lines of code	running time	lines of code	running time
SimSQL	126	16:34:39 (11:23:22)	129	4:52:06 (4:34:27)
Spark (Python)	NA	NA	188	≈15:45:00 (≈2:30:00)
Giraph	NA	NA	1358	22:22 (5:46)

- Interesting findings

- Only SimSQL can handle word-based imp, but really slow
- Only Giraph gives reasonable performance!
- Spark unable to join words-in-doc with topic-probs, hence an NA
- Giraph unable to load up word-based graph, hence an NA

# Example Three: LDA

- Results

LDA: Word-based and document-based implementations				
	Word-based, 5 machines		Document-based, 5 machines	
	lines of code	running time	lines of code	running time
SimSQL	126	16:34:39 (11:23:22)	129	4:52:06 (4:34:27)
Spark (Python)	NA	NA	188	≈15:45:00 (≈2:30:00)
Giraph	NA	NA	1358	22:22 (5:46)

- Interesting findings

- Only SimSQL can handle word-based imp, but really slow
- Only Giraph gives reasonable performance!
- Spark unable to join words-in-doc with topic-probs, hence an NA
- Giraph unable to load up word-based graph, hence an NA

- How about super vertex? (handle thousands of docs in a batch)



# Example Three: LDA

- Super vertex results

LDA: Super Vertex Implementations				
	lines of code	5 machines	20 machines	100 machines
Giraph	1406	18:49 (2:35)	20:02 (2:46)	Fail
GraphLab	517	39:27 (32:14)	Fail	Fail
Spark (Python)	220	≈3:56:00 (≈2:15:00)	≈3:57:00 (≈2:15:00)	Fail
SimSQL	117	1:00:17 (3:09)	1:06:59 (3:34)	1:13:58 (4:28)

- Interesting findings

— Only SimSQL can scale to 250M docs on 100 machines

# Example Three: LDA

- Super vertex results

LDA: Super Vertex Implementations				
	lines of code	5 machines	20 machines	100 machines
Giraph	1406	18:49 (2:35)	20:02 (2:46)	Fail
GraphLab	517	39:27 (32:14)	Fail	Fail
Spark (Python)	220	≈3:56:00 (≈2:15:00)	≈3:57:00 (≈2:15:00)	Fail
SimSQL	117	1:00:17 (3:09)	1:06:59 (3:34)	1:13:58 (4:28)

- Interesting findings

- Only SimSQL can scale to 250M docs on 100 machines

- Even super vertex can't help GraphLab here...

- 10K super vertices on 100 machines
- each broadcasts 100 different 10K vectors to each topic node
- 10K by 10K by 100 is 10 billion numbers...
- what if a machine gets 2 or three topic nodes?

# Example Three: LDA

- Super vertex results

LDA: Super Vertex Implementations				
	lines of code	5 machines	20 machines	100 machines
Giraph	1406	18:49 (2:35)	20:02 (2:46)	Fail
GraphLab	517	39:27 (32:14)	Fail	Fail
Spark (Python)	220	≈3:56:00 (≈2:15:00)	≈3:57:00 (≈2:15:00)	Fail
SimSQL	117	1:00:17 (3:09)	1:06:59 (3:34)	1:13:58 (4:28)

- Interesting findings

- Only SimSQL can scale to 250M docs on 100 machines
- Even super vertex can't help GraphLab here...
- Spark does quite poorly... might this be due to Python?

# Example Three: LDA

- Super vertex results

LDA: Super Vertex Implementations				
	lines of code	5 machines	20 machines	100 machines
Giraph	1406	18:49 (2:35)	20:02 (2:46)	Fail
GraphLab	517	39:27 (32:14)	Fail	Fail
Spark (Python)	220	≈3:56:00 (≈2:15:00)	≈3:57:00 (≈2:15:00)	Fail
SimSQL	117	1:00:17 (3:09)	1:06:59 (3:34)	1:13:58 (4:28)

- Interesting findings

- Only SimSQL can scale to 250M docs on 100 machines
- Even super vertex can't help GraphLab here...
- Spark does quite poorly... might this be due to Python?

LDA Spark Java Implementation			
lines of code	5 machines	20 machines	100 machines
377	9:47 (0:53)	19:36 (1:15)	Fail

# Summary of Findings

- Giraph can be made very fast
  - Mostly 'cause of distributed aggregation facilities
  - But it is still brittle, perhaps due to reliance on main memory

# Summary of Findings

- Giraph can be made very fast
  - Mostly 'cause of distributed aggregation facilities
  - But it is still brittle, perhaps due to reliance on main memory
- GraphLab codes are small and nice, especially considering C++
  - And it can be very fast
  - But lack of distributed agg is a killer... what does this even mean in asynch env?

# Summary of Findings

- Giraph can be made very fast
  - Mostly 'cause of distributed aggregation facilities
  - But it is still brittle, perhaps due to reliance on main memory
- GraphLab codes are small and nice, especially considering C++
  - And it can be very fast
  - But lack of distributed agg is a killer... what does this even mean in asynch env?
- Spark codes (Python) are startlingly beautiful. Wow!
  - But Spark was brittle, hard to tune, and often slow

# Summary of Findings

- Giraph can be made very fast
  - Mostly 'cause of distributed aggregation facilities
  - But it is still brittle, perhaps due to reliance on main memory
- GraphLab codes are small and nice, especially considering C++
  - And it can be very fast
  - But lack of distributed agg is a killer... what does this even mean in asynch env?
- Spark codes (Python) are startlingly beautiful. Wow!
  - But Spark was brittle, hard to tune, and often slow
- SimSQL codes fully declarative, and often competitive in speed
  - Only platform to run everything we threw at it
  - But lack of matrices and vectors really hurts



# Summary of Talk

- I've motivated a relational approach to large-scale ML
  - All about data independence!
  - Same code works for any data set, compute platform
  - Just drop in a new physical optimizer and runtime, keep application stack

# Summary of Talk

- I've motivated a relational approach to large-scale ML
  - All about data independence!
  - Same code works for any data set, compute platform
  - Just drop in a new physical optimizer and runtime, keep application stack
- I've briefly described SimSQL, our realization of the approach

# Summary of Talk

- I've motivated a relational approach to large-scale ML
  - All about data independence!
  - Same code works for any data set, compute platform
  - Just drop in a new physical optimizer and runtime, keep application stack
- I've briefly described SimSQL, our realization of the approach
- And I've given experimental evidence the approach is practical
  - Our Hadoop targeted optimizer and runtime competes well
  - And its the only platform to handle everything we threw at it

# That's It. Questions?

- Download SimSQL today
  - <http://cmj4.web.rice.edu/SimSQL/SimSQL.html>
- This presentation at
  - <http://cmj4.web.rice.edu/SimSQLNew.pdf>