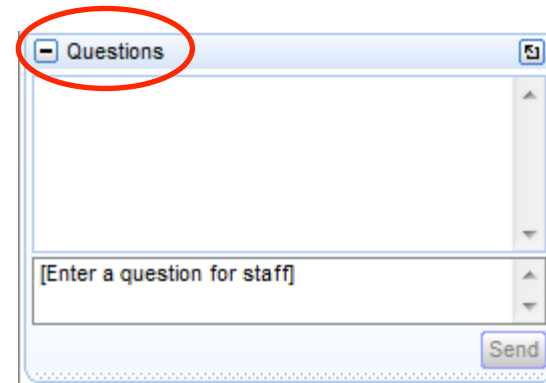





SparkR

The GoToWebinar Control Panel

1. The **orange** button
2. Audio Type
3. Close apps
4. Enlarge my screen
5. Headphones
6. Questions Pane



Today's Agenda

- Lecture
 - slides and/or video will be made available within one week
- Live Demonstration
- Q & A

"Intro to Spark" Training Course

To attend a **hands-on** Spark training course which runs every Saturday, please visit:

liondatasystems.com/courses

Thank You!

- This event has attracted nearly 900 registrants from various parts of the world.
- Thank you everyone for your support!

Today's Speaker

- Shivaram Venkataraman
- Co-Author of SparkR
- PhD Student @ UC Berkeley
- Former Google Engineer

Introduction to SparkR

Shivaram Venkataraman



Big Data & R

DataFrames
Visualization
Libraries



+

Data

Background



Engine for large-scale data processing

Fast, Easy to Use

Runs Everywhere - EC2, YARN, Mesos

SparkR



Interactive Shell

Batch Scripts

Outline

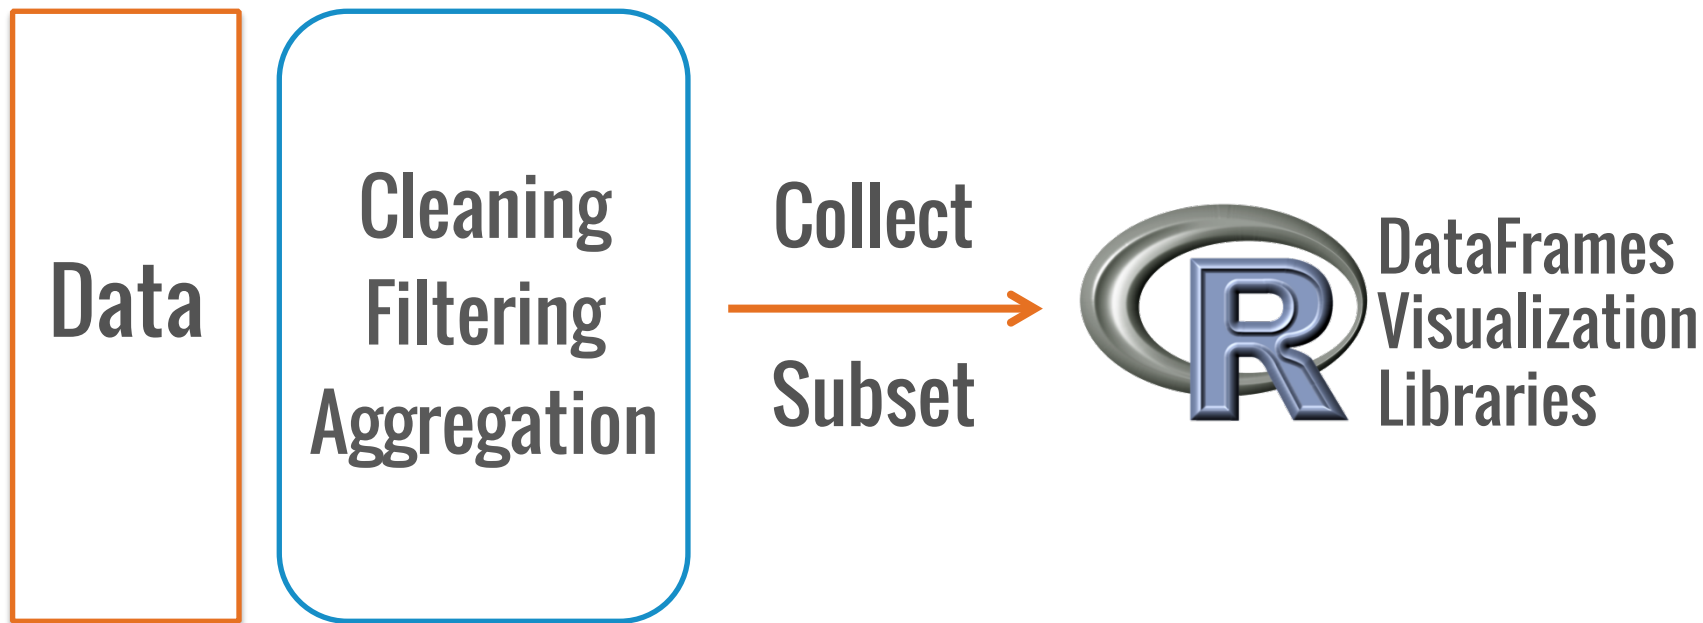
SparkR DataFrames

Architecture

Demo

SparkR Roadmap

Big Data Processing + R



SparkR DataFrames

High-level API for data manipulation

Read in CSV, JSON, JDBC etc.

dplyr-like syntax

Example

```
{"name": "Michael", "age": 29}  
  {"name": "Andy", "age": 30}  
  {"name": "Justin", "age": 19}  
    {"name": "Bob", "age": 22}  
  {"name": "Chris", "age": 28}  
  {"name": "Garth", "age": 36}  
  {"name": "Tasha", "age": 24}  
    {"name": "Mac", "age": 30}  
  {"name": "Neil", "age": 32}
```

Example

```
people <- read.df(  
  "hdfs://people.json",  
  "json")
```

Read input from HDFS

```
avgAge <- select(  
  df,  
  avg(df$age))
```

Collect to data.frame

```
collect(avgAge)
```

DataFrame API

Filtering Data

- select, `\$`, where, filter

Aggregating Data

- groupBy, summarize, arrange


Input/Output

- read.df, write.df, sql

SparkR DataFrames

Query Planning

SQL Optimizations



Architecture


Architecture

Local


Worker

Worker


Architecture



Architecture



Architecture



Demo

Demo Overview

Launching SparkR

- On your laptop
- On EC2

SparkR DataFrames

Running SparkR Locally

Download from <http://spark.apache.org/> (>1.4.0)

`./bin/sparkR` or RStudio

Useful for learning SparkR, demonstrations

SparkR on EC2

Launch cluster with Spark's EC2 scripts

```
./spark-ec2 -s 2 -t r3.xlarge -i <pem> -k <key> sparkr
```

Follow r-bloggers.com/spark-1-4-for-rstudio/

Thanks Vincent Warmerdam !

SparkR Future


Big Data & R

Big Data
Small Learning


Partition
Aggregate

Large Scale
Machine Learning


Big Data Processing + R




2(a). Partition Aggregate



2(b). Partition Aggregate



3. Large Scale Machine Learning



Big Data & R

Big Data

Small Learning

Partition

Aggregate

Large Scale

Machine Learning

SparkR:
Unified approach

Partition Aggregate

Upcoming feature:

Simple, parallel API for SparkR

Ex: Parameter tuning, Model Averaging

Integrated with DataFrames

Use existing R packages

Large Scale Machine Learning

Integration with MLlib

Support for GLM, KMeans etc.

```
model <- glm(  
  a ~ b + c,  
  data = df)
```

Large Scale Machine Learning

Key Features

DataFrame inputs

R-like formulas

Model statistics

```
model <- glm(  
  a ~ b + c,  
  data = df)
```

```
summary(model)
```

Developer Community

>20 contributors including
AMPLab, Databricks, Alteryx, Intel

New contributions welcome !

Big data processing from R

SparkR

DataFrames in Spark 1.4

Future: Large Scale ML & more

Local Demo

```
Sys.setenv(SPARK_HOME="/Users/shivaram/spark-1.4.1")
.libPaths(c(file.path(Sys.getenv("SPARK_HOME"), "R", "lib"), .libPaths()))
library(SparkR)
sc <- sparkR.init(master="local")
sqlContext <- sparkRSQL.init(sc)

df <- createDataFrame(sqlContext, faithful)
# Select one column
head(select(df, df$eruptions))

# Filter out rows
head(filter(df, df$waiting < 50))
```

EC2 Demo

```
# If you are using Spark 1.4, then launch SparkR with the command
# ./bin/sparkR --packages com.databricks:spark-csv_2.10:1.0.3
# as the `sparkPackages=` flag was only added in Spark 1.4.1.

## This will work in Spark 1.4.1.
sc <- sparkR.init(spark_link, sparkPackages = "com.databricks:spark-
csv_2.10:1.0.3")
sqlContext <- sparkRSQL.init(sc)

flights <- read.df(sqlContext, "s3n://sparkr-
data/nycflights13.csv", "com.databricks.spark.csv", header="true")

# Print the first few rows
head(flights)

# Run a query to print the top 5 most frequent destinations from JFK
jfk_flights <- filter(flights, flights$origin == "JFK")

# Group the flights by destination and aggregate by the number of flights
dest_flights <- agg(group_by(jfk_flights, jfk_flights$dest), count = n(jfk_flights$dest))

# Now sort by the `count` column and print the first few rows
head(arrange(dest_flights, desc(dest_flights$count)))

## dest count
##1 LAX 11262
##2 SFO 8204
##3 BOS 5898
```

```
# Running SQL Queries
registerTempTable(flights, "flightsTable")
delayDF <- sql(sqlContext, "SELECT dest, arr_delay FROM flightsTable")

# Creating new Columns, Deleting columns
flights$air_time_hr <- flights$air_time / 60
flights$air_time_hr <- NULL

# Combine the whole query into two lines using magrittr
library(magrittr)
dest_flights <- filter(flights, flights$origin == "JFK") %>%
  group_by(flights$dest) %>%
  summarize(count = n(flights$dest))

top_dests <- head(arrange(dest_flights, desc(dest_flights$count)))
barplot(top_dests$count, names.arg = top_dests$dest)
```