

XADATA

Sun
ORACLE

ORACLE®

Oracle In-Database Hadoop: When MapReduce Meets RDBMS

Oracle In-Database Hadoop

When MapReduce Meets RDBMS

Xueyuan Su

xueyuan.su@yale.edu

Yale University

Garret Swart

garret.swart@oracle.com

Oracle Corporation

*This work describes a prototype built on top of Oracle DB, which is not a feature of Oracle products.

Oracle In-Database Hadoop

Agenda

- Big Data, Hadoop MapReduce & SQL
- Oracle In-Database Hadoop
 - Architecture & Design
 - Implementation
- Summary



Why Big Data and Hadoop?

Big Data

Big Data Is A Phenomenon

- More commercial and social interactions are mediated by computing technology.
- The cost of storage continues to decrease.
- Cloud computing becomes popular.
- E.g., web logs, business records, social networks, search indexing, photo and video archives, and Internet documents.


Big Data

Why Is Processing Big Data Challenging?

- Low information density, often unstructured.
- Overwhelming scale.
- Varied formats and huge computational overhead.
- Tar sands - potentially valuable information resources but requiring huge inputs in energy and technology to exploit.

MapReduce Paradigm

You All Know This Stuff!



Map:
 $\langle K1, V1 \rangle \rightarrow$
 $\{ \langle K2, V2 \rangle, \dots \}$

Shuffle:
 $\{ \langle K2, V2 \rangle, \dots \} \rightarrow$
 $\{ \langle K2, \{ V2, \dots, V2 \} \rangle, \dots \}$

Reduce:
 $\langle K2, \{ V2, \dots, V2 \} \rangle$
 $\rightarrow \{ \langle K3, V3 \rangle, \dots \}$

Hadoop Implementation

API

- Model is similar to "Standard In / Standard Out": Mappers and Reducers rely on the environment to get the input from and to put output in the right place, but do not access File Systems or Data Stores directly.
- Some combination of the Configuration and the Driver program are used to specify the environment.
- This makes Mappers and Reducers more reusable -- a small application surface area that we have to reimplement to allow Mappers and Reducers to run over Tables of SQL types.



Why In-DB Hadoop?

Hadoop & SQL

Why In-DB Hadoop?

For customers who have already invested in a database infrastructure:

- Avoid additional investment into a Hadoop cluster.
- Reduce training and deployment time.
- Mix SQL and MapReduce processing for flexibility and efficiency.

Hadoop & SQL

Previous Efforts

- Oracle user-defined pipelined table functions and aggregation objects, SQL-MapReduce, Pig Latin, Hive, Tenzing, HadoopDB (Hadapt), etc.
- Limitations of previous efforts:
 - **Source compatibility**: Need to rewrite Hadoop programs in a different language;
 - **Dependency on Hadoop infrastructure**: Rely on Hadoop clusters for query execution.



Can We Do Better?


Oracle In-Database Hadoop

Design Goals

- **Source-compatibility**: Accept Hadoop programs written in Java.
- **Direct data access**: No need to move Oracle RDBMS resident data to a separate infrastructure.
- **Minimal dependency**: No Hadoop clusters.
- **Seamless integration**: MapReduce functionality embedded in Oracle SQL.
- **Java interface**: Allow Hadoop users to execute their applications in the traditional Java way.

Oracle In-Database Hadoop

Architecture



- Data storage: Table, external table, object view.
- PQ engine: Data partitioning & task scheduling.
- Oracle JVM: Task execution.
- TableReader, TableWriter: Data type mapping.

Hadoop SQL Queries

One Example

- Mix Hadoop with SQL.
- MapReduce steps as pipelined table functions.
- A link to retrieve configuration parameters.

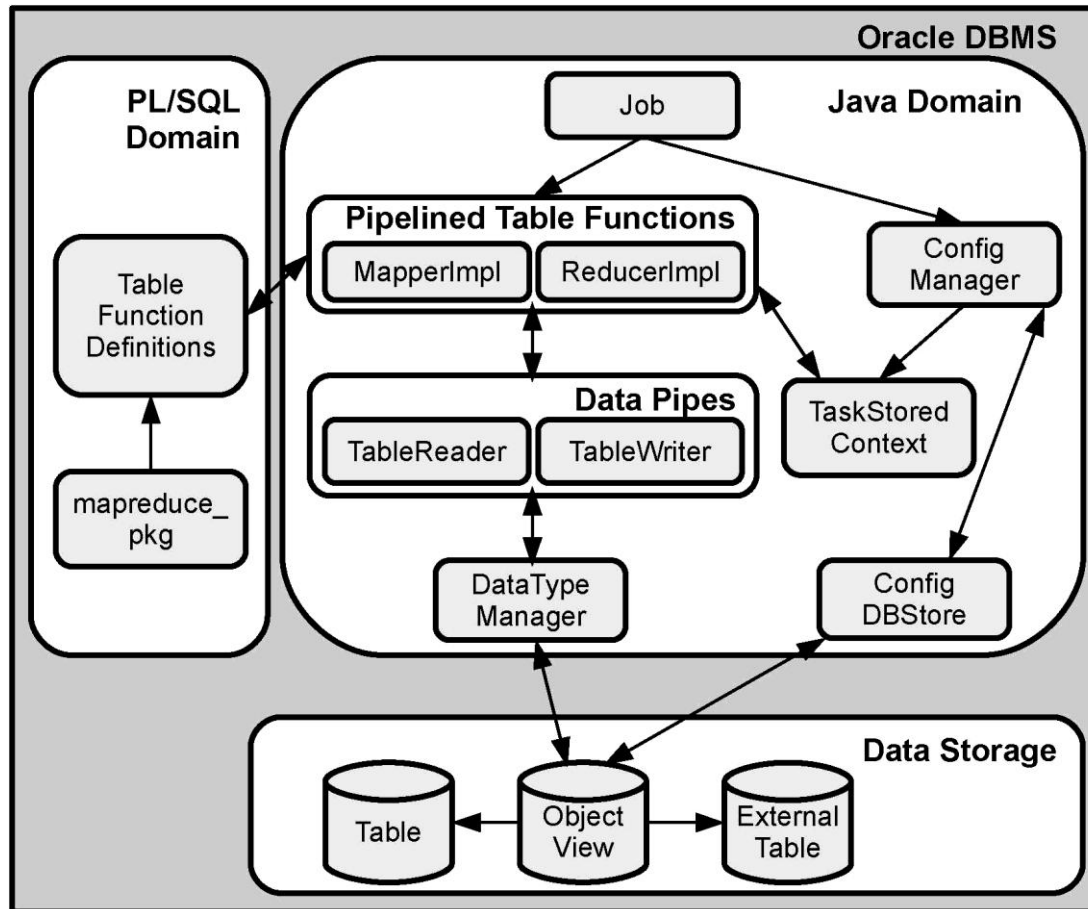
```
INSERT INTO OutTable
SELECT * FROM TABLE
  (Word_Count_Reduce(:ConfKey,
    CURSOR(SELECT * FROM TABLE
      (Word_Count_Map(:ConfKey,
        CURSOR(SELECT * FROM InTable))))))
```



What's Under the Hood?

Implementation


Software Components



- Oracle 11.2.
- Oracle JVM compatible with JDK 1.5.
- Documented public APIs.
- Hadoop 0.20.2.

Hadoop Job Configuration


Setup, Storage and Retrieval



- Context-switching between Java and SQL environments -- need a way to pass Hadoop configurations.
- Users setup configuration parameters in job drivers as in Hadoop.
- The framework stores and retrieves Hadoop configuration objects from a row in the database using our own specialization of the Hadoop ConfigManager.

Hadoop Job Configuration

Additional DB Parameters



```
Configuration conf = new Configuration();  
Job job = new Job(conf, "word count");
```

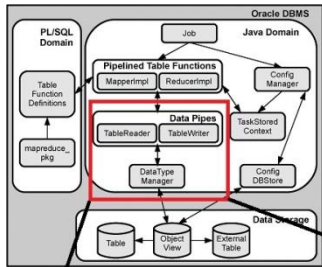
```
/* Set mapper and reducer classes. */  
job.setMapperClass(TokenizerMapper.class);  
job.setReducerClass(IntSumReducer.class);
```

```
/* Set input and output Hadoop types. */  
job.setInputKeyClass(Object.class);  
job.setInputValueClass(Text.class);  
job.setOutputKeyClass(Text.class);  
job.setOutputValueClass(IntWritable.class);
```


```
/* Set the output SQL types. */  
job.setOutputKeyDBType("VARCHAR2(100)");  
job.setOutputValueDBType("NUMBER");
```

SQL Data Types

Input Generics



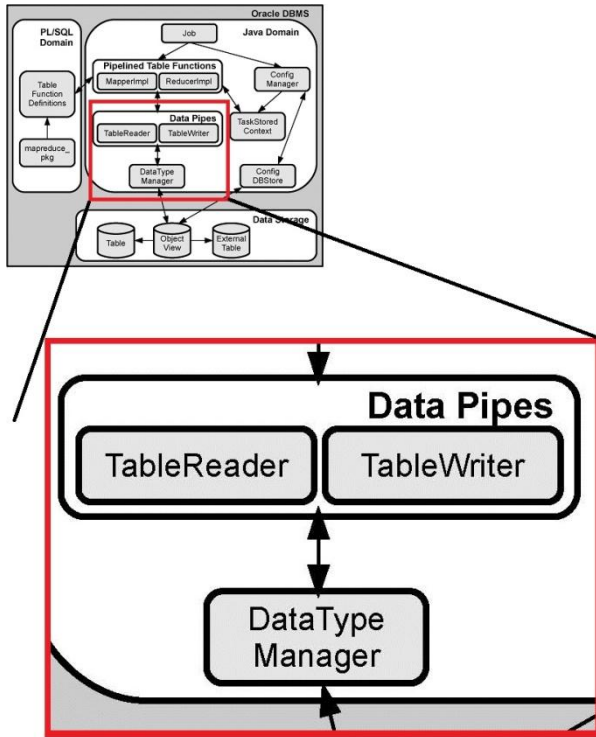
- Make use of the **ANYDATA** type defined in PL/SQL.
- Users can store data in any supported formats.



```
TYPE MapReduceInputType IS RECORD (  
    KEYIN SYS.ANYDATA,  
    VALUEIN SYS.ANYDATA  
);  
TYPE inputcur IS REF CURSOR  
    RETURN MapReduceInputType;
```

SQL Data Types

Type-Specific Output




- Users configure output key/value types as in Hadoop.
- The framework automatically generates corresponding object and table types for SQL.
- Name convention:

```
MAPOUT_<MAPOUTKEYTYPE>_<MAPOUTVALUETYPE>  
OUT_<REDOUTKEYTYPE>_<REDOUTVALUETYPE>
```

SQL Data Types


Mapping Between SQL and Hadoop Types



- Convert data between Hadoop Writables and SQL types.
- Automatically invoked in TableReader & TableWriter.
- Implicit mappings between SQL scalar types, VARRAY and corresponding Hadoop types.
- Composite types are handled by DBA creating SQL Object views that match the structure of the Hadoop classes.

Pipelined Table Functions


General Design



- Take a stream of rows as input and return a stream of rows as output.
- Implemented in Java. Invoked from the SQL domain and executed in the Java domain.
- Data pipelining, the avoidance of barriers and intermediate data materialization.

Pipelined Table Functions

Declaration



CREATE OR REPLACE FUNCTION

Reduce_<REDOUTKEYTYPE>_<REDOUTVALUETYPE>

(jobKey NUMBER, p mapreduce_pkg.inputcur)

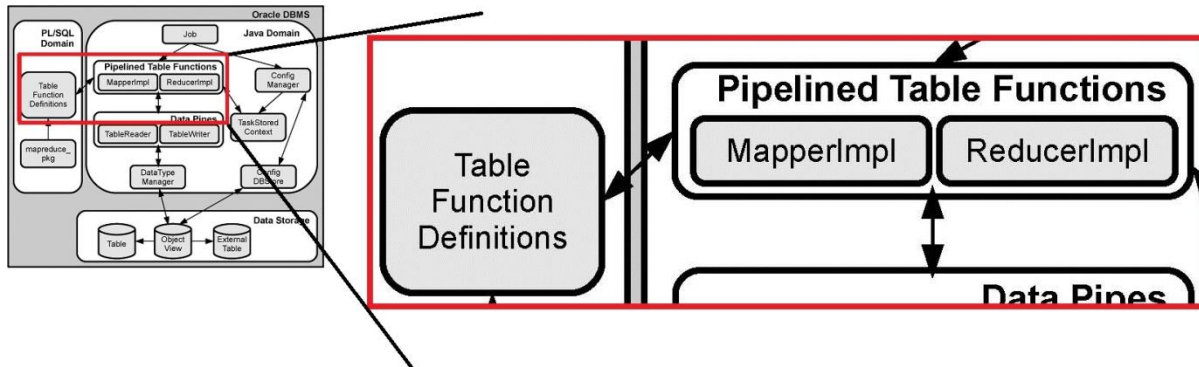
RETURN (OUTSET_<REDOUTKEYTYPE>_<REDOUTVALUETYPE>)

PIPELINED PARALLEL_ENABLE

(PARTITION p BY hash(KEYIN)) CLUSTER p BY (KEYIN)

USING ReducerImpl;

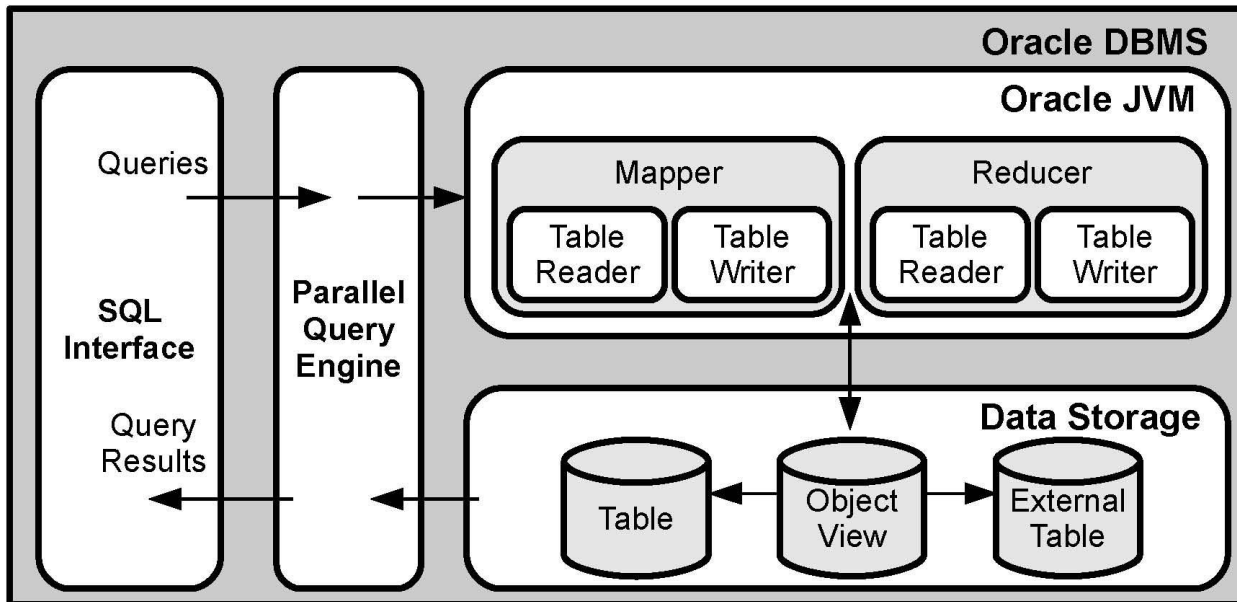
Pipelined Table Functions Interface



- **ODCITableStart:**
 - Instantiate user provided Hadoop Mapper/Reducer classes.
 - Gain access rights via Java reflection.
 - Accept configuration parameters and set up environment.
- **ODCITableFetch:**
 - Execute Mapper/Reducer program on input data.
- **ODCITableClose:**
 - Clean up and return.

Hadoop SQL Queries Revisited

Two Interfaces: SQL & Java



```
SELECT * FROM TABLE
(Reduce_VARCHAR2_NUMBER(:ConfKey,
  CURSOR(SELECT * FROM TABLE
(Map_VARCHAR2_NUMBER(:ConfKey,
  CURSOR(SELECT * from InTable))))))
```


```
public class WordCount {
  public static void main() throws Exception {
    /* Setup the parameters and run the job */
    .....
    job.init();
    job.run();
  }
}
```



What Have We Done?

Oracle In-Database Hadoop

Summary



A prototype of Oracle In-DB Hadoop:

- Source compatibility with Hadoop.
- Access to Oracle RDBMS resident data without the need to move the data to a separate infrastructure.
- Minimal dependency on the Apache Hadoop infrastructure.
- Greater efficiency in execution due to data pipelining, the avoidance of barriers and intermediate data materialization in the Oracle implementation.
- Seamless integration of MapReduce functionality with Oracle SQL, allowing mixing MapReduce steps with sophisticated SQL queries.



Where Can We Go From Here?

Potential Extensions

- SQL extensions: Allow table functions to compute their output type based on configuration parameters and the types of the input parameters.
- JDBC extensions: Map between Hadoop Writable types and SQL types, not PL/SQL types. PL/SQL is the Oracle specific stored procedure language.
- Support more Hadoop classes running inside the DB: InputFormat and OutputFormat classes would allow SQL to access any Hadoop data source or sink.

