



# Apache Tajo : A Big Data Warehouse System on Hadoop

김형준

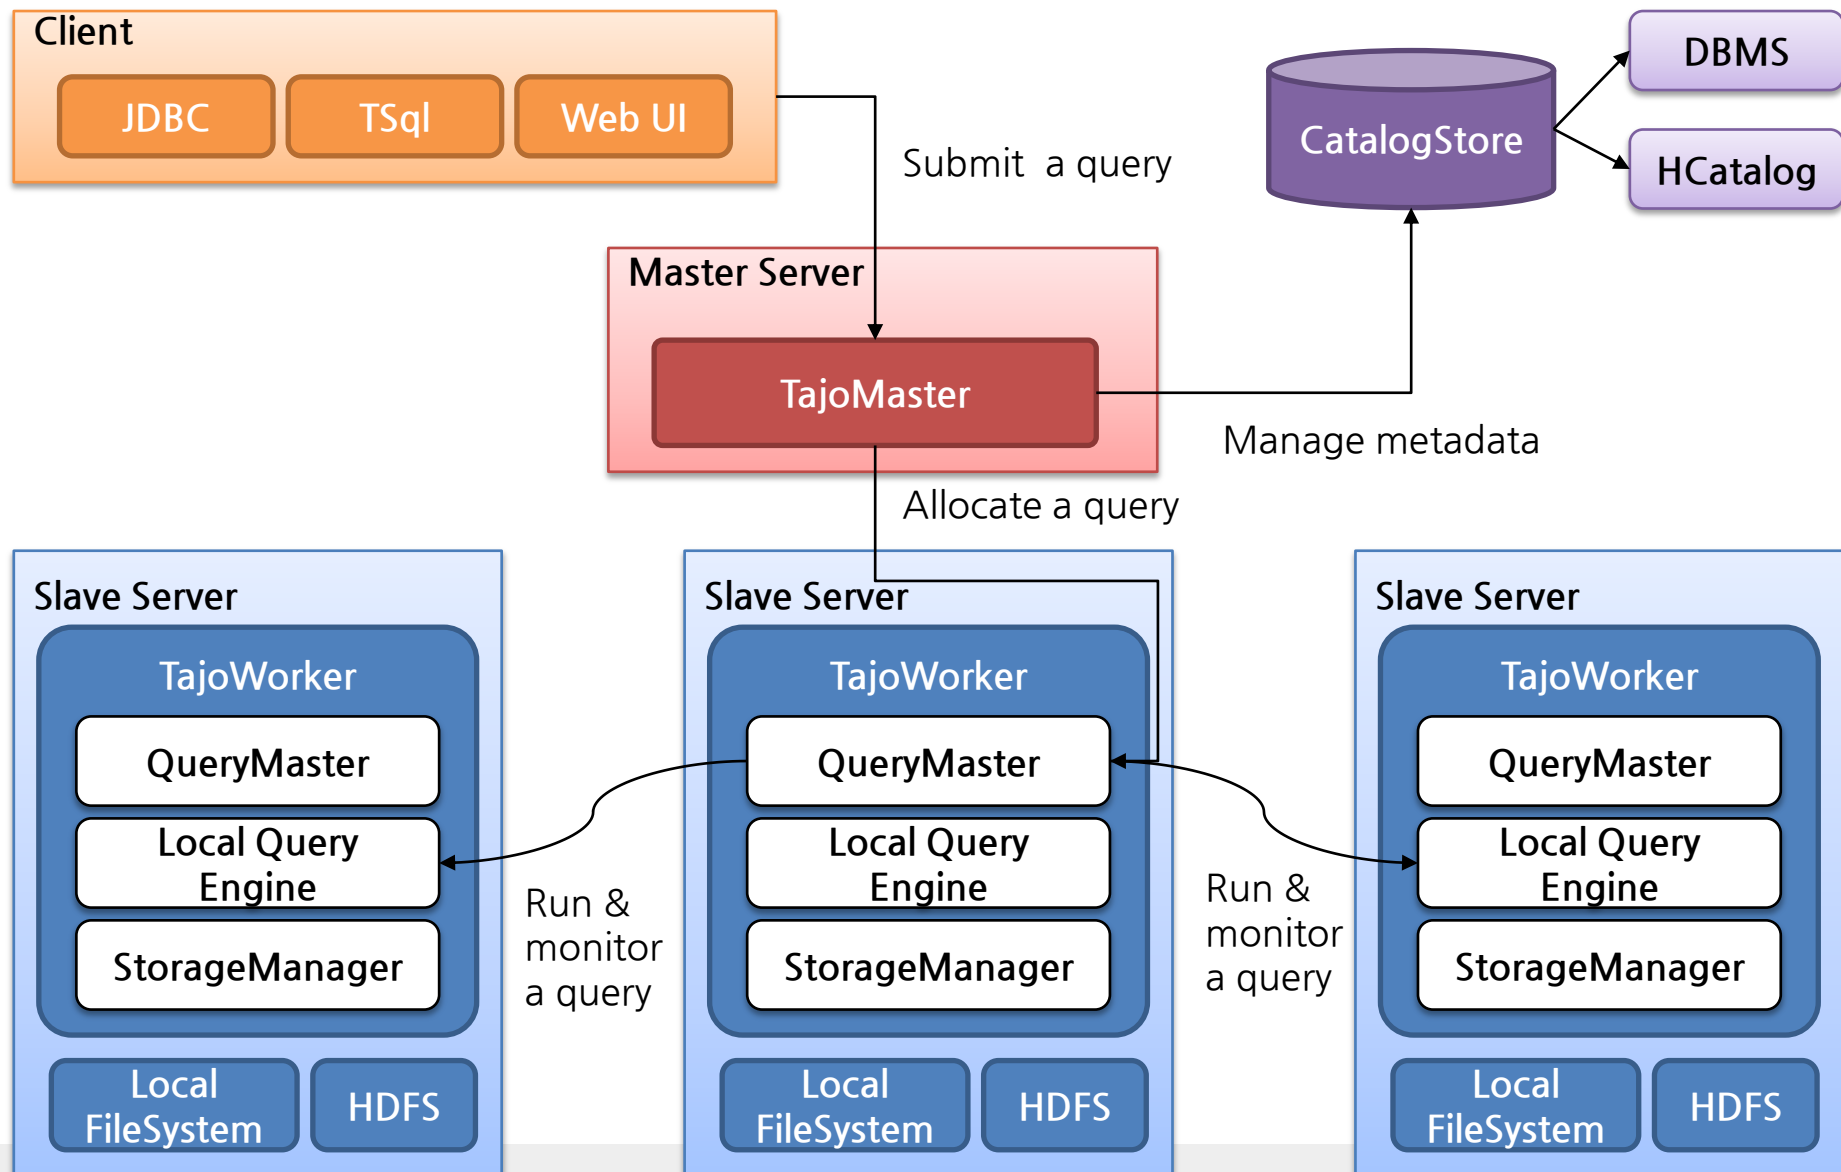
[babokim@gmail.com](mailto:babokim@gmail.com)

# Apache Tajo Overview

- A big data warehouse system on Hadoop
- Apache Top-level project since March 2014
- Supports SQL
- Features
  - Powerful (Not MapReduce)
  - Advanced query optimization techniques
  - Long running queries
  - Interactive analysis queries : from 100 milliseconds
- Recent 0.9.0 release



# Tajo Architecture



# Mature SQL Feature Set

- **Fully distributed query executions**
  - Inner join, and left/right/full outer join
  - Groupby, sort, multiple distinct aggregation
  - window function
- **SQL data types**
  - CHAR, BOOL, INT, DOUBLE, TEXT, DATE, Etc
- **Various file formats**
  - Text file (CSV), SequenceFile, RCFile, Parquet, Avro
- **SQL Standards**
  - Non standard features : PostgreSQL and Oracle

# Group by, Order by

- **Group by**
  - Multi-level distributed Group by
    - `tajo.dist-query.groupby.multi-level-aggr`
  - Multiple Count Distinct
    - `Select count(distinct col1), sum(distinct col2), ...`
  - Hash Aggregation or Sort Aggregation
- **Order by**
  - Fully distributed order by
  - Range partition

# Join

- **Join**
  - NATURAL, INNER, OUTER (LEFT, RIGHT, FULL)
  - SEMI, ANTI Join (planned for v0.9)
- **Join Predicates**
  - WHERE and ON predicates
  - de-factor standard outer join behavior with both predicates

```
SELECT * FROM t1 LEFT JOIN t2  
ON t1.num = t2.num WHERE t2.value = 'xxx';
```

```
SELECT * FROM t1 LEFT JOIN t2  
WHERE t1.num = t2.num and t2.value = 'xxx';
```

# Table Partitions

- **Column Value Partition**
  - Hive Compatible Partition

```
CREATE TABLE T1 (C1 INT, C2 TEXT)
  using PARQUET
  WITH ('parquet.compression' = 'SNAPPY')
  PARTITION BY COLUMN (C3 INT, C4 TEXT);
```

- **Range Partition (planned for 1.0)**
  - Table will be partitioned by disjoint ranges.
  - Will remove the partition granularity problem of Hive Partition

# Window Function

- **OVER clause**
  - row\_number() and rank()
  - Aggregation function support
  - PARTITION and ORDER BY clause

```
SELECT depname, empno, salary, enroll_date FROM ( SELECT
  depname, empno, salary, enroll_date,
  rank() OVER (PARTITION BY depname
    ORDER BY salary DESC, empno) AS pos
  FROM empsalary
) AS ss WHERE pos < 3;
```



## Currently Not Supported

- **IN/Exists SubQuery(2014, 4Q)**
  - Select \* from t1 where col1 in (select col1...)
- **Scalar SubQuery(2015, 1Q)**
  - Select col1, (select col2 from ...) from ...
- **Some analytic function(2015, 1Q or 요청시)**
  - ROLLUP, CUBE, some Window
- **Create/Drop Function(2014, 4Q)**
  - 실행 중에 UDF를 추가하는 기능 미제공
- **Alter meta property(2014, 4Q)**
  - Partition 정보 등


# Comparison with other platform

\* 버전 업 등에 따라 내용은 틀릴 수 있음

Function	Tajo	Hive	Impala	Spark
Computing	자체	MapReduce or Tez	자체	자체
Resource Management	자체 or YARN	YARN	자체	자체 or YARN
Scheduler	FIFO, Fair	FIFO, Fair, Capacity	FIFO, Fair	FIFO, Fair
Storage	HDFS, S3, HBase	HDFS, HBase, S3	HDFS, HBase	자체 RDD (HDSF 등)
File Format	CSV, RC, Parquet, Avro 등	CSV, RC, ORC, Parquet, Avro 등	CSV, RC, Parquet, Avro 등	CSV, RC, Parquet, Avro 등
Data Model	Relational	Nested Data Model	Relational	Nested Data Model
Query	ANSI-SQL	HiveQL	HiveQL	HiveQL
구현 언어	Java	Java	C++	Scala
Client	Java API, JDBC, CLI	CLI, JDBC, ODBC, Thrift Server API	CLI, JDBC, ODBC	Shark JDBC/ODBC, Scala, Java, Python API
Query Latency	Long run, Interactive	Long run, (Interactive-Tez)	Interactive	Interactive
컴퓨팅 특징	데이터는 Disk, 중간 데이터는 Memory/Disk 모두 사용	데이터는 Disk, 중간 데이터는 Memory/Disk 모두 사용	중간 데이터가 In-Memory (최근 On-Disk 지원)	분석 대상 데이터가 In-Memory에 로딩
License	Apache	Apache	Apache	Apache
Main Sponsor	Gruter	Hortonworks	Cloudera	Databricks

# Performance(1)

- 테스트 장비: 1 master + 6 workers  
CPU: Xeon 2.5GHz, E5, 24 Core, Memory: 64GB, Network:10Gb  
Disk: 3TB \* 6 SATA/HDD (7200 RPM)
- 테스트 데이터: TPC-H Scale 100 or 1000
- Version  
Hadoop: cdh-4.3.0, Hive: 0.10.0-cdh4.3.0, Impala: impalad\_version\_1.1.1\_RELEASE  
Tajo: 0.2-SNAPSHOT





## Performance(2)

- 실제 사용 데이터 및 질의
  - 데이터
    - 1.7TB (4.1B rows, Q1)
    - 8 or less GB (results of Q1, rest of Qs)
  - Query
    - Q1: scan using about 20 text pattern matching filters
    - Q2: 7 unions with joins
    - Q3: join
    - Q4: group by and order by
    - Q5: 30 text pattern matching filters with OR conditions, group by, having, and order by 4


<http://www.slideshare.net/gruter/tajo-case-studybayareahugmeetupevent20131105>

# Q1 – filter scan



Q1: scan using about 20 text pattern matching filters


## Q2 – unions, joins



# Q3 – join



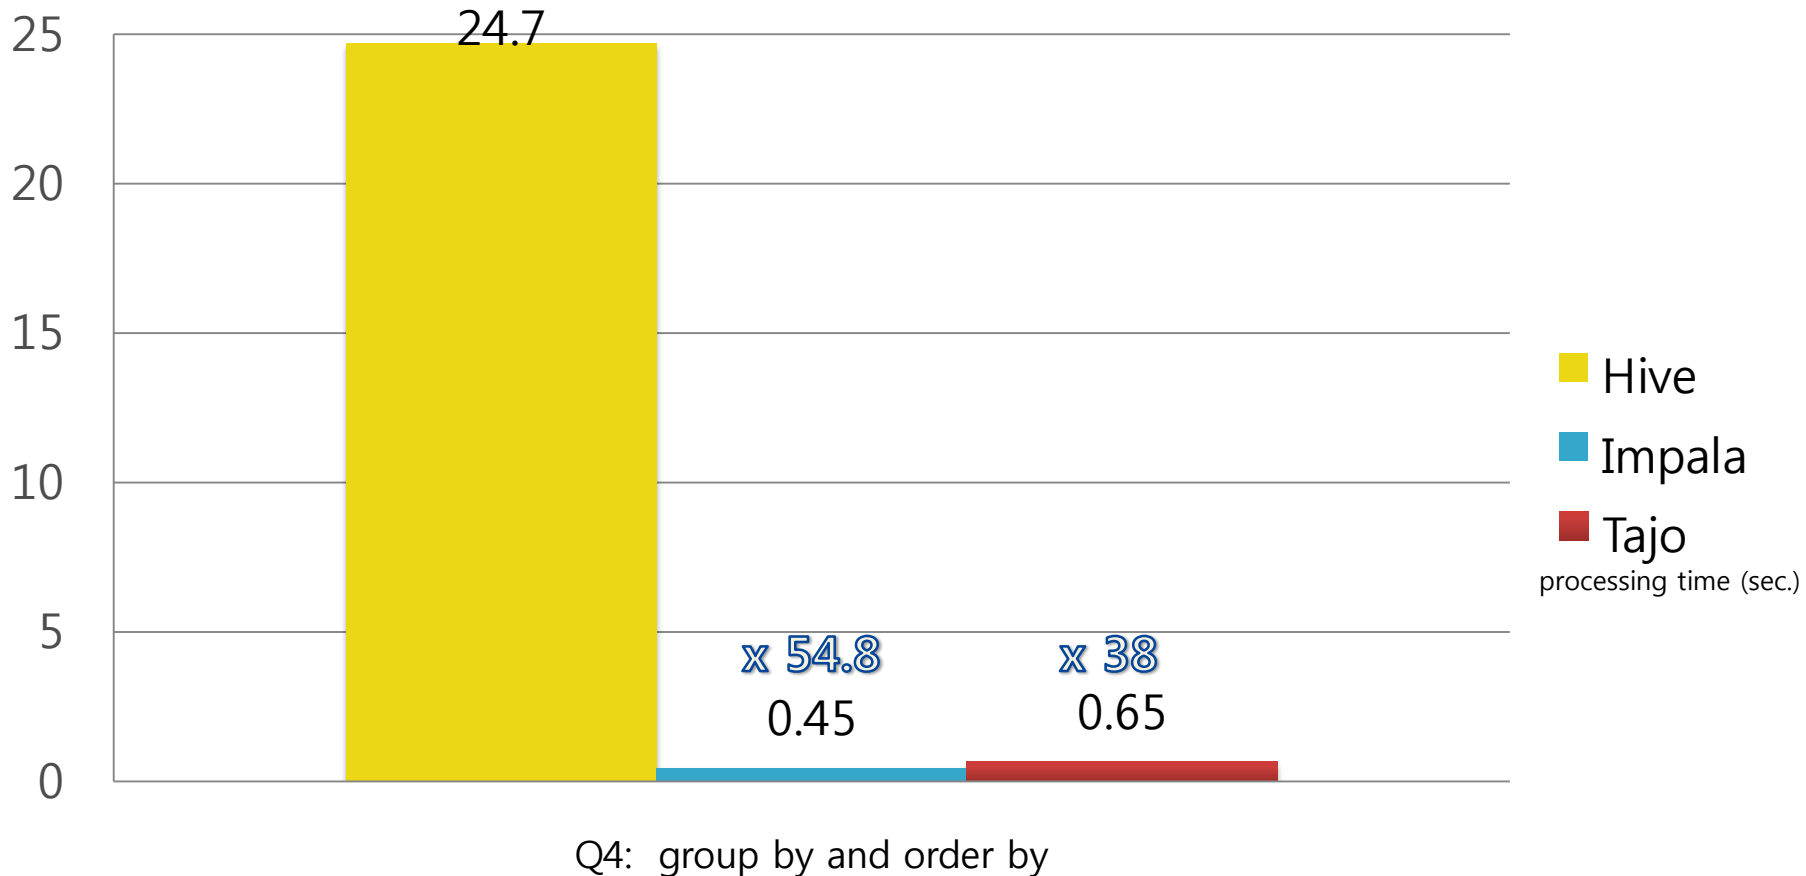
cloudera<sup>®</sup>  
IMPALA



## Q4 – group by and sort




cloudera<sup>®</sup>  
IMPALA





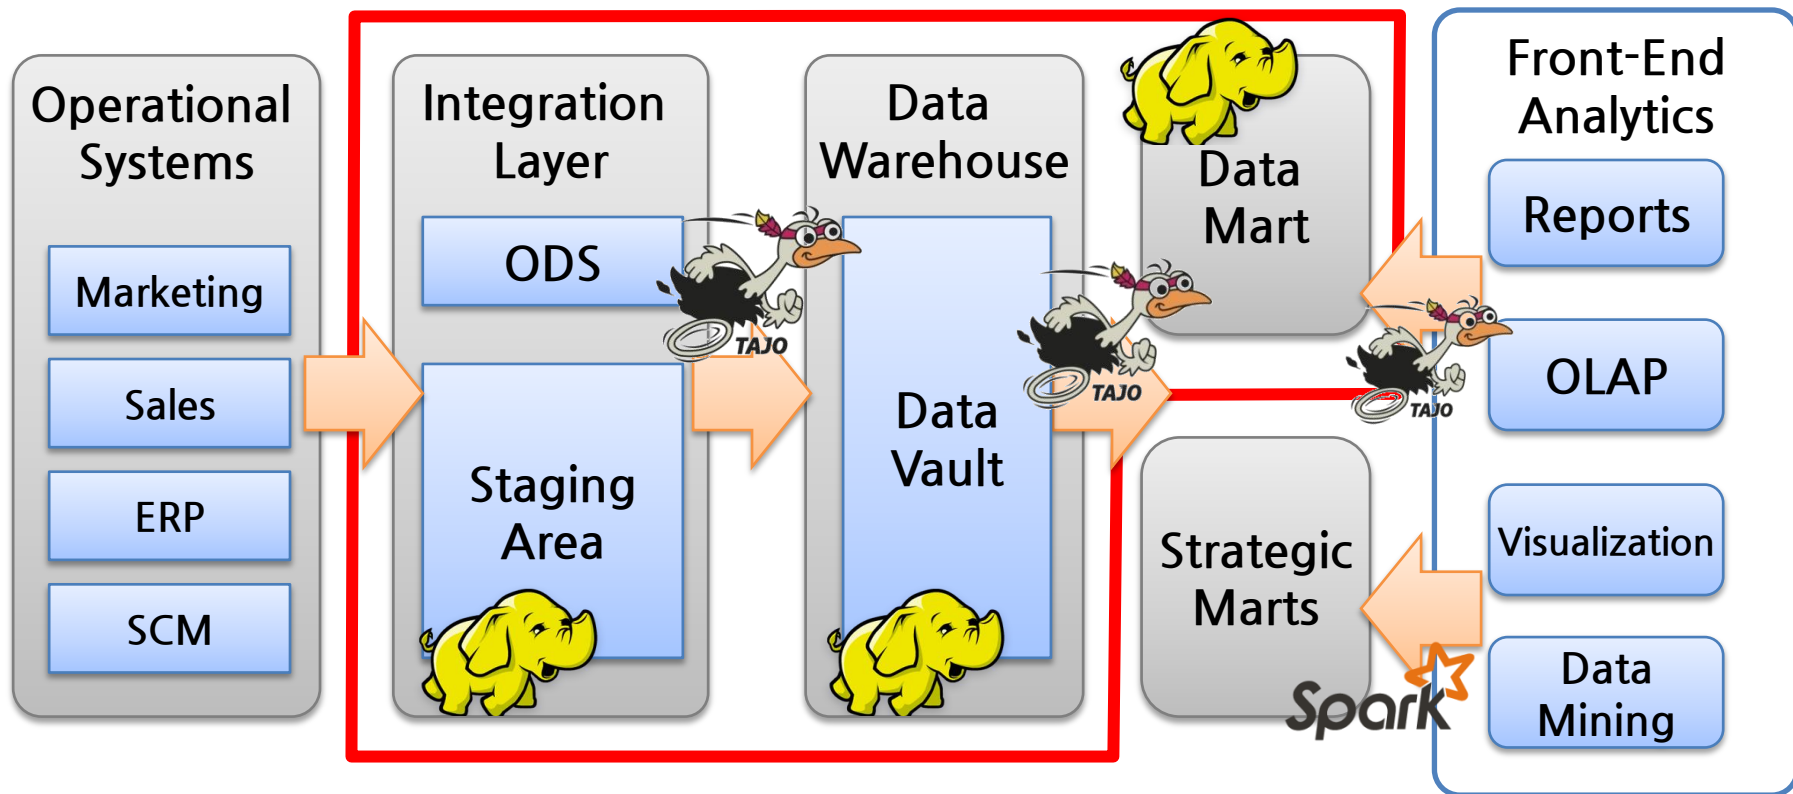
# Q5 – filters, group by, having and sort



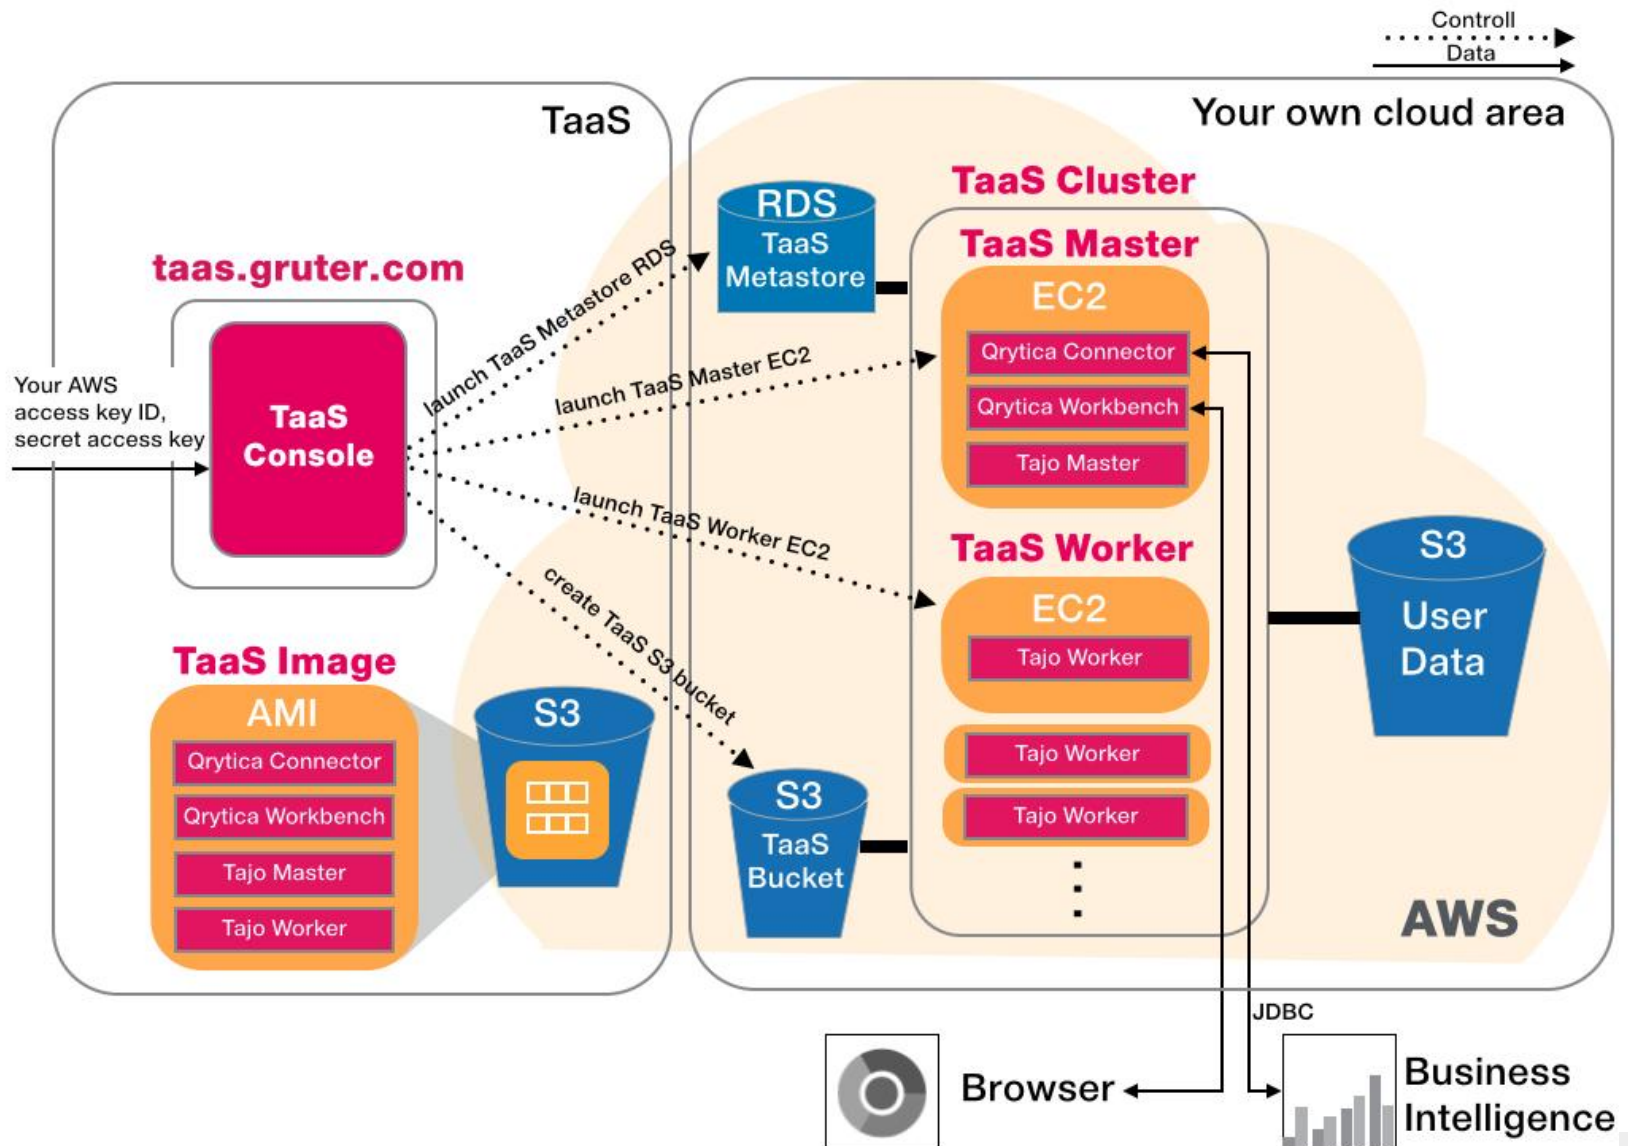
Q5: 30 Text pattern matching filters with OR conditions, group by, having, and order by resulting in smaller set of output

# Replace Commercial Data Warehouse

- ETL(Batch) Processing: 120+ queries, ~4TB read/day
- OLAP(Interactive) Processing: 500+ queries




# Tajo-as-a-Service on AWS



# Active Open Source Community

- Fully community-driven open source
- Stable development team
  - 17 committers + many contributors

## Issues: 30 Day Summary



Issues: **79** created and **92** resolved

# Future Works

- **2014 4Q**
  - HBase intergation
  - In/Exists SubQuery
  - User defined function
  - Multi-tenant Scheduler
- **2015 1Q**
  - Authentication and Standard Access Control
  - Scalar SubQuery
  - ROLLUP, CUBE
- **2015 2Q**
  - Vectorized Engine(C++ Operator)

*Gruter에서 진행하는 마일스톤으로  
Apache Tajo 커뮤니티 방향과는 다를 수 있음*


# **TAJO INTERNAL**

## Basic Operator

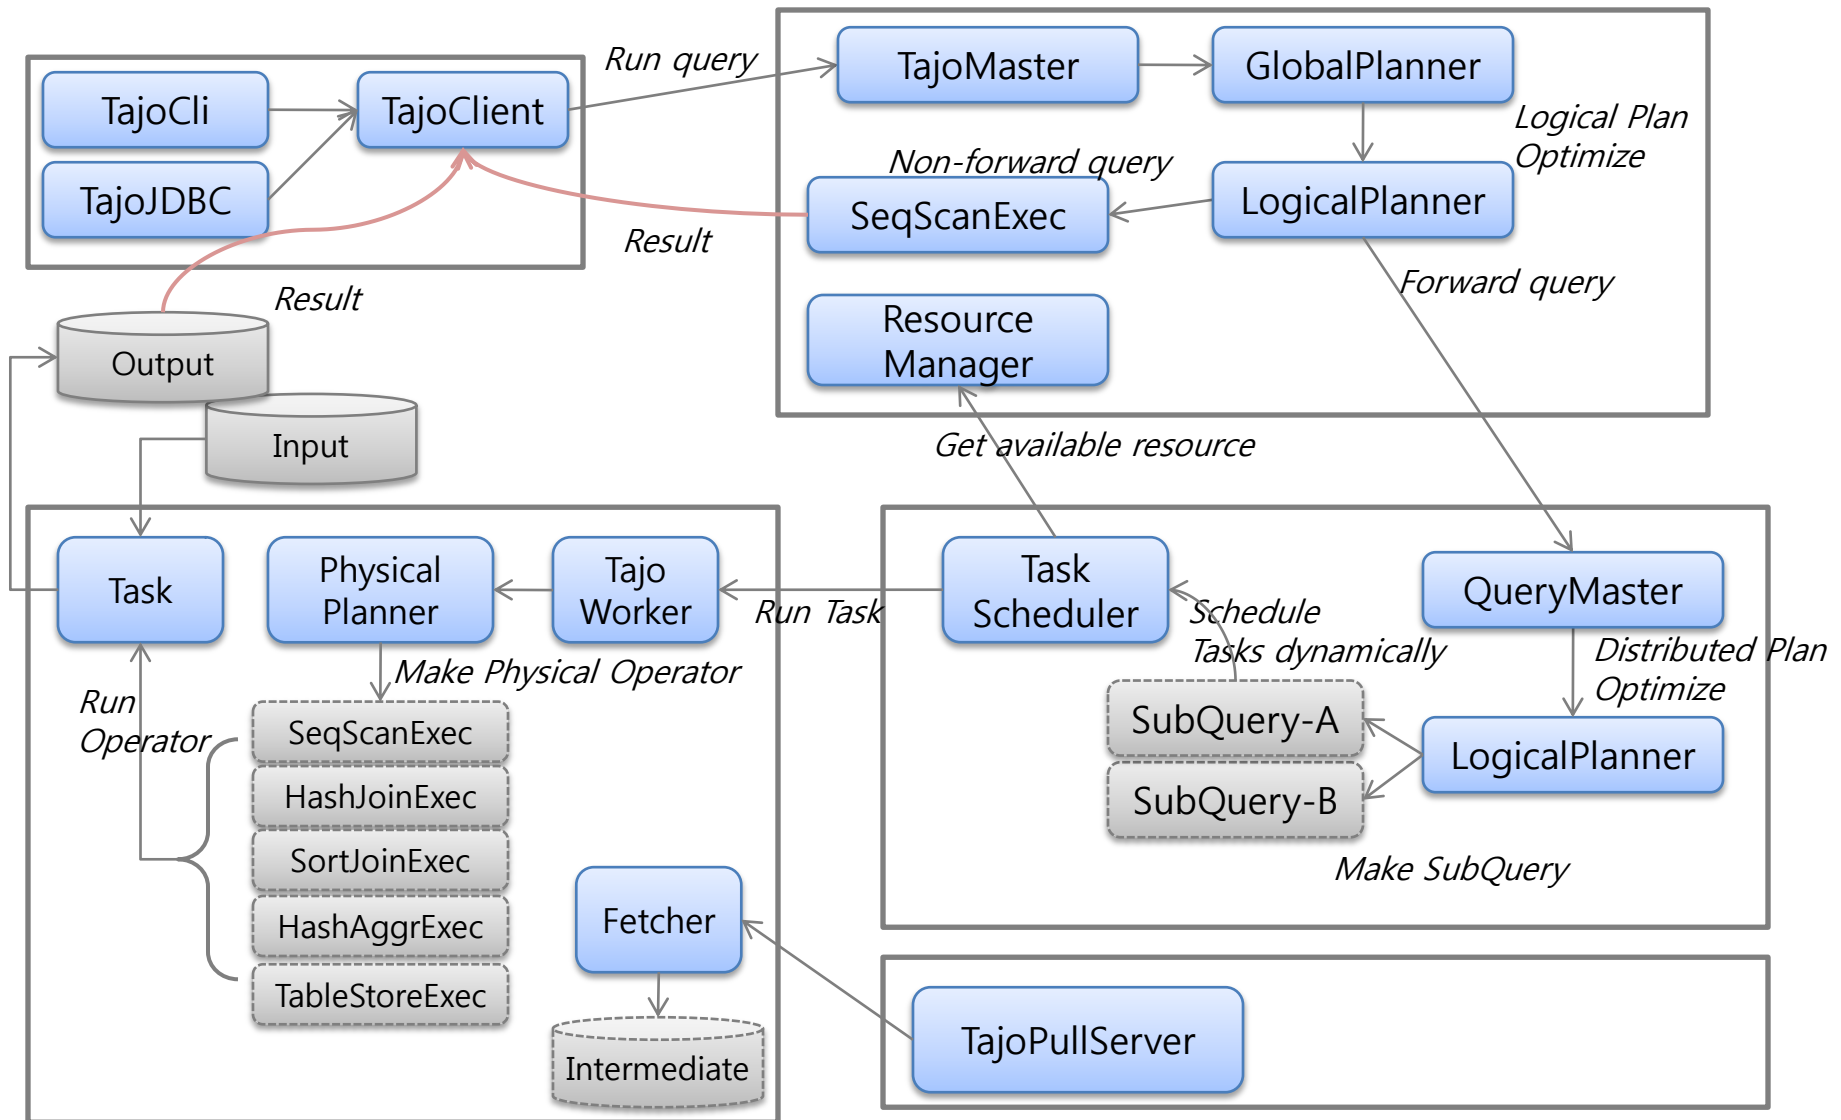


```
SLECT o_custkey, l_lineitem, count(*)  
FROM lineitem a  
JOIN orders b ON a.l_orderkey = b.o_orderkey  
WHERE a.l_shipdate > '2014-09-01'  
GROUP BY o_custkey, l_lineitem
```

*Operator 조합으로  
Tree 형태의  
PhysicalPlan 생성*




# Query Execution Flow

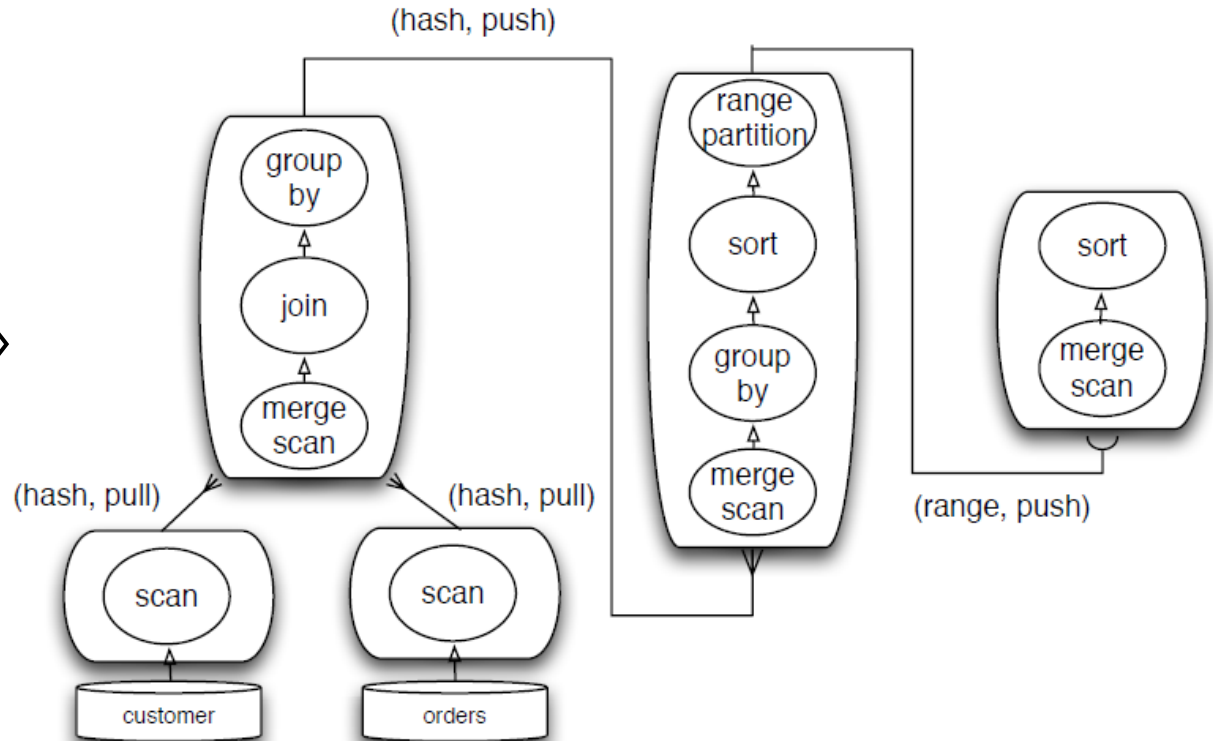




# Logical Plan/Distributed Plan



(A join-groupby-sort query plan)



(A distributed query execution plan)

select col1, sum(col2) as total, avg(col3) as average  
from r1, r2  
where r1.col1 = r2.col2 group by col1 order by average;

# Logical Plan Optimizer

- **Basic Rewrite Rule**
  - Common sub expression elimination
  - Constant folding (CF), and Null propagation
- **Projection Push Down (PPD)**
  - push expressions to operators lower as possible
  - narrow read columns
  - remove duplicated expressions
    - if some expressions has common expression
- **Filter Push Down (FPD)**
  - reduce rows to be processed earlier as possible
- **Extensible Rewrite Rule**
  - Allow developers to write their own rewrite rules

# Logical Plan Optimizer

```
SELECT
  item_id,
  order_id,
  sum_price * (1.2 * 0.3)
as total,
FROM (
  SELECT
    item_id,
    order_id,
    sum(price) as sum_price
  FROM
    ITEMS
  GROUP BY item_id, order_id
) a
WHERE item_id = 17234
```

**Original**

**CF + PPD**



**FPD**

```
SELECT
  item_id,
  order_id,
  sum(price) * (3.6)
FROM
  ITEMS
GROUP BY
  item_id,
  order_id
WHERE item_id = 17234
```

**Rewritten**

# Filter Push Down Rule(Outer Join)

	Preserved Row Table	Null Supplying Table
Join Predicate	Not Pushed(1)	Pushed(2)
Where Predicate	Pushed(3)	Not Pushed(4)

SELECT \* FROM lineitem a

OUTER JOIN orders b ON

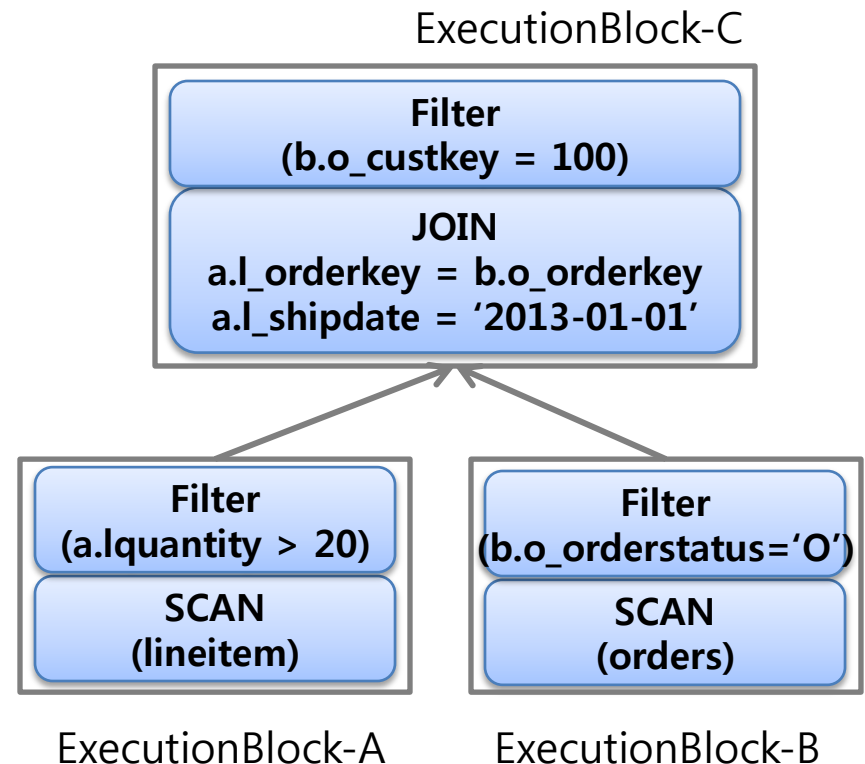
a.l\_orderkey = o\_orderkey AND

① a.l\_shipdate = '2013-01-01' AND

② b.o\_orderstatus = 'O'


WHERE a.l\_quantity > 20 ③

AND b.o\_custkey = 100 ④




# Logical Plan Optimizer

- **Cost-based Join Order**
  - Don't need to guess right join orders anymore
  - Greedy heuristic algorithm
    - Resulting in a bushy join tree instead of left-deep join tree



Left-deep Join Tree




Bush Join Tree

- **Union**
  - Filter push down
- **Repartitioner**
  - `tajo.dist-query.join.partition-volume-mb`
  - `tajo.dist-query.groupby.partition-volume-mb`
  - `tajo.dist-query.table-partition.task-volume-mb`
- **Broadcast Join**
  - `tajo.dist-query.join.broadcast.threshold-bytes`
- **Host/Disk aware scheduler**



# GRUTER: YOUR PARTNER IN THE BIG DATA REVOLUTION



Phone +82-70-8129-2950  
Fax +82-70-8129-2952



Phone +1-415-841-3345

E-mail [contact@gruter.com](mailto:contact@gruter.com)  
Web [www.gruter.com](http://www.gruter.com)