

Apache Tajo on Swift

Bringing SQL to the OpenStack World

Jihoon Son

Apache Tajo PMC member

Who am I

- Jihoon Son
 - Ph.D candidate (Computer Science & Engineering, 2010.3 ~)
 - Apache Tajo PMC and Committer (2014.5.1 ~)
 - Mentor of Google Summer of Code (2013)
 - Contacts
 - Email: jihoonson AT apache.org
 - LinkedIn: <https://www.linkedin.com/in/jihoonson>
-

Outline

- OpenStack Swift
 - Apache Tajo
 - Tajo on Swift
 - Demo
 - Our Roadmap
-

OpenStack Swift

- Popular object storage
 - Images, videos, logs, ...
- Enterprises store objects on Swift to provide their services
 - Usually private clusters

SQL on Swift

- Data analysis is important to improve the quality of their services
 - SQL is one of the most powerful and popular query language
 - Many enterprise data analysis tools relying on SQL
 - OLAP, visualization, data mining, ...
 - Need for using SQL on Swift
-


Apache Tajo

- Scalable, efficient, and fault-tolerant data warehouse system
 - Support SQL standards compliance
 - Efficient batch execution and interactive ad-hoc analysis
 - Low latency and high throughput
 - No use of MapReduce
 - No single point of failure
-

Apache Tajo


- Active open source project
 - 18 committers and 16 contributors
 - Activity summary

Issues: 30 Day Summary




Issues: **47** created and **50** resolved

Apache Tajo




Tajo on Swift



Tajo on Swift

- No need to modify code of Tajo and Swift
 - Tajo can access Swift with the Hadoop-openstack library
 - But, doesn't need to install or run Hadoop
 - Just use it



Tajo on Swift

- Configuration highlights
 - Swift configuration
 - Need the keystone authentication for the Hadoop
 - No additional configurations
 - HDFS configuration
 - Different cloud providers support
 - Key name pattern
 - `fs.swift.service.${provider}`

Tajo on Swift


- Configuration highlights
 - Swift configuration
 - Need the keystone
 - No additional co
 - HDFS configuration
 - Different cloud p
 - Key name pat

fs.swift.ser

```
<property>
  <name>fs.swift.impl</name>
  <value>org.apache.hadoop.fs.swift.snative.SwiftNativeFileSystem</value>
</property>
<property>
  <name>fs.swift.blocksize</name>
  <value>131072</value>
</property>
<property>
  <name>fs.swift.service.tajo.auth.url</name>
  <value>http://192.168.0.1:5000/v2.0/tokens</value>
</property>
<property>
  <name>fs.swift.service.tajo.tenant</name>
  <value>demo</value>
</property>
<property>
  <name>fs.swift.service.tajo.username</name>
  <value>hadoop</value>
</property>
<property>
  <name>fs.swift.service.tajo.password</name>
  <value>{password}</value>
</property>
```


Tajo on Swift

- Data locality problem



Tajo on Swift

- Data locality problem



Advanced Integration

- List endpoints middleware
 - Providing the location information of objects, accounts or containers
 - Tajo workers can directly access each object
 - Example

```
hadoop@t1 ~/tajo-0.10.0-SNAPSHOT $ curl -i -H "Accept: application/json" -H "Content-Type: application/json" -X GET http://192.168.0.1:8080/en
dpoints/v2/tajo/tpch/customer/customer.tbl.1
HTTP/1.1 200 OK
Content-Length: 278
Content-Type: application/json
X-Trans-Id: tx198d9332f1184d51881b1-0054cc763e
Date: Sat, 31 Jan 2015 06:29:19 GMT

{"headers": {"X-Backend-Storage-Policy-Index": "0"}, "endpoints": ["http://192.168.0.5:6000/sdb1/866/tajo/tpch/customer/customer.tbl.1", "http
://192.168.0.6:6000/sdb1/866/tajo/tpch/customer/customer.tbl.1", "http://192.168.0.10:6000/sdb1/866/tajo/tpch/customer/customer.tbl.1"]}hadoop
@t1 ~/tajo-0.10.0-SNAPSHOT $
```

Advanced Integration

- List endpoints middleware
 - Swift configuration

```
[pipeline:main]
pipeline = authToken cache healthcheck keystoneauth list_endpoints proxy-logging proxy-server
[filter:list_endpoints]
use = egg:swift#list_endpoints
```


- Hadoop configuration

```
<property>
  <name>fs.swift.service.tajo.auth.endpoint.prefix</name>
  <value>/endpoints/AUTH_</value>
</property>
<property>
  <name>fs.swift.service.tajo.location-aware</name>
  <value>true</value>
</property>
```


Advanced Integration


- Location-aware computing
 - Moving the processing close to the data
 - Avoiding the performance degradation due to the data transfer over the network
 - Important issue when Tajo and Swift share the same cluster
-

Location-aware Computing




Location-aware Computing

1. Getting object locations from the ring



Location-aware Computing

2. Assigning tasks based on object locations



Demo

Our Roadmap

- Storage layer specialized for Swift
 - Block storage support
 - Cinder and Ceph
 - Provisioning Tajo clusters
 - Sahara
 - Heat, TOSCA
-

Thanks!

<http://tajo.apache.org/>