

TEZ means fast



Wojciech Indyk



@woj_i



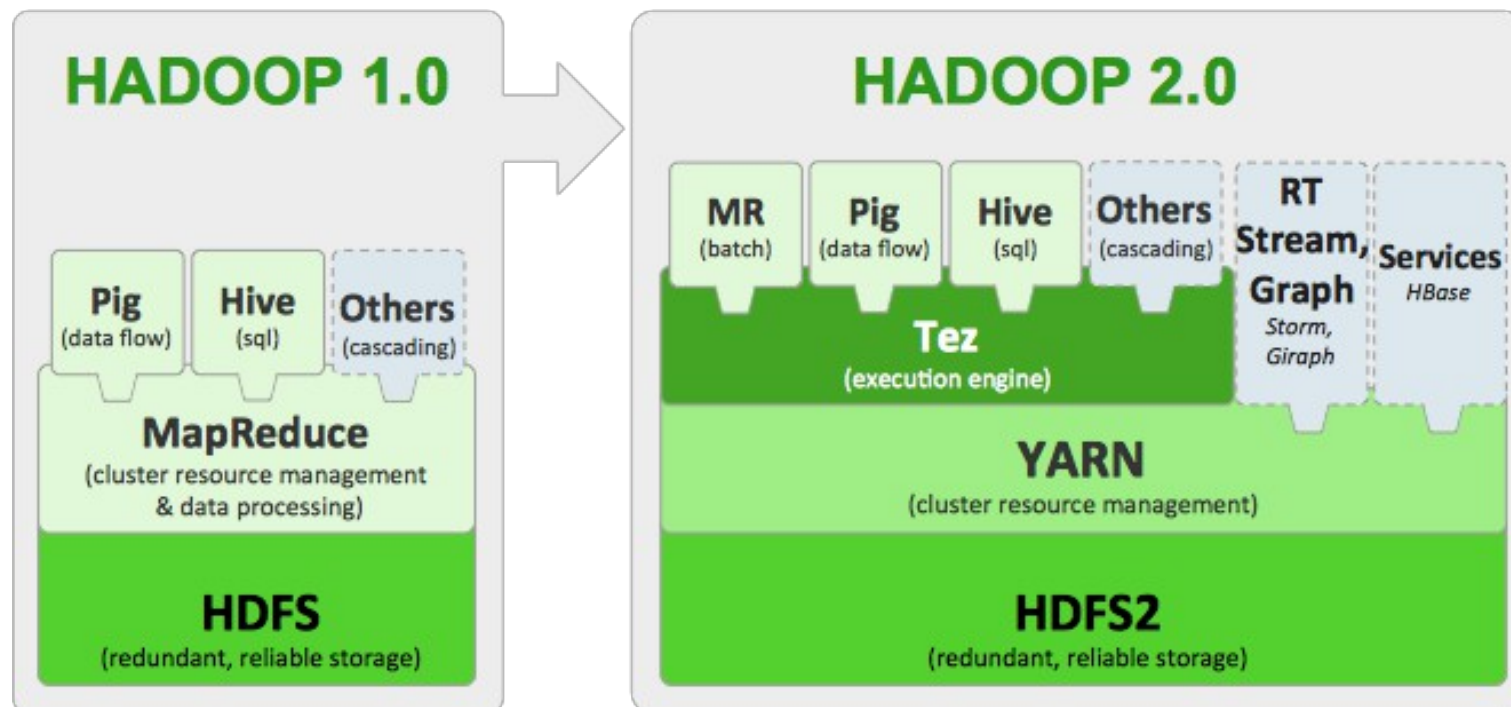
<https://www.linkedin.com/in/wojciechindyk>

Agenda

1. Introduction – evolution of Hadoop stack
2. Processing model- Directed Acyclic Graph
3. Architecture of TEZ
4. Example- MR implementation
5. Performance gain over MR- theory
6. TEZ Sessions
7. Benchmarks
8. Current status, next steps with TEZ


Introduction- evolution of Hadoop stack

- Decomposition of resource manager and computation engine (flexible processing approach)
- Why to interest in TEZ? It is the next generation of MapReduce execution engine



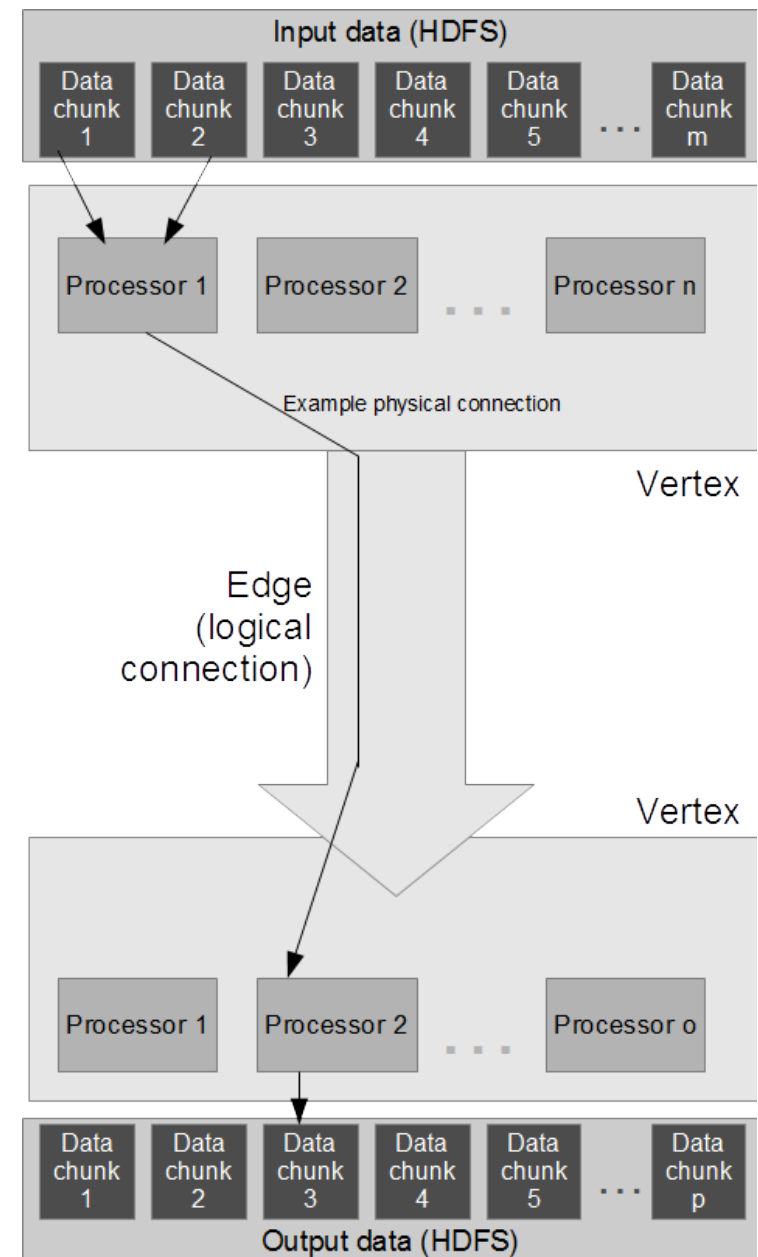
Directed Acyclic Graph

- Vertices and edges
- More elastic than MR
- Dedicated for complex data flow processing
- Intermediate phases can be stored in memory
- Optimizations of processing flow



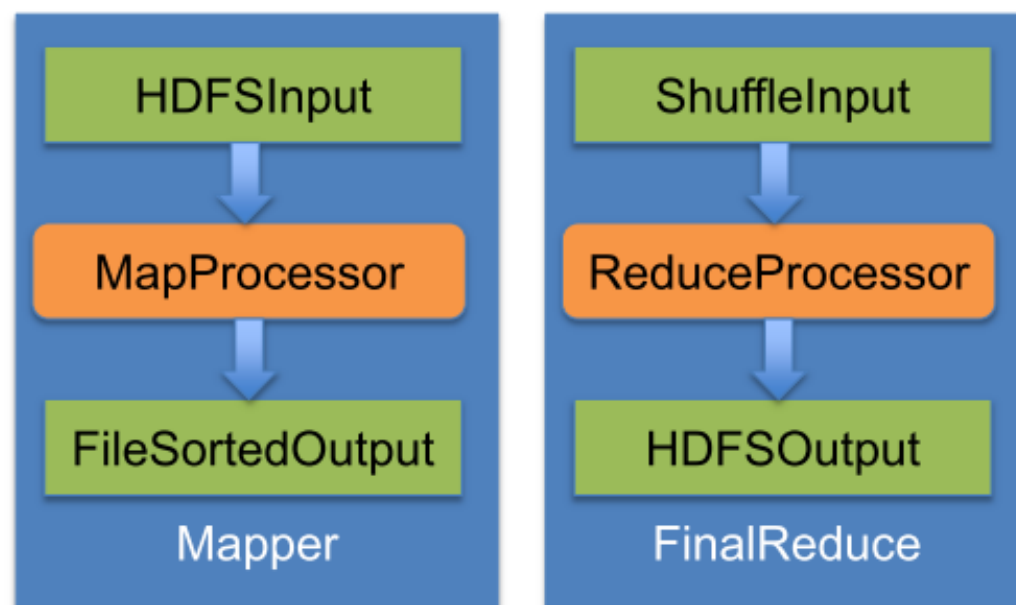
TEZ Architecture

- Vertex as a data processor
- Edges connect vertices
 - Routing (1-1, broadcast, scatter-gather)
 - Scheduling (sequential, concurrent)
 - Data source (persisted, ephemeral)
 - e.g. MapReduce: scatter-gather, sequential, persisted



Example- MR implementation

- Flexible Input-Processor-Output runtime model
- Data type agnostic
- API changes frequently... (TEZ 0.2-0.5) But the idea is still the same
- Let's see the code



ReduceProcessor code

```
public class ReduceProcessor extends MRTask implements LogicalIOProcessor {

    private static final Log LOG = LogFactory.getLog(ReduceProcessor.class);

    private Counter reduceInputKeyCounter;
    private Counter reduceInputValueCounter;

    public ReduceProcessor() {
        super(false);
    }

    @Override
    public void initialize(TezProcessorContext processorContext)
        throws IOException {
        ...
    }

    @Override
    public void handleEvents(List<Event> processorEvents) {
    }

    public void close() throws IOException {
    }

    @Override
    public void run(Map<String, LogicalInput> inputs,
        Map<String, LogicalOutput> outputs) throws Exception {
        THE CONTEXT OF PROCESSOR HERE
        ...
        LogicalInput in = inputs.values().iterator().next();
        ...
        ShuffledMergedInputLegacy shuffleInput = (ShuffledMergedInputLegacy)in; //data-type agnostic
    }
    ...
    @Override
    public void close(TaskAttemptContext context) throws IOException,
        InterruptedException {
    }
}
```

Create DAG

```
DAG dag = new DAG();
```


```
Vertex map1 = new Vertex(MapProcessor.class);  
Vertex map2 = new Vertex(MapProcessor.class);  
Vertex reduce1 = new Vertex(ReduceProcessor.class);  
Vertex reduce2 = new Vertex(ReduceProcessor.class);  
Vertex join1 = new Vertex(JoinProcessor.class);
```

...

```
Edge edge1 = Edge(map1, reduce1, SCATTER_GATHER,  
    PERSISTED, SEQUENTIAL, MOutput.class, RInput.class);  
Edge edge2 = Edge(map2, reduce2, SCATTER_GATHER, PERSISTED, SEQUENTIAL,  
    MOutput.class, RInput.class);  
Edge edge3 = Edge(reduce1, join1, SCATTER_GATHER, PERSISTED, SEQUENTIAL,  
    MOutput.class, RInput.class);  
Edge edge4 = Edge(reduce2, join1, SCATTER_GATHER, PERSISTED, SEQUENTIAL,  
    MOutput.class, RInput.class);
```

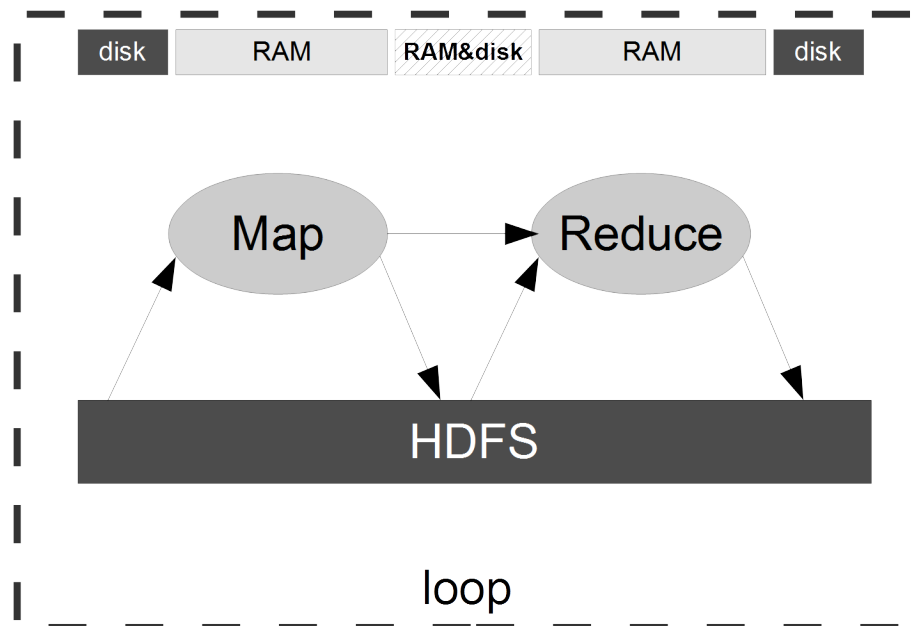
...

```
dag.addVertex(map1).addVertex(map2)  
    .addVertex(reduce1).addVertex(reduce2)  
    .addVertex(join1)  
    .addEdge(edge1).addEdge(edge2)  
    .addEdge(edge3).addEdge(edge4);
```



Performance gain over MR

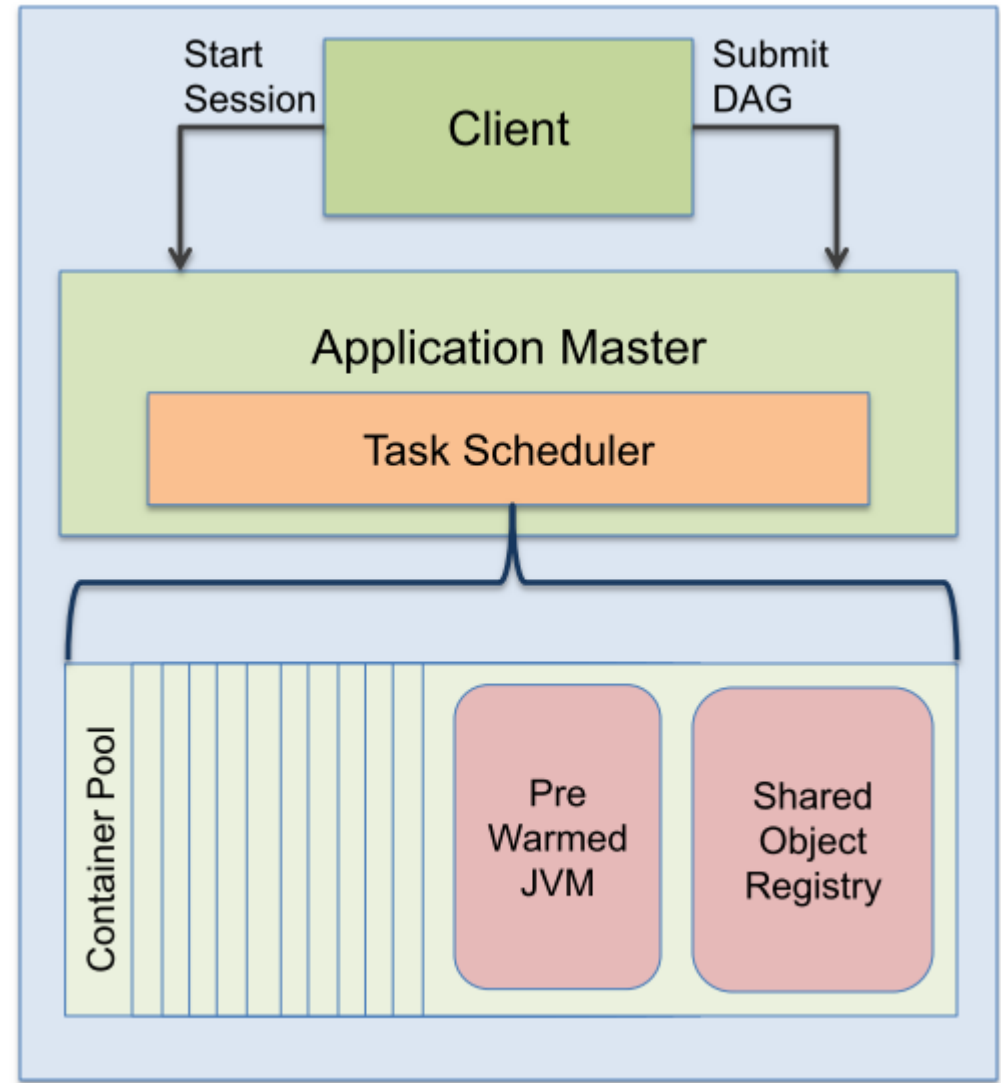
- Eliminate replicated write barrier between successive computations.
- Eliminate job launch overhead of workflow jobs.
- Eliminate extra stage of map reads in every workflow job.



TEZ Sessions

- One Application Master
- Re-use containers
- caching

```
tezSession = new  
TezSession("PageRankSession",  
            appId, sessionConfig);  
tezSession.start();  
...  
tezSession.submitDAG(dag);  
...  
tezSession.stop();
```



When to use sessions

- Analytic tasks on the same data (e.g. drill down)
- Data flows consisted with many sequential tasks
- But also by iterative jobs
 - Optimization of JVM for the same task
 - Near the same resource utilization over iterations
- The resources (memory, CPU) of the AM are fixed so please keep this in mind when configuring the AM for use in a session. For example, memory requirements may be higher for a very large DAG.

Experiment

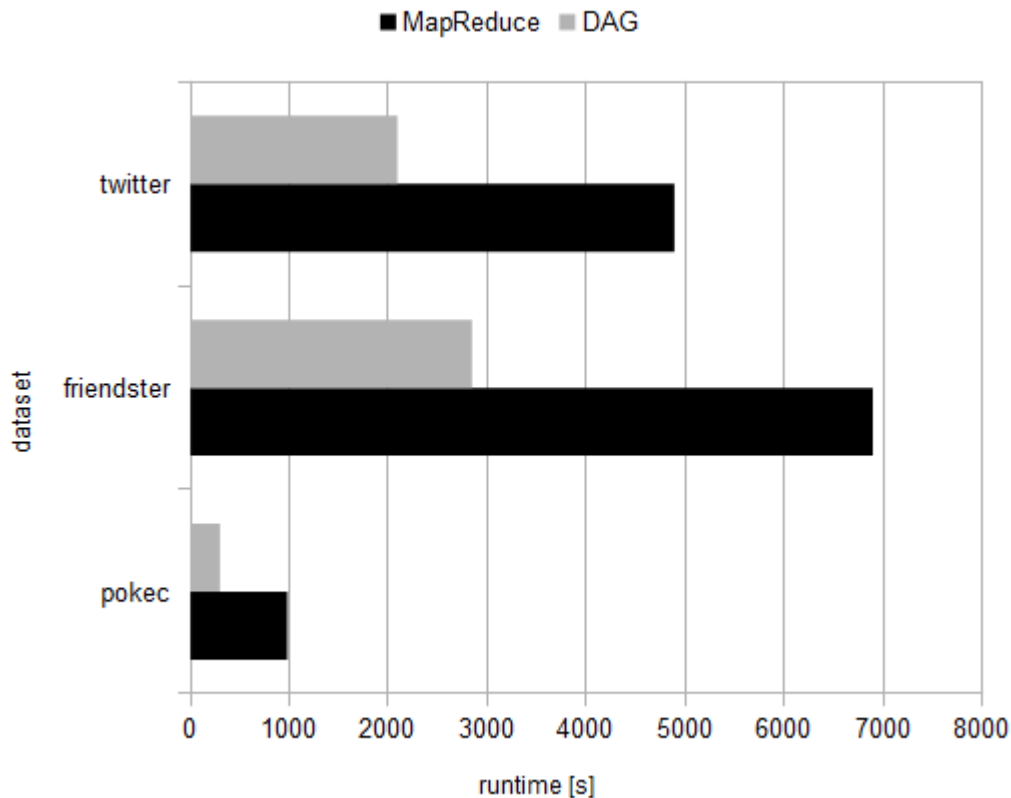
- Algorithms:
 - Single Source Shortest Paths
 - PageRank
- Datasets:

name	nodes	edges	size on disk
Friendster	65,6M	1806,1M	31GB
Twitter	41,7M	1470M	25GB
Pokec	1,6M	30,6M	0,5GB

Environment:


- 14 servers, Hadoop Hortonworks 2.1

Results- PageRank overview

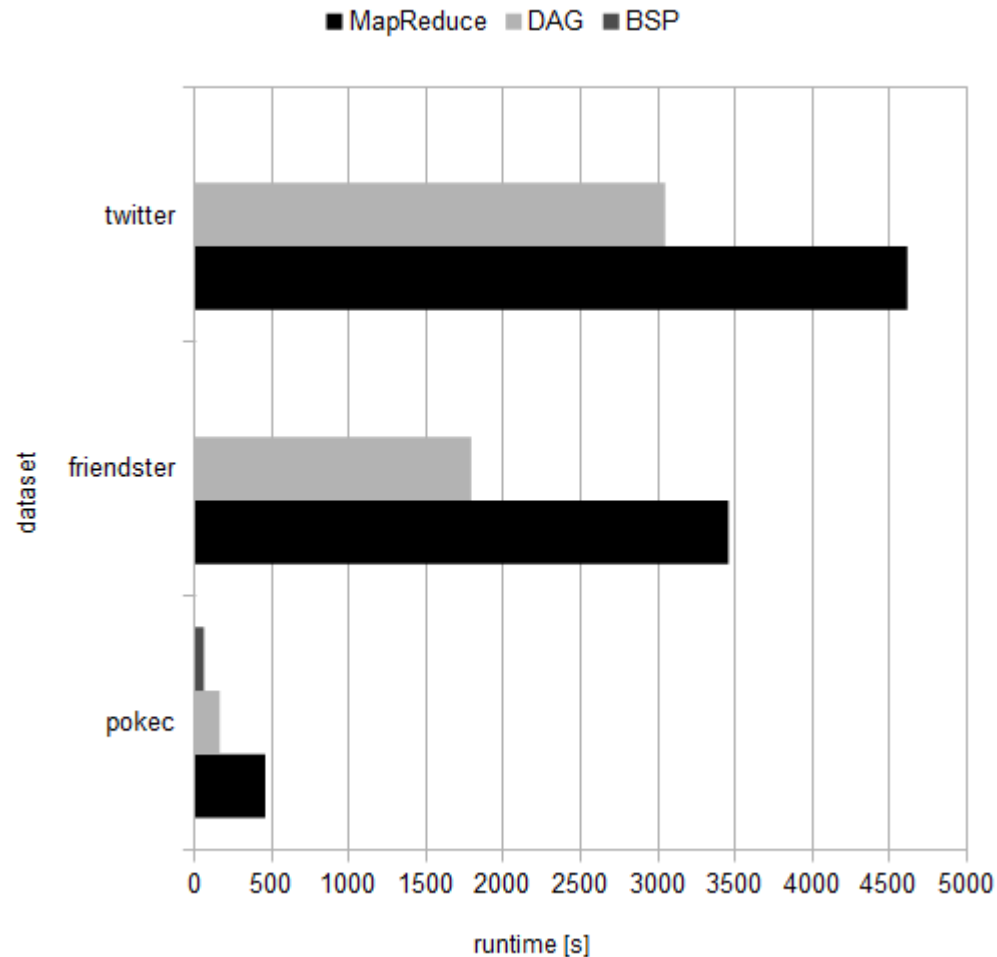


Results- PageRank for Twitter dataset

- Stable iteration results, using session (1-2 iteration) for TEZ is visible

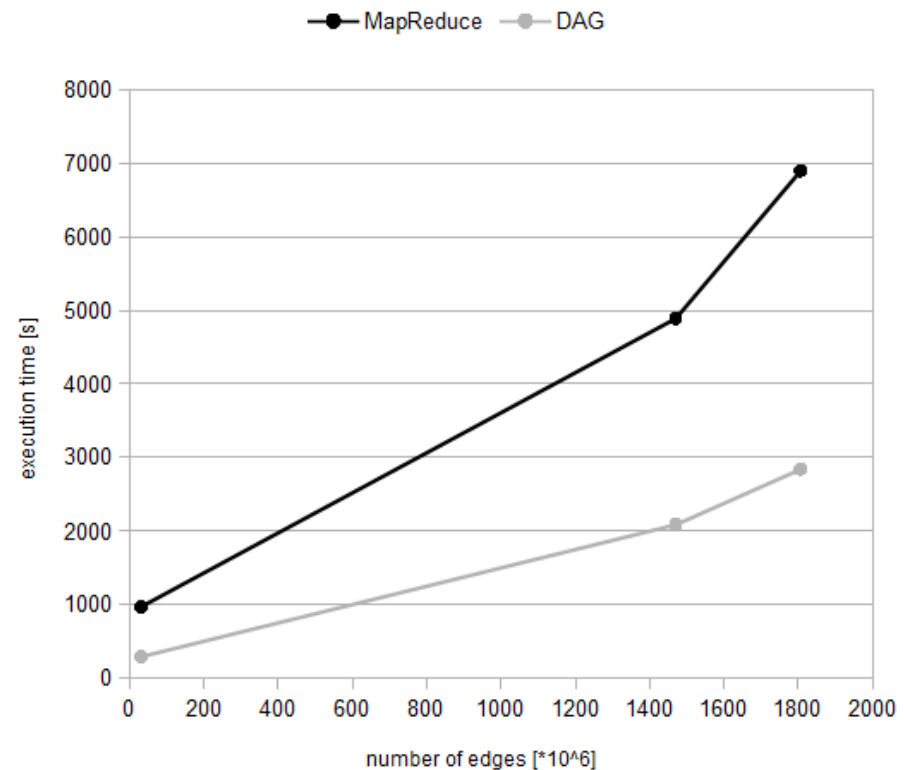
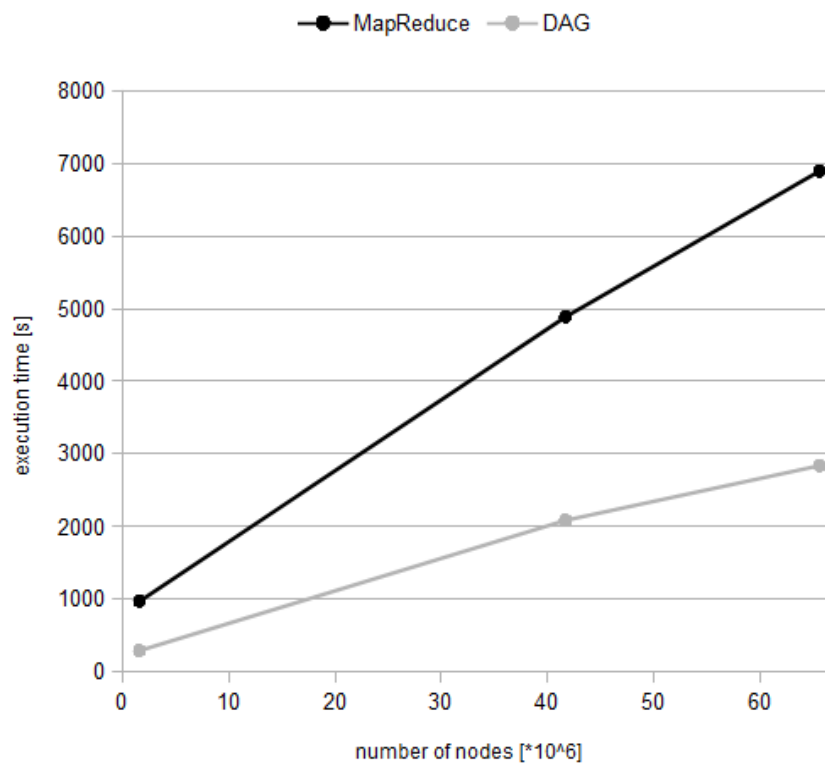


Results Single Source Shortest Paths overview



Results- Scalability of PageRank

- Near linear data-scalability, better scalability for TEZ than MapReduce (in terms of both edges and nodes),



Current status of TEZ

- Still in incubation phase
 - 4.07 – Decision of graduation to „top-level” project
- Development focused on stability of TEZ
- And richer DAG support
 - Better fault tolerance with checkpoints

Difference from Spark

- TEZ focused on large data rather than fast RAM iterations
- Oriented on support for Hive
- IMO TEZ is something between Spark and Impala
 - Oriented on Hive analytic jobs performance
 - Not only for queries
 - Support for complex flows as well as Spark

How to run TEZ immediately

- You can replace MR execution engine with TEZ
- In Hadoop configuration change execution engine „yarn” => „tez-yarn” (at least for Hortonworks)
- Just launch your MapReduce job (it will launch on TEZ)
 - Counters not implemented
 - Worse performance than pure TEZ job

Next steps with TEZ

- Use sandbox! :)
- Create your own DAG
 - At the beginning you can base on [Wordcount example](#)
- Stinger
 - Delivered with Hive 0.13
 - <http://hortonworks.com/labs/stinger/>



<http://lawkarezerwowych.blox.pl/2007/05/Beznadziejny-nawal-pracy.html>

Thank you!

- Questions?
- Comments?
- Ideas?
- Blames? :)