

SMART DATA FAST.™



REAL-TIME BIG DATA ANALYTICS IN THE IBM  
SOFTLAYER CLOUD WITH VOLTDB

# OUR SPEAKERS



John Hugg  
Founding Engineer  
VoltDB



Pethuru Raj Chelliah  
Infrastructure Architect  
IBM SoftLayer



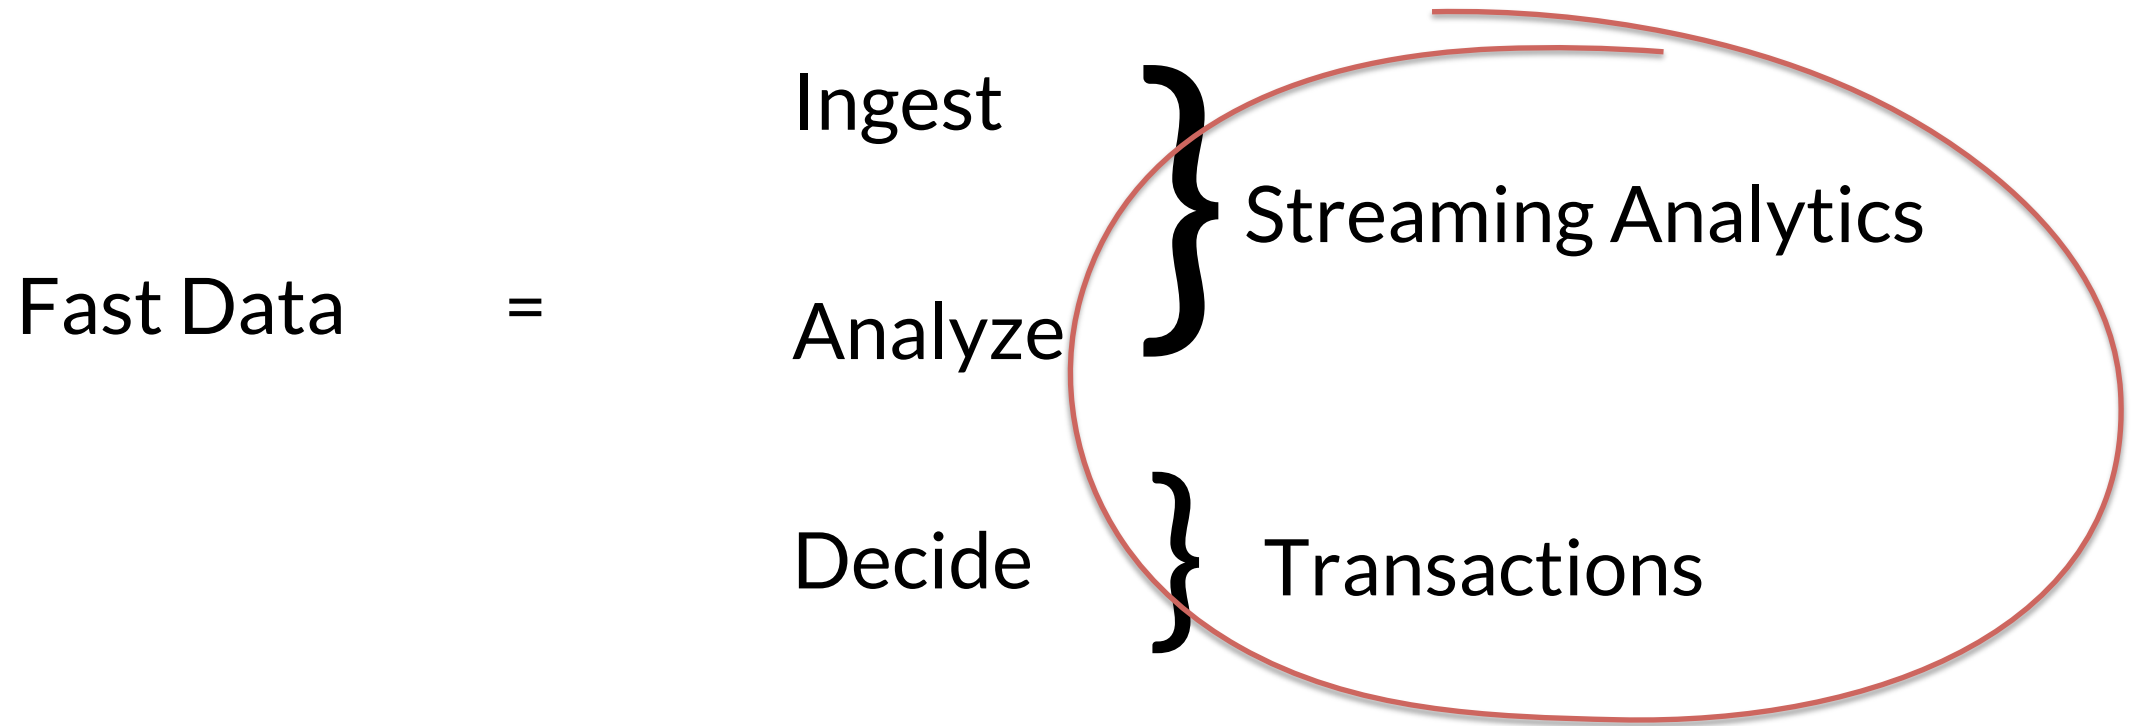
Skylab Vanga  
Infrastructure Architect  
IBM SoftLayer

# THE VOLTDB ARCHITECTURE FOR FAST DATA

- ✓ Per-machine performance + shared nothing scale out
- ✓ Millions of SQL statements per second per machine
- ✓ Up to 500k ACID-Transactional *write* proc calls per sec
- ✓ Full disk durability and sync cluster replication for HA
- ✓ Powerful tools for streaming analytics
  - ✓ Continuous queries, correlation, windowing, aggregation...
- ✓ Hadoop ecosystem integration

VoltDB is really different than everything else

# THE MODERN OPERATIONAL APPLICATION



# 1ST GENERATION FAST DATA: STREAMING ANALYTICS

- Examples:

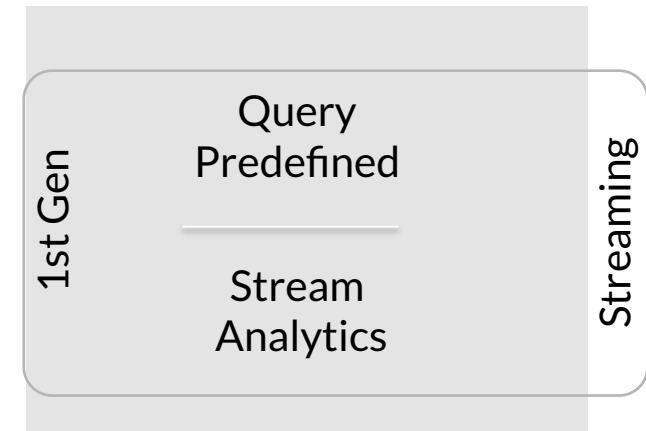
Spark Streaming, Storm, Kinesis, CEP (Streambase, Esper, etc...), et al

- Technical:

- Lacks integrated “state” for transaction processing (operational)
- Complex programming model
- No ability to do ad hoc queries

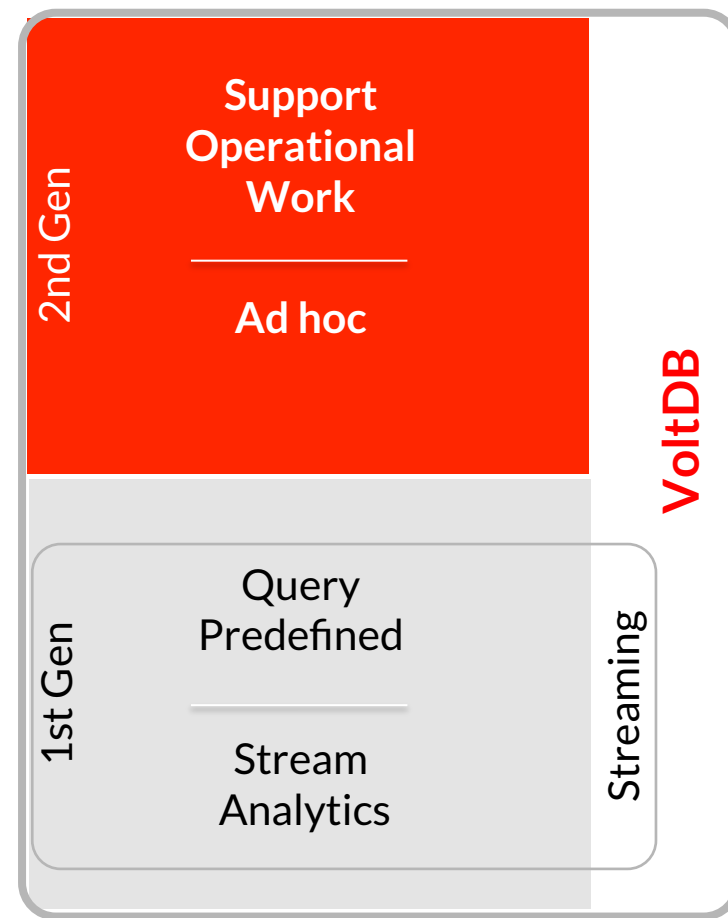
- Functional:

- 1st Gen. only offers streaming analytics on windows of events
- Separate database required for any meaningful work
- Proprietary interface is inconsistent with the rest of the data pipeline
- Does not support applications requirement for interaction



## 2ND GENERATION FAST DATA: COMBINE STREAMING ANALYSIS WITH THE OPERATIONAL STORE

- Operational workloads want to understand data as it changes
- Stream processing work needs state to enable richer analytics
- Operations and Stream Processing are set to converge
  - Convergence is necessary to use streaming analysis in real-time
  - Automated application interactions are informed by data
  - Brings the application into the “data analytics” world
- Streaming Analytics alone is *passive*, Fast Data is *interactive*



# FAST DATA PIPELINE

Data in Motion

Data at Rest

## Customer-Facing

- Personalization
- Customer experience



## Operations-Facing

- Network optimization
- API monitoring
- Sensors

### Streaming Analytics

- Filtering
- Windowing
- Aggregation
- Enrichment
- Correlations

+

### Operational Interaction/Transactions

- Context-aware
- Personal
- Real-time

**VoltDB**

Export

### Batch/Iterative Analytics

- Statistical correlations
- Multi-dimensional analysis
- Predictive analytics



Global Technology Services

# Real-time Analytics by VoltDB on IBM SoftLayer Cloud



# Agenda

- The IT Trends
- About VoltDB
- The Reference Architectures
- A Simple Demo

## The Top Disruptive Transformations

- **Digitization & Deeper Connectivity leads to massive amounts of Data**
- **Data Analytics technologies and tools enables data to information and to knowledge**
- **Knowledge-filled up Services and Applications for the Envisioned Smarter Planet**

## The Prominent Implications

- Trillions of Digitized Objects through digitization and edge technologies
- Billions of connected devices through deeper connectivity and service-enablement of physical, mechanical, electrical, electronics and IT systems
- Millions of software services (device-specific as well as agnostic)
- Zettabytes of multi-structured data
- Newer Analytical Capabilities – prognostic, predictive, prescriptive and personalized insights
- Cognitive Systems

## The Prominent Challenges

- Hyper-converged Data Analytics Platforms
- Analytics-aware IT Infrastructures

## What to do with the bulging data?

- Indulge in analytics to extract insights
- Implement Insights-driven Applications

# Why Data Analytics is vital?

Data analytics is a process for squeezing out actionable insights out of data towards tactical and strategic decision-enablement for individuals as well as institutions

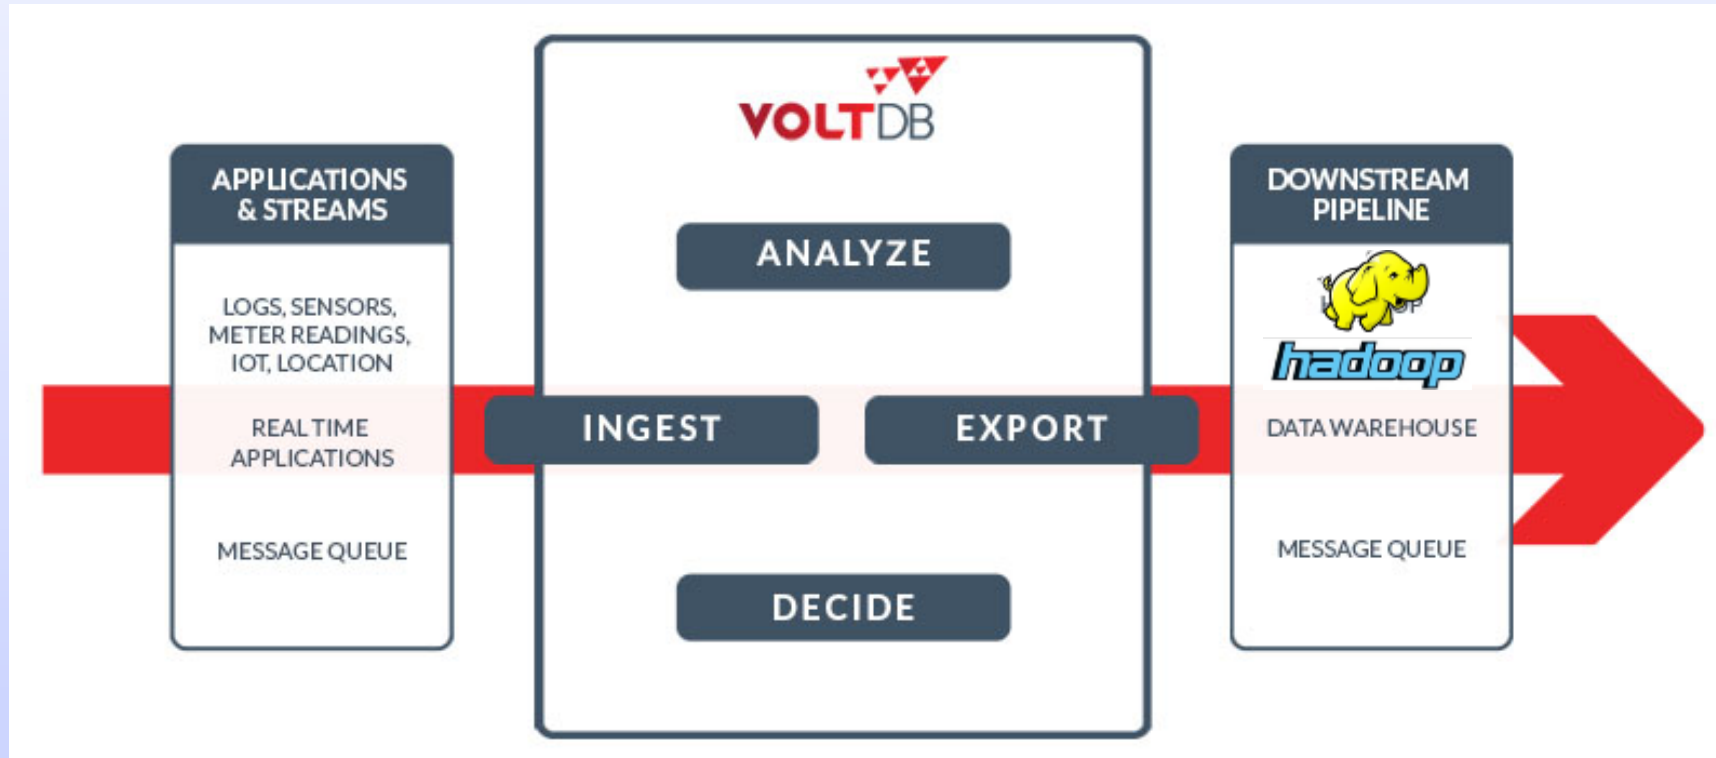
## Insights

1. Enabling business Process Innovation
2. Exploring fresh avenues for fresh revenues
3. Elevating business efficiency and value
4. Enhancing business operations, retaining customer loyalty, attaining newer customers,
5. Expounding outside-in thinking in product designs

## VoltDB for Real-time Data Analytics

- VoltDB is an in-memory scale-out database built to power the next generation of streaming data applications.
- VoltDB's modern architecture provides fast data ingestion and export with massive scalability, real-time analytics, and data enrichment.
- VoltDB is a relational database system (RDBMS) for high-throughput, operational applications requiring:
  - Orders of magnitude better performance than a conventional DBMS
  - SQL as the native data language
  - ACID transactions to ensure data consistency and integrity
  - Built-in high availability (database fault tolerance)
  - Database replication (for disaster recovery)

# The Big Picture

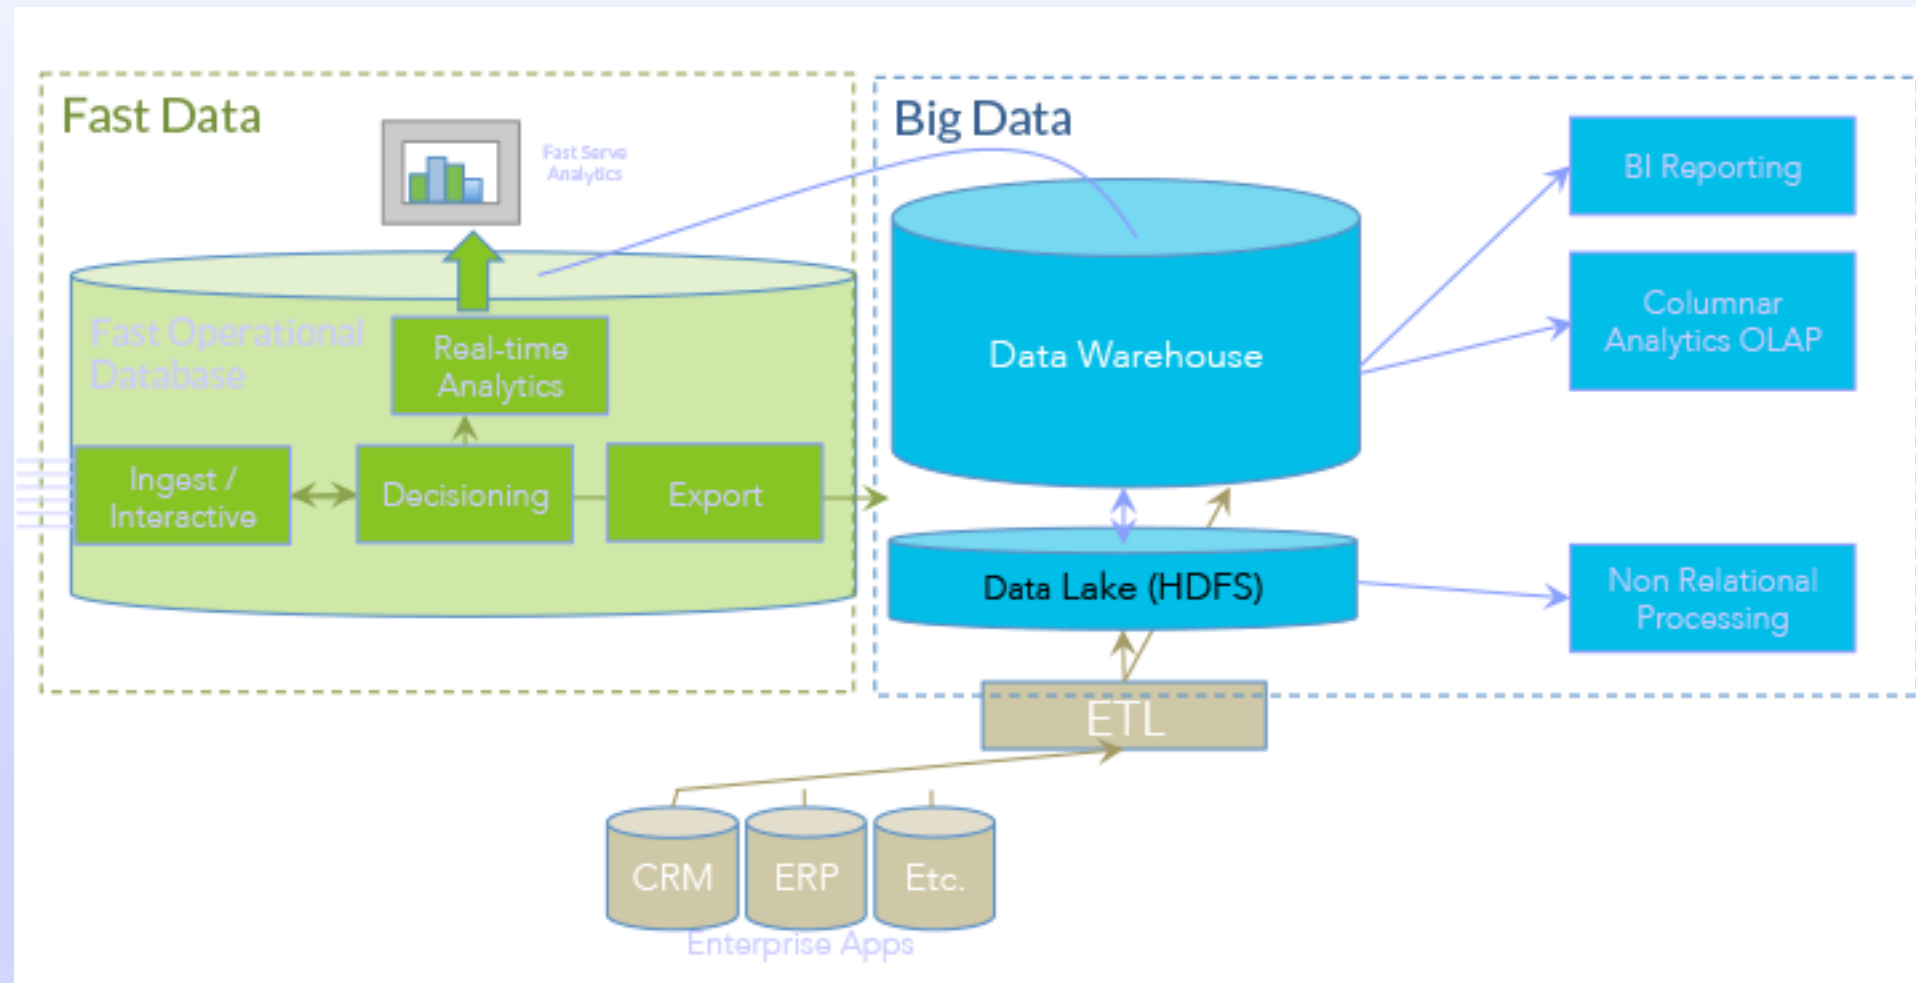


## How VoltDB eliminates the RDBMS Problems for real-time analytics

- Data and associated processing are partitioned together and distributed across the CPU cores (“partitions”) of a shared-nothing hardware cluster.
- Data is held in main memory, eliminating the need for buffer management.
- Transactions execute sequentially in memory, eliminating the need for locking and latching.
- Synchronous multi-master replication provides built-in high availability.
- Command logging provides a high performance replacement for “write-ahead” data logging.



## The Data Architecture for Fast and Big Data



## VoltDB on IBM SoftLayer Cloud

- Ideally suited to streaming applications in the cloud.
- In addition to providing extremely fast ingest, VoltDB exports data at high speed to long-term analytics stores.





Global Technology Services

## A Simple Demo on IBM SoftLayer Cloud

## VoltDB installation

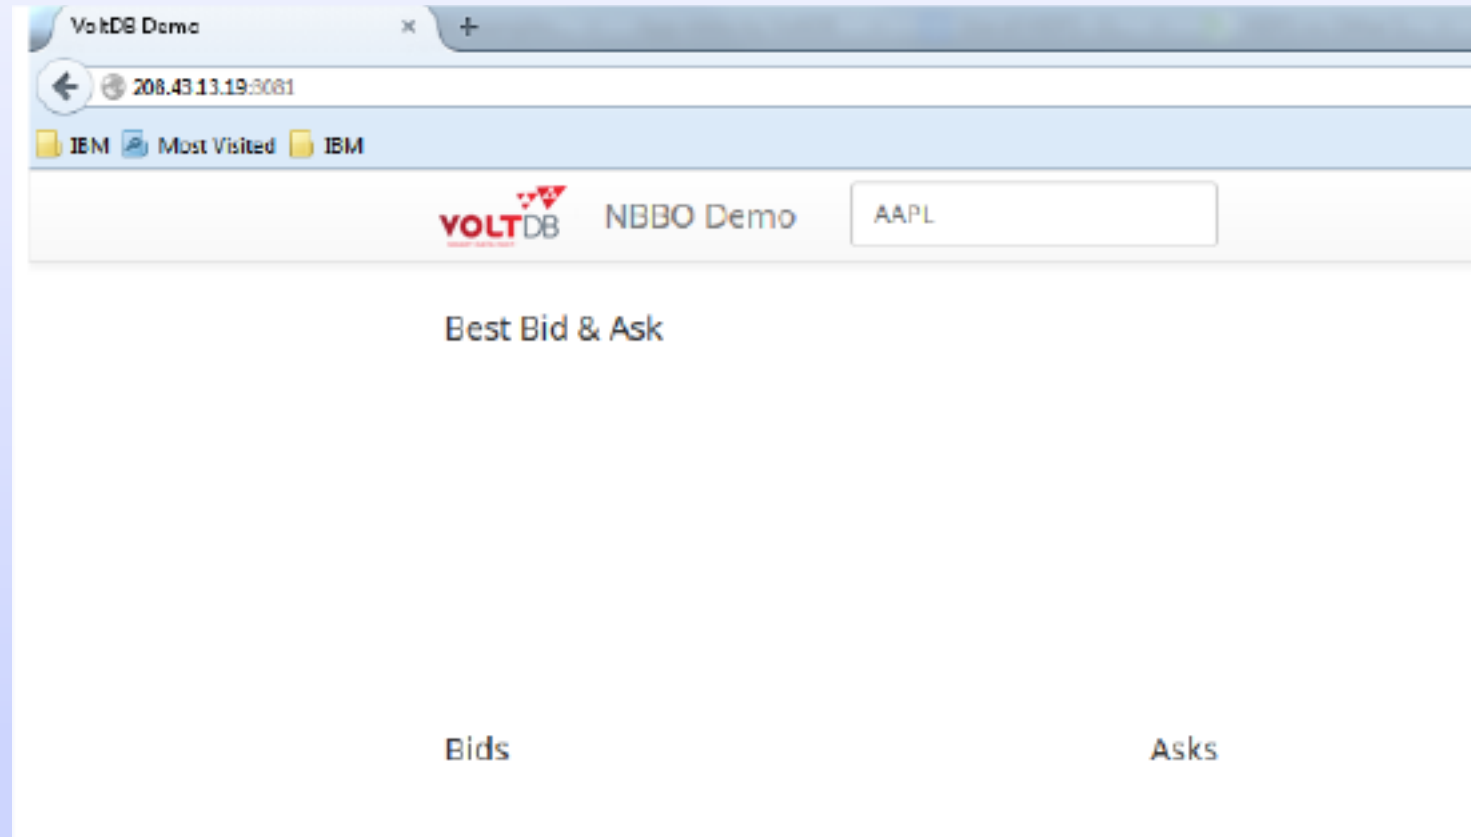
```
Initializing VoltDB...

  _____
 | |  / /  _  / /  / /  / /  \  _  |
 | |  / /  \  / /  / /  / /  _  |
 | |  / /  / /  / /  / /  / /  |
 | |  \  /  \  \  /  \  /  \  |
 |_____/_____/_____/_____/_____/

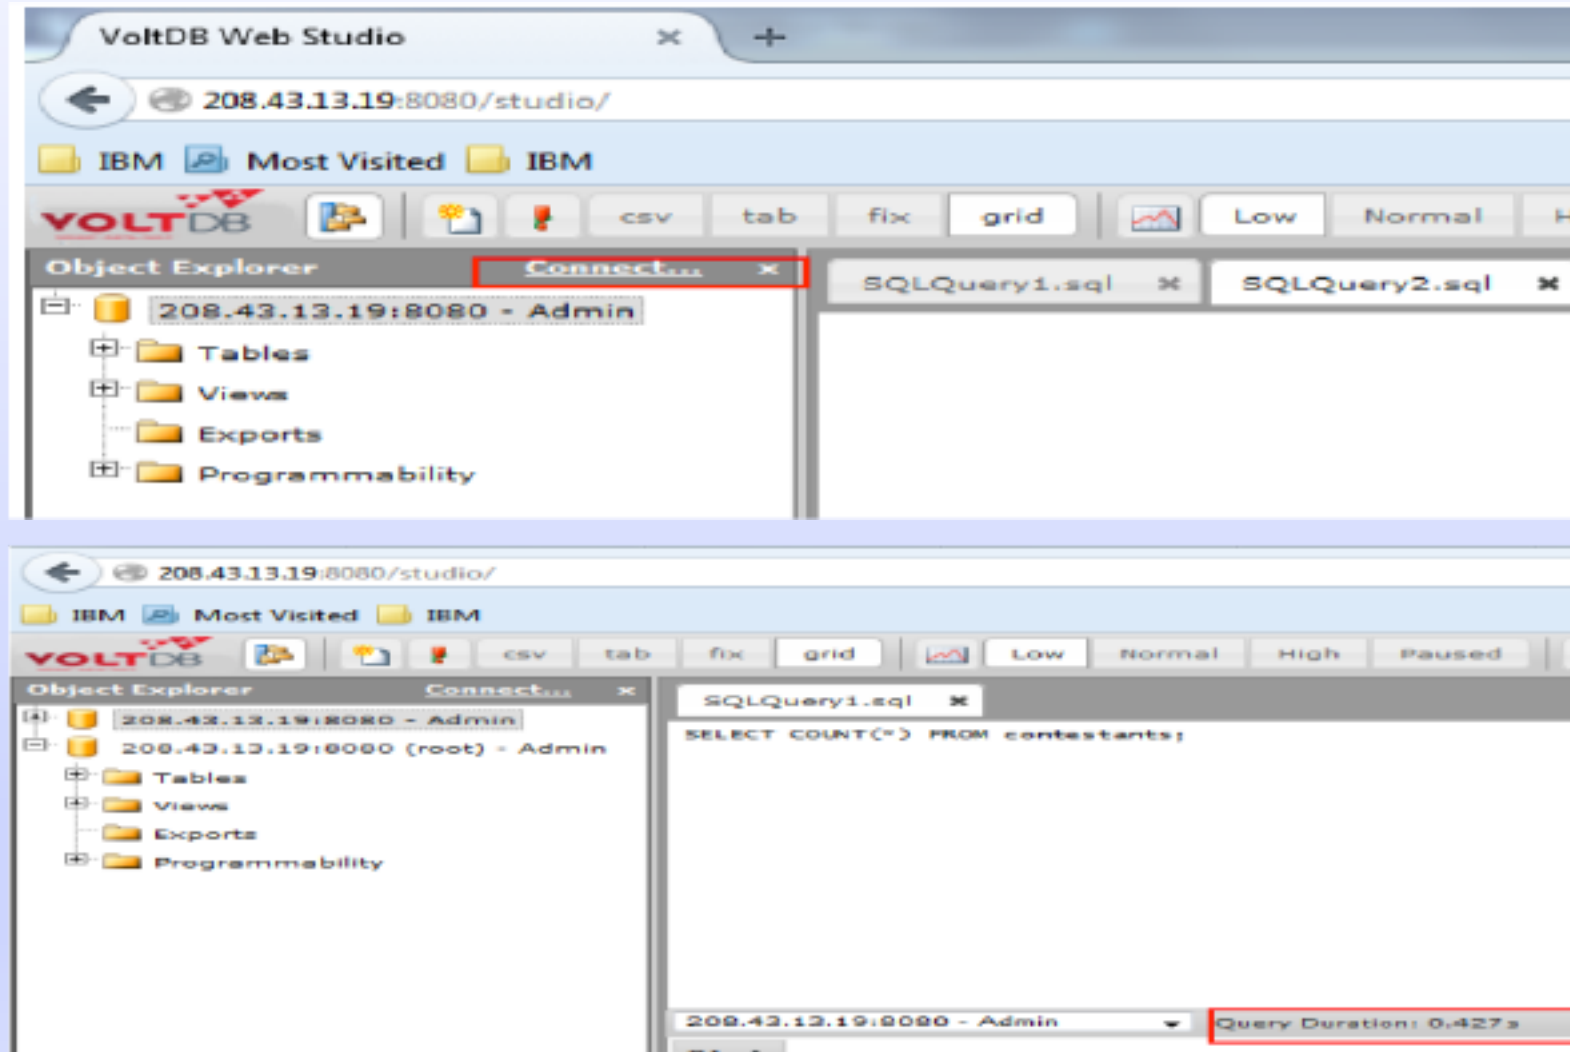
-----

Build: 4.7 voltdb-4.7-0-ge2560ad-local
Connecting to VoltDB cluster as the leader
Host id of this node is: 0
Starting VoltDB with trial license. License key:
Initializing the database and command interface
WARN: This is not a highly available cluster
Server completed initialization.
```

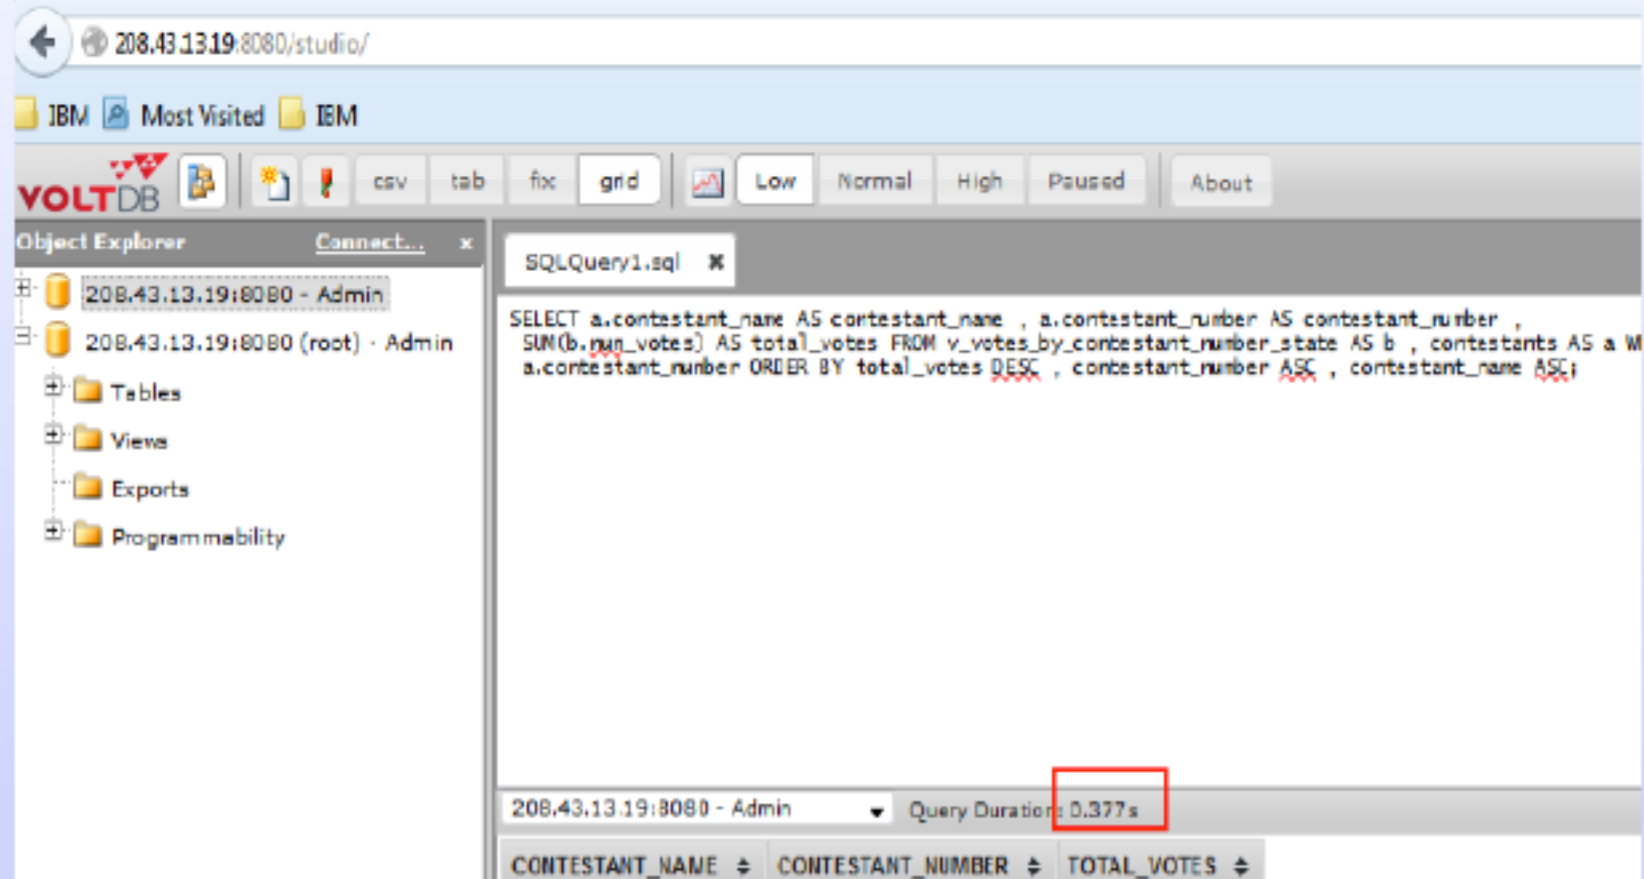
## Application main page



# Connecting VoltDB Server









Thanks for listening

# QUESTIONS?

- Use the chat window to type in your questions
- Try VoltDB yourself:
  - Free trial of the Enterprise Edition:
    - [www.voltodb.com/Download](http://www.voltodb.com/Download)
  - Open source version is available on github.com
- Email us your question: [askanengineer@voltodb.com](mailto:askanengineer@voltodb.com)