



SMART DATA FAST.™



VOLTDDB AND FLYTXT PRESENT:
BUILDING A SINGLE TECHNOLOGY PLATFORM FOR
REAL-TIME AND ITERATIVE ANALYTICS ON FAST + BIG
DATA

OUR SPEAKERS



Ryan Betts
CTO at VoltDB



Prateek Kapadia
CTO at Flytxt

VOLTDDB OVERVIEW

Founded by winner of the 2014 ACM Turing Award

Mike Stonebraker



Other Stonebraker Companies



Technology

- In-Memory (but data is durable to disk)
- Scale-Out shared-nothing architecture
- Reliability and fault tolerance
- SQL + Java with ACID
- Hadoop and data warehouse integration
- Open source and commercially licensed (24X7)


FAST

World Record Cloud Benchmark:

YCSB (Yahoo Cloud Serving Benchmark) - 2.4 million tps (transactions per second)

Customers





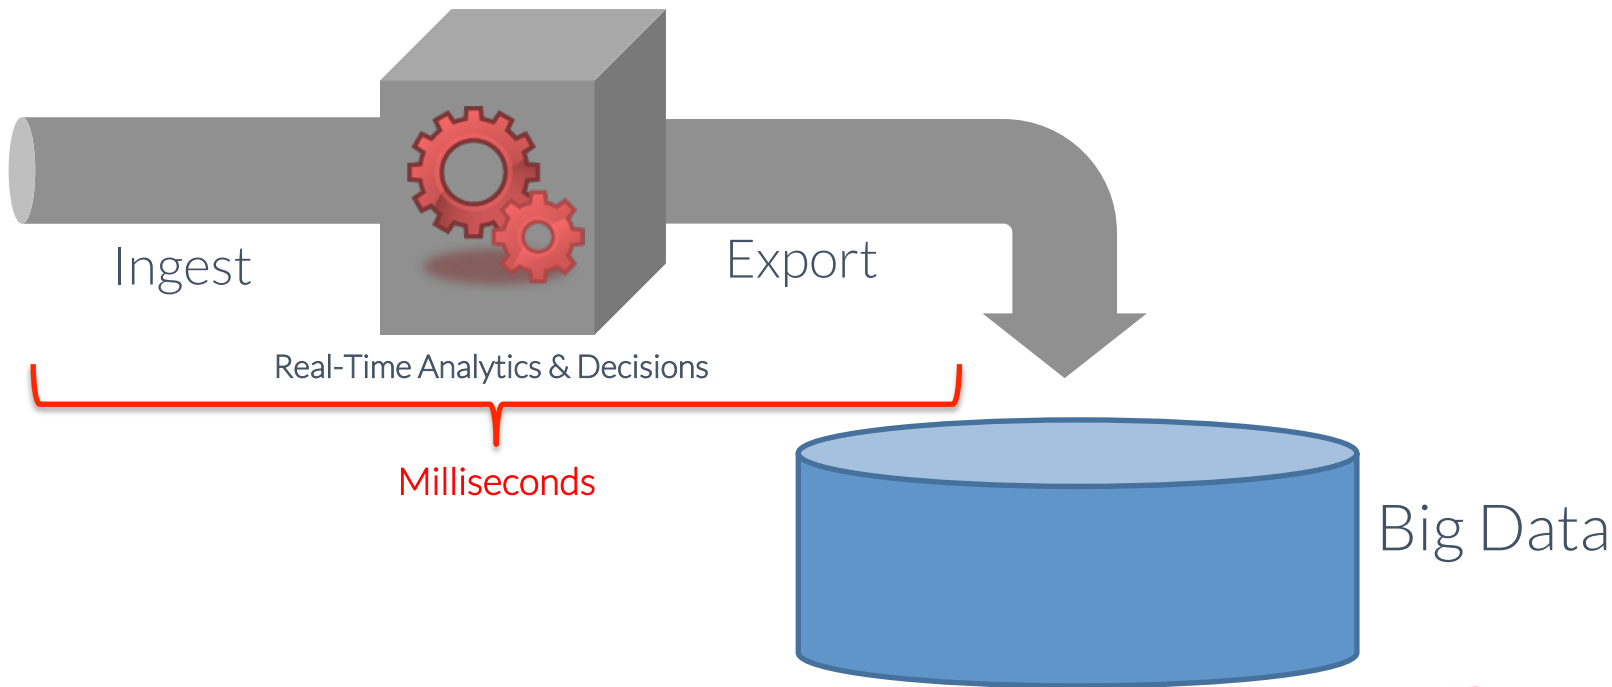
Big Data analytic results:

1. *Discoveries*: seasonal predictions, scientific results, long-term capacity planning
2. *Optimizations*: market segmentation, fraud heuristics, optimal customer journey

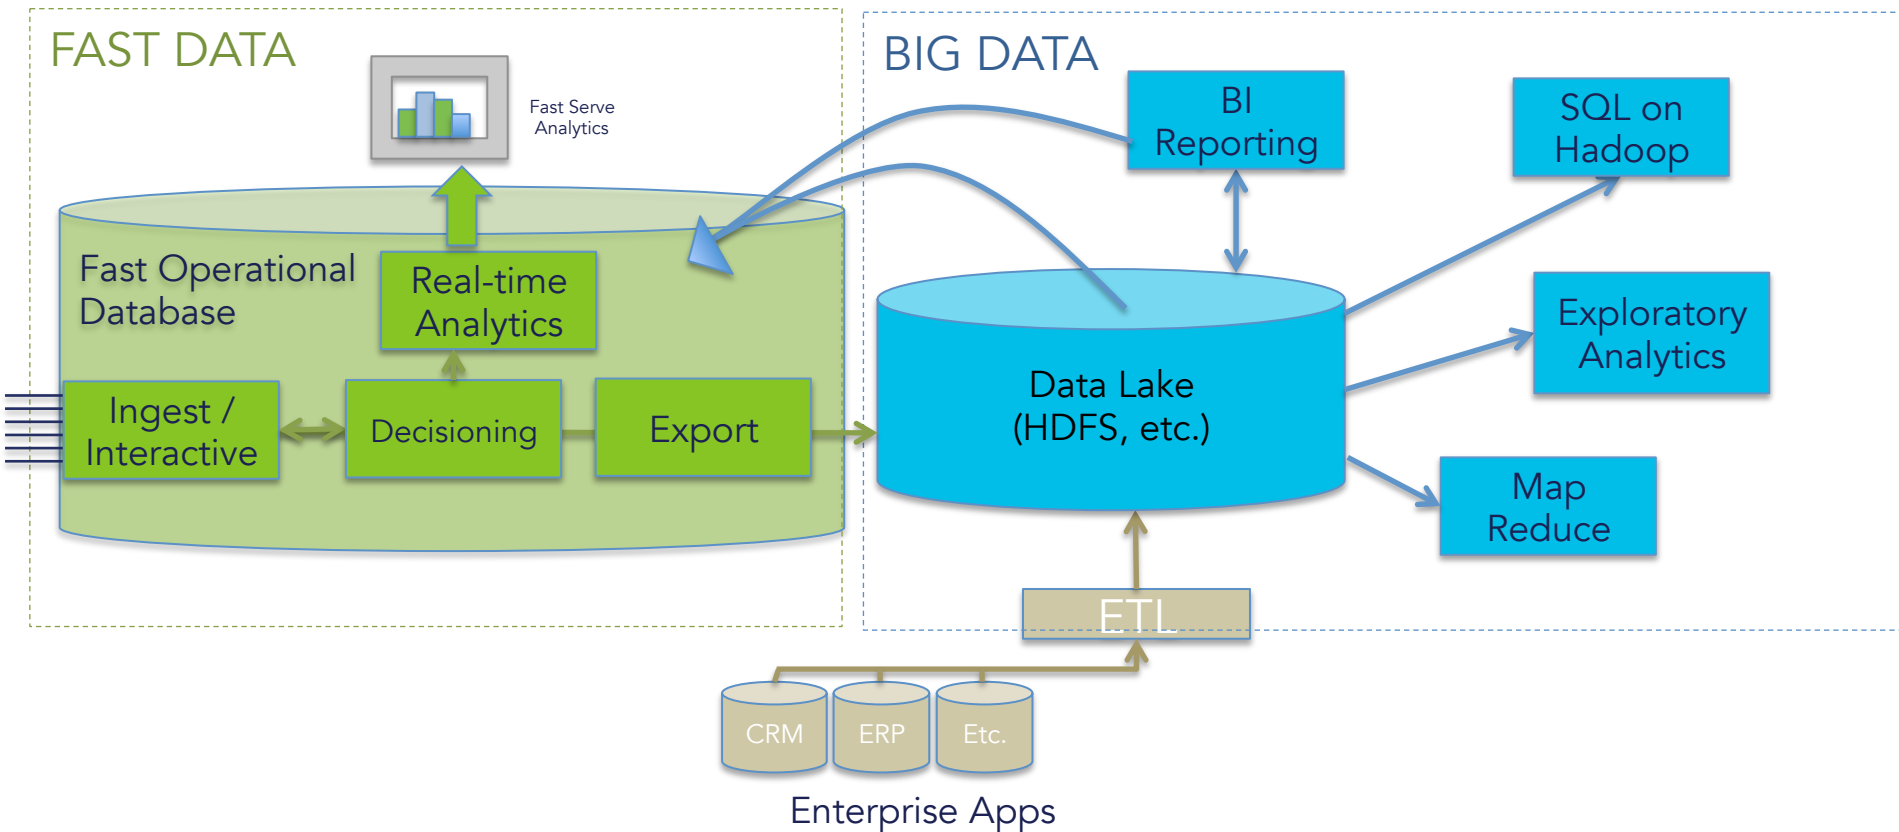
FAST DATA – BIG DATA

Fast Data Pipeline

Fast Data: the velocity side of Big Data



DATA ARCHITECTURE FOR FAST + BIG DATA



“89% of marketers surveyed plan to compete primarily on the basis of customer experience by 2016.”

Source: Gartner 2014 survey, Companies > \$50M in revenue

FAST DATA SOURCES AND DRIVERS

Mobile

IoT

Social

Sensors

Logs



Data is doubling every two years

Mobile

9.5 BILLION

mobile subscriptions
by the end of 2020

90%

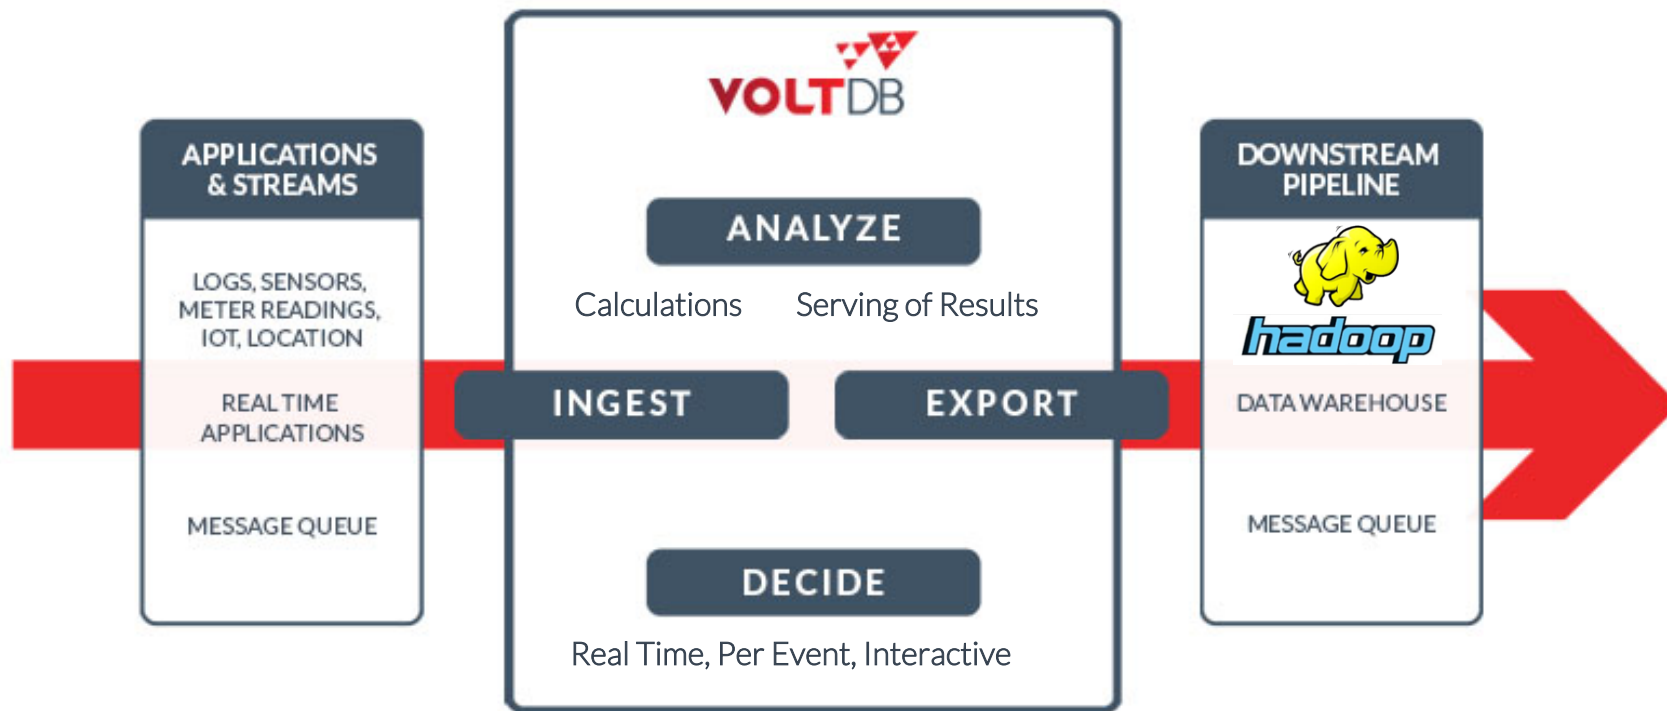
of the world's population
over 6 years old will have
a mobile phone by 2020



IoT

- 26 billion connected devices by 2020 (Gartner 2014)
- 37% of most data will be processed at the edge in milliseconds (Cisco IoT Study 12/11/14)

THE FAST DATA PIPELINE



STREAMING: REAL TIME ANALYTICS



- Operational analytics and monitoring
- RT analytics enabling user-facing applications
- KPI for internal BI/Dashboards

STREAMING OPERATORS NEED STATE

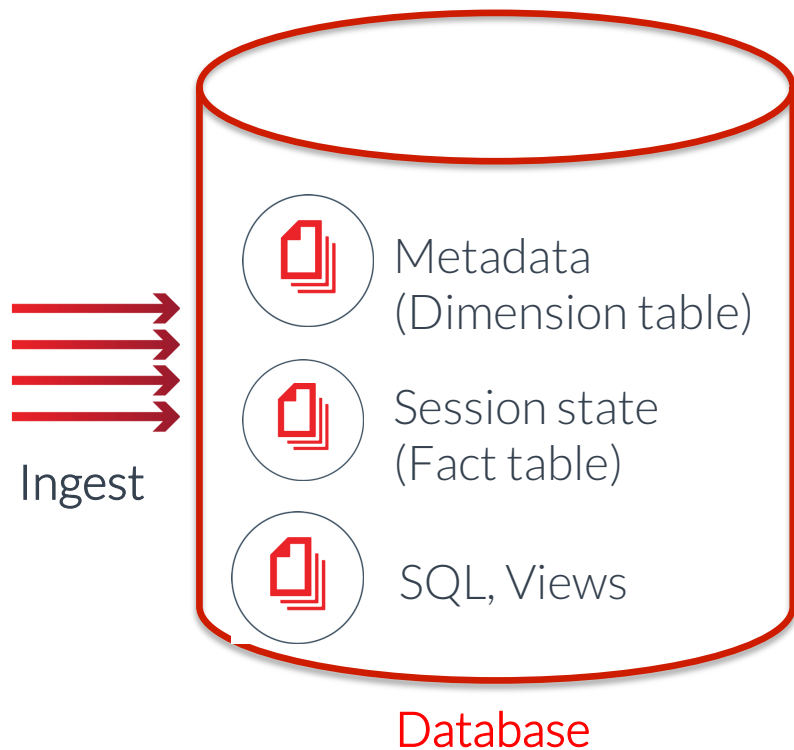
Require State

- Filter
- Join
- Aggregate
- Group By

Stateless


- Partition

REAL-TIME ANALYTICS



- In-memory MPP SQL over ODBC/JDBC
- Cheap + correct materialized views for streaming aggregations
- Operational analytics and monitoring
- RT analytics enabling user-facing applications
- KPI for internal BI/Dashboards

INTEGRATING WITH EXPORT TARGETS

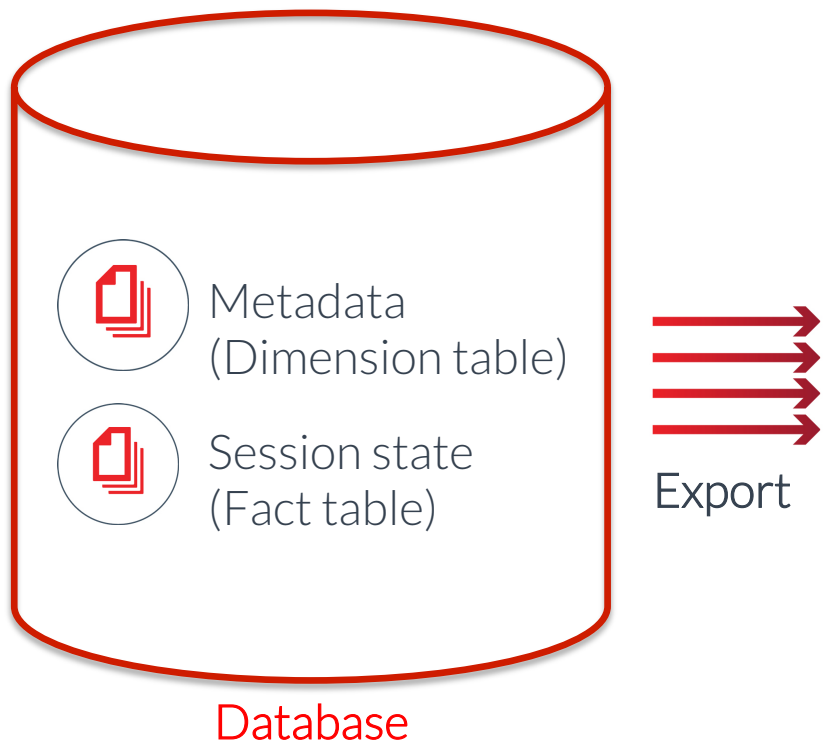


- Local file system export
- JDBC export
- Kafka export
- RabbitMQ export
- HDFS export
- HTTP export
- Extensible API

EXPORT FORMATS

- CSV
- TSV
- Avro container
- Raw data

DATA PIPELINES WITH EXPORT



- MPP streaming Export
 - Row data, Thrift messages, CSV
 - OLAP, HDFS and message queues
-
- Filtering (ex: only RFID / iBeacon readings that show change from previous location).
 - Sessionization
 - Common version re-writing
 - Data enrichment

FLYTXT AND VOLTDDB

FLYTXT OVERVIEW

Vision, Mission & Impact

- ▶ Vision: Create >10% measurable economic value for Communication Services Providers through Big Data Analytics
- ▶ Flytxt's internal and external monetization solutions increase revenue, reduce churn and improve customer experience
- ▶ Dutch company with corporate office in Dubai, global delivery centres in India and regional presence in Mexico City, Johannesburg, Singapore, Dhaka and Nairobi.

Awards & Achievements



Customers and Partners

Operators



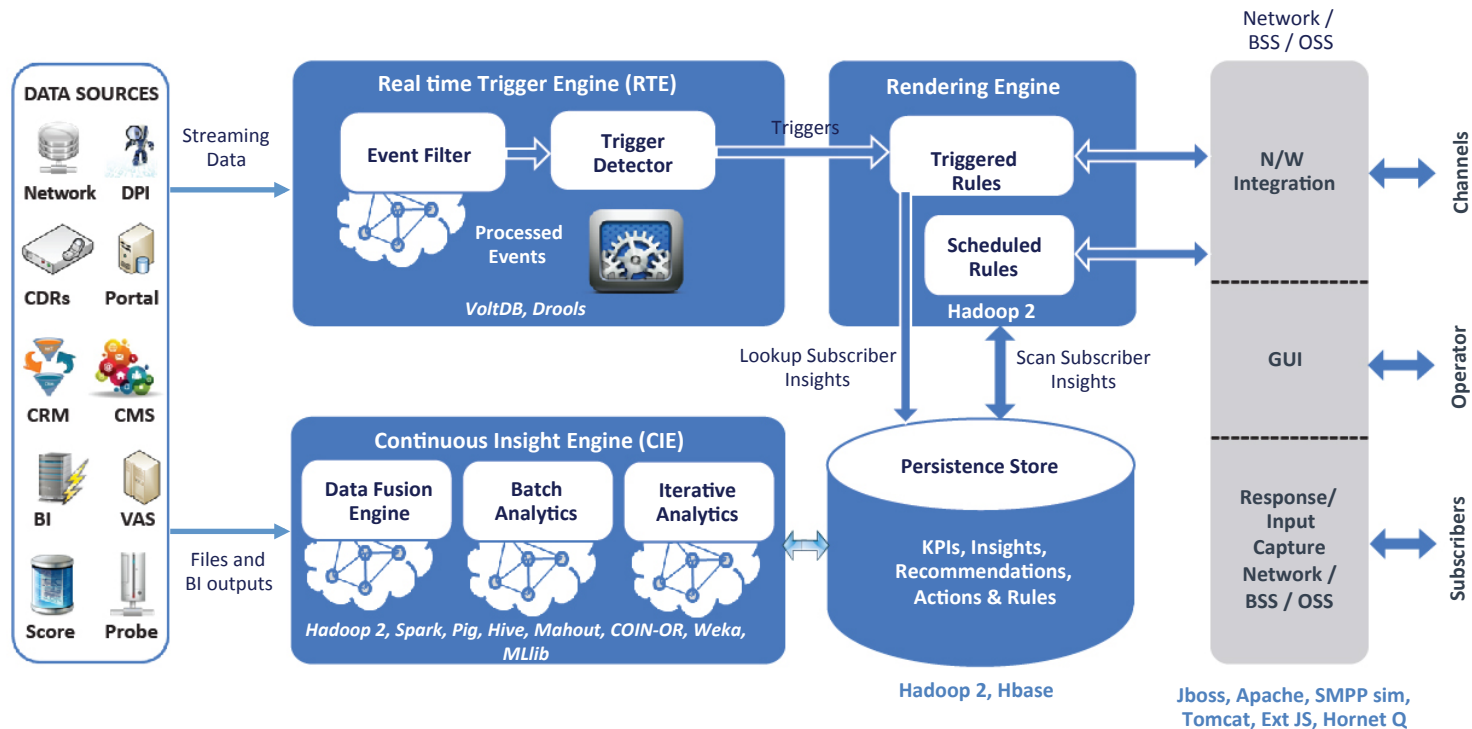
Brands



Partners

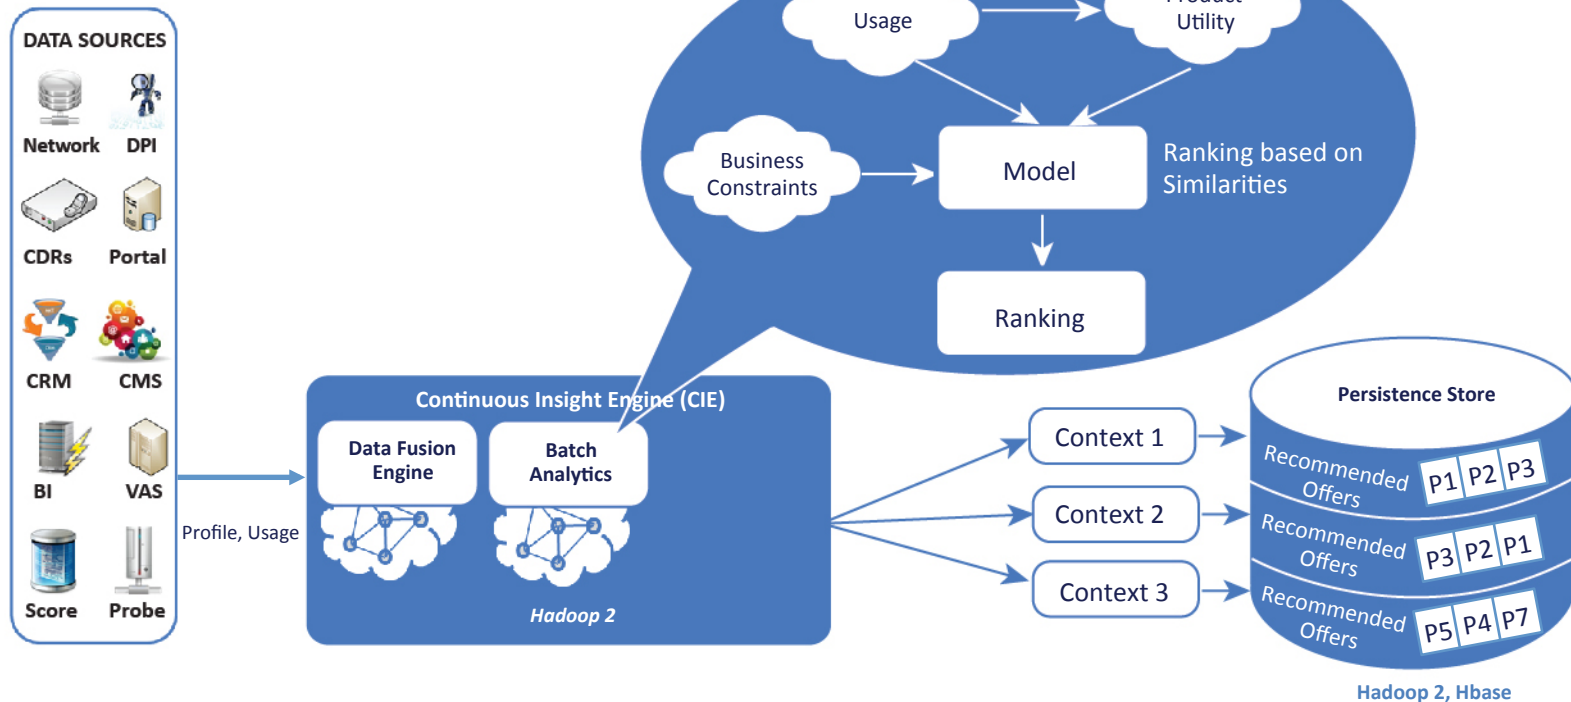


FLYTXT'S INTEGRATED ANALYTICS SOLUTION ARCHITECTURE: BIG DATA, ITERATIVE AND REAL-TIME ANALYTICS




BIG DATA ANALYTICS USE CASE: BEST FIT PRODUCT RECOMMENDATION

Objective: Recommend best fit product to subscribers based on usage and business objectives



Best Fit Product Recommendation

CASE STUDY: CONTEXTUAL PRODUCT RECOMMENDATION FOR TIER 1 OPERATOR



Recommendation Personas:

CLV (HVC, MVC, LVC), Volatile, Early Adopter, Frequent Handset Changer, Heavy Data user, Social Media Fan, Bollywood Fan, Music Fan, Sports Fan, potential iPad buyer, International Caller Etc.....

Objectives:

Cross sell, Upsell, Stimulate recharge/usage/Service adoption Etc...

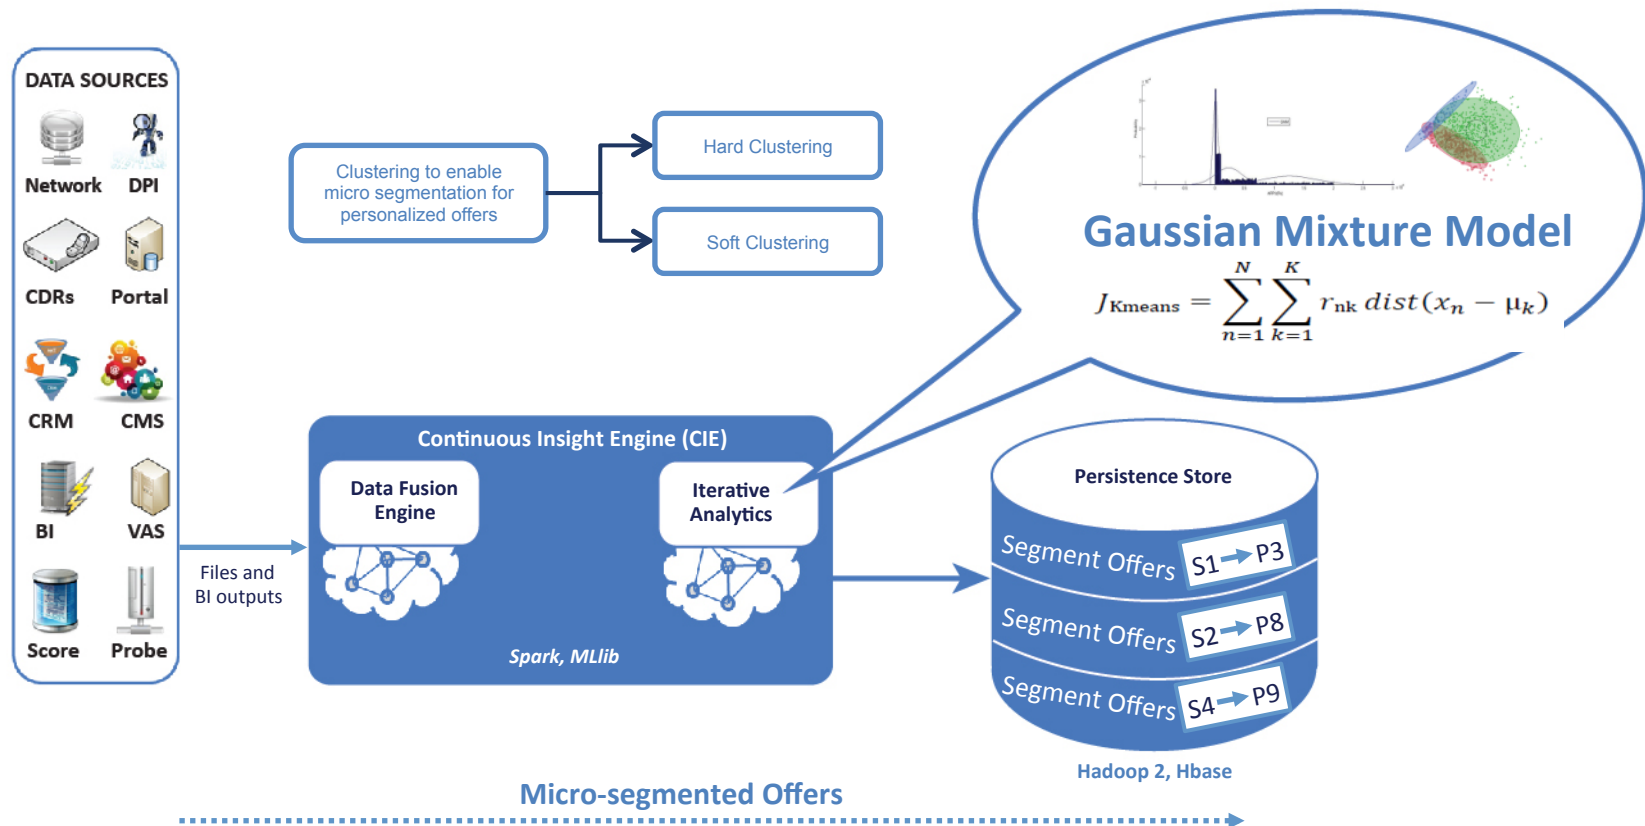
Offers:

Data Plan, 3G plan, VAS usage, International Calling packs, Bundle offers, Recharge stimulation, Seeding, ebill subscription etc.....

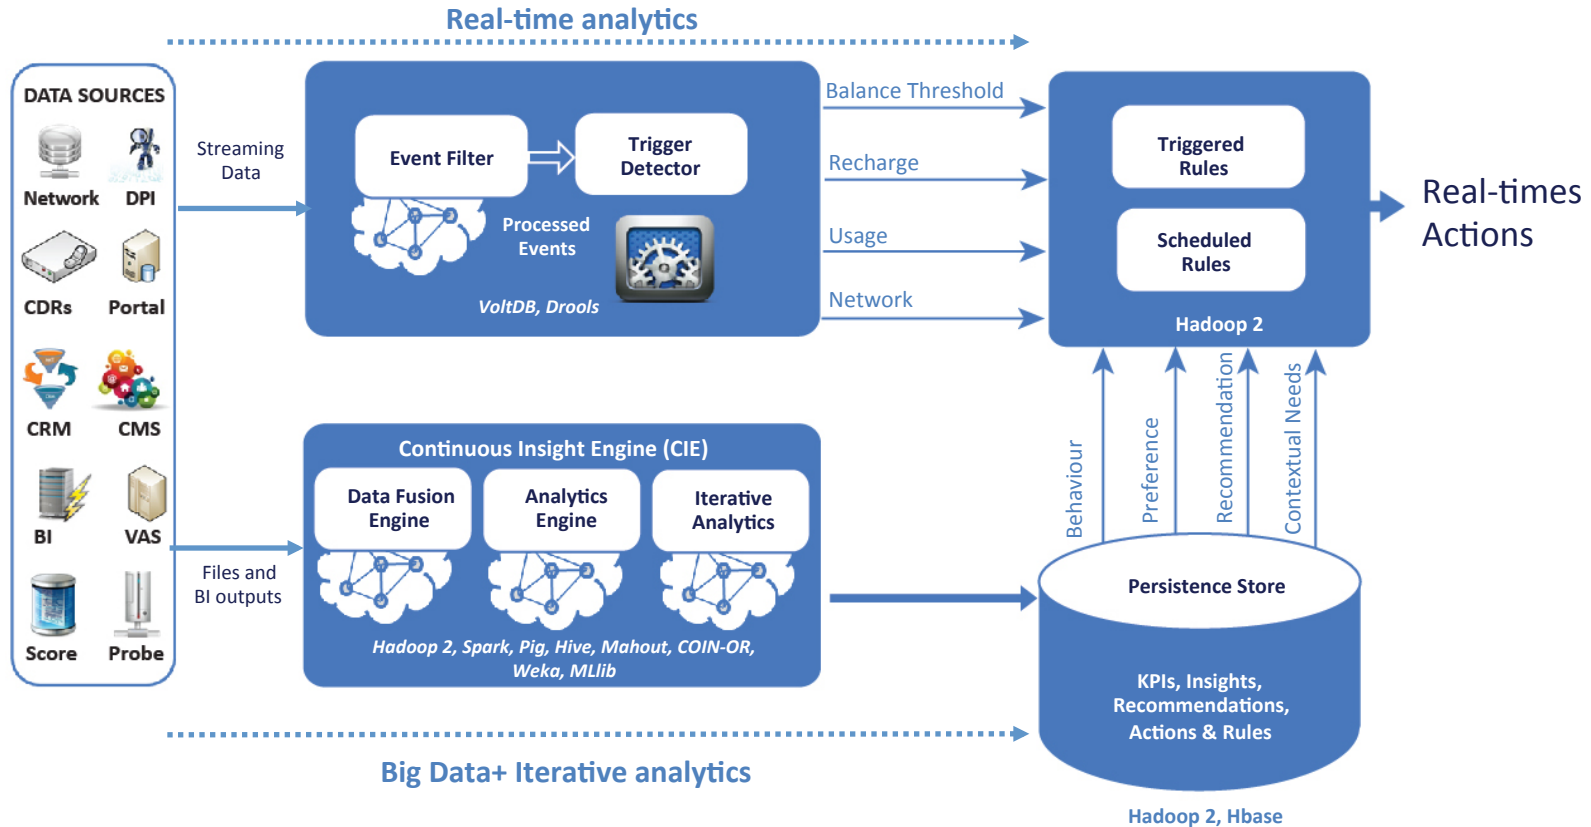
Channels:

IVR, In store, Retailer, WAP portal, Customer care portal

ITERATIVE ANALYTICS USE CASE: MICRO-SEGMENTATION



USE CASE: REAL-TIME ANALYTICS SUPPLEMENTED BY BIG DATA, ITERATIVE ANALYTICS



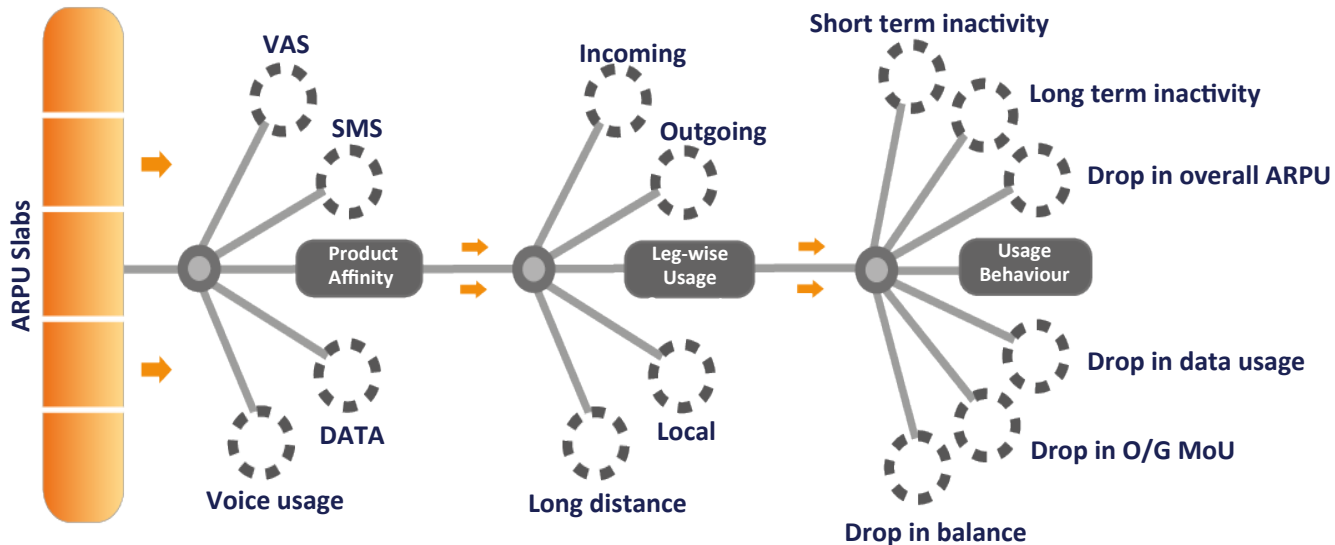
CASE STUDY: REAL-TIME TRIGGER BASED MICRO-SEGMENTED OFFERS

Client Objective

- Improve customer engagement for ARPU enhancement

Data Analyzed

- Customer usage history, ARPU charts, spend patterns and preferences



Solution

- Marketing Program based on usage behaviour driven micro-segmentation and tripwire monitoring

Impact

- 2% increase in month-on-month revenue
- 28% higher revenues & MOU

QUESTIONS?

- Use the chat window to type in your questions or hashtag #VoltDBFlytxt
- Know more about Flytxt
 - Visit www.Flytxt.com
- Try VoltDB yourself:
 - Free trial of the Enterprise Edition:
 - www.voltodb.com/download
 - Try VoltDB in the Cloud
 - Amazon's Cloud Formation
 - Open source version is available on github.com

THANK YOU!