

# A Survey of the State-of-the-art in Event Processing

Otávio M. de Carvalho, Eduardo Roloff, Philippe O. A. Navaux

Federal University of Rio Grande do Sul  
Parallel and Distributed Processing Group  
{omcarvalho,eroloff,navaux}@inf.ufrgs.br



# Motivation

- We currently live in a **world of 2.8 trillion GB**

- **Data production rate is growing 60% by year**

- Every **60 seconds, 600 new blog posts** are published and **34,000 tweets** are sent



# Applications of Information Flow Processing

Sensor networks, ...



traffic engineering, ...



social networking, ...



financial markets, ...



generate streams!



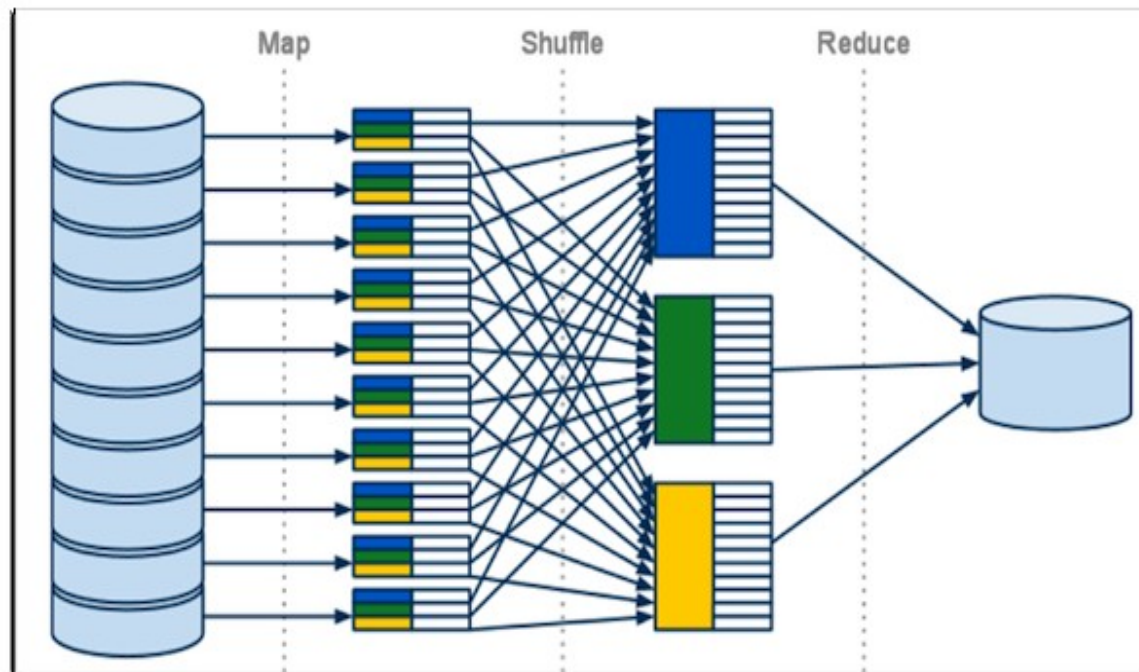
# The History Behind Event Processing Tools

- The **latest advance** in **Event Processing** can be characterized by two main sub-domains: **CEPs** and **DSMSs** systems that were developed in the early **90s**
- **Main problem:** Today's systems present **intersections** between these two sub-domains
- A new **naming convention** was proposed: **Information Flow Processing (IFP)**

Complex Event Processing Systems	Data Stream Management Systems
<ul style="list-style-type: none"><li>• Those systems <b>focuses</b> on <b>processing event</b> notifications, with a special attention to their ordering relationships, to capture <b>complex event patterns</b>.</li></ul>	<ul style="list-style-type: none"><li>• Those systems <b>focuses</b> mainly on <b>flowing data</b> and the application of <b>transformations</b> over data</li></ul>

# MapReduce is Still Important


- The Google's research on MapReduce model was fundamental to the evolution of event processing systems into distributed event processing systems
- The adoption of Apache Hadoop and his core tools has eased the process of developing distributed systems



# But MapReduce Model Isn't Fast Enough



**X**



# **State-of-the-art in Event Processing**

# Characterization of Tools

Name	Release Year	DBMS	DSMS	CEP	Distributed
Google Photon	2013		•		•
Walmart Muppet	2012		•		•
StreamDrill	2012			•	
SAP HANA	2011	•	•		•
Apache Storm	2011		•		•
Apache YARN	2011	•	•	•	•
Apache Flume	2011		•		•
Apache Kafka	2011		•		•
Apache S4	2011		•		•
Apache Chukwa	2010		•		•
HStreaming	2010		•		•
AMPLab Spark	2010	•	•		•
VoltDB	2010	•	•		•
Esper	2006			•	
StreamBase CEP	2003			•	•
SQLstream	2003		•		•



# Characterization of Tools

Name	Release Year	DBMS	DSMS	CEP	Distributed
Google Photon	2013		•		•
Walmart Muppet	2012		•		•
StreamDrill	2012			•	
SAP HANA	2011	•	•		•
Apache Storm	2011		•		•
Apache YARN	2011	•	•	•	•
Apache Flume	2011		•		•
Apache Kafka	2011		•		•
Apache S4	2011		•		•
Apache Chukwa	2010		•		•
HStreaming	2010		•		•
AMPLab Spark	2010	•	•		•
VoltDB	2010		•		•
<b>Esper</b>	<b>2006</b>			•	
<b>StreamBase CEP</b>	<b>2003</b>			•	•
<b>SQLstream</b>	<b>2003</b>		•		•

Legacy Tools

# Characterization of Tools

Name	Release Year	DBMS	DSMS	CEP	Distributed
Google Photon	2013		•		•
Walmart Muppet	2012		•		•
StreamDrill	2012			•	
<b>SAP HANA</b>	2011	•	•		•
Apache Storm	2011		•		•
Apache YARN	2011		•	•	•
Apache Flume	2011		•		•
Apache Kafka	2011		•		•
Apache S4	2011		•		•
Apache Chukwa	2010		•		•
HStreaming	2010		•		•
AMPLab Spark	2010	•	•		•
<b>VoltDB</b>	2010	•	•		•
Esper	2006			•	
StreamBase CEP	2003			•	•
SQLstream	2003		•		•

In-memory  
databases

# Characterization of Tools

Name	Release Year	DBMS	DSMS	CEP	Distributed
Google Photon	2013		•		•
Walmart Muppet	2012		•		•
StreamDrill	2012			•	
SAP HANA	2011	•	•		•
Apache Storm	2011		•		•
Apache YARN	2011			•	•
Apache Flume	2011				•
Apache Kafka	2011		•		•
Apache S4	2011		•		•
Apache Chukwa	2010		•		•
HStreaming	2010		•		•
AMPLab Spark	2010	•	•		•
VoltDB	2010	•	•		•
Esper	2006			•	
StreamBase CEP	2003			•	•
SQLstream	2003		•		•

# Characterization of Tools

Name	Release Year	DBMS	DSMS	CEP	Distributed
Google Photon	2013		•		•
Walmart Muppet	2012		•		•
StreamDrill	2012			•	
SAP HANA	2011	•	•		•
Apache Storm	2011		•		•
<b>Apache YARN</b>	2011	•	•	•	•
Apache Flume	2011				•
Apache Kafka	2011				•
Apache S4	2011				•
Apache Chukwa	2010		•		•
HStreaming	2010		•		•
<b>AMPLab Spark</b>	2010	•	•		•
VoltDB	2010	•	•		•
Esper	2006			•	
StreamBase CEP	2003			•	•
SQLstream	2003		•		•

General-purpose  
systems

# Characterization of Tools

Name	Release Year	DBMS	DSMS	CEP	Distributed
Google Photon	2013		•		•
Walmart Muppet	2012		•		•
<b>StreamDrill</b>	2012			•	
SAP HANA	2011				•
Apache Storm	2011				•
Apache YARN	2011			•	•
Apache Flume	2011		•		•
Apache Kafka	2011		•		•
Apache S4	2011		•		•
Apache Chukwa	2010		•		•
HStreaming	2010		•		•
AMPLab Spark	2010	•	•		•
VoltDB	2010	•	•		•
Esper	2006			•	
StreamBase CEP	2003			•	•
SQLstream	2003		•		•

Specific to solve one problem (top-k problem)

# Hadoop Influency

Name	Based on Hadoop and his core tools	Interact with Hadoop
Google Photon		
Walmart Muppet		
StreamDrill		
<b>SAP HANA</b>		•
Apache Storm		
<b>Apache YARN</b>	•	
<b>Apache Flume</b>		•
<b>Apache Kafka</b>		•
Apache S4		
<b>Apache Chukwa</b>	•	
<b>HStreaming</b>	•	
<b>AMPLab Spark</b>		•
<b>VoltDB</b>		•
Esper		
<b>StreamBase CEP</b>		•
<b>SQLstream</b>		•

# Challenges in Distributed Event Processing

- How to **program**?
- In **what applications** can it be **used**?
- **Now** it **scales**, but is the **throughput** the **same**?

# Conclusions

- **MapReduce** paper and later contributions have **collaborated** to the **changes** into **event processing systems**
- **Convergences** between the tools still are causing **naming problems**
- It is not possible to determine if these **implementations** will **converge** into greater sets of tools of general-purpose systems or more specialized systems.



# Future Works

- In our future works, we will **select a subset of these tools to analyze** further
- We will **evaluate the performance** of these selected tools
- And finally, we will **apply one** of them in a **real HPC application**

# A Survey of the State-of-the-art in Event Processing

## Thanks!

Otávio M. de Carvalho, Eduardo Roloff, Philippe O. A. Navaux

Federal University of Rio Grande do Sul  
Parallel and Distributed Processing Group  
{omcarvalho,eroloff,navaux}@inf.ufrgs.br

