

Spark Application 개발



■ ■ 챕터 시작 페이지

3장

Spark Streaming



■ spark/examples

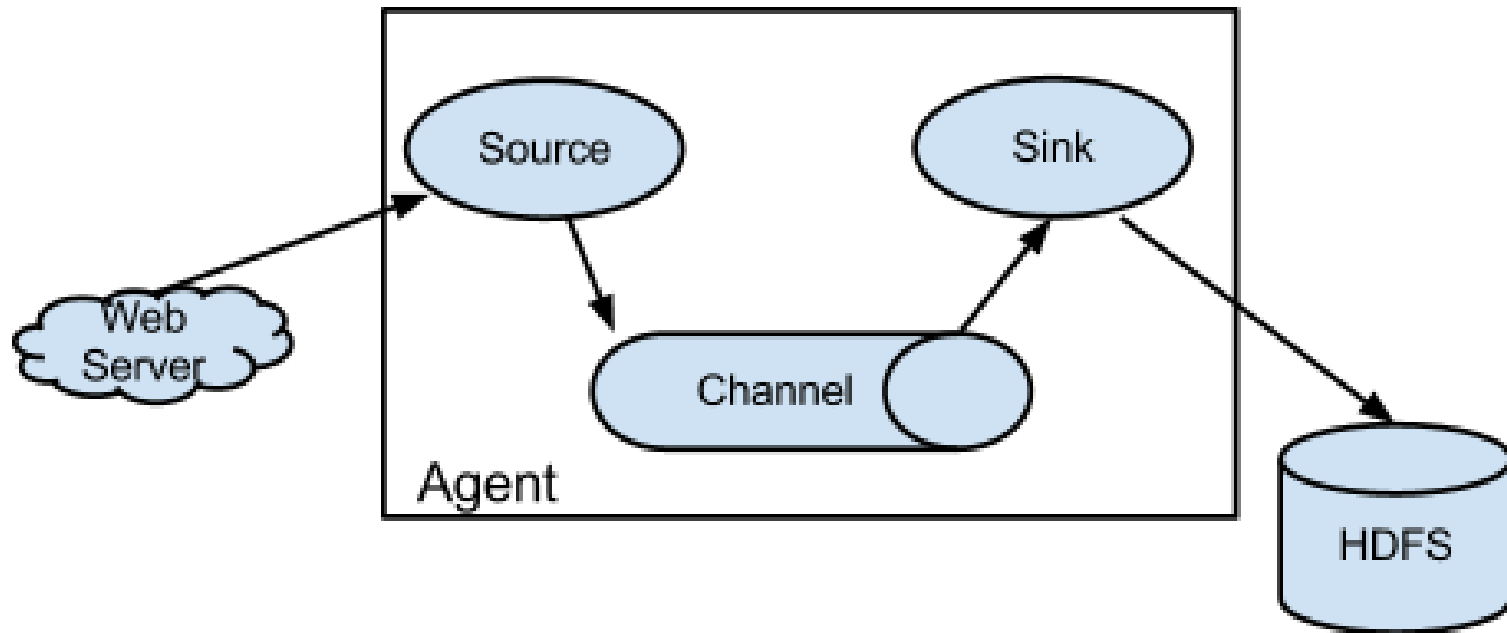
```
[hadoop@bds01 spark]$ ls
assembly      core          ec2           LICENSE       pom.xml       sbin          yarn
bagel         data          examples      logs          project       sbt
bin           dev           external      make-distribution.sh  python        streaming
CHANGES.txt  docker        extras        mllib         README.md     target
conf          docs          graphx        NOTICE       repl          tools
```

■ 예제 코드

```
[hadoop@bds01 examples]$ ls
ActorWordCount.scala      RawNetworkGrep.scala
clickstream               RecoverableNetworkWordCount.scala
FlumeEventCount.scala     StatefulNetworkWordCount.scala
HdfsWordCount.scala       StreamingExamples.scala
KafkaWordCount.scala      TwitterAlgebirdCMS.scala
MQTTWordCount.scala       TwitterAlgebirdHLL.scala
NetworkWordCount.scala    TwitterPopularTags.scala
QueueStream.scala         ZeroMQWordCount.scala
```



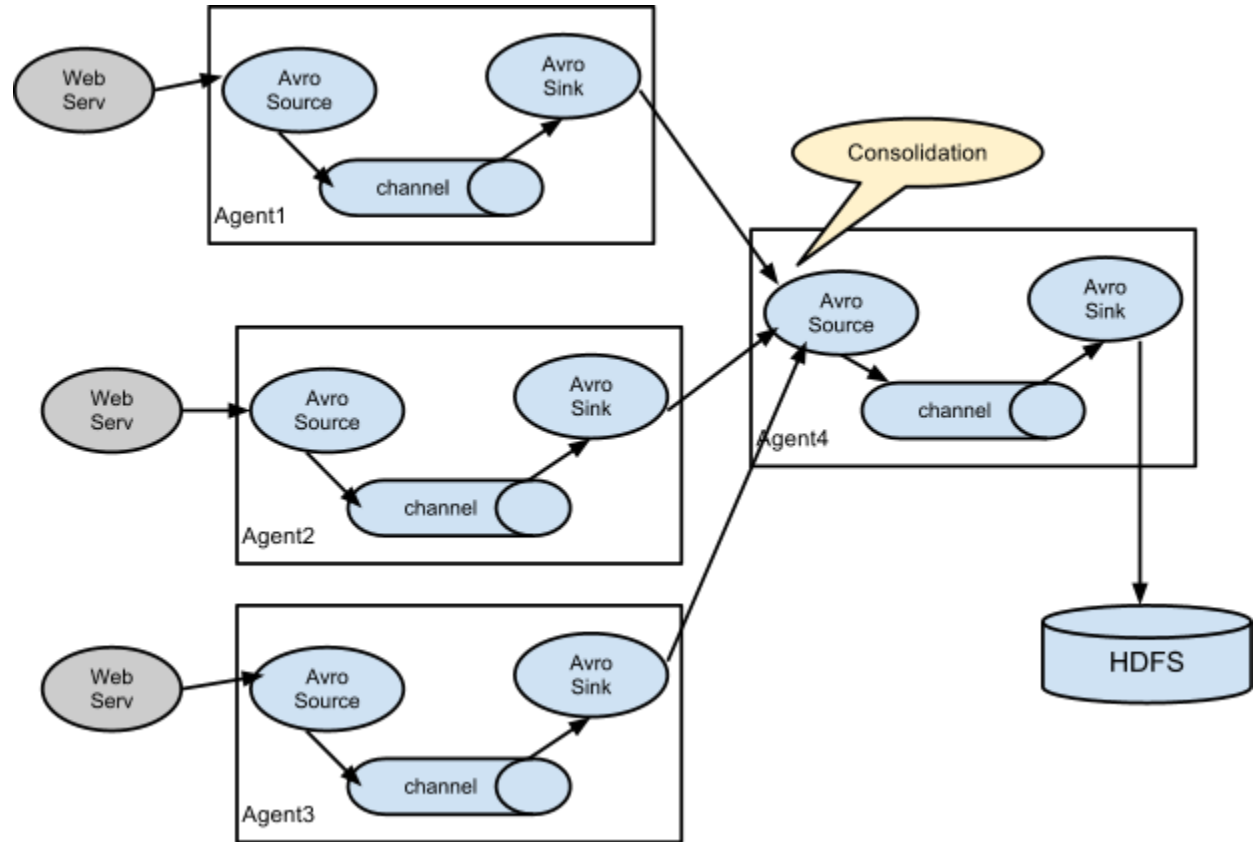
Apache Flume™



Flume-NG



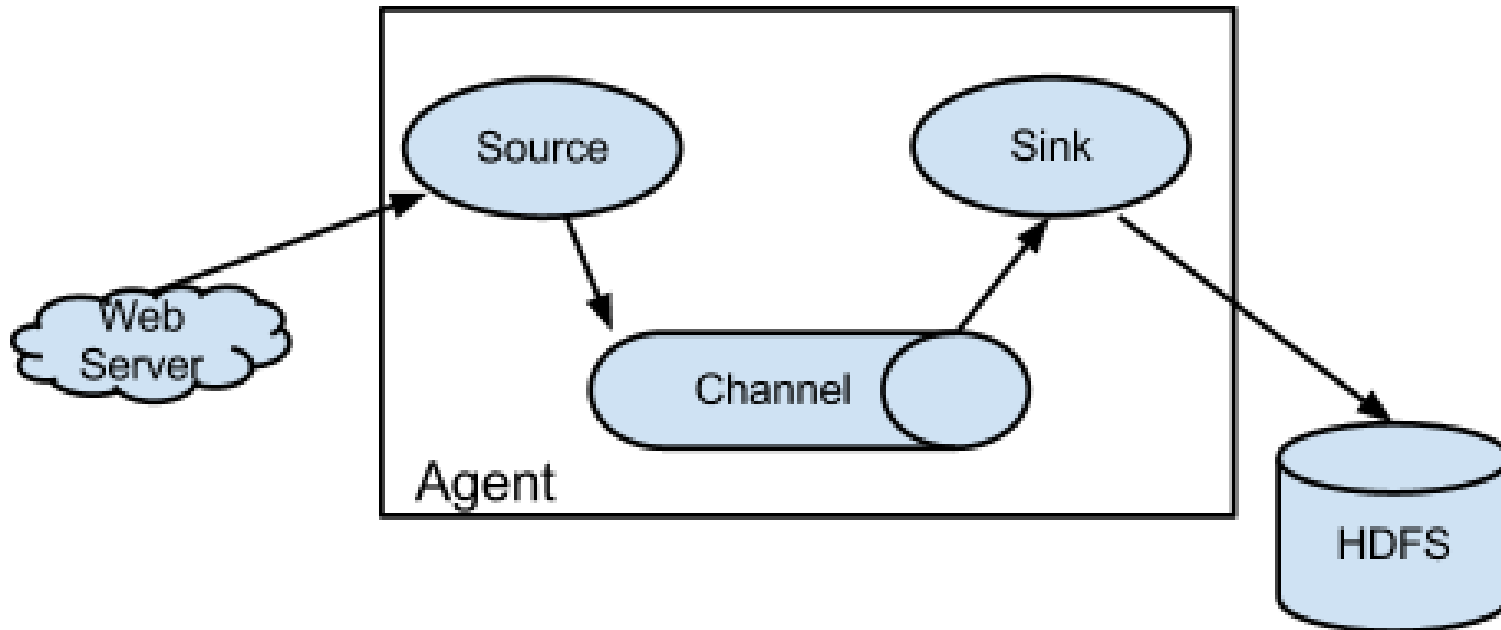
- 분산 데이터 수집/전송시스템
- 최초 설계 목적은 이벤트나 로그 구조의 데이터를 지속적으로 하둡 HDFS에 저장
- 에이전트
- 확장하여 다양한 분야에 활용 가능



Flume-NG



- ❑ 노드 : Flume이 구동되는 머신
- ❑ 모든 노드에는 “source”와 “sink”가 있음
 - source 예 : `tail -F /var/log/httpd/access_log`
 - sink 예 : `dfs(“hdfs://namenode/log/{host}%/%y%m%d”)`
- ❑ 데이터플로우 : 노드들의 체인



Flume-NG 설치



■ Flume 설치

```
$ tar -zxvf apache-flume-1.4.0-bin.tar.gz  
$ ln -s apache-flume-1.4.0-bin apache-flume  
cp flume-conf.properties.template flume-conf
```

■ conf/flume-env.sh 에 다음의 내용 추가

```
JAVA_HOME=/usr/java/java  
FLUME_CLASSPATH="/home/hadoop/apache-flume/lib/"  
export PATH=$PATH:/home/hadoop/hadoop/bin/
```

Flume-NG 설치



■ Starting an agent

```
$ bin/flume-ng agent -n $agent_name -c conf -f conf/flume-  
conf.properties.template
```

Flume-NG 설치



❑ conf/flume.conf 에 설정

agent의 각 요소에 이름을 부여

```
a1.sources = r1
```

```
a1.sinks = k1
```

```
a1.channels = c1
```

source 설정

```
a1.sources.r1.type = netcat
```

```
a1.sources.r1.bind = localhost
```

```
a1.sources.r1.port = 44444
```

sink 설정

```
a1.sinks.k1.type = logger
```

채널 설정

```
a1.channels.c1.type = memory
```

```
a1.channels.c1.capacity = 1000
```

```
a1.channels.c1.transactionCapacity = 100
```

source 와 sink 를 채널에 연결

```
a1.sources.r1.channels = c1
```

```
a1.sinks.k1.channel = c1
```


Flume-NG 설치



❑ conf/flume.conf 에 설정

```
# Name the components on this agent
a1.sources = r1
a1.sinks = k1
a1.channels = c1
### source 설정
a1.sources.r1.type = exec
a1.sources.r1.command = tail -F /home/hadoop/syslog/a.txt
a1.sources.r1.channels = c1
### sink 를 hdfs로 설정
a1.sinks.k1.type = hdfs
a1.sinks.k1.channel = c1
a1.sinks.k1.hdfs.path = hdfs://hadoop01:9000/user/hadoop/logdata/a.txt
a1.sinks.k1.hdfs.filePrefix = events-
a1.sinks.k1.hdfs.round = true
a1.sinks.k1.hdfs.roundValue = 10
a1.sinks.k1.hdfs.roundUnit = minute
### 채널 설정
a1.channels.c1.type = memory
a1.channels.c1.capacity = 1000
a1.channels.c1.transactionCapacity = 100
### Bind the source and sink to the channel
a1.sources.r1.channels = c1
a1.sinks.k1.channel = c1
```

Flume-NG 설치



예제 실행

```
$ bin/flume-ng agent --conf ./conf/ -f conf/flume.conf W  
-Dflume.root.logger=DEBUG,console -n agent
```

다음과 같이 실행 로그 출력

```
2013-06-18 14:00:49,784 (hdfs-hdfs-sink-call-runner-0) [INFO -  
org.apache.flume.sink.hdfs.BucketWriter.doOpen(BucketWriter.java:189)] Creating  
hdfs://localhost:54310/tmp/system.log//FlumeData.1371589249458.tmp
```



[Go to parent directory](#)

| Name | Type | Size | Replication | Block Size | Modification Time | Permission | Owner | Group |
|---------------------------------------|------|---------|-------------|------------|-------------------|------------|--------|------------|
| events-.1382271326590 | file | 0.41 KB | 3 | 64 MB | 2013-10-20 21:15 | rw-r--r-- | hadoop | supergroup |
| events-.1382271326591 | file | 0.31 KB | 3 | 64 MB | 2013-10-20 21:19 | rw-r--r-- | hadoop | supergroup |

Flume-NG 설치



Flume Sources

| | |
|---------------|---------------------------|
| Flume Sources | Avro Source |
| | Thrift Source |
| | Exec Source |
| | JMS Source |
| | Spooling Directory Source |
| | NetCat Source |
| | Syslog Sources |
| | Syslog UDP Source |
| | HTTP Source |
| | Legacy Sources |
| | Custom Source |

Flume-NG 설치



■ HDFS Sink

| Name | Default | Description |
|-------------------------------|-------------------|---|
| channel | – | |
| type | – | The component type name, needs to be <code>hdfs</code> |
| hdfs.path | – | HDFS directory path (eg <code>hdfs://namenode/flume/webdata/</code>) |
| hdfs.filePrefix | FlumeData | Name prefixed to files created by Flume in hdfs directory |
| hdfs.fileSuffix | – | Suffix to append to file (eg <code>.avro</code> - <i>NOTE: period is not automatically added</i>) |
| hdfs.inUsePrefix | – | Prefix that is used for temporal files that flume actively writes into |
| hdfs.inUseSuffix | <code>.tmp</code> | Suffix that is used for temporal files that flume actively writes into |
| hdfs.rollInterval | 30 | Number of seconds to wait before rolling current file (0 = never roll based on time interval) |
| hdfs.rollSize | 1024 | File size to trigger roll, in bytes (0: never roll based on file size) |
| hdfs.rollCount | 10 | Number of events written to file before it rolled (0 = never roll based on number of events) |
| hdfs.idleTimeout | 0 | Timeout after which inactive files get closed (0 = disable automatic closing of idle files) |
| hdfs.batchSize | 100 | number of events written to file before it is flushed to HDFS |
| hdfs.codeC | – | Compression codec. one of following : <code>gzip</code> , <code>bzip2</code> , <code>lzo</code> , <code>lzop</code> , <code>snappy</code> |
| hdfs.fileType | SequenceFile | File format: currently <code>SequenceFile</code> , <code>DataStream</code> or <code>CompressedStream</code> (1) <code>DataStream</code> will not compress output file and please don't set <code>codeC</code> (2) <code>CompressedStream</code> requires set <code>hdfs.codeC</code> with an available <code>codeC</code> |
| hdfs.maxOpenFiles | 5000 | Allow only this number of open files. If this number is exceeded, the oldest file is closed. |
| hdfs.minBlockReplicas | – | Specify minimum number of replicas per HDFS block. If not specified, it comes from the default Hadoop config in the classpath. |
| hdfs.writeFormat | – | Format for sequence file records. One of “Text” or “Writable” (the default). |
| hdfs.callTimeout | 10000 | Number of milliseconds allowed for HDFS operations, such as open, write, flush, close. This number should be increased if many HDFS timeout operations are occurring. |
| hdfs.threadsPoolSize | 10 | Number of threads per HDFS sink for HDFS IO ops (open, write, etc.) |
| hdfs.rollTimerPoolSize | 1 | Number of threads per HDFS sink for scheduling timed file rolling |

Spark Streaming



- 두 개의 터미널에서 실행
- 첫 번째 터미널

```
bin/run-example  
org.apache.spark.streaming.examples.FlumeEventCount bds02  
4545
```

- 두 번째 터미널

```
bin/flume-ng agent -n a1 -c conf -f conf/flume_avro.txt
```

Spark Streaming



■ 결과 확인

```
-----  
Time: 1407634802000 ms
```

```
-----  
Received 0 flume events.
```

```
-----  
Time: 1407634804000 ms
```

```
-----  
Received 10 flume events.
```

```
-----  
Time: 1407634806000 ms
```

```
-----  
Received 0 flume events.
```

Spark Streaming



FlumeEventCount

```
package org.apache.spark.streaming.examples

import org.apache.spark.storage.StorageLevel
import org.apache.spark.streaming._
import org.apache.spark.streaming.flume._
import org.apache.spark.util.IntParam

object FlumeEventCount {
  def main(args: Array[String]) {
    if (args.length != 3) {
      System.err.println(
        "Usage: FlumeEventCount <master> <host> <port>")
      System.exit(1)
    }
  }

  StreamingExamples.setStreamingLogLevels()
```

Spark Streaming



■ FlumeEventCount(계속)

```
val Array(master, host, IntParam(port)) = args

val batchInterval = Milliseconds(2000)
// Create the context and set the batch size
val ssc = new StreamingContext(master, "FlumeEventCount",
    batchInterval, System.getenv("SPARK_HOME"),
    StreamingContext.jarOfClass(this.getClass))

// Create a flume stream
val stream = FlumeUtils.createStream(ssc,
    host, port, StorageLevel.MEMORY_ONLY_SER_2)
```


Spark Streaming



FlumeEventCount(계속)

```
// Print out the count of events received events in each batch  
stream.count().map(cnt => "Received " + cnt + " flume  
    events." ).print()
```

```
ssc.start()  
ssc.awaitTermination()  
}  
}
```